

Distributed Data Framework Documentation

Version 2.11.3. Copyright (c) Codice Foundation

Table of Contents

License.....	1
Introduction.....	2
1. About DDF	2
1.1. Introducing DDF	2
1.2. Component Applications.....	2
1.3. Documentation Guide	3
1.3.1. Documentation Conventions	3
1.3.1.1. Customizable Values.....	3
1.3.1.2. Code Values.....	3
1.3.1.3. Hyperlinks.....	4
1.3.2. Support	4
1.3.2.1. Documentation Updates	4
1.4. Core Concepts	4
1.4.1. Introduction to Search.....	4
1.4.1.1. Search Types.....	4
1.4.2. Introduction to Metadata	5
1.4.3. Introduction to Ingest	5
1.4.4. Introduction to Resources.....	5
1.4.5. Introduction to the Catalog Framework	5
1.4.6. Introduction to Federation	6
1.4.7. Introduction to Events and Subscriptions.....	6
1.4.8. Introduction to Registries	6
1.4.9. Introduction to Endpoints.....	7
1.4.10. Introduction to High Availability	7
1.4.10.1. High Availability Supported Capabilities	8
2. Quick Start Tutorial	9
2.1. Quick Install Prerequisites	9
2.1.1. Hardware Requirements (Quick Install).....	9
2.1.2. Java Requirements (Quick Install)	9
2.2. Quick Install of DDF	11
2.3. Certificates (Quick Start).....	12
2.3.1. Demo Certificate Authority (CA)	12
2.3.1.1. Creating New Server Keystore Entry with the CertNew Scripts	12
2.3.2. Creating Self-Signed Certificates	15
2.3.2.1. Creating a custom CA Key and Certificate	15
2.3.2.2. Sign Certificates Using the custom CA	15
2.4. Configuring (Quick Start)	16
2.5. Ingesting (Quick Start).....	16

2.5.1. Ingesting Sample Data	16
Managing	17
3. Securing	17
3.1. Security Hardening	17
3.1.1. Hardening Checklist	18
3.2. Auditing	18
3.2.1. Enabling Fallback Audit Logging	19
4. Installing	20
4.1. Installation Prerequisites	20
4.2. Installing With the DDF Distribution Zip	23
4.2.1. Controlling File System Access	25
4.2.1.1. Setting Directory Permissions	25
4.3. Initial Startup	27
4.3.1. Verifying Startup	27
4.3.2. DDF Directory Contents after Installation and Initial Startup	28
4.3.3. Complete Installation	29
4.3.3.1. Completing Installation from the Admin Console	29
4.3.3.2. Completing Installation from the Command Console	31
4.3.3.2.1. Configure Guest Claim Attributes	31
4.3.3.2.2. System Configuration Settings	31
4.3.4. Firewall Port Configuration	32
4.3.5. Internet Explorer 11 Enhanced Security Configuration	32
4.4. High Availability Initial Setup	33
4.4.1. High Availability Initial Setup Exceptions	33
4.4.1.1. HTTP Proxy Integration	33
4.4.1.2. Identical Directory Structures	34
4.4.1.3. Highly Available Security Auditing	34
4.4.1.4. Shared Storage Provider	34
4.4.1.5. High Availability Certificates	34
4.4.1.6. High Availability Installation Profile	34
5. Configuring	34
5.1. Admin Console Tutorial	35
5.2. Console Command Reference	37
5.2.1. Application Commands	37
5.2.2. Feature Commands	38
5.2.2.1. Uninstalling Features from the Command Console	39
5.3. Configuration Files	39
5.3.1. Configuring Global Settings with system.properties	40
5.3.2. Configuring with .config Files	46
5.4. Managing Keystores and Certificates	47
5.4.1. Managing Keystores	48

5.4.1.1. Adding an Existing Server Keystore	49
5.4.1.2. Adding an Existing Server Truststore	49
5.4.1.3. Creating a New Keystore/Truststore with an Existing Certificate and Private Key ..	49
5.4.1.4. Updating Key Store / Trust Store via the Admin Console	50
5.4.2. Disallowing Login Without Certificates	51
5.4.3. Managing Certificate Revocation List (CRL)	52
5.4.3.1. Creating a Certificate Revocation List (CRL)	52
5.4.3.1.1. Revoke a Certificate and Create a New CRL that Contains the Revoked Certificat	52
5.4.3.1.2. Viewing a CRL	52
5.4.3.2. Enabling Certificate Revocation	52
5.4.3.2.1. Add Revocation to a Web Context	53
5.4.3.2.2. Adding Revocation to an Endpoint	53
5.4.3.2.3. Verifying Revocation	54
5.5. Configuring User Access	55
5.5.1. Configuring Guest Access	55
5.5.1.1. Denying Guest User Access	55
5.5.1.1.1. Configuring SOAP Guest Interceptor if Prohibiting Guest Users	55
5.5.1.1.2. Configuring Guest Validator if Prohibiting Guest Users	55
5.5.1.2. Allowing Guest User Access	55
5.5.1.2.1. Configuring Guest Interceptor if Allowing Guest Users	56
5.5.1.2.2. Configuring Guest Claim Attributes	56
5.5.2. Configuring REST Services for Users	56
5.5.2.1. Configuring Included IdP	56
5.5.2.1.1. Configuring Included STS	60
5.5.2.1.2. Connecting to External STS	63
5.5.2.2. Connecting to an External IdP	63
5.5.2.3. Configuring Without an IdP	64
5.5.2.3.1. Using STS without IdP	64
5.5.2.3.2. Connecting to External STS Without IdP	64
5.5.3. Configuring SOAP Services for Users	64
5.5.3.1. Connecting to Included STS with SOAP	64
5.5.3.2. Connecting to External STS with SOAP	65
5.5.4. Connecting to an LDAP Server	65
5.5.5. Updating System Users	65
5.5.6. Restricting Access to Admin Console	66
5.5.6.1. Restricting Feature, App, Service, and Configuration Access	67
5.5.6.2. Removing Default Users	67
5.6. Configuring Data Management	68
5.6.1. Configuring Solr	68
5.6.1.1. Configuring Solr Catalog Provider Data Directory	69
5.6.1.1.1. Changing an Empty Data Directory	69

5.6.1.1.2. Changing the Data Directory After Ingest	69
5.6.1.2. Hardening Solr Index	70
5.6.1.2.1. Configuring Solr Node Security	70
5.6.1.2.2. Configuring Solr Encryption	70
5.6.2. Changing Catalog Providers	70
5.6.3. Changing Hostname	71
5.6.4. Content Directory Monitor	71
5.6.4.1. Installing the Content Directory Monitor	71
5.6.4.2. Configuring the Content Directory Monitor	71
5.6.4.3. Using the Content Directory Monitor	72
5.7. Configuring Security Policies	73
5.7.1. Configuring the Web Context Policy Manager	73
5.7.1.1. Context Realms	74
5.7.1.2. Authentication Types	74
5.7.1.3. Required Attributes	74
5.7.1.4. White Listed Contexts	75
5.7.2. Configuring Catalog Filtering Policies	75
5.7.2.1. Setting Internal Policies	75
5.7.2.2. Setting XACML Policies	75
5.7.2.3. Catalog Filter Policy Plugins	76
5.8. Configuring Federation	76
5.8.1. Enable SSL for Clients	76
5.8.2. Configuring HTTP Port	77
5.8.3. Configuring HTTP Proxy	77
5.8.4. Federation Strategy	78
5.8.4.1. Configuring Federation Strategy	78
5.8.4.1.1. Catalog Federation Strategy	78
5.8.5. Connecting to Sources	79
5.8.5.1. Federated Source for Atlassian Confluence(R)	81
5.8.5.2. CSW Connected Source	82
5.8.5.3. CSW Federated Source	83
5.8.5.4. CSW Federation Profile Source	83
5.8.5.5. GMD CSW Source	84
5.8.5.6. OpenSearch Source	85
5.8.5.7. Registry Store	86
5.8.5.8. Solr Catalog Provider	87
5.8.5.9. WFS 1.0 Source	91
5.8.5.10. WFS 2.0 Source	93
5.9. Environment Hardening	95
5.10. Configuring for Special Deployments	97
5.10.1. Multiple Installations	97

5.10.1.1. Reusing Configurations	97
5.10.1.2. Isolating Solr Cloud and Zookeeper	98
5.10.1.3. Installing Multiple DDFs on the Same Host	99
5.10.2. Configuring for a Fanout Proxy	99
5.10.3. Standalone Security Token Service (STS) Installation	99
5.10.4. Configuring for a Highly Available Cluster	100
5.11. Miscellaneous Configurations	100
5.11.1. Configuring Thread Pools	100
5.11.2. Configuring Jetty ThreadPool Settings	100
5.11.3. Configuring Alerts	101
5.11.3.1. Configuring Decanter Service Level Agreement (SLA) Checker	101
5.11.3.2. Configuring Decanter Scheduler	101
5.11.4. Landing Page	102
5.11.4.1. Installing the Landing Page	102
5.11.4.2. Configuring the Landing Page	102
5.11.4.3. Customizing the Landing Page	102
6. Running	102
6.1. Starting	103
6.1.1. Starting from Startup Scripts	103
6.1.2. Starting as a Background Process	104
6.1.3. Stopping DDF	104
6.1.4. Starting as a Service with Automatic Start on System Boot	105
6.1.4.1. Karaf Documentation	107
6.2. Maintaining	107
6.2.1. Console Commands	107
6.2.1.1. Accessing the System Console	107
6.2.1.2. Available System Console Commands	107
6.2.1.3. System Console Command Help	108
6.2.1.4. Catalog Commands	109
6.2.1.5. catalog:dump Options	110
6.2.1.6. Date Syntax	110
6.2.1.6.1. catalog:dump Examples	111
6.2.1.7. Application Commands	112
6.2.1.7.1. Application Command Usage	112
6.2.1.7.2. Application in Failed State	114
6.2.1.8. Subscriptions Commands	116
6.2.1.9. subscriptions:list Command Usage Examples	116
6.2.1.10. subscriptions:delete Command Usage	118
6.2.1.11. Platform Commands	119
6.2.1.12. Persistence Store Commands	119
6.2.2. Command Scheduler	119

6.2.2.1. Using the Command Scheduler	120
6.2.2.1.1. Schedule a Command	120
6.2.2.2. Updating a Scheduled Command	120
6.2.2.2.1. Command Output	121
6.2.2.3. CQL Syntax	121
6.3. Monitoring	121
6.3.1. Metrics Reporting	121
6.3.2. Managing Logging	122
6.3.2.1. Configuring Logging	122
6.3.2.2. DDF log file	122
6.3.2.3. Controlling log level	122
6.3.2.4. Controlling the size of the log file	122
6.3.2.5. Number of backup log files to keep	122
6.3.2.6. Enabling logging of inbound and outbound SOAP messages for the DDF SOAP endpoints	122
6.3.2.7. Logging External Resources	123
6.3.2.8. Enabling HTTP Access Logging	123
6.3.2.9. Using the LogViewer	124
6.4. Troubleshooting	124
6.4.1. Deleted Records Are Being Displayed In The Standard Search UI's Search Results	126
7. Data Management	127
7.1. File Formats Supported	127
7.2. Ingesting Data	129
7.2.1. Methods of Ingest	129
7.2.1.1. Methods for ingesting small sets of records	129
7.2.1.1.1. Ingest Command (Command Console)	129
7.2.1.1.2. User Interface Ingest	129
7.2.1.2. Methods for ingesting large sets of records	129
7.2.1.2.1. Directory Monitor Ingest	130
7.2.1.2.2. External Methods	130
7.2.1.2.3. Other Methods	131
7.2.2. Verifying Ingest by REST Endpoint	131
7.3. Removing Expired Records from the Catalog	131
7.3.1. Manual Removal of Expired Records	132
7.3.2. Automated Removal of Expired Records	132
7.3.2.1. Explanation of Command to Remove Expired Records	132
7.3.3. Non-Universal or Catalog Specific Removal	132
7.3.4. Automatic Catalog Backup	133
7.3.5. DDF Data Migration	133
7.3.5.1. Exporting Metadata Out of a Catalog Provider	133
7.3.5.2. Ingesting Exported Metadata into a Catalog Provider	133

7.3.5.3. Translating Metadata from One Format to Another	133
Using	134
8. Using the Landing Page	134
8.1. Search DDF Button	134
8.2. Data Source Availability	134
8.3. Announcements	134
9. Using Intrigue	134
9.1. Accessing Intrigue	135
9.2. Workspaces in Intrigue	135
9.2.1. Creating a Workspace in Intrigue	136
9.2.2. Configuring a Workspace in Intrigue	137
9.2.3. Sharing Workspaces	138
9.3. Ingesting from Intrigue	138
9.4. Searching with Intrigue	138
9.4.1. Search Tab	139
9.4.1.1. Editing a Search	140
9.4.1.1.1. Editing Search Settings	141
9.4.1.1.2. Editing Search Notifications	141
9.4.1.1.3. Viewing Search Status	141
9.4.1.2. Refining Search Results	142
9.4.1.3. Search Result Options	143
9.4.2. Bookmarks Tab	143
9.5. Viewing Search Results	143
9.5.1. Adding Visuals	143
9.5.2. Editing Records	144
9.5.3. Editing Associations on a Record	144
9.5.4. Viewing Revision History	145
9.5.5. Viewing Metadata Quality	145
9.5.6. Exporting a Result	146
9.5.7. Archiving a Result	146
9.5.8. Restoring Archived Results	146
9.5.9. Overwriting a Resource	147
9.5.10. Intrigue Settings	147
9.5.11. Intrigue Notifications	148
10. Using Standard Search UI	149
10.1. Search	149
10.1.1. Search Criteria	149
10.1.2. Results	149
10.1.3. Enhanced Search	150
10.1.4. Record Summary	150
10.1.5. Record Details	151

10.2. Actions	152
10.2.1. Save a Search	152
10.3. Workspaces	153
10.3.1. Create a Workspace	153
10.4. Notifications	153
10.5. Activities	154
10.6. Downloads	154
10.7. Maps	154
11. Using the Simple Search	154
11.1. Search	154
11.1.1. Search Criteria	154
11.1.2. Results	155
11.1.2.1. Results Summary	155
11.1.2.2. Results Table	155
11.1.3. Result View	156
Integrating	156
12. Federation	156
13. Data and Metadata	157
13.1. Metacards	158
13.1.1. Metacard Type	158
13.1.1.1. Default Metacard Type and Attributes	158
13.1.1.2. Extensible Metacards	159
13.1.2. Metacard Type Registry	159
13.1.3. Attributes	160
13.1.3.1. Attribute Types	160
13.1.3.1.1. Attribute Format	161
13.1.3.1.2. Attribute Naming Conventions	161
13.1.3.2. Result	161
13.1.4. Creating Metacards	162
13.1.4.1. Limitations	162
13.1.4.2. Processing Metacards	162
13.1.4.3. Basic Types	162
13.2. JSON "Definition" Files	163
13.2.1. Definition File Format	164
13.2.2. Deploying Definition Files	164
13.3. Data Validation	164
13.3.1. Validation with Schematron	164
13.3.1.1. Configuring Schematron Services	165
14. Endpoints	165
14.1. Available Endpoints	165
14.1.1. Application Upload Endpoint	166

14.1.1.1. Installing Application Upload Endpoint	166
14.1.1.2. Application Upload Endpoint	166
14.1.1.3. Configuring Application Upload Endpoint	166
14.1.2. Catalog REST Endpoint	166
14.1.2.1. Installing the Catalog REST Endpoint	166
14.1.2.2. Installing the Catalog REST Endpoint	166
14.1.2.3. Configuring the Catalog REST Endpoint	167
14.1.2.4. Using the Catalog REST Endpoint	167
14.1.2.4.1. Sources Operation Example	169
14.1.2.5. Interacting with the REST CRUD Endpoint	169
14.1.2.6. Metacard Transforms with the REST CRUD Endpoint	170
14.1.2.6.1. POST Metadata	170
14.1.2.6.2. Example Responses for ReST Endpoint Error Conditions	171
14.1.3. CometD Endpoint	173
14.1.3.1. Installing CometD Endpoint	173
14.1.3.2. Configuring CometD Endpoint	174
14.1.3.3. Using CometD Endpoint	174
14.1.3.3.1. CometD Queries	174
14.1.3.3.2. Query Request Examples	175
14.1.3.3.3. Query Response Channel	175
14.1.3.3.4. Query Response Examples	176
14.1.3.3.5. CometD Notifications	178
14.1.3.3.6. Using CometD Notifications	178
14.1.3.3.7. Receive Notifications	179
14.1.3.3.8. Notification Events	179
14.1.3.3.9. Persistence of Notifications	179
14.1.3.3.10. Notification Operations Channel	179
14.1.3.3.11. Activity Events Channel	180
14.1.3.3.12. Activity Operations Channel	182
14.1.4. CSW Endpoint	182
14.1.4.1. Installing CSW Endpoint	182
14.1.4.2. Configuring CSW Endpoint	183
14.1.4.3. CSW Endpoint URL	183
14.1.4.3.1. CSW Endpoint Operations	183
14.1.4.3.2. Transaction Operations	217
14.1.4.3.3. Transaction Insert Sub-Operation HTTP POST	217
14.1.4.3.4. Transaction Insert Response	218
14.1.4.3.5. Transaction Update Sub-Operation HTTP POST	219
14.1.4.3.6. Transaction Delete Sub-Operation HTTP POST	223
14.1.4.3.7. Subscription GetRecords Operation	224
14.1.4.3.8. Subscription GetRecords HTTP GET	224

14.1.4.3.9. Subscription GetRecords HTTP POST	225
14.1.4.3.10. Subscription GetRecords HTTP PUT	225
14.1.4.3.11. Subscription HTTP GET or HTTP DELETE Request	227
14.1.4.3.12. Subscription HTTP GET or HTTP DELETE Response	227
14.1.4.3.13. Example Responses for CSW Endpoint Error Conditions	228
14.1.5. FTP Endpoint	236
14.1.5.1. Installing the FTP Endpoint	236
14.1.5.2. Configuring the FTP Endpoint	236
14.1.5.3. Using FTP Endpoint	237
14.1.5.3.1. From Code:	237
14.1.5.3.2. From an FTP client:	237
14.1.6. KML Endpoint	237
14.1.6.1. Installing the KML Endpoint	237
14.1.6.2. Configuring the KML Endpoint	237
14.1.6.3. Using the KML Endpoint	238
14.1.7. Metrics Endpoint	239
14.1.7.1. Installing the Metrics Endpoint	239
14.1.7.2. Configuring the Metrics Endpoint	239
14.1.7.3. Using the Metrics Endpoint	239
14.1.7.3.1. Metrics Data Supported Formats	240
14.1.7.3.2. Add Custom Metrics to the Metrics Tab	243
14.1.7.4. Usage Limitations of the Metrics Endpoint	243
14.1.8. OpenSearch Endpoint	243
14.1.8.1. Installing the OpenSearch Endpoint	243
14.1.8.2. Configuring the OpenSearch Endpoint	244
14.1.8.3. OpenSearch URL	244
14.1.8.3.1. From Code:	244
14.1.8.3.2. From a Web Browser:	244
14.1.8.3.3. Parameter List	244
14.1.8.3.4. Supported Complex Contextual Query Format	247
14.2. Endpoint Utility Services	248
14.2.1. Compression Services	248
14.2.1.1. Installing a Compression Service	248
14.2.1.2. Configuring Compression Services	248
15. Eventing	249
15.1. Eventing Intro	249
15.2. Eventing Components	249
15.3. Subscriptions	250
15.3.1. Subscriptions	250
15.3.1.1. Subscription Lifecycle	250
15.3.1.1.1. Creation	250

15.3.1.1.2. Evaluation	250
15.3.1.1.3. Update Evaluation	250
15.3.1.1.4. Durability	250
15.3.2. Creating a Subscription	251
15.3.2.1. Event Processing and Notification	251
15.3.2.1.1. Using DDF Implementation	252
15.3.2.2. Delivery Method	252
16. Security Services	252
16.1. Encryption Service	252
16.1.1. Encryption Command	253
16.2. Expansion Service	253
16.2.1. Configuring Expansion Service	253
16.2.1.1. Expansion Service Instances and Configuration	253
16.2.1.2. Expansion Command Examples and Explanation	254
16.2.1.2.1. security:expansions	254
16.2.1.2.2. security:expand	255
16.3. Security IdP	255
16.4. Security STS	256
16.4.1. Security STS Client Config	256
16.4.1.1. Installing the Security STS Client Config	256
16.4.1.2. Configuring the Security STS Client Config	256
16.4.2. External/WS-S STS Support	257
16.4.2.1. Security STS WSS	257
16.4.2.2. Security STS Address Provider	257
16.4.3. Security STS LDAP Login	257
16.4.3.1. Installing the Security STS LDAP Login	257
16.4.3.2. Configuring the Security STS LDAP Login	257
16.4.4. Security STS LDAP Claims Handler	258
16.4.4.1. Installing Security STS LDAP Claims Handler	258
16.4.4.2. Configuring the Security STS LDAP Claims Handler	258
16.4.5. Security STS Server	259
16.4.5.1. Installing the Security STS Server	259
16.4.5.2. Configuring the Security STS Server	259
16.4.6. Security STS Service	260
16.4.6.1. Installing the Security STS Realm	260
16.4.6.2. Configuring the Security STS Realm	260
Developing	260
17. Catalog Framework API	261
17.1. Catalog API Design	263
17.1.1. Ensuring Compatibility	263
17.1.2. Catalog Framework Sequence Diagrams	263

17.1.2.1. Error Handling	264
17.1.2.2. Query	264
17.1.2.3. Product Retrieval	265
17.1.2.4. Product Caching	265
17.1.2.5. Product Download Status	265
17.1.3. Catalog API	266
17.1.3.1. Catalog API Search Interfaces	266
17.1.3.2. Catalog Search Result Objects	266
17.1.3.3. Search Programmatic Flow	266
17.1.3.4. Sort Policies	267
17.1.3.5. Asynchronous Search & Retrieval	267
17.1.3.6. Product Retrieval	268
17.1.3.7. Notifications and Activities	269
17.1.3.7.1. Notifications	269
17.1.3.7.2. Activities	269
17.2. Included Catalog Frameworks, Associated Components, and Configurations	269
17.2.1. Standard Catalog Framework	270
17.2.1.1. Installing the Standard Catalog Framework	270
17.2.1.2. Configuring the Standard Catalog Framework	270
17.2.1.3. Known Issues with Standard Catalog Framework	271
17.2.2. Catalog Framework Camel Component	271
17.2.2.1. Message Headers	272
17.2.2.1.1. Catalog Framework Producer	272
17.2.2.2. Sending Messages to Catalog Framework Endpoint	272
17.2.2.2.1. Catalog Framework Producer	272
17.2.2.2.2. Samples	272
18. Transformers	273
18.1. Available Input Transformers	274
18.2. Available Metocard Transformers	275
18.3. Available Query Response Transformers	276
18.4. Transformers Details	276
18.4.1. Atom Query Response Transformer	276
18.4.1.1. Installing the Atom Query Response Transformer	276
18.4.1.2. Configuring the Atom Query Response Transformer	276
18.4.1.3. Using the Atom Query Response Transformer	277
18.4.2. CSW Query Response Transformer	280
18.4.2.1. Installing the CSW Query Response Transformer	280
18.4.2.2. Configuring the CSW Query Response Transformer	280
18.4.3. GeoJSON Input Transformer	280
18.4.3.1. Installing the GeoJSON Input Transformer	280
18.4.3.2. Configuring the GeoJSON Input Transformer	280

18.4.3.3. Using the GeoJSON Input Transformer	280
18.4.3.4. Conversion to a Metocard	280
18.4.3.4.1. Metocard Extensibility	281
18.4.3.5. Usage Limitations of the GeoJSON Input Transformer	282
18.4.4. GeoJSON Metocard Transformer	282
18.4.4.1. Installing the GeoJSON Metocard Transformer	282
18.4.4.2. Configuring the GeoJSON Metocard Transformer	283
18.4.4.3. Using the GeoJSON Metocard Transformer	283
18.4.5. GeoJSON Query Response Transformer	284
18.4.5.1. Installing the GeoJSON Query Response Transformer	284
18.4.5.2. Configuring the GeoJSON Query Response Transformer	284
18.4.6. HTML Metocard Transformer	285
18.4.6.1. Installing the HTML Metocard Transformer	285
18.4.6.2. Configuring the HTML Metocard Transformer	285
18.4.6.3. Using the HTML Metocard Transformer	285
18.4.7. KML Metocard Transformer	285
18.4.7.1. Installing the KML Metocard Transformer	285
18.4.7.2. Configuring the KML Metocard Transformer	285
18.4.7.3. Using the KML Metocard Transformer	285
18.4.8. KML Query Response Transformer	288
18.4.8.1. Installing the KML Query Response Transformer	288
18.4.8.2. Configuring the KML Query Response Transformer	288
18.4.8.3. Using the KML Query Response Transformer	289
18.4.9. KML Style Mapper	292
18.4.9.1. Installing the KML Style Mapper	292
18.4.9.2. Configuring the KML Style Mapper	292
18.4.10. Metadata Metocard Transformer	294
18.4.10.1. Installing the Metadata Metocard Transformer	294
18.4.10.2. Configuring the Metadata Metocard Transformer	294
18.4.10.3. Using the Metadata Metocard Transformer	294
18.4.11. PDF Input Transformer	294
18.4.11.1. Installing the PDF Input Transformer	294
18.4.11.2. Configuring the PDF Input Transformer	295
18.4.12. PPTX Input Transformer	295
18.4.12.1. Installing the PPTX Input Transformer	295
18.4.12.2. Configuring the PPTX Input Transformer	295
18.4.13. Query Response Transformer Consumer	295
18.4.13.1. Installing the Query Response Transformer Consumer	295
18.4.13.2. Configuring the Query Response Transformer Consumer	295
18.4.14. Registry Transformer	295
18.4.14.1. Installing the Registry Transformer	296

18.4.14.2. Configuring the Registry Transformer	296
18.4.15. Resource Metocard Transformer	296
18.4.15.1. Installing the Resource Metocard Transformer	296
18.4.15.2. Configuring the Resource Metocard Transformer	296
18.4.15.3. Using the Resource Metocard Transformer	296
18.4.16. SearchUI Transformer	296
18.4.16.1. Installing Simple SearchUI	296
18.4.16.2. Configuring Simple SearchUI	297
18.4.16.3. Using Simple SearchUI	297
18.4.16.3.1. Search Criteria	297
18.4.16.4. Usage Limitations of the Simple SearchUI	298
18.4.17. Thumbnail Metocard Transformer	298
18.4.17.1. Installing the Thumbnail Metocard Transformer	298
18.4.17.2. Configuring the Thumbnail Metocard Transformer	298
18.4.17.3. Using the Thumbnail Metocard Transformer	298
18.4.18. Tika Input Transformer	298
18.4.18.1. Installing the Tika Input Transformer	299
18.4.18.2. Configuring the Tika Input Transformer	299
18.4.19. Video Input Transformer	299
18.4.19.1. Installing the Video Input Transformer	299
18.4.19.1.1. Configuring the Video Input Transformer	299
18.4.20. XML Input Transformer	299
18.4.20.1. Installing the XML Input Transformer	299
18.4.20.2. Configuring the XML Input Transformer	299
18.4.21. XML Metocard Transformer	299
18.4.21.1. Installing the XML Metocard Transformer	299
18.4.21.2. Configuring the XML Metocard Transformer	300
18.4.21.3. Using the XML Metocard Transformer	300
18.4.22. XML Query Response Transformer	300
18.4.22.1. Installing the XML Query Response Transformer	301
18.4.22.2. Configuring the XML Query Response Transformer	301
18.4.22.3. Using the XML Query Response Transformer	301
18.5. Mime Type Mapper	302
18.5.1. DDF Mime Type Mapper	303
18.5.1.1. Installing the DDF Mime Type Mapper	303
18.5.1.2. Configuring DDF Mime Type Mapper	303
18.6. Mime Type Resolver	303
18.6.1. Custom Mime Type Resolver	303
18.6.1.1. Installing the Custom Mime Type Resolver	304
18.6.1.1.1. Configuring the Custom Mime Type Resolver	304
18.6.2. Tika Mime Type Resolver	304

18.6.2.1. Installing the Tika Mime Type Resolver	304
18.6.2.1.1. Configuring the Tika Mime Type Resolver	305
19. Catalog Plugins	305
19.1. Types of Plugins	305
19.1.1. Pre-Authorization Plugins	316
19.1.1.1. Available Pre-Authorization Plugins	316
19.1.2. Policy Plugins	316
19.1.2.1. Available Policy Plugins	317
19.1.3. Access Plugins	317
19.1.3.1. Available Access Plugins	317
19.1.4. Pre-Ingest Plugins	318
19.1.4.1. Available Pre-Ingest Plugins	318
19.1.5. Post-Ingest Plugins	319
19.1.5.1. Available Post-Ingest Plugins	319
19.1.6. Post-Process Plugins	320
19.1.6.1. Available Post-Process Plugins	320
19.1.7. Pre-Query Plugins	320
19.1.7.1. Available Pre-Query Plugins	320
19.1.8. Pre-Federated-Query Plugins	321
19.1.8.1. Available Pre-Federated-Query Plugins	321
19.1.9. Post-Query Plugins	321
19.1.9.1. Available Post-Query Plugins	321
19.1.10. Post-Federated-Query Plugins	322
19.1.10.1. Available Post-Federated-Query Plugins	322
19.1.11. Pre-Resource Plugins	322
19.1.11.1. Available Pre-Resource Plugins	322
19.1.12. Post-Resource Plugins	322
19.1.12.1. Available Post-Resource Plugins	323
19.1.13. Pre-Create Storage Plugins	323
19.1.13.1. Available Pre-Create Storage Plugins	323
19.1.14. Post-Create Storage Plugins	323
19.1.14.1. Available Post-Create Storage Plugins	323
19.1.15. Pre-Update Storage Plugins	323
19.1.15.1. Available Pre-Update Storage Plugins	324
19.1.16. Post-Update Storage Plugins	324
19.1.16.1. Available Post-Update Storage Plugins	324
19.1.17. Pre-Subscription Plugins	324
19.1.17.1. Available Pre-Subscription Plugins	324
19.1.18. Pre-Delivery Plugins	324
19.1.18.1. Available Pre-Delivery Plugins	325
19.2. Catalog Plugin Details	325

19.2.1. Catalog Backup Plugin	325
19.2.1.1. Installing the Catalog Backup Plugin	325
19.2.1.2. Configuring the Catalog Backup Plugin	325
19.2.1.3. Usage Limitations of the Catalog Backup Plugin	325
19.2.2. Catalog Metrics Plugin	325
19.2.2.1. Related Components to the Source Metrics Plugin	325
19.2.2.2. Installing the Catalog Metrics Plugin	326
19.2.2.3. Configuring the Catalog Metrics Plugin	326
19.2.3. Catalog Policy Plugin	326
19.2.3.1. Installing the Catalog Policy Plugin	326
19.2.3.2. Configuring the Catalog Policy Plugin	326
19.2.4. Checksum Plugin	326
19.2.4.1. Installing the Checksum Plugin	326
19.2.4.2. Configuring the Checksum Plugin	326
19.2.5. Client Info Plugin	327
19.2.5.1. Related Components to the Client Info Plugin	327
19.2.5.2. Installing the Client Info Plugin	327
19.2.5.3. Configuring the Client Info Plugin	327
19.2.6. Content URI Access Plugin	327
19.2.6.1. Installing the Content URI Access Plugin	327
19.2.6.2. Configuring the Content URI Access Plugin	327
19.2.7. Event Processor	327
19.2.7.1. Installing the Event Processor	327
19.2.7.2. Configuring the Event Processor	328
19.2.7.3. Usage Limitations of the Event Processor	328
19.2.8. Expiration Date Pre-Ingest Plugin	328
19.2.8.1. Installing the Expiration Date Pre-Ingest Plugin	328
19.2.8.2. Configuring the Expiration Date Pre-Ingest Plugin	328
19.2.9. Filter Plugin	328
19.2.9.1. Installing the Filter Plugin	330
19.2.9.2. Configuring the Filter Plugin	330
19.2.10. GeoCoder Plugin	330
19.2.10.1. Installing the GeoCoder Plugin	330
19.2.10.2. Configuring the GeoCoder Plugin	330
19.2.11. Historian Policy Plugin	330
19.2.11.1. Installing the Historian Policy Plugin	330
19.2.11.2. Configuring the Historian Policy Plugin	331
19.2.12. Identification Plugin	331
19.2.12.1. Installing the Identification Plugin	331
19.2.12.2. Configuring the Identification Plugin	331
19.2.13. JPEG2000 Thumbnail Converter	331

19.2.13.1. Installing the JPEG2000 Thumbnail Converter	331
19.2.13.2. Configuring the JPEG2000 Thumbnail Converter	331
19.2.14. Metacard Attribute Security Policy Plugin	331
19.2.14.1. Installing the Metacard Attribute Security Policy Plugin	332
19.2.14.2. Configuring the Metacard Attribute Security Policy Plugin	332
19.2.14.3. Sample configuration of the Metacard Attribute Security Policy Plugin	332
19.2.15. Metacard Backup File Storage Provider	332
19.2.15.1. Installing the Metacard Backup File Storage Provider	332
19.2.15.2. Configuring the Metacard Backup File Storage Provider	333
19.2.16. Metacard Backup S3 Storage Provider	333
19.2.16.1. Installing the Metacard S3 File Storage Provider	333
19.2.16.2. Configuring the Metacard S3 File Storage Provider	333
19.2.17. Metacard Groomer	334
19.2.17.1. Installing the Metacard Groomer	334
19.2.17.2. Configuring the Metacard Groomer	334
19.2.18. Metacard Resource Size Plugin	334
19.2.18.1. Installing the Metacard Resource Size Plugin	334
19.2.18.2. Configuring the Metacard Resource Size Plugin	334
19.2.19. Metacard Validity Filter Plugin	335
19.2.19.1. Related Components to the Metacard Validity Filter Plugin	335
19.2.19.2. Installing the Metacard Validity Filter Plugin	335
19.2.19.3. Configuring the Metacard Validity Filter Plugin	335
19.2.20. Metacard Validity Marker	335
19.2.20.1. Related Components to the Metacard Validity Marker	335
19.2.20.2. Installing Metacard Validity Marker	335
19.2.20.3. Configuring Metacard Validity Marker	336
19.2.20.4. Using Metacard Validity Marker	336
19.2.21. Metacard Ingest Network Plugin	336
19.2.21.1. Related Components to the Metacard Ingest Network Plugin	336
19.2.21.2. Installing the Metacard Ingest Network Plugin	336
19.2.21.3. Configuring the Metacard Ingest Network Plugin	336
19.2.21.3.1. Useful Attributes	337
19.2.21.4. Usage Limitations of the Metacard Ingest Network Plugin	338
19.2.22. Operation Plugin	338
19.2.22.1. Installing the Operation Plugin	338
19.2.22.2. Configuring the Operation Plugin	338
19.2.23. Point of Contact Policy Plugin	338
19.2.23.1. Related Components to Point of Contact Policy Plugin	338
19.2.23.2. Installing the Point of Contact Policy Plugin	338
19.2.23.3. Configuring the Point of Contact Policy Plugin	338
19.2.24. Processing Post-Ingest Plugin	338

19.2.24.1. Related Components to Processing Post-Ingest Plugin	339
19.2.24.2. Installing the Processing Post-Ingest Plugin	339
19.2.24.3. Configuring the Processing Post-Ingest Plugin	339
19.2.25. Registry Policy Plugin	339
19.2.25.1. Installing the Registry Policy Plugin	339
19.2.25.2. Configuring the Registry Policy Plugin	339
19.2.26. Resource URI Policy Plugin	339
19.2.26.1. Installing the Resource URI Policy Plugin	339
19.2.26.2. Configuring the Resource URI Policy Plugin	340
19.2.27. Resource Usage Plugin	340
19.2.27.1. Installing the Resource Usage Plugin	340
19.2.27.2. Configuring the Resource Usage Plugin	340
19.2.28. Security Audit Plugin	340
19.2.28.1. Installing the Security Audit Plugin	340
19.2.28.2. Configuring the Security Audit Plugin	340
19.2.29. Security Logging Plugin	341
19.2.29.1. Installing Security Logging Plugin	341
19.2.29.2. Configuring the Security Logging Plugin	341
19.2.30. Security Plugin	341
19.2.30.1. Installing the Security Plugin	341
19.2.30.2. Configuring the Security Plugin	341
19.2.31. Source Metrics Plugin	341
19.2.31.1. Related Components to the Source Metrics Plugin	341
19.2.31.2. Installing the Source Metrics Plugin	342
19.2.31.3. Configuring the Source Metrics Plugin	342
19.2.32. Tags Filter Plugin	342
19.2.32.1. Related Components to Tags Filter Plugin	342
19.2.32.2. Installing the Tags Filter Plugin	342
19.2.32.3. Configuring the Tags Filter Plugin	342
19.2.33. Video Thumbnail Plugin	342
19.2.33.1. Installing the Video Thumbnail Plugin	342
19.2.33.2. Configuring the Video Thumbnail Plugin	343
19.2.34. Workspace Access Plugin	343
19.2.34.1. Related Components to The Workspace Access Plugin	343
19.2.34.2. Installing the Workspace Access Plugin	343
19.2.34.3. Configuring the Workspace Access Plugin	343
19.2.35. Workspace Pre-Ingest Plugin	343
19.2.35.1. Related Components to The Workspace Pre-Ingest Plugin	343
19.2.35.2. Installing the Workspace Pre-Ingest Plugin	344
19.2.35.3. Configuring the Workspace Pre-Ingest Plugin	344
19.2.36. Workspace Sharing Policy Plugin	344

19.2.36.1. Related Components to The Workspace Sharing Policy Plugin	344
19.2.36.2. Installing the Workspace Sharing Policy Plugin	344
19.2.36.3. Configuring the Workspace Sharing Policy Plugin	344
19.2.37. XML Attribute Security Policy Plugin	344
19.2.37.1. Installing the XML Attribute Security Policy Plugin	344
19.2.37.2. Configuring the XML Attribute Security Policy Plugin	345
20. Operations	345
21. Resources	346
21.1. Content Item	347
21.1.1. Retrieving Resources	347
21.1.1.1. BinaryContent	348
21.1.2. Retrieving Resource Options	348
21.1.3. Storing Resources	349
21.2. Resource Components	349
21.3. Resource Readers	350
21.3.1. URL Resource Reader	350
21.3.1.1. Installing the URL Resource Reader	350
21.3.1.2. Configuring the URL Resource Reader	350
21.3.1.3. Using the URL Resource Reader	351
21.3.1.4. Usage Limitations of the URL Resource Reader	351
21.4. Resource Writers	351
21.5. Content Directory Monitor	352
21.5.1. Installing the Content Directory Monitor	352
21.5.2. Configuring the Content Directory Monitor	352
21.5.3. Using the Content Directory Monitor	352
22. Queries	354
22.1. Filters	354
22.1.1. FilterBuilder API	355
22.1.2. Boolean Operators	355
22.1.3. Attribute	355
22.1.4. XPath	355
23. Metrics	356
23.1. Metrics Collection Application	356
23.1.1. Installing Metrics Collection	356
23.1.2. Configuring Metrics Collection	357
23.1.3. Catalog Metrics	357
23.1.4. Source Metrics	358
23.2. Metrics Reporting Application	359
23.2.1. Metrics Aggregate Reports	361
23.2.2. Viewing Metrics	362
24. Action Framework	362

24.1. Action Providers	363
25. Asynchronous Processing Framework	363
26. Migration API	366
26.1. The Migration API Interfaces and Classes	367
26.1.1. Migratable	367
26.1.2. MigrationContext	369
26.1.3. ExportMigrationContext	369
26.1.4. ImportMigrationContext	369
26.1.5. MigrationEntry	370
26.1.6. ExportMigrationEntry	370
26.1.7. ImportMigrationEntry	371
26.1.8. MigrationOperation	372
26.1.9. MigrationReport	372
26.1.10. MigrationMessage	373
26.1.11. MigrationException	373
26.1.12. MigrationWarning	373
26.1.13. MigrationInformation	374
26.1.14. MigrationSuccessfulInformation	374
27. Application Service API	374
27.1. JMX Managed Bean	374
27.2. Initial Application Installation	374
27.3. Application Service Interfaces	375
27.3.1. ApplicationService Interface	375
27.3.2. Application Interface	376
27.3.3. ApplicationStatus Interface	376
27.3.4. ApplicationNode Interface	376
28. Security Framework	377
28.1. Subject	377
28.1.1. Security Manager	377
28.1.2. Realms	378
28.1.2.1. Authenticating Realms	378
28.1.2.2. Authorizing Realms	378
28.2. Security Core	380
28.2.1. Security Core API	380
28.2.1.1. Installing the Security Core API	380
28.2.1.2. Configuring the Security Core API	380
28.2.2. Security Core Implementation	380
28.2.2.1. Installing the Security Core Implementation	381
28.2.2.2. Configuring the Security Core Implementation	381
28.2.3. Security Core Commons	381
28.2.3.1. Configuring the Security Core Commons	381

28.3. Security Encryption	381
28.3.1. Security Encryption API	381
28.3.1.1. Installing Security Encryption API	381
28.3.1.2. Configuring the Security Encryption API	381
28.3.2. Security Encryption Implementation	381
28.3.2.1. Installing Security Encryption Implementation	382
28.3.2.2. Configuring Security Encryption Implementation	382
28.3.3. Security Encryption Commands	382
28.3.3.1. Installing the Security Encryption Commands	382
28.3.3.2. Configuring the Security Encryption Commands	382
28.4. Security LDAP	382
28.4.1. Embedded LDAP Server	382
28.4.1.1. Installing the Embedded LDAP Server	383
28.4.1.2. Configuring the Embedded LDAP	383
28.4.1.3. Connecting to Standalone LDAP Servers	383
28.4.1.4. Embedded LDAP Configuration	383
28.4.1.5. Schemas	384
28.4.1.6. Starting and Stopping the Embedded LDAP	385
28.4.1.7. Limitations of the Embedded LDAP	386
28.4.1.8. External Links for the Embedded LDAP	386
28.4.1.9. LDAP Administration	386
28.4.1.10. Downloading the Admin Tools	386
28.4.1.11. Using the Admin Tools	386
28.5. Security PDP	388
28.5.1. Security PDP AuthZ Realm	388
28.5.1.1. Configuring the Security PDP AuthZ Realm	389
28.5.2. Guest Interceptor	389
28.5.2.1. Installing Guest Interceptor	389
28.5.2.2. Configuring Guest Interceptor	389
28.6. Web Service Security Architecture	390
28.6.1. Securing REST	390
28.6.2. Securing SOAP	391
28.6.2.1. SOAP Secure Client	391
28.6.2.2. Policy-unaware SOAP Client	391
28.7. Security PEP	392
28.7.1. Security PEP Interceptor	392
28.7.1.1. Installing the Security PEP Interceptor	392
28.7.1.2. Configuring the Security PEP Interceptor	392
28.8. Filtering	392
28.9. Expansion Service	393
28.10. Security Token Service	394

28.10.1. STS Claims Handlers	395
29. Developing DDF Components	395
29.1. Developing Complementary Catalog Frameworks	395
29.1.1. Simple Catalog API Implementations	395
29.1.2. Use of the Whiteboard Design Pattern	395
29.1.3. Recommendations for Framework Development	396
29.1.4. Catalog Framework Reference	396
29.1.4.1. Methods	396
29.1.4.1.1. Create, Update, and Delete Methods	396
29.1.4.1.2. Query Methods	397
29.1.4.1.3. Resource Methods	397
29.1.4.1.4. Source Methods	397
29.1.4.1.5. Transform Methods	398
29.1.4.2. Implementing Catalog Methods	398
29.1.4.3. Dependency Injection	399
29.1.4.4. OSGi Service Registry	399
29.2. Developing Metocard Types	400
29.2.1. Metocard Type Definition File	400
29.3. Developing Global Attribute Validators	402
29.3.1. Global Attribute Validators File	402
29.4. Developing Attribute Types	405
29.4.1. Attribute Type Definition File	406
29.5. Developing Default Attribute Types	407
29.5.1. Default Attribute Values	407
29.6. Developing Attribute Injections	409
29.6.1. Attribute Injection Definition	409
29.7. Developing Endpoints	410
29.8. Developing Input Transformers	411
29.8.1. Create an XML Input Transformer using SaxEventHandlers	412
29.8.2. Create an Input Transformer Using Apache Camel	414
29.8.2.1. Input Transformer Design Pattern (Camel)	414
29.9. Developing Metocard Transformers	415
29.9.1. Creating a New Metocard Transformer	415
29.10. Developing Query Response Transformers	416
29.11. Developing Sources	418
29.11.1. Implement a Source Interface	418
29.11.1.1. Developing Catalog Providers	419
29.11.1.2. Developing Federated Sources	419
29.11.1.3. Developing Connected Sources	420
29.11.1.4. Exception Handling	420
29.11.1.4.1. Exception Examples	420

29.11.1.4.2. External Resources for Developing Sources	421
29.12. Developing Catalog Plugins	421
29.12.1. Implementing Catalog Plugins	422
29.12.1.1. Catalog Plugin Failure Behavior	423
29.12.1.2. Implementing Pre-Ingest Plugins	423
29.12.1.3. Implementing Post-Ingest Plugins	423
29.12.1.4. Implementing Pre-Query Plugins	424
29.12.1.5. Implementing Post-Query Plugins	424
29.12.1.6. Implementing Pre-Delivery Plugins	424
29.12.1.7. Implementing Pre-Subscription Plugins	425
29.12.1.8. Implementing Pre-Resource Plugins	425
29.12.1.9. Implementing Post-Resource Plugins	425
29.12.1.10. Implementing Policy Plugins	426
29.12.1.11. Implementing Access Plugins	426
29.13. Developing Token Validators	427
29.14. Developing STS Claims Handlers	427
29.14.1. Example Requests and Responses for SAML Assertions	435
29.14.2. BinarySecurityToken (CAS) SAML Security Token Samples	435
29.14.3. UsernameToken Bearer SAML Security Token Sample	441
29.14.4. X.509 PublicKey SAML Security Token Sample	447
29.14.5. X.509 PublicKey SAML Security Token Sample	449
29.15. Developing Registry Clients	454
29.16. Developing Resource Readers	454
29.16.1. Creating a New ResourceReader	454
29.16.1.1. Implementing the ResourceReader Interface	455
29.16.1.2. retrieveResource	455
29.16.1.3. Implement retrieveResource()	455
29.16.1.4. getSupportedSchemes	456
29.16.1.5. Export to OSGi Service Registry	456
29.17. Developing Resource Writers	457
29.17.1. Create a New ResourceWriter	457
29.18. Developing Filters	459
29.18.1. Units of Measure	459
29.18.2. Filter Examples	459
29.18.2.1. Contextual Searches	460
29.18.2.1.1. Tree View of Creating Filters	461
29.18.2.1.2. XML View of Creating Filters	462
29.18.2.2. Fuzzy Operations	462
29.18.3. Parsing Filters	463
29.18.3.1. Interpreting a Filter to Create SQL	463
29.18.3.2. Interpreting a Filter to Create XQuery	463

29.18.3.2.1. FilterAdapter/Delegate Process for Figure Parsing	464
29.18.3.2.2. FilterVisitor Process for Figure Parsing	465
29.18.4. Filter Profile	465
29.18.4.1. Role of the OGC Filter	465
29.18.4.2. Catalog Filter Profile	465
29.18.4.2.1. Comparison Operators	466
29.18.4.2.2. Logical Operators	467
29.18.4.2.3. Temporal Operators	468
29.18.4.2.4. Spatial Operators	468
29.19. Developing Filter Delegates	469
29.19.1. Creating a New Filter Delegate	469
29.19.1.1. Implementing the Filter Delegate	469
29.19.1.2. Throwing Exceptions	469
29.19.1.3. Using the Filter Adapter	469
29.19.1.4. Filter Support	470
29.20. Developing Action Components	471
29.20.1. Action Component Naming Convention	471
29.20.1.1. Action Component Taxonomy	472
29.21. Developing Query Options	472
29.21.1. Evaluating a query	473
29.21.2. Commons-DDF Utilities	474
29.21.2.1. FuzzyFunction	474
29.21.2.2. XPathHelper	474
29.22. Configuring Managed Service Factory Bundles	474
29.22.1. Configuring Managed Service Factory Bundles	474
29.22.1.1. File Format	474
29.23. Developing XACML Policies	478
29.23.1. XACML Policy Attributes	479
29.23.2. XACML Policy Subject	479
29.23.3. XACML Policy Resource	479
29.23.4. Using a XACML Policy	479
29.24. Assuring Authenticity of Bundles and Applications	479
29.24.1. Prerequisites	479
29.24.2. Signing a JAR/KAR	480
29.24.2.1. Verifying a JAR/KAR	480
29.25. WFS Services	481
30. Development Guidelines	483
30.1. Contributing	483
30.2. OSGi Basics	483
30.2.1. Packaging Capabilities as Bundles	484
30.2.1.1. Creating a Bundle	484

30.2.1.1.1. Bundle Development Recommendations	484
30.2.1.1.2. Maven Bundle Plugin	485
30.2.1.2. Third Party and Utility Bundles	486
30.2.1.3. Deploying a Bundle	486
30.2.1.4. Verifying Bundle State	487
30.2.2. Working with Features Files	487
30.2.2.1. Auto-starting an Application Feature	488
30.3. High Availability Guidance	488
Appendices	488
Appendix A: Application References	489
A.1. Admin Application Reference	489
A.1.1. Admin Application Prerequisites	489
A.1.2. Installing the Admin Application	489
A.1.3. Configuring the Admin Application	489
A.2. Message Broker Application Reference	491
A.2.1. Prerequisites for Message Broker Application	491
A.2.2. Installing Message Broker Application	491
A.2.3. Configuring the Message Broker Application	491
A.2.3.1. Configuring the Message Broker for a Highly Available Cluster	491
A.2.3.2. Securing the Message Broker Application	494
A.2.3.3. Artemis Broker Connection Configuration	495
A.2.4. Using the Message Broker Application	495
A.2.4.1. Undelivered Messages UI	495
A.2.4.2. Route Manager	496
A.3. Catalog Application Reference	496
A.3.1. Catalog Application Prerequisites	497
A.3.2. Installing the Catalog Application	497
A.3.3. Configuring the Catalog Application	497
A.4. GeoWebCache Application Reference	514
A.4.1. GeoWebCache Application Prerequisites	514
A.4.2. Installing GeoWebCache	514
A.4.3. Configuring GeoWebCache	515
A.4.3.1. Adding GeoWebCache Layers	515
A.4.3.2. Editing GeoWebCache Layers	515
A.4.3.3. Removing GeoWebCache Layers	515
A.4.3.4. Configuring GWC Disk Quota	515
A.4.3.5. Configuring the Standard Search UI for GeoWebCache	516
A.5. Intrigue Application Reference	517
A.5.1. Intrigue Prerequisites	517
A.5.2. Installing Intrigue	517
A.5.3. Configuring the Intrigue Application	517

A.6. Platform Application Reference	524
A.6.1. Platform Application Prerequisites	524
A.6.2. Installing Platform	524
A.6.3. Configuring the Platform Application	524
A.7. Registry Application Reference	529
A.7.1. Registry Prerequisites	529
A.7.2. Installing Registry	529
A.7.3. Customizing Registry Fields	529
A.7.4. Using Registry	531
A.7.4.1. Configuring Identity Node	531
A.7.4.1.1. Adding a Service Binding to a Node	533
A.7.4.2. Publishing to Others	534
A.7.4.3. Subscribing to Another Node	534
A.7.4.3.1. Editing a Subscription	534
A.7.4.3.2. Deleting a Subscription	535
A.7.5. Configuring the Registry Application	535
A.8. Resource Management Application Reference	537
A.8.1. Resource Management Prerequisites	538
A.8.2. Installing Resource Management	538
A.8.3. Configuring the Resource Management Application	538
A.9. Security Application Reference	539
A.9.1. Security Prerequisites	539
A.9.2. Installing Security	539
A.9.3. Configuring the Security Application	539
A.10. Solr Catalog Application Reference	553
A.10.1. Solr Catalog Prerequisites	553
A.10.2. Installing Solr Catalog	553
A.10.3. Configuring the Solr Catalog Application	553
A.11. Spatial Application Reference	553
A.11.1. Offline Gazetteer Service	554
A.11.1.1. Spatial GeoNames Console Commands	554
A.11.2. Spatial Prerequisites	555
A.11.3. Installing Spatial	555
A.11.4. Configuring the Spatial Application	555
A.12. Search UI Application Reference	569
A.12.1. Search UI Prerequisites	569
A.12.2. Installing Search UI	570
A.12.3. Configuring the Search UI Application	570
Appendix B: Application Whitelists	572
B.1. Admin Whitelist	572
B.2. Catalog Whitelist	572

B.3. Platform Whitelist	576
B.4. Registry Whitelist	577
B.5. Security Whitelist	577
B.6. Solr Catalog Whitelist	578
B.7. Spatial Whitelist	578
B.8. Search UI Whitelist	579
B.9. Packages Removed From Whitelist	579
Appendix C: Metadata Attributes	581
C.1. Format-specific Attribute Mappings	581
C.1.1. Mp4 Additional Attribute	581
C.2. Catalog Taxonomy	581
C.2.1. Core Attributes Table	581
C.2.2. Associations Attributes Table	584
C.2.3. Contact Attributes Table	585
C.2.4. DateTime Attributes Table	586
C.2.5. History Attributes Table	586
C.2.6. Location Attributes Table	587
C.2.7. Media Attributes Table	587
C.2.8. Metocard Attributes Table	589
C.2.9. Security Attributes Table	589
C.2.10. Topic Attributes Table	590
C.2.11. Validation Attributes Table	590
C.3. All Formats Supported	590
Appendix D: DDF Dependency List	604

License

Copyright (c) Codice Foundation.

This work is licensed under a [Creative Commons Attribution 4.0 International License](#).

This document last updated: 2017-12-12.

Introduction

1. About DDF

1.1. Introducing DDF

Distributed Data Framework (DDF) is a free and open-source common data layer that abstracts services and business logic from underlying data structures to enable rapid integration of new data sources.

Licensed under [LGPL](#), DDF is an interoperability platform that provides secure and scalable discovery and retrieval from a wide array of disparate sources.

DDF is:

- a flexible and modular integration framework.
- built to "unzip and run" even when scaled to large enterprise systems.
- primarily focused on data integration, enabling clients to insert, query, and transform information from disparate data sources via the DDF Catalog.

1.2. Component Applications

DDF is comprised of several modular applications, to be installed or uninstalled as needed.

Admin Application

Enhances administrative capabilities when installing and managing DDF. It contains various services and interfaces that allow administrators more control over their systems.

Catalog Application

Provides a framework for storing, searching, processing, and transforming information. Clients typically perform local and/or federated query, create, read, update, and delete (QCRUD) operations against the Catalog. At the core of the Catalog functionality is the **Catalog Framework**, which routes all requests and responses through the system, invoking additional processing per the system configuration.

Platform Application

The Core application of the distribution. The Platform application contains the fundamental building blocks to run the distribution.

Security Application

Provides authentication, authorization, and auditing services for the DDF. It is both a framework that developers and integrators can extend and a reference implementation that meets security requirements.

Solr Catalog Application

Includes the Solr Catalog Provider, an implementation of the Catalog Provider using [Apache Solr](#) as a data store.

Spatial Application

Provides OGC services, such as [CSW](#), [WCS](#), [WFS](#), and [KML](#).

Search UI

Allows a user to search for records in the local Catalog (provider) and federated sources. Results of the search are returned and displayed on a globe or map, providing a visual representation of where the records were found.

1.3. Documentation Guide

The DDF documentation is organized by audience.

Core Concepts

This introduction section is intended to give a high level overview of the concepts and capabilities of DDF.

Administrators

[Managing](#) | Administrators will be installing, maintaining, and supporting existing applications. Use this section to [prepare](#), [install](#), [configure](#), [run](#), and [monitor](#) a DDF.

Users

[Using](#) | Users interact with the system to search data stores. Use this section to navigate the various user interfaces available in DDF.

Integrators

[Integrating](#) | Integrators will use the existing applications to support their external frameworks. This section will provide details for finding, accessing and using the components of DDF.

Developers

[Developing](#) | Developers will build or extend the functionality of the applications.

1.3.1. Documentation Conventions

The following conventions are used within this documentation:

1.3.1.1. Customizable Values

Many values used in descriptions are customizable and should be changed for specific use cases. These values are denoted by `< >`, and by `[[]]` when within XML syntax. When using a real value, the placeholder characters should be omitted.

1.3.1.2. Code Values

Java objects, lines of code, or file properties are denoted with the [Monospace](#) font style. Example: `ddf.catalog.CatalogFramework`

1.3.1.3. Hyperlinks

Some hyperlinks (e.g., `/admin`) within the documentation assume a locally running installation of DDF. Simply change the hostname if accessing a remote host.

1.3.2. Support

Questions about DDF should be posted to the [ddf-users forum](#) or [ddf-developers forum](#), where they will be responded to quickly by a member of the DDF team.

1.3.2.1. Documentation Updates

The most current DDF documentation is available at [DDF Documentation](#).

1.4. Core Concepts

1.4.1. Introduction to Search

DDF provides the capability to search the Catalog for metadata. There are a number of different types of searches that can be performed on the Catalog, and these searches are accessed using one of several interfaces. This section provides a very high level overview of introductory concepts of searching with DDF. These concepts are expanded upon in later sections.

1.4.1.1. Search Types

There are four basic types of metadata search. Additionally, any of the types can be combined to create a compound search.

Text Search

A text search is used when searching for textual information. It searches all textual fields by default, although it is possible to refine searches to a text search on a single metadata attribute. Text searches may use wildcards, logical operators, and approximate matches.

Spatial Search

A spatial search is used for Area of Interest (AOI) searches. Polygon and point radius searches are supported.

Temporal Search

A temporal search finds information from a specific time range. Two types of temporal searches are supported: *relative* and *absolute*. Relative searches contain an offset from the current time, while absolute searches contain a start and an end timestamp. Temporal searches can use the `created` or `modified` date attributes.

Datatype Search

A datatype search is used to search for metadata based on the datatype of the resource. Wildcards (*) can be used in both the datatype and version fields. Metadata that matches any of the datatypes (and associated versions if specified) will be returned. If a version is not specified, then all metadata records for the specified datatype(s) regardless of version will be returned.

1.4.2. Introduction to Metadata

In DDF, [resources](#) are the data products, files, reports, or documents of interest to users of the system.

Metadata is information about those resources, organized into a schema to make search possible. The Catalog stores this metadata and allows access to it. Metacards are single instances of metadata, representing a single resource, in the Catalog. Metacards follow one of several schemas to ensure reliable, accurate, and complete metadata. Essentially, Metacards function as containers of metadata.

1.4.3. Introduction to Ingest

Ingest is the process of bringing data products, metadata, or both into the catalog to enable search, sharing, and discovery. Ingested files are [transformed](#) into a neutral format that can be searched against as well as migrated to other formats and systems. See [Ingesting Data](#) for the various methods of ingesting data.

Upon ingest, a transformer will read the metadata from the ingested file and populate the fields of a metocard. Exactly how this is accomplished depends on the origin of the data, but most fields (except id) are imported directly.

1.4.4. Introduction to Resources

The Catalog Framework can interface with [Storage Providers](#) to provide storage of resources to specific types of storage, e.g., file system, relational database, XML database. A default file system implementation is provided by default.

Storage providers act as a proxy between the Catalog Framework and the mechanism storing the content. Storage providers expose the storage mechanism to the Catalog Framework. Storage plugins provide pluggable functionality that can be executed either immediately before or immediately after content has been stored or updated.

Storage providers provide the capability to the Catalog Framework to create, read, update, and delete resources in the content repository.

See [Data Management](#) for more information on specific file types supported by DDF.

1.4.5. Introduction to the Catalog Framework

The Catalog Framework wires all the Catalog components together.

It is responsible for routing Catalog requests and responses to the appropriate source, destination, federated system, etc.

[Endpoints](#) send Catalog requests to the Catalog Framework. The Catalog Framework then invokes [Catalog Plugins](#), [Transformers](#), and [Resource Components](#) as needed before sending requests to the intended destination, such as one or more [Sources](#).

The Catalog Framework decouples clients from service implementations and provides integration

points for Catalog Plugins and convenience methods for Endpoint developers.

1.4.6. Introduction to Federation

Federation is the ability of the DDF to query other data sources, including other DDFs. By default, the DDF is able to federate using [OpenSearch](#) and [CSW](#) protocols. The minimum configuration necessary to configure those federations is a query address.

Federation enables constructing dynamic networks of data sources that can be queried individually, or aggregated into specific configuration to enable a wider range of accessibility for data and data products.

Federation provides the capability to extend the DDF enterprise to include [Remote Sources](#), which may include other instances of DDF. The Catalog handles all aspects of federated queries as they are sent to the Catalog Provider and Remote Sources, as they are processed, and as the query results are returned. Queries can be scoped to include only the local Catalog Provider (and any Connected Sources), only specific Federated Sources, or the entire enterprise (which includes all local and Remote Sources). If the query is federated, the Catalog Framework passes the query to a Federation Strategy, which is responsible for querying each federated source that is specified. The Catalog Framework is also responsible for receiving the query results from each federated source and returning them to the client in the order specified by the particular federation strategy used. After the federation strategy handles the results, the Catalog returns them to the client through the Endpoint. Query results are returned from a federated query as a list of metacards. The source ID in each metacard identifies the Source from which the metacard originated.

1.4.7. Introduction to Events and Subscriptions

DDF can be configured to receive notifications whenever metadata is created, updated, or deleted in any federated sources. Creations, updates, and deletions are collectively called **Events**, and the process of registering to receive them is called **Subscription**.

The behavior of these subscriptions is consistent, but the method of configuring them is specific to the [Endpoint](#) used.

1.4.8. Introduction to Registries

The Registry Application serves as an index of registry nodes and their information, including service bindings, configurations and supplemental details.

Each registry has the capability to serve as an index of information about a network of registries which, in turn, can be used to connect across a network of DDFs and other data sources. Registries communicate with each other through the CSW endpoint and each registry node is converted into a registry metacard to be stored in the catalog. When a registry is subscribed to or published from, it sends the details of one or more nodes to another registry.

Identity Node

The Registry is initially comprised of a single registry node, referred to as the **identity**, which represents the registry's primary configuration.

Subscription

Subscribing to a registry is the act of retrieving its information, specifically its identity information and any other registries it knows about. By default, subscriptions are configured to check for updates every 30 seconds.

Publication

Publishing is the act of sending a registry's information to another registry. Once publication has occurred, any updates to the local registry will be pushed out to the registries that have been published to.

1.4.9. Introduction to Endpoints

Endpoints expose the Catalog Framework to clients using protocols and formats that the clients understand.

Endpoint interface formats encompass a variety of protocols, including (but not limited to):

- SOAP Web services
- RESTful services
- JMS
- JSON
- OpenSearch

The endpoint may transform a client request into a compatible Catalog format and then transform the response into a compatible client format. Endpoints may use [Transformers](#) to perform these transformations. This allows an endpoint to interact with Source(s) that have different interfaces. For example, an OpenSearch Endpoint can send a query to the Catalog Framework, which could then query a federated source that has no OpenSearch interface.

Endpoints are meant to be the only client-accessible components in the Catalog.

1.4.10. Introduction to High Availability

DDF can be made highly available. In this context, High Availability is defined as the ability for DDF to be continuously operational with very little down time.


In a Highly Available Cluster, DDF has failover capabilities when a DDF node fails.

NOTE


The word "node", from a High Availability perspective, is one of the two DDF systems running within the Highly Available Cluster. Though there are multiple systems running with the Highly Available Cluster, it is still considered a single DDF from a user's perspective or from other DDFs' perspectives.

This setup consists of a Solr Cloud instance, 2 DDF nodes connected to that Solr Cloud, and a failover proxy that sits in front of those 2 nodes. One of the DDF nodes will be arbitrarily chosen to be the active node, and the other will be the "hot standby" node. It is called a "hot standby" node because it is ready to receive traffic even though it's not currently receiving any. The failover proxy

will route all traffic to the active node. If the active node fails for any reason, the standby node will become active and the failover proxy will route all traffic to the new active node. See the below diagrams for more detail.



Highly Available Cluster



Highly Available Cluster (after failover)

There are special procedures for initial setup and configuration of a highly available DDF. See [High Availability Initial Setup](#) and [High Availability Configuration](#) for those procedures.

1.4.10.1. High Availability Supported Capabilities

Only these capabilities are supported in a Highly Available Cluster. For a detailed list of features, look at the `ha.json` file located in `<DDF_HOME>/etc/profiles/`.

- User Interfaces:
 - Simple
 - Intrigue
- Catalog:
 - Validation
 - Plug-ins: Expiration Date, JPEG2000, Metocard Validation, Schematron, Versioning
 - Transformers
 - Content File System Storage Provider
- Platform:
 - Actions
 - Configuration

- Notifications
- Persistence
- Security: Audit, Encryption
- Solr
- Security
- Thirdy Party:
 - CXF
 - Camel
- Endpoints:
 - REST Endpoint
 - CSW Endpoint
 - OpenSearch Endpoint

2. Quick Start Tutorial

This quick tutorial will enable install, configuring and using a basic instance of DDF.

NOTE

This tutorial is intended for setting up a test, demonstration, or trial installation of DDF. For more complete installation and configuration steps, see [Installing](#).

These steps will demonstrate:

- [Prerequisites](#).
- [Quick Install of DDF](#).
- [Ingesting Data](#).

2.1. Quick Install Prerequisites

These are the basic requirements to set up the environment to run a DDF.

2.1.1. Hardware Requirements (Quick Install)

- At least 4096MB of memory for DDF.
 - This amount can be increased to support memory-intensive applications. See [Memory Considerations](#).

2.1.2. Java Requirements (Quick Install)

Set up Java to run DDF.

- Install/Upgrade to Java 8 x64 [J2SE 8 SDK](#)
 - The recommended version is [8u60](#) or later.

- Java Version and Build numbers must contain only number values.
- Supported platforms are Microsoft Windows, Linux, Solaris, and macOS.
- [JDK8x64](#) must be installed.
- The [JAVA_HOME](#) environment variable must be set to the location where the JDK is installed.

1. *Setting JAVA_HOME variable (Replace <JAVA_VERSION> with the version and build number installed.)*

1. Determine Java Installation Directory (This varies between operating system versions).

Find Java Path in Windows

```
for %i in (java.exe) do @echo. %~$PATH:i
```

*Find Java Path in *nix*

```
which java
```

1. Copy path to Java installation. (example: [/usr/java/<JAVA_VERSION>](#))

2. Set [JAVA_HOME](#) by replacing <PATH_TO_JAVA> with the copied path in this command:

Setting JAVA_HOME on Windows

```
set JAVA_HOME=<PATH_TO_JAVA><JAVA_VERSION> /m
```

Adding JAVA_HOME to PATH Environment Variable on Windows

```
setx PATH "%PATH%;%JAVA_HOME%\bin" /m
```

*Setting JAVA_HOME on *nix*

```
JAVA_HOME=<PATH_TO_JAVA><JAVA_VERSION>
export JAVA_HOME
```

**nix*

WARNING

Unlink [JAVA_HOME](#) if it is already linked to a previous version of the JRE: [unlink JAVA_HOME](#)

Verify that the `JAVA_HOME` was set correctly.

Windows

```
echo %JAVA_HOME%
```

TIP

*nix

```
echo $JAVA_HOME
```

File Descriptor Limit on Linux

- For Linux systems, increase the file descriptor limit by editing `/etc/sysctl.conf` to include:

```
fs.file-max = 6815744
```

NOTE

- (This file may need permissions changed to allow write access).
- For the change to take effect, a restart is required.
 1. *nix Restart Command

```
init 6
```

Check System Time

WARNING

Prior to installing DDF, ensure the system time is accurate to prevent federation issues.

2.2. Quick Install of DDF

1. Download the DDF [zip file](#).
2. Install DDF by unzipping the zip file.

Windows Zip Utility Warning

The Windows Zip implementation, which is invoked when a user double-clicks on a zip file in the Windows Explorer, creates a corrupted installation. This is a consequence of its inability to process long file paths. Instead, use the java jar command line utility to unzip the distribution (see example below) or use a third party utility such as 7-Zip.

WARNING

Use Java to Unzip in Windows(Replace <PATH_TO_JAVA> with correct path and <JAVA_VERSION> with current version.)

```
"<PATH_TO_JAVA>\jdk<JAVA_VERSION>\bin\jar.exe" xf ddf-2.11.3.zip
```

3. This will create an installation directory, which is typically created with the name and version of the application. This installation directory will be referred to as `<DDF_HOME>`. (Substitute the actual directory name.)
4. Start DDF by running the `<DDF_HOME>/bin/ddf` script (or `ddf.bat` on Windows).
5. Startup may take a few minutes.
 - a. Optionally, a `waitForReady` command (aliased to `wfr`) can be used to notify when setup is complete. A "System Ready" message will display. *nix users will see progress indicated before the message.
6. The Command Console will display.

Command Console Prompt

```
ddf@local>
```

2.3. Certificates (Quick Start)

DDF comes with a default keystore that contains certificates. This allows the distribution to be unzipped and run immediately. If these certificates are sufficient for testing purposes, proceed to [Configuring \(Quick Start\)](#).

To test federation using 2-way TLS, the default keystore certificates will need to be replaced, using either the included [Demo Certificate Authority](#) or by [Creating Self-signed Certificates](#).

If the installer was used to install the DDF and a hostname other than "localhost" was given, the user will be prompted to upload new trust/key stores.

If the hostname is `localhost` or, if the hostname was changed *after* installation, the default certificates will not allow access to the DDF instance from another machine over HTTPS (now the default for many services). The Demo Certificate Authority will need to be replaced with certificates that use the fully-qualified hostname of the server running the DDF instance.

2.3.1. Demo Certificate Authority (CA)

DDF comes with a populated truststore containing entries for many public certificate authorities, such as Go Daddy and Verisign. It also includes an entry for the DDF Demo Root CA. This entry is a self-signed certificate used for testing. It enables DDF to run immediately after unzipping the distribution. The keys and certificates for the DDF Demo Root CA are included as part of the DDF distribution. This entry must be removed from the truststore before DDF can operate securely.

2.3.1.1. Creating New Server Keystore Entry with the CertNew Scripts

To create a private key and certificate signed by the Demo Certificate Authority, use the provided scripts. To use the scripts, run them out of the `<DDF_HOME>/etc/certs` directory.

*NIX Demo CA Script

For *NIX, use the `CertNew.sh` script.

```
sh CertNew.sh [-cn <cn>|-dn <dn>] [-san <tag:name,tag:name, ...>]
```

where:

- **<cn>** represents a fully qualified common name (e.g. "<FQDN>", where <FQDN> could be something like cluster.yoyo.com)
 - **<dn>** represents a distinguished name as a comma-delimited string (e.g. "c=US, st=California, o=Yoyodyne, l=San Narciso, cn=<FQDN>")
 - **<tag:tag>** represents optional subject alternative names to be added to the generated certificate (e.g. "DNS:<FQDN>,DNS:node1.<FQDN>,DNS:node2.<FQDN>"). The format for subject alternative names is similar to the OpenSSL X509 configuration format. Supported tags are:
 - **email** - email subject
 - **URI** - uniformed resource identifier
 - **RID** - registered id
 - **DNS** - hostname
 - **IP** - ip address (V4 or V6)
 - **dirName** - directory name

If no arguments specified on the command line, `hostname -f` is used as the common-name for the certificate.

Windows Demo CA Script

For Windows, use the `CertNew.cmd` script.

```
CertNew (-cn <cn>|-dn <dn>) [-san "<tag: name, tag: name, ...>"]
```

where:

- `<cn>` represents a fully qualified common name (e.g. "<FQDN>", where <FQDN> could be something like cluster.yoyo.com)
- `<dn>` represents a distinguished name as a comma-delimited string (e.g. "c=US, st=California, o=Yoyodyne, l=San Narciso, cn=<FQDN>")
- `<tag: name, tag: name, ...>` represents optional subject alternative names to be added to the generated certificate (e.g. "DNS:<FQDN>, DNS:node1.<FQDN>, DNS:node2.<FQDN>"). The format for subject alternative names is similar to the OpenSSL X509 configuration format. Supported tags are:
 - `email` - email subject
 - `URI` - uniformed resource identifier
 - `RID` - registered id
 - `DNS` - hostname
 - `IP` - ip address (V4 or V6)
 - `dirName` - directory name

The `CertNew` scripts:

- Create a new entry in the server keystore.
- Use the hostname as the fully qualified domain name (FQDN) when creating the certificate.
- Adds the specified subject alternative names if any.
- Use the Demo Certificate Authority to sign the certificate so that it will be trusted by the default configuration.

To install a certificate signed by a different Certificate Authority, see [Managing Keystores](#).

Once the certificate has been updated the system user and `org.codice.ddf.system.hostname` property will also need to be updated to match the FQDN. Replace localhost with the FQDN in `<DDF_HOME>/etc/users.properties`, `<DDF_HOME>/etc/users.attributes` and `<DDF_HOME>/etc/system.properties`.

Finally, restart the DDF instance. Browse the Admin Console to test changes.

WARNING

If the server's fully qualified domain name is not recognized, the name may need to be added to the network's DNS server.

The DDF instance can be tested even if there is no entry for the FQDN in the DNS. First, test if the FQDN is already recognized. Execute this command:

```
ping <FQDN>
```

TIP

If the command responds with an error message such as unknown host, then modify the system's `hosts` file to point the server's FQDN to the loopback address. For example:

```
127.0.0.1 <FQDN>
```

2.3.2. Creating Self-Signed Certificates

If using the Demo CA is not desired, DDF supports creating self-signed certificates with a self-signed certificate authority. This is considered an advanced configuration.

Creating self-signed certificates involves creating and configuring the files that contain the certificates. In DDF, these files are generally Java Keystores (`jks`) and Certificate Revocation Lists (`crl`). This includes commands and tools that can be used to perform these operations.

For this example, the following tools are used:

- `openssl`
 - Windows users can use: `openssl` for windows.
- The standard Java `keytool` certificate management utility.
- `Portecle` can be used for `keytool` operations if a GUI is preferred over a command line interface.

2.3.2.1. Creating a custom CA Key and Certificate

The following steps demonstrate creating a root CA to sign certificates.

1. Create a key pair.

```
$> openssl genrsa -aes128 -out root-ca.key 1024
```

2. Use the key to sign the CA certificate.

```
$> openssl req -new -x509 -days 3650 -key root-ca.key -out root-ca.crt
```

2.3.2.2. Sign Certificates Using the custom CA

The following steps demonstrate signing a certificate for the `tokenissuer` user by a CA.

1. Generate a private key and a Certificate Signing Request (CSR).

```
$> openssl req -newkey rsa:1024 -keyout tokenissuer.key -out tokenissuer.req
```

2. Sign the certificate by the CA.

```
$> openssl ca -out tokenissuer.crt -infiles tokenissuer.req
```

These certificates will be used during system configuration to replace the default certificates.

2.4. Configuring (Quick Start)

Set the configurations needed to run DDF.

1. In a browser, navigate to the Admin Console at <https://{FQDN}:{PORT}/admin>.
 - a. The Admin Console may take a few minutes to start up.
2. Enter the default username of **admin** and the password of **admin**.
3. Follow the installer prompts for a standard installation.
 - a. Click **start** to begin the setup process.
 - b. Configure [guest claims attributes](#) or use defaults.
 - i. See [Configuring Guest Access](#) for more information about the Guest user.
 - ii. **All users will be automatically granted these permissions.**
 - iii. **Guest users will not be able to ingest data with more restrictive markings than the guest claims.**
 - iv. **Any data ingested that has more restrictive markings than these guest claims will not be visible to Guest users.**
 - c. Select **Standard Installation**.
 - i. This step may take several minutes to complete.
 - d. On the System Configuration page, configure any port or protocol changes desired and add any keystores/truststores needed.
 - i. See [Certificates \(Quick Start\)](#) for more details.
 - e. Click **Next**
 - f. Click **Finish**

2.5. Ingesting (Quick Start)

Now that DDF has been configured, ingest some sample data to demonstrate search capabilities.

This is one way to ingest into the catalog, for a complete list of the different methods, see [Ingesting Data](#).

2.5.1. Ingesting Sample Data

1. Download a sample valid [GeoJson file here](#).
2. Navigate in the browser to Intrigue at <https://{FQDN}:{PORT}/search>.
3.  Select the Menu icon (☰) in the upper left corner
4. Select **Upload**.
5. Drag and drop the sample file or click to navigate to it.
6. Select **Start** to begin upload.

NOTE

XML metadata for text searching is not automatically generated from GeoJson fields.

Querying from the Search UI (<https://{FQDN}:{PORT}/search>) will return the record for the file ingested:

1.



Select the Menu icon () and return to **Workspaces**.

2. Search for the ingested data.

NOTE

The sample data was selected as an example of well-formed metadata. Other data can and should be used to test other usage scenarios.

Managing

Administrators will be installing, maintaining, and supporting existing applications. Use this section to prepare, install, configure, run, and monitor a DDF.

3. Securing

Security is an important consideration for DDF, so it is imperative to update configurations away from the defaults to unique, secure settings.

Securing DDF Components

DDF is enabled with an Insecure Defaults Service which will warn users/admins if the system is configured with insecure defaults.

IMPORTANT

A banner is displayed on the admin console notifying "The system is insecure because default configuration values are in use."

A detailed view is available of the properties to update.

Security concerns will be highlighted in the configuration sections to follow.

3.1. Security Hardening

Security Hardening

To harden DDF, extra security precautions are required.

Where available, necessary mitigations to harden an installation of DDF are called out in the following configuration steps.

Refer to the [Hardening Checklist](#) for a compilation of these mitigations.

These security precautions are best performed as configuration is taking place, so hardening steps are integrated into configuration steps. This is to avoid setting an insecure configuration and having to revisit during hardening. Most configurations have a security component to them, and important considerations for hardening are labeled as such during configuration as well as provided in a checklist format.

Some of the items on the checklist are performed during [installation](#) and others during [configuration](#). Steps required for hardening are marked as **Required for Hardening** and are collected here for convenience. Refer to the checklist during system setup.

3.1.1. Hardening Checklist

The following list enumerates the required mitigations needed for hardening:

- [Configure Auditing](#)
- [Set Directory Permissions](#)
- [Configure Keystore and Certificates](#)
- [Disallow Login Without Certificates](#)
- [Configure Certificate Revocation List](#)
 - [Deny Guest User Access](#) (if denying Guest users)
 - [Allow Guest User Access](#) (if allowing Guest users)
- [Configure Guest Claim Attributes](#) (if allowing Guest users)
- [Configure Guest User Authentication](#)
- [Limit Access to the STS](#)
- [Restricting Access to Admin Console](#)
- [Restrict Feature, App, Service, and Configuration Access](#)
- [Remove Default Users](#)
- [Harden Solr Index](#)
- [Environment Hardening](#)
- [Isolate Solr Cloud and Zookeeper](#). (If using)

3.2. Auditing

- **Required Step for Security Hardening**

Audit logging captures security-specific system events for monitoring and review. DDF provides a [Audit Plugin](#) that logs all catalog transactions to the [security.log](#). Information captured includes user identity, query information, and resources retrieved.

Follow all operational requirements for the retention of the log files. This may include using cryptographic mechanisms, such as encrypted file volumes or databases, to protect the integrity of audit information.

NOTE The Audit Log default location is <DDF_HOME>/data/log/security.log

Audit Logging Best Practices

For the most reliable audit trail, it is recommended to configure the operational environment of the DDF to generate alerts to notify administrators of:

- NOTE**
- auditing software/hardware errors
 - failures in audit capturing mechanisms
 - audit storage capacity (or desired percentage threshold) being reached or exceeded.

WARNING The security audit logging function does not have any configuration for audit reduction or report generation. The logs themselves could be used to generate such reports outside the scope of DDF.

3.2.1. Enabling Fallback Audit Logging

• Required Step for Security Hardening

In the event the system is unable to write to the `security.log` file, DDF must be configured to fall back to report the error in the application log:

- edit <DDF_HOME>/etc/org.ops4j.pax.logging.cfg
 - uncomment the line (remove the `#` from the beginning of the line) for `log4j2 (org.ops4j.pax.logging.log4j2.config.file = ${karaf.etc}/log4j2.config.xml)`
 - delete all subsequent lines
- edit <DDF_HOME>/etc/log4j2.config.xml
 - find the entry for the `securityBackup` appender. (see example)
 - change value of `filename` and prefix of `filePattern` to the name/path of the desired failover security logs (<NEW_FILE_NAME>)

`securityBackup` Appender Before

```
<RollingFile name="securityBackup" append="true" ignoreExceptions="false"
              fileName="${sys:karaf.data}/log/securityBackup.log"
              filePattern="${sys:karaf.data}/log/securityBackup.log-%d{yyyy-MM-
dd-HH}-%i.log.gz">
```

`securityBackup` Appender After

```
<RollingFile name="securityBackup" append="true" ignoreExceptions="false"
              fileName="${sys:karaf.data}/log/<NEW_FILE_NAME>"
              filePattern="${sys:karaf.data}/log/<NEW_FILE_NAME>-%d{yyyy-MM-dd-
HH}-%i.log.gz">
```

WARNING

If the system is unable to write to the `security.log` file on system startup, fallback logging will be unavailable. Verify that the `security.log` file is properly configured and contains logs before configuring a fall back.

4. Installing

Set up a complete, secure instance of DDF. For simplified steps used for a testing, development, or demonstration installation, see the [DDF Quick Start](#).

IMPORTANT

Although DDF can be installed by any user, it is recommended for security reasons to have a non-`root` user execute the DDF installation.

NOTE

Hardening guidance assumes a Standard installation.

Adding other components does not have any security/hardening implications.

4.1. Installation Prerequisites

These are the system/environment requirements to configure *prior* to an installation.

Table 1. Using the Standard installation of the DDF application:

Minimum and Recommended Requirements for DDF Systems*		
Criteria	Minimum	Recommended
Operating System	Windows 10, Linux CentOS 5	Windows Server 2012, Linux CentOS 7.3
Additional Software	JDK 8 x64, IE 10	JDK 8 x64, Edge/FireFox/Chrome
CPU	Dual Core 1.6 GHz	Quad Core 2.6 GHz
RAM	4 GB**	32 GB
Disk Space	2 GB	4 GB
Video Card	WebGL capable GPU	WebGL capable GPU

NOTE

If installing a test, demo, or development installation, see the [Quick Start Installation Prerequisites](#).

*These requirements assume the use of a remote catalog. Extra resources will need to be allocated if the system will also be hosting a local catalog.

**The amount of RAM can be increased to support memory-intensive applications. See [Memory Considerations](#)

Java Requirements

- Supported platforms are *NIX - Unix/Linux/OSX, Solaris, and Windows.

- [JDK8](#) must be installed.
- The [JAVA_HOME](#) environment variable must be set to the location where the JDK is installed.
 1. Install/Upgrade to Java 8 x64 [J2SE 8 SDK](#)
 - a. The recommended version is [8u60](#) or later.
 - b. Java version must contain only number values.
 2. Install/Upgrade to [JDK8](#).
 3. Set the [JAVA_HOME](#) environment variable to the location where the JDK is installed.

Setting JAVA_HOME variable

*NIX

WARNING Unlink **JAVA_HOME** if it is already linked to a previous version of the JRE:

```
unlink JAVA_HOME
```

Replace **<JAVA_VERSION>** with the version and build number installed.

1. Open a terminal window(*NIX) or command prompt (Windows) with administrator privileges.
2. Determine Java Installation Directory (This varies between operating system versions).

*Find Java Path in *NIX*

```
which java
```

Find Java Path in Windows

The path to the JDK can vary between versions of Windows, so manually verify the path under:

```
C:\Program Files\Java\jdk<M.m.p_build>
```

3. Copy path of Java installation to clipboard. (example: **/usr/java/<JAVA_VERSION>**)
4. Set **JAVA_HOME** by replacing **<PATH_TO_JAVA>** with the copied path in this command:

*Setting JAVA_HOME on *NIX*

```
JAVA_HOME=<PATH_TO_JAVA><JAVA_VERSION>
export JAVA_HOME
```

Setting JAVA_HOME on Windows

```
set JAVA_HOME=<PATH_TO_JAVA><JAVA_VERSION>
setx JAVA_HOME "<PATH_TO_JAVA><JAVA_VERSION>" /m
```

Adding JAVA_HOME to PATH Environment Variable on Windows

```
setx PATH "%PATH%;%JAVA_HOME%\bin" /m
```

5. Restart Terminal (shell) or Command Prompt.
 - Verify that the **JAVA_HOME** was set correctly.

*NIX

```
echo $JAVA_HOME
```

Windows

```
echo %JAVA_HOME%
```

File Descriptor Limit on Linux

- For Linux systems, increase the file descriptor limit by editing [/etc/sysctl.conf](#) to include:

```
fs.file-max = 6815744
```

NOTE

- For the change to take effect, a restart is required.

**Nix Restart Command*

```
init 6
```

4.2. Installing With the DDF Distribution Zip

Check System Time

WARNING

Prior to installing DDF, ensure the system time is accurate to prevent federation issues.

To install the DDF distribution zip, perform the following:

1. Download the DDF [zip file](#).
2. After the [prerequisites](#) have been met, change the current directory to the desired install directory, creating a new directory if desired. This will be referred to as [`<DDF_HOME>`](#).

Windows Pathname Warning

WARNING

Do not use spaces in directory or file names of the [`<DDF_HOME>`](#) path. For example, do not install in the default [Program Files](#) directory.

*Example: Create a Directory (Windows and *NIX)*

```
mkdir new_installation
```

- a. Use a Non-[root](#) User on *NIX. (Windows users skip this step)

It is recommended that the **root** user create a new install directory that can be owned by a non-**root** user (e.g., DDF_USER). This can be a new or existing user. This DDF_USER can now be used for the remaining installation instructions.

- b. Create a new group or use an existing group (e.g., DDF_GROUP) (Windows users skip this step)

*Example: Add New Group on *NIX*

```
groupadd DDF_GROUP
```

*Example: Switch User on *NIX*

```
chown DDF_USER:DDF_GROUP new_installation
su - DDF_USER
```

3. Change the current directory to the location of the zip file (ddf-2.11.3.zip).

**NIX (Example assumes DDF has been downloaded to a CD/DVD)*

```
cd /home/user/cdrom
```

Windows (Example assumes DDF has been downloaded to the D drive)

```
cd D:\
```

4. Copy ddf-2.11.3.zip to <DDF_HOME>.

**NIX*

```
cp ddf-2.11.3.zip <DDF_HOME>
```

Windows

```
copy ddf-2.11.3.zip <DDF_HOME>
```

5. Change the current directory to the desired install location.

**NIX or Windows*

```
cd <DDF_HOME>
```

6. The DDF zip is now located within the <DDF_HOME>. Unzip ddf-2.11.3.zip.

*NIX

```
unzip ddf-2.11.3.zip
```

Windows Zip Utility Warning

The Windows Zip implementation, which is invoked when a user double-clicks on a zip file in the Windows Explorer, creates a corrupted installation. This is a consequence of its inability to process long file paths. Instead, use the java jar command line utility to unzip the distribution (see example below) or use a third party utility such as 7-Zip.

WARNING

Use Java to Unzip in Windows(Replace <PATH_TO_JAVA> with correct path and <JAVA_VERSION> with current version.)

```
"<PATH_TO_JAVA>\jdk<JAVA_VERSION>\bin\jar.exe" xf ddf-2.11.3.zip
```

The unzipping process may take time to complete. The command prompt will stop responding to input during this time.

4.2.1. Controlling File System Access

Restrict access to sensitive files by ensuring that the only users with access privileges are administrators.

Within the <DDF_HOME>, a directory is created named ddf-2.11.3. This directory will be referred to in the documentation as <DDF_HOME>.

1. Do not assume the deployment is from a trusted source; verify its origination.
2. Check the available storage space on the system to ensure the deployment will not exceed the available space.
3. Set maximum storage space on the <DDF_HOME>/deploy and <DDF_HOME>/system directories to restrict the amount of space used by deployments.

4.2.1.1. Setting Directory Permissions

- **Required Step for Security Hardening**

DDF relies on the Directory Permissions of the host platform to protect the integrity of the DDF during operation. System administrators MUST perform the following steps prior to deploying bundles added to the DDF.

IMPORTANT

The system administrator must restrict certain directories to ensure that the application (user) cannot access restricted directories on the system. For example the **DDF_USER** should have read-only access to <DDF_HOME>, except for the sub-directories **etc**, **data**, **solr** and **instances**.

Setting Directory Permissions on Windows

Set directory permissions on the `<DDF_HOME>`; all sub-directories except `etc`, `data`, and `instances`; and any directory intended to interact with the DDF to protect from unauthorized access.

1. Right-click on the `<DDF_HOME>` directory.
2. Select **Properties** -> **Security** -> **Advanced**.
3. Under **Owner**, select **Change**.
4. Enter `Creator Owner` into the **Enter the Object Name...** field.
5. Select **Check Names**.
6. Select **Apply**.
 - a. If prompted **Do you wish to continue**, select **Yes**.
7. Remove all Permission Entries for any groups or users with access to `<DDF_HOME>` other than **System**, **Administrators**, and **Creator Owner**.
 - a. Note: If prompted with a message such as: **You can't remove X because this object is inheriting permissions from its parent**. when removing entries from the Permission entries table:
 - i. Select **Disable Inheritance**.
 - ii. Select **Convert Inherited Permissions into explicit permissions on this object**.
 - iii. Try removing the entry again.
8. Select the option for **Replace all child object permission entries with inheritable permission entries from this object**.
9. Close the **Advanced Security Settings** window.

Setting Directory Permissions on *NIX

Set directory permissions to protect the DDF from unauthorized access.

- Change ownership of <DDF_HOME>
 - `chown -R ddf-user <DDF_HOME>`
- Create instances sub-directory if does not exist
 - `mkdir -p <DDF_HOME>/instances <DDF_HOME>/solr`
- Change group ownership on sub-directories
 - `chgrp -R DDF_GROUP <DDF_HOME>/etc <DDF_HOME>/data <DDF_HOME>/instances <DDF_HOME>/solr`
- Change group permissions
 - `chmod -R g-w <DDF_HOME>/etc <DDF_HOME>/data <DDF_HOME>/instances <DDF_HOME>/solr`
- Remove permissions for other users
 - `chmod -R o-rwx <DDF_HOME>/etc <DDF_HOME>/data <DDF_HOME>/instances`

4.3. Initial Startup

Run the DDF using the appropriate script.

*NIX

```
<DDF_HOME>/bin/ddf
```

Windows

```
<DDF_HOME>/bin/ddf.bat
```

The distribution takes a few moments to load depending on the hardware configuration.

TIP To run DDF as a service, see [Starting as a Service](#).

4.3.1. Verifying Startup

At this point, DDF should be configured and running with a Solr Catalog Provider. New features (endpoints, services, and sites) can be added as needed.

Verification is achieved by checking that all of the DDF bundles are in an **Active** state (excluding fragment bundles which remain in a **Resolved** state).

NOTE It may take a few moments for all bundles to start so it may be necessary to wait a few minutes before verifying installation.

Execute the following command to display the status of all the DDF bundles:

View Status

```
ddf@local>list | grep -i ddf
```

WARNING

Entries in the **Resolved** state are expected, they are OSGi bundle fragments. Bundle fragments are distinguished from other bundles in the command line console list by a field named **Hosts**, followed by a bundle number. Bundle fragments remain in the **Resolved** state and can never move to the **Active** state.

Example: Bundle Fragment in the Command Line Console

```
96 | Resolved | 80 | 2.10.0.SNAPSHOT | DDF :: Platform :: PaxWeb :: Jetty Config,  
Hosts: 90
```

After successfully completing these steps, the DDF is ready to be configured.

4.3.2. DDF Directory Contents after Installation and Initial Startup

During DDF installation, the major directories and files shown in the table below are created, modified, or replaced in the destination directory.

Table 2. DDF Directory Contents

Directory Name	Description
<code>bin</code>	Scripts to start, stop, and connect to DDF.
<code>data</code>	The working directory of the system – installed bundles and their data
<code>data/log/ddf.log</code>	Log file for DDF, logging all errors, warnings, and (optionally) debug statements. This log rolls up to 10 times, frequency based on a configurable setting (default=1 MB)
<code>data/log/ingest_error.log</code>	Log file for any ingest errors that occur within DDF.
<code>data/log/security.log</code>	Log file that records user interactions with the system for auditing purposes.
<code>data/log/solr.log</code>	Log file for any info coming from Solr.
<code>deploy</code>	Hot-deploy directory – KARs and bundles added to this directory will be hot-deployed (Empty upon DDF installation)
<code>documentation</code>	HTML and PDF copies of DDF documentation.
<code>etc</code>	Directory monitored for addition/modification/deletion of <code>.config</code> configuration files or third party <code>.cfg</code> configuration files.
<code>etc/templates</code>	Template <code>.config</code> files for use in configuring DDF sources, settings, etc., by copying to the etc directory.
<code>lib</code>	The system's bootstrap libraries. Includes the <code>ddf-branding.jar</code> file which is used to brand the system console with the DDF logo.
<code>licenses</code>	Licensing information related to the system.

Directory Name	Description
system	Local bundle repository. Contains all of the JARs required by DDF, including third-party JARs.

4.3.3. Complete Installation

Upon startup, complete installation from either the Admin Console or the command line.

4.3.3.1. Completing Installation from the Admin Console

Upon startup, the installation can be completed by navigating to the Admin Console at <https://FQDN:{PORT}/admin>.

Internet Explorer 10 TLS Warning

Internet Explorer 10 users may need to enable TLS 1.2 to access the Admin Console in the browser.

WARNING

Enabling TLS1.2 in IE10

1. Go to **Tools** -> Internet Options -> Advanced -> Settings -> Security.
2. Enable TLS1.2.

- Default user/password: **admin/admin**.

On the initial startup of the Admin Console, a series of prompts walks through essential configurations. These configurations can be changed later, if needed.

- Click **Start** to begin.

Configure Guest Claim Attributes Page

Setting the attributes on the **Configure Guest Claim Attributes** page determines the minimum claims attributes (and, therefore, permissions) available to a guest, or not signed-in, user.

To change this later, see [Configuring Guest Claim Attributes](#).

Setup Types

DDF is pre-configured with several installation profiles.

- Standard Installation: **Recommended**. Includes these applications by default:
 - [Admin](#)
 - [Catalog](#)
 - [Platform](#)
 - [Security](#)
 - [Solr Catalog](#)
 - [Spatial](#)
 - [Intrigue](#)
- Minimum Installation: Includes these applications for a minimum install:
 - [Admin](#)
 - [Platform](#)
 - [Security](#)
- Development: Includes all demo, beta, and experimental applications.
- Custom Installation: **Advanced**. Click **Customize** on either profile to add or remove applications to be installed.
 - If apps are preselected when the **Select Applications** page is reached, they will be uninstalled if unselected.

The Platform, Admin, and Security applications are required and CANNOT be selected or unselected.

WARNING

The Security Application appears to be unselected upon first view of the tree structure, but it is in fact automatically installed with a later part of the installation process.

System Configuration Settings

- System Settings: Set `hostname` and `ports` for this installation.
- Contact Info: Contact information for the point-of-contact or administrator for this installation.
- Certificates: Add PKI certificates for the Keystore and Truststore for this installation.
 - For a quick (test) installation, if the hostname/ports are not changed from the defaults, DDF includes self-signed certificates to use. Do not use in a working installation.
 - For more advanced testing, on initial startup of the **Admin Console** append the string `?dev=true` to the url (`https://{FQDN}:{PORT}/admin?dev=true`) to auto-generate self-signed certificates from a demo Certificate Authority(CA) which will enable change hostname and port settings. Do not use in a working installation.
 - NOTE: `?dev=true` generates certificates on initial installation only.
 - For more information about importing certificate from a Certificate Authority, see [Managing Keystores and Certificates](#).

Finished Page

Upon successful startup, the **Finish** page will redirect to the Admin Console to begin further configuration, ingest, or federation.

NOTE The redirect will only work if the certificates are configured in the browser. Otherwise the redirect link must be used.

4.3.3.2. Completing Installation from the Command Console

In order to install DDF from the Command Console, use the command `profile:install <profile-name>`. The `<profile-name>` should be the desired **Setup Type** in lowercase letters. To see the available profiles, use the command `profile:list`.

NOTE This only installs the desired Setup Type. There are other components that can be set up in the Admin Console Installer that cannot be setup on the Command Console. After installing the Setup Type, these other components can be set up as described below.

4.3.3.2.1. Configure Guest Claim Attributes

The Guest Claim Attributes can be configured via the Admin Console after running the `profile:install` command. See [Configuring Guest Claim Attributes](#).

4.3.3.2.2. System Configuration Settings

System Settings and Contact Info, as described in [System Configuration Settings](#), can be changed in `<DDF_HOME>/etc/system.properties`. The certificates must be set up manually as described in

Managing Keystores and Certificates.

NOTE The system will need to be restarted after changing any of these settings.

4.3.4. Firewall Port Configuration

Below is a table listing all of the default ports that DDF uses and a description of what they are used for. Firewalls will need to be configured to open these ports in order for external systems to communicate with DDF.

Table 3. Port List

Port	Usage description
8993	https access to DDF admin and search web pages.
8101	For administering DDF instances gives ssh access to the administration console.
61616	DDF broker port for JMS messaging over the OpenWire protocol.
5672	DDF broker port for JMS messaging over multiple protocols: Artemis CORE, AMQP and OpenWire by default .
5671	DDF broker port for JMS messaging over: AMQP by default.
1099	RMI Registry Port
44444	RMI Server Port

NOTE These are the default ports used by DDF. DDF can be configured to use different ports.

4.3.5. Internet Explorer 11 Enhanced Security Configuration

Below are steps listing all of the changes that DDF requires to run on Internet Explorer 11 and several additional considerations to keep in mind.

1. In the **IE11 Settings > Compatibility View Settings** dialog, un-check **Display intranet sites in Compatibility View**.
2. In the **Settings > Internet Options > Security tab**, **Local intranet** zone:
 - a. Click the **Sites > Advanced** button, add the current host name to the list, e.g., <https://windows-host-name.domain.edu>, and close the dialog.
 - b. Make sure the security level for the **Local intranet** zone is set to **Medium-low** in **Custom level...**
 - i. **Enable Protected Mode** is checked by default, but it may need to be disabled if the above changes do not fully resolve access issues.

3. Restart the browser.

NOTE

During installation, make sure to use the host name and not localhost when setting up the DDF's hostname, port, etc.

4.4. High Availability Initial Setup

This section describes how to complete the initial setup of DDF in a [Highly Available Cluster](#).

Prerequisites

- A failover proxy that can route HTTP traffic according to the pattern described in the [Introduction to High Availability](#). It is recommended that a hardware failover proxy be used in a production environment.
- Solr Cloud: See the [Solr Cloud section](#) for installation and configuration guidance to connect DDF nodes to Solr Cloud.

Once the prerequisites have been met, the below steps can be followed.

NOTE

Unless listed in the [High Availability Initial Setup Exceptions](#) section, the normal steps can be followed for installing, configuring, and hardening.

1. Install the first DDF node. See the [Installation Section](#).
2. Configure the first DDF node. See the [Configuring Section](#).
3. Optional: If hardening the first DDF node (excluding setting directory permissions). See the [Hardening Section](#).
4. Export the first DDF node's configurations, install the second DDF node, and import the exported configurations on that node. See [Reusing Configurations](#).
5. If hardening, set directory permissions on both DDF nodes. See [Setting Directory Permissions](#).

4.4.1. High Availability Initial Setup Exceptions

These steps are handled differently for the initial setup of a Highly Available Cluster.

4.4.1.1. HTTP Proxy Integration

In order to integrate with an HTTP proxy, the DDF node's system properties (in `<DDF_HOME>/etc/system.properties`) must be changed to publish the correct port to external systems and users. This must be done before installing the first DDF node.

There are two internal port properties that must be changed to whatever ports the DDF will use on its system. Then there are two external port properties that must be changed to whatever ports the HTTP proxy is forwarding traffic through.

WARNING

Make sure that the HTTP proxy is already running and forwarding traffic on the chosen ports before starting the DDF. There may be unexpected behavior otherwise.

In the below example, the HTTP proxy is forwarding https traffic via 8993 and http traffic via 8181. The DDF node will run on 1111 for https and 2222 for http (on the host that it's hosted on).

```
org.codice.ddf.system.httpsPort=8993
org.codice.ddf.system.httpPort=8181
org.codice.ddf.system.port=${org.codice.ddf.system.httpsPort}
org.codice.ddf.system.internalHttpsPort=1111
org.codice.ddf.system.internalHttpPort=2222
```

4.4.1.2. Identical Directory Structures

The two DDF nodes need to be under identical root directories on their corresponding systems. On Windows, this means they must be under the same drive.

4.4.1.3. Highly Available Security Auditing

A third party tool will have to be used to persist the logs in a highly available manner.

- Edit the `<DDF_HOME>/etc/org.ops4j.pax.logging.cfg` file to enable log4j2, following the steps in [Enabling Fallback Audit Logging](#).
- Then put the appropriate log4j2 appender in `<DDF_HOME>/etc/log4j2.config.xml` to send logs to the chosen third party tool. See [Log4j Appenders](#).

4.4.1.4. Shared Storage Provider

The storage provider must be in a location that is shared between the two DDF nodes and must be highly available. If hardening the Highly Available Cluster, this shared storage provider must be trusted/secured. One way to accomplish this is to use the default [Content File System Storage Provider](#) and configure it to point to a highly available shared directory.

4.4.1.5. High Availability Certificates

Certificates need to have Subject Alternative Names (SANs) which will include the host for the HTTP proxy and for both DDF nodes. A certificate with SANs signed by the Demo CA can be obtained by following the steps at [Creating New Server Keystore Entry with the CertNew Scripts](#). This certificate is intended for testing purposes only, and should not be used in production.

4.4.1.6. High Availability Installation Profile

Instead of having to manually turn features on and off, there is a High Availability installation profile. This profile will not show up in the UI Installer, but can be installed by executing `profile:install ha` on the command line **instead** of stepping through the UI Installer. This profile will install all of the High Availability supported features.

5. Configuring

DDF is highly configurable and many of the components of the system can be configured to use an included DDF implementation or replaced with an existing component of an integrating system.

Configuration Requirements

NOTE

Because components can easily be installed and uninstalled, it's important to remember that for proper DDF functionality, at least the Catalog API, one Endpoint, and one Catalog Framework implementation must be active.

Configuration Tools

DDF provides several tools for configuring the system. The [Admin Console](#) is a useful interface for configuring applications, their features, and important settings. Alternatively, many configurations can be updated through [console commands](#) entered into the Command Console. Finally, configurations are stored in [configuration files](#) within the `<DDF_HOME>` directory.

Configuration Outline

While many configurations can be set or changed in any order, for ease of use of this documentation, similar subjects have been grouped together sequentially.

See [Keystores and certificates](#) to set up the certificates needed for messaging integrity and authentication. Set up [Users](#) with security attributes, then configure [data](#) attribute handling, and finally, define the [Security Policies](#) that map between users and data and make decisions about access.

Connecting DDF to other data sources, including other instances of DDF is covered in the [Configuring Federation](#) section.

Lastly, see the [Configuring for Special Deployments](#) section for guidance on common specialized installations, such as [fanout](#) or [multiple identical configurations](#).

5.1. Admin Console Tutorial

The Admin Console is the centralized location for administering the system. The Admin Console allows an administrator to install and remove selected applications and their dependencies and access configuration pages to configure and tailor system services and properties. The default address for the Admin Console is <https://{{FQDN}}:{{PORT}}/admin>.

Managing Applications from Admin Console

The **Manage** button enables activation/deactivation and adding/removing applications.

Activating / Deactivating Applications

The **Deactivate** button stops individual applications and any dependent apps. Certain applications are central to overall functionality and cannot be deactivated. These will have the **Deactivate** button disabled. Disabled apps will be moved to a list at the bottom of the page, with an enable button to reactivate, if desired.

The **Add Application** button is at the end of the list of currently active applications.

Removing Applications

To remove an application, it must first be deactivated. This enables the **Remove Application** button.

Upgrading Applications

Each application tile includes an upgrade button to select a new version to install.

System Settings Tab

The configuration and features installed can be viewed and edited from the System tab as well; however, it is recommended that configuration be managed from the applications tab.

IMPORTANT In general, applications should be managed via the applications tab. Configuration via this page could result in an unstable system. Proceed with caution!

Managing Federation in the Admin Console

It is recommended to use the **Catalog App** → **Sources** tab to configure and manage sites/sources.

Viewing Currently Active Applications from Admin Console

DDF displays all active applications in the Admin Console. This view can be configured according to preference. Either view has an > arrow icon to view more information about the application as currently configured.

Table 4. Admin Console Views

View	Description
Tile View	The first view presented is the Tile View, displaying all active applications as individual tiles.
List View	Optionally, active applications can be displayed in a list format by clicking the list view button.

Application Detailed View

Each individual application has a detailed view to view information specific to that application, adjust configurations or enable/disable features. All applications have a standard set of tabs, although some apps may have additional ones with further information.

Table 5. Individual Application Views

Tab	Explanation
Configuration	The Configuration tab lists all bundles associated with the application as links to configure any configurable properties of that bundle.
Details	The Details tab gives a description, version, status, and list of other applications that are required by, or rely on, the current application.
Features	The features tab breaks down the individual features of the application that can be installed or uninstalled as configurable features.

Managing Features Using the Admin Console

DDF includes many components, packaged as *features*, that can be installed and/or uninstalled without restarting the system. Features are collections of OSGi bundles, configuration data, and/or other features.

Transitive Dependencies

- NOTE** Features may have dependencies on other features and will auto-install them as needed.

In the Admin Console, Features are found on the **Features** tab of each application.

1. Select the appropriate application.
2. Select the **Features** tab.
3. Uninstalled features are shown with a **play** arrow under the **Actions** column.
 - a. Select the **play** arrow for the desired feature.
 - b. The **Status** will change from **Uninstalled** to **Installed**.
4. Installed features are shown with a **stop** icon under the **Actions** column.
 - a. Select the **stop** icon for the desired feature.
 - b. The **Status** will change from **Installed** to **Uninstalled**.

Adding Feature Repositories

If needed, custom feature repositories can be added to extend DDF functionality.

1. Select the **Manage** button in the upper right.
2. Select the **Add an Application** tile
3. Select **File Upload** to add a new **.kar** OR **.jar** file.
4. Select the **Maven URL** tab and enter the URL of the feature repository.
 - a. Select the **Add URL** button.
5. Select the **Save Changes** button.

5.2. Console Command Reference

DDF provides access to a powerful Command Console to use to manage and configure the system.

5.2.1. Application Commands

Applications can be installed from the Command Console using the following commands:

Table 6. App Commands

Command	Effect
<code>app:add <appName></code>	Install an app.
<code>app:list</code>	List all installed apps and current status.
<code>app:remove <appName></code>	Uninstall an app.
<code>app:start</code>	Start an inactive app.
<code>app:status <appName></code>	Detailed view of application status
<code>app:stop <appName></code>	Stop an active app.

Command	Effect
app:tree	Dependency tree view of all installed apps.

5.2.2. Feature Commands

Individual features can also be added via the Command Console.

1. Determine which feature to install by viewing the available features on DDF.

`ddf@local>feature:list`

2. The console outputs a list of all features available (installed and uninstalled). A snippet of the list output is shown below (the versions may differ):

State	Version	Name	Repository
Description			
[installed] [2.11.3] security-handler-api			security-services-
app-2.11.3 API for authentication handlers for web applications.			
[installed] [2.11.3] security-core			security-services-
app-2.11.3 DDF Security Core			
[uninstalled] [2.11.3] security-expansion			security-services-
app-2.11.3 DDF Security Expansion			
[uninstalled] [2.11.3] security-cas-client			security-services-
app-2.11.3 DDF Security CAS Client.			
[uninstalled] [2.11.3] security-cas-tokenvalidator			security-services-
app-2.11.3 DDF Security CAS Validator for the STS.			
[uninstalled] [2.11.3] security-cas-cxfservletfilter			security-services-
app-2.11.3 DDF Security CAS Servlet Filter for CXF.			
[installed] [2.11.3] security-pdp-authz			security-services-
app-2.11.3 DDF Security PDP.			
[uninstalled] [2.11.3] security-pep-serviceauthz			security-services-
app-2.11.3 DDF Security PEP Service AuthZ			
[uninstalled] [2.11.3] security-expansion-user-attributes			security-services-
app-2.11.3 DDF Security Expansion User Attributes Expansion			
[uninstalled] [2.11.3] security-expansion-metocard-attributes			security-services-
app-2.11.3 DDF Security Expansion Metocard Attributes Expansion			
[installed] [2.11.3] security-sts-server			security-services-
app-2.11.3 DDF Security STS.			
[installed] [2.11.3] security-sts-realm			security-services-
app-2.11.3 DDF Security STS Realm.			
[uninstalled] [2.11.3] security-sts-ldaplogin			security-services-
app-2.11.3 DDF Security STS JAAS LDAP Login.			
[uninstalled] [2.11.3] security-sts-ldapclaimshandler			security-services-
app-2.11.3 Retrieves claims attributes from an LDAP store.			

1. Check the bundle status to verify the service is started.

`ddf@local>list`

The console output should show an entry similar to the following:

```
[ 117] [Active      ] [           ] [Started] [    75] DDF :: Catalog :: Source ::  
Dummy (<version>)
```

5.2.2.1. Uninstalling Features from the Command Console

1. Check the feature list to verify the feature is installed properly.

```
ddf@local>feature:list
```

State	Version	Name	Repository
Description			
[installed]	[2.11.3]] ddf-core	ddf-2.11.3
[uninstalled]	[2.11.3]] ddf-sts	ddf-2.11.3
[installed]	[2.11.3]] ddf-security-common	ddf-2.11.3
[installed]	[2.11.3]] ddf-resource-impl	ddf-2.11.3
[installed]	[2.11.3]] ddf-source-dummy	ddf-2.11.3

1. Uninstall the feature.

```
ddf@local>feature:uninstall ddf-source-dummy
```

WARNING Dependencies that were auto-installed by the feature are not automatically uninstalled.

1. Verify that the feature has uninstalled properly.

```
ddf@local>feature:list
```

State	Version	Name	Repository	Description
[installed]	[2.11.3]] ddf-core	ddf-2.11.3	
[uninstalled]	[2.11.3]] ddf-sts	ddf-2.11.3	
[installed]	[2.11.3]] ddf-security-common	ddf-2.11.3	
[installed]	[2.11.3]] ddf-resource-impl	ddf-2.11.3	
[uninstalled]	[2.11.3]] ddf-source-dummy	ddf-2.11.3	

5.3. Configuration Files

Many important configuration settings are stored in the `<DDF_HOME>` directory.

NOTE

Depending on the environment, it may be easier for integrators and administrators to configure DDF using the Admin Console prior to disabling it for hardening purposes. The Admin Console can be re-enabled for additional configuration changes.

In an environment hardened for security purposes, access to the Admin Console or the Command Console might be denied and using the latter in such an environment may cause configuration errors. It is necessary to configure DDF (e.g., providers, Schematron rulesets, etc.) using `.config` files.

A template file is provided for some configurable DDF items so that they can be copied/renamed then modified with the appropriate settings.

WARNING

If the Admin Console is enabled again, all of the configuration done via `.config` files will be loaded and displayed. However, note that the name of the `.config` file is not used in the Admin Console. Rather, a universally unique identifier (UUID) is added when the DDF item was created and displays this UUID in the console (e.g., `OpenSearchSource.112f298e-26a5-4094-befc-79728f216b9b`)

5.3.1. Configuring Global Settings with `system.properties`

Global configuration settings are configured via the properties file `system.properties`. These properties can be manually set by editing this file or set via the initial configuration from the Admin Console.

NOTE

Any changes made to this file require a restart of the system to take effect.

IMPORTANT

The passwords configured in this section reflect the passwords used to decrypt JKS (Java KeyStore) files. Changing these values without also changing the passwords of the JKS causes undesirable behavior.

Table 7. Global Settings

Title	Property	Type	Description	Default Value	Required
Keystore and truststore java properties					
Keystore	<code>javax.net.ssl.keyStore</code>	String	Path to server keystore	<code>etc/keystores/serverKeystore.jks</code>	Yes
Keystore Password	<code>javax.net.ssl.keyStorePassword</code>	String	Password for accessing keystore	<code>changeit</code>	Yes
Truststore	<code>javax.net.ssl.trustStore</code>	String	The trust store used for SSL/TLS connections. Path is relative to <code><DDF_HOME></code> .	<code>etc/keystores/serverTruststore.jks</code>	Yes
Truststore Password	<code>javax.net.ssl.trustStorePassword</code>	String	Password for server Truststore	<code>changeit</code>	Yes
Keystore Type	<code>javax.net.ssl.keyStoreType</code>	String	File extension to use with server keystore	<code>jks</code>	Yes
Truststore Type	<code>javax.net.ssl.trustStoreType</code>	String	File extension to use with server truststore	<code>jks</code>	Yes
Headless mode					

Title	Property	Type	Description	Default Value	Required
Headless Mode	<code>java.awt.headless</code>	Boolean	Force java to run in headless mode for when the server doesn't have a display device	<code>true</code>	No
Global URL Properties					
Default Protocol	<code>org.codice.ddf.system.protocol</code>	String	Default protocol that should be used to connect to this machine.	<code>https://</code>	Yes
Host	<code>org.codice.ddf.system.hostname</code>	String	<p>The hostname or IP address used to advertise the system. Do not enter <code>localhost</code>. Possibilities include the address of a single node or that of a load balancer in a multi-node deployment.</p> <p>If the hostname is changed during the install to something other than <code>localhost</code> a new keystore and truststore must be provided. See Managing Keystores and Certificates for details.</p> <p>NOTE: Does not change the address the system runs on.</p>	<code>localhost</code>	Yes
HTTPS Port	<code>org.codice.ddf.system.httpsPort</code>	String	<p>The https port used to advertise the system.</p> <p>NOTE: This does not change the port the system runs on.</p>	<code>8993</code>	Yes

Title	Property	Type	Description	Default Value	Required
HTTP Port	<code>org.codice.ddf.system.httpPort</code>	String	<p>The http port used to advertise the system.</p> <p>NOTE: This does not change the port the system runs on.</p>	8181	Yes
Default Port	<code>org.codice.ddf.system.port</code>	String	<p>The default port used to advertise the system. This should match either the above http or https port.</p> <p>NOTE: Does not change the port the system runs on.</p>	8993	Yes
Internal HTTPS Port	<code>org.codice.ddf.system.internalHttpsPort</code>	String	<p>The https port that the system uses.</p> <p>NOTE: This DOES change the port the system runs on.</p>	8993	Yes
Internal HTTP Port	<code>org.codice.ddf.system.internalHttpPort</code>	String	<p>The http port that the system uses.</p> <p>NOTE: This DOES change the port the system runs on.</p>	8181	Yes
Root Context	<code>org.codice.ddf.system.rootContext</code>	String	<p>The the base or root context that services will be made available under.</p>	/services	Yes

System Information Properties

Site Name	<code>org.codice.ddf.system.siteName</code>	String	The site name for DDF.	<code>ddf.distribution</code>	Yes
Site Contact	<code>org.codice.ddf.system.siteContact</code>	String	The email address of the site contact.		No
Version	<code>org.codice.ddf.system.version</code>	String	<p>The version of DDF that is running.</p> <p>This value should not be changed from the factory default.</p>	2.11.3	Yes

Title	Property	Type	Description	Default Value	Required
Organization	<code>org.codice.ddf.system.organization</code>	String	The organization responsible for this installation of DDF.	Codice Foundation	Yes
Registry ID	<code>org.codice.ddf.system.registry-id</code>	String	The registry id for this installation of DDF.		No
Thread Pool Settings					
Thread Pool Size	<code>org.codice.ddf.system.threadPoolSize</code>	Integer	Size of thread pool used for handling UI queries, federating requests, and downloading resources. See Configuring Thread Pools	128	Yes
HTTPS Specific Settings					
Cipher Suites	<code>https.cipherSuites</code>	String	Cipher suites to use with secure sockets. If using the JCE unlimited strength policy, use this list in place of the defaults: .	<code>TLS_DHE_RSA_WITH_AES_128_GCM_SHA256,</code> <code>TLS_DHE_RSA_WITH_AES_128_CBC_SHA256,</code> <code>TLS_DHE_RSA_WITH_AES_128_CBC_SHA,</code> <code>TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256,</code> <code>TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256</code>	No
Https Protocols	<code>https.protocols</code>	String	Protocols to allow for secure connections	<code>TLSv1.1, TLSv1.2</code>	No
Allow Basic Auth Over Http	<code>org.codice.allowBasicAuthOverHttp</code>	Boolean	Set to true to allow Basic Auth credentials to be sent over HTTP unsecurely. This should only be done in a test environment. These events will be audited.	<code>false</code>	Yes

Title	Property	Type	Description	Default Value	Required
Restrict the Security Token Service to allow connections only from DNs matching these patterns	<code>ws-security.subject.cert.constraints</code>	String	Set to a comma separated list of regex patterns to define which hosts are allowed to connect to the STS	<code>.*</code>	Yes
XML Settings					
Parse XML documents into DOM object trees	<code>javax.xml.parsers.DocumentBuilderFactory</code>	String	Enables Xerces-J implementation of <code>DocumentBuilderFactory</code>	<code>org.apache.xerces.jaxp.DocumentBuilderFactoryImpl</code>	Yes
Catalog Source Retry Interval					
Initial Endpoint Contact Interval	<code>org.codice.ddf.platform.util.http.initialRetryInterval</code>	Integer	If a Catalog Source is unavailable, try to connect to it after the initial interval has elapsed. After every retry, the interval doubles, up to a given maximum interval. The interval is measured in seconds.	<code>10</code>	Yes
Maximum Endpoint Contact Interval	<code>Maximum seconds between attempts to establish contact with unavailable Catalog Source.</code>	Integer	Do not wait longer than the maximum interval to attempt to establish a connection with an unavailable Catalog Source. Smaller values result in more current information about the status of Catalog Sources, but cause more network traffic. The interval is measured in seconds.	<code>300</code>	Yes
File Upload Settings					

Title	Property	Type	Description	Default Value	Required
File extensions flagged as potentially dangerous to the host system or external clients	bad.file.extensions	String	<p>Files uploaded with these bad file extensions will have their file names sanitized before being saved</p> <p>E.g. sample_file.exe will be renamed to sample_file.bin upon ingest</p>	.exe, .jsp, .html, .js, .php, .phtml, .php3, .php4, .php5, .phps, .shtml, .jhtml, .pl, .py, .cgi, .msi, .com, .scr, .gadget, .application, .pif, .hta, .cpl, .msc, .jar, .kar, .bat, .cmd, .vb, .vbs, .vbe, .jse, .ws, .wsf, .wsc, .wsh, .ps1, .ps1xml, .ps2, .ps2xml, .psc1, .psc2, .msh, .msh1, .msh2, .mshxml, .msh1xml, .msh2xml, .scf, .lnk, .inf, .reg, .dll, .vxd, .cpl, .cfg, .config, .crt, .cert, .pem, .jks, .p12, .p7b, .key, .der, .csr, .jsb, .mhtml, .mht, .xhtml, .xht	Yes
File names flagged as potentially dangerous to the host system or external clients	bad.files	String	<p>Files uploaded with these bad file names will have their file names sanitized before being saved</p> <p>E.g. crossdomain.xml will be renamed to file.bin upon ingest</p>	crossdomain.xml, clientaccesspolicy.xml, .htaccess, .htpasswd, hosts, passwd, group, resolv.conf, nfs.conf, ftpd.conf, ntp.conf, web.config, robots.txt	Yes
Mime types flagged as potentially dangerous to external clients	bad.mime.types	String	Files uploaded with these mime types will be rejected from the upload	text/html, text/javascript, text/x-javascript, application/x-shellscript, text/scriptlet, application/x-msdownload, application/x-msmetafile	Yes
File names flagged as potentially dangerous to external clients	ignore.files	String	Files uploaded with these file names will be rejected from the upload	.DS_Store, Thumbs.db	Yes
Solr Catalog Provider Properties					

Title	Property	Type	Description	Default Value	Required
Solr Client	<code>solr.client</code>	String	Solr Catalog Provider selected.	<code>HttpSolrClient, CloudSolrClient, EmbeddedSolrServer</code>	Yes
Solr Data Directory	<code>solr.data.dir</code>	String	Solr Data Directory	<code> \${ddf.home}/data/solr</code>	Yes
Solr Catalog Provider Url	<code>solr.http.url</code>	String	Url for a HTTP Solr server (required for HTTP Solr)	<code> \${org.codice.ddf.system.protocol}\${org.codice.ddf.system.hostname}: \${org.codice.ddf.system.httpsPort}/solr</code>	No
Zookeeper Nodes	<code>solr.cloud.zookeeper</code>	String	Zookeeper hostnames and port numbers (required for Solr Cloud)	<code>zookeeperhost1:2181, zookeeperhost2:2181, zookeeperhost3:2181</code>	No

These properties are available to be used as variable parameters in input url fields within the Admin Console. For example, the url for the local csw service (`https://{FQDN}:{PORT}/services/csw`) could be defined as:

```
 ${org.codice.ddf.system.protocol}${org.codice.ddf.system.hostname}: ${org.codice.ddf.system.port} ${org.codice.ddf.system.rootContext}/csw
```

This variable version is more verbose, but will not need to be changed if the system `host`, `port` or `root` context changes.

WARNING Only root can access ports < 1024 on Unix systems.

5.3.2. Configuring with .config Files

The DDF is configured using `.config` files. Like the Karaf `.cfg` files, these configuration files must be located in the `<DDF_HOME>/etc/` directory. Unlike the Karaf `.cfg` files, `.config` files must follow the naming convention that includes the *configuration persistence ID* (PID) that they represent. The filenames must be the pid with a `.config` extension. This type of configuration file also supports lists within configuration values (metatype `cardinality` attribute greater than 1) and String, Boolean, Integer, Long, Float, and Double values.

IMPORTANT

This new configuration file format **must** be used for any configuration that makes use of lists. Examples include Web Context Policy Manager (`org.codice.ddf.security.policy.context.impl.PolicyManager.config`) and Security STS Guest Claims Handler (`ddf.security.sts.guestclaims.config`).

WARNING

Only one configuration file should exist for any given PID. The result of having both a `.cfg` and a `.config` file for the same PID is undefined and could cause the application to fail.

The main purpose of the configuration files is to allow administrators to pre-configure DDF without having to use the Admin Console. In order to do so, the configuration files need to be copied to the `<DDF_HOME>/etc` directory after DDF zip has been extracted.

Upon start up, all the `.config` files located in `<DDF_HOME>/etc` are automatically read and processed. DDF monitors the `<DDF_HOME>/etc` directory for any new `.config` file that gets added. As soon as a new file is detected, it is read and processed. Changes to these configurations from the Admin Console or otherwise are persisted in the original configuration file in the `<DDF_HOME>/etc` directory.


5.4. Managing Keystores and Certificates

- **Required Step for Security Hardening**

DDF uses certificates in two ways:

1. Ensuring the privacy and integrity of messages sent or received over a network.
2. Authenticating an incoming user request.

To ensure proper configuration of keystore, truststore, and certificates, follow the options below according to situation.



Configuring Certificates Workflow

Jump to the steps referenced in the diagram:

Certificate Workflow Steps

1. [Adding an Existing Keystore](#)
2. [Creating a New Keystore/Truststore with an Existing Certificate and Private Key](#)
3. [Adding an Existing Truststore](#)
4. [Creating a Server Keystore](#)
 - a. [Creating a Server Truststore](#)

5.4.1. Managing Keystores

Certificates, and sometimes their associated private keys, are stored in keystore files. DDF includes two default keystore files, the server key store and the server trust store. The server keystore holds the certificates and private keys that DDF uses to identify itself to other nodes on the network. The truststore holds the certificates of nodes or other entities that DDF needs to trust.

5.4.1.1. Adding an Existing Server Keystore

If provided an existing keystore for use with DDF, follow these steps to replace the default keystore.

1. Remove the default keystore at `etc/keystore/serverKeystore.jks`.
2. Add the desired keystore file to the `etc/keystore` directory.
3. Edit `system.properties` file to set filenames and passwords.
 - a. If using a type of keystore other than `jks` (such as `pkcs12`), change the `javax.net.ssl.keyStoreType` property as well.
4. If the truststore has the correct certificates, restart server to complete configuration.
 - a. If provided with an existing server truststore, continue to [Adding an Existing Server Truststore](#).
 - b. Otherwise, [create a server truststore](#).

5.4.1.2. Adding an Existing Server Truststore

1. Remove the default truststore at `etc/keystore/serverTruststore.jks`.
2. Add the desired truststore file to the `etc/keystore` directory.
3. Edit `system.properties` file to set filenames and passwords.
 - a. If using a type of truststore other than `jks` (such as `pkcs12`), change the `javax.net.ssl.trustStoreType` property as well.

If the provided server keystore does not include the CA certificate that was used to sign the server's certificate, [add the CA certificate into the `serverKeystore` file](#).

5.4.1.3. Creating a New Keystore/Truststore with an Existing Certificate and Private Key

If provide an existing certificate create a new keystore and truststore with it.

NOTE DDF requires that the keystore contains both the private key and the CA.

1. Using the private key, certificate, and CA certificate, create a new keystore containing the data from the new files.

```
cat client.crt >> client.key
openssl pkcs12 -export -in client.key -out client.p12
keytool -importkeystore -srckeystore client.p12 -destkeystore serverKeystore.jks
-srcstoretype pkcs12 -alias 1
keytool -changealias -alias 1 -destalias client -keystore serverKeystore.jks
keytool -importcert -file ca.crt -keystore serverKeystore.jks -alias "ca"
keytool -importcert -file ca-root.crt -keystore serverKeystore.jks -alias "ca-root"
```

2. Create the truststore using only the CA certificate. Based on the concept of CA signing, the CA should be the only entry needed in the truststore.

```
keytool -import -trustcacerts -alias "ca" -file ca.crt -keystore truststore.jks
keytool -import -trustcacerts -alias "ca-root" -file ca-root.crt -keystore
truststore.jks
```

3. Create a PEM file using the certificate, as some applications require that format.

```
openssl x509 -in client.crt -out client.der -outform DER
openssl x509 -in client.der -inform DER -out client.pem -outform PEM
```

5.4.1.4. Updating Key Store / Trust Store via the Admin Console

Certificates (and certificates with keys) can be managed in the Admin Console.

1. Navigate to the Admin Console.
2. Select the **Security** application.
3. Select the **Certificates** tab.
4. Add and remove certificates and private keys as necessary.
5. Restart DDF.

IMPORTANT

The default trust store and key store files for DDF included in [etc/keystores](#) use self-signed certificates. Self-signed certificates should never be used outside of development/testing areas.

This view shows the alias (name) of every certificate in the trust store and the key store. It also displays if the entry includes a private key ("Is Key") and the encryption scheme (typically "RSA" or "EC").

This view allows administrators remove certificates from DDF's key and trust stores. It also allows administrators to import certificates and private keys into the keystores with the "+" button. The import function has two options: import from a file or import over HTTPS. The file option accepts a Java Keystore file or a PKCS12 keystore file. Because keystores can hold many keys, the import dialog asks the administrator to provide the alias of the key to import. Private keys are typically encrypted and the import dialog prompts the administrator to enter the password for the private key. Additionally, keystore files themselves are typically encrypted and the dialog asks for the keystore ("Store") password.

The name and location of the DDF trust and key stores can be changed by editing the system properties files, [etc/system.properties](#). Additionally, the password that DDF uses to decrypt (unlock) the key and trust stores can be changed here.

IMPORTANT

DDF assumes that password used to unlock the keystore is the same password that unlocks private keys in the keystore.

The location, file name, passwords and type of the server and trust key stores can be set in the [system.properties](#) file:

1. Setting the Keystore and Truststore Java Properties

```
javax.net.ssl.keyStore=etc/keystores/serverKeystore.jks
javax.net.ssl.keyStorePassword=changeit
javax.net.ssl.trustStore=etc/keystores/serverTruststore.jks
javax.net.ssl.trustStorePassword=changeit
javax.net.ssl.keyStoreType=jks
javax.net.ssl.trustStoreType=jks
```

NOTE

If the server's fully qualified domain name is not recognized, the name may need to be added to the network's DNS server.

The DDF instance can be tested even if there is no entry for the FQDN in the DNS. First, test if the FQDN is already recognized. Execute this command:

```
ping <FQDN>
```

TIP

If the command responds with an error message such as unknown host, then modify the system's `hosts` file to point the server's FQDN to the loopback address. For example:

```
127.0.0.1 <FQDN>
```

Changing Default Passwords

This step is not required for a hardened system.

- The default password in `system.properties` for `serverKeystore.jks` is `changeit`. This needs to be modified.
 - `ds-cfg-key-store-file: ../../keystores/serverKeystore.jks`
 - `ds-cfg-key-store-type: JKS`
 - `ds-cfg-key-store-pin: password`
 - `cn: JKS`
- The default password in `system.properties` for `serverTruststore.jks` is `changeit`. This needs to be modified.
 - `ds-cfg-trust-store-file: ../../keystores/serverTruststore.jks`
 - `ds-cfg-trust-store-pin: password`
 - `cn: JKS`

NOTE

5.4.2. Disallowing Login Without Certificates

DDF can be configured to prevent login without a valid PKI certificate.

- Navigate to **Admin Console**
- Under **Security**, select → **Web Context Policy Manager**
- Add a policy for each context requiring restriction

- For example: `/search=SAML|PKI` will disallow login without certificates to the Search UI.
- The format for the policy should be: `/<CONTEXT>=SAML|PKI`
- Click **Save**

NOTE Ensure certificates comply with organizational hardening policies.

5.4.3. Managing Certificate Revocation List (CRL)

- **Required Step for Security Hardening**

For hardening purposes, it is recommended to implement a way to verify the CRL at least daily.

A Certificate Revocation List is a collection of formerly-valid certificates that should explicitly *not* be accepted.

5.4.3.1. Creating a Certificate Revocation List (CRL)

Create a CRL in which the token issuer's certificate is valid. The example uses OpenSSL.

```
$> openssl ca -gencrl -out crl-tokenissuer-valid.pem
```

Windows and OpenSSL

NOTE Windows does not include OpenSSL by default. For Windows platforms, a additional download of [OpenSSL](#) or an alternative is required.

5.4.3.1.1. Revoke a Certificate and Create a New CRL that Contains the Revoked Certificate

```
$> openssl ca -revoke tokenissuer.crt
$> openssl ca -gencrl -out crl-tokenissuer-revoked.pem
```

5.4.3.1.2. Viewing a CRL

1. Use the following command to view the serial numbers of the revoked certificates: `$> openssl crl -inform PEM -text -noout -in crl-tokenissuer-revoked.pem`

5.4.3.2. Enabling Certificate Revocation

NOTE Enabling CRL revocation or modifying the CRL file will require a restart of DDF to apply updates.

1. Place the CRL in <DDF_HOME>/etc/keystores.
2. Add the line `org.apache.ws.security.crypto.merlin.x509crl.file=etc/keystores/<CRL_FILENAME>` to the following files (Replace <CRL_FILENAME> with the URL or file path of the CRL location):
 - a. `<DDF_HOME>/etc/ws-security/server/encryption.properties`
 - b. `<DDF_HOME>/etc/ws-security/issuer/encryption.properties`
 - c. `<DDF_HOME>/etc/ws-security/server/signature.properties`

- d. <DDF_HOME>/etc/ws-security/issuer/signature.properties
3. (Replace <CRL_FILENAME> with the file path or URL of the CRL file used in previous step.)

Adding this property will also enable CRL revocation for any context policy implementing PKI authentication. For example, adding an authentication policy in the Web Context Policy Manager of `/search=SAML|PKI` will disable basic authentication, require a certificate for the search UI, and allow a SAML SSO session to be created. If a certificate is not in the CRL, it will be allowed through, otherwise it will get a 401 error. If no certificate is provided, the guest handler will grant guest access.

This also enables CRL revocation for the STS endpoint. The STS CRL Interceptor monitors the same `encryption.properties` file and operates in an identical manner to the PKI Authentication's CRL handler. Enabling the CRL via the `encryption.properties` file will also enable it for the STS, and also requires a restart.

5.4.3.2.1. Add Revocation to a Web Context

The `PKIHandler` implements CRL revocation, so any web context that is configured to use PKI authentication will also use CRL revocation if revocation is enabled.

1. After enabling revocation (see above), open the **Web Context Policy Manager**.
2. Add or modify a Web Context to use PKI in authentication. For example, enabling CRL for the search ui endpoint would require adding an authorization policy of `/search=SAML|PKI`
3. If guest access is required, add `GUEST` to the policy. Ex, `/search=SAML|PKI|GUEST`.

With guest access, a user with a revoked certificate will be given a 401 error, but users without a certificate will be able to access the web context as the guest user.

The STS CRL interceptor does not need a web context specified. The CRL interceptor for the STS will become active after specifying the CRL file path, or the URL for the CRL, in the `encryption.properties` file and restarting DDF.

NOTE Disabling or enabling CRL revocation or modifying the CRL file will require a restart of DDF to apply updates. If CRL checking is already enabled, adding a new context via the **Web Context Policy Manager** will not require a restart.

5.4.3.2.2. Adding Revocation to an Endpoint

NOTE This section explains how to add CXF's CRL revocation method to an endpoint and not the CRL revocation method in the `PKIHandler`.

This guide assumes that the endpoint being created uses CXF and is being started via Blueprint from inside the OSGi container. If other tools are being used the configuration may differ.

Add the following property to the `jasws` endpoint in the endpoint's `blueprint.xml`:

```
<entry key="ws-security.enableRevocation" value="true"/>
```

Example xml snippet for the `jaxws:endpoint` with the property:

```
<jaxws:endpoint id="Test" implementor="#testImpl"
    wsdlLocation="classpath: META-INF/wsdl/TestService.wsdl"
    address="/TestService">

    <jaxws:properties>
        <entry key="ws-security.enableRevocation" value="true"/>
    </jaxws:properties>
</jaxws:endpoint>
```

5.4.3.2.3. Verifying Revocation

A **Warning** similar to the following will be displayed in the logs of the source and endpoint showing the exception encountered during certificate validation:

```
11:48:00,016 | WARN  | tp2085517656-302 | WSS4JInInterceptor           |
security.wss4j.WSS4JInInterceptor 330 | 164 - org.apache.cxf.cxf-rt-ws-security -
2.7.3 |
org.apache.ws.security.WSSecurityException: General security error (Error during
certificate path validation: Certificate has been revoked, reason: unspecified)
    at
org.apache.ws.security.components.crypto.Merlin.verifyTrust(Merlin.java:838)[161:org.a
pache.ws.security.wss4j:1.6.9]
    at
org.apache.ws.security.validate.SignatureTrustValidator.verifyTrustInCert(SignatureTru
stValidator.java:213)[161:org.apache.ws.security.wss4j:1.6.9]

[ ... section removed for space]

Caused by: java.security.cert.CertPathValidatorException: Certificate has been
revoked, reason: unspecified
    at
sun.security.provider.certpath.PKIXMasterCertPathValidator.validate(PKIXMasterCertPath
Validator.java:139)[:1.6.0_33]
    at
sun.security.provider.certpath.PKIXCertPathValidator.doValidate(PKIXCertPathValidator.
java:330)[:1.6.0_33]
    at
sun.security.provider.certpath.PKIXCertPathValidator.engineValidate(PKIXCertPathValida
tor.java:178)[:1.6.0_33]
    at
java.security.cert.CertPathValidator.validate(CertPathValidator.java:250)[:1.6.0_33]
    at
org.apache.ws.security.components.crypto.Merlin.verifyTrust(Merlin.java:814)[161:org.a
pache.ws.security.wss4j:1.6.9]
    ... 45 more
```

5.5. Configuring User Access

DDF does not define accounts or types of accounts to support access. DDF uses an *attribute based access control (ABAC)* model. For reference, ABAC systems control access by evaluating rules against the attributes of the entities (*subject* and *object*), actions, and the environment relevant to a request.

DDF can be configured to access many different types of user stores to manage and monitor user access.

5.5.1. Configuring Guest Access

Unauthenticated access to a secured DDF system is provided by the **Guest** user. By default, DDF allows guest access as part of the karaf security realm.

Because DDF does not know the identity of a Guest user, it cannot assign security attributes to the Guest. The administrator must configure the attributes and values (i.e. the "claims") to be assigned to Guests. The Guest Claims become the default minimum attributes for every user, both authenticated and unauthenticated. Even if a user claim is more restrictive, the guest claim will grant access, so ensure the guest claim is only as permissive as necessary.

- Make sure that all the default logical names for locations of the security services are defined.

5.5.1.1. Denying Guest User Access

By default, the Search UI allows guest access as part of the karaf security realm. To disable guest access for a context, use the [Web Context Policy Manager](#) configuration to remove **Guest** from the **Authentication Type** for that context. Only authorized users are then allowed to continue to the Search UI page.

5.5.1.1.1. Configuring SOAP Guest Interceptor if Prohibiting Guest Users

- **Required Step for Security Hardening**

Uninstall the [Guest Interceptor](#).

5.5.1.1.2. Configuring Guest Validator if Prohibiting Guest Users

Uninstall the [Guest Validator](#).

5.5.1.2. Allowing Guest User Access

Guest authentication must be enabled and configured to allow guest users. Once the guest user is configured, redaction and filtering of metadata is done for the guest user the same way it is done for normal users.

To enable guest authentication for a context, use the [Web Context Policy Manager](#) configuration to change the **Authentication Type** for that context to **Guest**.

1. Navigate to the **Admin Console**.

2. Select the **Security** application.
3. Select the **Configuration** tab.
4. Select **Web Context Policy Manager**.
5. Select the desired context (`/`, `/search`, `/admin`, etc.).
6. Add **Guest** to the **Authentication Type** list.
 - a. Separate entries with a `|` symbol (eg. `/=SAML|Guest`).

5.5.1.2.1. Configuring Guest Interceptor if Allowing Guest Users

- **Required Step for Security Hardening**

If a legacy client requires the use of the secured SOAP endpoints, the [guest interceptor](#) should be configured. Otherwise, the guest interceptor and [public](#) endpoints should be uninstalled for a hardened system.

5.5.1.2.2. Configuring Guest Claim Attributes

A guest user's attributes define the most permissive set of claims for an unauthenticated user.

A guest user's claim attributes are stored in configuration, not in the LDAP as normal authenticated users' attributes are.

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Select the **Configuration** tab.
4. Select the **Security Guest Claims Handler**.
5. Add any additional attributes desired for the guest user.
6. Save changes.

5.5.2. Configuring REST Services for Users

If using REST services or connecting to REST sources, several configuration options are available.

DDF includes a [Identity Provider\(IdP\)](#), but can also be configured to support an [external IdP](#) or [no IdP](#) at all.

5.5.2.1. Configuring Included IdP

Installing the IdP from the Admin Console

Install the included IdP implementation.

1. Navigate to the **Admin Console**.
2. Select **Features**.
3. Install **security-idp** feature.

Installing the IdP from the Command Console

Install the included IdP implementation.

1. Run the command `features:install security-idp` from the Command Console.

Configuring the IdP Server

1. Navigate to **Admin Console**.
2. Select the **Security** application.
3. Select the **Configuration** tab.
4. Select **IdP Server**.
5. Click the **+** next to SP Metadata to add a new entry.
6. Populate the new entry:
 - a. with an HTTPS URL (<https://>),
 - b. file URL (file:), or
 - c. XML block to refer to desired metadata, e.g. (<https://localhost:8993/services/saml/sso/metadata>)

IdP Server example.xml

```
<md:EntityDescriptor xmlns:md="urn:oasis:names:tc:SAML:2.0:metadata" entityID="https://localhost:8993/services/saml">
  <md:SPSSODescriptor protocolSupportEnumeration="urn:oasis:names:tc:SAML:2.0:protocol">
    <md:KeyDescriptor use="signing">
      <ds:KeyInfo xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
        <ds:X509Data>
          <ds:X509Certificate>
MIIDEzCCAnygAwIBAgIJJAIZc4FYrIp9mMA0GCSqGSIB3DQEBBQUAMHcxCzAJBgNVBAYTA1VTMQswCQYDVQQIDA
JBWjEMMAoGA1UECgwDRERGMQwwCgYDVQQLDANEZXYxGTAXBgNVBAMMEERERiBEZW1vIFJvb3QgQ0ExJDAiBjkq
hkig9w0BCQEWFWRkZnJvb3RjYUBleGFtcGx1Lm9yZzAeFw0xNDEyMTAyMTU4MThaFw0xNTEyMTAyMTU4MThaMI
GDMQswCQYDVQQGEwJVUzELMAkGA1UECAwCQVoxETAPBgNVBAcMCEdvb2R5ZWFrMQwwCgYDVQQKDANEREYxDDAK
BgNVBAsMA0RldjESMBAGA1UEAwwJbG9jYWxob3N0MSQwIgYJKoZIhvcNAQkBFhVsb2NhbGhvc3RAZXhhbXBsZS
5vcmcwgZ8wDQYJKoZIhvcNAQEBBQADgY0AMIGJAoGBAMeCyNzbCTZphHQfB5g8FrgBq1RYzV7ikVw/pVGkz8gx
313A99s8WtA4mRAeb6n0vTR9yNBOekW4nY0iE0q//YTi/frI1kz0QbEH1s2cI5nFButabD3PYGxUSuapbc+AS7
+Pk1r0TDI4MRzPPkkTp4w10RQ/a6CfVsNr/mVgL2CfAgMBAAGjgZkwgZYwCQYDVR0TBAIwADAnBglghkgBhvC
AQ0EGhYYRk9S1FRFU1RJTkcgUFVSUE9TRSBPTkxZMB0GA1UdDgQWBBSA95QIMyBAHRsd0R4s7C3BreFrDAfBg
NVHSMEGDAWgBThVMeX3wrcv61feF47CvkSBe9xjAgBqNVHREEGTAXgRVsb2NhbGhvc3RAZXhhbXBsZS5vcmcw
DQYJKoZIhvcNAQEFBQADgYEAtRUp7fAxU/E6JD2Kj/+CTWqu8Elx13S0TxoIqv3gMoBW0ehyzEKjJi0bb1gUx0
7n1Sm0ESp5sE3jGTnh0GtYV0D219z/09n90cd/imAEhknJlavyd0SjpnaL9JUd8uYxJexy8TJ2sMhsGAZ6EMTZ
CfT9m07XduxjsmDz0h1SGV0=
        </ds:X509Certificate>
      </ds:X509Data>
    </ds:KeyInfo>
  </md:KeyDescriptor>
  <md:KeyDescriptor use="encryption">
    <ds:KeyInfo xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
      <ds:X509Data>
```

<ds:X509Certificate>

```
MIIDExCCAnygAwIBAgIJAIzc4FYrIp9mMA0GCSqGSIB3DQEBBQUAMHcxCzAJBgNVBAYTA1VTMQswCQYDVQQIDA
JBWjEMMAoGA1UECgwDRERGMQwwCgYDVQQLDANEZXYxGTAXBgNVBAMMEERERiBEZW1vIFJvb3QgQ0ExJDAiBpkq
hkiG9w0BCQEWFRkZnJvb3RjYUB1eGFtcGx1Lm9yZzAeFw0xNDEyMTAyMTU4MThaFw0xNTEyMTAyMTU4MThaMI
GDMQswCQYDVQQGEwJVUzELMAkGA1UECAwCQVoxETAPBgNVBAcMCEdvb2R5ZWFrMQwwCgYDVQQKDANEREYxDDAK
BgNVBAsMA0RldjESMBAGA1UEAwJbG9jYWxob3N0MSQwIgYJKoZIhvcNAQkBFhVsb2NhbGhvc3RAZXhhXBsZS
5vcmcwgZ8wDQYJKoZIhvcNAQEBBQADgY0AMIGJAoGBAMeCyNZbCTZphHQfB5g8FrgBq1RYzV7ikVw/pVGkz8gx
313A99s8WtA4mAeb6n0vTR9yNBOekW4nY0iE0q//YTi/frI1kz0QbEH1s2cI5nFButabD3PYGxUSuapbc+AS7
+Pk1r0TDI4MRzPPkkTp4w10RQ/a6CfVsNr/mVgL2CfAgMBAAGjgZkwgZYwCQYDVR0TBAIwADAnBglghkgBvhvC
AQ0EGhYYRk9S1FRFU1RJTkcgUFVSUE9TRSBPTkxZMB0GA1UdDgQWBBSA95QIMyBAHRsd0R4s7C3BreFrsDAfBg
NVHSMEGDAWgBThVMeX3wrCv61feF47CvkSBe9xjAgBqNVHREEGTAXgRVsb2NhbGhvc3RAZXhhXBsZS5vcmcw
DQYJKoZIhvcNAQEFBQADgYEAtRUp7fAxU/E6JD2Kj/+CTWqu8Elx13S0TxoIqv3gMoBW0ehyzEKjJi0bb1gUx0
7n1Sm0ESp5sE3jGTnh0GtYV0D219z/09n90cd/imAEhknJlayyd0Sjpnal9JUd8uYxJexy8TJ2sMhsGAZ6EMTz
CfT9m07XduxjsmDz0h1SGV0=
```

</ds:X509Certificate>

</ds:X509Data>

</ds:KeyInfo>

</md:KeyDescriptor>

<md:SingleLogoutService Binding="urn:oasis:names:tc:SAML:2.0:bindings:HTTP-Redirect" Location="https://localhost:8993/logout"/>

<md:SingleLogoutService Binding="urn:oasis:names:tc:SAML:2.0:bindings:HTTP-POST" Location="https://localhost:8993/logout"/>

<md:AssertionConsumerService Binding="urn:oasis:names:tc:SAML:2.0:bindings:HTTP-Redirect" Location="https://localhost:8993/services/saml/sso"/>

<md:AssertionConsumerService Binding="urn:oasis:names:tc:SAML:2.0:bindings:HTTP-POST" Location="https://localhost:8993/services/saml/sso"/>

</md:SPSSODescriptor>

</md:EntityDescriptor>

Configuring IdP Authentication Types

Set the authentication types that will be accepted by the IdP.

1. Navigate to the Admin Console.
2. Select the Security application.
3. Select the Configuration tab.
4. Select Web Context Policy Manager
5. Under Authentication Types, set the IDP authentication type as necessary. Note that it should only be used on context paths that will be accessed by users via web browsers. For example:
 - `/search=IDP`

Other authentication types can also be used in conjunction with the IDP type. For example, if you wanted to secure the entire system with the IDP, but still allow legacy clients that don't understand the SAML ECP specification to connect, you could set `/=IDP|PKI`. With that configuration, any clients that failed to connect using either the SAML 2.0 Web SSO Profile or the SAML ECP spec would fall back to 2-way TLS for authentication.

NOTE

If you have configured /search to use IDP, ensure to select the "External Authentication" checkbox in Search UI standard settings.

Configuring Security IdP Service Provider (SP)

The IdP client that interacts with the specified Identity Provider.

1. Navigate to Admin Console.
2. Select the Security application.
3. Select the Configuration tab.
4. Select IdP Client.
5. Populate IdP Metadata field through one of the following:
 - a. An HTTPS URL (<https://>)
 - b. A file URL (file:)
 - c. An XML block to refer to desired metadata, such as:
 - i. e.g., <https://localhost:8993/services/idp/login/metadata>

IdP Client (SP) example.xml

```
<md:EntityDescriptor xmlns:md="urn:oasis:names:tc:SAML:2.0:metadata" entityID="https://localhost:8993/services/idp/login">
  <md:IDPSSODescriptor WantAuthnRequestsSigned="true" protocolSupportEnumeration="urn:oasis:names:tc:SAML:2.0:protocol">
    <md:KeyDescriptor use="signing">
      <ds:KeyInfo xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
        <ds:X509Data>
          <ds:X509Certificate>
MIIDEZCCAnygAwIBAgIJIAIzc4FYrIp9mMA0GCSqGSIb3DQEBCUAMCxszAJBgNVBAYTA1VTMQswCQYDVQQIDA
JBWjEMMAoGA1UECgwDRERGMQwwCgYDVQQLDANEZXYxGTAXBgNVBAMMEERERiBEZW1vIFJvb3QgQ0ExJDAiBkgq
hkiG9w0BCQEWFWRkZnJvb3RjYUB1eGFtcGx1Lm9yZzAeFw0xNDEyMTAyMTU4MThaFw0xNTEyMTAyMTU4MThaMI
GDMQswCQYDVQQGEwJVUzELMAkGA1UECAwCQVoxETAPBgNVBAcMCEdvb2R5ZWFrMQwwCgYDVQQKDANEREYxDDAK
BgNVBAsMA0RldjESMBAGA1UEAwvJbG9jYWxob3N0MSQwIgYJKoZIhvcNAQkBFhVsb2NhbGhvc3RAZXhhbXBsZS
5vcmcwgZ8wDQYJKoZIhvcNAQEBBQADgY0AMIGJAoGBAMeCyNZbCTZphHQfb5g8FrgBq1RYzV7ikVw/pVGkz8gx
313A99s8WtA4mRAeb6n0vTR9yNB0ekW4nY0iE0q//YTi/frI1kz0QbEH1s2cI5nFButabD3PYGxUSapbc+AS7
+Pk1r0TDI4MRzPPkkTp4w10RQ/a6CfVsNr/mVgL2CfAgMBAAGjgZkwgZYwCQYDVR0TBAIwADAnBglghkgBvhvC
AQ0EGhYYRk9SIFRFU1RJTkgUFVSUE9TRSBPTkxZMB0GA1UdDgQWBBSA95QIMyBAHRsd0R4s7C3BreFrDAfBg
NVHSMEGDAWgBThVMwX3wrCv61feF47CyvkSBe9xjAgBgNVHREEGTAXgRVsb2NhbGhvc3RAZXhhbXBsZS5vcmcw
DQYJKoZIhvcNAQEFBQADgYEAtRUp7fAxU/E6JD2Kj/+CTWqu8Elx13S0TxoIqv3gMoBW0ehyzEKjJi0bb1gUx0
7n1Sm0ESp5sE3jGTnh0GtYV0D219z/09n90cd/imAEhknJlayyd0SjpnaL9JUd8uYxJexy8TJ2sMhsGAZ6EMTz
CfT9m07XduxjsmDz0h1SGV0=
        </ds:X509Certificate>
      </ds:X509Data>
    </ds:KeyInfo>
  </md:KeyDescriptor>
  <md:KeyDescriptor use="encryption">
    <ds:KeyInfo xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
      <ds:X509Data>
```

```

<ds:X509Certificate>
MIIDEZCCAnygAwIBAgIJIAIzc4FYrIp9mMA0GCSqGSIB3DQEBBQUAMhcxCzAJBgNVBAYTA1VTMQswCQYDVQQIDA
JBWjEMMAoGA1UECgwDRERGMQwwCgYDVQQLDANEZXYxGTAXBgNVBAMMEERERiBEZW1vIFJvb3QgQ0ExJDAiBkgq
hkiG9w0BCQEWFRkZnJvb3RjYUB1eGFtcGx1Lm9yZzAeFw0xNDEyMTAyMTU4MThaFw0xNTEyMTAyMTU4MThaMI
GDMQswCQYDVQQGEwJVUzELMAkGA1UECAwCQVoxETAPBgNVBAcMCEdvb2R5ZWFrMQwwCgYDVQQKDANEREYxDDAK
BgNVBAsMA0RldjESMBAGA1UEAwjB9jYWxob3N0MSQwIgYJKoZIhvcNAQkBFhVsb2NhbGhvc3RAZXhhXBsZS
5vcmcwgZ8wDQYJKoZIhvcNAQEBBQADgY0AMIGJAoGBAMeCyNZbCTphHQfB5g8FrgBq1RYzV7ikVw/pVGkz8gx
313A99s8WtA4mRAeb6n0vTR9yNBOekW4nY0iE0q//YTi/frI1kz0QbEH1s2cI5nFButabD3PYGxUSuapbc+AS7
+Pk1r0TDI4MRzPPkkTp4w10RQ/a6CfVsNr/mVgL2CfAgMBAAGjgZkwgZYwCQYDVR0TBAIwADAnBglghkgBvhvC
AQ0EGhYYRk9SIFRFU1RJTkcgUFVSUE9TRSBPTkxZMB0GA1UdDgQWBBSA95QIMyBAHRsd0R4s7C3BreFrsDAfBg
NVHSMEGDAwgbThVMex3wrCv61feF47CvkSBe9xjAgBgnVHREEGTAXgRVsb2NhbGhvc3RAZXhhXBsZS5vcmcw
DQYJKoZIhvcNAQEFBQADgYEAtRUp7fAxU/E6JD2Kj/+CTWqu8Elx13S0TxoIqv3gMoBW0ehyzEKjJi0bb1gUx0
7n1Sm0ESp5sE3jGTnh0GtYV0D219z/09n90cd/imAEhknJlayyd0Sjpnal9JUd8uYxJexy8TJ2sMhsGAZ6EMTz
CfT9m07XduxjsmDz0h1SGV0=
</ds:X509Certificate>
</ds:X509Data>
</ds:KeyInfo>
</md:KeyDescriptor>
<md:SingleLogoutService Binding="urn:oasis:names:tc:SAML:2.0:bindings:HTTP-Redirect" Location="https://localhost:8993/logout"/>
<md:SingleLogoutService Binding="urn:oasis:names:tc:SAML:2.0:bindings:HTTP-POST" Location="https://localhost:8993/logout"/>
<md:NameIDFormat>
  urn:oasis:names:tc:SAML:2.0:nameid-format:persistent
</md:NameIDFormat>
<md:NameIDFormat>
  urn:oasis:names:tc:SAML:1.1:nameid-format:unspecified
</md:NameIDFormat>
<md:NameIDFormat>
  urn:oasis:names:tc:SAML:1.1:nameid-format:X509SubjectName
</md:NameIDFormat>
<md:SingleSignOnService Binding="urn:oasis:names:tc:SAML:2.0:bindings:HTTP-Redirect" Location="https://localhost:8993/services/idp/login"/>
<md:SingleSignOnService Binding="urn:oasis:names:tc:SAML:2.0:bindings:HTTP-POST" Location="https://localhost:8993/services/idp/login"/>
</md:IDPSSODescriptor>
</md:EntityDescriptor>

```

When using the included IdP, DDF can be configured to use the [included Security Token Service \(STS\)](#) or an [external STS](#).

5.5.2.1.1. Configuring Included STS

An LDAP server can be used to maintain a list of DDF users and the attributes associated with them. The [Security Token Service \(STS\)](#) can use an LDAP server as an attribute store and convert those attributes to SAML claims.

DDF includes a demo [LDAP server](#), but an external LDAP is required for actual installation.

The STS is installed by default in DDF.

Configuring STS

1. Verify that the `serverKeystores.jks` file in `<DDF_HOME>/etc/keystores` trusts the hostnames used in your environment (the hostnames of LDAP, and any DDF users that make use of this STS server).
2. Navigate to the **Admin Console**.
3. Select the **Security** application.
4. Select the **Features** tab.
5. Start the `security-sts-ldaplogin` and `security-sts-ldapclaimshandler` features.
6. Select the **Configuration** tab.
7. Select the **Security STS LDAP Login** configuration.
8. Verify that the **LDAP URL**, **LDAP Bind User DN**, and **LDAP Bind User Password** fields match your LDAP server's information.
 - a. The default DDF LDAP settings will match up with the default settings of the OpenDJ embedded LDAP server. Change these values to map to the location and settings of the LDAP server being used.
9. Select the **Save changes** button if changes were made.
10. Open the **Security STS LDAP and Roles Claims Handler** configuration.
11. Populate the same URL, user, and password fields with your LDAP server information.
12. Select the **Save Changes** button.

Configuring DDF Authentication Scheme

Configure the DDF to use this authentication scheme.

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Open the **Web Context Policy Manager** configuration.
 - a. Under **Context Realms** add the contexts that should be protected under the ldap realm.
 - i. The default setting is `/=karaf`, the `karaf` realm refers to the `user.properties` user store file located in the `<DDF_HOME>/etc` directory. This can be changed to `/=ldap`, if it is desired that the entire container be protected under ldap. If the `/admin` context is changed to something other than the default (`karaf`), it will be required that you refresh the page in order to log in again, or your changes may not be saved. This includes changing the root context to something other than `karaf`, without specifically setting `/admin` to a realm. The policies for all contexts will roll up, for example: the `/admin` context policy will roll up to the `karaf` realm with the default configuration because `/` is higher in the context hierarchy than `/admin` and no realm is specifically set for `/admin`.
 - b. Under **Authentication Types**, make any desired authentication changes to contexts.
 - i. In order to use the SAML 2.0 Web SSO profile against a context, you must specify only the IdP authentication type.

Security STS Client

Configure the client connecting to the STS.

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Open the **Security STS Client** configuration.
4. Verify that the host/port information in the **STS Address** field points to your STS server. If you are using the default bundled STS, this information will already be correct.

See [Security STS Client](#) table for all configuration options.

The DDF should now use the SSO/STS/LDAP servers when it attempts to authenticate a user upon an attempted log in.

STS Server

Connect to the server hosting the STS.

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Select the **Security STS Server** configuration.
4. Verify the hostname and usernames are correct.

See [Security STS Server](#) table for all configuration options.

SAML Name ID

Set up alternatives to displaying the username of the logged in user.

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Select the **SAML NameID Policy** configuration.
4. Add any desired attributes to display instead of the username. (The first matching attribute will be used.)

Limiting Access to the STS

Be sure to limit the hosts that are allowed to connect to the STS:

- **Required Step for Security Hardening**
- Open the `<DDF_HOME>/etc/system.properties` file.
- Edit the line `ws-security.subject.cert.constraints = .*`.
 - Remove the `.*` and replace with a comma-delimited list of regular expressions for the DNs of all desired hosts (`<MY_HOST_DN>`):
 - `ws-security.subject.cert.constraints = <MY_HOST_DN>,<OTHER_HOST_DN>`

5.5.2.1.2. Connecting to External STS

Configure DDF to connect to an external WSS STS.

Security STS Address Provider

Configure the STS address provider to use WSS.

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Select **Configuration**.
4. Select the **Security STS Address Provider**.
5. Enable the option **Use WSS STS**.

Security STS WSS

Configure the location and credentials for the STS.

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Select **Configuration**.
4. Select the **Security STS WSS** configuration.
5. Update the **Address**, **Endpoint Name**, and **Service Name** properties.

Disable Security STS Client Configuration

Disable the client configuration for the Security STS

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Select **Features**.
4. Uninstall the **Security STS Client** feature.

5.5.2.2. Connecting to an External IdP

Service Provider Metadata

See [Configuring Security IdP Service Provider](#) for details on configuring the IdP client.

Disable Included IdP

Remove the included IdP.

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Select **Features**.
4. Uninstall **security-idp** feature.

5.5.2.3. Configuring Without an IdP

To configure DDF to not use an Identity Provider (IdP), simply disable the IdP feature.

1. Navigate to the **{admin-console}**.
 2. Select the **Security** application.
 3. Select **Features**.
 4. Uninstall the **security-idp** feature.
-

5.5.2.3.1. Using STS without IdP

To configure DDF to use the included Security Token Service (STS) without an IdP, follow the same [Configuring STS](#) steps, with one additional configuration to make via the [Web Context Policy Manager](#).

Configuring Authentication Types for STS

1. Navigate to the **{admin-console}**.
 2. Select the **Security** application.
 3. Select **Configuration**.
 4. Select the **Web Context Policy Manager**.
 5. Add any needed authentication types to the Authentication Types list, such as **PKI**, **Basic**, etc.
-

5.5.2.3.2. Connecting to External STS Without IdP

The process for [connecting to an external STS](#) is the same with or without an IdP.

5.5.3. Configuring SOAP Services for Users

If using SOAP services, DDF can be configured to use the included [Security Token Service \(STS\)](#), or [connected to an external STS](#).

5.5.3.1. Connecting to Included STS with SOAP

DDF includes a STS implementation that can be used for user authentication over SOAP services.

Configure the STS WSS

Configure the STS WSS.

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Select **Configuration**.
4. Select **Security STS WSS**.

5. Update the **Claims** that should be requested by the STS.
-

5.5.3.2. Connecting to External STS with SOAP

See [connecting to an external STS](#) for initial STS setup.

5.5.4. Connecting to an LDAP Server

WARNING The configurations for Security STS LDAP and Roles Claims Handler and Security STS LDAP Login contain plain text default passwords for the embedded LDAP, which is insecure to use in production.

Use the [Encryption Service](#), from the Command Console to set passwords for your LDAP server. Then change the LDAP Bind User Password in the [Security STS LDAP and Roles Claims Handler](#) configurations to use the encrypted password.

A claim is an additional piece of data about a principal that can be included in a token along with basic token data. A claims manager provides hooks for a developer to plug in claims handlers to ensure that the STS includes the specified claims in the issued token.

Claims handlers convert incoming user credentials into a set of attribute claims that will be populated in the SAML assertion. For example, the [LDAPClaimsHandler](#) takes in the user's credentials and retrieves the user's attributes from a backend LDAP server. These attributes are then mapped and added to the SAML assertion being created. Integrators and developers can add more claims handlers that can handle other types of external services that store user attributes.

See the [Security STS LDAP and Roles Claims Handler](#) for all possible configurations.

5.5.5. Updating System Users

By default, all system users are located in the `<DDF_HOME>/etc/users.properties` and `<DDF_HOME>/etc/users.attributes` files. The default users included in these two files are "admin" and "localhost". The `users.properties` file contains username, password, and role information; while the `users.attributes` file is used to mix in additional attributes. The `users.properties` file must also contain the user corresponding to the fully qualified domain name (FQDN) of the system where DDF is running. This FQDN user represents this host system internally when making decisions about what operations the system is capable of performing. For example, when performing a DDF Catalog Ingest, the system's attributes will be checked against any security attributes present on the metocard, prior to ingest, to determine whether or not the system should be allowed to ingest that metocard.

Additionally, the `users.attributes` file can contain user entries in a regex format. This allows an administrator to mix in attributes for external systems that match a particular regex pattern. The FQDN user within the `users.attributes` file should be filled out with attributes sufficient to allow the system to ingest the expected data. The `users.attributes` file uses a JSON format as shown below:

```
{
  "admin" : {
    "test" : "testValue",
    "test1" : [ "testing1", "testing2", "testing3" ]
  },
  "localhost" : {

  },
  ".*host.*" : {
    "reg" : "ex"
  }
}
```

For this example, the "admin" user will end up with two additional claims of "test" and "test1" with values of "testValue" and ["testing1", "testing2", "testing3"] respectively. Also, any host matching the regex ".host." would end up with the claim "reg" with the single value of "ex". The "localhost" user would have no additional attributes mixed in.

WARNING It is possible for a regex in `users.attributes` to match users as well as a system, so verify that the regex pattern's scope will not be too great when using this feature.

WARNING If your data will contain security markings, and these markings are being parsed out into the METACARD.security attribute via a PolicyPlugin, then the FQDN user **MUST** be updated with attributes that would grant the privileges to ingest that data. Failure to update the FQDN user with sufficient attributes will result in an error being returned for any ingest request.

5.5.6. Restricting Access to Admin Console

- **Required Step for Security Hardening**

If you have integrated DDF with your existing security infrastructure, then you may want to limit access to parts of the DDF based on user roles/groups.

Limit access to the Admin Console to those users who need access. To set access restrictions on the Admin Console, consult the organization's security architecture to identify specific realms, authentication methods, and roles required.

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Select the **Configuration** tab.
4. Select the **Web Context Policy Manager**.
 - a. A dialogue will pop up that allows you to edit DDF access restrictions.
 - b. Once you have configured your `realms` in your security infrastructure, you can associate them with DDF contexts.

- c. If your infrastructure supports multiple **authentication methods**, they may be specified on a per-context basis.
- d. Role requirements may be enforced by configuring the **required attributes** for a given context.
- e. The **white listed contexts** allows child contexts to be excluded from the authentication constraints of their parents.

5.5.6.1. Restricting Feature, App, Service, and Configuration Access

- **Required Step for Security Hardening**

Limit access to the individual applications, features, or services to those users who need access. Organizational requirements should dictate which applications are restricted and the extent to which they are restricted.

1. Navigate to the **Admin Console**.
2. Select the **Admin** application.
3. Select the **Configuration** tab.
4. Select the **Admin Configuration Policy**.
5. To add a feature or app permission:

- a. Add a new field to "Feature and App Permissions" in the format of:

```
<feature name>/<app name> = "attribute name=attribute value","attribute name2=attribute value2", ...
```

- b. For example, to restrict access of any user without an admin role to the catalog-app:

```
catalog-app = "http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=admin", ...
```

6. To add a configuration permission:

- a. Add a new field to "Configuration Permissions" in the format of:

```
configuration id = "attribute name=attribute value","attribute name2=attribute value2", ...
```

- b. For example, to restrict access of any user without an admin role to the Web Context Policy Manager:

```
org.codice.ddf.security.policy.context.impl.PolicyManager="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=admin"
```

If a permission is specified, any user without the required attributes will be unable to see or modify the feature, app, or configuration.

5.5.7. Removing Default Users

- **Required Step for Security Hardening**

The default security configuration uses a property file located at `<DDF_HOME>/etc/users.properties` to

store users and passwords. A hardened system will remove this file and manage all users externally, via an LDAP server or by other means.

Default Users are an Insecure Default

NOTE The Admin Console has an insecure default warning if the default users are not removed.

Once DDF is configured to use an external user (such as LDAP), remove the `users.properties` file from the `<DDF_HOME>/etc` directory. Use of a `users.properties` file should be limited to emergency recovery operations and replaced as soon as effectively possible.

The deletion of the default users in the `users.properties` file can be done automatically after 72 hours. This feature can be found at **Admin Console** → **Admin** → **Default Users Deletion Scheduler** → **Enable default users automatic deletion**.

If SSH access to the Karaf shell is to be supported, edit the file `org.apache.karaf.shell.cfg` in the `<INSTALL_HOME>/etc` directory, changing the value of the `sshRealm` property from `karaf` to `ldap`.

Emergency Use of `users.properties` file

Typically, the DDF does not manage passwords. Authenticators are stored in an external identity management solution. However, administrators may temporarily use a `users.properties` file for emergencies.

If a system recovery account is configured in `users.properties`, ensure:

NOTE

- The use of this account should be for as short a time as possible.
- The default username/password of “admin/admin” should not be used.
- All organizational standards for password complexity should still apply.
- The password should be encrypted. For steps on how, see the section “Passwords Encryption” at <https://karaf.apache.org/manual/latest/security>.

Compliance Reviews

NOTE It is recommended to perform yearly reviews of accounts for compliance with organizational account management requirements.

5.6. Configuring Data Management

Data ingested into DDF has security attributes that can be mapped to users' permissions to ensure proper access. This section covers configurations that ensure only the appropriate data is contained in or exposed by DDF.

5.6.1. Configuring Solr

The default catalog provider for DDF is [Solr](#).

5.6.1.1. Configuring Solr Catalog Provider Data Directory

The HTTP Solr Catalog Provider and the Embedded Solr Catalog Provider writes index files to the file system. By default, these files are stored under `<DDF_HOME>/data/solr/catalog/data`. If there is inadequate space in `<DDF_HOME>`, or if it is desired to maintain backups of the indexes only, this directory can be changed.

In order to change the Data Directory, the `system.properties` file in `<DDF_HOME>/etc` must be edited prior to starting DDF.

5.6.1.1.1. Changing an Empty Data Directory

Edit the `system.properties` file

```
# Uncomment the following line and set it to the desired path
# solr.data.dir = ${karaf.home}/data/solr
```

5.6.1.1.2. Changing the Data Directory After Ingest

It may become necessary to change the data directory after DDF has ingested data.

1. Shut down the DDF.
2. Create the new directory to hold the indexes.

Make new Data Directory

```
mkdir -p /path/to/new/data/dir
```

3. Copy the indexes to the new directory.

Copy the indexes to the new Directory.

```
cp /path/to/old/data/dir/* /path/to/new/data/dir/.
```

4. Set the `system.properties` file to use the new directory.

Update `system.properties` file

```
solr.data.dir = /path/to/new/data/dir
```

5. Restart the DDF.

Changes Require a Distribution Restart

WARNING

If the Data Directory File Path property is changed, no changes will occur to the Solr Catalog Provider until the distribution has been restarted.

NOTE

If data directory file path property is changed to a new directory, and the previous data is not moved into that directory, no data will exist in Solr. Instead, Solr will create an empty index. Therefore, it is possible to have multiple places where Solr files are stored, and a user can toggle between those locations for different sets of data.

5.6.1.2. Hardening Solr Index

- **Required Step for Security Hardening**

The DDF's design includes support for pluggable indexes. The default installation contains a Solr index to be used as the Metadata Catalog, but this implementation can be replaced with an alternate 3rd party index implementation.

The following sections provide hardening guidance for Solr; however, they are to serve only as reference and other an additional security requirements may be added.

The Solr Admin user interface uses basic authentication as part of the server configuration.

5.6.1.2.1. Configuring Solr Node Security

The Solr server is protected by the built in REST security architecture. The configuration can be changed by editing the [Web Context Policy Manager](#) configuration for the `/solr` web context.

1. Navigate to the [Admin Console](#).
2. Select the **Security** application.
3. Select the **Configuration** tab.
4. Select **Web Context Policy Manager**.

By default, the configuration is set to `/solr=SAML|PKI|BASIC`. This allows a user or another system to connect to Solr using any of those authentication methods.

5.6.1.2.2. Configuring Solr Encryption

While it is possible to encrypt the Solr index, it decreases performance significantly. An encrypted Solr index also can only perform exact match queries, not relative or contextual queries. As this drastically reduces the usefulness of the index, this configuration is not recommended. The recommended approach is to encrypt the entire drive through the Operating System of the server on which the index is located.

5.6.2. Changing Catalog Providers

This scenario describes how to reconfigure DDF to use a different catalog provider.

This scenario assumes DDF is already running.

Uninstall Catalog Provider (if installed).

1. Navigate to the **Admin Console**.
2. Select the **Solr Catalog** application.
3. Select the **Features** tab.
4. Find and Stop the installed Catalog Provider

Install the new Catalog Provider

1. Navigate to the **Admin Console**.
2. Select the **Solr Catalog** application.
3. Select the **Features** tab.
4. Find and Start the desired Catalog Provider.

5.6.3. Changing Hostname

By default Solr, STS server, STS client and the rest of the services use the system property `org.codice.ddf.system.hostname` which is defaulted to 'localhost' and not to the fully qualified domain name of the DDF instance. Assuming the DDF instance is providing these services, the configuration must be updated to use the **fully qualified domain name** as the service provider.

This can be changed during **Initial Configuration** or later by editing the `<DDF_HOME>/etc/system.properties` file.

5.6.4. Content Directory Monitor

The Content Directory Monitor (CDM) provides the capability to easily add content and metacards into the Catalog by placing a file in a directory.

5.6.4.1. Installing the Content Directory Monitor

The Content Directory Monitor is installed by default with a standard installation of the Catalog application.

5.6.4.2. Configuring the Content Directory Monitor

Configure the CDM from the Admin Console:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Configuration** tab.
4. Select **Catalog Content Directory Monitor**.

Content Directory Monitor Permissions

IMPORTANT

When setting monitored directory, make sure to set permissions on the new directory to allow DDF to access it.

See [Content Directory Monitor configurations](#) for all possible configurations.

5.6.4.3. Using the Content Directory Monitor

The CDM processes files in a directory, and all of its sub-directories. The CDM offers three options:

- Delete
- Move
- Monitor in place

Regardless of the option, the DDF takes each file in a monitored directory structure and creates a metocard for it. The metocard is linked to the file. The behavior of each option is given below.

Delete

- Copies the file into the Content Repository.
- Creates a metocard in the Catalog from the file.
- **Erases** the original file from the monitored directory.

Move

- Copies the file into the directory **.\ingested** (**this will double the disk space used**)
- Copies the file into the Content Repository.
- Creates a metocard in the Catalog from the file.
- **Erases** the original file from the monitored directory.

Monitor in place

- Creates a metocard in the Catalog from the file.
- Creates a reference from the metocard to the original file in the monitored directory.
- If the original file is deleted, the metocard is removed from the Catalog.
- If the original file is modified, the metocard is updated to reflect the new content.
- If the original file is renamed, the old metocard is deleted and a new metocard is created.

Parallel Processing

The CDM supports parallel processing of files (up to 8 files processed concurrently). This is configured by setting the number of **Maximum Concurrent Files** in the configuration. A maximum of 8 is imposed to protect system resources.

Read Lock

When the CDM is set up, the directory specified is continuously scanned, and files are locked for processing based on the **ReadLock Time Interval**. This does not apply to the **Monitor in place** processing directive. Files will not be ingested without having a ReadLock that has observed no change in the file size. This is done so that files that are in transit will not be ingested prematurely. The interval should be dependent on the speed of the copy to the directory monitor (ex. network drive vs local disk). For local files, the default value of 500 milliseconds is recommended. The recommended interval for network drives is 1000 - 2000 milliseconds. If the value provided is less

than 100, 100 milliseconds will be used. It is also recommended that the **ReadLock Time Interval** be set to a lower amount of time when the **Maximum Concurrent Files** is set above 1 so that files are locked in a timely manner and processed as soon as possible. When a higher **ReadLock Time Interval** is set, the time it takes for files to be processed is increased.

Attribute Mapper

The CDM supports setting metocard attributes directly when DDF ingests a file. Custom mappings are entered in the form:

attribute-name=attribute-value

For example, to set the contact email for all metacards, add the custom mapping:

contact.point-of-contact-email=doctor@clinic.com

Each mapping sets the value of a single metocard attribute. To set the value of an additional attribute, select the "plus" icon in the UI. This creates an empty line for the entry.

To set multi-valued attributes, use a separate override for each value. For example, to add the keywords *PPI* and *radiology* to each metocard, add the custom attribute mappings:

topic.keyword=PPI
topic.keyword=radiology

Blacklist

The CDM blacklist uses the "bad.files" and "bad.file.extensions" properties from the system.properties file in "etc/" in order to prevent malicious or unwanted data from being ingested into DDF. While the CDM automatically omits hidden files, this is particularly useful when an operating system automatically generates files that should not be ingested. One such example of this is "thumbs.db" in Windows. This file type and any temporary files are included in the blacklist.

Errors

If the CDM fails to read the file, an error will be logged in the ingest log. If the directory monitor is configured to **Delete** or **Move**, the original file is also moved to the **\.errors** directory.

Other

- Multiple directories can be monitored. Each directory has an independent configuration.
- To support the monitoring in place behavior, DDF indexes the files to track their names and modification timestamps. This enables the Content Directory Monitor to take appropriate action when files are changed or deleted.
- The Content Directory Monitor recursively processes all subdirectories.

5.7. Configuring Security Policies

User attributes and Data attributes are matched by security policies defined within DDF.

5.7.1. Configuring the Web Context Policy Manager

The Web Context Policy Manager defines all security policies for REST endpoints within DDF. It

defines:

- the realms a context should authenticate against.
- the type of authentication that a context requires.
- any user attributes required for authorization.

See [Web Context Policy Manager Configurations](#) for detailed descriptions of all fields.

5.7.1.1. Context Realms

The karaf realm is the only realm available by default and it authenticates against the `users.properties` file. As JAAS authentication realms are added to the STS, more realms become available to authenticate against.

For example, installing the `security-sts-ldaplogin` feature adds an ldap realm. Contexts can then be pointed to the ldap realm for authentication and STS will be instructed to authenticate them against ldap.

5.7.1.2. Authentication Types

As you add REST endpoints, you may need to add different types of authentication through the Web Context Policy Manager.

Any web context that allows or requires specific authentication types should be added here with the following format:

```
/<CONTEXT>=<AUTH_TYPE>|<AUTH_TYPE>|...
```

Table 8. Default Types of Authentication

Authentication Type	Description
<code>saml</code>	Activates single-sign on (SSO) across all REST endpoints that use SAML.
<code>basic</code>	Activates basic authentication.
<code>PKI</code>	Activates public key infrastructure authentication.
<code>IdP</code>	Activates SAML Web SSO authentication support. Additional configuration is necessary.
<code>CAS</code>	Enables SSO through a Central Authentication Server
<code>guest</code>	provides guest access

5.7.1.3. Required Attributes

The fields for required attributes allows configuring certain contexts to only be accessible to users with pre-defined attributes. For example, the default required attribute for the `/admin` context is `role=system-admin`, limiting access to the Admin Console to system administrators

5.7.1.4. White Listed Contexts

White listed contexts are trusted contexts which will bypass security. Any sub-contexts of a white listed context will be white listed as well, unless they are specifically assigned a policy.

5.7.2. Configuring Catalog Filtering Policies

Filtering is the process of evaluating security markings on data products, comparing them to the users permissions and protecting resources from inappropriate access.

There are two options for processing filtering policies: internally, or through the use of a policy formatted in eXtensible Access Control Markup Language (XACML). The procedure for setting up a policy differs depending on whether that policy is to be used internally or by the external XACML processing engine.

5.7.2.1. Setting Internal Policies

1. Navigate to the **Admin Console**.
2. Select the **Security** application.
3. Click the **Configuration** tab.
4. Click on the **Security AuthZ Realm** configuration.
5. Add any attribute mappings necessary to map between subject attributes and the attributes to be asserted.
 - a. For example, the above example would require two Match All mappings of `subjectAttribute1=assertedAttribute1` and `subjectAttribute2=assertedAttribute2`.
 - b. Match One mappings would contain `subjectAttribute3=assertedAttribute3` and `subjectAttribute4=assertedAttribute4`.

With the `security-pdp-authz` feature configured in this way, the above Metocard would be displayed to the user. Note that this particular configuration would not require any XACML rules to be present. All of the attributes can be matched internally and there is no reason to call out to the external XACML processing engine. For more complex decisions, it might be necessary to write a XACML policy to handle certain attributes.

5.7.2.2. Setting XACML Policies

To set up a XACML policy, place the desired XACML policy in the `<distribution root>/etc/pdp/policies` directory and update the included `access-policy.xml` to include the new policy. This is the directory in which the PDP will look for XACML policies every 60 seconds.

See [Developing XACML Policies](#) for more information about custom XACML policies.

5.7.2.3. Catalog Filter Policy Plugins

Several Policy Plugins for catalog filtering exist currently: [Metocard Attribute Security Policy Plugin](#) and [XML Attribute Security Policy Plugin](#). These Policy Plugin implementations allow an administrator to easily add filtering capabilities to some standard Metocard types for all Catalog operations. These plugins will place policy information on the Metocard itself that allows the [Filter Plugin](#) to restrict unauthorized users from viewing content they are not allowed to view.

5.8. Configuring Federation

DDF is able to [federate](#) to other data sources, including other instances of DDF, with some simple configuration.

5.8.1. Enable SSL for Clients

In order for outbound secure connections (HTTPS) to be made from components like Federated Sources and Resource Readers configuration may need to be updated with keystores and security properties. These values are configured in the `<DDF_HOME>/etc/system.properties` file. The following values can be set:

Property	Sample Value	Description
<code>javax.net.ssl.trustStore</code>	<code>etc/keystores/serverTruststore.jks</code>	The java keystore that contains the trusted public certificates for Certificate Authorities (CA's) that can be used to validate SSL Connections for outbound TLS/SSL connections (e.g. HTTPS). When making outbound secure connections a handshake will be done with the remote secure server and the CA that is in the signing chain for the remote server's certificate must be present in the trust store for the secure connection to be successful.
<code>javax.net.ssl.trustStorePassword</code>	<code>changeit</code>	This is the password for the truststore listed in the above property
<code>javax.net.ssl.keyStore</code>	<code>etc/keystores/serverKeystore.jks</code>	The keystore that contains the private key for the local server that can be used for signing, encryption, and SSL/TLS.
<code>javax.net.ssl.keyStorePassword</code>	<code>changeit</code>	The password for the keystore listed above
<code>javax.net.ssl.keyStoreType</code>	<code>jks</code>	The type of keystore

Property	Sample Value	Description
<code>https.cipherSuites</code>	<code>TLS_DHE_RSA_WITH_AES_128_GCM_SHA256, TLS_DHE_RSA_WITH_AES_128_CBC_SHA256, TLS_DHE_RSA_WITH_AES_128_CBC_SHA, TLS_ECDHE_ECDSA_WITH_AES_128_GCM_SHA256, TLS_ECDHE_RSA_WITH_AES_128_GCM_SHA256</code>	The cipher suites that are supported when making outbound HTTPS connections
<code>https.protocols</code>	<code>TLSv1.1,TLSv1.2</code>	The protocols that are supported when making outbound HTTPS connections
<code>jdk.tls.client.protocols</code>	<code>TLSv1.1,TLSv1.2</code>	The protocols that are supported when making inbound HTTPS connections
<code>jdk.tls.ephemeralDHKeySize</code>	'matched'	For X.509 certificate based authentication (of non-exportable cipher suites), the DH key size matching the corresponding authentication key is used, except that the size must be between 1024 bits and 2048 bits. For example, if the public key size of an authentication certificate is 2048 bits, then the ephemeral DH key size should be 2048 bits unless the cipher suite is exportable. This key sizing scheme keeps the cryptographic strength consistent between authentication keys and key-exchange keys.

NOTE

`<DDF_HOME> Directory`

DDF is installed in the `<DDF_HOME>` directory.

5.8.2. Configuring HTTP Port

IMPORTANT

Do not use the Admin Console to change the HTTP port. While the Admin Console's Pax Web Runtime offers this configuration option, it has proven to be unreliable and may crash the system. Use the Command Console instead.

5.8.3. Configuring HTTP Proxy

The `platform-http-proxy` feature proxies https to http for clients that cannot use HTTPS and should not have HTTP enabled for the entire container via the `etc/org.ops4j.pax.web.cfg` file.

Enabling the HTTP Proxy from the Admin Console

1. Navigate to the **Admin Console**.
2. Select the **Platform** application tile.
3. Select the **Features** tab.
4. Select `platform-http-proxy`.
5. Select the **Play** button to the right of the word “Uninstalled”

Enabling the HTTP Proxy from the Command Console

- Type the command `feature:install platform-http-proxy`

Configuring HTTP Proxy Hostname

1. Select **Configuration** tab.
2. Select **HTTP to HTTPS Proxy Settings**
 - a. Enter the Hostname to use for HTTPS connection in the proxy.
3. Click **Save changes**.

HTTP Proxy and Hostname

NOTE The hostname should be set by default. Only configure the proxy if this is not working.

5.8.4. Federation Strategy

A federation strategy federates a query to all of the Remote Sources in the query's list, processes the results in a unique way, and then returns the results to the client. For example, implementations can choose to halt processing until all results return and then perform a mass sort or return the results back to the client as soon as they are received back from a Federated Source.

An endpoint can optionally specify the federation strategy to use when it invokes the query operation. Otherwise, the Catalog provides a default federation strategy that will be used: the Catalog Federation Strategy.

5.8.4.1. Configuring Federation Strategy

The Catalog Federation Strategy configuration can be found in the Admin Console.

1. Navigate to Admin Console.
2. Select **Catalog**
3. Select **Configuration**
4. Select **Catalog Federation Strategy**.

See [Federation Strategy configurations](#) for all possible configurations.

5.8.4.1.1. Catalog Federation Strategy

The Catalog Federation Strategy is the default federation strategy and is based on sorting metacards by the sorting parameter specified in the federated query.

The possible sorting values are:

- metacard's effective date/time
- temporal data in the query result
- distance data in the query result
- relevance of the query result

The supported sorting orders are ascending and descending.

The default sorting value/order automatically used is relevance descending.

WARNING

The Catalog Federation Strategy expects the results returned from the Source to be sorted based on whatever sorting criteria were specified. If a metadata record in the query results contains null values for the sorting criteria elements, the Catalog Federation Strategy expects that result to come at the end of the result list.

5.8.5. Connecting to Sources

A **source** is a system consisting of a catalog containing Metacards.

Catalog sources are used to connect Catalog components to data sources, local and remote. Sources act as proxies to the actual external data sources, e.g., a RDBMS database or a NoSQL database.

Types of Sources

Remote Source

Read-only data sources that support query operations but cannot be used to create, update, or delete metacards.

Federated Sources

A federated source is a remote source that can be included in federated queries by request or as part of an enterprise query. Federated sources support query and site information operations only. Catalog modification operations, such as create, update, and delete, are not allowed. Federated sources also expose an event service, which allows the Catalog Framework to subscribe to event notifications when metacards are created, updated, and deleted.

Catalog instances can also be federated to each other. Therefore, a Catalog can also act as a federated source to another Catalog.

Connected Sources

A Connected Source is a local or remote source that is always included in every local and enterprise query, but is hidden from being queried individually. A connected source's identifier is removed in all query results by replacing it with DDF's source identifier. The Catalog Framework does not reveal a connected source as a separate source when returning source information responses.

Catalog Providers

A Catalog Provider is used to interact with data providers, such as files systems or databases, to query, create, update, or delete data. The provider also translates between DDF objects and native data formats.

All sources, including federated source and connected source, support queries, but a Catalog provider also allows metacards to be created, updated, and deleted. A Catalog provider typically connects to an external application or a storage system (e.g., a database), acting as a proxy for all catalog operations.

Catalog Stores

A Catalog Store is an editable store that is either local or remote.

Available Federated Sources

The following Federated Sources are available in a standard installation of DDF:

Federated Source for Atlassian Confluence®

Retrieve pages, comments, and attachments from an Atlassian Confluence® REST API.

CSW Federated Source

Queries a CSW version 2.0.2 compliant service.

CSW Federation Profile Source

Queries a CSW version 2.0.2 compliant service.

GMD CSW Source

Queries a GMD CSW APISO compliant service.

OpenSearch Source

Performs OpenSearch queries for metadata.

WFS 1.0 Source

Allows for requests for geographical features across the web.

WFS 2.0 Source

Allows for requests for geographical features across the web.

Available Connected Sources

The following Connected Sources are available in a standard installation of DDF:

CSW Connected Source

Searches collections of descriptive information (metadata) for data, services, and related information objects.

WFS 1.0 Source

Allows for requests for geographical features across the web.

WFS 2.0 Source

Allows for requests for geographical features across the web.

Available Catalog Stores

The following Catalog Stores are available in a standard installation of DDF:

Registry Store

Allows CSW messages to be turned into usable Registry metacards and for those metacards to be turned back into CSW messages.

Available Catalog Providers

The following Catalog Providers are available in a standard installation of DDF:

Solr Catalog Provider

Uses Solr as a catalog.

Available Storage Providers

The following Storage Providers are available in a standard installation of DDF:

None.

Sources Details

Availability and configuration details of available sources.

5.8.5.1. Federated Source for Atlassian Confluence(R)

The Confluence source provides a Federated Source to retrieve pages, comments, and attachments from an Atlassian Confluence® REST API and turns the results into Metacards the system can use. The Confluence source does provide a Connected Source interface but its functionality has not been verified.

Confluence Source has been tested against the following versions of Confluence with REST API v2

- Confluence 1000.444.5 (Cloud)
- Confluence 5.10.6 (Server)
- Confluence 5.10.7 (Server)

Installing the Confluence Federated Source

The Confluence Federated Source is installed by default with a standard installation in the Catalog application.

Add a New Confluence Federated Source through the Admin Console:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Sources** tab.
4. Add a New source.
5. Name the New source.
6. Select **Confluence Federated Source** from **Binding Configurations**.

Configuring the Confluence Federated Source

Configure an Existing Confluence Federated Source through the Admin Console:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Sources** tab.
4. Select the name of the source to edit.

See [Confluence Federated Source configurations](#) for all possible configurations.

Usage Limitations of the Confluence Federated Source

Most of the fields that can be queried on Confluence have some sort of restriction on them. Most of the fields do not support the `like` aka `~` operation so the source will convert `like` queries to `equal` queries for attributes that don't support `like`. If the source receives a query with attributes it doesn't understand, it will just ignore them. If the query doesn't contain any attributes that map to Confluence search attributes, an empty result set will be returned.

Depending on your version of Confluence, when downloading attachments you might get redirected to a different download URL. If this happens the default `URLResourceReader` configuration does not allow redirects and will fail the download. This can be fixed by enabling redirects in the [URLResourceReader configuration](#).

5.8.5.2. CSW Connected Source

The CSW Connected Source supports the ability to search collections of descriptive information (metadata) for data, services, and related information objects.

Use the CSW source if querying a CSW version 2.0.2 compliant service.

Installing the CSW Connected Source

The CSW Connected Source is not installed by default with a standard installation in the Spatial application.

Add a CSW Connected Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Add a New source.
- Name the New source.
- Select **CSW Connected Source** from **Binding Configurations**.

Configuring the CSW Connected Source

Configure an Existing CSW Connected Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the name of the source to edit.

See [CSW Connected Source configurations](#) for all possible configurations.

Usage Limitations of the CSW Connected Source

- The CSW Connected Source does not support text path searches.
- Nearest neighbor spatial searches are not supported.

5.8.5.3. CSW Federated Source

The CSW Federated Source supports the ability to search collections of descriptive information (metadata) for data, services, and related information objects.

Use the CSW source if querying a CSW version 2.0.2 compliant service.

Installing the CSW Federated Source

The CSW Federated Source is installed by default with a standard installation in the Spatial application.

Add a New CSW Federated Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Add a New source.
- Name the New source.
- Select **CSW Federation Profile Source** from **Binding Configurations**.

Configuring the CSW Federated Source

Configure an Existing CSW Federated Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Select the name of the source to edit.

See [CSW Federated Source configurations](#) for all possible configurations.

Usage Limitations of the CSW Federated Source

- The CSW Federated Source does not support text path searches.
- Nearest neighbor spatial searches are not supported.

5.8.5.4. CSW Federation Profile Source

The CSW Federation Profile Source supports the ability to search collections of descriptive information (metadata) for data, services, and related information objects.

Use the CSW source if querying a CSW version 2.0.2 compliant service.

Installing the CSW Federation Profile Source

The CSW Federation Profile Source is installed by default with a standard installation in the Spatial application.

Configure the CSW Federation Profile Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Add a New source.
- Name the New source.
- Select **CSW Specification Profile Federated Source** from **Binding Configurations**.

Configuring the CSW Federation Profile Source

Configure an Existing CSW Federated Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Select the name of the source to edit.

See [CSW Federation Profile Source configurations](#) for all possible configurations.

Usage Limitations of the CSW Federation Profile Source

- The CSW Federation Profile Source does not support text path searches.
- Nearest neighbor spatial searches are not supported.

5.8.5.5. GMD CSW Source

The GMD CSW source supports the ability to search collections of descriptive information (metadata) for data, services, and related information objects, based on the [Application Profile ISO 19115/ISO19119](#).

Use the GMD CSW source if querying a GMD CSW APISO compliant service.

Installing the GMD CSW APISO v2.0.2 Source

The GMD CSW source is installed by default with a standard installation in the Spatial application.

Configure a new GMD CSW APISO v2.0.2 Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Add a New source.
- Name the New source.
- Select **GMD CSW ISO Federated Source** from **Binding Configurations**.

Configuring the GMD CSW APISO v2.0.2 Source

Configure an existing GMD CSW APISO v2.0.2 Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Select the name of the source to edit.

See [GMD CSW APISO v2.0.2 Source configurations](#) for all possible configurations.

5.8.5.6. OpenSearch Source

The OpenSearch source provides a [Federated Source](#) that has the capability to do [OpenSearch](#) queries for metadata from Content Discovery and Retrieval (CDR) Search V1.1 compliant sources. The OpenSearch source does not provide a [Connected Source](#) interface.

Installing an OpenSearch Source

The OpenSearch Source is installed by default with a standard installation in the Catalog application.

Configure a new OpenSearch Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Add a New source.
- Name the New source.
- Select **OpenSearch Source** from **Binding Configurations**.

Configuring an OpenSearch Source

Configure an existing OpenSearch Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Select the name of the source to edit.

See [OpenSearch Source configurations](#) for all possible configurations.

Using OpenSearch Source

Use the OpenSearch source if querying a CDR-compliant search service is desired.

Table 9. OpenSearch Parameter to DDF Query Mapping

OpenSearch/CDR Parameter	DDF Data Location
q={searchTerms}	Pulled verbatim from DDF query.
src={fs:routeTo?}	Unused
mr={fs:maxResults?}	Pulled verbatim from DDF query.
count={count?}	Pulled verbatim from DDF query.
mt={fs:maxTimeout?}	Pulled verbatim from DDF query.
dn={idn:userDN?}	DDF Subject
lat={geo:lat?}	Pulled verbatim from DDF query.
lon={geo:lon?}	Pulled verbatim from DDF query.
radius={geo:radius?}	Pulled verbatim from DDF query.
bbox={geo:box?}	Converted from Point-Radius DDF query.
polygon={geo:polygon?}	Pulled verbatim from DDF query.
dtstart={time:start?}	Pulled verbatim from DDF query.
dtend={time:end?}	Pulled verbatim from DDF query.
dateName={cat:dateName?}	Unused
filter={fsa:filter?}	Unused
sort={fsa:sort?}	Translated from DDF query. Format: "relevance" or "date" Supports "asc" and "desc" using colon as delimiter.

Usage Limitations of the OpenSearch Source

The OpenSearch source does not provide a [Connected Source](#) interface.

5.8.5.7. Registry Store

NOTE

The Registry Store is currently marked **Experimental**. While functional and tested, it is subject to change or removal during the incubation period.

The Registry Store is the interface that allows CSW messages to be turned into usable Registry metacards and for those metacards to be turned back into CSW messages.

Installing Registry Store

The Registry Store is installed by default with the Registry application.

Configuring Registry Store

To configure the Registry store:

1. Navigate to the Admin Console.
2. Select **Registry**.
3. Select the **Remote Registries** Tab and click the **Add** button.
 - a. ALTERNATIVELY: Select the **Configuration** Tab and select **Registry Store**.

5.8.5.8. Solr Catalog Provider

The Solr Catalog Provider is included with a standard installation of DDF. The following are the three ways to configure the Solr Catalog Provider:

HTTP Solr (default)

Used with an external (HTTP) Solr server.

Embedded Solr

Used with an embedded Solr server with no HTTP interface.

Solr Cloud

Used with a cluster of Solr servers.

If not using Solr Cloud, whenever the DDF Solr ingests data, it creates an index for it in the data directory. This data indexing enables the data to be searchable by the DDF Solr. The data in the data directory can then be persisted by performing a `commit` on the DDF Solr. When the DDF Solr is restarted, it reloads the data to get back the index.

HTTP Solr


The HTTP Solr is an external Apache Solr server instance. There are two possible installations for the HTTP Solr configuration:

Internal (default)


Points to an Apache Solr server that runs in the same container as DDF.

External

Points to an Apache Solr server in a different container or on a different system.



HTTP Solr Internal Deployment



HTTP Solr External Deployment

When to Use HTTP Solr

The HTTP Solr configuration is good for most use cases. See below for details.

HTTP Solr Tradeoffs

- Solr Admin UI can be used.
- Backup can be performed manually or using the Backup command.
- External installation allows isolation of data storage.
- Requires additional protection to be secured.
- Not scalable.
- External installation has network overhead

Installing HTTP Solr

By default, the Solr Catalog application installs a Solr server inside the same container as DDF. To install an External Solr server, refer to [Apache Solr Reference Guide](#).

Configuring HTTP Solr

1. Edit the `<DDF_HOME>/etc/system.properties`

- Comment out the Solr Client Configuration for **Cloud Solr Client** and **Embedded Solr Client** sections.
- Uncomment the section for the **Http Solr Client**.

NOTE To point to an External Solr server instead of the default Internal one, put the url in `solr.http.url` (see below for example). Also, two jar files need to be added to the Solr installation as described in [Installing Solr Cloud](#).

system.properties

```


solr.client = HttpSolrClient
solr.http.url = https://solr-host-name:8993/solr
solr.data.dir = ${karaf.home}/data/solr

```

Embedded Solr

The Embedded Solr is an embedded, local Apache Solr server instance. It is a local instance that is a

lightweight Catalog Provider solution. However, it does not provide a Solr Admin UI or a REST-like HTTP/XML and JSON API. If that is necessary, see [HTTP Solr](#) or [Solr Cloud](#).



Embedded Solr Deployment

When to Use the Embedded Solr

The Embedded Solr configuration is great for demonstrations, training exercises, or for sparse querying and ingesting. See below for details.

Embedded Solr Tradeoffs

- Requires no installation and very little configuration.
- No HTTP connection required. As a result, there is no network overhead.
- Not scalable.
- There is no Solr Admin UI.
- Backup must be performed manually.

Installing Embedded Solr

No installation required. Follow the Configuring section below and restart the DDF.

Configuring Embedded Solr

1. Edit the `<DDF_HOME>/etc/system.properties`
 - a. Comment out the Solr Client Configuration for **Cloud Solr Client** and **Http Solr Client** sections.
 - b. Uncomment the section for the **Embedded Solr Client**:


system.properties

```
solr.client = EmbeddedSolrServer
solr.data.dir = ${karaf.home}/data/solr
```

Solr Cloud

Solr Cloud is a cluster of Solr Server instances that are fault tolerant and highly available. Each Solr Server instance in Solr Cloud provides a Solr Admin UI.

Configuration shared between Solr Server instances is managed by Zookeeper. Zookeeper helps manage the overall structure.



Solr Cloud Deployment

NOTE

Solr Cloud is currently in Beta version. Bugs and feature changes may affect the final release.

When to Use Solr Cloud

If the DDF needs to be accessible for a higher than normal period with low downtime, then Solr Cloud should be used. See below for details.

Solr Cloud Tradeoffs

- Scalable. Can exceed ~2 billion indexed documents.
- Solr Admin UI can be used.
- Has network overhead and requires additional protection to be secure.
- Installation is more involved (requires Zookeeper)
- Configuration and administration is more complex due to replicate, sharding, etc.
- No way to backup currently, but will automatically recover from system failure.

NOTE

The instructions on setting up Solr Cloud for DDF only include setup in a *NIX environment.

Solr Cloud Prerequisites

Before Solr Cloud can be installed:

- ZooKeeper 3.4.5 (Refer to https://zookeeper.apache.org/doc/r3.1.2/zookeeperStarted.html#sc_Download for installation instructions.)
- *NIX environment
- JDK 8 or greater

NOTE

A minimum of three Zookeeper nodes required. Three Zookeeper nodes are needed to form a quorum. A three Zookeeper ensemble allows for a single server to fail and the service will still be available. More Zookeeper nodes can be added to achieve greater fault tolerance. The total number of nodes must always be an odd number.

Installing Solr Cloud

Repeat the following procedure for each Solr server instance that will be part of the Solr Cloud cluster:

1. Refer to <https://cwiki.apache.org/confluence/display/solr/Apache+Solr+Reference+Guide> for installation instructions.
2. Download jar files. The jars are needed to support geospatial and xpath queries and need to be installed on every Solr server instance after the Solr Cloud installation instructions have been followed.
 - a. http://artifacts.codice.org/service/local/repositories/releases/content/org/codice/thirdparty/jts/1.12_1/jts-1.12_1.jar
 - b. <http://artifacts.codice.org/service/local/artifact/maven/content?r=public&g=ddf.platform.solr&a=solr-xpath&v=2.11.3>
3. Copy downloaded jar files to: <SOLR_INSTALL_DIR>/server/solr-webapp/webapp/WEB-INF/lib/

NOTE

A minimum of two Solr server instances is required. Each Solr server instance must have a minimum of two shards. Having two Solr server instances guarantees that at least one Solr server is available if one fails. The two shards enables the document mapping to be restored if one shard becomes unavailable.

Configuring Solr Cloud

1. On the DDF server, edit <DDF_HOME>/etc/system.properties:
 - a. Comment out the Solr Client Configuration for **Http Solr Client** and **Embedded Solr Client** sections.
 - b. Uncomment the section for the **Cloud Solr Client**:
 - c. Set **solr.cloud.zookeeper** to <ZOOKEEPER_1_HOSTNAME>:<PORT_NUMBER>, <ZOOKEEPER_2_HOSTNAME>:<PORT_NUMBER>, <ZOOKEEPER_n_HOSTNAME>:<PORT_NUMBER>
 - d. Set **solr.data.dir** to the desired data directory.

system.properties

```
solr.client = CloudSolrClient
solr.data.dir = ${karaf.home}/data/solr
solr.cloud.zookeeper = zk1:2181,zk2:2181,zk3:2181
```

5.8.5.9. WFS 1.0 Source

The WFS Source allows for requests for geographical features across the web using platform-

independent calls.

A Web Feature Service (WFS) source is an implementation of the [FederatedSource](#) interface provided by the DDF Framework.

Use the WFS Source if querying a WFS version 1.0.0 compliant service.

Installing the WFS v1.0.0 Source

The WFS v1.0.0 Source is installed by default with a standard installation in the Spatial application.

Configure a new WFS v1.0.0 Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Add a New source.
- Name the New source.
- Select **WFS v1.0.0 Source** from **Binding Configurations**.

Configuring the WFS v1.0.0 Source

Configure an existing WFS v1.0.0 Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Select the name of the source to edit.

See [WFS v.1.0 Federated Source configurations](#) or [WFS v1.0 Connected Source configurations](#) for all possible configurations.a

WFS URL

The WFS URL must match the endpoint for the service being used. The type of service and version are added automatically, so they do not need to be included. Some servers will throw an exception if they are included twice, so do not include those.

The syntax depends on the server. However, in most cases, the syntax will be everything before the **?** character in the URL that corresponds to the [GetCapabilities](#) query.

Example GeoServer 2.5 Syntax

```
http://www.example.org:8080/geoserver/ows?service=wfs&version=1.0.0&request=GetCapabil  
ities
```

In this case, the WFS URL would be: <http://www.example.org:8080/geoserver/ows>

5.8.5.10. WFS 2.0 Source

The WFS 2.0 Source allows for requests for geographical features across the web using platform-independent calls.

Use the WFS Source if querying a WFS version 2.0.0 compliant service. Also see [Working with WFS Sources](#).

Installing the WFS v2.0.0 Source

The WFS v2.0.0 Source is installed by default with a standard installation in the Spatial application.

Configure a new WFS v2.0.0 Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Add a New source.
- Name the New source.
- Select **WFS v2.0.0 Source** from **Binding Configurations**.

Configuring the WFS v2.0.0 Source

Configure an existing WFS v2.0.0 Source through the Admin Console:

- Navigate to the **Admin Console**.
- Select the **Catalog** application.
- Select the **Sources** tab.
- Select the name of the source to edit.

See [WFS v.2.0 Federated source configurations](#) or [WFS v2.0 Connected source configurations](#) for all possible configurations.

WFS URL

The WFS URL must match the endpoint for the service being used. The type of service and version is added automatically, so they do not need to be included. Some servers will throw an exception if they are included twice, so do not include those.

The syntax depends on the server. However, in most cases, the syntax will be everything before the **?** character in the URL that corresponds to the **GetCapabilities** query.

Example GeoServer 2.5 Syntax

```
http://www.example.org:8080/geoserver/ows?service=wfs&version=2.0.0&request=GetCapabil  
ities
```

In this case, the WFS URL would be

<http://www.example.org:8080/geoserver/ows>

Mapping WFS Feature Properties to Metocard Attributes

The WFS 2.0 Source allows for virtually any schema to be used to describe a feature. A feature is relatively equivalent to a metocard. The **MetocardMapper** was added to allow an administrator to configure which feature properties map to which metocard attributes.

Using the WFS MetocardMapper

Use the WFS **MetocardMapper** to configure which feature properties map to which metocard attributes when querying a WFS version 2.0.0 compliant service. When feature collection responses are returned from WFS sources, a default mapping occurs which places the feature properties into metocard attributes, which are then presented to the user via DDF. There can be situations where this automatic mapping is not optimal for your solution. Custom mappings of feature property responses to metocard attributes can be achieved through the **MetocardMapper**. The **MetocardMapper** is set by creating a feature file configuration which specifies the appropriate mapping. The mappings are specific to a given feature type.

Installing the WFS MetocardMapper

The WFS **MetocardMapper** is not installed by default with a standard application in the Spatial application.

Configuring the WFS MetocardMapper

There are two ways to configure the **MetocardMapper**, one is to use the Configuration Admin available via the Admin Console. Additionally, a **feature.xml** file can be created and copied into the "deploy" directory.

Example WFS MetocardMapper Configuration

The following shows how to configure the **MetocardMapper** to be used with the sample data provided with GeoServer. This configuration shows a custom mapping for the feature type 'states'. For the given type, we are taking the feature property 'states.STATE_NAME' and mapping it to the metocard attribute 'title'. In this particular case, since we mapped the state name to title in the metocard, it will now be fully searchable. More mappings can be added to the **featurePropToMetocardAttrMap** line through the use of comma as a delimiter.

*Example MetocardMapper Configuration Within a **feature.xml** file:*

```
<feature name="geoserver-states" version="2.11.3"
  description="WFS Feature to Metocard mappings for GeoServer Example
  {http://www.openplans.org/topp}states">
  <config name="org.codice.ddf.spatial.ogc.wfs.catalog.mapper.MetocardMapper-
  geoserver.http://www.openplans.org/topp.states">
    featureType = {http://www.openplans.org/topp}states
    featurePropToMetocardAttrMap = states.STATE_NAME=title
  </config>
</feature>
```

5.9. Environment Hardening

- Required Step for Security Hardening

IMPORTANT It is recommended to apply the following security mitigations to the DDF.

Protocol/ Type	Risk	Mitigation
JMX	tampering, information disclosure, and unauthorized access	<ul style="list-style-type: none">• Stop the management feature using the command line console: <code>feature:stop management</code>.
File System Access	tampering, information disclosure, and denial of service	<p>Set OS File permissions under the <code><DDF_HOME></code> directory (e.g. <code>/deploy</code>, <code>/etc</code>) to ensure unauthorized viewing and writing is not allowed.</p> <p>If Caching is installed:</p> <ul style="list-style-type: none">• Set permissions for the installation directory <code>/data/product-cache</code> such that only the DDF process and users with the appropriate permissions can view any stored product.• Caching can be turned off as well to mitigate this risk.<ul style="list-style-type: none">◦ To disable caching, navigate to Admin Console.◦ Select the Catalog application.◦ Select Resource Download Settings.◦ Uncheck the <code>Enable Product Caching</code> box.• Install Security to ensure only the appropriate users are accessing the products.<ul style="list-style-type: none">◦ Navigate to the Admin Console◦ Select Manage.◦ Install the Security application, if applicable.• Cached files are written by the user running the DDF <code>process/application</code>. <p>On system: ensure that not everyone can change ACLs on your object.</p>

SSH	<p>tampering, information disclosure, and denial of service</p>	<p>By default, SSH access to DDF is only enabled to connections originating from the same host running DDF. For remote access to DDF, first establish an SSH session with the host running DDF. From within that session, initiate a new SSH connection (to localhost), and use the <code>sshPort</code> as configured in the file <code><DDF_HOME>/etc/org.apache.karaf.shell.cfg</code>.</p> <p>To allow direct remote access to the DDF shell from any host, change the value of the <code>sshHost</code> property to <code>0.0.0.0</code> in the <code><DDF_HOME>/etc/org.apache.karaf.shell.cfg</code> file.</p> <p>SSH can also be authenticated and authorized through an external Realm, such as LDAP. This can be accomplished by editing the <code><DDF_HOME>/etc/org.apache.karaf.shell.cfg</code> file and setting the value for <code>sshRealm</code>, e.g. to <code>ldap</code>. No restart of DDF is necessary after this change.</p> <p>By definition, all connections over SSH will be authenticated and authorized and secure from eavesdropping.</p> <p>WARNING</p> <p>Enabling SSH will expose your file system such that any user with access to your DDF shell will have read/write/execute access to all directories and files accessible to your installation user.</p> <p>Because of this, SSH is not recommended in a secure environment and should be turned off in a fully hardened system.</p> <p>Set <code>karaf.shutdown.port=-1</code> in <code><DDF_HOME>/etc/custom.properties</code> or <code><DDF_HOME>/etc/config.properties</code>.</p>
-----	---	--

SSL/TLS	man-in-the-middle, information disclosure	Update the <code><DDF_HOME>/etc/org.ops4j.pax.web.cfg</code> file to add the entry <code>org.ops4j.pax.web.ssl.clientauthneeded=true</code> .	
		WARNING	Setting this configuration may break compatibility to legacy systems that do not support two-way SSL.
		WARNING	Setting this configuration will require a certificate to be installed in the browser.

5.10. Configuring for Special Deployments

In addition to standard configurations, several specialized configurations are possible for specific uses of DDF.

5.10.1. Multiple Installations

One common specialized configuration is installing multiple instances of DDF.

5.10.1.1. Reusing Configurations

The Migration Export/Import capability allows administrators to export the current DDF configuration and use it to restore the same state for either a brand new installation or a second node for a Highly Available Cluster.

To export the current configuration settings:

1. Run the command `migration:export` from the Command Console.
2. A file named `ddf-2.11.3.dar` will be created in the `exported` directory underneath `<DDF_HOME>`.
Copy this file to a secure place.

To import previously exported configuration settings:

1. Install DDF

2. Restore all external files, softlinks, and directories that would not have been exported and for which warnings would have been generated during export. This could include (but is not limited to) external certificates or monitored directories.
3. Launch the newly installed DDF.
4. Step through the installation process.
5. Make sure to re-enable [automatic startup and shutdown](#) of DDF on the new system if it was enabled on the originating system.
6. Copy the previously exported file from your secure place to the [exported](#) directory underneath [`<DDF_HOME>`](#). Create the directory if it does not exist.
7. Run the command `migration:import` from the Command Console.
8. DDF will automatically restart if the command is successful. Otherwise address any generated warnings before manually restarting DDF.
9. Finish [hardening](#) the system (e.g. file and directory permissions).

- The following is currently not supported when importing configuration files:
 - importing from a different DDF version
 - importing from a system installed on a different OS
 - importing from a system installed in a different directory location
- To keep the export/import process simple and consistent, all system configuration files are required to be under the [`<DDF_HOME>`](#) directory and not be softlinks. Presence of external files or symbolic links during export will not fail the export; they will yield warnings. It will be up to the administrator to manually copy these files over to the new system before proceeding with the import. The import process will verify their presence and consistency and yield warnings if they don't match the original files.
- The import process will re-import all configurations done on the original system as part of the [hardening process](#) including changes to starting scripts and certificates.

IMPORTANT

5.10.1.2. Isolating Solr Cloud and Zookeeper

- **Required Step for Security Hardening** (if using Solr Cloud/Zookeeper)

Zookeeper cannot use secure (SSL/TLS) connection. The configuration information that Zookeeper sends and receives is vulnerable to network sniffing. Also, the connections between the local Solr Catalog service and the Solr Cloud is not necessarily secure. The connections between Solr Cloud nodes are not necessarily secure. Any unencrypted network traffic is vulnerable to sniffing attacks. To use Solr Cloud and Zookeeper securely, these processes must be isolated on the network, or their communications must be encrypted by other means. The DDF process must be visible on the network to allow authorized parties to interact with it.

Examples of Isolation:

- Create a private network for Solr Cloud and Zookeeper. Only DDF is allowed to contact devices inside the private network.
 - Use IPsec to encrypt the connections between DDF, Solr Cloud nodes, and Zookeeper nodes.
 - Put DDF, Solr Cloud and Zookeeper behind a firewall that only allows access to DDF.
-

5.10.1.3. Installing Multiple DDFs on the Same Host

To have multiple DDF instances on the same host, it is necessary to edit the port numbers in the files in the DDF install folder.

File to Edit	Property(ies)	Original Value	Example of New Value
bin/karaf.bat	address	5005	5006
etc/org.apache.karaf.management.cfg	rmiRegistryPort	1099	1199
	rmiServerPort	44444	44445
etc/system.properties	httpsPort, port	8993	8994
	httpPort	8181	8281

5.10.2. Configuring for a Fanout Proxy

This scenario describes how to configure DDF as a fanout proxy such that only queries and resource retrieval requests are processed and create/update/delete requests are rejected. All queries are enterprise queries and no catalog provider needs to be configured.

1. Reconfigure DDF in fanout proxy mode by going to the Features tab in the
2. Navigate to the **Admin Console**.
3. Select the **Catalog** application.
4. Select the **Configuration** tab.
5. Select **Catalog Standard Framework**.
6. Select **Enable Fanout Proxy**.
7. Save changes.

DDF is now operating as a fanout proxy. Only queries and resource retrieval requests will be allowed. All queries will be federated. Create, update, and delete requests will throw an **UnsupportedOperationException**, even if a Catalog Provider was configured prior to the reconfiguration to fanout.

5.10.3. Standalone Security Token Service (STS) Installation

To run a STS-only DDF installation, uninstall the catalog components that are not being used. The following list displays the features that can be uninstalled to minimize the runtime size of DDF in an STS-only mode. This list is not a comprehensive list of every feature that can be uninstalled; it is

a list of the larger components that can be uninstalled without impacting the STS functionality.

Unnecessary Features for Standalone STS

- `catalog-core-standardframework`
- `catalog-opensearch-endpoint`
- `catalog-opensearch-souce`
- `catalog-rest-endpoint`

5.10.4. Configuring for a Highly Available Cluster

This section describes how to make configuration changes after the initial setup for a DDF in a [Highly Available Cluster](#).

In a Highly Available Cluster, configuration changes must be made on both DDF nodes. The changes can still be made in the standard ways via the [Admin Console](#), the [Command Line](#), or the [file system](#).

Changes made in the Admin Console must be made through the HTTP proxy. This means that the below steps should be followed to make a change in the Admin Console:

NOTE

- Make a configuration change on the currently active DDF node
- Shut down the active DDF node, making the failover proxy switch to the standby DDF node
- Make the same configuration change on the newly active DDF node
- Start the DDF node that was just shut down

5.11. Miscellaneous Configurations

The optional configurations in this section cover minor changes that can be made to optimize DDF.

5.11.1. Configuring Thread Pools

The `org.codice.ddf.system.threadPoolSize` property can be used to specify the size of thread pools used by:

- Federating requests between DDF systems
- Downloading resources
- Handling asynchronous queries, such as queries from the UI

By default, this value is set to 128. It is not recommended to set this value extremely high. If unsure, leave this setting at its default value of 128.

5.11.2. Configuring Jetty ThreadPool Settings

To prevent resource shortages in the event of concurrent requests, DDF allows configuring Jetty

ThreadPool settings to specify the minimum and maximum available threads.

1. The settings can be changed at `etc/org.ops4j.pax.web.cfg` under Jetty Server ThreadPool Settings
2. Specify the maximum thread amount with `org.ops4j.pax.web.server.maxThreads`
3. Specify the minimum thread amount with `org.ops4j.pax.web.server.minThreads`
4. Specify the allotted time for a thread to complete with `org.ops4j.pax.web.server.idleTimeout`

DDF does not support changing ThreadPool settings from the Command Console or the Admin Console.

5.11.3. Configuring Alerts

By default, DDF uses two services provided by Karaf Decanter for alerts that can be configured by configuration file. Further information on Karaf Decanter services and configurations can be found [here](#).

5.11.3.1. Configuring Decanter Service Level Agreement (SLA) Checker

The Decanter SLA Checker provides a way to create alerts based on configurable conditions in events posted to `decanter/collect/*` and can be configured by editing the file `<DDF_HOME>/etc/org.apache.karaf.decanter.sla.checker.cfg`. By default there are only two checks that will produce alerts, and they are based on the SystemNotice event property of `priority`.

Table 10. Decanter SLA Configuration

Property	Alert Level	Expression	Description
priority	warn	equal:1,2,4	Produce a warn level alert if priority is important (3)
priority	error	equal:1,2,3	Produce an error level alert if priority is critical (4)

5.11.3.2. Configuring Decanter Scheduler

The Decanter Scheduler looks up services implementing the Runnable interface with the service-property `decanter.collector.name` and executes the Runnable periodically. The Scheduler can be configured by editing the file `<DDF_HOME>/etc/org.apache.karaf.decanter.scheduler.simple.cfg`.

Table 11. Decanter Scheduler Configuration

Property Name	Description	Default Value
period	Decanter simple scheduler period (milliseconds)	300000 (5 minutes)
threadIdleTimeout	The time to wait before stopping an idle thread (milliseconds)	60000 (1 minute)

Property Name	Description	Default Value
threadInitCount	Initial number of threads created by the scheduler	5
threadMaxCount	Maximum number of threads created by the scheduler	200

5.11.4. Landing Page

The Landing Page is the first page presented to users of DDF. It is customizable to allow adding organizationally-relevant content.

5.11.4.1. Installing the Landing Page

The Landing Page is installed by default with a standard installation.

5.11.4.2. Configuring the Landing Page

The DDF landing page offers a starting point and general information for a DDF node. It is accessible at `/(index|home|landing(.htm|html))`.

5.11.4.3. Customizing the Landing Page

Configure the Landing Page from the Admin Console:

1. Navigate to the **Admin Console**.
2. Select **Platform** Application.
3. Select **Configuration** tab.
4. Select **Landing Page**.

See [Landing Page configurations](#) for all possible configurations.

6. Running

Find directions here for running an installation of DDF.

Starting

Getting an instance of DDF up and running.

Maintaining

Keeping DDF running with useful tasks.

Monitoring

Tracking system health and usage.

Troubleshooting

Common tips for unexpected behavior.

6.1. Starting

Follow the below steps to start and stop DDF.

Memory Considerations

If the DDF will be running memory-intensive applications, consider increasing the java memory. The java memory can be changed by updating the setenv script:

*Setenv Scripts: *NIX*

<DDF_HOME>/bin/setenv
Update the JAVA_OPTS -Xmx value
<DDF_HOME>/bin/setenv-wrapper.conf
Update the wrapper.java.maxmemory value

NOTE

Setenv Scripts: Windows

<DDF_HOME>/bin/setenv.bat
Update the JAVA_OPTS -Xmx value
<DDF_HOME>/bin/setenv-windows-wrapper.conf
Update the wrapper.java.maxmemory value

JMX Connectivity Considerations

By default the DDF prevents connections to JMX. To monitor the DDF via tools such as JConsole, consider removing the no attach mechanism by updating the setenv script:

*Setenv Scripts: *NIX*

<DDF_HOME>/bin/setenv
Remove -XX:+DisableAttachMechanism from JAVA_OPTS
<DDF_HOME>/bin/setenv-wrapper.conf
Comment out the -XX:+DisableAttachMechanism line and re-number remainder lines appropriately

NOTE

Setenv Scripts: Windows

<DDF_HOME>/bin/setenv.bat
Remove -XX:+DisableAttachMechanism from JAVA_OPTS
<DDF_HOME>/bin/setenv-windows-wrapper.conf
Comment out the -XX:+DisableAttachMechanism line and re-number remainder lines appropriately

6.1.1. Starting from Startup Scripts

Run one of the start scripts from a command shell to start the distribution and open a local console:

*Start Script: *NIX*

```
<DDF_HOME>/bin/ddf
```

Start Script: Windows

```
<DDF_HOME>/bin/ddf.bat
```

6.1.2. Starting as a Background Process

Alternatively, to run DDF as a background process, run the **start** script:

**NIX*

```
<DDF_HOME>/bin/start
```

Windows

```
<DDF_HOME>/bin/start.bat
```

If console access is needed while running as a service, run the **client** script on the host where the DDF is running:

**NIX*

```
<DDF_HOME>/bin/client
```

NOTE

Windows

```
<DDF_HOME>/bin/client.bat -h <FQDN>
```

Use the **-h** option followed by the name (**<FQDN>**) or IP of the host where DDF is running. U

6.1.3. Stopping DDF

There are two options to stop a running instance:

- Call shutdown from the console:

Shut down with a prompt

```
ddf@local>shutdown
```

Force Shutdown without prompt

```
ddf@local>shutdown -f
```

- Keyboard shortcut for shutdown
 - **Ctrl-D**
 - **Cmd-D**
- Or run the stop script:

**NIX*

```
<DDF_HOME>/bin/stop
```

Windows

```
<DDF_HOME>/bin/stop.bat
```

Shut Down

IMPORTANT

Do not shut down by closing the window (Windows, Unix) or using the **kill -9 <pid>** command (Unix). This prevents a clean shutdown and can cause significant problems when DDF is restarted. Always use the shutdown command or the shortcut from the command line console.

6.1.4. Starting as a Service with Automatic Start on System Boot

Because DDF is built on top of Apache Karaf, DDF can use the Karaf Wrapper to run DDF as a service and enable automatic startup and shutdown. When DDF is started using Karaf Wrapper, new **wrapper.log** and **wrapper.log.n** (where n goes from 1 to 5 by default) log files will be generated to include wrapper and console specific information.

1. Create the Karaf wrapper.

Within the DDF console:

**NIX*

```
ddf@local> feature:install -r wrapper
ddf@local> wrapper:install -i setenv-wrapper.conf -s AUTO_START -n ddf -d ddf -D
"DDF Service"
```

Windows

```
ddf@local> feature:install -r wrapper
ddf@local> wrapper:install -i setenv-windows-wrapper.conf -s AUTO_START -n ddf -d
ddf -D "DDF Service"
```

2. (Windows users skip to next step) (All *NIX) If DDF was installed to run as a non-root user (recommended,) edit `<DDF_HOME>/bin/ddf-service`.

Change:

```
<DDF_HOME>/bin/ddf-service
```

```
#RUN_AS_USER=
```

to (<ddf-user> is the intended username):

```
<DDF_HOME>/bin/ddf-service
```

```
RUN_AS_USER=<ddf-user>
```

3. (Windows users skip to next step) (All *NIX) Edit `<DDF_HOME>/bin/ddf-service`. Add LimitNOFILE to the [Service] section.

```
<DDF_HOME>/bin/ddf.service
```

```
LimitNOFILE=6815744
```

4. Install the wrapper startup/shutdown scripts.

Windows

Run the following command in a console window. The command must be run with elevated permissions.

```
<DDF_HOME>\bin\ddf-service.bat install
```

Startup and shutdown settings can then be managed through **Services** → **MMC Start** → **Control Panel** → **Administrative Tools** → **Services**.

Redhat, Fedora, or Centos (SystemV)

```
root@localhost# ln -s <DDF_HOME>/bin/ddf-service /etc/init.d/
root@localhost# chkconfig ddf-service --add
root@localhost# chkconfig ddf-service on
```

Redhat 7 or Centos 7 (systemd)

```
root@localhost# systemctl enable <DDF_HOME>/bin/ddf.service
```

Ubuntu

```
root@localhost# ln -s <DDF_HOME>/bin/ddf-service /etc/init.d/
root@localhost# update-rc.d -f ddf-service defaults
```

Solaris

```
root@localhost# ln -s <DDF_HOME>/bin/ddf-service /etc/init.d/
root@localhost# ln -s /etc/init.d/ddf-service /etc/rc0.d/K20ddf-service
root@localhost# ln -s /etc/init.d/ddf-service /etc/rc1.d/K20ddf-service
root@localhost# ln -s /etc/init.d/ddf-service /etc/rc2.d/K20ddf-service
root@localhost# ln -s /etc/init.d/ddf-service /etc/rc3.d/S20ddf-service
```

NOTE

If the system fails to start automatically on boot with any messages about the JVM, ensure that JAVA_HOME is properly set. [Java Requirements \(Quick Install\)](#)

While it is not a necessary step, information on how to convert the System V init scripts to the Solaris System Management Facility can be found at <http://www.oracle.com/technetwork/articles/servers-storage-admin/scripts-to-smf-1641705.html>

6.1.4.1. Karaf Documentation

Because DDF is built on top of Apache Karaf, more information on operating DDF can be found in the [Karaf documentation](#).

6.2. Maintaining

6.2.1. Console Commands

Once the distribution has started, users will have access to a powerful command line console, the Command Console. This Command Console can be used to manage services, install new features and applications, and manage the state of the system.

6.2.1.1. Accessing the System Console

The Command Line Console is the console that is available to the user when the distribution is started manually. It may also be accessed by using the `bin/client.bat` or `bin/client.sh` scripts.

NOTE

The majority of functionality and information available on the Admin Console is also available on the Command Line Console.

6.2.1.2. Available System Console Commands

To get a list of commands, type in the namespace of the desired extension then press the **Tab** key.

For example, type `catalog`, then press **Tab**.

6.2.1.3. System Console Command Help

For details on any command, type `help` then the command. For example, `help search` (see results of this command in the example below).

Example Help

```
ddf@local>help search
DESCRIPTION
  catalog:search
  Searches records in the catalog provider.

SYNTAX
  catalog:search [options] SEARCH_PHRASE [NUMBER_OF_ITEMS]

ARGUMENTS
  SEARCH_PHRASE
    Phrase to query the catalog provider.
  NUMBER_OF_ITEMS
    Number of maximum records to display.
    (defaults to -1)

OPTIONS
  --help
    Display this help message
  case-sensitive, -c
    Makes the search case sensitive
  -p, -provider
    Interacts with the provider directly instead of the framework.
```

The `help` command provides a description of the provided command, along with the syntax in how to use it, arguments it accepts, and available options.

Table 12. Available Console Commands

Title	Namesp ace	Description
DDF:: Catalog :: Core :: Commands	catalog	The Catalog Shell Commands are meant to be used with any CatalogProvider implementations. They provide general useful queries and functions against the Catalog API that can be used for debugging, printing, or scripting.
DDF :: Admin :: Application Service	app	The Admin Application Service contains operations to work with applications.
DDF :: Catalog :: Core :: PubSub Commands	subscription	The DDF PubSub shell commands provide functions to list the registered subscriptions in DDF and to delete subscriptions.

Title	Namespace	Description
DDF Platform Commands	platform	The DDF Platform Shell Commands provide generic platform management functions
DDF:: Persistence :: Core :: Commands	store	The Persistence Shell Commands are meant to be used with any PersistentStore implementations. They provide the ability to query and delete entries from the persistence store.

6.2.1.4. Catalog Commands

WARNING

Most commands can bypass the Catalog framework and interact directly with the Catalog provider if given the `--provider` option, if available. No pre/post plugins are executed and no message validation is performed if the `--provider` option is used.

Table 13. Catalog Command Descriptions

Command	Description	
<code>catalog:describe</code>	Provides a basic description of the Catalog implementation.	
<code>catalog:dump</code>	Exports metacards from the local Catalog. Does not remove them. See date filtering options below.	
<code>catalog:envlist</code>	IMPORTANT	Deprecated as of ddf-catalog 2.5.0. Please use <code>platform:envlist</code> .
	Provides a list of environment variables.	
<code>catalog:ingest</code>	Ingests data files into the Catalog.	
<code>catalog:inspect</code>	Provides the various fields of a metocard for inspection.	
<code>catalog:latest</code>	Retrieves the latest records from the Catalog based on the Metocard.MODIFIED date.	
<code>catalog:migrate</code>	Allows two <code>CatalogProvider</code> s to be configured and migrates the data from the primary to the secondary.	
<code>catalog:range</code>	Searches by the given range arguments (exclusively).	
<code>catalog:remove</code>	Deletes a record from the local Catalog.	

Comma nd	Description
<code>catalog: removeal</code>	Attempts to delete all records from the local Catalog.
<code>catalog: replicat e</code>	Replicates data from a federated source into the local Catalog.
<code>catalog: search</code>	Searches records in the local Catalog.
<code>catalog: spatial</code>	Searches spatially the local Catalog.
<code>catalog: transfor mers</code>	Provides information on available transformers.
<code>catalog: validate</code>	Validates an XML file against all installed validators and prints out human readable errors and warnings.

6.2.1.5. `catalog:dump` Options

The `catalog:dump` command was extended in DDF version 2.5.0 to provide selective export of metacards based on date ranges. The `--created-after` and `--created-before` options allow filtering on the date and time that the metocard was created, while `--modified-after` and `--modified-before` options allow filtering on the date and time that the metocard was last modified (which is the created date if no other modifications were made). These date ranges are exclusive (i.e., if the date and time match exactly, the metocard will not be included). The date filtering options (`--created-after`, `--created-before`, `--modified-after`, and `--modified-before`) can be used in any combination, with the export result including only metacards that match all of the provided conditions.

If no date filtering options are provided, created and modified dates are ignored, so that all metacards match.

6.2.1.6. Date Syntax

Supported dates are taken from the common subset of ISO8601, matching the datetime from the following syntax:

```

datetime      = time | date-opt-time
time         = 'T' time-element [offset]
date-opt-time = date-element ['T' [time-element] [offset]]
date-element  = std-date-element | ord-date-element | week-date-element
std-date-element = yyyy ['-' MM ['-' dd]]
ord-date-element = yyyy ['-' DDD]
week-date-element = xxxx '-W' ww ['-' e]
time-element  = HH [minute-element] | [fraction]
minute-element = ':' mm [second-element] | [fraction]
second-element = ':' ss [fraction]
fraction      = ('.' | ',') digit+
offset        = 'Z' | (( '+' | '-' ) HH [':' mm [':' ss [('.') | ',' ) SSS]])
```

6.2.1.6.1. catalog:dump Examples

```

ddf@local>// Given we've ingested a few metacards
ddf@local>catalog:latest
#           ID           Modified Date      Title
1   a6e9ae09c792438e92a3c9d7452a449f 2014-06-13T09:56:18+10:00
2   b4aced45103a400da42f3b319e58c3ed 2014-06-13T09:52:12+10:00
3   a63ab22361e14cee9970f5284e8eb4e0 2014-06-13T09:49:36+10:00  myTitle

ddf@local>// Filter out older files
ddf@local>catalog:dump --created-after 2014-06-13T09:55:00+10:00 /home/bradh/ddf-
catalog-dump
1 file(s) dumped in 0.015 seconds

ddf@local>// Filter out new file
ddf@local>catalog:dump --created-before 2014-06-13T09:55:00+10:00 /home/bradh/ddf-
catalog-dump
2 file(s) dumped in 0.023 seconds

ddf@local>// Choose middle file
ddf@local>catalog:dump --created-after 2014-06-13T09:50:00+10:00 --created-before
2014-06-13T09:55:00+10:00 /home/bradh/ddf-catalog-dump
1 file(s) dumped in 0.020 seconds

ddf@local>// Modified dates work the same way
ddf@local>catalog:dump --modified-after 2014-06-13T09:50:00+10:00 --modified-before
2014-06-13T09:55:00+10:00 /home/bradh/ddf-catalog-dump
1 file(s) dumped in 0.015 seconds

ddf@local>// Can mix and match, most restrictive limits apply
ddf@local>catalog:dump --modified-after 2014-06-13T09:45:00+10:00 --modified-before
2014-06-13T09:55:00+10:00 --created-before 2014-06-13T09:50:00+10:00 /home/bradh/ddf-
catalog-dump
1 file(s) dumped in 0.024 seconds

ddf@local>// Can use UTC instead of (or in combination with) explicit timezone offset
```

```

ddf@local>catalog:dump --modified-after 2014-06-13T09:50:00+10:00 --modified-before
2014-06-13T09:55:00Z /home/bradh/ddf-catalog-dump
2 file(s) dumped in 0.020 seconds
ddf@local>catalog:dump --modified-after 2014-06-13T09:50:00+10:00 --modified-before
2014-06-12T23:55:00Z /home/bradh/ddf-catalog-dump
1 file(s) dumped in 0.015 seconds

ddf@local>// Can leave off timezone, but default (local time on server) may not match
what you expect!
ddf@local>catalog:dump --modified-after 2014-06-13T09:50:00 --modified-before 2014-06-
13T09:55:00 /home/bradh/ddf-catalog-dump
1 file(s) dumped in 0.018 seconds

ddf@local>// Can leave off trailing minutes / seconds
ddf@local>catalog:dump --modified-after 2014-06-13T09 --modified-before 2014-06-
13T09:55 /home/bradh/ddf-catalog-dump
2 file(s) dumped in 0.024 seconds

ddf@local>// Can use year and day number
ddf@local>catalog:dump --modified-after 2014-164T09:50:00 /home/bradh/ddf-catalog-dump
2 file(s) dumped in 0.027 seconds

```

6.2.1.7. Application Commands

Application commands are used from the DDF Admin application to manage applications in the DDF.

NOTE The Application Commands are installed automatically with the Admin Application.

Table 14. App Command Descriptions

Command	Syntax	Description
add	app:add appUri	Adds an application with the given uri.
remove	app:remove appName	Removes an application with the given name.
start	app:start appName	Starts an application with the given name.
stop	app:stop appName	Stops an application with the given name.
list	app:list	Lists the applications that are in the system and gives their current state.
status	app:status appName	Shows status of an application. Gives information on the current state, features within the application, what required features are not started and what required bundles are not started.
tree	app:tree	Creates a hierarchy tree of all of the applications.

6.2.1.7.1. Application Command Usage

Listing All Applications

```
ddf@local>app:list
State      Name
[ACTIVE]  catalog-app-<VERSION>
[ACTIVE]  distribution-<VERSION>
[ACTIVE]  platform-app-<VERSION>
```

```
[...]
```

This list shows all of the applications installed in DDF. From here, use the name of an application to get more information on its status.

Get Status for a Specific Application

```
ddf@local>app:status catalog-app-<VERSION>
catalog-app-<VERSION>
```

Current State is: ACTIVE

Features Located within this Application:

```
catalog-security-filter
catalog-transformer-resource
catalog-rest-endpoint
abdera
catalog-transformer-xml
catalog-transformer-thumbnail
catalog-transformer-metadata
catalog-transformer-xsltengine
catalog-core-fanoutframework
catalog-transformer-tika
catalog-core-api
catalog-opensearch-source
catalog-plugin-federationreplication
catalog-opensearch-endpoint
catalog-schematron-plugin
catalog-transformer-geoformatter
catalog-transformer-atom
catalog-core-sourcetricsplugin
catalog-core-metricsplugin
catalog-app
catalog-transformer-json
catalog-core-standardframework
catalog-core
```

Required Features Not Started

NONE

Required Bundles Not Started

NONE

6.2.1.7.2. Application in Failed State

If an application is in a 'FAILED' state, it means that there is a required feature or bundle that is not started.

App in Failed State

```
ddf@local>app:list
State      Name
[FAILED   ] catalog-app-<VERSION>
[ACTIVE   ] distribution-<VERSION>
[ACTIVE   ] platform-app-<VERSION>
```

In the above case, the catalog app is in a failed state. Checking the status of that application will show what did not start correctly.

Check Application Status

```
ddf@local>app:status catalog-app-<VERSION>
catalog-app-<VERSION>
```

Current State is: FAILED

Features Located within this Application:

```
catalog-security-filter
catalog-transformer-resource
catalog-rest-endpoint
abdera
catalog-transformer-xml
catalog-transformer-thumbnail
catalog-transformer-metadata
catalog-transformer-xsltengine
catalog-core-fanoutframework
catalog-transformer-tika
catalog-core-api
catalog-opensearch-source
catalog-plugin-federationreplication
catalog-opensearch-endpoint
catalog-schematron-plugin
catalog-transformer-geoformatter
catalog-transformer-atom
catalog-core-sourcetricsplugin
catalog-core-metricsplugin
catalog-app
catalog-transformer-json
catalog-core-standardframework
catalog-core
```

Required Features Not Started

NONE

Required Bundles Not Started

[261] catalog-opensearch-endpoint

This status shows that bundle &261, the catalog-opensearch-endpoint, did not start. Performing a 'list' on the console verifies this:

```
[ 261] [Resolved ] [ ] [ ] [ 80] DDF :: Catalog :: OpenSearch ::  
Endpoint (<VERSION>)
```

Once that bundle is started by fixing its error, the catalog application will show as being in an ACTIVE state.

6.2.1.8. Subscriptions Commands

NOTE The subscriptions commands are installed when the Catalog application is installed.

Table 15. Subscription Command Descriptions

Command	Description
<code>subscription:delete</code>	Deletes the subscription(s) specified by the search phrase or LDAP filter.
<code>subscription:list</code>	List the subscription(s) specified by the search phrase or LDAP filter.

6.2.1.9. `subscriptions:list` Command Usage Examples

Note that no arguments are required for the `subscriptions:list` command. If no argument is provided, all subscriptions will be listed. A count of the subscriptions found matching the list command's search phrase (or LDAP filter) is displayed first followed by each subscription's ID.

List All Subscriptions

```
ddf@local>subscriptions:list

Total subscriptions found: 3

Subscription ID
my.contextual.id.v20|http://172.18.14.169:8088/mockCatalogEventConsumerBinding?WSDL
my.contextual.id.v30|http://172.18.14.169:8088/mockEventConsumerBinding?WSDL
my.contextual.id.json|http://172.18.14.169:8088/services/json/local/event/notification
```

List a Specific Subscription by ID

```
ddf@local>subscriptions:list
"my.contextual.id.v20|http://172.18.14.169:8088/mockCatalogEventConsumerBinding?WSDL"

Total subscriptions found: 1

Subscription ID
my.contextual.id.v20|http://172.18.14.169:8088/mockCatalogEventConsumerBinding?WSDL
```

WARNING

It is recommended to always quote the search phrase (or LDAP filter) argument to the command so that any special characters are properly processed.

List Subscriptions Using Wildcards

```
ddf@local>subscriptions:list "my*"
```

Total subscriptions found: 3

Subscription ID

```
my.contextual.id.v20|http://172.18.14.169:8088/mockCatalogEventConsumerBinding?WSDL
my.contextual.id.v30|http://172.18.14.169:8088/mockEventConsumerBinding?WSDL
my.contextual.id.json|http://172.18.14.169:8088/services/json/local/event/notification
```

```
ddf@local>subscriptions:list "*json*"
```

Total subscriptions found: 1

Subscription ID

```
my.contextual.id.json|http://172.18.14.169:8088/services/json/local/event/notification
```

```
ddf@local>subscriptions:list "*WSDL"
```

Total subscriptions found: 2

Subscription ID

```
my.contextual.id.v20|http://172.18.14.169:8088/mockCatalogEventConsumerBinding?WSDL
my.contextual.id.v30|http://172.18.14.169:8088/mockEventConsumerBinding?WSDL
```

The example below illustrates searching for any subscription that has "json" or "v20" anywhere in its subscription ID.

List Subscriptions Using an LDAP Filter

```
ddf@local>subscriptions:list -f "((subscription-id=*json*) (subscription-id=*v20*))"
```

Total subscriptions found: 2

Subscription ID

```
my.contextual.id.v20|http://172.18.14.169:8088/mockCatalogEventConsumerBinding?WSDL
my.contextual.id.json|http://172.18.14.169:8088/services/json/local/event/notification
```

The example below illustrates searching for any subscription that has **json** and **172.18.14.169** in its subscription ID. This could be a handy way of finding all subscriptions for a specific site.

```
ddf@local>subscriptions:list -f "(&(subscription-id=*json*) (subscription-id=*172.18.14.169*))"

Total subscriptions found: 1

Subscription ID
my.contextual.id.json|http://172.18.14.169:8088/services/json/local/event/notification
```

6.2.1.10. `subscriptions:delete` Command Usage

The arguments for the `subscriptions:delete` command are the same as for the `list` command, except that a search phrase or LDAP filter must be specified. If one of these is not specified an error will be displayed. When the `delete` command is executed it will display each subscription ID it is deleting. If a subscription matches the search phrase but cannot be deleted, a message in red will be displayed with the ID. After all matching subscriptions are processed, a summary line is displayed indicating how many subscriptions were deleted out of how many matching subscriptions were found.

Delete a Specific Subscription Using Its Exact ID

```
ddf@local>subscriptions:delete
"my.contextual.id.json|http://172.18.14.169:8088/services/json/local/event/notification"

Deleted subscription for ID =
my.contextual.id.json|http://172.18.14.169:8088/services/json/local/event/notification

Deleted 1 subscriptions out of 1 subscriptions found.
```

Delete Subscriptions Using Wildcards

```
ddf@local>subscriptions:delete "my*"

Deleted subscription for ID =
my.contextual.id.v20|http://172.18.14.169:8088/mockCatalogEventConsumerBinding?WSDL
Deleted subscription for ID =
my.contextual.id.v30|http://172.18.14.169:8088/mockEventConsumerBinding?WSDL

Deleted 2 subscriptions out of 2 subscriptions found.

ddf@local>subscriptions:delete "*json*"

Deleted subscription for ID =
my.contextual.id.json|http://172.18.14.169:8088/services/json/local/event/notification

Deleted 1 subscriptions out of 1 subscriptions found.
```

Delete All Subscriptions

```
ddf@local>subscriptions:delete *

Deleted subscription for ID =
my.contextual.id.v30|http://172.18.14.169:8088/mockEventConsumerBinding?WSDL
Deleted subscription for ID =
my.contextual.id.v20|http://172.18.14.169:8088/mockCatalogEventConsumerBinding?WSDL
Deleted subscription for ID =
my.contextual.id.json|http://172.18.14.169:8088/services/json/local/event/notification

Deleted 3 subscriptions out of 3 subscriptions found.
```

Delete Subscriptions Using an LDAP Filter

```
ddf@local>subscriptions:delete -f "(&(subscription-id=*WSDL) (subscription-
id=*172.18.14.169*))"

Deleted subscription for ID =
my.contextual.id.v20|http://172.18.14.169:8088/mockCatalogEventConsumerBinding?WSDL
Deleted subscription for ID =
my.contextual.id.v30|http://172.18.14.169:8088/mockEventConsumerBinding?WSDL

Deleted 2 subscriptions out of 2 subscriptions found.
```

6.2.1.11. Platform Commands

NOTE The Platform Commands are installed when the Platform application is installed.

Table 16. Platform Command Descriptions

Command	Description
describe	Shows the current platform configuration.
envlist	Provides a list of environment variables.

6.2.1.12. Persistence Store Commands

Table 17. Persistence Store Command Descriptions

Command	Description
store:delete	Delete entries from the persistence store that match a given CQL statement
store:list	Lists entries that are stored in the persistence store.

6.2.2. Command Scheduler

The Command Scheduler allows administrators to schedule Command Line Commands to be run at specified intervals.

6.2.2.1. Using the Command Scheduler

The Command Scheduler allows administrators to schedule Command Line Shell Commands to be run in a platform-independent way. For instance, if an administrator wanted to use the Catalog commands to export all records of a Catalog to a directory, the administrator could write a cron job or a scheduled task to remote into the container and execute the command. Writing these types of scripts are specific to the administrator's operating system and also requires extra logic for error handling if the container is up. The administrator can also create a Command Schedule, which currently requires only two fields. The Command Scheduler only runs when the container is running, so there is no need to verify if the container is up. In addition, when the container is restarted, the commands are rescheduled and executed again. A command will be repeatedly executed indefinitely according to the configured interval until the container is shutdown or the Scheduled Command is deleted.

NOTE

There will be further attempts to execute the command according to the configured interval even if an attempt fails. See the log for details about failures.

6.2.2.1.1. Schedule a Command

Configure the Command Scheduler to execute a command at specific intervals.

1. Navigate to the **Admin Console** (<https://{FQDN}:{PORT}/admin>).
2. Select the **Platform** application.
3. Click on the **Configuration** tab.
4. Select **Platform Command Scheduler**.
5. Enter the command or commands to be executed in the **Command** text field. Commands can be separated by a semicolon and will execute in order from left to right.
6. Enter an interval in the **Interval** field. This can either be a Quartz Cron expression or a positive integer (seconds) (e.g. `0 0 0 1/1 * ? *` or `12`).
7. Select the interval type in the **Interval Type** drop-down.
8. Click the **Save changes** button.

NOTE

Scheduling commands will be delayed by 1 minute to allow time for bundles to load when DDF is starting up.

6.2.2.2. Updating a Scheduled Command

Change the timing, order, or execution of scheduled commands.

1. Navigate to the **Admin Console**.
2. Click on the **Platform** application.
3. Click on the **Configuration** tab.
4. Under the **Platform Command Scheduler** configuration are all of the scheduled commands. Scheduled commands have the following syntax: `ddf.platform.scheduler.Command.{GUID}` such as `ddf.platform.scheduler.Command.4d60c917-003a-42e8-9367-1da0f822ca6e`.

5. Find the desired configuration to modify, and update fields.
6. Click the **Save changes** button.

6.2.2.2.1. Command Output

Commands that normally write out to the console will write out to the log. For example, if an `echo "Hello World"` command is set to run every five seconds, the log contains the following:

Sample Command Output in the Log

```
16:01:32,582 | INFO | heduler_Worker-1 | ddf.platform.scheduler.CommandJob
68 | platform-scheduler | Executing command [echo Hello World]
16:01:32,583 | INFO | heduler_Worker-1 | ddf.platform.scheduler.CommandJob
70 | platform-scheduler | Execution Output: Hello World
16:01:37,581 | INFO | heduler_Worker-4 | ddf.platform.scheduler.CommandJob
68 | platform-scheduler | Executing command [echo Hello World]
16:01:37,582 | INFO | heduler_Worker-4 | ddf.platform.scheduler.CommandJob
70 | platform-scheduler | Execution Output: Hello World
```

In short, administrators can view the status of a run within the log as long as `INFO` was set as the status level.

6.2.2.3. CQL Syntax

The CQL syntax used with console commands should follow the OGC CQL format. GeoServer provides a description of the grammar and examples at [CQL Tutorial](#).

CQL Syntax Examples

```
Finding all notifications that were sent due to a download:
ddf@local>store:list --cql "application='Downloads'" --type notification
```

```
Deleting a specific notification:
ddf@local>store:delete --cql "id='fdc150b157754138a997fe7143a98cfa'" --type notification
```

6.3. Monitoring

The DDF contains many tools to monitor system functionality, usage, and overall system health.

6.3.1. Metrics Reporting

Metrics are available in several formats and levels of detail.

Complete the following procedure now that several queries have been executed.

1. Select **Platform**
2. Select **Metrics** tab

3. For individual metrics, choose the format desired from the desired timeframe column:
 - a. PNG
 - b. CSV
 - c. XLS
4. For a detailed report of all metrics, at the bottom of the page are selectors to choose time frame and summary level. A report is generated in xls format.

6.3.2. Managing Logging

The DDF supports a dynamic and customizable logging system including log level, log format, log output destinations, roll over, etc.

6.3.2.1. Configuring Logging

Edit the configuration file `<DDF_HOME>/etc/org.ops4j.pax.logging.cfg`]

6.3.2.2. DDF log file

The name and location of the log file can be changed with the following setting:

```
log4j.appender.out.file=<KARAF.DATA>/log/ddf.log
```

6.3.2.3. Controlling log level

A useful way to debug and detect issues is to change the log level:

```
log4j.rootLogger=DEBUG, out, osgi:VmLogAppender
```

6.3.2.4. Controlling the size of the log file

Set the maximum size of the log file before it is rolled over by editing the value of this setting:

```
log4j.appender.out.maxFileSize=20MB
```

6.3.2.5. Number of backup log files to keep

Adjust the number of backup files to keep by editing the value of this setting:

```
log4j.appender.out.maxBackupIndex=10
```

6.3.2.6. Enabling logging of inbound and outbound SOAP messages for the DDF SOAP endpoints

By default, the DDF start scripts include a system property enabling logging of inbound and outbound SOAP messages.

```
-Dcom.sun.xml.ws.transport.http.HttpAdapter.dump=true
```

In order to see the messages in the log, one must set the logging level for `org.apache.cxf.services` to `INFO`. By default, the logging level for `org.apache.cxf` is set to `WARN`.

```
ddf@local>log:set INFO org.apache.cxf.services
```

6.3.2.7. Logging External Resources

Other appenders can be selected and configured.

For more detail on configuring the log file and what is logged to the console see: [Karaf Documentation: Log](#).

6.3.2.8. Enabling HTTP Access Logging

To enable access logs for the current DDF, do the following:

- Update the `jetty.xml` file located in `etc/` adding the following xml:

```
<Get name="handler">
  <Call name="addHandler">
    <Arg>
      <New class="org.eclipse.jetty.server.handler.RequestLogHandler">
        <Set name="requestLog">
          <New id="RequestLogImpl" class="org.eclipse.jetty.server.NCSARequestLog">
            <Arg><SystemProperty name="jetty.logs" default="data/log/">
            <Arg>/yyyy_mm_dd.request.log</Arg>
            <Set name="retainDays">90</Set>
            <Set name="append">true</Set>
            <Set name="extended">false</Set>
            <Set name="LogTimeZone">GMT</Set>
          </New>
        </Set>
      </New>
    </Arg>
  </Call>
</Get>
```

Change the location of the logs to the desired location. In the settings above, location will default to `data/log` (same place where the log is located).

The log is using *National Center for Supercomputing Association Applications (NCSA)* or Common format (hence the class 'NCSARequestLog'). This is the most popular format for access logs and can be parsed by many web server analytics tools. Here is a sample output:

```
127.0.0.1 - - [14/Jan/2013:16:21:24 +0000] "GET /favicon.ico HTTP/1.1" 200 0
127.0.0.1 - - [14/Jan/2013:16:21:33 +0000] "GET /services/ HTTP/1.1" 200 0
127.0.0.1 - - [14/Jan/2013:16:21:33 +0000] "GET /services//?stylesheet=1 HTTP/1.1"
200 0
127.0.0.1 - - [14/Jan/2013:16:21:33 +0000] "GET /favicon.ico HTTP/1.1" 200 0
```

6.3.2.9. Using the LogViewer

Monitor system logs with the **LogViewer**, a convenient "viewing portal" for incoming logs.

- Navigate to the LogViewer at <https://{{FQDN}}:{{PORT}}/admin/logviewer>

or

- Navigate to the Admin Console
- Navigate to the **System** tab
- Select **Logs**

The LogViewer displays the most recent 500 log messages by default, but will grow to a maximum of 5000 messages. To view incoming logs, select the **PAUSED** button to toggle it to **LIVE** mode. Switching this back to **PAUSED** will prevent any new logs from being displayed in the LogViewer. Note that this only affects the logs displayed by the LogViewer and does not affect the underlying log.

Log events can be filtered by:

- Log level (**ERROR**, **WARNING**, etc).
 - The LogViewer will display at the currently configured log level for the Karaf logs.
 - See [Controlling Log Level](#) to change log level.
- Log message text.
- Bundle generating the message.

It is not recommended to use the LogViewer if the system logger is set to a low reporting level such as **TRACE**. The volume of messages logged will exceed the polling rate, and incoming logs may be missed.

WARNING

The actual logs being polled by the LogViewer can still be accessed at `<DDF_HOME>/data/log`

NOTE

The LogViewer settings don't change any of the underlying logging settings, only which messages are displayed. It does not affect the logs generated or events captured by the system logger.

6.4. Troubleshooting

If, after configuration, a DDF is not performing as expected, consult this table of common fixes and workarounds.

Table 18. General Troubleshooting

Issue	Solution
Unable to unzip distribution on Windows platform	The default Windows zip utility is not compatible with the DDF distribution zip file. Use Java or a third-party zip utility.
Unable to federate on Windows Platform	Windows default firewall is not compatible with DDF.
Ingesting more than 200,000 data files stored NFS shares may cause Java Heap Space error (Linux-only issue).	<p>This is an NFS bug where it creates duplicate entries for some files when doing a file list. Depending on the OS, some Linux machines can handle the bug better and able get a list of files but get an incorrect number of files. Others would have a Java Heap Space error because there are too many file to list.</p> <p>As a workaround, ingest files in batches smaller than 200,000.</p>
Ingesting serialized data file with scientific notation in WKT string causes RuntimeException.	WKT string with scientific notation such as POINT (-34.8932113039107 -4.77974239601E-5) won't ingest. This occurs with serialized data format only.
<p>Exception Starting DDF (Windows)</p> <p>An exception is sometimes thrown starting DDF on a Windows machine (x86).</p> <p>If using an unsupported terminal, <code>java.lang.NoClassDefFoundError: Could not initialize class org.fusesource.jansi.internal.Kernel32</code> is thrown.</p>	<p>Install missing Windows libraries.</p> <p>Some Windows platforms are missing libraries that are required by DDF. These libraries are provided by the Microsoft Visual C++ 2008 Redistributable Package x64.</p>
<p>CXF BusException</p> <p>The following exception is thrown: <code>org.apache.cxf.BusException: No conduit initiator</code></p>	<p>Restart DDF. Shut down DDF: <code>ddf@local>shutdown</code>. Start up DDF: <code>./ddf</code></p>

Issue	Solution
<p>Distribution Will Not Start</p> <p>DDF will not start when calling the start script defined during installation.</p>	<p>Complete the following procedure.</p> <ol style="list-style-type: none"> Verify that Java is correctly installed. <code>java -version</code> This should return something similar to: <code>java version "1.8.0_45" Java™ SE Runtime Environment (build 1.8.0_45-b14) Java HotSpot™ Server VM (build 25.45-b02, mixed mode)</code> If running *nix, verify that bash is installed. <code>echo \$SHELL</code> This should return: <code>/bin/bash</code>
<p>Multiple <code>java.exe</code> processes running, indicating more than one DDF instance is running.</p> <p>This can be caused when another DDF is not properly shut down.</p>	<p>Perform one or all of the following recommended solutions, as necessary.</p> <ul style="list-style-type: none"> Wait for proper shutdown of DDF prior to starting a new instance. Verify running <code>java.exe</code> are not DDF (e.g., kill/close if necessary). Utilize automated start/stop scripts to run DDF as a service.

6.4.1. Deleted Records Are Being Displayed In The Standard Search UI's Search Results

When queries are issued by the Standard Search UI, the query results that are returned are also cached in an internal Solr database for faster retrieval when the same query may be issued in the future. As records are deleted from the catalog provider, this Solr cache is kept in sync by also deleting the same records from the cache if they exist.

Sometimes the cache may get out of sync with the catalog provider such that records that should have been deleted are not. When this occurs, users of the Standard Search UI may see stale results since these records that should have been deleted are being returned from the cache. Records in the cache can be manually deleted using the URL commands listed below from a browser. In these command URLs, `metocard_cache` is the name of the Solr query cache.

- To delete all of the records in the Solr cache:

Deletion of all records in Solr query cache

```
\https://{{FQDN}}:{{PORT}}/solr/metocard_cache/update?stream.body=<delete><query>*:*</query></delete>&commit=true
```

- To delete a specific record in the Solr cache by ID (specified by the `original_id_txt` field):

Deletion of record in Solr query cache by ID

```
\https://{{FQDN}}:{{PORT}}/solr/metocard_cache/update?stream.body=<delete><query>original_id_txt:50ffd32b21254c8a90c15fccfb98f139</query></delete>&commit=true
```

- To delete record(s) in the Solr cache using a query on a field in the record(s) - in this example, the `title_txt` field is being used with wildcards to search for any records with word `remote` in the title:

Deletion of records in Solr query cache using search criteria

```
\https://{{FQDN}}:{{PORT}}/solr/metocard_cache/update?stream.body=<delete><query>title_txt :*remote*</query></delete>&commit=true
```

7. Data Management

7.1. File Formats Supported

DDF supports a wide variety of file types and data types for ingest. The DDF's internal Input Transformers extract the necessary data into a [generalized format](#). DDF supports ingest of many datatypes and commonly used file formats, such as Microsoft office products: Word documents, Excel spreadsheets, and PowerPoint presentations as well as .pdf files, GeoJson and others. See [complete list](#). Many of these file types support additional [format-specific attribute mappings](#) from which additional metadata can be extracted.

NOTE

These attributes will be available in all the specified file formats; however, values will only be present if present in the original document/resource.

These attributes are supported by any file type ingested into DDF:

Common Attributes in All File Types

- metadata
- id
- modified (date)
- title (filename)
- metadata content type (mime type)
- effective (date)

- created (date)

These 'media' file types have support for additional attributes to be available when ingested into DDF:

File Types Supporting Additional Attributes

- Video Types
 - WMV
 - AVI
 - MP4
 - MOV
 - h.264 MPEG2
- Image Types
 - JPEG-2000
- Document Types
 - .DOC, .DOCX, .DOTX, .DOCM
 - .PPT, .PPTX
 - .XLS, .XLSX
 - .PDF

These are the attributes common to any of the media file types which support additional attributes:

Additional Possible Attributes Common to 'Media' File Types

- `media.format-version`
- `media.format`
- `media.bit-rate`
- `media.bits-per-sample`
- `media.compression`
- `media.encoding`
- `media.frame-center`
- `media.frame-rate`
- `media.height-pixels`
- `media.number-of-bands`
- `media.scanning-mode`
- `media.type`
- `media.duration`
- `media.page-count`
- `datatype`
- `description`
- `contact.point-of-contact-name`
- `contact.contributor-name`

- `contact.creator-name`
- `contact.publisher-name`
- `contact.point-of-contact-phone`
- `topic.keyword`

7.2. Ingesting Data

Ingesting is the process of getting metocard(s) into the Catalog Framework. Ingested files are "transformed" into a neutral format that can be searched against as well as migrated to other formats and systems. There are multiple methods available for ingesting files into the DDF.

Guest Claims Attributes and Ingest

NOTE Ensure that appropriate [Guest Claims](#) are configured to allow guest users to ingest data and query the catalog.

7.2.1. Methods of Ingest

There are many methods of ingest supported by the DDF to support different needs.

7.2.1.1. Methods for ingesting small sets of records

These methods are the simplest for small data sets or for testing purposes.

7.2.1.1.1. Ingest Command (Command Console)

The DDF console application has a command-line option for ingesting files

The syntax for the ingest command is `ingest -t <transformer type> <file path>` relative to the installation path.

For more information on available transformers see the [Input Transformers](#) or run the `catalog:transformers` command.

For metocard XML data, run this command:

```
ingest -t xml examples/metacards/xml
```

NOTE Ingesting with the console ingest command will only create and store a metocard and not the resource itself.

7.2.1.1.2. User Interface Ingest

Files can also be ingested directly from [Intrigue](#).

7.2.1.2. Methods for ingesting large sets of records

These methods are useful for slightly larger data sets or infrequent ingestions.

7.2.1.2.1. Directory Monitor Ingest

The Catalog application contains a Directory Monitor feature that allows files placed in a single directory to be monitored and ingested automatically. For more information about configuring a directory to be monitored, see [Configuring the Content Directory Monitor](#).

Simply place the desired files in the monitored directory and it will be ingested automatically. If, for any reason, the files cannot be ingested, they will be moved to an automatically created sub-folder named `.errors`. Optionally, ingested files can be automatically moved to a sub-folder called `.ingested`.

7.2.1.2.2. External Methods

Several third-party tools, such as [cURL.exe](#) and the [Chrome Advanced Rest Client](#), can be used to send files and other types of data to DDF for ingest.

Windows Example

```
curl -H "Content-type: application/json;id=geojson" -i -X POST -d  
@"C:\path\to\geojson_valid.json" \https://{{FQDN}}:{{PORT}}/services/catalog
```

**NIX Example*

```
curl -H "Content-type: application/json;id=geojson" -i -X POST -d @geojson_valid.json  
\https://{{FQDN}}:{{PORT}}/services/catalog
```

Where:

-H adds an HTTP header. In this case, Content-type header `application/json;id=geojson` is added to match the data being sent in the request.

-i requests that HTTP headers are displayed in the response.

-X specifies the type of HTTP operation. For this example, it is necessary to POST (ingest) data to the server.

-d specifies the data sent in the POST request. The `@` character is necessary to specify that the data is a file.

The last parameter is the URL of the server that will receive the data.

This should return a response similar to the following (the actual catalog ID in the id and Location URL fields will be different):

Sample Response

```
HTTP/1.1 201 Created  
Content-Length: 0  
Date: Mon, 22 Apr 2015 22:02:22 GMT  
id: 44dc84da101c4f9d9f751e38d9c4d97b  
Location: \https://{{FQDN}}:{{PORT}}/services/catalog/44dc84da101c4f9d9f751e38d9c4d97b  
Server: Jetty(7.5.4.v20111024)
```

1. Verify the entry was successfully ingested by entering in a browser the URL returned in the POST response's HTTP header. For instance in our example, it was </services/catalog/44dc84da101c4f9d9f751e38d9c4d97b>. This should display the catalog entry in XML within the browser.
2. Verify the catalog entry exists by executing a query via the OpenSearch endpoint.
3. Enter the following URL in a browser `/services/catalog/query?q=ddf`. A single result, in Atom format, should be returned.

A resource can also be ingested with metocard metadata associated with it using the multipart/mixed content type.

Example

```
curl -k -X POST -i -H "Content-Type: multipart/mixed" -F
parse.resource=@/path/to/resource -F parse.metadata=@/path/to/metocard
\https://{{FQDN}}:{PORT}/services/catalog
```

7.2.1.2.3. Other Methods

The DDF provides endpoints for both REST and SOAP services, allowing integration with other data systems and the ability to further automate ingesting data into the catalog. See [Available Endpoints](#) for more information.

7.2.2. Verifying Ingest by REST Endpoint

Verify file was stored using the Content REST endpoint.

- Send a GET command to read the content from the content repository using the Content REST endpoint. This can be done using [cURL](#) command below. Note that the GUID will be different for each ingest. The GUID can be determined by going to the `<DDF_HOME>/content/store` directory and copying the sub-directory in this folder (there should only be one).

Windows Example

```
curl -X GET \https://{{FQDN}}:{PORT}/services/content/c90147bf86294d46a9d35ebbd44992c5
```

**NIX Example*

```
curl -X GET \https://{{FQDN}}:{PORT}/services/content/c90147bf86294d46a9d35ebbd44992c5
```

The response to the GET command will be the contents of the file originally ingested.

7.3. Removing Expired Records from the Catalog

DDF has many ways to remove expired records from the underlying Catalog data store. Nevertheless, the benefits of data standardization is that an attempt can be made to remove records without the need to know any vendor-specific information. Whether the data store is a search

server, a No-SQL database, or a relational database, expired records can be removed universally using the Catalog API and the Catalog Commands.

7.3.1. Manual Removal of Expired Records

To manually remove expired records from the Catalog, execute in the Command Console:

```
catalog:removeall --expired
```

When prompted, type yes to remove all expired records.

TIP For help on the removeall command, execute
`help catalog:removeall`

7.3.2. Automated Removal of Expired Records

By default, the DDF runs a scheduled command every 24 hours to remove expired records. The command is executed and scheduled [Using the Command Scheduler](#). To change the configuration out of the box, follow the [Updating a Scheduled Command](#) instructions. If an administrator wants to create additional scheduled tasks to remove records from the local Catalog, the administrator can follow the steps provided in the Scheduling a Command section. In the Command text field, type the following:

```
catalog:removeall --force --expired
```

If it is intended to have this run daily, type 86400 for the amount of seconds. (60 seconds/min x 60 minutes/hr x 24 hours/day = 86400 seconds for one day)

7.3.2.1. Explanation of Command to Remove Expired Records

The `catalog:removeall` command states you want to remove records from the local Catalog.

The `--force` option is used to suppress the confirmation message which asks a user if the user intentionally wants to permanently remove records from the Catalog.

The `--expired` option is to remove only expired records.

IMPORTANT If the `--expired` option is omitted, then all records will be removed from the Catalog.

7.3.3. Non-Universal or Catalog Specific Removal

Using the Catalog Commands is convenient for achieving many goals such as removing expired records, but is not always the most efficient since not all Catalog implementation details are known. The Catalog API does not allow for every special nuance of a specific data store. Therefore, whether an administrator's data store is from Oracle, Solr, or any other vendor, the administrator should consult the specific Catalog implementation's documentation on the best method to remove metadata.

7.3.4. Automatic Catalog Backup

To backup local catalog records, a [Catalog Backup Plugin](#) is available. It is not installed by default for performance reasons.

It can be installed and configured in the **Admin Console** (see [Catalog Backup Plugin](#) for instructions).

7.3.5. DDF Data Migration

Data migration is the process of moving metadata from one catalog provider to another. It is also the process of translating metadata from one format to another. Data migration is necessary when a user decides to use metadata from one catalog provider in another catalog provider.

The process for changing catalog providers involves first exporting the metadata from the original catalog provider and ingesting it into another.

7.3.5.1. Exporting Metadata Out of a Catalog Provider

1. Configure a desired catalog provider.
2. From the command line of DDF console, use the command to export all metadata from the catalog provider into serialized data files dump. The following example shows a command for running on Linux and a command for running on Windows.

ddf@local

```
dump "/myDirectory/exportFolder"  
or  
dump "C:/myDirectory/exportFolder"
```

7.3.5.2. Ingesting Exported Metadata into a Catalog Provider

1. Configure a different catalog provider.
2. From the command line of DDF console, use the **ingest** command to import exported metadata from serialized data files into catalog provider. The following example shows a command for running on Linux and a command for running on Windows.

ddf@local

```
ingest -p "/myDirectory/exportFolder"  
or  
ingest -p "C:/myDirectory/exportFolder"
```

7.3.5.3. Translating Metadata from One Format to Another

Metadata can be converted from one data format to another format. Only the data format changes, but the content of the metadata does not, as long as **option -p** is used with the **ingest** command. The process for converting metadata is performed by ingesting a data file into a catalog provider in one

format and dumping it out into a file in another format.

Using

These user interfaces are available in DDF.

Using the Landing Page

Using the Landing Page.

Using Intrigue

Using Intrigue.

Using Standard Search UI

Using Standard Search UI.

Using the Simple Search

Using the Simple Search user interface.

8. Using the Landing Page

The DDF Landing Page is the starting point for using DDF. It is accessible at <https://{{FQDN}}:{{PORT}}>.

8.1. Search DDF Button

The search button navigates to the Search UI, enabling catalog queries.

8.2. Data Source Availability

The data source availability pane provides a quick glance at the status of configured data sources.

8.3. Announcements

The announcements pane contains messages from system administrators.

9. Using Intrigue

Introduction: Intrigue represents the most advanced search interface available with DDF. It provides metadata search and discovery, resource retrieval, and workspace management with a 3D or optional 2D map visualization.

NOTE

For more detail on any feature or button within Intrigue, click the **?** icon in the upper right of the screen; then, hover over any item on the screen and a contextual tooltip will be displayed to define its purpose. To exit this mode, click the **?** again or press **escape**.

9.1. Accessing Intrigue

The default URL for Intrigue is <https://{}{FQDN}:{PORT}/search/catalog>

Catalog UI Guest Users

NOTE

If Guest access has been enabled, users not signed in to DDF (i.e. guest users) will have access to search functions, but all workspace configuration and settings will only exist locally and will not be available for sharing.

The default view for Intrigue is the **Workspaces** view. For other views or to return to the **Workspaces** view, click the **Navigation** menu in the upper-left corner of Intrigue and select the desired view.




Figure 1. Select the desired view from the Navigation menu.

9.2. Workspaces in Intrigue


Within Intrigue, workspaces are collections of settings, searches, and bookmarks that can be shared between users and stored for repeated access.

9.2.1. Creating a Workspace in Intrigue

Before searching in DDF, at least one workspace must be created.

Start new workspace

1. From the **Workspaces** view, enter search terms into the **Start new workspace** search field and



click the magnifying glass () icon. This will create a new workspace and perform a search based on the entered search terms.




Figure 2. Start new workspace

Using a template

1. From the **Workspaces** view, click on an existing template.
2. Change the workspace title by clicking on the temporary workspace title in the upper left corner and entering a new title.
3. Click the save () icon next to the workspace title in the upper left corner.




Figure 3. Default Workspace Templates

Blank

A blank workspace with no default search.

Local

An example of a local search.

Federated

An example of a search across all federated sources and the local Catalog.

Location

An example of a geographically constrained search.

Temporal

An example of a time-range search.

9.2.2. Configuring a Workspace in Intrigue

Configure each workspace with searches and share options.

Adding searches

1. From the default **Workspaces** view, select the workspace to add a search to.
2. Click **Search DDF Intrigue** in the upper left corner, enter search terms, and click **Search** to add a search. This step can be repeated to add additional searches. Each workspace can have up to ten searches.
 - a. Select **Basic** Search to select simple search criteria, such as **text**, **time**, and **location**.
 - b. Select **Advanced** Search to access a query builder for more complex queries.
3. Click the save () icon next to the workspace title in the upper left corner.

Navigation Menu Options

- **Workspaces**: View all available workspaces.
- **Upload**: Add new metadata and resources to the catalog.
- **Sources**: Lists all sources and their statuses.
- **Open Workspaces**: Lists open workspaces.

Workspace Menu Options

- To view a workspace's options from the **Workspaces** view, press the **Options** button () for the workspace.
 - **Save**: Save changes to the workspace.
 - **Run All Searches**: Start all saved searches within this workspace.
 - **Cancel All Searches**: Cancel all running searches.
 - **Open in New Tab**: Opens this workspace in a separate tab.
 - **View Sharing**: View and edit settings for sharing this workspace. Users must be signed in to share workspaces or view shared workspaces.
 - **View Details**: View the current details for a cloud-based workspace. Users must be signed in to view workspace details.
 - **Duplicate**: Create a copy of this workspace.
 - **Subscribe/Unsubscribe**: Selecting **Subscribe** will enable email notifications for search results on this workspace. Selecting **Unsubscribe** will disable email notifications for search results on this workspace.
 - **Move to Trash**: Delete (archive) this workspace.

9.2.3. Sharing Workspaces

Workspaces can be shared between users at different levels of access as needed.

Share a Workspace

1. From the **Workspaces** view, select the **Options** menu (⋮) for the workspace in which sharing will be modified.
2. Select **View Sharing**.
 - a. To share by user role, set the drop-down menu to **Can Access** for each desired role. All users with that role will be able to view the workspace.
 - b. To share with an individual user, add his/her email to the email list.
3. Click **Apply**.

Remove Sharing on a Workspace

1. From the **Workspaces** view, select the **Options** menu (⋮) for the workspace in which sharing will be modified.
2. Select **View Sharing**.
 - a. To remove the workspace from users with specific roles, set the drop-down menu to **No Access** for those roles.
 - b. To remove individual users, remove the users' email addresses from the email list.
3. Click **Apply**.

9.3. Ingesting from Intrigue

Data can be ingested via Intrigue.

1. Select the Menu icon (☰) in the upper left corner.
2. Select **Upload**.
3. Drag and drop file(s) or click to open a navigation window.
4. After selecting the file(s) to be uploaded, select **Start** to begin uploading.

9.4. Searching with Intrigue

The Search pane has two tabs: **Search** and **Bookmarks**.




Figure 4. Search Pane Tabs

9.4.1. Search Tab

View and edit searches from the **Search** tab.

The available searches for a workspace can be viewed by clicking on the drop-down on the **Search** tab.




Figure 5. Viewing available searches.

Search Menu Options

At the bottom of each search is a list of options for the search.

- **Run:** Trigger this search to begin immediately.
- **Edit:** Edits the search criteria.
- **Settings:** Edits the search settings, such as sorting.
- **Notifications:** Allows setting up search notifications.
- **Stop:** Stop this search.
- **Delete:** Remove this search.
- **Duplicate:** Create a copy of this search as a starting point.
- **Search Archived:** Execute this search, but specifically for archived results.
- **Search Historical:** Execute this search, but specifically for historical results.

9.4.1.1. Editing a Search

An existing search can be updated by selecting the search in the **Search** tab of a workspace and by clicking the **Edit** () icon.

- **Basic:** Define a [Text](#), [Spatial](#), [Temporal](#), or [Type](#) Search.

Text Search Details

Searches across all textual data of the targeted data source. Text search capabilities include:

Search for an exact word, such as `Text = apple` : Returns items containing the word "apple" but not "apples". Matching occurs on word boundaries.

Search for the existence of items containing multiple words, such as `Text = apple orange` : Returns items containing both "apple" and "orange" words. Words can occur anywhere in an item's metadata.

Search using wildcards, such as `Text = foo*` : Returns items containing words like "food", "fool", etc..

Text searches are by default case insensitive, but case sensitive searches are an option.

Temporal Search Details

Search based on absolute time of the created, modified, or effective date.

Any: Search without any time restrictions (default).

After: Search records after a specified time.

Before: Search records before a specified time.

Between: Set a beginning and end time to search between.

Relative: Search records relative to the current time.

Spatial Search Details

Search by latitude/longitude, USNG/MGRS, or UTM using a line, polygon, point-radius, or bounding box. Spatial criteria can also be defined by entering a **Keyword** for a region, country, or city in the **Location** section of the query builder.

Type Search Details

Search for specific content types.

- **Advanced:** Advanced query builder to create more specific searches.

Advanced Query Builder Details

Operator: If 'AND' is used, all the filters in the branch have to be true for this branch to be true. If 'OR' is used, only one of the filters in this branch has to be true for this branch to be true.

Property: Property to compare against.

Comparison: How to compare the value for this property against the provided value.

Search Terms: The value for the property to use during comparison.

Sorting: Sort results by relevance, distance, created time, modified time or effective time.

Sources: Perform an enterprise search (the local Catalog and all federated sources) or search specific sources.

- **Settings:** Set preferences for **Sorting** and **Federation**.
- **Notifications:** Define the frequency at which to run this search and receive notifications for new results.
- **Status:** View details of recent instances of this search.

9.4.1.1.1. Editing Search Settings

An existing search's settings can be modified by selecting the search in the **Search** tab of a workspace and by clicking the **Settings** (⚙️) icon. Sorting and sources can be customized here.

9.4.1.1.2. Editing Search Notifications

An existing search's notifications can be modified by selecting the search in the **Search** tab of a workspace and by clicking the **Notifications** (🔔) icon. Notification frequency can be customized here.

9.4.1.1.3. Viewing Search Status

An existing search's status can be viewed by selecting the search in the **Search** tab of a workspace and by clicking the **Status** (心脏病) icon. The **Status** view for a search displays information about the sources searched.

Intersecting Polygon Searches

If a self intersecting polygon is used to perform a geographic search, the polygon will be converted into a non-intersection one via a convex hull conversion. In the example below the blue line shows the original self intersecting search polygon and the red line shows the converted polygon that will be used for the search. The blue dot shows a search result that was not within the original polygon but was returned because it was within the converted polygon.

NOTE



Figure 6. Self Intersecting Polygon Conversion Example

9.4.1.2. Refining Search Results

Returned search results can be refined further, bookmarked, and/or downloaded from the **Search** tab. Result sets are color-coded by source as a visual aid. There is no semantic meaning to the colors assigned.



Figure 7. Search Results Options

1. On the **Search** tab, select a search from the drop-down list.
2. Perform any of these actions on the results list of the selected search:
 - a. Filter the result set locally. This does not re-execute the search.

- b. Customize results sorting. The default sort is by title in ascending order.
- c. Toggle results view between **List** and **Gallery**.

9.4.1.3. Search Result Options

Options for each individual search result

- **Download:** Downloads the result's associated product directly to the local machine. This option is only available for results that have products.
- **Bookmark:** Adds/removes the results to/from the saved bookmarks.
- **Hide from Future Searches:** Adds to a list of results that will be hidden from future searches.
- **Expand Metocard View:** Navigates to a view that only focuses on this particular result.
- **Create Search from Location:** Searches for all records that intersect the current result's location geometry.

9.4.2. Bookmarks Tab



Bookmarked results are denoted with a yellow star () icon.

1. Perform any of these actions on bookmarked results:
 - a. Filter the result set locally (does not re-execute the search),
 - b. Customize results sorting (Default: Title in Ascending Order).
 - c. Toggle results view between **List** and **Gallery**.

Bookmarked Result Options


- **Download:** Downloads the result's associated product directly to the local machine. This option is only available for results that have products.
- **Remove from Bookmarks:** Removes the result from bookmarks.
- **Hide from Future Searches:** Adds to a list of results that will be hidden from future searches.
- **Expand Metocard View:** Navigates to a view that only focuses on this particular result.
- **Create Search from Location:** Searches for all records that intersect the current result's location geometry.


9.5. Viewing Search Results

9.5.1. Adding Visuals

Visuals are different ways to view search results.

1.



Click the **Add Visual** () icon in the bottom right corner of Intrigue.

2. Select a visual to add.

- a. **2D Map:** A 2 dimensional map view.
- b. **3D Map:** A 3 dimensional map view.
- c. **Inspector:** In depth details and actions for the results of a search.
- d. **Histogram:** A configurable histogram view for the results of a search.
- e. **Table:** A configurable table view for the results of a search.

The **Search** tab displays a list of all of the search results for the selected search. The **Inspector** visual provides in depth information and actions for each search result.

Summary

A summarized view of the result.

Details

A detailed view of the result.

History

View revision history of this record.

Associations

View or edit the relationship(s) between this record and others in the catalog.

Quality

View the completeness and accuracy of the metadata for this record.

Actions

Export the metadata/resource to a specific format.

Archive

Remove the selected result from standard search results.

Overwrite

Overwrite a resource.

9.5.2. Editing Records

Results can be edited from the **Summary** or **Details** tabs in the **Inspector** visual.

9.5.3. Editing Associations on a Record

Update relationships between records through **Associations**.

1. Select the desired result from the **Search** tab.
2. Select the **Inspector** visual.
3. Select the **Associations** tab.
4. Select **Edit**.
5. For a new association, select **Add Association**. Only items in the current result set can be added

as associations.

- a. Select the related result from either the **Parent** or **Child** drop-down.
 - b. Select the type of relationship from the **Relationship** drop-down.
 - c. Select **Save**.
6. To edit an existing association, update the selections from the appropriate drop-downs and select **Save**.

View a graphical representation of the associations by selecting **Graph** icon from the **Associations** menu.




Figure 8. Associations menu.

9.5.4. Viewing Revision History

View the complete revision history of a record.

1. Select the desired result from the **Search** tab.
2. Select the **Inspector** visual.
3. Select the **History** tab.
 - a. Select a previous version from the list.
 - b. Select **Revert to Selected Version** to undo changes made after that revision.

9.5.5. Viewing Metadata Quality

View and fix issues with metadata quality in a record.

NOTE Correcting metadata issues may require administrative permissions.

1. Select the desired result from the **Search** tab.
2. Select the **Inspector** visual.
3. Select the **Quality** tab.
4. A report is displayed showing any issues:
 - a. Metacard Validation Issues.
 - b. Attribute Validation Issues.

9.5.6. Exporting a Result

Export a result's metadata and/or resource.

1. Select the desired result from the **Search** tab.
2. Select the **Inspector** visual.
3. Select **Actions** tab.
4. Select the desired export format.
5. Export opens in a new browser tab. Save, if desired.

9.5.7. Archiving a Result

To remove a result from the active search results, archive it.

1. Select the desired result from the **Search** tab.
2. Select the **Inspector** visual.
3. Select the **Archive** tab.
4. Select **Archive item(s)**.
5. Select **Archive**.


9.5.8. Restoring Archived Results

Restore an archived result to return it to the active search results.

1. Select the **Search Archived** option from the **Search Results Options** menu.
2. Select the desired result from the **Search** tab.
3. Select the **Inspector** visual.
4. Select the **Archive** tab.
5. Select **Restore item(s)**.
6. Select **Restore**.

Restore hidden results to the active search results.

1. Select the **Settings** () icon on navigation bar.
2. Select **Hidden**.
3. Click on the eye () icon next to each result to be unhidden.
 - a. Or select **Unhide All** to clear the list.



9.5.9. Overwriting a Resource

Replace a resource.

1. Select the desired result from the **Search** tab.
2. Select the **Inspector** visual.
3. Select the **Overwrite** tab.
4. Select **Overwrite content**.
5. Select **Overwrite**.
6. Navigate to the new content via the navigation window.

9.5.10. Intrigue Settings

Customize the look and feel of Intrigue using the **Settings** (⚙️) menu on the navigation bar.




Figure 9. Settings Menu Options

- **Theme:** Visual options for page layout.
- **Notifications:** Select if notifications persist across sessions.
- **Map:** Select options for map layers.
- **Query:** Customize the number of search results returned.
- **Time:** Set the time format (24 hour or 12 hour).
- **Hidden:** View or edit a list of results that have been hidden from the current search results.

9.5.11. Intrigue Notifications



Notifications can be checked/dismissed by clicking the **Notifications** icon () on the navigation bar.

10. Using Standard Search UI

The DDF Standard Search UI application allows a user to search for records in the local Catalog (provider) and federated sources. Results of the search are returned in HTML format and are displayed on a globe, providing a visual representation of where the records were found.

Located at the bottom of the left pane of the Search UI are two tabs: [Search](#) and [Workspaces](#). The Search tab contains basic fields to query the Catalog and other sources. The workspaces feature uses the same search criteria that are provided in the Search tab, and it also allows the user to create custom searches that can be saved for later execution. The right-side pane displays a map that, when records are found, shows the location of where the record was retrieved.

10.1. Search

The Standard Search UI allows users to search for records in the local Catalog and federated sources based on the criteria entered in the Search tab. After a search is executed, the UI provides results based on the defined criteria and detailed information about each result. Additionally, a user can save individual records that were returned by the query to a workspace, so they can be referenced in the future. The user can also save the search criteria to a workspace so the query can be run again.

10.1.1. Search Criteria

The Search UI queries a Catalog using the following criteria.

Criteria	Description
Text	Search by free text using the grammar of the underlying endpoint/Catalog.
Time	Search based on relative or absolute time of the created, modified, or effective date.
Location	Search by latitude/longitude or the USNG using a point-radius or bounding box.
Type	Search for specific content types.
Sorting	Sort results by relevance, distance, created time, modified time or effective time.
Additional Sources	Select All Sources or Specific Sources .
All Sources	Create an enterprise search. All federations are queried.
Specific Sources	Search a specific source(s). If a source is unavailable, it is displayed in red text.

10.1.2. Results

After a query is executed, the records matching the search criteria are automatically displayed in the Results pane.

Item	Description
Search criteria	Enhanced search toggle. Enables the enhanced search menu (see below). Allows users to filter and refine search and build more sophisticated queries.

Item	Description
Results	The number of records that were returned as a result of the query. Only the top 250 results are displayed, with the most relevant records displayed at the top of the list. If more than 250 results are returned, try narrowing the search criteria.
Search button	Navigates back to the Search pane.
Save button	Allows the user to select individual records to save.
Records list	Shows the results of the search. The following information is displayed for each record: <ul style="list-style-type: none"> Title – The title of the record is displayed in blue text. Select the title of the record to view more details. Source – The gray text displayed below the record title is the data source (e.g., ddf.distribution) and the amount of time that has passed since the record was modified (e.g., an hour ago).

10.1.3. Enhanced Search

The enhanced search menu allows more granular filtering of results and the ability to construct sophisticated queries.

Item	Description
Source	List of all sources searched with check boxes to allow users to refine searches to the most relevant sources.
Metadata Content Type	List of metadata content types found in the search, with the ability to select and deselect content type.
Query	The Query builder enables users to construct a very granular search query. The first drop-down menu contains the metadata elements in the search results, and the second contains operators based on the field selected (greater than, equals, contains, matchcase, before, after, etc.) Click the + to add further constraints to the query, or x to remove. Click Search to use the new query.
Search button	Executes an enhanced search on any new parameters that were specified or the query built above.

10.1.4. Record Summary

When an individual record is selected in the results list, the Record pane opens. When the Summary button is selected in the Record pane, the following information is displayed.

Item	Description
Results button	Navigates back to the original query results list.
Up and down arrows	Navigate through the records that were returned in the search. When the end or the beginning of the search results list is reached, the respective up or down arrow is disabled.

Item	Description
Details button	Opens the Details tab, which displays more information about the record.
Title	The title of the record is displayed in white font.
Source	The location that the metadata came from, which could be the local provider or a federated source.
Created time	When the record was created.
Modified time	Time since the record was last modified.
Locate button	Centers the map on the record's originating location.
Thumbnail	Depicts a reduced-size image of the original artifact for the current record, if available.
Download	A link to download the record. The size, if known, is indicated.

10.1.5. Record Details

When an individual record is selected in the results list, the Record pane opens. When the Details button is selected in the Record pane, the following information is displayed.

Item	Description
Results button	Navigates back to the original query results list.
Up and down arrows	Navigate through the records that were returned in the search. When the end or the beginning of the search results list is reached, the respective up or down arrow is disabled.
Summary button	Opens the Summary tab, which provides a high level overview of the result set.
Id	The record's unique identifier.
Source Id	Where the metadata was retrieved from, which could be the local provider or a federated source.
Title	The title of the record is displayed in white font.
Thumbnail	Depicts a reduced size image of the original artifact for the current record, if available.
Resource URI	Identifies the stored resource within the server.
Created time	When the record was created.
Metocard Content Type version	The version of the metadata associated with the record.
Metocard Type	The type of metocard associated with the record.

Item	Description
Metocard Content Type	The type of the metadata associated with the record.
Resource size	The size of the resource, if available.
Modified	Time since the record was last modified.
Download	When applicable, a download link for the product associated with the record is displayed. The size of the product is also displayed, if available. If the size is not available, N/A is displayed.
Metadata	Shows a representation of the metadata XML, if available.

10.2. Actions

Depending on the contents of the metocard, various actions will be available to perform on the metadata.

Troubleshooting: if no actions are available, ensure IP address is configured correctly under global configuration in Admin Console.

10.2.1. Save a Search

Saved searches are search criteria that are created and saved by a user. Each saved search has a name that was defined by the user, and the search can be executed at a later time or be scheduled for execution. Bookmarked records that the user elected to save for future use are returned as part of a search. These queries can be saved to a [workspace](#), which is a collection of searches and records created by a user. Complete the following procedure to create a saved search.

1. Select the Search tab at the bottom of the left pane.
2. Use the fields provided to define the [Search Criteria](#) for the query to be saved.
3. Select the **Save** button. The Select Workspace pane opens.
4. Type a name for the query in the **ENTER NAME FOR SEARCH** field.
5. Select a workspace in which to save the query, or create a workspace by typing a title for the new workspace in the **New Workspace** field.
6. Select the **Save** button.

The size of the product is based on the value in the associated metocard's resource-size attribute. This is defined when the metocard was originally created and may or may not be accurate. Often it will be set to N/A, indicating that the size is unknown or not applicable.

NOTE

However, if the administrator has enabled caching on DDF, and has installed the [catalog-core-resourcesizeplugin](#) PostQuery Plugin, and if the product has been retrieved, it has been cached and the size of the product can be determined based on the cached file's size. Therefore, subsequent query results that include that product will display an accurate size under the download link.

10.3. Workspaces

Each user can create multiple workspaces and assign each of them a descriptive name. Each workspace can contain multiple [saved searches](#) and contain multiple saved records. Workspaces are saved for each user and are loaded when the user logs in. Workspaces and their contents are persisted, so they survive if DDF is restarted. Within the Standard Search UI, workspaces are private and cannot be viewed by other users.

10.3.1. Create a Workspace

1. Select the Workspaces tab at the bottom of the Search UI's left pane. The Workspaces pane opens, which displays the existing workspaces that were created by the user. At the top of the pane, an option to **Add** and an option to **Edit** are displayed.
2. Select the **Add** button at the top of the left pane. A new workspace is created.
3. In the **Workspace Name** field, enter a descriptive name for the workspace.
4. Select the **Add** button. The Workspaces pane opens, which now displays the new workspace and any existing workspaces.
5. Select the name of the new workspace. The data (i.e., saved searches and records) for the selected workspace is displayed in the Workspace pane.
6. Select the + icon near the top of the Workspace pane to begin adding queries to the workspace. The Add/Edit Search pane opens.
7. Enter a name for the new query to be saved in the **QUERY NAME** field.
8. Complete the rest of the [Search Criteria](#).
9. Select the **Save & Search** button. The Search UI begins searching for records matching the criteria, and the new query is saved to the workspace. When the search is complete, the Workspace pane opens.
10. Select the name of the search to view the query results.
11. If necessary, in the Workspace pane, select the **Edit** button then select the pencil icon next to the name of a query to change the search criteria.
12. If necessary, in the Workspace pane, select the delete icon next to the name of a query to delete the query from the workspace.

10.4. Notifications

The Standard Search UI receives all notifications from DDF. These notifications appear as pop-up windows inside the Search UI to alert the user of an event of interest. To view all notifications, select the notification icon.

Currently, the notifications provide information about product retrieval only. After a user initiates a resource download, they receive periodic notifications that provide the progress of the download (e.g., the download completed, failed, or is being retried).

NOTE

A notification pop-up remains visible until it is dismissed or the browser is refreshed. Once a notification is dismissed, it cannot be retrieved again.

10.5. Activities

Similar to notifications, activities appear as pop-up windows inside the Search UI. Activity events include the status and progress of actions that are being performed by the user, such as searches and downloads. To view all activities, select the activity icon in the top-right corner of the window. A list of all activities opens in a drop-down menu, from which activities can be read and deleted. If a download activity is being performed, the Activity drop-down menu provides the link to retrieve the product.

If caching is enabled, a progress bar is displayed in the Activity (Product Retrieval) drop-down menu until the action being performed is complete.

10.6. Downloads

Downloads from the UI are currently managed by the user-specific browser's download manager. The UI itself does not have a built-in download manager utility.

10.7. Maps

The right side of the Search UI contains a map to locate search results on. There are three views for this map, 3D, 2D, and Columbus View. To choose a different view, select the map icon in the upper right corner. (The icon will change depending on current view selected.)

11. Using the Simple Search

The DDF Simple Search UI application provides a low-bandwidth option for searching records in the local Catalog (provider) and federated sources. Results are returned in HTML format.

11.1. Search

The **Input** form allows the user to specify keyword, geospatial, temporal, and type query parameters. It also allows the user to select the sources to search and the number of results to return.

11.1.1. Search Criteria

Enter one or more of the available search criteria to execute a query:

Keyword Search

A text box allowing the user to enter a textual query. This supports the use of (*) wildcards. If blank, the query will contain a contextual component.

Temporal Query

Select from **any**, **relative**, or **absolute**. Selecting **Any** results in no temporal restrictions on the query, selecting **relative** allows the user to query a period from some length of time in the past until now, and selecting **absolute** allows the user to specify a **start** and **stop** date range.

Spatial Search

Select from **any**, **point-radius**, and **bounding box**. Selecting **Any** results in no spatial restrictions on the query, selecting **point-radius** allows the user to specify a **lat/lon** and **radius** to search, and selecting a **bounding box** allows the user to specify an **eastern**, **western**, **southern** and **northern** boundary to search within.

Type Search

Select from **any**, or a specific type. Selecting **Any** results in no type restrictions on the query, and Selecting **Specific Types** shows a list of known content types on the federation, and allows the user to select a specific type to search for.

Sources

Select from **none**, **all sources**, or **specific sources**. Selecting **None** results in querying only the local provider, Selecting **All Sources** results in an enterprise search where all federations are queried, and selecting **Specific Sources** allows the user to select which sources are queried.

Results per Page

Select the number of results to be returned by a single query.

11.1.2. Results

The table of results shows the details of the results found, as well as a link to download the product if applicable.

11.1.2.1. Results Summary

Total Results

Total Number of Results available for this query. If there are more results than the number displayed per page then a page navigation links will appear to the right.

Pages

Provides page navigation, which generate queries for requesting additional pages of results.

11.1.2.2. Results Table

The Results table provides a preview of and links to the results. The table consists of these columns:

Title

Displays title of the metocard. This will be a link which can be clicked to view the metocard in the Metocard View.

Source

Displays where the metadata came from, which could be the local provider or a federated

source.

Location

Displays the WKT Location of the metocard, if available.

Time

Shows the Received (Created) and Effective times of the metocard, if available.

Thumbnail

Shows the thumbnail of the metocard, if available.

Download

A download link to retrieve the product associated with the metocard, when applicable, if available.

11.1.3. Result View

This view shows more detailed look at a result.

Back to Results Button

Returns the view back to the Results Table.

Previous & Next

Navigation to page through the results one by one.

Result Table

Provides the list of properties and associated values of a single search result.

Metadata

The metadata, when expanded, displays a tree structure representing the result's custom metadata.


Integrating

WARNING

If integrating with a Highly Available Cluster of DDF, see [High Availability Guidance](#).


Integrators will use the existing applications to support their external frameworks. This section will provide details for finding, accessing and using the components of DDF.

12. Federation




Federation

The Catalog normalizes incoming queries into an OGC Filter format. When the query is disseminated by the Catalog Framework to the sources, each source is responsible for denormalizing the OGC-Filter-formatted query into the format understood by the external store being proxied by that source. This normalization/denormalization is what allows any endpoint to interface with any type of source. For example, a query received by the OpenSearch Endpoint can be executed against an OpenSearch Source.



Federation Data Flow

13. Data and Metadata



Catalog Architecture Diagram: Data

The Catalog stores and translates Metadata, which can be transformed into many data formats, shared, and queried. The primary form of this metadata is the metocard. A **Metocard** is a container for metadata. **CatalogProviders** accept **Metacards** as input for ingest, and **Sources** search for metadata and return matching **Results** that include **Metacards**.

13.1. Metacards

A metocard is a single instance of metadata in the Catalog (an instance of a metocard type) which generally contains general information about the product, such as the title of the product, the product's geo-location, the date the product was created and/or modified, the owner or producer, and/or the security classification.

13.1.1. Metocard Type

A metocard type indicates the attributes available for a particular metocard. It is a model used to define the attributes of a metocard, much like a schema.

A metocard type indicates the attributes available for a particular type of data. For example, an image may have different attributes than a PDF document, so each could be defined to have their own metocard type.

13.1.1.1. Default Metocard Type and Attributes

Most metacards within the system are created using the default metocard type or a metocard type based on the default type. The default metocard type of the system can be programmatically retrieved by calling `ddf.catalog.data.BasicTypes.BASIC_METACARD`. The name of the default **MetacardType** can be retrieved from `ddf.catalog.data.MetacardType.DEFAULT_METACARD_TYPE_NAME`.

The default metocard type has the following required attributes. Though the following attributes are required on all metocard types, setting their values is optional except for ID.

Core Attributes

NOTE

It is highly recommended when referencing a default attribute name to use the `ddf.catalog.data.types.*` interface constants whenever possible. Mapping to a normalized taxonomy allows for higher quality transformations between different formats and for improved federation. This neutral profile facilitates improved search and discovery across disparate data types.

WARNING

Every `Source` should at the very least return an ID attribute according to Catalog API. Other fields may or may not be applicable, but a unique ID must be returned by a source.

13.1.1.2. Extensible Metacards

Metacard extensibility is achieved by creating a new `MetacardType` that supports attributes in addition to the required attributes listed above.

Required attributes must be the base of all extensible metacard types.

WARNING

Not all `Catalog Providers` support extensible metacards. Nevertheless, each Catalog Provider should at least have support for the default `MetacardType`; i.e., it should be able to store and query on the attributes and attribute formats specified by the default metacard type. Catalog providers are neither expected nor required to store attributes that are not in a given metacard's type.

Consult the documentation of the Catalog Provider in use for more information on its support of extensible metacards.

Often, the `BASIC_METACARD MetacardType` does not provide all the functionality or attributes necessary for a specific task. For performance or convenience purposes, it may be necessary to create custom attributes even if others will not be aware of those attributes. One example could be if a user wanted to optimize a search for a date field that did not fit the definition of `CREATED`, `MODIFIED`, `EXPIRATION`, or `EFFECTIVE`. The user could create an additional `java.util.Date` attribute in order to query the attribute separately.

`Metacard` objects are extensible because they allow clients to store and retrieve standard and custom key/value `Attributes` from the `Metacard`. All `Metacards` must return a `MetacardType` object that includes an `AttributeDescriptor` for each `Attribute`, indicating its key and value type. `AttributeType` support is limited to those types defined by the Catalog.

New `MetacardType` implementations can be made by implementing the `MetacardType` interface.

13.1.2. Metacard Type Registry

WARNING

The `MetacardTypeRegistry` is experimental. While this component has been tested and is functional, it may change as more information is gathered about what is needed and as it is used in more scenarios.

The `MetacardTypeRegistry` allows DDF components, primarily catalog providers and sources, to make available the `MetacardTypes` that they support. It maintains a list of all supported `MetacardTypes` in the `CatalogFramework`, so that other components such as `Endpoints`, `Plugins`, and `Transformers` can make use of those `MetacardTypes`. The `MetacardType` is essential for a component in the `CatalogFramework` to understand how it should interpret a metocard by knowing what attributes are available in that metocard.

For example, an endpoint receiving incoming metadata can perform a lookup in the `MetacardTypeRegistry` to find a corresponding `MetacardType`. The discovered `MetacardType` will then be used to help the endpoint populate a metocard based on the specified attributes in the `MetacardType`. By doing this, all the incoming metadata elements can then be available for processing, cataloging, and searching by the rest of the `CatalogFramework`.

`MetacardTypes` should be registered with the `MetacardTypeRegistry`. The `MetacardTypeRegistry` makes those `MetacardTypes` available to other DDF `CatalogFramework` components. Other components that need to know how to interpret metadata or metacards should look up the appropriate `MetacardType` from the registry. By having these `MetacardTypes` available to the `CatalogFramework`, these components can be aware of the custom attributes.

The `MetacardTypeRegistry` is accessible as an OSGi service. The following blueprint snippet shows how to inject that service into another component:

MetacardTypeRegistry Service Injection

```
<bean id="sampleComponent" class="ddf.catalog.SampleComponent">
    <argument ref="metacardTypeRegistry" />
</bean>

<!-- Access MetacardTypeRegistry -->
<reference id="metacardTypeRegistry" interface="ddf.catalog.data.MetacardTypeRegistry"
"/>
```

The reference to this service can then be used to register new `MetacardTypes` or to lookup existing ones.

Typically, new `MetacardTypes` will be registered by `CatalogProviders` or sources indicating they know how to persist, index, and query attributes from that type. Typically, `Endpoints` or `InputTransformers` will use the lookup functionality to access a `MetacardType` based on a parameter in the incoming metadata. Once the appropriate `MetacardType` is discovered and obtained from the registry, the component will know how to translate incoming raw metadata into a DDF Metocard.

13.1.3. Attributes

An attribute is a single field of a metocard, an instance of an attribute type. Attributes are typically indexed for searching by a source or catalog provider.

13.1.3.1. Attribute Types

An attribute type indicates the attribute format of the value stored as an attribute. It is a model for

an attribute.

13.1.3.1.1. Attribute Format

An enumeration of attribute formats are available in the catalog. Only these attribute formats may be used.

Table 19. Attribute Formats

AttributeFormat	Description
BINARY	Attributes of this attribute format must have a value that is a Java <code>byte[]</code> and <code>AttributeType.getBinding()</code> should return <code>Class<Array>of byte</code> .
BOOLEAN	Attributes of this attribute format must have a value that is a Java boolean.
DATE	Attributes of this attribute format must have a value that is a Java date.
DOUBLE	Attributes of this attribute format must have a value that is a Java double.
FLOAT	Attributes of this attribute format must have a value that is a Java float.
GEOMETRY	Attributes of this attribute format must have a value that is a WKT-formatted Java string.
INTEGER	Attributes of this attribute format must have a value that is a Java integer.
LONG	Attributes of this attribute format must have a value that is a Java long.
OBJECT	Attributes of this attribute format must have a value that implements the serializable interface.
SHORT	Attributes of this attribute format must have a value that is a Java short.
STRING	Attributes of this attribute format must have a value that is a Java string and treated as plain text.
XML	Attributes of this attribute format must have a value that is a XML-formatted Java string.

13.1.3.1.2. Attribute Naming Conventions

Catalog taxonomy elements follow the naming convention of `group-or-namespace.specific-term`, except for extension fields outside of the core taxonomy. These follow the naming convention of `ext.group-or-namespace.specific-term` and must be namespaced. Nesting is not permitted.

13.1.3.2. Result

A single "hit" included in a query response.

A result object consists of the following:

- a metocard.
- a relevance score if included.
- distance in meters if included.

13.1.4. Creating Metacards

The quickest way to create a `Metocard` is to extend or construct the `MetocardImpl` object. `MetocardImpl` is the most commonly used and extended `Metocard` implementation in the system because it provides a convenient way for developers to retrieve and set `Attributes` without having to create a new `MetocardType` (see below). `MetocardImpl` uses `BASIC_METACARD` as its `MetocardType`.

13.1.4.1. Limitations

A given developer does not have all the information necessary to programmatically interact with any arbitrary source. Developers hoping to query custom fields from extensible `Metacards` of other sources cannot easily accomplish that task with the current API. A developer cannot question a source for all its *queryable* fields. A developer only knows about the `MetocardTypes` which that individual developer has used or created previously.

The only exception to this limitation is the `Metocard.ID` field, which is required in every `Metocard` that is stored in a source. A developer can always request `Metacards` from a source for which that developer has the `Metocard.ID` value. The developer could also perform a wildcard search on the `Metocard.ID` field if the source allows.

13.1.4.2. Processing Metacards

As `Metocard` objects are created, updated, and read throughout the Catalog, care should be taken by all catalog components to interrogate the `MetocardType` to ensure that additional `Attributes` are processed accordingly.

13.1.4.3. Basic Types

The Catalog includes definitions of several basic types all found in the `ddf.catalog.data.BasicTypes` class.

Table 20. Basic Types

Name	Type	Description
<code>BASIC_METACARD</code>	<code>MetocardType</code>	Represents all required Metocard Attributes.
<code>BINARY_TYPE</code>	<code>AttributeType</code>	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.BINARY</code> .
<code>BOOLEAN_TYPE</code>	<code>AttributeType</code>	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.BOOLEAN</code> .

Name	Type	Description
DATE_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.DATE</code> .
DOUBLE_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.DOUBLE</code> .
FLOAT_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.FLOAT</code> .
GEO_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.GEOMETRY</code> .
INTEGER_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.INTEGER</code> .
LONG_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.LONG</code> .
OBJECT_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.OBJECT</code> .
SHORT_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.SHORT</code> .
STRING_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.STRING</code> .
XML_TYPE	AttributeType	A Constant for an <code>AttributeType</code> with <code>AttributeType.AttributeFormat.XML</code> .

13.2. JSON "Definition" Files.

WARNING

This section concerns capabilities that are considered experimental. The features described in this section may change or be removed in a future version of the application.

DDF supports adding new attribute types, metocard types, validators, and more using json-formatted definition files.

The following may be defined in a JSON definition file:

- [Attribute Types](#)
- [Metocard Types](#)

- [Global Attribute Validators](#)
- [Default Attribute Values](#)
- [Attribute Injections](#)

13.2.1. Definition File Format

A definition file follows the JSON format as specified in [ECMA-404](#). All definition files must be valid JSON in order to be parsed.

A single definition file may define as many of the types as needed. This means that types can be defined across multiple files for grouping or clarity.

13.2.2. Deploying Definition Files

The file must have a `.json` extension in order to be picked up by the deployer. Once the definition file is ready to be deployed, put the definition file `<filename>.json` into the `etc/definitions` folder.

Definition files can be added, updated, and/or deleted in the `etc/definitions` folder. The changes are applied dynamically and no restart is required.

If a definition file is removed from the `etc/definitions` folder, the changes that were applied by that file will be undone.

13.3. Data Validation

DDF can be configured to perform validation on ingested documents to verify the integrity of the metadata brought into the catalog.

Available Validation Services

Validation with Schematron

Introduction to validation with schematron.

13.3.1. Validation with Schematron

DDF uses [Schematron Validation](#) to validate metadata ingested into the catalog.

Custom schematron rulesets can be used to validate metocard metadata. Multiple services can be created, and each service can have multiple rulesets associated with it. Namespaces are used to distinguish services. The root schematron files may be placed anywhere on the file system as long as they are configured with an absolute path. Any root schematron files with a relative path are assumed to be relative to `<DDF_HOME>/schematron`.

TIP

Schematron files may reference other schematron files using an include statement with a relative path. However, when using the document function within a schematron ruleset to reference another file, the path must be absolute or relative to the DDF installation home directory.


13.3.1.1. Configuring Schematron Services

Schematron validation services are configured with a namespace and one or more schematron rulesets. Additionally, warnings may be suppressed so that only errors are reported.

To create a new service:

- Navigate to the **Admin Console**.
- Select the **Catalog**.
- Select **Configuration**.
- Ensure that **catalog-schematron-plugin** is started.
- Select **Schematron Validation Services**.

14. Endpoints



Endpoints

Endpoints act as a proxy between the client and the [Catalog Framework](#).

14.1. Available Endpoints

The following endpoints are available in a standard installation of DDF:

[Application Upload Endpoint](#)

Enables uploading new and upgraded applications to the system.

[Catalog REST Endpoint](#)

Allows clients to perform CRUD operations on the Catalog using REST, a simple architectural style that performs communication using HTTP.

CometD Endpoint

Enables asynchronous search capabilities.

CSW Endpoint

Searches collections of descriptive information (metadata) about geospatial data and services.

FTP Endpoint

Ingests files directly into the DDF Catalog using the FTP protocol.

KML Endpoint

Generates a view-based KML Query Results Network Link.

Metrics Endpoint

Reports on system metrics.

OpenSearch Endpoint

Sends query parameters and receives search results.

14.1.1. Application Upload Endpoint

The Application Upload Endpoint enables uploading new and upgraded applications to the system.

14.1.1.1. Installing Application Upload Endpoint

The Application Upload Endpoint is installed by default with a standard installation as part of the Admin application.

14.1.1.2. Application Upload Endpoint

The Application endpoint is available at <<https://{{FQDN}}:{{PORT}}/services/application>>.

14.1.1.3. Configuring Application Upload Endpoint

No configuration necessary.

14.1.2. Catalog REST Endpoint

The Catalog REST Endpoint allows clients to perform CRUD operations on the Catalog using REST, a simple architectural style that performs communication using HTTP.

14.1.2.1. Installing the Catalog REST Endpoint

The Catalog REST Endpoint allows clients to perform CRUD operations on the Catalog using REST, a simple architectural style that performs communication using HTTP.

The URL exposing the REST functionality is located at <https://{{FQDN}}:{{PORT}}/services/catalog>.

14.1.2.2. Installing the Catalog REST Endpoint

The Catalog REST Endpoint is installed by default with a standard installation in the Catalog

application.

14.1.2.3. Configuring the Catalog REST Endpoint

The RESTful CRUD Endpoint has no configurable properties. It can only be installed or uninstalled.

14.1.2.4. Using the Catalog REST Endpoint

The Catalog REST Endpoint provides the capability to query, create, update, and delete metacards and associated resource in the catalog provider as follows:

Operation	HTTP Request	Details	Example URL
create	HTTP POST	<p>HTTP request body contains the input to be ingested.</p> <p><code><input transformer></code> is the name of the transformer to use when parsing metadata (optional).</p>	<code>\http://{FQDN}:{PORT}/services/catalog?transformer=<input transformer></code>
update	HTTP PUT	<p>The ID of the Metacard to be updated is appended to the end of the URL. The updated metadata is contained in the HTTP body.</p> <p><code><metacardId></code> is the Metacard.ID of the metacard to be updated and <code><input transformer></code> is the name of the transformer to use when parsing an override metadata attribute (optional).</p>	<code>\http://{FQDN}:{PORT}/services/catalog/<metacardId>?transformer=<input transformer></code>

Operation	HTTP Request	Details	Example URL
<code>delete</code>	HTTP DELETE	<p>The ID of the Metacard to be deleted is appended to the end of the URL.</p> <p><code><metacardId></code> is the Metacard.ID of the metacard to be deleted.</p>	<code>\http://{FQDN}:{PORT}/services/catalog/<metacardId></code>
<code>read</code>	HTTP GET	<p>The ID of the Metacard to be retrieved is appended to the end of the URL.</p> <p>By default, the response body will include the XML representation of the Metacard.</p> <p><code><metacardId></code> is the Metacard.ID of the metacard to be retrieved.</p>	<code>\http://{FQDN}:{PORT}/services/catalog/<metacardId></code>
<code>federated read</code>	HTTP GET	<p>The SOURCE ID of a federated source is appended to the URL before the ID of the Metacard to be retrieved is appended to the end.</p> <p><code><sourceId></code> is the FEDERATED SOURCE ID and <code><metacardId></code> is the Metacard.ID of the Metacard to be retrieved.</p>	<code>\http://{FQDN}:{PORT}/services/catalog/sources/<sourceId>/<metacardId></code>
<code>sources</code>	HTTP GET	<p>Retrieves information about federated sources, including <code>sourceId</code>, <code>availability</code>, <code>contentTypes</code>, and <code>version</code>.</p>	<code>\http://{FQDN}:{PORT}/services/catalog/sources/</code>

14.1.2.4.1. Sources Operation Example

In the example below there is the local DDF distribution and a DDF OpenSearch federated source with id "DDF-OS".

Sources Response Example

```
[  
  {  
    "id" : "DDF-OS",  
    "available" : true,  
    "contentTypes" :  
      [  
        ],  
    "version" : "2.0"  
  },  
  {  
    "id" : "ddf.distribution",  
    "available" : true,  
    "contentTypes" :  
      [  
        ],  
    "version" : "2.5.0-SNAPSHOT"  
  }  
]
```

Note that for all RESTful CRUD commands only one metocard ID is supported in the URL, i.e., bulk operations are not supported.

14.1.2.5. Interacting with the REST CRUD Endpoint

Any web browser can be used to perform a REST read. Various other tools and libraries can be used to perform the other HTTP operations on the REST endpoint (e.g., soapUI, cURL, etc.)

The REST endpoint can be used to upload resources as attachments. The [create](#) and [update](#) methods both support the multipart mime format. If only a single attachment exists, it will be interpreted as a resource to be parsed, which will result in a metocard and resource being stored in the system.

If multiple attachments exist, then the REST endpoint will assume that 1 attachment is the actual resource (attachment should be named [parse.resource](#)) and the other attachments are overrides of metocard attributes (attachment names should follow metocard attribute names). In the case of the metadata attribute, it is possible to also have the system transform that metadata and use the results of that to override the metocard that would be generated from the resource (attachment should be named [parse.metadata](#)).

For example:

```

POST /services/catalog?transform=xml HTTP/1.1
Host: <FQDN>:<PORT>
Content-Type: multipart/form-data; boundary=----WebKitFormBoundary7MA4YWxkTrZu0gW
Cache-Control: no-cache

-----WebKitFormBoundary7MA4YWxkTrZu0gW
Content-Disposition: form-data; name="parse.resource"; filename=""
Content-Type:

-----WebKitFormBoundary7MA4YWxkTrZu0gW
Content-Disposition: form-data; name="parse.metadata"; filename=""
Content-Type:

-----WebKitFormBoundary7MA4YWxkTrZu0gW--

```

14.1.2.6. Metocard Transforms with the REST CRUD Endpoint

The **read** operation can be used to retrieve metadata in different formats.

1. Install the appropriate feature for the desired transformer. If desired transformer is already installed such as those that come out of the box (**xml, html, etc**), then skip this step.
2. Make a read request to the REST URL specifying the catalog id.
3. Add a transform query parameter to the end of the URL specifying the shortname of the transformer to be used (e.g., **transform=kml**).

Example Metocard Transform

```
\http://{FQDN}:{PORT}/services/catalog/<metacardId>?transform=<TRANSFORMER_ID>
```

TIP Transforms also work on read operations for metacards in federated sources.
`https://{FQDN}:{PORT}/services/catalog/sources/<sourceId>/<metacardId>?transform=<TRANSFORMER_ID>`

See [Metocard Transformers](#) for details on metocard transformers.

14.1.2.6.1. POST Metadata

The following is a successful post of well-formed XML data sent to the Catalog ReST endpoint.

Example Metacard

```
<?xml version="1.0" encoding="UTF-8"?>
<metacard xmlns="urn:catalog:metacard" xmlns:gml="http://www.opengis.net/gml"
  xmlns:xlink="http://www.w3.org/1999/xlink" xmlns:smil="http://www.w3.org/2001/SMIL20/"
  xmlns:smillang="http://www.w3.org/2001/SMIL20/Language"
  gml:id="3a59483ba44e403a9f0044580343007e">
  <type>ddf.metacard</type>
  <string name="title">
    <value>Test REST Metacard</value>
  </string>
  <string name="description">
    <value>Vestibulum quis ipsum sit amet metus imperdiet vehicula. Nulla scelerisque
    cursus mi.</value>
  </string>
</metacard>
```

14.1.2.6.2. Example Responses for ReST Endpoint Error Conditions

The following are example data and expected errors responses that will be returned for each error condition.

HTTP error codes are also returned. https://en.wikipedia.org/wiki/List_of_HTTP_status_codes#4xx_Client_errors

Malformed XML

The following request with malformed XML data sent to the Catalog ReST endpoint.

Malformed XML Example

```
<?xml version="1.0" encoding="UTF-8"?>
<metacard xmlns="urn:catalog:metacard" xmlns:gml="http://www.opengis.net/gml"
  xmlns:xlink="http://www.w3.org/1999/xlink" xmlns:smil="http://www.w3.org/2001/SMIL20/"
  xmlns:smillang="http://www.w3.org/2001/SMIL20/Language"
  gml:id="3a59483ba44e403a9f0044580343007e">
  <type>ddf.metacard</type>
  <string name="title">
    <value>Test REST Metacard</value>
  </string>
  <string name="description">
    <value>Vestibulum quis ipsum sit amet metus imperdiet vehicula. Nulla scelerisque
    cursus mi.</value>
  </string>
</document>
```

A HTTP 400 is returned and the following response body is returned. The specific error is logged in the error log.

Malformed XML Response

```
<pre>Error while storing entry in catalog: </pre>
```

Request with Unknown Schema

The following is a malformed XML document sent to the Catalog ReST endpoint.

Malformed XML Example

```
<?xml version="1.0" encoding="UTF-8"?>
<mydoc xmlns="http://example.com/unknown" xmlns:gml="http://www.opengis.net/gml"
  xmlns:xlink="http://www.w3.org/1999/xlink" xmlns:smil="http://www.w3.org/2001/SMIL20/"
  xmlns:smillang="http://www.w3.org/2001/SMIL20/Language"
  gml:id="3a59483ba44e403a9f0044580343007e">
  <type>ddf.metacard</type>
  <string name="title">
    <value>Test REST Metacard</value>
  </string>
  <string name="description">
    <value>Vestibulum quis ipsum sit amet metus imperdiet vehicula. Nulla scelerisque
    cursus mi.</value>
  </string>
</mydoc>
```

Creates a generic resource metacard with the provided XML as content for the [metadata](#) XML field in the metacard.

Request with Missing XML Prologue

The following is an example request with a missing XML prologue sent to the Catalog ReST endpoint.

Missing XML Tag Example

```
<metacard xmlns="urn:catalog:metacard" xmlns:gml="http://www.opengis.net/gml"
  xmlns:xlink="http://www.w3.org/1999/xlink" xmlns:smil="http://www.w3.org/2001/SMIL20/"
  xmlns:smillang="http://www.w3.org/2001/SMIL20/Language"
  gml:id="3a59483ba44e403a9f0044580343007e">
  <type>ddf.metacard</type>
  <string name="title">
    <value>Test REST Metacard</value>
  </string>
  <string name="description">
    <value>Vestibulum quis ipsum sit amet metus imperdiet vehicula. Nulla scelerisque
    cursus mi.</value>
  </string>
</metacard>
```

Metacard is created successfully

Request with Non-XML Data

The following is an example request with non-XML data sent to the Catalog ReST endpoint.

Non-XML data Example

```
title: Non-XML title
id: abc123
```

Metocard will be created and the content will be stored in the **metadata** field.

Request with Invalid Transform

Testing valid data with an invalid `?transform=invalid` appended to the POST URL: `{public_url}/services/catalog?transform=blah`

Valid data with an invalid ?transform=invalid

```
<?xml version="1.0" encoding="UTF-8"?>
<metocard xmlns="urn:catalog:metocard" xmlns:gml="http://www.opengis.net/gml"
  xmlns:xlink="http://www.w3.org/1999/xlink" xmlns:smil="http://www.w3.org/2001/SMIL20/"
  xmlns:smillang="http://www.w3.org/2001/SMIL20/Language"
  gml:id="3a59483ba44e403a9f0044580343007e">
  <type>ddf.metocard</type>
  <string name="title">
    <value>Test REST Metocard</value>
  </string>
  <string name="description">
    <value>Vestibulum quis ipsum sit amet metus imperdiet vehicula. Nulla scelerisque
    cursus mi.</value>
  </string>
</metocard>
```

A HTTP 400 is returned and the following response body is returned. The specific error is logged in the error log.

Malformed XML Response

```
<pre>Error while storing entry in catalog: </pre>
```

14.1.3. CometD Endpoint

The [CometD](#) endpoint enables asynchronous search capabilities. The CometD protocol is used to execute searches, retrieve results, and receive notifications.

For an example of using CometD within a webapp see: [distribution/sdk/sample-cometd/](#)

14.1.3.1. Installing CometD Endpoint

The CometD Endpoint is installed by default with a standard installation in the Search UI application.

14.1.3.2. Configuring CometD Endpoint

The CometD endpoint has no configurable properties.

14.1.3.3. Using CometD Endpoint

CometD Endpoint URL

```
\https://{{FQDN}}:{{PORT}}/search/cometd
```

14.1.3.3.1. CometD Queries

Queries can be executed over CometD using the `/service/query` channel. Query messages are JSON-formatted and use CQL alongside several other parameters.

Table 21. Query Parameters

Parameter Name	Description	Required
<code>src</code>	Comma-delimited list of federated sites to search over	No
<code>cql</code>	CQL query. See OpenGIS® Catalogue Services Specification for more information about CQL.	Yes
<code>sort</code>	Sort Type The format for the sort options is <code><Sort Field>:<Sort Order></code> , where <code><Sort Order></code> can be either <code>asc</code> or <code>desc</code> .	No
<code>id</code>	Query ID (Must be a unique ID, such as a UUID). This determines the channel that the query results will be returned on.	Yes
<code>count</code>	Number of entries to return in the response. Default is 10.	No
<code>start</code>	Specifies the number of the first result that should be returned.	No
<code>timeout</code>	Time (in milliseconds) to wait for response.	No

Before a query is published the client should subscribe to the channel that will be passed in to the `id` field in order to receive query results once the query is executed.

For example if the following id was generated `3b19bc9c-2155-4ca6-bae8-65a9c8e373f6`, the client should subscribe to `/3b19bc9c-2155-4ca6-bae8-65a9c8e373f6`

Then the following example query could be executed:

```
/service/query
```

```
{
  "src": "ddf.distribution",
  "cql": "(\"anyText\" ILIKE 'foo')",
  "id": "3b19bc9c-2155-4ca6-bae8-65a9c8e373f6",
  "sort": "asc"
}
```

This would return any results matching the text `foo` on the `/3b19bc9c-2155-4ca6-bae8-65a9c8e373f6` channel

14.1.3.3.2. Query Request Examples

Enterprise Contextual Query

```
"data": {  
  "count": 250,  
  "format": "geojson",  
  "id": "4303ba5d-21af-4878-9a4c-808e80052e6c",  
  "cql": "anyText LIKE '*'",  
  "start": 1  
}
```

Multiple Site Temporal Absolute Query

```
"data": {  
  "count": 250,  
  "format": "geojson",  
  "id": "4303ba5d-21af-4878-9a4c-808e80052e6c",  
  "cql": "modified DURING 2014-09-01T00:00:00Z/2014-09-30T00:00:00Z",  
  "src": "DDF-OS,ddf.distribution",  
  "start": 1  
}
```

Enterprise Spatial Bounding Box Query

```
"data": {  
  "count": 250,  
  "format": "geojson",  
  "id": "4303ba5d-21af-4878-9a4c-808e80052e6c",  
  "cql": "INTERSECTS(anyGeo, POLYGON ((-112.7786 32.2159, -112.7786 45.6441, -83.7297  
45.6441, -83.7297 32.2159, -112.7786 32.2159)))",  
  "start": 1  
}
```

14.1.3.3.3. Query Response Channel

The query responses are returned on the `/<id>` channel, which should be subscribed to in order to retrieve the results. Replace `<id>` with the id that was used in the request. The [Subscribing to Notifications](#) section details how to subscribe to a CometD channel.

The response is returned as a data map that contains an internal map with the following keys:

Table 22. Query Response Message Format

Map Key	Description	Value Type
<code>id</code>	ID that corresponds to the request.	String

Map Key	Description	Value Type
<code>hits</code>	Total number of query hits that were found by the server. Depending on the 'count' in the request, not all of the results may be returned.	Integer ≥ 0
<code>results</code>	Array of metocard results, formatted as defined by the GeoJSON Metocard Transformer.	Array of Maps
<code>results/metocard/is-resource-local</code>	A property indicating whether a metocard's associated resource is cached.	Boolean
<code>results/metocard/actions</code>	An array of actions that applies to each metocard, injected into each metocard containing an id, title, description, and url.	Array of Maps
<code>status</code>	Array of status for each source queried.	Array
<code>status.state</code>	Specifies the state of the query: SUCCEEDED, FAILED, ACTIVE.	String
<code>status.elapsed</code>	Time in milliseconds that it took for the source to complete the request.	Integer ≥ 0
<code>status.hits</code>	Number of records that were found on the source that matched the query	Integer ≥ 0
<code>status.id</code>	ID of the federated source	String
<code>status.results</code>	Number of results that were returned in this response from the source	Integer ≥ 0
<code>types</code>	A Map mapping a metocard-type's name to a map about that metocard-type. Only metocard-types represented by the metacards returned in the query are represented. The Map defining a particular <code>metocard-type</code> maps the fields supported by that <code>metocard-type</code> to the datatype for that particular field.	Map of Maps

14.1.3.3.4. Query Response Examples

Example Query Response

```
{
  "data": {
    "hits": 1,
    "metocard-types": {
      "ddf.metocard": {...}
    },
    "id": "6f0e04e9-acd1-4935-b9dd-c83e770a36d5",
    "results": [
      {
        "metocard": {
          "is-resource-local": false,
          "cached": "2016-07-13T19:22:18.220+0000",
          "geometry": {
            "coordinates": [
              -84.415337,
              37.774844
            ]
          }
        }
      }
    ]
  }
}
```

```

        42.729925
    ],
    "type": "Point"
},
"type": "Feature",
"actions": [...],
"properties": {
    "thumbnail": "...",
    "metadata": "<?xml version=\"1.0\" encoding=\"UTF-8
\"?><metadata>...</metadata>",
    "resource-size": "362417",
    "created": "2010-06-10T12:07:26.000+0000",
    "resource-uri": "content:faade630a2a247468ca9a9b57303b437",
    "metocard-tags": [
        "resource"
    ],
    "checksum-algorithm": "Adler32",
    "metadata-content-type": "image/jpeg",
    "metocard-type": "ddf.metocard",
    "resource-download-url": "
https://{FQDN}:{PORT}services/catalog/sources/ddf.distribution/faade630a2a247468ca9a9
b57303b437?transform=resource",
        "title": "example.jpg",
        "source-id": "ddf.distribution",
        "effective": "2016-07-13T19:22:06.966+0000",
        "point-of-contact": "",
        "checksum": "dc7337c5",
        "modified": "2010-06-10T12:07:26.000+0000",
        "id": "faade630a2a247468ca9a9b57303b437"
    }
}
},
],
"status": [
{
    "hits": 1,
    "elapsed": 453,
    "reasons": [],
    "id": "ddf.distribution",
    "state": "SUCCEEDED",
    "results": 1
}
],
"successful": true
},
"channel": "/6f0e04e9-acd1-4935-b9dd-c83e770a36d5"
},
{
    "successful": true
},
{

```

```

"channel": "/service/query",
"id": "142",
"successful": true
}

```

14.1.3.3.5. CometD Notifications

Notifications are messages that are sent to clients to inform them of some significant event happening. Clients must subscribe to a notification channel to receive these messages.

Notifications are published by the server on several notification channels depending on the type.

- subscribing to `/ddf/notifications/**` will cause the client to receive all notifications.
- subscribing to `/ddf/notifications/catalog/downloads` will cause the client to only receive notifications of downloads.

14.1.3.3.6. Using CometD Notifications

NOTE

The DDF Search UI serves as a reference implementation of how clients can use notifications.

Notifications are currently being utilized in the Catalog application for resource retrieval. When a user initiates a resource retrieval, the channel `/ddf/notification/catalog/downloads` is opened, where notifications indicating the progress of that resource download are sent. Any client interested in receiving these progress notifications must subscribe to that channel.

DDF starts downloading the resource to the client that requested it, a notification with a status of "Started" will be broadcast. If the resource download fails, a notification with a status of "Failed" will be broadcast. Or, if the resource download is being attempted again after a failure, "Retry" will be broadcast. When a notification is received, DDF Search UI displays a popup containing the contents of the notification, so a user is made aware of how their downloads are proceeding. Behind the scenes, the DDF Search UI invokes the REST endpoint to retrieve a resource.

In this request, it adds the query parameter "user" with the CometD session ID or the unique User ID as the value. This allows the CometD server to know which subscriber is interested in the notification.

For

example,

`https://[FQDN]:[PORT]/services/catalog/sources/ddf.distribution/2f5db9e5131444279a1293c541c106cd? transform=resource&user=1w1ql079j6tscii19jszwp9s2i55` notifications contain the following information:

Table 23. Notification Contents

Property Name	Description	Always Included with Notification
<code>application</code>	Name of the application that caused the notification to be sent.	Yes
<code>id</code>	ID of the notification "thread" – Notifications about the same event should use the same id to allow clients to filter out notifications that may be outdated.	Yes

Property Name	Description	Always Included with Notification
<code>title</code>	Resource/file name for resource retrieval.	Yes
<code>message</code>	Human-readable message containing status details.	Yes
<code>timestamp</code>	Timestamp in milliseconds when notification was sent.	Yes
<code>session</code>	CometD Session ID or unique User ID.	Yes

Example: Notification Message

```
"data": {
  "application": "Downloads",
  "title": "Product retrieval successful",
  "message": "The requested product was retrieved successfully
    and is available for download.",
  "id": "27ec3222af1144ff827a351b1962a236",
  "timestamp": "1403734355420",
  "user": "admin"
}
```

14.1.3.3.7. Receive Notifications

- If interested in retrieve resource notifications, a client must subscribe to the CometD channel `/ddf/notification/catalog/downloads`.
- If interested in all notification types, a client must subscribe to the CometD channel `/ddf/notification/**`
- A client will only receive notifications for resources they have requested.
- Standard UI is subscribed to all notifications of interest to that user/browser session: `/ddf/notification/**`
- See [Notification Contents](#) for the data that a notification contains.

14.1.3.3.8. Notification Events

Notifications are messages sent to clients to inform them of a significant event happening. Clients must subscribe to a notification channel to receive these messages.

14.1.3.3.9. Persistence of Notifications

Notifications are persisted between sessions, however due to the nature of CometD communications, they will not be visible at first connection/subscription to `/ddf/notifications/**`.

In order to retrieve notifications that were persisted or may have occurred since the previous session a client simply must publish an empty json message, `{}` to `/ddf/notifications`. This will return all existing notifications to the user.

14.1.3.3.10. Notification Operations Channel

Notification Operations are commands that change the behavior of future notifications. A notification operation is performed by publishing a list of commands to the CometD endpoint at

Table 24. Operation Format

Map Key	Description	Value Type
action	Type of action to request. If a client publishes with the <code>remove</code> action, it dismisses the notification and makes it unavailable again when notifications are retrieved. "remove" is currently only used action.	String
id	ID of the notification to which the action relates	String

Example: Notification Operation Request

```
"data": [ {
  "action": "remove",
  "id": "27ec3222af1144ff827a351b1962a236"
} ]
```

14.1.3.3.11. Activity Events Channel

To receive all activity updates, follow the instructions at [Subscribing to Notifications](#) and subscribe to `/ddf/activities/**`

Activity update messages follow a specific format when they are published to the activities channel. These messages contain a data map that encapsulates the activity information.

Table 25. CometD Activity Format

Property	Description	Value Type
category	Category of the activity	String
id	ID that uniquely identifies the activity that sent out the update. Not required to be unique per update.	String
message	User-readable message that explains the current activity status	String

Property	Description	Value Type
<code>operations</code>	<p>Map of operations that can be performed on this activity. If the value is a URL, the client should invoke the URL as a result of the user invoking the activity operation.</p> <p>If the value is not a URL, the client should send a message back to the server on the same topic with the operation name.</p> <p>Note: the DDF UI will interpret several values with special icons:</p> <ul style="list-style-type: none"> * <code>cancel</code> * <code>download</code> * <code>remove</code> 	JSON Map
<code>progress</code>	Percentage value of activity completion	String (Integer between 0 - 100 followed by a %)
<code>status</code>	Enumerated value that displays the current state of the activity	String + * <code>STARTED</code> * <code>RUNNING</code> * <code>COMPLETED</code> * <code>STOPPED</code> * <code>PAUSED</code> * <code>FAILED</code>
<code>timestamp</code>	Time that the activity update was sent	Date-Time
<code>title</code>	User-readable title for the activity update	String
<code>subject</code>	User who started the activity	String
<code>bytes</code>	Number of bytes the activity consumed (upload or download)	Positive Integer
<code>session</code>	The session ID of the user/subject	String
<code>Custom Value</code>	Additional keys can be inserted by the component sending the activity notification	Any JSON Type

Example: Activity update with custom 'bytes' field

```
data: {  
  "category": "Product Retrieval",  
  "id": "a62f6666-fc41-4a19-91f1-485e73a564b5",  
  "message": "The requested product is being retrieved. Standby.",  
  "operations": {  
    "cancel" : true  
  },  
  "progress": "45",  
  "status": "RUNNING",  
  "timestamp": "1403801920875",  
  "title": "Product downloading",  
  "user": "admin",  
  "bytes": 635084800  
}
```

14.1.3.3.12. Activity Operations Channel

Different operations can be performed on activities through the `/service/action` channel.

Table 26. CometD Activity Format

Map Key	Description	Value Type	action
The requested action. This value is based on the operations map that comes in from an activity event.	String * "cancel" * "download" * "remove"	id	ID of the activity to which the requested operation relates

Example: Activity Operation Request Message

```
"data": [ {  
  "action": "cancel",  
  "id": "a62f6666-fc41-4a19-91f1-485e73a564b5"  
} ]
```

14.1.4. CSW Endpoint

The CSW endpoint enables a client to search collections of descriptive information (metadata) about geospatial data and services.

14.1.4.1. Installing CSW Endpoint

The CSW Endpoint is installed by default with a standard installation in the Spatial application.

14.1.4.2. Configuring CSW Endpoint

The CSW endpoint has no configurable properties. It can only be installed or uninstalled.

14.1.4.3. CSW Endpoint URL

The CSW endpoint is accessible from <https://{{FQDN}}:{{PORT}}/services/csw>.

14.1.4.3.1. CSW Endpoint Operations

GetCapabilities Operation

The `GetCapabilities` operation is meant to describe the operations the catalog supports and the URLs used to access those operations. The CSW endpoint supports both `HTTP GET` and `HTTP POST` requests for the `GetCapabilities` operation. The response to either request will always be a `csw:Capabilities` XML document. This XML document is defined by the [CSW-Discovery XML Schema](#).

GetCapabilities HTTP GET

The `HTTP GET` form of `GetCapabilities` uses query parameters via the following URL:

GetCapabilities KVP (Key-Value Pairs) Encoding

```
\https://{{FQDN}}:{{PORT}}/services/csw?service=CSW&version=2.0.2&request=GetCapabilities
```

GetCapabilities HTTP POST

The `HTTP POST` form of `GetCapabilities` operates on the root CSW endpoint URL (<https://{{FQDN}}:{{PORT}}/services/csw>) with an XML message body that is defined by the `GetCapabilities` element of the [CSW-Discovery XML Schema](#).

GetCapabilities XML Request

```
<?xml version="1.0" ?>
<csw:GetCapabilities
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
  service="CSW"
  version="2.0.2" >
</csw:GetCapabilities>
```

GetCapabilities Sample Response (application/xml)

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Capabilities xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:gml=
"http://www.opengis.net/gml" xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
  xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
  "http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
  instance" version="2.0.2" ns10:schemaLocation="http://www.opengis.net/csw
  /ogc/csw/2.0.2/CSW-publication.xsd">
```

```

<ows:ServiceIdentification>
  <ows:Title>Catalog Service for the Web</ows:Title>
  <ows:Abstract>DDF CSW Endpoint</ows:Abstract>
  <ows:ServiceType>CSW</ows:ServiceType>
  <ows:ServiceTypeVersion>2.0.2</ows:ServiceTypeVersion>
</ows:ServiceIdentification>
<ows:ServiceProvider>
  <ows:ProviderName>DDF</ows:ProviderName>
  <ows:ProviderSite/>
  <ows:ServiceContact/>
</ows:ServiceProvider>
<ows:OperationsMetadata>
  <ows:Operation name="GetCapabilities">
    <ows:DCP>
      <ows:HTTP>
        <ows:Get ns2:href="\https://\{FQDN\}:\{PORT\}/services/csw"/>
        <ows:Post ns2:href="\https://\{FQDN\}:\{PORT\}/services/csw">
          <ows:Constraint name="PostEncoding">
            <ows:Value>XML</ows:Value>
          </ows:Constraint>
        </ows:Post>
      </ows:HTTP>
    </ows:DCP>
    <ows:Parameter name="sections">
      <ows:Value>ServiceIdentification</ows:Value>
      <ows:Value>ServiceProvider</ows:Value>
      <ows:Value>OperationsMetadata</ows:Value>
      <ows:Value>Filter_Capabilities</ows:Value>
    </ows:Parameter>
  </ows:Operation>
  <ows:Operation name="DescribeRecord">
    <ows:DCP>
      <ows:HTTP>
        <ows:Get ns2:href="\https://\{FQDN\}:\{PORT\}/services/csw"/>
        <ows:Post ns2:href="\https://\{FQDN\}:\{PORT\}/services/csw">
          <ows:Constraint name="PostEncoding">
            <ows:Value>XML</ows:Value>
          </ows:Constraint>
        </ows:Post>
      </ows:HTTP>
    </ows:DCP>
    <ows:Parameter name="typeName">
      <ows:Value>csw:Record</ows:Value>
      <ows:Value>gmd:MD_Metadata</ows:Value>
    </ows:Parameter>
    <ows:Parameter name="OutputFormat">
      <ows:Value>application/xml</ows:Value>
      <ows:Value>application/json</ows:Value>
      <ows:Value>application/atom+xml</ows:Value>
      <ows:Value>text/xml</ows:Value>
    </ows:Parameter>
  </ows:Operation>
</ows:OperationsMetadata>

```

```

<ows:Parameter name="schemaLanguage">
  <ows:Value>http://www.w3.org/XMLSchema</ows:Value>
  <ows:Value>http://www.w3.org/XML/Schema</ows:Value>
  <ows:Value>http://www.w3.org/2001/XMLSchema</ows:Value>
  <ows:Value>http://www.w3.org/TR/xmlschema-1/</ows:Value>
</ows:Parameter>
</ows:Operation>
<ows:Operation name="GetRecords">
  <ows:DCP>
    <ows:HTTP>
      <ows:Get ns2:href="\https://\{FQDN\}\{PORT\}/services/csw"/>
      <ows:Post ns2:href="\https://\{FQDN\}\{PORT\}/services/csw">
        <ows:Constraint name="PostEncoding">
          <ows:Value>XML</ows:Value>
        </ows:Constraint>
      </ows:Post>
    </ows:HTTP>
  </ows:DCP>
  <ows:Parameter name="ResultType">
    <ows:Value>hits</ows:Value>
    <ows:Value>results</ows:Value>
    <ows:Value>validate</ows:Value>
  </ows:Parameter>
  <ows:Parameter name="OutputFormat">
    <ows:Value>application/xml</ows:Value>
    <ows:Value>application/json</ows:Value>
    <ows:Value>application/atom+xml</ows:Value>
    <ows:Value>text/xml</ows:Value>
  </ows:Parameter>
  <ows:Parameter name="OutputSchema">
    <ows:Value>urn:catalog:metacard</ows:Value>
    <ows:Value>http://www.isotc211.org/2005/gmd</ows:Value>
    <ows:Value>http://www.opengis.net/cat/csw/2.0.2</ows:Value>
  </ows:Parameter>
  <ows:Parameter name="typeNames">
    <ows:Value>csw:Record</ows:Value>
    <ows:Value>gmd:MD_Metadata</ows:Value>
  </ows:Parameter>
  <ows:Parameter name="ConstraintLanguage">
    <ows:Value>Filter</ows:Value>
    <ows:Value>CQL_Text</ows:Value>
  </ows:Parameter>
  <ows:Constraint name="FederatedCatalogs">
    <ows:Value>Source1</ows:Value>
    <ows:Value>Source2</ows:Value>
  </ows:Constraint>
</ows:Operation>
<ows:Operation name="GetRecordById">
  <ows:DCP>
    <ows:HTTP>
      <ows:Get ns2:href="\https://\{FQDN\}\{PORT\}/services/csw"/>

```

```

<ows:Post ns2:href="\https://\{FQDN\}:\{PORT\}/services/csw">
    <ows:Constraint name="PostEncoding">
        <ows:Value>XML</ows:Value>
    </ows:Constraint>
</ows:Post>
</ows:HTTP>
</ows:DCP>
<ows:Parameter name="OutputSchema">
    <ows:Value>urn:catalog:metacard</ows:Value>
    <ows:Value>http://www.isotc211.org/2005/gmd</ows:Value>
    <ows:Value>http://www.opengis.net/cat/csw/2.0.2</ows:Value>
    <ows:Value>http://www.iana.org/assignments/media-
types/application/octet-stream</ows:Value>
</ows:Parameter>
<ows:Parameter name="OutputFormat">
    <ows:Value>application/xml</ows:Value>
    <ows:Value>application/json</ows:Value>
    <ows:Value>application/atom+xml</ows:Value>
    <ows:Value>text/xml</ows:Value>
    <ows:Value>application/octet-stream</ows:Value>
</ows:Parameter>
<ows:Parameter name="ResultType">
    <ows:Value>hits</ows:Value>
    <ows:Value>results</ows:Value>
    <ows:Value>validate</ows:Value>
</ows:Parameter>
<ows:Parameter name="ElementSetName">
    <ows:Value>brief</ows:Value>
    <ows:Value>summary</ows:Value>
    <ows:Value>full</ows:Value>
</ows:Parameter>
</ows:Operation>
<ows:Operation name="Transaction">
    <ows:DCP>
        <ows:HTTP>
            <ows:Post ns2:href="\https://\{FQDN\}:\{PORT\}/services/csw">
                <ows:Constraint name="PostEncoding">
                    <ows:Value>XML</ows:Value>
                </ows:Constraint>
            </ows:Post>
        </ows:HTTP>
    </ows:DCP>
    <ows:Parameter name="typeNames">
        <ows:Value>xml</ows:Value>
        <ows:Value>appxml</ows:Value>
        <ows:Value>csw:Record</ows:Value>
        <ows:Value>gmd:MD_Metadata</ows:Value>
        <ows:Value>tika</ows:Value>
    </ows:Parameter>
    <ows:Parameter name="ConstraintLanguage">
        <ows:Value>Filter</ows:Value>

```

```

        <ows:Value>CQL_Text</ows:Value>
    </ows:Parameter>
</ows:Operation>
<ows:Parameter name="service">
    <ows:Value>CSW</ows:Value>
</ows:Parameter>
<ows:Parameter name="version">
    <ows:Value>2.0.2</ows:Value>
</ows:Parameter>
</ows:OperationsMetadata>
<ogc:Filter_Capabilities>
    <ogc:Spatial_Capabilities>
        <ogc:GeometryOperands>
            <ogc:GeometryOperand>gml:Point</ogc:GeometryOperand>
            <ogc:GeometryOperand>gml:LineString</ogc:GeometryOperand>
            <ogc:GeometryOperand>gml:Polygon</ogc:GeometryOperand>
        </ogc:GeometryOperands>
        <ogc:SpatialOperators>
            <ogc:SpatialOperator name="BBOX"/>
            <ogc:SpatialOperator name="Beyond"/>
            <ogc:SpatialOperator name="Contains"/>
            <ogc:SpatialOperator name="Crosses"/>
            <ogc:SpatialOperator name="Disjoint"/>
            <ogc:SpatialOperator name="DWithin"/>
            <ogc:SpatialOperator name="Intersects"/>
            <ogc:SpatialOperator name="Overlaps"/>
            <ogc:SpatialOperator name="Touches"/>
            <ogc:SpatialOperator name="Within"/>
        </ogc:SpatialOperators>
    </ogc:Spatial_Capabilities>
    <ogc:Scalar_Capabilities>
        <ogc:LogicalOperators/>
        <ogc:ComparisonOperators>
            <ogc:ComparisonOperator>Between</ogc:ComparisonOperator>
            <ogc:ComparisonOperator>NullCheck</ogc:ComparisonOperator>
            <ogc:ComparisonOperator>Like</ogc:ComparisonOperator>
            <ogc:ComparisonOperator>EqualTo</ogc:ComparisonOperator>
            <ogc:ComparisonOperator>Greater Than</ogc:ComparisonOperator>
            <ogc:ComparisonOperator>Greater Than Equal To</ogc:ComparisonOperator>
            <ogc:ComparisonOperator>Less Than</ogc:ComparisonOperator>
            <ogc:ComparisonOperator>Less Than Equal To</ogc:ComparisonOperator>
            <ogc:ComparisonOperator>Equal To</ogc:ComparisonOperator>
            <ogc:ComparisonOperator>Not Equal To</ogc:ComparisonOperator>
        </ogc:ComparisonOperators>
    </ogc:Scalar_Capabilities>
    <ogc:Id_Capabilities>
        <ogc:EID/>
    </ogc:Id_Capabilities>
</ogc:Filter_Capabilities>
</csw:Capabilities>

```

DescribeRecord Operation

The **describeRecord** operation retrieves the type definition used by metadata of one or more registered resource types. There are two request types one for **GET** and one for **POST**. Each request has the following common data parameters:

Namespace

In **POST** operations, namespaces are defined in the xml. In **GET** operations, namespaces are defined in a comma separated list of the form: `xmlns([prefix=]namespace-url)(,xmlns([prefix=]namespace-url))*`

Service

The service being used, in this case it is fixed at CSW.

Version

The version of the service being used (2.0.2).

OutputFormat

The requester wants the response to be in this intended output. Currently, only one format is supported (application/xml). If this parameter is supplied, it is validated against the known type. If this parameter is not supported, it passes through and returns the XML response upon success.

SchemaLanguage

The schema language from the request. This is validated against the known list of schema languages supported (refer to <http://www.w3.org/XML/Schema>).

DescribeRecord HTTP GET

The **HTTP GET** request differs from the **POST** request in that the **typeName** is a comma-separated list of namespace prefix qualified types as strings (e.g., `csw:Record,xyz:MyType`). These prefixes are then matched against the prefix qualified namespaces in the request. This is converted to a list of QName(s). In this way, it behaves exactly as the post request that uses a list of QName(s) in the first place.

DescribeRecord KVP (Key-Value Pairs) Encoding

```
\https://{{FQDN}}:{{PORT}}/services/csw?service=CSW&version=2.0.2&request=DescribeRecord&NAMESPACES=xmlns(http://www.opengis.net/cat/csw/2.0.2)&outputFormat=application/xml&schemaLanguage=http://www.w3.org/XML/Schema
```

DescribeRecord HTTP POST

The **HTTP POST** request **DescribeRecordType** has the **typeName** as a List of QName(s). The QNames are matched against the namespaces by prefix, if prefixes exist.

DescribeRecord XML Request

```
<?xml version="1.0" ?>
<DescribeRecord
  version="2.0.2"
  service="CSW"
  outputFormat="application/xml"
  schemaLanguage="http://www.w3.org/XMLSchema"
  xmlns="http://www.opengis.net/cat/csw/2.0.2">
</DescribeRecord>
```

DescribeRecord Sample Response (application/xml)

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:DescribeRecordResponse xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:gml=
"http://www.opengis.net/gml" xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
"http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
instance" ns10:schemaLocation="http://www.opengis.net/csw /ogc/csw/2.0.2/CSW-
publication.xsd">
  <csw:SchemaComponent targetNamespace="http://www.opengis.net/cat/csw/2.0.2"
  schemaLanguage="http://www.w3.org/XMLSchema">
    <xsd:schema xmlns:xsd="http://www.w3.org/2001/XMLSchema" elementFormDefault=
"qualified" id="csw-record" targetNamespace="http://www.opengis.net/cat/csw/2.0.2"
  version="2.0.2">
      <xsd:annotation>
        <xsd:appinfo>
          <dc:identifier>
http://schemas.opengis.net/csw/2.0.2/record.xsd</dc:identifier>

          </xsd:appinfo>
          <xsd:documentation xml:lang="en">
This schema defines the basic record types that must be supported
by all CSW implementations. These correspond to full, summary, and
brief views based on DCMI metadata terms.
</xsd:documentation>

        </xsd:annotation>
        <xsd:import namespace="http://purl.org/dc/terms/" schemaLocation="rec-
dcterms.xsd"/>
        <xsd:import namespace="http://purl.org/dc/elements/1.1/" schemaLocation=
"rec-dcemes.xsd"/>
        <xsd:import namespace="http://www.opengis.net/ows" schemaLocation=
"../../ows/1.0.0/owsAll.xsd"/>
        <xsd:element abstract="true" id="AbstractRecord" name="AbstractRecord"
type="csw:AbstractRecordType"/>
        <xsd:complexType abstract="true" id="AbstractRecordType" name=
"AbstractRecordType"/>
    </xsd:schema>
  </csw:SchemaComponent>
</csw:DescribeRecordResponse>
```

```

<xsd:element name="DCMIRecord" substitutionGroup="csw:AbstractRecord"
type="csw:DCMIRecordType"/>
  <xsd:complexType name="DCMIRecordType">
    <xsd:annotation>
      <xsd:documentation xml:lang="en">
        This type encapsulates all of the standard DCMI metadata terms,
        including the Dublin Core refinements; these terms may be mapped
        to the profile-specific information model.
    </xsd:documentation>

    </xsd:annotation>
    <xsd:complexContent>
      <xsd:extension base="csw:AbstractRecordType">
        <xsd:sequence>
          <xsd:group ref="dct:DCMI-terms"/>

        </xsd:sequence>
      </xsd:extension>
    </xsd:complexContent>

  </xsd:complexType>
  <xsd:element name="BriefRecord" substitutionGroup="csw:AbstractRecord"
type="csw:BriefRecordType"/>
    <xsd:complexType final="#all" name="BriefRecordType">
      <xsd:annotation>
        <xsd:documentation xml:lang="en">
          This type defines a brief representation of the common record
          format. It extends AbstractRecordType to include only the
          dc:identifier and dc:type properties.
      </xsd:documentation>

      </xsd:annotation>
      <xsd:complexContent>
        <xsd:extension base="csw:AbstractRecordType">
          <xsd:sequence>
            <xsd:element maxOccurs="unbounded" minOccurs="1" ref=
"dc:identifier"/>
            <xsd:element maxOccurs="unbounded" minOccurs="1" ref=
"dc:title"/>
            <xsd:element minOccurs="0" ref="dc:type"/>
            <xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"ows:BoundingBox"/>

          </xsd:sequence>
        </xsd:extension>
      </xsd:complexContent>
    </xsd:complexType>
  </xsd:element>

```

```

</xsd:complexType>
<xsd:element name="SummaryRecord" substitutionGroup="csw:AbstractRecord"
type="csw:SummaryRecordType"/>
<xsd:complexType final="#all" name="SummaryRecordType">
<xsd:annotation>
<xsd:documentation xml:lang="en">
This type defines a summary representation of the common record
format. It extends AbstractRecordType to include the core
properties.
</xsd:documentation>

</xsd:annotation>
<xsd:complexContent>
<xsd:extension base="csw:AbstractRecordType">
<xsd:sequence>
<xsd:element maxOccurs="unbounded" minOccurs="1" ref=
"dc:identifier"/>
<xsd:element maxOccurs="unbounded" minOccurs="1" ref=
"dc:title"/>
<xsd:element minOccurs="0" ref="dc:type"/>
<xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:subject"/>
<xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:format"/>
<xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:relation"/>
<xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:modified"/>
<xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:abstract"/>
<xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:spatial"/>
<xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"ows:BoundingBox"/>

</xsd:sequence>
</xsd:extension>
</xsd:complexContent>

</xsd:complexType>
<xsd:element name="Record" substitutionGroup="csw:AbstractRecord" type=
"csw:RecordType"/>
<xsd:complexType final="#all" name="RecordType">
<xsd:annotation>
<xsd:documentation xml:lang="en">
This type extends DCMIRecordType to add ows:BoundingBox;
it may be used to specify a spatial envelope for the
catalogued resource.
</xsd:documentation>

```

```

        </xsd:annotation>
        <xsd:complexContent>
            <xsd:extension base="csw:DCMIRecordType">
                <xsd:sequence>
                    <xsd:element maxOccurs="unbounded" minOccurs="0" name="AnyText" type="csw:EmptyType"/>
                    <xsd:element maxOccurs="unbounded" minOccurs="0" ref="ows:BoundingBox"/>
                </xsd:sequence>
            </xsd:extension>
        </xsd:complexContent>
    </xsd:complexType>
    <xsd:complexType name="EmptyType"/>
</xsd:schema>
</csw:SchemaComponent>
<csw:SchemaComponent targetNamespace="http://www.isotc211.org/2005/gmd" schemaLanguage="http://www.w3.org/XMLSchema">
    <xs:schema xmlns:xs="http://www.w3.org/2001/XMLSchema" xmlns:gco="http://www.isotc211.org/2005/gco" xmlns:gmd="http://www.isotc211.org/2005/gmd" xmlns:xlink="http://www.w3.org/1999/xlink" elementFormDefault="qualified" targetNamespace="http://www.isotc211.org/2005/gmd" version="2012-07-13">
        <xs:annotation>
            <xs:documentation>
                Geographic MetaData (GMD) extensible markup language is a component of the XML Schema Implementation of Geographic Information Metadata documented in ISO/TS 19139:2007. GMD includes all the definitions of http://www.isotc211.org/2005/gmd namespace. The root document of this namespace is the file gmd.xsd. This identification.xsd schema implements the UML conceptual schema defined in A.2.2 of ISO 19115:2003. It contains the implementation of the following classes: MD_Identification, MD_BrowseGraphic, MD_DataIdentification, MD_ServiceIdentification, MD_RepresentativeFraction, MD_Usage, MD_Keywords, DS_Association, MD_AggregateInformation, MD_CharacterSetCode, MD_SpatialRepresentationTypeCode, MD_TopicCategoryCode, MD_ProgressCode, MD_KeywordTypeCode, DS_AssociationTypeCode, DS_InitiativeTypeCode, MD_ResolutionType.
            </xs:documentation>
        </xs:annotation>
        <xs:import namespace="http://www.isotc211.org/2005/gco" schemaLocation="http://schemas.opengis.net/iso/19139/20070417/gco/gco.xsd"/>
        <xs:include schemaLocation="gmd.xsd"/>
        <xs:include schemaLocation="constraints.xsd"/>
        <xs:include schemaLocation="distribution.xsd"/>
        <xs:include schemaLocation="maintenance.xsd"/>
        <xs:complexType abstract="true" name="AbstractMD_Identification_Type">
            <xs:annotation>
                <xs:documentation>Basic information about data</xs:documentation>

```

```

</xs:annotation>
<xs:complexContent>
  <xs:extension base="gco:AbstractObject_Type">
    <xs:sequence>
      <xs:element name="citation" type=
"gmd:CI_Citation_PropertyType"/>
      <xs:element name="abstract" type=
"gco:CharacterString_PropertyType"/>
      <xs:element minOccurs="0" name="purpose" type=
"gco:CharacterString_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"credit" type="gco:CharacterString_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"status" type="gmd:MD_ProgressCode_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"pointOfContact" type="gmd:CI_ResponsibleParty_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"resourceMaintenance" type="gmd:MD_MaintenanceInformation_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"graphicOverview" type="gmd:MD_BrowseGraphic_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"resourceFormat" type="gmd:MD_Format_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"descriptiveKeywords" type="gmd:MD_Keywords_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"resourceSpecificUsage" type="gmd:MD_Usage_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"resourceConstraints" type="gmd:MD_Constraints_PropertyType"/>
      <xs:element maxOccurs="unbounded" minOccurs="0" name=
"aggregationInfo" type="gmd:MD_AggregateInformation_PropertyType"/>
    </xs:sequence>
  </xs:extension>
</xs:complexContent>
</xs:complexType>
<xs:element abstract="true" name="AbstractMD_Identification" type=
"gmd:AbstractMD_Identification_Type"/>
<xs:complexType name="MD_Identification_PropertyType">
  <xs:sequence minOccurs="0">
    <xs:element ref="gmd:AbstractMD_Identification"/>
  </xs:sequence>
  <xs:attributeGroup ref="gco:ObjectReference"/>
  <xs:attribute ref="gco:nilReason"/>
</xs:complexType>
<xs:complexType name="MD_BrowseGraphic_Type">

```

```

<xs:annotation>
  <xs:documentation>
    Graphic that provides an illustration of the dataset (should include a
    legend for the graphic)
  </xs:documentation>

  </xs:annotation>
  <xs:complexContent>
    <xs:extension base="gco:AbstractObject_Type">
      <xs:sequence>
        <xs:element name="fileName" type=
        "gco:CharacterString_PropertyType"/>
        <xs:element minOccurs="0" name="fileDescription" type=
        "gco:CharacterString_PropertyType"/>
        <xs:element minOccurs="0" name="fileType" type=
        "gco:CharacterString_PropertyType"/>
      </xs:sequence>
    </xs:extension>
  </xs:complexContent>

  </xs:complexType>
  <xs:element name="MD_BrowseGraphic" type="gmd:MD_BrowseGraphic_Type"/>
  <xs:complexType name="MD_BrowseGraphic_PropertyType">
    <xs:sequence minOccurs="0">
      <xs:element ref="gmd:MD_BrowseGraphic"/>
    </xs:sequence>
    <xs:attributeGroup ref="gco:ObjectReference"/>
    <xs:attribute ref="gco:nilReason"/>
  </xs:complexType>
  <xs:complexType name="MD_DataIdentification_Type">
    <xs:complexContent>
      <xs:extension base="gmd:AbstractMD_Identification_Type">
        <xs:sequence>
          <xs:element maxOccurs="unbounded" minOccurs="0" name=
          "spatialRepresentationType" type="gmd:MD_SpatialRepresentationTypeCode_PropertyType"/>
          <xs:element maxOccurs="unbounded" minOccurs="0" name=
          "spatialResolution" type="gmd:MD_Resolution_PropertyType"/>
          <xs:element maxOccurs="unbounded" name="language" type=
          "gco:CharacterString_PropertyType"/>
          <xs:element maxOccurs="unbounded" minOccurs="0" name=
          "characterSet" type="gmd:MD_CharacterSetCode_PropertyType"/>
          <xs:element maxOccurs="unbounded" minOccurs="0" name=
          "topicCategory" type="gmd:MD_TopicCategoryCode_PropertyType"/>
          <xs:element minOccurs="0" name="environmentDescription" type=
          "gco:CharacterString_PropertyType"/>
          <xs:element maxOccurs="unbounded" minOccurs="0" name=

```

```

"extent" type="gmd:EX_Extent_PropertyType"/>
    <xs:element minOccurs="0" name="supplementalInformation"
type="gco:CharacterString_PropertyType"/>

    </xs:sequence>

</xs:extension>

</xs:complexContent>

</xs:complexType>
<xs:element name="MD_DataIdentification" substitutionGroup=
"gmd:AbstractMD_Identification" type="gmd:MD_DataIdentification_Type"/>
<xs:complexType name="MD_DataIdentification_PropertyType">
    <xs:sequence minOccurs="0">
        <xs:element ref="gmd:MD_DataIdentification"/>

    </xs:sequence>
    <xs:attributeGroup ref="gco:ObjectReference"/>
    <xs:attribute ref="gco:nilReason"/>

</xs:complexType>
<xs:complexType name="MD_ServiceIdentification_Type">
    <xs:annotation>
        <xs:documentation>See 19119 for further info</xs:documentation>

    </xs:annotation>
    <xs:complexContent>
        <xs:extension base="gmd:AbstractMD_Identification_Type"/>

    </xs:complexContent>

</xs:complexType>
<xs:element name="MD_ServiceIdentification" substitutionGroup=
"gmd:AbstractMD_Identification" type="gmd:MD_ServiceIdentification_Type"/>
<xs:complexType name="MD_ServiceIdentification_PropertyType">
    <xs:sequence minOccurs="0">
        <xs:element ref="gmd:MD_ServiceIdentification"/>

    </xs:sequence>
    <xs:attributeGroup ref="gco:ObjectReference"/>
    <xs:attribute ref="gco:nilReason"/>

</xs:complexType>
<xs:complexType name="MD_RepresentativeFraction_Type">
    <xs:complexContent>
        <xs:extension base="gco:AbstractObject_Type">
            <xs:sequence>
                <xs:element name="denominator" type=
"gco:Integer_PropertyType"/>

```

```

        </xs:sequence>

    </xs:extension>

</xs:complexContent>

</xs:complexType>
<xs:element name="MD_RepresentativeFraction" type=
"gmd:MD_RepresentativeFraction_Type"/>
<xs:complexType name="MD_RepresentativeFraction_PropertyType">
    <xs:sequence minOccurs="0">
        <xs:element ref="gmd:MD_RepresentativeFraction"/>

    </xs:sequence>
    <xs:attributeGroup ref="gco:ObjectReference"/>
    <xs:attribute ref="gco:nilReason"/>

</xs:complexType>
<xs:complexType name="MD_Usage_Type">
    <xs:annotation>
        <xs:documentation>
            Brief description of ways in which the dataset is currently used.
        </xs:documentation>

        </xs:annotation>
        <xs:complexContent>
            <xs:extension base="gco:AbstractObject_Type">
                <xs:sequence>
                    <xs:element name="specificUsage" type=
"gco:CharacterString_PropertyType"/>
                    <xs:element minOccurs="0" name="usageDateTime" type=
"gco:DateTime_PropertyType"/>
                    <xs:element minOccurs="0" name="userDeterminedLimitations" type="gco:CharacterString_PropertyType"/>
                    <xs:element maxOccurs="unbounded" name="userContactInfo" type="gmd:CI_ResponsibleParty_PropertyType"/>

                </xs:sequence>
            </xs:extension>
        </xs:complexContent>
    </xs:complexType>
    <xs:element name="MD_Usage" type="gmd:MD_Usage_Type"/>
    <xs:complexType name="MD_Usage_PropertyType">
        <xs:sequence minOccurs="0">
            <xs:element ref="gmd:MD_Usage"/>

        </xs:sequence>
        <xs:attributeGroup ref="gco:ObjectReference"/>

```

```

<xs:attribute ref="gco:nilReason"/>

</xs:complexType>
<xs:complexType name="MD_Keywords_Type">
  <xs:annotation>
    <xs:documentation>Keywords, their type and reference
source</xs:documentation>

  </xs:annotation>
  <xs:complexContent>
    <xs:extension base="gco:AbstractObject_Type">
      <xs:sequence>
        <xs:element maxOccurs="unbounded" name="keyword" type=
"gco:CharacterString_PropertyType"/>
        <xs:element minOccurs="0" name="type" type=
"gmd:MD_KeywordTypeCode_PropertyType"/>
        <xs:element minOccurs="0" name="thesaurusName" type=
"gmd:CI_Citation_PropertyType"/>

      </xs:sequence>

    </xs:extension>

  </xs:complexContent>

</xs:complexType>
<xs:element name="MD_Keywords" type="gmd:MD_Keywords_Type"/>
<xs:complexType name="MD_Keywords_PropertyType">
  <xs:sequence minOccurs="0">
    <xs:element ref="gmd:MD_Keywords"/>

  </xs:sequence>
  <xs:attributeGroup ref="gco:ObjectReference"/>
  <xs:attribute ref="gco:nilReason"/>

</xs:complexType>
<xs:complexType name="DS_Association_Type">
  <xs:complexContent>
    <xs:extension base="gco:AbstractObject_Type">
      <xs:sequence/>

    </xs:extension>

  </xs:complexContent>

</xs:complexType>
<xs:element name="DS_Association" type="gmd:DS_Association_Type"/>
<xs:complexType name="DS_Association_PropertyType">
  <xs:sequence minOccurs="0">
    <xs:element ref="gmd:DS_Association"/>

```

```

</xs:sequence>
<xs:attributeGroup ref="gco:ObjectReference"/>
<xs:attribute ref="gco:nilReason"/>

</xs:complexType>
<xs:complexType name="MD_AggregateInformation_Type">
    <xs:annotation>
        <xs:documentation>Encapsulates the dataset aggregation
information</xs:documentation>
    </xs:annotation>
    <xs:complexContent>
        <xs:extension base="gco:AbstractObject_Type">
            <xs:sequence>
                <xs:element minOccurs="0" name="aggregateDataSetName"
type="gmd:CI_Citation_PropertyType"/>
                <xs:element minOccurs="0" name="
aggregateDataSetIdentifier" type="gmd:MD_Identifier_PropertyType"/>
                <xs:element name="associationType" type=
"gmd:DS_AssociationTypeCode_PropertyType"/>
                <xs:element minOccurs="0" name="initiativeType" type=
"gmd:DS_InitiativeTypeCode_PropertyType"/>
            </xs:sequence>
        </xs:extension>
    </xs:complexContent>
</xs:complexType>
<xs:element name="MD_AggregateInformation" type=
"gmd:MD_AggregateInformation_Type"/>
<xs:complexType name="MD_AggregateInformation_PropertyType">
    <xs:sequence minOccurs="0">
        <xs:element ref="gmd:MD_AggregateInformation"/>
    </xs:sequence>
    <xs:attributeGroup ref="gco:ObjectReference"/>
    <xs:attribute ref="gco:nilReason"/>

</xs:complexType>
<xs:complexType name="MD_Resolution_Type">
    <xs:choice>
        <xs:element name="equivalentScale" type=
"gmd:MD_RepresentativeFraction_PropertyType"/>
        <xs:element name="distance" type="gco:Distance_PropertyType"/>
    </xs:choice>
</xs:complexType>
<xs:element name="MD_Resolution" type="gmd:MD_Resolution_Type"/>

```

```

<xs:complexType name="MD_Resolution_PropertyType">
  <xs:sequence minOccurs="0">
    <xs:element ref="gmd:MD_Resolution"/>

  </xs:sequence>
  <xs:attribute ref="gco:nilReason"/>

</xs:complexType>
<xs:simpleType name="MD_TopicCategoryCode_Type">
  <xs:annotation>
    <xs:documentation>
      High-level geospatial data thematic classification to assist in the
      grouping and search of available geospatial datasets
    </xs:documentation>

    </xs:annotation>
    <xs:restriction base="xs:string">
      <xs:enumeration value="farming"/>
      <xs:enumeration value="biota"/>
      <xs:enumeration value="boundaries"/>
      <xs:enumeration value="climatologyMeteorologyAtmosphere"/>
      <xs:enumeration value="economy"/>
      <xs:enumeration value="elevation"/>
      <xs:enumeration value="environment"/>
      <xs:enumeration value="geoscientificInformation"/>
      <xs:enumeration value="health"/>
      <xs:enumeration value="imageryBaseMapsEarthCover"/>
      <xs:enumeration value="inlandWaters"/>
      <xs:enumeration value="intelligenceMilitary"/>
      <xs:enumeration value="location"/>
      <xs:enumeration value="oceans"/>
      <xs:enumeration value="planningCadastre"/>
      <xs:enumeration value="society"/>
      <xs:enumeration value="structure"/>
      <xs:enumeration value="transportation"/>
      <xs:enumeration value="utilitiesCommunication"/>

    </xs:restriction>
  </xs:simpleType>
  <xs:element name="MD_TopicCategoryCode" substitutionGroup=
  "gco:CharacterString" type="gmd:MD_TopicCategoryCode_Type"/>
  <xs:complexType name="MD_TopicCategoryCode_PropertyType">
    <xs:sequence minOccurs="0">
      <xs:element ref="gmd:MD_TopicCategoryCode"/>

    </xs:sequence>
    <xs:attribute ref="gco:nilReason"/>

  </xs:complexType>
  <xs:element name="MD_CharacterSetCode" substitutionGroup=

```

```

" gco:CharacterString" type="gco:CodeListValue_Type"/>
  <xs:complexType name="MD_CharacterSetCode_PropertyType">
    <xs:sequence minOccurs="0">
      <xs:element ref="gmd:MD_CharacterSetCode"/>

    </xs:sequence>
    <xs:attribute ref="gco:nilReason"/>

  </xs:complexType>
  <xs:element name="MD_SpatialRepresentationTypeCode" substitutionGroup=
" gco:CharacterString" type="gco:CodeListValue_Type"/>
  <xs:complexType name="MD_SpatialRepresentationTypeCode_PropertyType">
    <xs:sequence minOccurs="0">
      <xs:element ref="gmd:MD_SpatialRepresentationTypeCode"/>

    </xs:sequence>
    <xs:attribute ref="gco:nilReason"/>

  </xs:complexType>
  <xs:element name="MD_ProgressCode" substitutionGroup="gco:CharacterString"
type="gco:CodeListValue_Type"/>
  <xs:complexType name="MD_ProgressCode_PropertyType">
    <xs:sequence minOccurs="0">
      <xs:element ref="gmd:MD_ProgressCode"/>

    </xs:sequence>
    <xs:attribute ref="gco:nilReason"/>

  </xs:complexType>
  <xs:element name="MD_KeywordTypeCode" substitutionGroup=
" gco:CharacterString" type="gco:CodeListValue_Type"/>
  <xs:complexType name="MD_KeywordTypeCode_PropertyType">
    <xs:sequence minOccurs="0">
      <xs:element ref="gmd:MD_KeywordTypeCode"/>

    </xs:sequence>
    <xs:attribute ref="gco:nilReason"/>

  </xs:complexType>
  <xs:element name="DS_AssociationTypeCode" substitutionGroup=
" gco:CharacterString" type="gco:CodeListValue_Type"/>
  <xs:complexType name="DS_AssociationTypeCode_PropertyType">
    <xs:sequence minOccurs="0">
      <xs:element ref="gmd:DS_AssociationTypeCode"/>

    </xs:sequence>
    <xs:attribute ref="gco:nilReason"/>

  </xs:complexType>
  <xs:element name="DS_InitiativeTypeCode" substitutionGroup=
" gco:CharacterString" type="gco:CodeListValue_Type"/>
  <xs:complexType name="DS_InitiativeTypeCode_PropertyType">
    <xs:sequence minOccurs="0">

```

```

<xs:element ref="gmd:DS_InitiativeTypeCode"/>
</xs:sequence>
<xs:attribute ref="gco:nilReason"/>
</xs:complexType>
</xs:schema>
</csw:SchemaComponent>
</csw:DescribeRecordResponse>

```

DescribeRecord *HTTP POST With TypeNames*

The HTTP POST request **DescribeRecordType** has the **typeName** as a List of QName(s). The QNames are matched against the namespaces by prefix, if prefixes exist. [.DescribeRecord XML Request](#)

```

<?xml version="1.0" ?>
<DescribeRecord
  version="2.0.2"
  service="CSW"
  schemaLanguage="http://www.w3.org/XMLSchema"
  xmlns="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2">
  <TypeName>csw:Record</TypeName>
</DescribeRecord>

```

DescribeRecord *Sample Response (application/xml)*

```

<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:DescribeRecordResponse xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:gml=
"http://www.opengis.net/gml" xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
"http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
instance" ns10:schemaLocation="http://www.opengis.net/csw /ogc/csw/2.0.2/CSW-
publication.xsd">
  <csw:SchemaComponent targetNamespace="http://www.opengis.net/cat/csw/2.0.2"
  schemaLanguage="http://www.w3.org/XMLSchema">
    <xsd:schema xmlns:xsd="http://www.w3.org/2001/XMLSchema" elementFormDefault=
"qualified" id="csw-record" targetNamespace="http://www.opengis.net/cat/csw/2.0.2"
  version="2.0.2">
      <xsd:annotation>
        <xsd:appinfo>
          <dc:identifier>
            http://schemas.opengis.net/csw/2.0.2/record.xsd</dc:identifier>
          </xsd:appinfo>
        <xsd:documentation xml:lang="en">

```

This schema defines the basic record types that must be supported by all CSW implementations. These correspond to full, summary, and brief views based on DCMI metadata terms.

```
</xsd:documentation>
```

```

</xsd:annotation>
<xsd:import namespace="http://purl.org/dc/terms/" schemaLocation="rec-
dcterms.xsd"/>
    <xsd:import namespace="http://purl.org/dc/elements/1.1/" schemaLocation=
"rec-dcmes.xsd"/>
        <xsd:import namespace="http://www.opengis.net/ows" schemaLocation=
"../../ows/1.0.0/owsAll.xsd"/>
            <xsd:element abstract="true" id="AbstractRecord" name="AbstractRecord"
type="csw:AbstractRecordType"/>
                <xsd:complexType abstract="true" id="AbstractRecordType" name=
"AbstractRecordType"/>
                    <xsd:element name="DCMIRecord" substitutionGroup="csw:AbstractRecord"
type="csw:DCMIRecordType"/>
                        <xsd:complexType name="DCMIRecordType">
                            <xsd:annotation>
                                <xsd:documentation xml:lang="en">
This type encapsulates all of the standard DCMI metadata terms,
including the Dublin Core refinements; these terms may be mapped
to the profile-specific information model.
                            </xsd:documentation>
                        </xsd:complexType>
                        <xsd:complexContent>
                            <xsd:extension base="csw:AbstractRecordType">
                                <xsd:sequence>
                                    <xsd:group ref="dct:DCMI-terms"/>
                                </xsd:sequence>
                            </xsd:extension>
                        </xsd:complexContent>
                    </xsd:complexType>
                    <xsd:element name="BriefRecord" substitutionGroup="csw:AbstractRecord"
type="csw:BriefRecordType"/>
                        <xsd:complexType final="#all" name="BriefRecordType">
                            <xsd:annotation>
                                <xsd:documentation xml:lang="en">
This type defines a brief representation of the common record
format. It extends AbstractRecordType to include only the
dc:identifier and dc:type properties.
                            </xsd:documentation>
                        </xsd:complexType>
                        <xsd:complexContent>
                            <xsd:extension base="csw:AbstractRecordType">
                                <xsd:sequence>
                                    <xsd:element maxOccurs="unbounded" minOccurs="1" ref=
"dc:identifier"/>

```

```

        <xsd:element maxOccurs="unbounded" minOccurs="1" ref=
"dc:title"/>
        <xsd:element minOccurs="0" ref="dc:type"/>
        <xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"ows:BoundingBox"/>

    </xsd:sequence>

</xsd:extension>

</xsd:complexContent>

</xsd:complexType>
<xsd:element name="SummaryRecord" substitutionGroup="csw:AbstractRecord"
type="csw:SummaryRecordType"/>
<xsd:complexType final="#all" name="SummaryRecordType">
    <xsd:annotation>
        <xsd:documentation xml:lang="en">
This type defines a summary representation of the common record
format. It extends AbstractRecordType to include the core
properties.
        </xsd:documentation>

    </xsd:annotation>
    <xsd:complexContent>
        <xsd:extension base="csw:AbstractRecordType">
            <xsd:sequence>
                <xsd:element maxOccurs="unbounded" minOccurs="1" ref=
"dc:identifier"/>
                <xsd:element maxOccurs="unbounded" minOccurs="1" ref=
"dc:title"/>
                <xsd:element minOccurs="0" ref="dc:type"/>
                <xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:subject"/>
                <xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:format"/>
                <xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:relation"/>
                <xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:modified"/>
                <xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:abstract"/>
                <xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"dc:spatial"/>
                <xsd:element maxOccurs="unbounded" minOccurs="0" ref=
"ows:BoundingBox"/>

            </xsd:sequence>

        </xsd:extension>
    </xsd:complexContent>
</xsd:complexType>

```

```

        </xsd:complexContent>

        </xsd:complexType>
        <xsd:element name="Record" substitutionGroup="csw:AbstractRecord" type="csw:RecordType"/>
            <xsd:complexType final="#all" name="RecordType">
                <xsd:annotation>
                    <xsd:documentation xml:lang="en">
                        This type extends DCMIRecordType to add ows:BoundingBox;
                        it may be used to specify a spatial envelope for the
                        catalogued resource.
                </xsd:documentation>

                </xsd:annotation>
                <xsd:complexContent>
                    <xsd:extension base="csw:DCMIRecordType">
                        <xsd:sequence>
                            <xsd:element maxOccurs="unbounded" minOccurs="0" name="AnyText" type="csw:EmptyType"/>
                            <xsd:element maxOccurs="unbounded" minOccurs="0" ref="ows:BoundingBox"/>
                        </xsd:sequence>
                    </xsd:extension>
                </xsd:complexContent>
            </xsd:complexType>
            <xsd:complexType name="EmptyType"/>
        </xsd:schema>
    </csw:SchemaComponent>
</csw:DescribeRecordResponse>

```

GetRecords Operation

The **GetRecords** operation is the principal means of searching the catalog. The matching entries may be included with the response. The client may assign a **requestId** (absolute URI). A distributed search is performed if the **DistributedSearch** element is present and the catalog is a member of a federation. Profiles may allow alternative query expressions. There are two types of request types: one for **GET** and one for **POST**. Each request has the following common data parameters:

Namespace

In POST operations, namespaces are defined in the XML. In GET operations, namespaces are defined in a comma-separated list of the form `xmlns([prefix=]namespace-url),xmlns([pref:=]namespace-url))*`.

Service

The service being used, in this case it is fixed at CSW.

Version

The version of the service being used (2.0.2).

OutputFormat

The requester wants the response to be in this intended output. Currently, only one format is supported (application/xml). If this parameter is supplied, it is validated against the known type. If this parameter is not supported, it passes through and returns the XML response upon success.

OutputSchema

The OutputSchema indicates which schema shall be used to generate the response to the GetRecords operation. The supported output schemas are listed in the GetCapabilities response.

ElementSetName

CodeList with allowed values of "brief", "summary", or "full". The default value is "summary". The predefined set names of "brief", "summary", and "full" represent different levels of detail for the source record. "Brief" represents the least amount of detail, and "full" represents all the metadata record elements.

IMPORTANT

The CSW Endpoint expects all geospatial filters using the EPSG:4326 CRS to use "longitude then latitude" coordinate ordering. Similarly, unless the output schema explicitly states otherwise, the GetRecordsResponse will use the same coordinate ordering.

GetRecords HTTP GET

The **HTTP GET** request differs from the **POST** request in that it has the "typeNames" as a comma-separated list of namespace prefix qualified types as strings. For example **csw:Record,xyz:MyType**. These prefixes are then matched against the prefix qualified namespaces in the request. This is converted to a list QName(s). In this way, it behaves exactly as the post request that uses a list of QName(s) in the first place.

GetRecords KVP (Key-Value Pairs) Encoding

```
\https://{{FQDN}}:{{PORT}}/services/csw?service=CSW&version=2.0.2&request=GetRecords&outputFormat=application/xml&outputSchema=http://www.opengis.net/cat/csw/2.0.2&NAME_SPACE=xmns(csw=http://www.opengis.net/cat/csw/2.0.2)&resultType=results&typeNames=csw:Record&ElementSetName=brief&ConstraintLanguage=CQL_TEXT&constraint=AnyText Like '%25'
```

GetRecords HTTP POST

The **HTTP POST** request GetRecords has the **typeNames** as a List of QName(s). The QNames are matched against the namespaces by prefix, if prefixes exist.

GetRecords XML Request

```
<?xml version="1.0" ?>
<GetRecords xmlns="http://www.opengis.net/cat/csw/2.0.2"
    xmlns:ogc="http://www.opengis.net/ogc"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    service="CSW"
    version="2.0.2"
    maxRecords="4"
    startPosition="1"
    resultType="results"
    outputFormat="application/xml"
    outputSchema="http://www.opengis.net/cat/csw/2.0.2"
    xsi:schemaLocation="http://www.opengis.net/cat/csw/2.0.2
    ../../../../../../csw/2.0.2/CSW-discovery.xsd">
    <Query typeNames="Record">
        <ElementSetName>summary</ElementSetName>
        <Constraint version="1.1.0">
            <ogc:Filter>
                <ogc:PropertyIsLike wildCard="#" singleChar="_" escapeChar="\">
                    <ogc:PropertyName>AnyText</ogc:PropertyName>
                    <ogc:Literal>%</ogc:Literal>
                </ogc:PropertyIsLike>
            </ogc:Filter>
        </Constraint>
    </Query>
</GetRecords>
```

GetRecords Specific Source

It is possible to query a **Specific Source** by specifying a query for that source-id. The valid **source-id**'s will be listed in the **FederatedCatalogs** section of the **GetCapabilities** Response. The example below shows how to query for a specific source.

NOTE

The **DistributedSearch** element must be specific with a **hopCount** greater than 1 to identify it as a federated query, otherwise the **source-id**'s will be ignored.

GetRecords XML Request

```
<?xml version="1.0" ?>
<csw:GetRecords resultType="results"
  outputFormat="application/xml"
  outputSchema="urn:catalog:metacard"
  startPosition="1"
  maxRecords="10"
  service="CSW"
  version="2.0.2"
  xmlns:ns2="http://www.opengis.net/ogc" xmlns:csw=
  "http://www.opengis.net/cat/csw/2.0.2" xmlns:ns4="http://www.w3.org/1999/xlink"
  xmlns:ns3="http://www.opengis.net/gml" xmlns:ns9=
  "http://www.w3.org/2001/SMIL20/Language" xmlns:ns5="http://www.opengis.net/ows"
  xmlns:ns6="http://purl.org/dc/elements/1.1/" xmlns:ns7="http://purl.org/dc/terms/"
  xmlns:ns8="http://www.w3.org/2001/SMIL20/">
  <csw:DistributedSearch hopCount="2" />
  <ns10:Query typeNames="csw:Record" xmlns="" xmlns:ns10=
  "http://www.opengis.net/cat/csw/2.0.2">
    <ns10:ElementSetName>full</ns10:ElementSetName>
    <ns10:Constraint version="1.1.0">
      <ns2:Filter>
        <ns2:And>
          <ns2:PropertyIsEqualTo wildCard="*" singleChar="#" escapeChar="!">
            <ns2:PropertyName>source-id</ns2:PropertyName>
            <ns2:Literal>Source1</ns2:Literal>
          </ns2:PropertyIsEqualTo>
          <ns2:PropertyIsLike wildCard="*" singleChar="#" escapeChar="!">
            <ns2:PropertyName>title</ns2:PropertyName>
            <ns2:Literal>*</ns2:Literal>
          </ns2:PropertyIsLike>
        </ns2:And>
      </ns2:Filter>
    </ns10:Constraint>
  </ns10:Query>
</csw:GetRecords>
```

GetRecords Sample Response (application/xml)

```
<csw:GetRecordsResponse version="2.0.2" xmlns:dc="http://purl.org/dc/elements/1.1/"
  xmlns:dct="http://purl.org/dc/terms/" xmlns:ows="http://www.opengis.net/ows" xmlns:xs
  ="http://www.w3.org/2001/XMLSchema" xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">
  <csw:SearchStatus timestamp="2014-02-19T15:33:44.602-05:00"/>
  <csw:SearchResults numberOfRecordsMatched="41" numberOfRecordsReturned="4"
  nextRecord="5" recordSchema="http://www.opengis.net/cat/csw/2.0.2" elementSet="summary">
    <csw:SummaryRecord>
      <dc:identifier>182fb33103414e5cbb06f8693b526239</dc:identifier>
      <dc:title>Product10</dc:title>
```

```

<dc:type>pdf</dc:type>
<dct:modified>2014-02-19T15:22:51.563-05:00</dct:modified>
<ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">
  <ows:LowerCorner>20.0 10.0</ows:LowerCorner>
  <ows:UpperCorner>20.0 10.0</ows:UpperCorner>
</ows:BoundingBox>
</csw:SummaryRecord>
<csw:SummaryRecord>
  <dc:identifier>c607440db9b0407e92000d9260d35444</dc:identifier>
  <dc:title>Product03</dc:title>
  <dc:type>pdf</dc:type>
  <dct:modified>2014-02-19T15:22:51.563-05:00</dct:modified>
  <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">
    <ows:LowerCorner>6.0 3.0</ows:LowerCorner>
    <ows:UpperCorner>6.0 3.0</ows:UpperCorner>
  </ows:BoundingBox>
</csw:SummaryRecord>
<csw:SummaryRecord>
  <dc:identifier>034cc757abd645f0abeb6acaccfe194de</dc:identifier>
  <dc:title>Product03</dc:title>
  <dc:type>pdf</dc:type>
  <dct:modified>2014-02-19T15:22:51.563-05:00</dct:modified>
  <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">
    <ows:LowerCorner>6.0 3.0</ows:LowerCorner>
    <ows:UpperCorner>6.0 3.0</ows:UpperCorner>
  </ows:BoundingBox>
</csw:SummaryRecord>
<csw:SummaryRecord>
  <dc:identifier>5d6e987bd6084bd4919d06b63b77a007</dc:identifier>
  <dc:title>Product01</dc:title>
  <dc:type>pdf</dc:type>
  <dct:modified>2014-02-19T15:22:51.563-05:00</dct:modified>
  <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">
    <ows:LowerCorner>2.0 1.0</ows:LowerCorner>
    <ows:UpperCorner>2.0 1.0</ows:UpperCorner>
  </ows:BoundingBox>
</csw:SummaryRecord>
</csw:SearchResults>
</csw:GetRecordsResponse>

```

GetRecords GMD OutputSchema

It is possible to receive a response to a **GetRecords** query that conforms to the GMD specification.

GetRecords XML Request

```
<?xml version="1.0" ?>
<GetRecords xmlns="http://www.opengis.net/cat/csw/2.0.2"
    xmlns:ogc="http://www.opengis.net/ogc"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xmlns:gmd="http://www.isotc211.org/2005/gmd"
    xmlns:gml="http://www.opengis.net/gml"
    service="CSW"
    version="2.0.2"
    maxRecords="8"
    startPosition="1"
    resultType="results"
    outputFormat="application/xml"
    outputSchema="http://www.isotc211.org/2005/gmd"
    xsi:schemaLocation="http://www.opengis.net/cat/csw/2.0.2
    ../.././././csw/2.0.2/CSW-discovery.xsd">
    <Query typeNames="gmd:MD_Metadata">
        <ElementSetName>summary</ElementSetName>
        <Constraint version="1.1.0">
            <ogc:Filter>
                <ogc:PropertyIsLike wildCard="#" singleChar="_" escapeChar="\">
                    <ogc:PropertyName>apiso:Title</ogc:PropertyName>
                    <ogc:Literal>prod%</ogc:Literal>
                </ogc:PropertyIsLike>
            </ogc:Filter>
        </Constraint>
    </Query>
</GetRecords>
```

GetRecords Sample Response (application/xml)

```
<?xml version='1.0' encoding='UTF-8'?>
<csw:GetRecordsResponse xmlns:dct="http://purl.org/dc/terms/" xmlns:xml=
"http://www.w3.org/XML/1998/namespace" xmlns:csw="http://www.opengis.net/cat/csw/2.0.2" xmlns:ows="http://www.opengis.net/ows" xmlns:xs=
"http://www.w3.org/2001/XMLSchema" xmlns:xsi="http://www.w3.org/2001/XMLSchema-
instance" xmlns:dc="http://purl.org/dc/elements/1.1/" version="2.0.2">
    <csw:SearchStatus timestamp="2016-03-23T11:31:34.531-06:00"/>
    <csw:SearchResults numberOfRecordsMatched="7" numberOfRecordsReturned="1"
    nextRecord="2" recordSchema="http://www.isotc211.org/2005/gmd" elementSet="summary">
        <MD_Metadata xmlns="http://www.isotc211.org/2005/gmd" xmlns:gco=
"http://www.isotc211.org/2005/gco">
            <fileIdentifier>
                <gco:CharacterString>
d5f6acd5ccf34d18af5192c38a276b12</gco:CharacterString>
            </fileIdentifier>
            <hierarchyLevel>
                <MD_ScopeCode codeListValue="nitf" codeList="urn:catalog:metacard"/>
            </hierarchyLevel>
```

```

<contact/>
<dateStamp>
    <gco:DateTime>2015-03-04T17:23:42.332-07:00</gco:DateTime>
</dateStamp>
<identificationInfo>
    <MD_DataIdentification>
        <citation>
            <CI_Citation>
                <title>
                    <gco:CharacterString>product.ntf</gco:CharacterString>
                </title>
                <date>
                    <CI_Date>
                        <date>
                            <gco:DateTime>2015-03-04T17:23:42.332-
07:00</gco:DateTime>
                        </date>
                    <dateType>
                        <CI_DateTypeCode codeList=
"urn:catalog:metacard" codeListValue="created"/>
                    </dateType>
                </CI_Date>
                </date>
            </CI_Citation>
        </citation>
        <abstract>
            <gco:CharacterString></gco:CharacterString>
        </abstract>
        <pointOfContact>
            <CI_ResponsibleParty>
                <organisationName>
                    <gco:CharacterString></gco:CharacterString>
                </organisationName>
                <role/>
            </CI_ResponsibleParty>
        </pointOfContact>
        <language>
            <gco:CharacterString>en</gco:CharacterString>
        </language>
        <extent>
            <EX_Extent>
                <geographicElement>
                    <EX_GeographicBoundingBox>
                        <westBoundLongitude>
                            <gco:Decimal>32.975277</gco:Decimal>
                        </westBoundLongitude>
                        <eastBoundLongitude>
                            <gco:Decimal>32.996944</gco:Decimal>
                        </eastBoundLongitude>
                        <southBoundLatitude>
                            <gco:Decimal>32.305</gco:Decimal>
                        </southBoundLatitude>
                    </EX_GeographicBoundingBox>
                </geographicElement>
            </EX_Extent>
        </extent>
    </MD_DataIdentification>
</identificationInfo>

```

```

        </southBoundLatitude>
        <northBoundLatitude>
            <gco:Decimal>32.323333</gco:Decimal>
        </northBoundLatitude>
    </EX_GeographicBoundingBox>
</geographicElement>
</EX_Extent>
</extent>
</MD_DataIdentification>
</identificationInfo>
<distributionInfo>
<MD_Distribution>
<distributor>
<MD_Distributor>
<distributorContact/>
<distributorTransferOptions>
<MD_DigitalTransferOptions>
<onLine>
<CI_OnlineResource>
<linkage>
<URL>http://example.com</URL>
</linkage>
</CI_OnlineResource>
</onLine>
</MD_DigitalTransferOptions>
</distributorTransferOptions>
</MD_Distributor>
</distributor>
</MD_Distribution>
</distributionInfo>
</MD_Metadata>
</csw:SearchResults>
</csw:GetRecordsResponse>

```

GetRecords XML Request using UTM Coordinates

UTM coordinates can be used when making a CSW GetRecords request using an **ogc:Filter**. UTM coordinates should use EPSG:326XX as the srsName where XX is the zone within the northern hemisphere. UTM coordinates should use EPSG:327XX as the srsName where XX is the zone within the southern hemisphere.

Note: UTM coordinates are only supported with requests providing an **ogc:Filter**, but not with CQL as there isn't a way to specify the UTM srsName in CQL.

GetRecords XML Request - UTM Northern Hemisphere Zone 36

```
<GetRecords xmlns="http://www.opengis.net/cat/csw/2.0.2"
    xmlns:ogc="http://www.opengis.net/ogc"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xmlns:gml="http://www.opengis.net/gml"
    service="CSW"
    version="2.0.2"
    maxRecords="4"
    startPosition="1"
    resultType="results"
    outputFormat="application/xml"
    outputSchema="http://www.opengis.net/cat/csw/2.0.2"
    xsi:schemaLocation="http://www.opengis.net/cat/csw/2.0.2
    ../../csw/2.0.2/CSW-discovery.xsd">
    <Query typeNames="Record">
        <ElementSetName>summary</ElementSetName>
        <Constraint version="1.1.0">
            <ogc:Filter>
                <ogc:Intersects>
                    <ogc:PropertyName>ows:BoundingBox</ogc:PropertyName>
                    <gml:Envelope srsName="EPSG:32636">
                        <gml:lowerCorner>171070 1106907</gml:lowerCorner>
                        <gml:upperCorner>225928 1106910</gml:upperCorner>
                    </gml:Envelope>
                </ogc:Intersects>
            </ogc:Filter>
        </Constraint>
    </Query>
</GetRecords>
```

GetRecordById Operation

The **GetRecordById** operation request retrieves the default representation of catalog records using their identifier. This operation presumes that a previous query has been performed in order to obtain the identifiers that may be used with this operation. For example, records returned by a **GetRecords** operation may contain references to other records in the catalog that may be retrieved using the **GetRecordById** operation. This operation is also a subset of the **GetRecords** operation and is included as a convenient short form for retrieving and linking to records in a catalog.

Clients can also retrieve products from the catalog using the **GetRecordById** operation. The client sets the output schema to <http://www.iana.org/assignments/media-types/application/octet-stream> and the output format to [application/octet-stream](http://www.iana.org/assignments/media-types/application/octet-stream) within the request. The endpoint will do the following: check that only one Id is provided, otherwise an error will occur as multiple products cannot be retrieved. If both output format and output schema are set to values mentioned above, the catalog framework will retrieve the resource for that Id. The HTTP content type is then set to the resource's MIME type and the data is sent out. The endpoint also supports the resumption of partial downloads. This would typically occur at the request of a browser when a download was prematurely terminated.

There are two request types: one for **GET** and one for **POST**. Each request has the following common data parameters:

Namespace

In POST operations, namespaces are defined in the XML. In GET operations namespaces are defined in a comma separated list of the form: `xmlns([prefix=]namespace-url),xmlns([prefix=]namespace-url))*`.

Service

The service being used, in this case it is fixed at "CSW".

Version

The version of the service being used (2.0.2).

OutputFormat

The requester wants the response to be in this intended output. Currently, two output formats are supported: `application/xml` for retrieving records, and `application/octet-stream` for retrieving a product. If this parameter is supplied, it is validated against the known type. If this parameter is not supported, it passes through and returns the XML response upon success.

OutputSchema

The OutputSchema indicates which schema shall be used to generate the response to the GetRecordById operation. The supported output schemas are listed in the GetCapabilities response.

ElementSetName

CodeList with allowed values of "brief", "summary", or "full". The default value is "summary". The predefined set names of "brief", "summary", and "full" represent different levels of detail for the source record. "Brief" represents the least amount of detail, and "full" represents all the metadata record elements.

Id

The Id parameter is a comma-separated list of record identifiers for the records that CSW returns to the client. In the XML encoding, one or more `<Id>` elements may be used to specify the record identifier to be retrieved.

GetRecordById *HTTP GET KVP (Key-Value Pairs) Encoding*

```
\https://{{FQDN}}:{{PORT}}/services/csw?service=CSW&version=2.0.2&request=GetRecordById&NA  
MESPACE=xmlns="http://www.opengis.net/cat/csw/2.0.2"&ElementSetName=full&outputFormat=  
application/xml&outputSchema=http://www.opengis.net/cat/csw/2.0.2&id=fd7ff1535dfe47db8  
793b550d4170424,ba908634c0eb439b84b5d9c42af1f871
```

GetRecordById *HTTP POST*

```
<GetRecordById xmlns="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:ogc="http://www.opengis.net/ogc"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  service="CSW"
  version="2.0.2"
  outputFormat="application/xml"
  outputSchema="http://www.opengis.net/cat/csw/2.0.2"
  xsi:schemaLocation="http://www.opengis.net/cat/csw/2.0.2
  ../../../../../../csw/2.0.2/CSW-discovery.xsd">
  <ElementSetName>full</ElementSetName>
  <Id>182fb33103414e5cbb06f8693b526239</Id>
  <Id>c607440db9b0407e92000d9260d35444</Id>
</GetRecordById>
```

GetRecordByIdResponse *Sample Response (application/xml)*

```
<csw:GetRecordByIdResponse xmlns:dc="http://purl.org/dc/elements/1.1/"  
    xmlns:dct="http://purl.org/dc/terms/" xmlns:ows="http://www.opengis.net/ows"  
    xmlns:xs="http://www.w3.org/2001/XMLSchema"  
    xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"  
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">  
    <csw:Record>  
        <dc:identifier>182fb33103414e5cbb06f8693b526239</dc:identifier>  
        <dct:bibliographicCitation>182fb33103414e5cbb06f8693b526239</dct:bibliographicCitation>  
        <dc:title>Product10</dc:title>  
        <dct:alternative>Product10</dct:alternative>  
        <dc:type>pdf</dc:type>  
        <dc:date>2014-02-19T15:22:51.563-05:00</dc:date>  
        <dct:modified>2014-02-19T15:22:51.563-05:00</dct:modified>  
        <dct:created>2014-02-19T15:22:51.563-05:00</dct:created>  
        <dct:dateAccepted>2014-02-19T15:22:51.563-05:00</dct:dateAccepted>  
        <dct:dateCopyrighted>2014-02-19T15:22:51.563-05:00</dct:dateCopyrighted>  
        <dct:dateSubmitted>2014-02-19T15:22:51.563-05:00</dct:dateSubmitted>  
        <dct:issued>2014-02-19T15:22:51.563-05:00</dct:issued>  
        <dc:source>ddf.distribution</dc:source>  
        <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">  
            <ows:LowerCorner>20.0 10.0</ows:LowerCorner>  
            <ows:UpperCorner>20.0 10.0</ows:UpperCorner>  
        </ows:BoundingBox>  
    </csw:Record>  
    <csw:Record>  
        <dc:identifier>c607440db9b0407e92000d9260d35444</dc:identifier>  
        <dct:bibliographicCitation>c607440db9b0407e92000d9260d35444</dct:bibliographicCitation>  
        <dc:title>Product03</dc:title>  
        <dct:alternative>Product03</dct:alternative>  
        <dc:type>pdf</dc:type>  
        <dc:date>2014-02-19T15:22:51.563-05:00</dc:date>  
        <dct:modified>2014-02-19T15:22:51.563-05:00</dct:modified>  
        <dct:created>2014-02-19T15:22:51.563-05:00</dct:created>  
        <dct:dateAccepted>2014-02-19T15:22:51.563-05:00</dct:dateAccepted>  
        <dct:dateCopyrighted>2014-02-19T15:22:51.563-05:00</dct:dateCopyrighted>  
        <dct:dateSubmitted>2014-02-19T15:22:51.563-05:00</dct:dateSubmitted>  
        <dct:issued>2014-02-19T15:22:51.563-05:00</dct:issued>  
        <dc:source>ddf.distribution</dc:source>  
        <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">  
            <ows:LowerCorner>6.0 3.0</ows:LowerCorner>  
            <ows:UpperCorner>6.0 3.0</ows:UpperCorner>  
        </ows:BoundingBox>  
    </csw:Record>  
</csw:GetRecordByIdResponse>
```

Table 27. CSW Record to Metocard Mapping

CSW Record Field	Metacard Field	Brief Record	Summary Record	Record
dc:title	title	1-n	1-n	0-n
dc:creator				0-n
dc:subject			0-n	0-n
dc:description				0-n
dc:publisher				0-n
dc:contributor				0-n
dc:date	modified			0-n
dc:type	metadata-content-type	0-1	0-1	0-n
dc:format			0-n	0-n
dc:identifier	id	1-n	1-n	0-n
dc:source	source-id			0-n
dc:language				0-n
dc:relation			0-n	0-n
dc:coverage				0-n
dc:rights				0-n
dct:abstract			0-n	0-n
dct:accessRights				0-n
dct:alternative	title			0-n
dct:audience				0-n
dct:available				0-n
dct:bibliographicCitation	id			0-n
dct:conformsTo				0-n
dct:created	created			0-n
dct:dateAccepted	effective			0-n
dct:Copyrighted	effective			0-n
dct:dateSubmitted	modified			0-n
dct:educationLevel				0-n
dct:extent				0-n
dct:hasFormat				0-n
dct:hasPart				0-n
dct:hasVersion				0-n

CSW Record Field	Metacard Field	Brief Record	Summary Record	Record
dct:isFormatOf				0-n
dct:isPartOf				0-n
dct:isReferencedBy				0-n
dct:isReplacedBy				0-n
dct:isRequiredBy				0-n
dct:issued	modified			0-n
dct:isVersionOf				0-n
dct:license				0-n
dct:mediator				0-n
dct:medium				0-n
dct:modified	modified		0-n	0-n
dct:provenance				0-n
dct:references				0-n
dct:replaces				0-n
dct:requires				0-n
dct:rightsHolder				0-n
dct:spatial	location		0-n	0-n
dct:tableOfContents				0-n
dct:temporal	effective + " - " + expiration			0-n
dct:valid	expiration			0-n
ows:BoundingBox		0-n	0-n	0-n

14.1.4.3.2. Transaction Operations

Transactions define the operations for creating, modifying, and deleting catalog records. The supported sub-operations for the Transaction operation are Insert, Update, and Delete.

The CSW Transactions endpoint only supports [HTTP POST](#) requests since there are no KVP (Key-Value Pairs) operations.

14.1.4.3.3. Transaction Insert Sub-Operation [HTTP POST](#)

The Insert sub-operation is a method for one or more records to be inserted into the catalog. The schema of the record needs to conform to the schema of the information model that the catalog supports as described using the [DescribeRecord](#) operation.

The following example shows a request for a record to be inserted.

Sample XML Transaction Insert Request

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction
  service="CSW"
  version="2.0.2"
  verboseResponse="true"
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2">
  <csw:Insert typeName="csw:Record">
    <csw:Record
      xmlns:ows="http://www.opengis.net/ows"
      xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
      xmlns:dc="http://purl.org/dc/elements/1.1/"
      xmlns:dct="http://purl.org/dc/terms/"
      xmlns:xsd="http://www.w3.org/2001/XMLSchema">
      <dc:identifier></dc:identifier>
      <dc:title>Aliquam fermentum purus quis arcu</dc:title>
      <dc:type>http://purl.org/dc/dcmitype/Text</dc:type>
      <dc:subject>Hydrography--Dictionaries</dc:subject>
      <dc:format>application/pdf</dc:format>
      <dc:date>2006-05-12</dc:date>
      <dct:abstract>Vestibulum quis ipsum sit amet metus imperdiet vehicula.
      Nulla scelerisque cursus mi.</dct:abstract>
      <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">
        <ows:LowerCorner>44.792 -6.171</ows:LowerCorner>
        <ows:UpperCorner>51.126 -2.228</ows:UpperCorner>
      </ows:BoundingBox>
    </csw:Record>
  </csw:Insert>
</csw:Transaction>
```

14.1.4.3.4. Transaction Insert Response

The following is an example of an **application/xml** response to the Transaction Insert sub-operation:

Note that you will only receive the **InsertResult** element if you specify **verboseResponse="true"**.

Sample XML Transaction Insert Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:TransactionResponse xmlns:ogc="http://www.opengis.net/ogc"
                           xmlns:gml="http://www.opengis.net/gml"
                           xmlns:ns3="http://www.w3.org/1999/xlink"
                           xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
                           xmlns:ns5="http://www.w3.org/2001/SMIL20/"
                           xmlns:dc="http://purl.org/dc/elements/1.1/"
                           xmlns:ows="http://www.opengis.net/ows"
                           xmlns:dct="http://purl.org/dc/terms/"
                           xmlns:ns9="http://www.w3.org/2001/SMIL20/Language"
                           xmlns:ns10="http://www.w3.org/2001/XMLSchema-instance"
                           version="2.0.2"
                           ns10:schemaLocation="http://www.opengis.net/csw
/ogc/csw/2.0.2/CSW-publication.xsd">
  <csw:TransactionSummary>
    <csw:totalInserted>1</csw:totalInserted>
    <csw:totalUpdated>0</csw:totalUpdated>
    <csw:totalDeleted>0</csw:totalDeleted>
  </csw:TransactionSummary>
  <csw:InsertResult>
    <csw:BriefRecord>
      <dc:identifier>2dbcfa3f3e24e3e8f68c50f5a98a4d1</dc:identifier>
      <dc:title>Aliquam fermentum purus quis arcu</dc:title>
      <dc:type>http://purl.org/dc/dcmitype/Text</dc:type>
      <ows:BoundingBox crs="EPSG:4326">
        <ows:LowerCorner>-6.171 44.792</ows:LowerCorner>
        <ows:UpperCorner>-2.228 51.126</ows:UpperCorner>
      </ows:BoundingBox>
    </csw:BriefRecord>
  </csw:InsertResult>
</csw:TransactionResponse>
```

14.1.4.3.5. Transaction Update Sub-Operation [HTTP POST](#)

The Update sub-operation is a method to specify values used to change existing information in the catalog. If individual record property values are specified in the [Update](#) element, using the [RecordProperty](#) element, then those individual property values of a catalog record are replaced. The [RecordProperty](#) contains a [Name](#) and [Value](#) element. The [Name](#) element is used to specify the name of the record property to be updated. The [Value](#) element contains the value that will be used to update the record in the catalog. The values in the [Update](#) will completely replace those that are already in the record. A property is removed only if the [RecordProperty](#) contains a [Name](#) but not a [Value](#).

The number of records affected by an Update operation is determined by the contents of the [Constraint](#) element, which contains a filter for limiting the update to a specific record or group of records.

The following example shows how the newly inserted record could be updated to modify the date field. If your update request contains a [<csw:Record>](#) rather than a set of [<RecordProperty>](#) elements

plus a <Constraint>, the existing record with the same ID will be replaced with the new record.

Sample XML Transaction Update Request

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction
  service="CSW"
  version="2.0.2"
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2">
  <csw:Update>
    <csw:Record
      xmlns:ows="http://www.opengis.net/ows"
      xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
      xmlns:dc="http://purl.org/dc/elements/1.1/"
      xmlns:dct="http://purl.org/dc/terms/"
      xmlns:xsd="http://www.w3.org/2001/XMLSchema">
      <dc:identifier>2dbcfba3f3e24e3e8f68c50f5a98a4d1</dc:identifier>
      <dc:title>Aliquam fermentum purus quis arcu</dc:title>
      <dc:type>http://purl.org/dc/dcmitype/Text</dc:type>
      <dc:subject>Hydrography--Dictionaries</dc:subject>
      <dc:format>application/pdf</dc:format>
      <dc:date>2008-08-10</dc:date>
      <dct:abstract>Vestibulum quis ipsum sit amet metus imperdiet vehicula.
      Nulla scelerisque cursus mi.</dct:abstract>
      <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">
        <ows:LowerCorner>44.792 -6.171</ows:LowerCorner>
        <ows:UpperCorner>51.126 -2.228</ows:UpperCorner>
      </ows:BoundingBox>
    </csw:Record>
  </csw:Update>
</csw:Transaction>
```

The following example shows how the newly inserted record could be updated to modify the date field while using a filter constraint with title equal to **Aliquam fermentum purus quis arcu**.

Sample XML Transaction Update Request with filter constraint

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction
  service="CSW"
  version="2.0.2"
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:ogc="http://www.opengis.net/ogc">
  <csw:Update>
    <csw:RecordProperty>
      <csw:Name>title</csw:Name>
      <csw:Value>Updated Title</csw:Value>
    </csw:RecordProperty>
    <csw:RecordProperty>
      <csw:Name>date</csw:Name>
      <csw:Value>2015-08-25</csw:Value>
    </csw:RecordProperty>
    <csw:RecordProperty>
      <csw:Name>format</csw:Name>
      <csw:Value></csw:Value>
    </csw:RecordProperty>
    <csw:Constraint version="2.0.0">
      <ogc:Filter>
        <ogc:PropertyIsEqualTo>
          <ogc:PropertyName>title</ogc:PropertyName>
          <ogc:Literal>Aliquam fermentum purus quis arcu</ogc:Literal>
        </ogc:PropertyIsEqualTo>
      </ogc:Filter>
    </csw:Constraint>
  </csw:Update>
</csw:Transaction>
```

The following example shows how the newly inserted record could be updated to modify the date field while using a CQL filter constraint with title equal to **Aliquam fermentum purus quis arcu**.

Sample XML Transaction **Update** Request with CQL filter constraint

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction
  service="CSW"
  version="2.0.2"
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:ogc="http://www.opengis.net/ogc">
  <csw:Update>
    <csw:RecordProperty>
      <csw:Name>title</csw:Name>
      <csw:Value>Updated Title</csw:Value>
    </csw:RecordProperty>
    <csw:RecordProperty>
      <csw:Name>date</csw:Name>
      <csw:Value>2015-08-25</csw:Value>
    </csw:RecordProperty>
    <csw:RecordProperty>
      <csw:Name>format</csw:Name>
      <csw:Value></csw:Value>
    </csw:RecordProperty>
    <csw:Constraint version="2.0.0">
      <ogc:CqlText>
        title = 'Aliquam fermentum purus quis arcu'
      </ogc:CqlText>
    </csw:Constraint>
  </csw:Update>
</csw:Transaction>
```

Sample XML Transaction Update Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:TransactionResponse xmlns:ogc="http://www.opengis.net/ogc"
                           xmlns:gml="http://www.opengis.net/gml"
                           xmlns:ns3="http://www.w3.org/1999/xlink"
                           xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
                           xmlns:ns5="http://www.w3.org/2001/SMIL20/"
                           xmlns:dc="http://purl.org/dc/elements/1.1/"
                           xmlns:ows="http://www.opengis.net/ows"
                           xmlns:dct="http://purl.org/dc/terms/"
                           xmlns:ns9="http://www.w3.org/2001/SMIL20/Language"
                           xmlns:ns10="http://www.w3.org/2001/XMLSchema-instance"
                           ns10:schemaLocation="http://www.opengis.net/csw
/ogc/csw/2.0.2/CSW-publication.xsd"
                           version="2.0.2">
  <csw:TransactionSummary>
    <csw:totalInserted>0</csw:totalInserted>
    <csw:totalUpdated>1</csw:totalUpdated>
    <csw:totalDeleted>0</csw:totalDeleted>
  </csw:TransactionSummary>
</csw:TransactionResponse>
```

14.1.4.3.6. Transaction Delete Sub-Operation HTTP POST

The Delete sub-operation is a method to identify a set of records to be deleted from the catalog.

The following example shows a delete request for all records with a SpatialReferenceSystem name equal to [WGS-84](#).

Sample XML Transaction Delete Request with filter constraint

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction service="CSW" version="2.0.2"
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:gml="http://www.opengis.net/gml"
  xmlns:ogc="http://www.opengis.net/ogc">
  <csw:Delete typeName="csw:Record" handle="something">
    <csw:Constraint version="2.0.0">
      <ogc:Filter>
        <ogc:PropertyIsEqualTo>
          <ogc:PropertyName>SpatialReferenceSystem</ogc:PropertyName>
          <ogc:Literal>WGS-84</ogc:Literal>
        </ogc:PropertyIsEqualTo>
      </ogc:Filter>
    </csw:Constraint>
  </csw:Delete>
</csw:Transaction>
```

The following example shows a delete operation specifying a CQL constraint to delete all records

with a title equal to **Aliquam fermentum purus quis arcu**

Sample XML Transaction Delete Request with CQL filter constraint

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction service="CSW" version="2.0.2"
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:gml="http://www.opengis.net/gml"
  xmlns:ogc="http://www.opengis.net/ogc">
  <csw:Delete typeName="csw:Record" handle="something">
    <csw:Constraint version="2.0.0">
      <ogc:CqlText>
        title = 'Aliquam fermentum purus quis arcu'
      </ogc:CqlText>
    </csw:Constraint>
  </csw:Delete>
</csw:Transaction>
```

Sample XML Transaction Delete Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:TransactionResponse
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:ns10="http://www.w3.org/2001/XMLSchema-instance"
  ns10:schemaLocation="http://www.opengis.net/csw
/ogc/csw/2.0.2/CSW-publication.xsd"
  version="2.0.2">
  <csw:TransactionSummary>
    <csw:totalInserted>0</csw:totalInserted>
    <csw:totalUpdated>0</csw:totalUpdated>
    <csw:totalDeleted>1</csw:totalDeleted>
  </csw:TransactionSummary>
</csw:TransactionResponse>
```

14.1.4.3.7. Subscription **GetRecords** Operation

The subscription **GetRecords** operation is very similar to the **GetRecords** operation used to search the catalog but it subscribes to a search and sends events to a **ResponseHandler** endpoint as metacards are ingested matching the GetRecords request used in the subscription. The **ResponseHandler** must use the https protocol and receive a HEAD request to poll for availability and POST/PUT/DELETE requests for creation, updates, and deletions. The response to a **GetRecords** request on the subscription url will be an acknowledgement containing the original GetRecords request and a requestId. The client will be assigned a requestId (URN). A Subscription listens for events from federated sources if the **DistributedSearch** element is present and the catalog is a member of a federation.

14.1.4.3.8. Subscription **GetRecords** HTTP GET

GetRecords KVP (Key-Value Pairs) Encoding

```
\https://:{FQDN}:{PORT}/services/csw/subscription?service=CSW&version=2.0.2&request=GetRecords&outputFormat=application/xml&outputSchema=http://www.opengis.net/cat/csw/2.0.2&NAMESPACE=xmlns(csw=http://www.opengis.net/cat/csw/2.0.2)&resultType=results&typeName=s=csw:Record&elementSetName=brief&responseHandler=https%3A%2F%2Fsome.ddf%2Fservices%2Fcsw%2Fsubscription%2Fevent&constraintLanguage=CQL_TEXT&constraint=Text Like '%25'
```

14.1.4.3.9. Subscription GetRecords HTTP POST

Subscription GetRecords XML Request

```
<?xml version="1.0" ?>
<GetRecords xmlns="http://www.opengis.net/cat/csw/2.0.2"
    xmlns:ogc="http://www.opengis.net/ogc"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    service="CSW"
    version="2.0.2"
    maxRecords="4"
    startPosition="1"
    resultType="results"
    outputFormat="application/xml"
    outputSchema="http://www.opengis.net/cat/csw/2.0.2"
    xsi:schemaLocation="http://www.opengis.net/cat/csw/2.0.2
    ../../csw/2.0.2/CSW-discovery.xsd">
    <ResponseHandler>
        https://some.ddf/services/csw/subscription/event</ResponseHandler>
        <Query typeName="Record">
            <ElementSetName>summary</ElementSetName>
            <Constraint version="1.1.0">
                <ogc:Filter>
                    <ogc:PropertyIsLike wildCard "%" singleChar "_" escapeChar="\">\>
                        <ogc:PropertyName>xml</ogc:PropertyName>
                        <ogc:Literal>%</ogc:Literal>
                    </ogc:PropertyIsLike>
                </ogc:Filter>
            </Constraint>
        </Query>
    </GetRecords>
```

14.1.4.3.10. Subscription GetRecords HTTP PUT

The **HTTP PUT** request **GetRecords** is used to update an existing subscription. It is the same as the **POST**, except the **requestid** URN is appended to the url.

Subscription GetRecords XML Request

```
\https://:{FQDN}:{PORT}/services/csw/subscription/urn:uuid:4d5a5249-be03-4fe8-afea-6115021dd62f
```

Subscription **GetRecords** XML Response

```
<?xml version="1.0" ?>
<Acknowledgement timeStamp="2008-09-28T18:49:45" xmlns=
"http://www.opengis.net/cat/csw/2.0.2"
xmlns:ogc="http://www.opengis.net/ogc"
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
xsi:schemaLocation="http://www.opengis.net/cat/csw/2.0.2 ../../../../../../csw/2.0.2/CSW-
discover.xsd">
  <EchoedRequest>
    <GetRecords
      requestId="urn:uuid:4d5a5249-be03-4fe8-afea-6115021dd62f"
      service="CSW"
      version="2.0.2"
      maxRecords="4"
      startPosition="1"
      resultType="results"
      outputFormat="application/xml"
      outputSchema="urn:catalog:metacard">
      <ResponseHandler>
        https://some.ddf/services/csw/subscription/event</ResponseHandler>
      <Query typeNames="Record">
        <ElementSetName>summary</ElementSetName>
        <Constraint version="1.1.0">
          <ogc:Filter>
            <ogc:PropertyIsLike wildCard "%" singleChar "_" escapeChar="\\">
              <ogc:PropertyName>xml</ogc:PropertyName>
              <ogc:Literal>%</ogc:Literal>
            </ogc:PropertyIsLike>
          </ogc:Filter>
        </Constraint>
      </Query>
    </GetRecords>
  </EchoedRequest>
  <RequestId>urn:uuid:4d5a5249-be03-4fe8-afea-6115021dd62f</RequestId>
</Acknowledgement>
```

Subscription **GetRecords** event Response

The following is an example of an **application/xml** event sent to a subscribers **ResponseHandler** using an **HTTP POST** for a create, **HTTP PUT** for an update, and **HTTP DELETE** for a delete using the default **outputSchema** of <http://www.opengis.net/cat/csw/2.0.2> if you specified another supported schema format in the subscription it will be returned in that format.

Subscription **GetRecords** event XML Response

```
<csw:GetRecordsResponse version="2.0.2" xmlns:dc="http://purl.org/dc/elements/1.1/"  
  xmlns:dct="http://purl.org/dc/terms/" xmlns:ows="http://www.opengis.net/ows" xmlns:xs  
  ="http://www.w3.org/2001/XMLSchema" xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"  
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">  
  <csw:SearchStatus timestamp="2014-02-19T15:33:44.602-05:00"/>  
  <csw:SearchResults numberOfRecordsMatched="1" numberOfRecordsReturned="1"  
  nextRecord="5" recordSchema="http://www.opengis.net/cat/csw/2.0.2" elementSet="  
  summary">  
    <csw:SummaryRecord>  
      <dc:identifier>182fb33103414e5cbb06f8693b526239</dc:identifier>  
      <dc:title>Product10</dc:title>  
      <dc:type>pdf</dc:type>  
      <dct:modified>2014-02-19T15:22:51.563-05:00</dct:modified>  
      <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">  
        <ows:LowerCorner>20.0 10.0</ows:LowerCorner>  
        <ows:UpperCorner>20.0 10.0</ows:UpperCorner>  
      </ows:BoundingBox>  
    </csw:SummaryRecord>  
  </csw:SearchResults>  
</csw:GetRecordsResponse>
```

14.1.4.3.11. Subscription HTTP GET or HTTP DELETE Request

The following is an example **HTTP GET** Request to retrieve an active subscription

Subscription **HTTP GET or HTTP DELETE**

```
\https://:{FQDN}:{PORT}/services/csw/subscription/urn:uuid:4d5a5249-be03-4fe8-afea-  
6115021dd62f
```

14.1.4.3.12. Subscription **HTTP GET or HTTP DELETE Response**

The following is an example **HTTP GET** Response retrieving an active subscription

Subscription **HTTP GET or HTTP DELETE** XML Response

```
<?xml version="1.0" ?>
<Acknowledgement timeStamp="2008-09-28T18:49:45" xmlns=
"http://www.opengis.net/cat/csw/2.0.2"
                                         xmlns:ogc=""
                                         xmlns:xsi=
"http://www.w3.org/2001/XMLSchema-instance"
                                         xsi:schemaLocation=
"http://www.opengis.net/cat/csw/2.0.2 ../../../../../../csw/2.0.2/CSW-discovery.xsd">
  <EchoedRequest>
    <GetRecords
      requestId="urn:uuid:4d5a5249-be03-4fe8-afea-6115021dd62f"
      service="CSW"
      version="2.0.2"
      maxRecords="4"
      startPosition="1"
      resultType="results"
      outputFormat="application/xml"
      outputSchema="urn:catalog:metocard">
      <ResponseHandler>
        https://some.ddf/services/csw/subscription/event</ResponseHandler>
      <Query typeNames="Record">
        <ElementSetName>summary</ElementSetName>
        <Constraint version="1.1.0">
          <ogc:Filter>
            <ogc:PropertyIsLike wildCard "%" singleChar "_" escapeChar="\\">
              <ogc:PropertyName>xml</ogc:PropertyName>
              <ogc:Literal>%</ogc:Literal>
            </ogc:PropertyIsLike>
          </ogc:Filter>
        </Constraint>
      </Query>
    </GetRecords>
  </EchoedRequest>
  <RequestId>urn:uuid:4d5a5249-be03-4fe8-afea-6115021dd62f</RequestId>
</Acknowledgement>
```

14.1.4.3.13. Example Responses for CSW Endpoint Error Conditions

The following are example data and expected errors responses that will be returned for each error condition.

HTTP error codes are also returned. https://en.wikipedia.org/wiki/List_of_HTTP_status_codes#4xx_Client_errors

No Transaction Contents

This will not generate an error, but the response will tell you that nothing was processed as part of the transaction. For security purposes the **ows:ExceptionText** on invalid data is generic. The log file

should be consulted for more information.

Example CSW Request with no payload

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction service="CSW" verboseResponse="true" version="2.0.2" xmlns:csw=
"http://www.opengis.net/cat/csw/2.0.2">
</csw:Transaction>
```

No Payload CSW Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:TransactionResponse xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:csw=
"http://www.opengis.net/cat/csw/2.0.2" xmlns:gml="http://www.opengis.net/gml"
xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
"http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
instance" version="2.0.2" ns10:schemaLocation="http://www.opengis.net/csw
/ogc/csw/2.0.2/CSW-publication.xsd">
  <csw:TransactionSummary>
    <csw:totalInserted>0</csw:totalInserted>
    <csw:totalUpdated>0</csw:totalUpdated>
    <csw:totalDeleted>0</csw:totalDeleted>
  </csw:TransactionSummary>
</csw:TransactionResponse>
```

Malformed XML

The follow example sends malformed XML to the CSW Endpoint.

Example Malformed XML request

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction
  service="CSW"
  version="2.0.2"
  verboseResponse="true"
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2">
  <csw:Insert typeName="csw:Record">
    <csw:Record
      xmlns:ows="http://www.opengis.net/ows"
      xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
      xmlns:dc="http://purl.org/dc/elements/1.1/"
      xmlns:dct="http://purl.org/dc/terms/"
      xmlns:xsd="http://www.w3.org/2001/XMLSchema">
      <dc:identifier></dc:identifier>
      <dc:title>Aliquam fermentum purus quis arcu</dc:title>
      <dc:type>http://purl.org/dc/dcmitype/Text</dc:type>
      <dc:subject>Hydrography--Dictionaries</dc:subject>
      <dc:format>application/pdf</dc:format>
      <dc:date>2006-05-12</dc:date>
      <dct:abstract>Vestibulum quis ipsum sit amet metus imperdiet vehicula.
      Nulla scelerisque cursus mi.</dct:abstract>
      <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">
        <ows:LowerCorner>44.792 -6.171</ows:LowerCorner>
        <ows:UpperCorner>51.126 -2.228</ows:UpperCorner>
      </ows:BoundingBox>
    </csw:Record>
  </csw:Update>
</csw:Transaction>
```

An HTTP 400 Bad request response is returned. The error is logged in the log file and the following response body is returned.

Malformed XML CSW Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ows:ExceptionReport xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:csw=
"http://www.opengis.net/cat/csw/2.0.2" xmlns:gml="http://www.opengis.net/gml"
xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
"http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
instance" version="1.2.0" ns10:schemaLocation="http://www.opengis.net/csw
/ogc/csw/2.0.2/CSW-publication.xsd">
  <ows:Exception exceptionCode="MissingParameterValue">
    <ows:ExceptionText>Error parsing the request. XML parameters may be missing
    or invalid.</ows:ExceptionText>
  </ows:Exception>
</ows:ExceptionReport>
```

Non-CSW Request

The following example sends a non-CSW request to the CSW endpoint.

Example Non-CSW request

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<document>
  <title>boucle dampish caulkers</title>
  <id>abc123</id>
</document>
```

An HTTP 400 Bad request response is returned, and the following response body is returned.

Non-CSW Data Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ows:ExceptionReport xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:csw=
"http://www.opengis.net/cat/csw/2.0.2" xmlns:gml="http://www.opengis.net/gml"
xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
"http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
instance" version="1.2.0" ns10:schemaLocation="http://www.opengis.net/csw
/ogc/csw/2.0.2/CSW-publication.xsd">
  <ows:Exception exceptionCode="ParameterValue" locator="service">
    <ows:ExceptionText>Unknown Service</ows:ExceptionText>
  </ows:Exception>
</ows:ExceptionReport>
```

Request with Unknown Schema

This type of request will succeed and attribute names that match the expected names for the typeName (e.g. csw:Record) will get mapped into the metocard. In the example, the **title** attribute will get mapped to the metocard **title** attribute since it's the same attribute name as **<dc:title>** that **csw:Record** is configured to parse.

Example Unknown Schema

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction service="CSW" verboseResponse="true" version="2.0.2" xmlns:csw=
"http://www.opengis.net/cat/csw/2.0.2">
  <csw:Insert typeName="csw:Record">
    <csw:Record
      xmlns:ows="http://www.opengis.net/ows"
      xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
      xmlns:dc="http://purl.org/dc/elements/1.1/"
      xmlns:dct="http://purl.org/dc/terms/"
      xmlns:xsd="http://www.w3.org/2001/XMLSchema"
      xmlns:unk="http://example.com/unknown">
      <unk:id>123</unk:id>
      <unk:title>Aliquam fermentum purus quis arcu</unk:title>
    </csw:Record>
  </csw:Insert>
</csw:Transaction>
```

Metocard is created successfully.

Example Successful Unknown Schema Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:TransactionResponse xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:gml=
"http://www.opengis.net/gml" xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
  xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
  xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
  "http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
  instance" version="2.0.2" ns10:schemaLocation="http://www.opengis.net/csw
  /ogc/csw/2.0.2/CSW-publication.xsd">
  <csw:TransactionSummary>
    <csw:totalInserted>1</csw:totalInserted>
    <csw:totalUpdated>0</csw:totalUpdated>
    <csw:totalDeleted>0</csw:totalDeleted>
  </csw:TransactionSummary>
  <csw:InsertResult>
    <csw:BriefRecord>
      <dc:identifier>4ec3ec03f75344a7b4404773f97e5a03</dc:identifier>
      <dc:title>Aliquam fermentum purus quis arcu</dc:title>
      <dc:type/>
    </csw:BriefRecord>
  </csw:InsertResult>
</csw:TransactionResponse>
```

Well-formed, but Invalid TypeName

The `typeName` on the `csw:Insert` specifies the transformer to use when parsing the data. If the name specified is not configured, an error response is returned.

Example Invalid typeName

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:Transaction service="CSW" verboseResponse="true" version="2.0.2" xmlns:csw=
"http://www.opengis.net/cat/csw/2.0.2">
  <csw:Insert typeName="invalid-data">
    <root>
      <id>abcd16df29413796b388b02ee017a315</id>
    </document>
  </csw:Insert>
</csw:Transaction>
```

An HTTP 400 Bad request response is returned. The error is logged in the log file and the following response body is returned.

Invalid typeName Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ows:ExceptionReport xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:csw=
"http://www.opengis.net/cat/csw/2.0.2" xmlns:gml="http://www.opengis.net/gml"
xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
"http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
instance" version="1.2.0" ns10:schemaLocation="http://www.opengis.net/csw
/ogc/csw/2.0.2/CSW-publication.xsd">
  <ows:Exception exceptionCode="MissingParameterValue">
    <ows:ExceptionText>Error parsing the request. XML parameters may be missing
or invalid.</ows:ExceptionText>
  </ows:Exception>
</ows:ExceptionReport>
```

Request with Missing XML Prologue

The following example sends XML data to the CSW Endpoint without the XML prologue.

Example Missing XML Tag

```
<csw:Transaction
  service="CSW"
  version="2.0.2"
  verboseResponse="true"
  xmlns:csw="http://www.opengis.net/cat/csw/2.0.2">
  <csw:Insert typeName="csw:Record">
    <csw:Record
      xmlns:ows="http://www.opengis.net/ows"
      xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
      xmlns:dc="http://purl.org/dc/elements/1.1/"
      xmlns:dct="http://purl.org/dc/terms/"
      xmlns:xsd="http://www.w3.org/2001/XMLSchema">
      <dc:identifier></dc:identifier>
      <dc:title>Aliquam fermentum purus quis arcu</dc:title>
      <dc:type>http://purl.org/dc/dcmitype/Text</dc:type>
      <dc:subject>Hydrography--Dictionaries</dc:subject>
      <dc:format>application/pdf</dc:format>
      <dc:date>2006-05-12</dc:date>
      <dct:abstract>Vestibulum quis ipsum sit amet metus imperdiet vehicula.
      Nulla scelerisque cursus mi.</dct:abstract>
      <ows:BoundingBox crs="urn:x-ogc:def:crs:EPSG:6.11:4326">
        <ows:LowerCorner>44.792 -6.171</ows:LowerCorner>
        <ows:UpperCorner>51.126 -2.228</ows:UpperCorner>
      </ows:BoundingBox>
    </csw:Record>
  </csw:Insert>
</csw:Transaction>
```

Metocard is created successfully.

Example Missing XML Tag Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<csw:TransactionResponse xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:gml=
"http://www.opengis.net/gml" xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
"http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
instance" version="2.0.2" ns10:schemaLocation="http://www.opengis.net/csw
/ogc/csw/2.0.2/CSW-publication.xsd">
  <csw:TransactionSummary>
    <csw:totalInserted>1</csw:totalInserted>
    <csw:totalUpdated>0</csw:totalUpdated>
    <csw:totalDeleted>0</csw:totalDeleted>
  </csw:TransactionSummary>
  <csw:InsertResult>
    <csw:BriefRecord>
      <dc:identifier>c318d32e9c9a4bb5b1cd00bc1aaf704</dc:identifier>
      <dc:title>Aliquam fermentum purus quis arcu</dc:title>
      <dc:type>http://purl.org/dc/dcmitype/Text</dc:type>
      <ows:BoundingBox crs="EPSG:4326">
        <ows:LowerCorner>44.792 -6.171</ows:LowerCorner>
        <ows:UpperCorner>51.126 -2.228</ows:UpperCorner>
      </ows:BoundingBox>
    </csw:BriefRecord>
  </csw:InsertResult>
</csw:TransactionResponse>
```

Request with Non-XML Data

The following is a non-XML request sent to the CSW Endpoint.

Non-XML data Example

```
title: Non-XML title
id: abc123
```

An HTTP 400 Bad request response is returned. The error is logged in the log file and the following response body is returned.

Non-XML Response

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ows:ExceptionReport xmlns:ows="http://www.opengis.net/ows" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ogc="http://www.opengis.net/ogc" xmlns:gml=
"http://www.opengis.net/gml" xmlns:csw="http://www.opengis.net/cat/csw/2.0.2"
xmlns:ns6="http://www.w3.org/2001/SMIL20/" xmlns:dc="http://purl.org/dc/elements/1.1/"
xmlns:dct="http://purl.org/dc/terms/" xmlns:ns9=
"http://www.w3.org/2001/SMIL20/Language" xmlns:ns10="http://www.w3.org/2001/XMLSchema-
instance" version="1.2.0" ns10:schemaLocation="http://www.opengis.net/csw
/ogc/csw/2.0.2/CSW-publication.xsd">
    <ows:Exception exceptionCode="MissingParameterValue">
        <ows:ExceptionText>Error parsing the request. XML parameters may be missing
        or invalid.</ows:ExceptionText>
    </ows:Exception>
</ows:ExceptionReport>
```

14.1.5. FTP Endpoint

The FTP Endpoint provides a method for ingesting files directly into the DDF Catalog using the FTP protocol. The files sent over FTP are not first written to the file system, like the Directory Monitor, but instead the FTP stream of the file is ingested directly into the DDF catalog, thus avoiding extra I/O overhead.

14.1.5.1. Installing the FTP Endpoint

The FTP Endpoint is not installed by default with a standard installation.

To install:

- Navigate to Admin Console.
- Navigate to **Catalog** application.
- Navigate to **Features** tab.
- Install the **catalog-core-ftp** feature.

14.1.5.2. Configuring the FTP Endpoint

Once installed, the configurable properties for the FTP Endpoint are accessed from the **FTP Endpoint** Configuration:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **FTP Endpoint**.

See [FTP Endpoint configurations](#) for all possible configurations.

14.1.5.3. Using FTP Endpoint

FTP Endpoint URL

```
ftp://<FQDN>:8021/
```

The FTP endpoint supports the **PUT**, **MPUT**, **DELE**, **RETR**, **RMD**, **APPE**, **RNTO**, **STOU**, and **SITE** operations.

The FTP endpoint supports files being uploaded as a dot-file (e.g., `.foo`) and then being renamed to the final filename (e.g., `some-file.pdf`). The endpoint will complete the ingest process when the rename command is sent.

14.1.5.3.1. From Code:

Custom Ftplets can be implemented by extending the `DefaultFtplet` class provided by Apache FTP Server. Doing this will allow custom handling of various FTP commands by overriding the methods of the `DefaultFtplet`. Refer to <https://mina.apache.org/ftpserver-project/ftplet.html> for available methods that can be overridden. After creating a custom Ftplet, it needs to be added to the FTP server's Ftplets before the server is started. Any Ftplets that are registered to the FTP server will execute the FTP command in the order that they were registered.

14.1.5.3.2. From an FTP client:

The FTP endpoint can be accessed from any FTP client of choice. Some common clients are FileZilla, PuTTY, or the FTP client provided in the terminal. The default port number is **8021**. If FTPS is enabled with 2-way TLS, a client that supports client authentication is required.

14.1.6. KML Endpoint

Keyhole Markup Language (*KML*) is an XML notation for describing geographic annotation and visualization for 2- and 3- dimensional maps.

The KML Network Link endpoint allows a user to generate a view-based KML Query Results Network Link. This network link can be opened with Google Earth, establishing a dynamic connection between Google Earth and DDF. The root network link will create a network link for each configured source, including the local catalog. The individual source network links will perform a query against the OpenSearch Endpoint periodically based on the current view in the KML client. The query parameters for this query are obtained by a bounding box generated by Google Earth. The root network link will refresh every 12 hours or can be forced to refresh. As a user changes their current view, the query will be re-executed with the bounding box of the new view. (This query gets re-executed two seconds after the user stops moving the view.)

14.1.6.1. Installing the KML Endpoint

The KML Network Link Endpoint is installed by default with a standard installation as part of the Spatial application.

14.1.6.2. Configuring the KML Endpoint

This KML Network Link endpoint has the ability to serve up custom KML style documents and Icons

to be used within that document. The KML style document must be a valid XML document containing a KML style. The KML Icons should be placed in a single level directory and must be an image type (png, jpg, tif, etc.). The Description will be displayed as a pop-up from the root network link on Google Earth. This may contain the general purpose of the network and URLs to external resources.

See [KML Endpoint configurations](#) for all possible configurations.

14.1.6.3. Using the KML Endpoint

Once installed, the KML Network Link endpoint can be accessed at:

```
\https://{{FQDN}}:{{PORT}}/services/catalog/kml
```

After the above request is sent, a KML Network Link document is returned as a response to download or open. This KML Network Link can then be opened in Google Earth.

Example Output from KML Endpoint

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<kml xmlns="http://www.opengis.net/kml/2.2" xmlns:ns2=
"http://www.google.com/kml/ext/2.2"
      xmlns:ns3="http://www.w3.org/2005/Atom" xmlns:ns4=
"urn:oasis:names:tc:ciq:xsdschema:xAL:2.0">
  <NetworkLink>
    <name>DDF</name>
    <open xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" xmlns:xs=
"http://www.w3.org/2001/XMLSchema" xsi:type="xs:boolean">true</open>
    <Snippet maxLines="0"/>
    <Link>
      <href>http://0.0.0.0:8181/services/catalog/kml/sources</href>
      <refreshMode>onInterval</refreshMode>
      <refreshInterval>43200.0</refreshInterval>
      <viewRefreshMode>never</viewRefreshMode>
      <viewRefreshTime>0.0</viewRefreshTime>
      <viewBoundScale>0.0</viewBoundScale>
    </Link>
  </NetworkLink>
</kml>
```

The KML endpoint can also serve up Icons to be used in conjunction with the KML style document. The request below shows the format to return an icon.

NOTE `<icon-name>` must be the name of an icon contained in the directory being served.

Return KML Icon

```
\https://{{FQDN}}:{{PORT}}/services/catalog/kml/icons?<icon-name>
```

14.1.7. Metrics Endpoint

NOTE EXPERIMENTAL

WARNING Note that the Metrics endpoint URL is marked "internal." This indicates that this endpoint is intended for internal use by the DDF code. This endpoint is subject to change in future versions.

The Metrics Endpoint is used by the [Metrics Collection Application](#) to report on system metrics.

14.1.7.1. Installing the Metrics Endpoint

The Metrics Endpoint is installed by default with a standard installation in the Platform application.

14.1.7.2. Configuring the Metrics Endpoint

No configuration can be made for the Metrics Endpoint. All of the metrics that it collects data on are pre-configured in DDF.

14.1.7.3. Using the Metrics Endpoint

Metrics Endpoint URL

```
\https://:{FQDN}:{PORT}/services/internal/metrics/catalogQueries.png?startDate=2013-03-31T06:00:00-07:00&endDate=2013-04-01T11:00:00-07:00
```

The table below lists all of the options for the Metrics endpoint URL to execute custom metrics data requests:

Table 28. Metrics Endpoint URL Options

Parameter	Description	Example	Required
<code>startDate</code>	Specifies the start of the time range of the search on the metric's data (RFC-3339 - Date and Time format, i.e. YYYY-MM-DDTHH:mm:ssZ). Date/time must be earlier than the endDate. <i>This parameter cannot be used with the dateOffset parameter.</i>	<code>startDate=2013-03-31T06:00:00-07:00</code>	true
<code>endDate</code>	Specifies the end of the time range of the search on the metric's data (RFC-3339 - Date and Time format, i.e. YYYY-MM-DDTHH:mm:ssZ). Date/time must be later than the startDate. <i>This parameter cannot be used with the dateOffset parameter.</i>	<code>endDate=2013-04-01T11:00:00-07:00</code>	true

Parameter	Description	Example	Required
<code>dateOffset</code>	Specifies an offset, backwards from the current time, to search on the modified time field for entries. Defined in seconds and must be a positive Integer. <i>This parameter cannot be used with the <code>startDate</code> or <code>endDate</code> parameters.</i>	<code>dateOffset=1800</code>	true
<code>yAxisLabel</code>	The label to apply to the graph's y-axis. Will default to the metric's name, e.g., Catalog Queries. <i>This parameter is only applicable for the metric's graph display format.</i>	Catalog Query Count	false
<code>title</code>	The title to be applied to the graph. + Will default to the metric's name plus the time range used for the graph. + <i>This parameter is only applicable for the metric's graph display format.</i>	Catalog Query Count for the last 15 minutes	false

14.1.7.3.1. Metrics Data Supported Formats

The metric's historical data can be displayed in several formats, including PNG, a CSV file, an Excel .xls file, a PowerPoint .ppt file, an XML file, and a JSON file. The PNG, CSV, and XLS formats are accessed via hyperlinks provided in the Metrics tab web page. The PPT, XML, and JSON formats are accessed by specifying the format in the custom URL, e.g., <http://{FQDN}:{PORT}/services/internal/metrics/catalogQueries.json?dateOffset=1800>.

The table below describes each of the supported formats, how to access them, and an example where applicable. (NOTE: all example URLs begin with <https://{FQDN}:{PORT}> which is omitted in the table for brevity.)

Table 29. Metrics Formats

Display Format	Description	How To Access	Example URL
PNG	Displays the metric's data as a PNG-formatted graph, where the x-axis is time and the y-axis is the metric's sampled data values.	Via hyperlink on the Metrics tab or directly via custom URL.	<p>Accessing Catalog Queries metric data for last 8 hours (8 * 60 * 60 = 28800 seconds):</p> <p>/services/internal/metrics/catalogQueries.png?dateOffset=28800&yAxisLabel=mylabel&title=mygraphtitle</p> <p>Accessing Catalog Queries metric data between 6:00 am on March 10, 2013, and 10:00 am on April 2, 2013:</p> <p>/services/internal/metrics/catalogQueries.png?startDate=2013-03-10T06:00:00-07:00&endDate=2013-04-02T10:00:00-07:00&yAxisLabel=mylabel&title=mygraphtitle</p> <p><i>Note that the <code>yAxisLabel</code> and <code>title</code> parameters are optional.</i></p>

Display Format	Description	How To Access	Example URL
CSV	<p>Displays the metric's data as a Comma-Separated Value (CSV) file, which can be auto-displayed in Excel based on browser settings.</p> <p>The generated CSV file will consist of two columns of data: Timestamp and Value, where the first row contains the column headers and the remaining rows contain the metric's sampled data over the specified time range.</p>	Via hyperlink on the Metrics tab or directly via custom URL.	<p>Accessing Catalog Queries metric data for last 8 hours (8 * 60 * 60 = 28800 seconds):</p> <p>/services/internal/metrics/catalogQueries.csv?dateOffset=28800</p> <p>Accessing Catalog Queries metric data between 6:00 am on March 10, 2013, and 10:00 am on April 2, 2013:</p> <p>/services/internal/metrics/catalogQueries.csv?startDate=2013-03-10T06:00:00-07:00&endDate=2013-04-02T10:00:00-07:00</p>
XLS	<p>Displays the metric's data as an Excel (XLS) file, which can be auto-displayed in Excel based on browser settings. The generated XLS file will consist of: Title in first row based on metric's name and specified time range Column headers for Timestamp and Value; Two columns of data containing the metric's sampled data over the specified time range; The total count, if applicable, in the last row</p>	Via hyperlink on the Metrics tab or directly via custom URL.	<p>Accessing Catalog Queries metric data for last 8 hours (8 * 60 * 60 = 28800 seconds):</p> <p>/services/internal/metrics/catalogQueries.xls?dateOffset=28800</p> <p>Accessing Catalog Queries metric data between 6:00 am on March 10, 2013, and 10:00 am on April 2, 2013:</p> <p>/services/internal/metrics/catalogQueries.xls?startDate=2013-03-10T06:00:00-07:00&endDate=2013-04-02T10:00:00-07:00</p>
PPT	<p>Displays the metric's data as a PowerPoint (PPT) file, which can be auto-displayed in PowerPoint based on browser settings. The generated PPT file will consist of a single slide containing: A title based on the metric's name; The metric's PNG graph embedded as a picture in the slide The total count, if applicable</p>	Via custom URL only	<p>Accessing Catalog Queries metric data for last 8 hours (8 * 60 * 60 = 28800 seconds):</p> <p>/services/internal/metrics/catalogQueries.ppt?dateOffset=28800</p> <p>Accessing Catalog Queries metric data between 6:00 am on March 10, 2013, and 10:00 am on April 2, 2013:</p> <p>/services/internal/metrics/catalogQueries.ppt?startDate=2013-03-10T06:00:00-07:00&endDate=2013-04-02T10:00:00-07:00</p>

Display Format	Description	How To Access	Example URL
XML	Displays the metric's data as an XML-formatted file.	via custom URL only	<p>Accessing Catalog Queries metric data for last 8 hours ($8 * 60 * 60 = 28800$ seconds):</p> <p>/services/internal/metrics/catalogQueries.xml?dateOffset=28800</p> <p>Accessing Catalog Queries metric data between 6:00 am on March 10, 2013, and 10:00 am on April 2, 2013:</p> <p>/services/internal/metrics/catalogQueries.xml?startDate=2013-03-10T06:00:00-07:00&endDate=2013-04-02T10:00:00-07:00</p> <p>See Sample XML-formatted output.</p>
JSON	Displays the metric's data as an JSON-formatted file.	via custom URL only	<p>Accessing Catalog Queries metric data for last 8 hours ($8 * 60 * 60 = 28800$ seconds):</p> <p>/services/internal/metrics/catalogQueries.json?dateOffset=28800</p> <p>Accessing Catalog Queries metric data between 6:00 am on March 10, 2013, and 10:00 am on April 2, 2013:</p> <p>/services/internal/metrics/catalogQueries.json?startDate=2013-03-10T06:00:00-07:00&endDate=2013-04-02T10:00:00-07:00</p> <p>See Sample JSON-Formatted Output.</p>

Sample XML-Formatted Output

```

<catalogQueries>
  <title>Catalog Queries for Apr 15 2013 08:45:53 to Apr 15 2013 09:00:53</title>
  <data>
    <sample>
      <timestamp>Apr 15 2013 08:45:00</timestamp>
      <value>361</value>
    </sample>
    <sample>
      <timestamp>Apr 15 2013 09:00:00</timestamp>
      <value>353</value>
    </sample>
    <totalCount>5721</totalCount>
  </data>
</catalogQueries>

```

Sample JSON-formatted Output

```
{  
  "title": "Query Count for Jul 9 1998 09:00:00 to Jul 9 1998 09:50:00",  
  "totalCount": 322,  
  "data": [  
    {  
      "timestamp": "Jul 9 1998 09:20:00",  
      "value": 54  
    },  
    {  
      "timestamp": "Jul 9 1998 09:45:00",  
      "value": 51  
    }  
  ]  
}
```

14.1.7.3.2. Add Custom Metrics to the Metrics Tab

It is possible to add custom (or existing, but non-collected) metrics to the Metrics tab by writing an application. Refer to the SDK example source code for Sample Metrics located in the DDF source code at [sdk/sample-metrics](#) and [sdk/sdk-app](#).

WARNING

The Metrics framework is not an open API, but rather a closed, internal framework that can change at any time in future releases. Be aware that any custom code written may not work with future releases.

14.1.7.4. Usage Limitations of the Metrics Endpoint

The Metrics Collecting Application uses a “round robin” database. It uses one that does not store individual values but, instead, stores the rate of change between values at different times. Due to the nature of this method of storage, along with the fact that some processes can cross time frames, small discrepancies (differences in values of one or two have been experienced) may appear in values for different time frames. These will be especially apparent for reports covering shorter time frames such as 15 minutes or one hour. These are due to the averaging of data over time periods and should not impact the values over longer periods of time.

14.1.8. OpenSearch Endpoint

The OpenSearch Endpoint enables a client to send query parameters and receive search results. This endpoint uses the input query parameters to create an OpenSearch query. The client does not need to specify all of the query parameters, only the query parameters of interest.

14.1.8.1. Installing the OpenSearch Endpoint

The OpenSearch Endpoint is installed by default with a standard installation in the Catalog application.

14.1.8.2. Configuring the OpenSearch Endpoint

The OpenSearch Endpoint has no configurable properties. It can only be installed or uninstalled.

14.1.8.3. OpenSearch URL

```
\https://{{FQDN}}:{{PORT}}/services/catalog/query
```

14.1.8.3.1. From Code:

The OpenSearch specification defines a file format to describe an OpenSearch endpoint. This file is XML-based and is used to programmatically retrieve a site's endpoint, as well as the different parameter options a site holds. The parameters are defined via the [OpenSearch](#) and [CDR IPT](#) Specifications.

14.1.8.3.2. From a Web Browser:

Many modern web browsers currently act as OpenSearch clients. The request call is an HTTP GET with the query options being parameters that are passed.

Example of an OpenSearch request:

```
\http://{{FQDN}}:{{PORT}}/services/catalog/query?q=Predator
```

This request performs a full-text search for the phrase 'Predator' on the DDF providers and provides the results as Atom-formatted XML for the web browser to render.

14.1.8.3.3. Parameter List

Table 30. Main OpenSearch Standard

OS Element	HTTP Parameter	Possible Values	Comments
searchTerms	q	URL-encoded string	Complex contextual search string.
count	count	integer ≥ 0	Maximum # of results to retrieve default: 10
startIndex	start	integer ≥ 1	Index of first result to return. default: 1 This value uses a one based index for the results.

OS Element	HTTP Parameter	Possible Values	Comments
format	format	<p>requires a transformer shortname as a string, possible values include, when available</p> <p>atom</p> <p>html</p> <p>kml</p> <p>see Included Query Response Transformers for more possible values.</p>	default: atom

Table 31. Temporal Extension

OS Element	HTTP Parameter	Possible Values	Comments
start	dtstart	RFC-3399-defined value	yyyy-MM-dd'T'HH:mm:ss.SSSZ
end	dtend	RFC-3399-defined value	yyyy-MM-dd'T'HH:mm:ss.SSSZ

The start and end temporal criteria must be of the format specified above. Other formats are currently not supported. Example:

NOTE `2011-01-01T12:00:00.111-04:00.`

The start and end temporal elements are based on modified timestamps for a metocard.

Geospatial Extension

These geospatial query parameters are used to create a geospatial **INTERSECTS** query, where **INTERSECTS** means geometries that are not **DISJOINT** of the given geospatial parameter.

OS Element	HTTP Parameter	Possible Values	Comments
lat	lat	EPSG:4326 decimal degrees	Expects a latitude and a radius to be specified.
lon	lon	EPSG:4326 decimal degrees	Expects a longitude and a radius to be specified.
radius	radius	Meters along the Earth's surface > 0	Used in conjunction with lat and lon query parameters.

OS Element	HTTP Parameter	Possible Values	Comments
polygon	polygon	clockwise <code>lat lon</code> pairs ending at the first one	example: <code>-80, -170, 0, -170, 80, -170, 80, 170, 0, 170, -80, 170, -80, -170</code> According to the OpenSearch Geo Specification this is deprecated . Use <code>geometry</code> instead.
box	bbox	4 comma-separated EPSG:4326 decimal degrees	west, south, east, north
geometry	geometry	WKT Geometries: <code>POINT</code> , <code>POLYGON</code> , <code>MULTIPOINT</code> , <code>MULTIPOLYGON</code>	Examples: <code>POINT(10 20)</code> where 10 is the longitude and 20 is the latitude. <code>POLYGON ((30 10, 10 20, 20 40, 40 40, 30 10)).</code> 30 is longitude and 10 is latitude for the first point. Make sure to repeat the starting point as the last point to close the polygon.

Table 32. Extensions

OS Element	HTTP Parameter	Possible Values	Comments
sort	sort	<code>sbfield</code> : 'date' or 'relevance' <code>sborder</code> : 'asc' or 'desc'	<code>sort=<sbfield>:<sborder></code> default: <code>relevance:desc</code> Sorting by date will sort the effective date.
maxResults	mr	Integer ≥ 0	Maximum # of results to return. If count is also specified, the count value will take precedence over the <code>maxResults</code> value

OS Element	HTTP Parameter	Possible Values	Comments
maxTimeout	mt	Integer > 0	Maximum timeout (milliseconds) for query to respond default: 300000 (5 minutes)

Table 33. Federated Search

OS Element	HTTP Parameter	Possible Values	Comments
routeTo	src	(varies depending on the names of the sites in the federation)	comma delimited list of site names to query. Also can specify <code>src=local</code> to query the local site. If src is not provided, the default behavior is to execute an enterprise search to the entire federation.

Table 34. DDF Extensions

OS Element	HTTP Parameter	Possible Values	Comments
dateOffset	dtoffset	integer > 0	Specifies an offset, backwards from the current time, to search on the modified time field for entries. Defined in milliseconds.
type	type	Any valid datatype	Specifies the type of data to search for.
version	version	20,30	Comma-delimited list of version values to search for.
selector	selector	//namespace:example,//example	Comma-delimited list of XPath string selectors that narrow down the search.

14.1.8.3.4. Supported Complex Contextual Query Format

The OpenSearch Endpoint supports the following operators: `AND`, `OR`, and `NOT`. These operators are case sensitive. Implicit `ANDs` are also supported.

Using parentheses to change the order of operations is supported. Using quotes to group keywords

into literal expressions is supported.

See the [OpenSearch](#) specification for more syntax specifics.

14.2. Endpoint Utility Services

DDF also provides a variety of services to enhance the function of the endpoints.

14.2.1. Compression Services

DDF supports compression of outgoing and incoming messages through the Compression Services. These compression services are based on [CXF](#) message encoding.

The formats supported in DDF are:

gzip

Adds GZip compression to messages through CXF components. Code comes with CXF.

exi

Adds [Efficient XML Interchange \(EXI\)](#) support to outgoing responses. EXI is an W3C standard for XML encoding that shrinks xml to a smaller size than normal GZip compression.

14.2.1.1. Installing a Compression Service

The compression services are not installed by default with a standard installation.

To Install:

- Navigate to Admin Console.
- Select the **Platform Application**.
- Select the **Features** Tab.
- Start the service for the desired compression format:
 - **compression-exi**
 - **compression-gzip**

WARNING

The compression services either need to be installed BEFORE the desired CXF service is started or the CXF service needs to be refreshed / restarted after the compression service is installed.

14.2.1.2. Configuring Compression Services

None.

Compression Imported Services


None.

Table 35. Compression Exported Services

Registered Interface	Implemented Class(es)	Service Property	Value
org.apache.cxf.feature.Feature	ddf.compression.exi.EXIFeature org.apache.cxf.transport.common.gzip.GZIPFeature	N/A	N/A

15. Eventing

15.1. Eventing Intro



Eventing Architecture

The Eventing capability of the Catalog allows endpoints (and thus external users) to create a "standing query" and be notified when a matching metocard is created, updated, or deleted.

Notably, the Catalog allows event evaluation on both the previous value (if available) and new value of a Metocard when an update occurs.

Eventing allows DDFs to receive events on operations (e.g. create, update, delete) based on particular queries or actions. Once subscribed, users will receive notifications of events such as update or create on any source.

15.2. Eventing Components

The key components of DDF Eventing include:

- Subscription
- Delivery Method
- Event Processor

After reading this section, you will be able to:

- Create new subscriptions
- Register subscriptions
- Perform operations on event notification
- Remove a subscription

15.3. Subscriptions

15.3.1. Subscriptions

Subscriptions represent "standing queries" in the Catalog. Like a query, subscriptions are based on the OGC Filter specification.

15.3.1.1. Subscription Lifecycle

A Subscription itself is a series of events during which various plugins or transformers can be called to process the subscription.

15.3.1.1.1. Creation

- Subscriptions are created directly with the [Event Processor](#) or declaratively through use of the Whiteboard Design Pattern.
- The Event Processor will invoke each Pre-Subscription Plugin and, if the subscription is not rejected, the subscription will be activated.

15.3.1.1.2. Evaluation

- When a metocard matching the subscription is created, updated, or deleted in any Source, each Pre-Delivery Plugin will be invoked.
- If the delivery is not rejected, the associated Delivery Method callback will be invoked.

15.3.1.1.3. Update Evaluation

Notably, the Catalog allows event evaluation on both the previous value (if available) and new value of a Metocard when an update occurs.

15.3.1.1.4. Durability

Subscription durability is not provided by the Event Processor. Thus, all subscriptions are transient and will not be recreated in the event of a system restart. It is the responsibility of Endpoints using subscriptions to persist and re-establish the subscription on startup. This decision was made for the sake of simplicity, flexibility, and the inability of the Event Processor to recreate a fully-configured Delivery Method without being overly restrictive.

IMPORTANT

Subscriptions are not persisted by the Catalog itself.

Subscriptions must be explicitly persisted by an endpoint and are not persisted by the Catalog. The Catalog Framework, or more specifically the Event Processor itself, does not persist subscriptions. Certain endpoints, however, can persist the subscriptions on their own and recreate them on system startup.

15.3.2. Creating a Subscription

Currently, the Catalog reference implementation does not contain a subscription endpoint. Therefore, an endpoint that exposes a web service interface to create, update, and delete subscriptions would provide a client's subscription filtering criteria to be used by Catalog's Event Processor to determine which events are of interest to the client. The endpoint client also provides the callback URL of the event consumer to be called when an event matching the subscription's criteria is found. This callback to the event consumer is made by a Delivery Method implementation that the client provides when the subscription is created. Whenever an event occurs in the Catalog matching the subscription, the Delivery Method implementation will be called by the Event Processor. The Delivery Method will, in turn, send the event notification out to the event consumer. As part of the subscription creation process, the Catalog verifies that the event consumer at the specified callback URL is available to receive callbacks. Therefore, the client must ensure the event consumer is running prior to creating the subscription. The Catalog completes the subscription creation by executing any pre-subscription Catalog Plugins, and then registering the subscription with the OSGi Service Registry. The Catalog does not persist subscriptions by default.

15.3.2.1. Event Processing and Notification

If an event matches a subscription's criteria, any pre-delivery plugins that are installed are invoked, the subscription's `DeliveryMethod` is retrieved, and its operation corresponding to the type of ingest event is invoked. For example, the `DeliveryMethod created()` function is called when a metocard is created. The `DeliveryMethod` operations subsequently invoke the corresponding operation in the client's event consumer service, which is specified by the callback URL provided when the `DeliveryMethod` was created. An internal subscription tracker monitors the OSGi registry, looking for subscriptions to be added (or deleted). When it detects a subscription being added, it informs the Event Processor, which sets up the subscription's filtering and is responsible for posting event notifications to the subscriber when events satisfying their criteria are met.

The Standard Event Processor is an implementation of the Event Processor and provides the ability to create/delete subscriptions. Events are generated by the CatalogFramework as metacards are created/updated/deleted and the Standard Event Processor is called since it is also a Post-Ingest Plugin. The Standard Event Processor checks each event against each subscription's criteria.

When an event matches a subscription's criteria the Standard Event Processor:

- invokes each pre-delivery plugin on the metocard in the event.
- invokes the `DeliveryMethod` operation corresponding to the type of event being processed, e.g., `created()` operation for the creation of a metocard.

- Standard Event Processor
- Fanout Event Processor

15.3.2.1.1. Using DDF Implementation

If applicable, the implementation of `Subscription` that comes with DDF should be used. It is available at `ddf.catalog.event.impl.SubscriptionImpl` and offers a constructor that takes in all of the necessary objects. Specifically, all that is needed is a `Filter`, `DeliveryMethod`, `Set<String>` of source IDs, and a `boolean` for enterprise.

The following is an example code stub showing how to create a new instance of `Subscription` using the DDF implementation.

Creating a Subscription

```
// Create a new filter using an imported FilterBuilder
Filter filter = filterBuilder.attribute(Metocard.ANY_TEXT).like().text("*");

// Create a implementation of DeliveryMethod
DeliveryMethod deliveryMethod = new MyCustomDeliveryMethod();

// Create a set of source ids
// This set is empty as the subscription is not specific to any sources
Set<String> sourceIds = new HashSet<String>();

// Set the isEnterprise boolean value
// This subscription example should notifications from all sources (not just local)
boolean isEnterprise = true;

Subscription subscription = new SubscriptionImpl(filter, deliveryMethod, sourceIds
, isEnterprise);
```

15.3.2.2. Delivery Method

A Delivery Method provides the operation (created, updated, deleted) for how an event's metocard can be delivered.

A Delivery Method is associated with a subscription and contains the callback URL of the event consumer to be notified of events. The Delivery Method encapsulates the operations to be invoked by the Event Processor when an event matches the criteria for the subscription. The Delivery Method's operations are responsible for invoking the corresponding operations on the event consumer associated with the callback URL.

16. Security Services

16.1. Encryption Service

DDF includes an encryption service to encrypt plain text such as passwords.

16.1.1. Encryption Command

An encrypt security command is provided with DDF to encrypt text. This is useful when displaying password fields to users.

Below is an example of the `security:encrypt` command used to encrypt the plain text `myPasswordToEncrypt`. The output is the encrypted value.

security:encrypt Command Example

```
ddf@local>security:encrypt myPasswordToEncrypt
```

security:encrypt Command Output

```
ddf@local>bR9mJpDV08bTRwqGwIFxHJ5yFJzatKwjXjIo/8USWm8=
```

16.2. Expansion Service

The expansion service defines rulesets to map metocard attributes and user attributes to more complete sets of values. For example if a user has an attribute "alphabet" that contained the value "full", the expansion service can be configured to expand the "full" value out to ["a","b","c",...].

16.2.1. Configuring Expansion Service

To use the expansion service, modify the following two files within the `<DDF_HOME>/etc` directory:

- `<DDF_HOME>/etc/pdp/ddf-metocard-attribute-ruleset.cfg`
- `<DDF_HOME>/etc/pdp/ddf-user-attribute-ruleset.cfg`

Within these files, the following configuration details will be defined.

16.2.1.1. Expansion Service Instances and Configuration

It is expected that multiple instances of the expansion service will be running at the same time. Each instance of the service defines a unique property that is useful for retrieving specific instances of the expansion service. The following table lists the two pre-defined instances used by DDF for expanding user attributes and metocard attributes respectively.

Property Name	Value	Description
mapping	<code>security.user.attribute.mapping</code>	This instance is configured with rules that expand the user's attribute values for security checking.
mapping	<code>security.metocard.attribute.mapping</code>	This instance is configured with rules that expand the metocard's security attributes before comparing with the user's attributes.

Each instance of the expansion service can be configured using a configuration file. The configuration file can have three different types of lines: comments:: any line prefixed with the `#`

character is ignored as a comment (for readability, blank lines are also ignored) attribute separator:: a line starting with **separator=** defines the attribute separator string. rule:: all other lines are assumed to be rules defined in a string format **<key>:<original value>:<new value>**

The following configuration file defines the rules shown above in the example table (using the space as a separator):

```
# This defines the separator that will be used when the expansion string contains
multiple
# values - each will be separated by this string. The expanded string will be split at
the
# separator string and each resulting attributed added to the attribute set
(duplicates are
# suppressed). No value indicates the defualt value of ' ' (space).
separator=

# The following rules define the attribute expansion to be performed. The rules are of
the
# form:
#      <attribute name>:<original value>:<expanded value>
# The rules are ordered, so replacements from the first rules may be found in the
original
# values of subsequent rules.
Location:Goodyear:Goodyear AZ
Location:AZ:AZ USA
Location:CA:CA USA
Title:VP-Sales:VP-Sales VP Sales
Title:VP-Engineering:VP-Engineering VP Engineering
```

Table 36. Expansion Commands

Title	Namespace	Description
DDF::Security::Expansion::Commands	security	The expansion commands provide detailed information about the expansion rules in place and the ability to see the results of expanding specific values against the active ruleset.

Command	Description
security:expand	Runs the expansion service on the provided data returning the expanded value.
security:expansions	Dumps the ruleset for each active expansion service.

16.2.1.2. Expansion Command Examples and Explanation

16.2.1.2.1. **security:expansions**

The **security:expansions** command dumps the ruleset for each active expansion service. It takes no arguments and displays each rule on a separate line in the form: **<attribute name> : <original string> : <expanded string>**. The following example shows the results of executing the expansions

command with no active expansion service.

```
ddf@local>security:expansions
No expansion services currently available.
```

After installing the expansions service and configuring it with an appropriate set of rules, the expansions command will provide output similar to the following:

```
ddf@local>security:expansions
Location : Goodyear : Goodyear AZ
Location : AZ : AZ USA
Location : CA : CA USA
Title : VP-Sales : VP-Sales VP Sales
Title : VP-Engineering : VP-Engineering VP Engineering
```

16.2.1.2.2. security:expand

The security:expand command runs the expansion service on the provided data. It takes an attribute and an original value, expands the original value using the current expansion service and ruleset and dumps the results. For the ruleset shown above, the expand command produces the following results:

```
ddf@local>security:expand Location Goodyear
[Goodyear, USA, AZ]

ddf@local>security:expand Title VP-Engineering
[VP-Engineering, Engineering, VP]

ddf@local>expand Title "VP-Engineering Manager"
[VP-Engineering, Engineering, VP, Manager]
```

16.3. Security IdP

The Security IdP application provides service provider handling that satisfies the [SAML 2.0 Web SSO profile](#) in order to support external IdPs (Identity Providers) or SPs (Service Providers). This capability allows use of DDF as the SSO solution for an entire enterprise.

Table 37. Security IdP Components

Bundle Name	Located in Feature	Description
<code>security-idp-sp</code>	<code>security-idp</code>	The IdP client that interacts with the specified Identity Provider.
<code>security-idp-server</code>	<code>security-idp</code>	An internal Identity Provider solution.

Limitations

NOTE The internal Identity Provider solution should be used in favor of any external solutions until the IdP Service Provider fully satisfies the [SAML 2.0 Web SSO profile](#).

16.4. Security STS

The Security STS application contains the bundles and services necessary to run and talk to a Security Token Service (STS). It builds off of the Apache CXF STS code and adds components specific to DDF functionality.

Table 38. Security STS Components

Bundle Name	Located in Feature	Description/Link to Bundle Page
security-sts-realm	security-sts-realm	Security STS Realm
security-sts-ldaplogin	security-sts-ldaplogin	Security STS LDAP Login
security-sts-ldapclaimshandler	security-sts-ldapclaimshandler	Security STS LDAP Claims Handler
security-sts-server	security-sts-server	Security STS Server
security-sts-samlvalidator	security-sts-server	Contains the default CXF SAML validator and exposes it as a service for the STS.
security-sts-x509validator	security-sts-server	Contains the default CXF x509 validator and exposes it as a service for the STS.

16.4.1. Security STS Client Config

The Security STS Client Config bundle keeps track and exposes configurations and settings for the CXF STS client. This client can be used by other services to create their own STS client. Once a service is registered as a watcher of the configuration, it will be updated whenever the settings change for the sts client.

16.4.1.1. Installing the Security STS Client Config

This bundle is installed by default.

16.4.1.2. Configuring the Security STS Client Config

Configure the Security STS Client Config from the Admin Console:

1. Navigate to the Admin Console.
2. Select **Security** Application.
3. Select **Configuration** tab.
4. Select **Security STS Client**.

See [Security STS Client configurations](#) for all possible configurations.

16.4.2. External/WS-S STS Support

16.4.2.1. Security STS WSS

This configuration works just like the STS Client Config for the internal STS, but produces standard requests instead of the custom DDF ones. It supports two new auth types for the context policy manager, WSSBASIC and WSSPKI. Use these auth types when connecting to a non-DDF STS or if ignoring realms.

16.4.2.2. Security STS Address Provider

This allows one to select which STS address will be used (e.g. in SOAP sources) for clients of this service. Default is off (internal).

16.4.3. Security STS LDAP Login

The Security STS LDAP Login bundle enables functionality within the STS that allows it to use an LDAP to perform authentication when passed a `UsernameToken` in a `RequestSecurityToken` SOAP request.

16.4.3.1. Installing the Security STS LDAP Login

This bundle is not installed by default but can be added by installing the `security-sts-ldaplogin` feature.

16.4.3.2. Configuring the Security STS LDAP Login

Configure the Security STS LDAP Login from the Admin Console:

1. Navigate to the Admin Console.
2. Select **Security** Application.
3. Select **Configuration** tab
4. Select **Security STS LDAP Login**.

Table 39. Security STS LDAP Login Settings

Configuration Name	Default Value	Additional Information
LDAP URL	<code>ldaps://\${org.codice.ddf.system.hostname}:1636</code>	
StartTLS	<code>false</code>	Ignored if the URL uses ldaps.
LDAP Bind User DN	<code>cn=admin</code>	This user should have the ability to verify passwords and read attributes for any user.
LDAP Bind User Password	<code>secret</code>	This password value is encrypted by default using the Security Encryption application.

Configuration Name	Default Value	Additional Information
LDAP Group User Membership Attribute	uid	Attribute used as the membership attribute for the user in the group. Usually this is uid, cn, or something similar.
LDAP User Login Attribute	uid	Attribute used as the login username. Usually this is uid, cn, or something similar.
LDAP Base User DN	ou=users,dc=example,dc=com	
LDAP Base Group DN	ou=groups,dc=example,dc=com	

16.4.4. Security STS LDAP Claims Handler

The Security STS LDAP Claims Handler bundle adds functionality to the STS server that allows it to retrieve claims from an LDAP server. It also adds mappings for the LDAP attributes to the STS SAML claims.

NOTE All claims handlers are queried for user attributes regardless of realm. This means that two different users with the same username in different LDAP servers will end up with both of their claims in each of their individual assertions.

16.4.4.1. Installing Security STS LDAP Claims Handler

This bundle is not installed by default and can be added by installing the `security-sts-ldapclaimshandler` feature.

16.4.4.2. Configuring the Security STS LDAP Claims Handler

Configure the Security STS LDAP Claims Handler from the Admin Console:

1. Navigate to the Admin Console.
2. Select **Security Application**
3. Select **Configuration** tab.
4. Select **Security STS LDAP and Roles Claims Handler**.

Table 40. Security STS LDAP Claims Handler Settings

Configuration Name	Default Value	Additional Information
LDAP URL	<code>ldaps://\${org.codice.ddf.system.hostname}:1636</code>	
StartTLS	<code>false</code>	Ignored if the URL uses ldaps.
LDAP Bind User DN	<code>cn=admin</code>	This user should have the ability to verify passwords and read attributes for any user.
LDAP Bind User Password	<code>secret</code>	This password value is encrypted by default using the Security Encryption application.

Configuration Name	Default Value	Additional Information
LDAP Username Attribute	<code>uid</code>	
LDAP Base User DN	<code>ou=users,dc=example,dc=com</code>	
LDAP Group ObjectClass	<code>groupOfNames</code>	<code>ObjectClass</code> that defines structure for group membership in LDAP. Usually this is <code>groupOfNames</code> or <code>groupOfUniqueNames</code>
LDAP Membership Attribute	<code>member</code>	Attribute used to designate the user's name as a member of the group in LDAP. Usually this is <code>member</code> or <code>uniqueMember</code>
LDAP Base Group DN	<code>ou=groups,dc=example,dc=com</code>	
User Attribute Map File	<code>etc/ws-security/attributeMap.properties</code>	Properties file that contains mappings from Claim=LDAP attribute.

Table 41. Security STS LDAP Claims Handler Imported Services

Registered Interface	Availability	Multiple
<code>ddf.security.encryption.EncryptionService</code>	optional	false

Table 42. Security STS LDAP Claims Handler Exported Services

Registered Interface	Implementation Class	Properties Set
<code>org.apache.cxf.sts.claims.ClaimsHandler</code>	<code>ddf.security.sts.claimsHandler.LdapClaimsHandler</code>	Properties from the settings
<code>org.apache.cxf.sts.claims.claimsHandler</code>	<code>ddf.security.sts.claimsHandler.RoleClaimsHandler</code>	Properties from the settings

16.4.5. Security STS Server

The Security STS Server is a bundle that starts up an implementation of the CXF STS. The STS obtains many of its configurations (Claims Handlers, Token Validators, etc.) from the OSGi service registry as those items are registered as services using the CXF interfaces. The various services that the STS Server imports are listed in the Implementation Details section of this page.

NOTE

The WSDL for the STS is located at the `security-sts-server/src/main/resources/META-INF/sts/wsdl/ws-trust-1.4-service.wsdl` within the source code.

16.4.5.1. Installing the Security STS Server

This bundle is installed by default and is required for DDF to operate.

16.4.5.2. Configuring the Security STS Server

Configure the Security STS Server from the Admin Console:

1. Navigate to the Admin Console.
2. Select **Security Application**
3. Select **Configuration** tab.
4. Select **Security STS Server**.

Table 43. Security STS Server Settings

Configuration Name	Default Value	Additional Information
SAML Assertion Lifetime	1800	
Token Issuer	localhost	The name of the server issuing tokens. Generally this is the cn or hostname of this machine on the network.
Signature Username	localhost	Alias of the private key in the STS Server's keystore used to sign messages.
Encryption Username	localhost	Alias of the private key in the STS Server's keystore used to encrypt messages.

16.4.6. Security STS Service

The Security STS Service performs authentication of a user by delegating the authentication request to an STS. This is different than the services located within the Security PDP application as those ones only perform authorization and not authentication.

16.4.6.1. Installing the Security STS Realm

This bundle is installed by default and should not be uninstalled.

16.4.6.2. Configuring the Security STS Realm

The Security STS Realm has no configurable properties.

Table 44. Security STS Realm Imported Services

Registered Interface	Availability	Multiple
ddf.security.encryption.EncryptionService	optional	false

Table 45. Security STS Realm Exported Services

Registered Interfaces	Implementation Class	Properties Set
org.apache.shiro.realm.Realm	ddf.security.realm.sts.StsRealm	None


Developing

Developers will build or extend the functionality of the applications.

DDF includes several extension points where external developers can add functionality to support

individual use cases.

DDF is written in Java and uses many open source libraries. DDF uses OSGi to provide modularity, lifecycle management, and dynamic services. OSGi services can be installed and uninstalled while DDF is running. DDF development typically means developing new OSGi bundles and deploying them to the running DDF. A complete description of OSGi is outside the scope of this documentation. For more information about OSGi, see the [OSGi Alliance website](#).




Architecture Diagram


IMPORTANT

If developing for a Highly Available Cluster of DDF, see [High Availability Guidance](#).

17. Catalog Framework API



Catalog Architecture



Catalog Framework Architecture

The **CatalogFramework** is the routing mechanism between catalog components that provides integration points for the Catalog Plugins. An **endpoint** invokes the active Catalog Framework, which calls any configured **Pre-query** or **Pre-ingest** plug-ins. The selected **federation strategy** calls the active **Catalog Provider** and any connected or federated sources. Then, any Post-query or Post-ingest plug-ins are invoked. Finally, the appropriate response is returned to the calling endpoint.

The Catalog Framework wires all Catalog components together.

It is responsible for routing Catalog requests and responses to the appropriate target.

[Endpoints](#) send Catalog requests to the Catalog Framework. The Catalog Framework then invokes [Catalog Plugins](#), [Transformers](#), and [Resource Components](#) as needed before sending requests to the intended destination, such as one or more [Sources](#).

The Catalog Framework decouples clients from service implementations and provides integration points for Catalog Plugins and convenience methods for Endpoint developers.

17.1. Catalog API Design


The Catalog is composed of several components and an API that connects them together. The Catalog API is central to DDF's architectural qualities of extensibility and flexibility. The Catalog API consists of Java interfaces that define Catalog functionality and specify interactions between components. These interfaces provide the ability for components to interact without a dependency on a particular underlying implementation, thus allowing the possibility of alternate implementations that can maintain interoperability and share developed components. As such, new capabilities can be developed independently, in a modular fashion, using the Catalog API interfaces and reused by other DDF installations.

17.1.1. Ensuring Compatibility

The Catalog API will evolve, but great care is taken to retain backwards compatibility with developed components. Compatibility is reflected in version numbers.

17.1.2. Catalog Framework Sequence Diagrams

Because the Catalog Framework plays a central role to Catalog functionality, it interacts with many different Catalog components. To illustrate these relationships, high level sequence diagrams with notional class names are provided below. These examples are for illustrative purposes only and do not necessarily represent every step in each procedure.



Ingest Request Data Flow

The Ingest Service Endpoint, the Catalog Framework, and the Catalog Provider are key components of the Reference Implementation. The Endpoint bundle implements a Web service that allows clients to create, update, and delete metacards. The Endpoint calls the [CatalogFramework](#) to execute the operations of its specification. The [CatalogFramework](#) routes the request through optional

[PreIngest](#) and [PostIngest](#) Catalog Plugins, which may modify the ingest request/response before/after the Catalog Provider executes the ingest request and provides the response. Note that a [CatalogProvider](#) must be present for any ingest requests to be successfully processed, otherwise a fault is returned.

This process is similar for updating catalog entries, with update requests calling the [update\(UpdateRequest\)](#) methods on the Endpoint, [CatalogFramework](#), and Catalog Provider. Similarly, for deletion of catalog entries, the delete requests call the [delete\(DeleteRequest\)](#) methods on the Endpoint, [CatalogFramework](#), and [CatalogProvider](#).

17.1.2.1. Error Handling


Any ingest attempts that fail inside the Catalog Framework (whether the failure comes from the Catalog Framework itself, pre-ingest plugin failures, or issues with the Catalog Provider) will be logged to a separate log file for ease of error handling. The file is located at [data/log/ingest_error.log](#) and will log the Metacards that fail, their ID and Title name, and the stack trace associated with their failure. By default, successful ingest attempts are not logged. However, that functionality can be achieved by setting the log level of the [ingestLogger](#) to DEBUG (note that enabling DEBUG can cause a non-trivial performance hit).

To turn off logging failed ingest attempts into a separate file, execute the following via the command line console

TIP

```
log:set
  ERROR ingestLogger
```

17.1.2.2. Query



Query Request Data Flow

The [Query Service Endpoint](#), the Catalog Framework, and the [CatalogProvider](#) are key components for processing a query request as well. The Endpoint bundle contains a Web service that exposes the interface to query for [Metacards](#). The Endpoint calls the [CatalogFramework](#) to execute the operations of its specification. The [CatalogFramework](#) relies on the [CatalogProvider](#) to execute the actual query. Optional PreQuery and PostQuery Catalog Plugins may be invoked by the [CatalogFramework](#) to modify the query request/response prior to the Catalog Provider processing the query request and providing the query response. If a [CatalogProvider](#) is not configured and no other remote Sources are configured, a fault will be returned. It is possible to have only remote Sources configured and no local [CatalogProvider](#) configured and be able to execute queries to

specific remote Sources by specifying the site name(s) in the query request.

17.1.2.3. Product Retrieval

The Query Service Endpoint, the Catalog Framework, and the [CatalogProvider](#) are key components for processing a retrieve product request. The Endpoint bundle contains a Web service that exposes the interface to retrieve products, also referred to as Resources. The Endpoint calls the [CatalogFramework](#) to execute the operations of its specification. The [CatalogFramework](#) relies on the Sources to execute the actual product retrieval. Optional [PreResource](#) and [PostResource](#) Catalog Plugins may be invoked by the [CatalogFramework](#) to modify the product retrieval request/response prior to the Catalog Provider processing the request and providing the response. It is possible to retrieve products from specific remote Sources by specifying the site name(s) in the request.

17.1.2.4. Product Caching

The Catalog Framework optionally provides caching of products, so future requests to retrieve the same product will be serviced much quicker. If caching is enabled, each time a retrieve product request is received, the Catalog Framework will look in its cache (default location `<DDF_HOME>/data/product-cache`) to see if the product has been cached locally. If it has, the product is retrieved from the local site and returned to the client, providing a much quicker turnaround because remote product retrieval and network traffic was avoided. If the requested product is not in the cache, the product is retrieved from the Source (local or remote) and cached locally while returning the product to the client. The caching to a local file of the product and the streaming of the product to the client are done simultaneously so that the client does not have to wait for the caching to complete before receiving the product. If errors are detected during the caching, caching of the product will be abandoned, and the product will be returned to the client.

The Catalog Framework attempts to detect any network problems during the product retrieval, e.g., long pauses where no bytes are read implying a network connection was dropped. (The amount of time defined as a "long pause" is configurable, with the default value being five seconds.) The Catalog Framework will attempt to retrieve the product up to a configurable number of times (default = three), waiting for a configurable amount of time (default = 10 seconds) between each attempt, trying to successfully retrieve the product. If the Catalog Framework is unable to retrieve the product, an error message is returned to the client.

If the admin has enabled the **Always Cache When Canceled** option, caching of the product will occur even if the client cancels the product retrieval so that future requests will be serviced quickly. Otherwise, caching is canceled if the user cancels the product download.

17.1.2.5. Product Download Status

As part of the caching of products, the Catalog Framework also posts events to the OSGi notification framework. Information includes when the product download started, whether the download is retrying or failed (after the number of retrieval attempts configured for product caching has been exhausted), and when the download completes. These events are retrieved by the Search UI and presented to the user who initiated the download.

17.1.3. Catalog API

The Catalog API is an OSGi bundle ([catalog-core-api](#)) that contains the Java interfaces for the Catalog components and implementation classes for the Catalog Framework, Operations, and Data components.

17.1.3.1. Catalog API Search Interfaces

The Catalog API includes two different search interfaces.

Search UI Application Search Interface

The DDF Search UI application provides a graphic interface to return results and locate them on an interactive globe or map.

SSH Search Interface

Additionally, it is possible to use a client script to remotely access DDF via SSH and send console commands to search and ingest data.

17.1.3.2. Catalog Search Result Objects

Data is returned from searches as Catalog Search **Result** objects. This is a subtype of Catalog **Entry** that also contains additional data based on what type of sort policy was applied to the search. Because it is a subtype of Catalog **Entry**, a Catalog Search **Result** has all Catalog **Entry**'s fields such as metadata, effective time, and modified time. It also contains some of the following fields, depending on type of search, that are populated by DDF when the search occurs:

Distance

Populated when a point-radius spatial search occurs. Numerical value that indicates the result's distance from the center point of the search.

Units

Populated when a point-radius spatial search occurs. Indicates the units (kilometer, mile, etc.) for the distance field.

Relevance

Populated when a contextual search occurs. Numerical value that indicates how relevant the text in the result is to the text originally searched for.

17.1.3.3. Search Programmatic Flow

Searching the catalog involves three basic steps:

1. Define the search criteria (contextual, spatial, or temporal).
 - a. Optionally define a sort policy and assign it to the criteria.
 - b. For contextual search, optionally set the **fuzzy** flag to **true** or **false** (the default value for the **Metadata Catalog fuzzy** flag is **true**, while the **portal** default value is **false**).
 - c. For contextual search, optionally set the **caseSensitive** flag to true (the default is that **caseSensitive** flag is NOT set and queries are not case sensitive). Doing so enables case

sensitive matching on the search criteria. For example, if `caseSensitive` is set to true and the phrase is “Baghdad” then only metadata containing “Baghdad” with the same matching case will be returned. Words such as “baghdad”, “BAGHDAD”, and “baghDad” will not be returned because they do not match the exact case of the search term.

2. Issue a search.
3. Examine the results.

17.1.3.4. Sort Policies

Searches can also be sorted according to various built-in policies. A sort policy is applied to the search criteria after its creation but before the search is issued. The policy specifies to the DDF the order the Catalog search results should be in when they are returned to the requesting client. Only one sort policy may be defined per search.

There are three policies available.

Table 46. Sort Policies

Sort Policy	Sorts By	Default Order	Available for
Temporal	The catalog search result’s effective time field	Newest to oldest	All Search Types
Distance	The catalog search result’s distance field	Nearest to farthest	Point-Radius Spatial searches
Relevance	The catalog search result’s relevance field	Most to least relevant	Contextual

If no sort policy is defined for a particular search, the temporal policy will automatically be applied.

17.1.3.5. Asynchronous Search & Retrieval

Asynchronous Search & Retrieval allows a requestor to execute multiple queries at once, begin multiple product downloads while query results are being returned, cancel queries and downloads, and receive status on the state of incoming query results and product downloads.

Table 47. Important Terms for Asynchronous Search

Capability	Description	Endpoint Integration
Asynchronous Search	Search multiple sources simultaneously	Search UI
Product caching	Allows quick retrieval of already downloaded products	Catalog
Canceling Product Downloads	The ability to cancel a download in progress	Catalog
Activities	Activities * <code>download</code> * <code>retry</code> * <code>cancel</code> * <code>pause</code> * <code>remove</code> * <code>resume</code>	Catalog, CometD endpoint

Capability	Description	Endpoint Integration
Notifications	Time-stamped messages of an action	Catalog, Search UI/CometD endpoint
Workspaces	Ability to save and manage queries and save metacards	Platform, Search UI/CometD endpoint
3D Map support	Ability to execute a geospatial search using a 3D map	N/A

17.1.3.6. Product Retrieval

The DDF is used to catalog resources. A Resource is a URI-addressable entity that is represented by a Metocard. Resources may also be known as products or data. Resources may exist either locally or on a remote data store.

Examples of Resources

- NITF image
- MPEG video
- Live video stream
- Audio recording
- Document

Product Retrieval Services

- SOAP Web services
- DDF JSON
- DDF REST

The Query Service Endpoint, the Catalog Framework, and the [CatalogProvider](#) are key components for processing a retrieve product request. The Endpoint bundle contains a Web service that exposes the interface to retrieve products, also referred to as Resources. The Endpoint calls the [CatalogFramework](#) to execute the operations of its specification. The [CatalogFramework](#) relies on the Sources to execute the actual product retrieval. Optional PreResource and PostResource Catalog Plugins may be invoked by the [CatalogFramework](#) to modify the product retrieval request/response prior to the Catalog Provider processing the request and providing the response. It is possible to retrieve products from specific remote Sources by specifying the site name(s) in the request.

Product Caching

NOTE

Existing DDF clients are able to leverage product caching due to the product cache being implemented in the DDF. Enabling the product cache is an administrator function.


Product Caching is enabled by default.

To configure product caching:

1. Navigate to the **Admin Console**.

2. Select Catalog.
3. Select **Configuration**.
4. Select **Resource Download Settings**.

See [Resource Download Settings configurations](#) for all possible configurations.



Product Retrieval Request

17.1.3.7. Notifications and Activities

DDF can send/receive notifications of "Activities" occurring in the system.

17.1.3.7.1. Notifications

Currently, the notifications provide information about product retrieval only. For example, in the DDF Search UI, after a user initiates a resource download, they receive notifications when the download completed, failed, canceled, or is being retried.

17.1.3.7.2. Activities

Activity events include the status and progress of actions that are being performed by the user, such as searches and downloads. Activities can be enabled by selecting "Show tasks" in the Standard Search UI configuration. A list of all activities opens in a drop-down menu, from which activities can be read and deleted. If a download activity is being performed, the Activity drop-down menu provides the link to retrieve the product. If caching is enabled, a progress bar is displayed in the Activity (Product Retrieval) drop-down menu until the action being performed is complete.

17.2. Included Catalog Frameworks, Associated Components, and Configurations

These catalog frameworks are available in a standard DDF installation:

Standard Catalog Framework

Reference implementation of a Catalog Framework that implements all requirements of the Catalog API.

Catalog Framework Camel Component

Supports creating, updating, and deleting metacards using the Catalog Framework from a Camel

route.

17.2.1. Standard Catalog Framework

The Standard Catalog Framework provides the reference implementation of a Catalog Framework that implements all requirements of the Catalog API. `CatalogFrameworkImpl` is the implementation of the DDF Standard Catalog Framework.

The Standard Catalog Framework is the core class of DDF. It provides the methods for create, update, delete, and resource retrieval (CRUD) operations on the `Sources`. By contrast, the Fanout Catalog Framework only allows for query and resource retrieval operations, no catalog modifications, and all queries are enterprise-wide.

Use this framework if:

- access to a catalog provider is required to create, update, and delete catalog entries.
- queries to specific sites are required.
- queries to only the local provider are required.

It is possible to have only remote Sources configured with no local `CatalogProvider` configured and be able to execute queries to specific remote sources by specifying the site name(s) in the query request.

The Standard Catalog Framework also maintains a list of `ResourceReaders` for resource retrieval operations. A resource reader is matched to the scheme (i.e., protocol, such as `file://`) in the URI of the resource specified in the request to be retrieved.

Site information about the catalog provider and/or any federated source(s) can be retrieved using the Standard Catalog Framework. Site information includes the source's name, version, availability, and the list of unique content types currently stored in the source (e.g., NITF). If no local catalog provider is configured, the site information returned includes site info for the catalog framework with no content types included.

17.2.1.1. Installing the Standard Catalog Framework

The Standard Catalog Framework is bundled as the `catalog-core-standardframework` feature and can be installed and uninstalled using the normal processes described in Configuration.

When this feature is installed, the Catalog Fanout Framework App feature `catalog-core-fanoutframework` should be uninstalled, as both catalog frameworks should not be installed simultaneously.

17.2.1.2. Configuring the Standard Catalog Framework

These are the configurable properties on the Standard Catalog Framework.

See [Catalog Standard Framework configurations](#) for all possible configurations.

Table 48. Standard Catalog Framework Exported Services

Registered Interface	Service Property	Value
<code>ddf.catalog.federation.FederationStrategy</code>	<code>shortname</code>	<code>sorted</code>
<code>org.osgi.service.event.EventHandler</code>	<code>event.topics</code>	<code>ddf/catalog/event/CREATED, ddf/catalog/event/UPDATED, ddf/catalog/event/DELETED</code>
<code>ddf.catalog.CatalogFramework</code>		
<code>ddf.catalog.event.EventProcessor</code>		
<code>ddf.catalog.plugin.PostIngestPlugin</code>		

Table 49. Standard Catalog Framework Imported Services

Registered Interface	Availability	Multiple
<code>ddf.catalog.plugin.PostFederatedQueryPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PostIngestPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PostQueryPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PostResourcePlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PreDeliveryPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PreFederatedQueryPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PreIngestPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PreQueryPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PreResourcePlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PreSubscriptionPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.PolicyPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.plugin.AccessPlugin</code>	optional	<code>true</code>
<code>ddf.catalog.resource.ResourceReader</code>	optional	<code>true</code>
<code>ddf.catalog.source.CatalogProvider</code>	optional	<code>true</code>
<code>ddf.catalog.source.ConnectedSource</code>	optional	<code>true</code>
<code>ddf.catalog.source.FederatedSource</code>	optional	<code>true</code>
<code>ddf.cache.CacheManager</code>		<code>false</code>
<code>org.osgi.service.event.EventAdmin</code>		<code>false</code>

17.2.1.3. Known Issues with Standard Catalog Framework

None.

17.2.2. Catalog Framework Camel Component

The Catalog Framework Camel Component supports creating, updating, and deleting metacards using the Catalog Framework from a Camel route.

```
catalog:framework
```

17.2.2.1. Message Headers

17.2.2.1.1. Catalog Framework Producer

Header	Description
operation	the operation to perform using the Catalog Framework (possible values are CREATE UPDATE DELETE)

17.2.2.2. Sending Messages to Catalog Framework Endpoint

17.2.2.2.1. Catalog Framework Producer

In Producer mode, the component provides the ability to provide different inputs and have the Catalog framework perform different operations based upon the header values.

For the CREATE and UPDATE operation, the message body can contain a list of metacards or a single metacard object.

For the DELETE operation, the message body can contain a list of strings or a single string object. The string objects represent the IDs of metacards to be deleted. The exchange's "in" message will be set with the affected metacards. In the case of a CREATE, it will be updated with the created metacards. In the case of the UPDATE, it will be updated with the updated metacards and with the DELETE it will contain the deleted metacards.

Table 50. Catalog Framework Camel Component Operations

Header	Message Body (Input)	Exchange Modification (Output)
operation = CREATE	List<Metocard> or Metocard	exchange.getIn().getBody() updated with List of Metacards created
operation = UPDATE	List<Metocard> or Metocard	exchange.getIn().getBody() updated with List of Metacards updated
operation = DELETE	List<String> or String (representing metocard IDs)	exchange.getIn().getBody() updated with List of Metacards deleted

17.2.2.2.2. Samples

This example demonstrates:


1. Reading in some sample data from the file system.
2. Using a Java bean to convert the data into a metocard.
3. Setting a header value on the Exchange.
4. Sending the Metocard to the Catalog Framework component for ingesting.

```

<route>
  <from uri="file:data/sampleData?noop=true"/>
    <bean ref="sampleDataToMetocardConverter" method="convertToMetocard"/> \
    <setHeader headerName="operation">
      <constant>CREATE</constant>
    </setHeader>
    <to uri="catalog:framework"/>
</route>

```

18. Transformers




Transformers

Transformers transform data to and from various formats. Transformers are categorized by when they are invoked and used. The [existing types](#) are [Input transformers](#), [Metocard transformers](#), and [Query Response transformers](#). Additionally, XSLT transformers are provided to aid in developing custom, lightweight Metocard and Query Response transformers.

Transformers are utility objects used to transform a set of standard DDF components into a desired format, such as into PDF, GeoJSON, XML, or any other format. For instance, a transformer can be used to convert a set of query results into an easy-to-read GeoJSON format ([GeoJSON Transformer](#)) or convert a set of results into a RSS feed that can be easily published to a URL for RSS feed subscription. Transformers can be registered in the OSGi Service Registry so that any other developer can access them based on their standard interface and self-assigned identifier, referred to as its "shortname." Transformers are often used by endpoints for data conversion in a system standard way. Multiple endpoints can use the same transformer, a different transformer, or their own published transformer.

WARNING

The current transformers only work for UTF-8 characters and do not support Non-Western Characters (for example, Hebrew). It is recommended not to use international character sets, as they may not be displayed properly.

*Communication Diagram*

Transformers are used to alter the format of a resource or its metadata to or from the catalog's metocard format.

*Types of Transformers***Input Transformers**

Input Transformers create metacards from input. Once converted to a Metocard, the data can be used in a variety of ways, such as in an [UpdateRequest](#), [CreateResponse](#), or within Catalog Endpoints or Sources. For instance, an input transformer could be used to receive and translate XML into a Metocard so that it can be placed within a [CreateRequest](#) in order to be ingested within the Catalog. Input transformers should be registered within the Service Registry with the interface [ddf.catalog.transform.InputTransformer](#) in order to notify Catalog components of any new transformers.

Metocard Transformers

Metocard Transformers translate a metocard from catalog metadata to a specific data format.

Query Response Transformers

Query Response transformers convert query responses into other data formats.

18.1. Available Input Transformers

The following input transformers are available in a standard installation of DDF:

GeoJSON Input Transformer

Translates GeoJSON into a Catalog metocard.

PDF Input Transformer

Translates a PDF document into a Catalog Metocard.

PPTX Input Transformer

Translates Microsoft PowerPoint (OOXML only) documents into Catalog Metacards.

Registry Transformer

Creates Registry metacards from **ebrim** messages and translates a Registry metocard. (used by the Registry application)

Tika Input Transformer

Translates Microsoft Word, Microsoft Excel, Microsoft PowerPoint, OpenOffice Writer, and PDF documents into Catalog Metacards.

Video Input Transformer

Creates Catalog metacards from certain video file types.

XML Input Transformer

Translates an XML document into a Catalog Metocard.

18.2. Available Metocard Transformers

The following metocard transformers are available in a standard installation of DDF:

GeoJSON Metocard Transformer

Translates a metocard into GeoJSON.

HTML Metocard Transformer

Translates a metocard into an HTML-formatted document.

KML Metocard Transformer

Translates a metocard into a KML-formatted document.

KML Style Mapper

Maps a KML Style URL to a metocard based on that metocard's attributes.

Metadata Metocard Transformer

returns the **Metocard.METADATA** attribute when given a metocard.

Registry Transformer

Creates Registry metacards from **ebrim** messages and translates a Registry metocard. (used by the Registry application)

Resource Metocard Transformer

Retrieves the resource bytes of a metocard by returning the product associated with the metocard.

Thumbnail Metocard Transformer

Retrieves the thumbnail bytes of a Metocard by returning the `Metocard.THUMBNAIL` attribute value.

XML Metocard Transformer

Translates a metocard into an XML-formatted document.

18.3. Available Query Response Transformers

The following query response transformers are available in a standard installation of DDF:

Atom Query Response Transformer

Transforms a query response into an [Atom 1.0](#) feed.

CSW Query Response Transformer

Transforms a query response into a [CSW-formatted](#) document.

GeoJSON Query Response Transformer

Translates a query response into a GeoJSON-formatted document.

KML Query Response Transformer

Translates a query response into a KML-formatted document.

Query Response Transformer Consumer

Translates a query response into a Catalog Metocard.

SearchUI Transformer

Provides results in html format and a convenient, simple querying user interface.

XML Query Response Transformer

Translates a query response into an XML-formatted document.

18.4. Transformers Details

Availability and configuration details of available transformers.

18.4.1. Atom Query Response Transformer

The Atom Query Response Transformer transforms a query response into an [Atom 1.0](#) feed. The Atom transformer maps a `QueryResponse` object as described in the Query Result Mapping.

18.4.1.1. Installing the Atom Query Response Transformer

The Atom Query Response Transformer is installed by default with a standard installation.

18.4.1.2. Configuring the Atom Query Response Transformer

The Atom Query Response Transformer has no configurable properties.

18.4.1.3. Using the Atom Query Response Transformer

Use this transformer when Atom is the preferred medium of communicating information, such as for feed readers or federation. An integrator could use this with an endpoint to transform query responses into an Atom feed.

For example, clients can use the [OpenSearch Endpoint](#). The client can query with the format option set to the shortname, `atom`.

Sample OpenSearch Query with Atom Specified as Return Format

```
\http://{FQDN}:{PORT}/services/catalog/query?q=ddf?format=atom
```

Developers could use this transformer to programmatically transform `QueryResponse` objects on the fly.

Sample Atom Feed from `QueryResponse` object

```
<feed xmlns="http://www.w3.org/2005/Atom" xmlns:os="http://a9.com/-/spec/opensearch/1.1/">
  <title type="text">Query Response</title>
  <updated>2017-01-31T23:22:37.298Z</updated>
  <id>urn:uuid:a27352c9-f935-45f0-9b8c-5803095164bb</id>
  <link href="#" rel="self" />
  <author>
    <name>Organization Name</name>
  </author>
  <generator version="2.1.0.20130129-1341">ddf123</generator>
  <os:totalResults>1</os:totalResults>
  <os:itemsPerPage>10</os:itemsPerPage>
  <os:startIndex>1</os:startIndex>
  <entry xmlns:relevance="http://a9.com/-/opensearch/extensions/relevance/1.0/" xmlns:fs="http://a9.com/-/opensearch/extensions/federation/1.0/"
    xmlns:georss="http://www.georss.org/georss">
    <fs:resultSource fs:sourceId="ddf123" />
    <relevance:score>0.19</relevance:score>
    <id>urn:catalog:id:ee7a161e01754b9db1872bfe39d1ea09</id>
    <title type="text">F-15 lands in Libya; Crew Picked Up</title>
    <updated>2013-01-31T23:22:31.648Z</updated>
    <published>2013-01-31T23:22:31.648Z</published>
    <link href="http://123.45.67.123:8181/services/catalog/ddf123/ee7a161e01754b9db1872bfe39d1ea09"
      rel="alternate" title="View Complete Metocard" />
    <category term="Resource" />
    <georss:where xmlns:gml="http://www.opengis.net/gml">
      <gml:Point>
        <gml:pos>32.8751900768792 13.1874561309814</gml:pos>
      </gml:Point>
    </georss:where>
    <content type="application/xml">
      <ns3:metocard xmlns:ns3="urn:catalog:metocard" xmlns:ns2=
```

```

"http://www.w3.org/1999/xlink" xmlns:ns1="http://www.opengis.net/gml"
    xmlns:ns4="http://www.w3.org/2001/SMIL20/" xmlns:ns5=
"http://www.w3.org/2001/SMIL20/Language" ns1:id="4535c53fc8bc4404a1d32a5ce7a29585">
    <ns3:type>ddf.metacard</ns3:type>
    <ns3:source>ddf.distribution</ns3:source>
    <ns3:geometry name="location">
        <ns3:value>
            <ns1:Point>
                <ns1:pos>32.8751900768792 13.1874561309814</ns1:pos>
            </ns1:Point>
        </ns3:value>
    </ns3:geometry>
    <ns3:dateTime name="created">
        <ns3:value>2013-01-31T16:22:31.648-07:00</ns3:value>
    </ns3:dateTime>
    <ns3:dateTime name="modified">
        <ns3:value>2013-01-31T16:22:31.648-07:00</ns3:value>
    </ns3:dateTime>
    <ns3:stringxml name="metadata">
        <ns3:value>
            <ns6:xml xmlns:ns6="urn:sample:namespace" xmlns=
"urn:sample:namespace">Example description.</ns6:xml>
        </ns3:value>
    </ns3:stringxml>
    <ns3:string name="metadata-content-type-version">
        <ns3:value>myVersion</ns3:value>
    </ns3:string>
    <ns3:string name="metadata-content-type">
        <ns3:value>myType</ns3:value>
    </ns3:string>
    <ns3:string name="title">
        <ns3:value>Example title</ns3:value>
    </ns3:string>
    </ns3:metacard>
</content>
</entry>
</feed>

```

Table 51. Atom Query Response Transformer Result Mapping

XPath to Atom XML	Value
/feed/title	"Query Response"
/feed/updated	ISO 8601 dateTime of when the feed was generated
/feed/id	Generated (UUID URN
/feed/author/name	Platform Global Configuration organization

XPath to Atom XML	Value
/feed/generator	Platform Global Configuration site name
/feed/generator/@version	Platform Global Configuration version
/feed/os:totalResults	SourceResponse Number of Hits
/feed/os:itemsPerPage	Request's Page Size
/feed/os:startIndex	Request's Start Index
/feed/entry/fs:resultSource/@fs:sourceId	Source Id from which the Result came. <code>Metacard.getSourceId()</code>
/feed/entry/relevance:score	Result's relevance score if applicable. <code>Result.getRelevanceScore()</code>
/feed/entry/id	<code>urn:catalog:id:<Metacard.ID></code>
/feed/entry/title	<code>Metacard.TITLE</code>
/feed/entry/updated	ISO 8601 dateTIme of <code>Metacard.MODIFIED</code>
/feed/entry/published	ISO 8601 dateTIme of <code>Metacard.CREATED</code>
/feed/entry/link[@rel='related']	URL to retrieve underlying resource (if applicable and link is available)
/feed/entry/link[@rel='alternate']	Link to alternate view of the Metacard (if a link is available)
/feed/entry/category	<code>Metacard.CONTENT_TYPE</code>
/feed/entry//georss:where	GeoRSS GML of every Metacard attribute with format <code>AttributeFormat.GEOMETRY</code>
/feed/entry/content	Metacard XML generated by <code>DDF.catalog.transform.MetacardTransformer</code> with <code>shortname=xml</code> . If no transformer found, <code>/feed/entry/content/@type</code> will be text and <code>Metacard.ID</code> is displayed <code><content type="text">4e1f38d1913b4e93ac622e6c1b258f89</content></code>

18.4.2. CSW Query Response Transformer

The CSW Query Response Transformer transforms a query response into a [CSW-formatted](#) document.

18.4.2.1. Installing the CSW Query Response Transformer

The CSW Query Response Transformer is installed by default with a standard installation in the Spatial application.

18.4.2.2. Configuring the CSW Query Response Transformer

The CSW Query Response Transformer has no configurable properties.

18.4.3. GeoJSON Input Transformer

The GeoJSON input transformer is responsible for translating GeoJSON into a Catalog metocard.

18.4.3.1. Installing the GeoJSON Input Transformer

The GeoJSON Input Transformer is installed by default with a standard installation.

18.4.3.2. Configuring the GeoJSON Input Transformer

The GeoJSON Input Transformer has no configurable properties.

18.4.3.3. Using the GeoJSON Input Transformer

Using the REST Endpoint, for example, HTTP POST a GeoJSON metocard to the Catalog. Once the REST Endpoint receives the GeoJSON Metocard, it is converted to a Catalog metocard.

Example HTTP POST of a Local `metocard.json` File Using the Curl Command

```
curl -X POST -i -H "Content-Type: application/json" -d "@metocard.json"  
\https://{FQDN}:{PORT}/services/catalog
```

18.4.3.4. Conversion to a Metocard

A [GeoJSON object](#) consists of a single JSON object. This can be a geometry, a feature, or a [FeatureCollection](#). The GeoJSON input transformer only converts "feature" objects into metacards because feature objects include geometry information and a list of properties. A geometry object alone does not contain enough information to create a metocard. Additionally, the input transformer currently does not handle [FeatureCollections](#).

Cannot create Metacard from this limited GeoJSON

IMPORTANT

```
{ "type": "LineString",
  "coordinates": [ [100.0, 0.0], [101.0, 1.0] ]
}
```

The following sample *will* create a valid metacard:

Sample Parseable GeoJson (Point)

```
{
  "properties": {
    "title": "myTitle",
    "thumbnail": "CA==",
    "resource-uri": "http://example.com",
    "created": "2012-09-01T00:09:19.368+0000",
    "metadata-content-type-version": "myVersion",
    "metadata-content-type": "myType",
    "metadata": "<xml></xml>",
    "modified": "2012-09-01T00:09:19.368+0000"
  },
  "type": "Feature",
  "geometry": {
    "type": "Point",
    "coordinates": [
      30.0,
      10.0
    ]
  }
}
```

In the current implementation, **Metacard.LOCATION** is not taken from the properties list as WKT, but instead interpreted from the **geometry** JSON object. The geometry object is formatted according to the **GeoJSON** standard. Dates are in the ISO 8601 standard. White space is ignored, as in most cases with JSON. Binary data is accepted as Base64. XML must be properly escaped, such as what is proper for normal JSON.

Currently, only **Required Attributes** are recognized in the properties.

18.4.3.4.1. Metacard Extensibility

GeoJSON supports custom, extensible properties on the incoming GeoJSON using DDF's extensible metacard support. To have those customized attributes understood by the system, a corresponding **MetacardType** must be registered with the **MetacardTypeRegistry**. That **MetacardType** must be specified by name in the metacard-type property of the incoming GeoJSON. If a **MetacardType** is specified on the GeoJSON input, the customized properties can be processed, cataloged, and indexed.

Sample GeoJSON input

```
{  
  "properties": {  
    "title": "myTitle",  
    "thumbnail": "CA==",  
    "resource-uri": "http://example.com",  
    "created": "2012-09-01T00:09:19.368+0000",  
    "metadata-content-type-version": "myVersion",  
    "metadata-content-type": "myType",  
    "metadata": "<xml></xml>",  
    "modified": "2012-09-01T00:09:19.368+0000",  
    "min-frequency": "10000000",  
    "max-frequency": "20000000",  
    "metacard-type": "ddf.metacard.custom.type"  
  },  
  "type": "Feature",  
  "geometry": {  
    "type": "Point",  
    "coordinates": [  
      30.0,  
      10.0  
    ]  
  }  
}
```

When the GeoJSON Input Transformer gets GeoJSON with the `MetacardType` specified, it will perform a lookup in the `MetacardTypeRegistry` to obtain the specified `MetacardType` in order to understand how to parse the GeoJSON. If no `MetacardType` is specified, the GeoJSON Input Transformer will assume the default `MetacardType`. If an unregistered `MetacardType` is specified, an exception will be returned to the client indicating that the `MetacardType` was not found.

18.4.3.5. Usage Limitations of the GeoJSON Input Transformer

The GeoJSON Input Transformer does not handle multiple geometries.

18.4.4. GeoJSON Metacard Transformer

GeoJSON Metacard Transformer translates a metacard into GeoJSON.

18.4.4.1. Installing the GeoJSON Metacard Transformer

The GeoJSON Metacard Transformer is not installed by default with a standard installation.

To install:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.

3. Select the **Features** tab.
4. Install the `catalog-transformer-json` feature.

18.4.4.2. Configuring the GeoJSON Metacard Transformer

The GeoJSON Metacard Transformer has no configurable properties.

18.4.4.3. Using the GeoJSON Metacard Transformer

The GeoJSON Metacard Transformer can be used programmatically by requesting a `MetacardTransformer` with the id `geojson`. It can also be used within the REST Endpoint by providing the transform option as `geojson`.

Example REST GET Method with the GeoJSON Metacard Transformer

```
\https://{{FQDN}}:{{PORT}}/services/catalog/0123456789abcdef0123456789abcdef?transform=geojson
```

```
{  
  "properties":{  
    "title": "myTitle",  
    "thumbnail": "CA==",  
    "resource-uri": "http://example.com",  
    "created": "2012-08-31T23:55:19.518+0000",  
    "metadata-content-type-version": "myVersion",  
    "metadata-content-type": "myType",  
    "metadata": "<xml>text</xml>",  
    "modified": "2012-08-31T23:55:19.518+0000",  
    "metacard-type": "ddf.metacard"  
  },  
  "type": "Feature",  
  "geometry":{  
    "type": "LineString",  
    "coordinates": [  

```

18.4.5. GeoJSON Query Response Transformer

The GeoJSON Query Response Transformer translates a query response into a GeoJSON-formatted document.

18.4.5.1. Installing the GeoJSON Query Response Transformer

The GeoJSON Query Response Transformer is installed by default with a standard installation in the Catalog application.

18.4.5.2. Configuring the GeoJSON Query Response Transformer

The GeoJSON Query Response Transformer has no configurable properties.

18.4.6. HTML Metacard Transformer

The HTML metacard transformer is responsible for translating a metacard into an HTML-formatted document.

18.4.6.1. Installing the HTML Metacard Transformer

The HTML Metacard Transformer is installed by default with a standard installation in the Search UI application.

18.4.6.2. Configuring the HTML Metacard Transformer

The HTML Metacard Transformer has no configurable properties.

18.4.6.3. Using the HTML Metacard Transformer

Using the REST Endpoint for example, request a metacard with the transform option set to the HTML shortname.

```
\https://{{FQDN}}:{{PORT}}/services/catalog/0123456789abcdef0123456789abcdef?transform=htm  
l
```

HTML Metacard Transformer Example Output

```
html metacard.png
```

18.4.7. KML Metacard Transformer

The KML Metacard Transformer is responsible for translating a metacard into a KML-formatted document. The KML will contain an HTML description that will display in the pop-up bubble in Google Earth. The HTML contains links to the full metadata view as well as the product.

18.4.7.1. Installing the KML Metacard Transformer

The KML Metacard Transformer is installed by default with a standard installation in the Spatial Application.

18.4.7.2. Configuring the KML Metacard Transformer

The KML Metacard Transformer has no configurable properties.

18.4.7.3. Using the KML Metacard Transformer

Using the REST Endpoint for example, request a metacard with the transform option set to the KML shortname.

KML Metacard Transformer Example Output

```

<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<kml xmlns:ns2="http://www.google.com/kml/ext/2.2" xmlns=
"http://www.opengis.net/kml/2.2" xmlns:ns4="urn:oasis:names:tc:ciq:xsdschema:xAL:2.0"
xmlns:ns3="http://www.w3.org/2005/Atom">
  <Placemark id="Placemark-0103c77e66d9428d8f48fab939da528e">
    <name>MultiPolygon</name>
    <description>&lt;!DOCTYPE html&gt;
      &lt;html&gt;
        &lt;head&gt;
          &lt;meta content="text/html; charset=windows-1252" http-equiv="content-
type"&gt;
            &lt;style media="screen" type="text/css"&gt;
              .label {
                font-weight: bold
              }
              .linkTable {
                width: 100%
              }
              .thumbnailDiv {
                text-align: center
              }
              img {
                max-width: 100px;
                max-height: 100px;
                border-style:none
              }
            &lt;/style&gt;
          &lt;/head&gt;
          &lt;body&gt;
            &lt;div class="thumbnailDiv"&gt;&lt;a
href="http://[FQDN]:[PORT]/services/catalog/sources/ddf.distribution/0103c77e66d9428d
8f48fab939da528e?transform=resource"&gt;&lt;img alt="Thumbnail"
src="data:image/jpeg; charset=utf-8; base64, CA=="&gt;&lt;/a&gt;&lt;/div&gt;
            &lt;table&gt;
              &lt;tr&gt;
                &lt;td class="label"&gt;Source:&lt;/td&gt;
                &lt;td&gt;ddf.distribution&lt;/td&gt;
              &lt;/tr&gt;
              &lt;tr&gt;
                &lt;td class="label"&gt;Created:&lt;/td&gt;
                &lt;td&gt;Wed Oct 30 09:46:29 MDT 2013&lt;/td&gt;
              &lt;/tr&gt;
              &lt;tr&gt;
                &lt;td class="label"&gt;Effective:&lt;/td&gt;
                &lt;td&gt;2014-01-07T14:58:16-0700&lt;/td&gt;
              &lt;/tr&gt;
            &lt;/table&gt;
            &lt;table class="linkTable"&gt;
              &lt;tr&gt;
                &lt;td&gt;&lt;a
href="http://[FQDN]:[PORT]/services/catalog/sources/ddf.distribution/0103c77e66d9428d
8f48fab939da528e?transform=html"&gt;View Details...&lt;/a&gt;&lt;/td&gt;
              &lt;/tr&gt;
            &lt;/table&gt;
          &lt;/body&gt;
        &lt;/html&gt;
      &lt;/description>
    </Placemark>

```

```

<td><a href="http://{FQDN}:{PORT}/services/catalog/sources/ddf.distribution/0103c77e66d9428d8f48fab939da528e?transform=resource">Download...</a></td>
</tr>
</table>
</body>
</html>
</description>
<TimeSpan>
  <begin>2014-01-07T21:58:16</begin>
</TimeSpan>
<Style id="bluenormal">
  <LabelStyle>
    <scale>0.0</scale>
  </LabelStyle>
  <LineStyle>
    <color>33ff0000</color>
    <width>3.0</width>
  </LineStyle>
  <PolyStyle>
    <color>33ff0000</color>
    <fill xsi:type="xs:boolean" xmlns:xs="http://www.w3.org/2001/XMLSchema"
      xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">true</fill>
  </PolyStyle>
  <BalloonStyle>
<text>&lt;h3&gt;&lt;b&gt;${name}&lt;/b&gt;&lt;/h3&gt;&lt;table&gt;&lt;tr&gt;&lt;td
width="400"&gt;${description}&lt;/td&gt;&lt;/tr&gt;&lt;/table&gt;</text>
  </BalloonStyle>
</Style>
<Style id="bluehighlight">
  <LabelStyle>
    <scale>1.0</scale>
  </LabelStyle>
  <LineStyle>
    <color>99ff0000</color>
    <width>6.0</width>
  </LineStyle>
  <PolyStyle>
    <color>99ff0000</color>
    <fill xsi:type="xs:boolean" xmlns:xs="http://www.w3.org/2001/XMLSchema"
      xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">true</fill>
  </PolyStyle>
  <BalloonStyle>
<text>&lt;h3&gt;&lt;b&gt;${name}&lt;/b&gt;&lt;/h3&gt;&lt;table&gt;&lt;tr
&gt;&lt;td width="400"&gt;${description}&lt;/td&gt;&lt;/tr&gt;&lt;/table&gt;</text>
  </BalloonStyle>
</Style>
<StyleMap id="default">
  <Pair>
    <key>normal</key>
    <styleUrl>#bluenormal</styleUrl>

```

```

</Pair>
<Pair>
  <key>highlight</key>
  <styleUrl>#bluehighlight</styleUrl>
</Pair>
</StyleMap>
<MultiGeometry>
  <Point>
    <coordinates>102.0,2.0</coordinates>
  </Point>
  <MultiGeometry>
    <Polygon>
      <outerBoundaryIs>
        <LinearRing>
          <coordinates>102.0,2.0 103.0,2.0 103.0,3.0 102.0,3.0 102.0,2.0<
/coordinates>
        </LinearRing>
      </outerBoundaryIs>
    </Polygon>
    <Polygon>
      <coordinates>100.8,0.2
        <outerBoundaryIs>
          <LinearRing>
            <coordinates>100.0,0.0 101.0,0.0 101.0,1.0 100.0,1.0 100.0,0.0 100.2,0.2
      </coordinates>
      </LinearRing>
    </outerBoundaryIs>
  </Polygon>
</MultiGeometry>
</Placemark>
</kml>

```

18.4.8. KML Query Response Transformer

The KML Query Response Transformer translates a query response into a KML-formatted document. The KML will contain an HTML description for each metocard that will display in the pop-up bubble in Google Earth. The HTML contains links to the full metadata view as well as the product.

18.4.8.1. Installing the KML Query Response Transformer

The `spatial-kml-transformer` feature is installed by default in the Spatial Application.

18.4.8.2. Configuring the KML Query Response Transformer

The KML Query Response Transformer has no configurable properties.

18.4.8.3. Using the KML Query Response Transformer

Using the OpenSearch Endpoint, for example, query with the format option set to the KML shortname: `kml`.

KML Query Response Transformer URL

```
\http://{FQDN}:{PORT}/services/catalog/query?q=schematypesearch&format=kml
```

KML Query Response Transformer Example Output

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<kml xmlns:ns2="http://www.google.com/kml/ext/2.2" xmlns=
"http://www.opengis.net/kml/2.2" xmlns:ns4="urn:oasis:names:tc:ciq:xsdschema:xAL:2.0"
xmlns:ns3="http://www.w3.org/2005/Atom">
  <Document id="f0884d8c-cf9b-44a1-bb5a-d3c6fb9a96b6">
    <name>Results (1)</name>
    <open xsi:type="xs:boolean" xmlns:xs="http://www.w3.org/2001/XMLSchema" xmlns:xsi
="http://www.w3.org/2001/XMLSchema-instance">false</open>
    <Style id="bluenormal">
      <LabelStyle>
        <scale>0.0</scale>
      </LabelStyle>
      <LineStyle>
        <color>33ff0000</color>
        <width>3.0</width>
      </LineStyle>
      <PolyStyle>
        <color>33ff0000</color>
        <fill xsi:type="xs:boolean" xmlns:xs="http://www.w3.org/2001/XMLSchema"
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">true</fill>
      </PolyStyle>
      <BalloonStyle>
        <text>&lt;h3&gt;&lt;b&gt;[$name]&lt;/b&gt;&lt;/h3&gt;&lt;table&gt;&lt;tr
&gt;&lt;td width="400"&gt;[$description]&lt;/td&gt;&lt;/tr&gt;&lt;/table&gt;</text>
      </BalloonStyle>
    </Style>
    <Style id="bluehighlight">
      <LabelStyle>
        <scale>1.0</scale>
      </LabelStyle>
      <LineStyle>
        <color>99ff0000</color>
        <width>6.0</width>
      </LineStyle>
      <PolyStyle>
        <color>99ff0000</color>
        <fill xsi:type="xs:boolean" xmlns:xs="http://www.w3.org/2001/XMLSchema"
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">true</fill>
      </PolyStyle>
      <BalloonStyle>
```

```

<text>&lt;h3&gt;&lt;b&gt;${name}&lt;/b&gt;&lt;/h3&gt;&lt;table&gt;&lt;tr
&gt;&lt;td width="400"&gt;${description}&lt;/td&gt;&lt;/tr&gt;&lt;/table&gt;</text>
</BalloonStyle>
</Style>
<StyleMap id="default">
  <Pair>
    <key>normal</key>
    <styleUrl>#bluenormal</styleUrl>
  </Pair>
  <Pair>
    <key>highlight</key>
    <styleUrl>#bluehighlight</styleUrl>
  </Pair>
</StyleMap>
<Placemark id="Placemark-0103c77e66d9428d8f48fab939da528e">
  <name>MultiPolygon</name>
  <description>&lt;!DOCTYPE html&gt;
&lt;html&gt;
  &lt;head&gt;
    &lt;meta content="text/html; charset=windows-1252" http-equiv="content-type"&gt;
    &lt;style media="screen" type="text/css"&gt;
      .label {
        font-weight: bold
      }
      .linkTable {
        width: 100%
      }
      .thumbnailDiv {
        text-align: center
      }
      img {
        max-width: 100px;
        max-height: 100px;
        border-style:none
      }
    &lt;/style&gt;
  &lt;/head&gt;
  &lt;body&gt;
    &lt;div class="thumbnailDiv"&gt;&lt;a
      href="http://{FQDN}:{PORT}/services/catalog/sources/ddf.distribution/0103c77e66d9428d
      8f
      48fab939da528e?transform=resource"&gt;&lt;img alt="Thumbnail"
      src="data:image/jpeg; charset=utf-8; base64, CA=="&gt;&lt;/a&gt;&lt;/div&gt;
    &lt;table&gt;
      &lt;tr&gt;
        &lt;td class="label"&gt;Source:&lt;/td&gt;
        &lt;td&gt;ddf.distribution&lt;/td&gt;
      &lt;/tr&gt;
      &lt;tr&gt;
        &lt;td class="label"&gt;Created:&lt;/td&gt;
        &lt;td&gt;Wed Oct 30 09:46:29 MDT 2013&lt;/td&gt;
      &lt;/tr&gt;
      &lt;tr&gt;

```

```

<td class="label">Effective:</td>
<td>2014-01-07T14:48:47-0700</td>
</tr>
</table>
<table class="linkTable">
<tr>
<td>
View Details...
Download...
</td>
</tr>
</table>
</body>
</html>
</description>
<TimeSpan>
<begin>2014-01-07T21:48:47</begin>
</TimeSpan>
<styleUrl>#default</styleUrl>
<MultiGeometry>
<Point>
<coordinates>102.0,2.0</coordinates>
</Point>
<MultiGeometry>
<Polygon>
<outerBoundaryIs>
<LinearRing>
<coordinates>102.0,2.0 103.0,2.0 103.0,3.0 102.0,3.0 102.0,2.0</coordinates>
</LinearRing>
100.8,0.2
</outerBoundaryIs>
</Polygon>
<Polygon>
<outerBoundaryIs>
<LinearRing>
<coordinates>100.0,0.0 101.0,0.0 101.0,1.0 100.0,1.0 100.0,0.0 100.2,0.2 100.8,0.8 100.2,0.8 100.2,0.2</coordinates>
</LinearRing>
</outerBoundaryIs>
</Polygon>
</MultiGeometry>
</MultiGeometry>
</Placemark>
</Document>
</kml>

```

18.4.9. KML Style Mapper

The KML Style Mapper provides the ability for the [KmlTransformer](#) to map a KML Style URL to a metocard based on that metocard's attributes. For example, if a user wanted all JPEGs to be blue, the KML Style Mapper provides the ability to do so. This would also allow an administrator to configure metacards from each source to be different colors.

The configured style URLs are expected to be HTTP URLs. For more information on style URL's, refer to the [KML Reference](#).

The KML Style Mapper supports all basic and extended metocard attributes. When a style mapping is configured, the resulting transformed KML contain a [`<styleUrl>`](#) tag pointing to that style, rather than the default KML style supplied by the [KmlTransformer](#).

18.4.9.1. Installing the KML Style Mapper

The KML Style Mapper is installed by default with a standard installation in the [Spatial Application](#) in the [spatial-kml-transformer](#) feature.

18.4.9.2. Configuring the KML Style Mapper

The properties below describe how to configure a style mapping. The configuration name is [Spatial KML Style Map Entry](#).

See [KML Style Mapper configurations](#) for all possible configurations.

KML Style Mapper Example Values

```
xmlns="http://www.opengis.net/kml/2.2"
xmlns:ns4="urn:oasis:names:tc:ciq:xsdschema:xAL:2.0"
xmlns:ns3="http://www.w3.org/2005/Atom">
  <Placemark id="Placemark-0103c77e66d9428d8f48fab939da528e">
    <name>MultiPolygon</name>
    <description>&lt;!DOCTYPE html&gt;
&lt;html&gt;
  &lt;head&gt;
    &lt;meta content="text/html; charset=windows-1252" http-equiv="content-type"&gt;
    &lt;style media="screen" type="text/css"&gt;
      .label {
        font-weight: bold
      }
      .linkTable {
        width: 100%
      }
      .thumbnailDiv {
        text-align: center
      }
      img {
        max-width: 100px;
        max-height: 100px;
        border-style:none
      }
    &lt;/style&gt;
  &lt;/head&gt;
  &lt;body&gt;
    <h1>MultiPolygon</h1>
    <div>
      <table border="1">
        <caption>Link Table</caption>
        <thead>
          <tr>
            <th>Link ID</th>
            <th>Link Description</th>
          </tr>
        </thead>
        <tbody>
          <tr>
            <td>1</td>
            <td>Link 1 Description</td>
          </tr>
          <tr>
            <td>2</td>
            <td>Link 2 Description</td>
          </tr>
        </tbody>
      </table>
    </div>
    <div>
      
    </div>
  &lt;/body&gt;
</Placemark>
```

```

<?xml version="1.0" encoding="UTF-8"?>
<?kml version="2.2"?>
<Document>
  <?xhtml:head?gt;
  <?xhtml:body?gt;
    <?div class="thumbnailDiv"?gt;<a href="http://{FQDN}:{PORT}/services/catalog/sources/ddf.distribution/0103c77e66d9428d8f48fab939da528e?transform=resource"?gt;</div>
    <?table?gt;
      <?tr?gt;
        <?td class="label"?gt;Source:</td>
        <?td>ddf.distribution</td>
      </tr>
      <?tr?gt;
        <?td class="label"?gt;Created:</td>
        <?td>Wed Oct 30 09:46:29 MDT 2013</td>
      </tr>
      <?tr?gt;
        <?td class="label"?gt;Effective:</td>
        <?td>2014-01-07T14:58:16-0700</td>
      </tr>
    </table>
    <?table class="linkTable"?gt;
      <?tr?gt;
        <?td><a href="http://{FQDN}:{PORT}/services/catalog/sources/ddf.distribution/0103c77e66d9428d8f48fab939da528e?transform=html"?gt;View Details...</a></td>
        <?td><a href="http://{FQDN}:{PORT}/services/catalog/sources/ddf.distribution/0103c77e66d9428d8f48fab939da528e?transform=resource"?gt;Download...</a></td>
      </tr>
    </table>
  </body>
</html>
</body>
</Document>
</kml>

```

```
100.8,0.2
<outerBoundaryIs>
  <LinearRing>
    <coordinates>100.0,0.0 101.0,0.0 101.0,1.0 100.0,1.0 100.0,0.0 100.2,0.2
      100.8,0.8 100.2,0.8 100.2,0.2</coordinates>
  </LinearRing>
  </outerBoundaryIs>
</Polygon>
</MultiGeometry>
</MultiGeometry>
</Placemark>
</kml>
```

18.4.10. Metadata Metacard Transformer

The Metadata Metacard Transformer returns the `Metacard.METADATA` attribute when given a metacard. The MIME Type returned is `text/xml`.

18.4.10.1. Installing the Metadata Metacard Transformer

The Metadata Metacard Transformer is installed by default in a standard installation with the Catalog application.

18.4.10.2. Configuring the Metadata Metacard Transformer

The Metadata Metacard Transformer has no configurable properties.

18.4.10.3. Using the Metadata Metacard Transformer

The Metadata Metacard Transformer can be used programmatically by requesting a metacard transformer with the id `metadata`. It can also be used within the REST Endpoint by providing the transform option as `metadata`.

Example REST GET method with the Metadata Metacard Transformer

```
\http://{FQDN}:{PORT}/services/catalog/0123456789abcdef0123456789abcdef?transform=meta
data
```

18.4.11. PDF Input Transformer

The PDF Input Transformer is responsible for translating a PDF document into a Catalog Metacard.

18.4.11.1. Installing the PDF Input Transformer

The PDF Transformer is installed by default with a standard installation in the Catalog application.

18.4.11.2. Configuring the PDF Input Transformer

To configure the PDF Input Transformer:

1. Navigate to the **Catalog** application.
2. Select the **Configuration** tab.
3. Select the **PDF Input Transformer**.

These configurations are available for the PDF Input Transformer:

See [PDF Input Transformer configurations](#) for all possible configurations.

18.4.12. PPTX Input Transformer

The PPTX Input Transformer translates Microsoft PowerPoint (OOXML only) documents into Catalog Metacards, using [Apache Tika](#) for basic metadata and [Apache POI](#) for thumbnail creation. The PPTX Input Transformer ingests PPTX documents into the DDF Content Repository and the Metadata Catalog, and adds a thumbnail of the first page in the PPTX document.

The PPTX Input Transformer will take precedence over the Tika Input Transformer for PPTX documents.

18.4.12.1. Installing the PPTX Input Transformer

This transformer is installed by default with a standard installation in the Catalog application.

18.4.12.2. Configuring the PPTX Input Transformer

The PPTX Input Transformer has no configurable properties. ""

18.4.13. Query Response Transformer Consumer

The Query Response Transformer Consumer is responsible for translating a query response into a Catalog Metacard.

18.4.13.1. Installing the Query Response Transformer Consumer

The Query Response Transformer Consumer is installed by default with a standard installation in the Catalog application.

18.4.13.2. Configuring the Query Response Transformer Consumer

The Query Response Transformer Consumer has no configurable properties.

18.4.14. Registry Transformer

The Registry Transformer creates Registry metacards from **ebrim** messages. It also returns the **ebrim** message from the metacard metadata.

18.4.14.1. Installing the Registry Transformer

The Registry Transformer is not installed by default with a standard installation, but with the Registry application.

To install:

1. [Install Registry](#) application.

18.4.14.2. Configuring the Registry Transformer

The Registry Transformer has no configurable properties.

18.4.15. Resource Metocard Transformer

The Resource Metocard Transformer retrieves a resource associated with a metocard.

18.4.15.1. Installing the Resource Metocard Transformer

The Resource Metocard Transformer is installed by default in a standard installation with the Catalog application as the feature [catalog-transformer-resource](#).

18.4.15.2. Configuring the Resource Metocard Transformer

The Resource Metocard Transformer has no configurable properties.

18.4.15.3. Using the Resource Metocard Transformer

Endpoints or other components can retrieve an instance of the Resource Metocard Transformer using its `id` resource.

Sample Resource Metocard Transformer Blueprint Reference Snippet

```
<reference id="metocardTransformer" interface="ddf.catalog.transform.MetocardTransformer" filter="(id=resource)"/>
```

18.4.16. SearchUI Transformer

The Simple SearchUI is a [QueryResponseTransformer](#) that provides results in html format and a convenient, simple querying user interface. It is primarily used as a test tool and verification of configuration. The left pane of the Simple SearchUI contains basic fields to query the Catalog and other Sources. The right pane consists of the results returned from the query.

18.4.16.1. Installing Simple SearchUI

This transformer is installed by default with a standard installation in the Catalog application as the feature, [catalog-transformer-ui](#).

18.4.16.2. Configuring Simple SearchUI

From the Admin Console, the Simple SearchUI can be configured under the **Catalog HTML Query Response Transformer**.

See [Simple Search UI configurations](#) for all possible configurations.

18.4.16.3. Using Simple SearchUI

In order to access the Simple SearchUI, a user must use the transformer with an endpoint that queries the Catalog, such as the OpenSearch Endpoint. If a distribution is running locally, <https://{}:{}/search/simple> should bring up the Simple Search UI. After the page has loaded, enter the desired search criteria in the appropriate fields. Then click the "**Search**" button in order to execute the search on the Catalog.

The "**Clear**" button will reset the query criteria specified.

Table 52. Simple SearchUI Query Response Result Mapping

Simple SearchUI Column Title	Catalog Result	Notes
Title	<code>Metocard.TITLE</code>	The title may be hyperlinked to view the full Metocard.
Source	<code>Metocard.getSourceId()</code>	Source where the Metocard was discovered.
Location	<code>Metocard.LOCATION</code>	Geographical location of the Metocard.
Time	<code>Metocard.CREATED or Metocard.EFFECTIVE</code>	Time received/created.
Thumbnail	<code>Metocard.THUMBNAIL</code>	Column not shown if no results have a thumbnail.
Resource	<code>Metocard.RESOURCE_URI</code>	Column not shown if no results have a resource.

18.4.16.3.1. Search Criteria

The Simple SearchUI allows for querying a Catalog in the following methods:

Keyword Search

searching with keywords using the grammar of the underlying endpoint/Catalog.

Temporal Search

searching based on relative or absolute time.

Spatial search

searching spatially with a Point-Radius or Bounding Box.

Content Type Search

searching for specific `Metocard.CONTENT_TYPE` values

18.4.16.4. Usage Limitations of the Simple SearchUI

- If the Simple SearchUI results do not provide usable links on the metocard results, verify that a valid host has been entered in the Platform Global Configuration.

18.4.17. Thumbnail Metocard Transformer

The Thumbnail Metocard Transformer retrieves the thumbnail bytes of a Metocard by returning the [Metocard.THUMBNAIL](#) attribute value.

18.4.17.1. Installing the Thumbnail Metocard Transformer

This transformer is installed by default with a standard installation in the Catalog application.

18.4.17.2. Configuring the Thumbnail Metocard Transformer

The Thumbnail Metocard Transformer has no configurable properties.

18.4.17.3. Using the Thumbnail Metocard Transformer

Endpoints or other components can retrieve an instance of the Thumbnail Metocard Transformer using its id [thumbnail](#).

Sample Blueprint Reference Snippet

```
<reference id="metocardTransformer" interface="ddf.catalog.transform.MetocardTransformer" filter="(id=thumbnail)"/>
```

The Thumbnail Metocard Transformer returns a [BinaryContent](#) object of the [Metocard.THUMBNAIL](#) bytes and a MIME Type of [image/jpeg](#).

18.4.18. Tika Input Transformer

The Tika Input Transformer is the default input transformer responsible for translating Microsoft Word, Microsoft Excel, Microsoft PowerPoint, OpenOffice Writer, and PDF documents into Catalog Metacards. This input transformer utilizes [Apache Tika](#) to provide basic support for these mime types. As such, the metadata extracted from these types of documents is the metadata that is common across all of these document types, e.g., creation date, author, last modified date, etc. The Tika Input Transformer's main purpose is to ingest these types of content into the Metadata Catalog.

The Tika input transformer is given a service ranking (priority) of -1 so that it is guaranteed to be the last input transformer that is invoked. This allows any registered input transformer that are more specific for any of these document types to be invoked instead of this rudimentary input transformer.

18.4.18.1. Installing the Tika Input Transformer

This transformer is installed by default with a standard installation in the Catalog.

18.4.18.2. Configuring the Tika Input Transformer

The properties below describe how to configure the tika input transformer.

See [Tika Input Transformer configurations](#) for all possible configurations.

18.4.19. Video Input Transformer

The video input transformer Creates Catalog metacards from certain video file types. Currently, it handles MPEG-2 transport streams as well as MPEG-4, AVI, MOV, and WMV videos. This input transformer uses [Apache Tika](#) to extract basic metadata from the video files and applies more sophisticated methods to extract more meaningful metadata from these types of video.

18.4.19.1. Installing the Video Input Transformer

This transformer is installed by default with a standard installation in the Catalog application.

18.4.19.1.1. Configuring the Video Input Transformer

The Video Input Transformer has no configurable properties.

18.4.20. XML Input Transformer

The XML Input Transformer is responsible for translating an XML document into a Catalog Metacard.

18.4.20.1. Installing the XML Input Transformer

The XML Input Transformer is installed by default with a standard installation in the Catalog application.

18.4.20.2. Configuring the XML Input Transformer

The XML Input Transformer has no configurable properties.

18.4.21. XML Metacard Transformer

The XML metacard transformer is responsible for translating a metacard into an XML-formatted document. The metacard element that is generated is an extension of [gml:AbstractFeatureType](#), which makes the output of this transformer GML 3.1.1 compatible.

18.4.21.1. Installing the XML Metacard Transformer

This transformer comes installed by default with a standard installation in the Catalog application.

To install or uninstall manually, use the [catalog-transformer-xml](#) feature.

18.4.21.2. Configuring the XML Metacard Transformer

The XML Metacard Transformer has no configurable properties.

18.4.21.3. Using the XML Metacard Transformer

Using the REST Endpoint for example, request a metacard with the transform option set to the XML shortname.

XML Metacard Transformer URL

```
\https://{{FQDN}}:{{PORT}}/services/catalog/ac0c6917d5ee45bfb3c2bf8cd2ebaa67?transform=xml
```

Table 53. Metacard to XML Mappings

Metacard Variables	XML Element
<code>id</code>	<code>metacard/gml:id</code>
<code>metacardType</code>	<code>metacard/type</code>
<code>sourceId</code>	<code>metacard/source</code>
<code>all other attributes</code>	<code>metacard/<AttributeType>[name='<AttributeName>']/value</code> For instance, the value for the metacard attribute named "title" would be found at <code>metacard/string[@name='title']/value</code>

XML Adapted Attributes (AttributeTypes)

- `boolean`
- `base64Binary`
- `dateTime`
- `double`
- `float`
- `geometry`
- `int`
- `long`
- `object`
- `short`
- `string`
- `stringxml`

18.4.22. XML Query Response Transformer

The XML Query Response Transformer is responsible for translating a query response into an XML-formatted document. The metacard element generated is an extension of `gml:AbstractFeatureCollectionType`, which makes the output of this transformer GML 3.1.1 compatible.

18.4.22.1. Installing the XML Query Response Transformer

This transformer is installed by default with a standard installation in the Catalog application. To uninstall, uninstall the **catalog-transformer-xml** feature.

18.4.22.2. Configuring the XML Query Response Transformer

To configure the XML Query Response Transformer:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Configuration** tab.
4. Select the XML Query Response Transformer.

See [XML Query Response Transformer configurations](#) for all possible configurations.

18.4.22.3. Using the XML Query Response Transformer

Using the OpenSearch Endpoint, for example, query with the format option set to the XML shortname **xml**.

XML Query Response Transformer Query Example

```
\http://{FQDN}:{PORT}/services/catalog/query?q=input?format=xml
```

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<ns3:metacards xmlns:ns1="http://www.opengis.net/gml" xmlns:ns2=
"http://www.w3.org/1999/xlink" xmlns:ns3="urn:catalog:metacard" xmlns:ns4=
"http://www.w3.org/2001/SMIL20/" xmlns:ns5="http://www.w3.org/2001/SMIL20/Language">
  <ns3:metacard ns1:id="000ba4dd7d974e258845a84966d766eb">
    <ns3:type>ddf.metacard</ns3:type>
    <ns3:source>southwestCatalog1</ns3:source>
    <ns3:dateTime name="created">
      <ns3:value>2013-04-10T15:30:05.702-07:00</ns3:value>
    </ns3:dateTime>
    <ns3:string name="title">
      <ns3:value>Input 1</ns3:value>
    </ns3:string>
  </ns3:metacard>
  <ns3:metacard ns1:id="00c0eb4ba9b74f8b988ef7060e18a6a7">
    <ns3:type>ddf.metacard</ns3:type>
    <ns3:source>southwestCatalog1</ns3:source>
    <ns3:dateTime name="created">
      <ns3:value>2013-04-10T15:30:05.702-07:00</ns3:value>
    </ns3:dateTime>
    <ns3:string name="title">
      <ns3:value>Input 2</ns3:value>
    </ns3:string>
  </ns3:metacard>
</ns3:metacards>
```

18.5. Mime Type Mapper

The **MimeTypeMapper** is the entry point in DDF for resolving file extensions to mime types, and vice versa.

MimeTypeMappers are used by the **ResourceReader** to determine the file extension for a given mime type in aid of retrieving a product. **MimeTypeMappers** are also used by the **FileSystemProvider** in the Catalog Framework to read a file from the content file repository.

The **MimeTypeMapper** maintains a list of all of the **MimeTypeResolvers** in DDF.

The **MimeTypeMapper** accesses each **MimeTypeResolver** according to its priority until the provided file extension is successfully mapped to its corresponding mime type. If no mapping is found for the file extension, **null** is returned for the mime type. Similarly, the **MimeTypeMapper** accesses each **MimeTypeResolver** according to its priority until the provided mime type is successfully mapped to its corresponding file extension. If no mapping is found for the mime type, **null** is returned for the file extension.

For files with no file extension, the **MimeTypeMapper** will attempt to determine the mime type from the contents of the file. If it is unsuccessful, the file will be ingested as a binary file.

DDF Mime Type Mapper

Core implementation of the DDF Mime API.

18.5.1. DDF Mime Type Mapper

The DDF Mime Type Mapper is the core implementation of the DDF Mime API. It provides access to all [MimeTypeResolvers](#) within DDF, which provide mapping of mime types to file extensions and file extensions to mime types.

18.5.1.1. Installing the DDF Mime Type Mapper

The DDF Mime Type Mapper is installed by default with a standard installation in the Platform application.

18.5.1.2. Configuring DDF Mime Type Mapper

The DDF Mime Type Mapper has no configurable properties.

18.6. Mime Type Resolver

A [MimeTypeResolver](#) is a DDF service that can map a file extension to its corresponding mime type and, conversely, can map a mime type to its file extension.

[MimeTypeResolvers](#) are assigned a priority (0-100, with the higher the number indicating the higher priority). This priority is used to sort all of the [MimeTypeResolvers](#) in the order they should be checked to map a file extension to a mime type (or vice versa). This priority also allows custom [MimeTypeResolvers](#) to be invoked before default [MimeTypeResolvers](#) by setting custom resolver's priority higher than the default.

[MimeTypeResolvers](#) are not typically invoked directly. Rather, the [MimeTypeMapper](#) maintains a list of [MimeTypeResolvers](#) (sorted by their priority) that it invokes to resolve a mime type to its file extension (or to resolve a file extension to its mime type).

Custom Mime Type Resolver

The Custom Mime Type Resolver is a [MimeTypeResolver](#) that defines the custom mime types that DDF will support.

Tika Mime Type Resolver

Provides support for resolving over 1300 mime types.

18.6.1. Custom Mime Type Resolver

These are mime types not supported by the default [TikaMimeTypeResolver](#).

Table 54. Custom Mime Type Resolver Default Supported Mime Types

File Extension	Mime Type
nitf	image/nitf

File Extension	Mime Type
ntf	image/nitf
json	json=application/json;id=geojson

As a `MimeTypeResolver`, the Custom Mime Type Resolver will provide methods to map the file extension to the corresponding mime type, and vice versa.

18.6.1.1. Installing the Custom Mime Type Resolver

One Custom Mime Type Resolver is configured and installed for the `image/nitf` mime type. This custom resolver is bundled in the `mime-core-app` application and is part of the `mime-core` feature.

Additional Custom Mime Type Resolvers can be added for other custom mime types.

18.6.1.1.1. Configuring the Custom Mime Type Resolver

The configurable properties for the Custom Mime Type Resolver are accessed from the **MIME Custom Types** configuration in the Admin Console.

- Navigate to the Admin Console.
- Select the **Platform** application.
- Select **Configuration**.
- Select **MIME Custom Types**.

Managed Service Factory PID

- `Ddf_Custom_Mime_Type_Resolver`

See [Custom Mime Type Resolver configurations](#) for all possible configurations.

18.6.2. Tika Mime Type Resolver

The `TikaMimeTypeResolver` is a `MimeTypeResolver` that is implemented using the [Apache Tika](#) open source product.

Using the Apache Tika content analysis toolkit, the `TikaMimeTypeResolver` provides support for resolving over 1300 mime types, but not all mime types yield the same quality metadata.

The `TikaMimeTypeResolver` is assigned a default priority of `-1` to insure that it is always invoked last by the `MimeTypeMapper`. This insures that any custom `MimeTypeResolvers` that may be installed will be invoked before the `TikaMimeTypeResolver`.

The `TikaMimeTypeResolver` provides the bulk of the default mime type support for DDF.

18.6.2.1. Installing the Tika Mime Type Resolver


The `TikaMimeTypeResolver` is bundled as the `mime-tika-resolver` feature in the `mime-tika-app` application.

This feature is installed by default.

18.6.2.1.1. Configuring the Tika Mime Type Resolver

The Tika Mime Type Resolver has no configurable properties.

19. Catalog Plugins



Catalog Architecture: Catalog Plugins

Plugins are additional tools to use to add additional business logic at certain points, depending on the type of plugin.

The Catalog Framework calls Catalog Plugins to process requests and responses as they enter and leave the Framework.

19.1. Types of Plugins

Plugins can be designed to run before or after certain processes. They are often used for validation, optimization, or logging. Many plugins are designed to be called at more than one time. See [Catalog Plugin Compatibility](#).

Pre-Authorization Plugins

Perform any changes needed before security rules are applied.

Policy Plugins

Allows or denies access to the Catalog operation or response.

Access Plugins

Used to build policy information for requests.

Pre-Ingest Plugins

Perform any changes to a metocard prior to ingest.

Post-Ingest Plugins

Perform actions after ingest is completed.

Post-Process Plugins

Performs additional processing after ingest.

Pre-Query Plugins

Perform any changes to a query before execution.

Pre-Federated-Query Plugins

Perform any changes to a federated query before execution.

Post-Query Plugins

Perform any changes to a response after query completes.

Post-Federated-Query Plugins

Perform any changes to a response after federated query completes.

Pre-Resource Plugins

Perform any changes to a request associated with a metocard prior to download.

Post-Resource Plugins

Perform any changes to a resource after download.

Pre-Create Storage Plugins

Perform any changes before creating a resource.

Post-Create Storage Plugins

Perform any changes after creating a resource.

Pre-Update Storage Plugins

Perform any changes before updating a resource.

Post-Update Storage Plugins

Perform any changes after updating a resource.

Pre-Subscription Plugins

Perform any changes before creating a subscription.


Pre-Delivery Plugins

Perform any changes before delivering a subscribed event.


Plugins are called in a specific order during different operations. [Custom Plugins](#) can be added to the chain for special use cases.




Query Request Plugin Call Order




Create Request Plugin Call Order




Update Request Plugin Call Order




Delete Request Plugin Call Order



Resource Request Plugin Call Order



Storage Create Request Plugin Call Order



Storage Update Request Plugin Call Order

Table 55. Catalog Plugin Compatibility

Plugin	Pre-Authorization Plugins	Policy Plugins	Access Plugins	Pre-Ingest Plugins	Post-Ingest Plugins	Pre-Query Plugins	Post-Query Plugins	Post-Process Plugins
Catalog Backup Plugin					X			
Catalog Metrics Plugin					X	X	X	
Catalog Policy Plugin		X						
Client Info Plugin	X							
Content URI Access Plugin			X					
Event Processor					X			
Expiration Date Pre-Ingest Plugin				X				
Filter Plugin			X					
GeoCoder Plugin				X				
Historian Policy Plugin		X						
Identification Plugin				X	X			
JPEG2000 Thumbnail Converter							X	
Metocard Attribute Security Policy Plugin		X						

Plugin	Pre-Authorization Plugins	Policy Plugins	Access Plugins	Pre-Ingest Plugins	Post-Ingest Plugins	Pre-Query Plugins	Post-Query Plugins	Post-Process Plugins
Metocard Backup File Storage Provider					x			
Metocard Backup S3 Storage Provider					x			
Metocard Groomer				x				
Metocard Resource Size Plugin							x	
Metocard Validity Filter Plugin		x						
Metocard Validity Marker				x				
Metocard Ingest Network Plugin	x							
Operation Plugin			x					
Point of Contact Policy Plugin		x						
Processing Post-Ingest Plugin					x			
Registry Policy Plugin		x						

Plugin	Pre-Authorization Plugins	Policy Plugins	Access Plugins	Pre-Ingest Plugins	Post-Ingest Plugins	Pre-Query Plugins	Post-Query Plugins	Post-Process Plugins
Resource URI Policy Plugin		X						
Security Audit Plugin			X					
Security Logging Plugin				X	X	X	X	
Security Plugin			X					
Source Metrics Plugin					X	X	X	
Workspace Access Plugin			X					
Workspace Pre-Ingest Plugin				X				
Workspace Sharing Policy Plugin		X						
XML Attribute Security Policy Plugin		X						

Table 56. Catalog Plugin Compatibility, Cont.

Plugin	Pre-Federated-Query Plugins	Post-Federated-Query Plugins	Pre-Resource Plugins	Post-Resource Plugins	Pre-Create Storage Plugins	Post-Create Storage Plugins	Pre-Update Storage Plugins	Post-Update Storage Plugins	Pre-Subscription Plugins	Pre-Delivery Plugins
Catalog Metrics Plugin				X						

Plugin	Pre-Federated-Query Plugins	Post-Federated-Query Plugins	Pre-Resource Plugins	Post-Resource Plugins	Pre-Create Storage Plugins	Post-Create Storage Plugins	Pre-Update Storage Plugins	Post-Update Storage Plugins	Pre-Subscription Plugins	Pre-Delivery Plugins
Checksum Plugin					X		X			
Resource Usage Plugin			X	X						
Security Logging Plugin	X!	X	X	X	X	X	X	X		
Source Metrics Plugin				X						
Video Thumbnail Plugin						X		X		

19.1.1. Pre-Authorization Plugins

Pre-delivery plugins are invoked before any security rules are applied. This is an opportunity to take any action before authorization, including but not limited to:

- logging.
- adding network-specific information.
- adding user-identifying information.

19.1.1.1. Available Pre-Authorization Plugins

Client Info Plugin

Injects request-specific network information into a request.

Metocard Ingest Network Plugin

Adds attributes for network info from ingest request.

19.1.2. Policy Plugins

Policy plugins are invoked to set up the policy for a request/response. This provides an opportunity to attach custom requirements on operations or individual metacards. All the 'requirements' from each Policy plugin will be combined into a single policy that will be included in the request/response. Access plugins will be used to act on this combined policy.

19.1.2.1. Available Policy Plugins

Catalog Policy Plugin

Configures user attributes required for catalog operations.

Historian Policy Plugin

Protects metocard history from being edited by users without the history role.

Metocard Attribute Security Policy Plugin

Collects attributes into a security field for the metocard.

Metocard Validity Filter Plugin

Determines whether to filter metacards with validation errors or warnings.

Point of Contact Policy Plugin

Adds a policy if Point of Contact is updated.

Registry Policy Plugin

Defines user access polices for registry operations.

Resource URI Policy Plugin

Configures required user attributes for setting or altering a resource URI.

Workspace Sharing Policy Plugin

Collects attributes for a workspace to identify the appropriate policy to allow sharing.

XML Attribute Security Policy Plugin

Find security attributes contained in a metocard's metadata.

19.1.3. Access Plugins

Access plugins are invoked directly after the [Policy plugins](#) have been successfully executed. This is an opportunity to either stop processing or modify the request/response based on policy information.

19.1.3.1. Available Access Plugins

Content URI Access Plugin

Prevents a Metocard's resource URI from being overridden by an incoming UpdateRequest.

Filter Plugin

Performs filtering on query responses as they pass through the framework.

Operation Plugin

Validates a user or subject's security attributes.

Security Audit Plugin

Audits specific metocard attributes.


Security Plugin

Identifies the subject for an operation.

Workspace Access Plugin

Prevents non-owner users from changing workspace permissions.

19.1.4. Pre-Ingest Plugins



Ingest Plugin Flow

Pre-ingest plugins are invoked before an ingest operation is sent to the catalog. They are not run on a query. This is an opportunity to take any action on the ingest request, including but not limited to:

- validation.
- logging.
- auditing.
- optimization.
- security filtering.

19.1.4.1. Available Pre-Ingest Plugins

Expiration Date Pre-Ingest Plugin

Adds or updates expiration dates for the resource.

GeoCoder Plugin

Populates the `Location.COUNTRY_CODE` attribute if the Metocard has an associated location.

Identification Plugin

Manages IDs on registry metacards.

Metocard Groomer

Modifies metacards when created or updated.

Metocard Validity Marker

Modifies metacards when created or ingested according to metocard validator services.


Security Logging Plugin

Logs operations to the security log.

Workspace Pre-Ingest Plugin

Verifies that a workspace has an associated email to enable sharing.

19.1.5. Post-Ingest Plugins



Query Plugin Flow

Post-ingest plugins are invoked after data has been created, updated, or deleted in a Catalog Provider.

19.1.5.1. Available Post-Ingest Plugins

Catalog Backup Plugin

Enables backup of the catalog and its metacards.

Catalog Metrics Plugin

Captures metrics on catalog operations.

Event Processor

Creates, updates, and deletes subscriptions.

Identification Plugin

Manages IDs on registry metacards.

Metocard Backup File Storage Provider

Stores backed-up metacards.

Metocard Backup S3 Storage Provider

Stores backed-up metacards in a specified S3 bucket and key.

Processing Post-Ingest Plugin

Submits catalog Create, Update, or Delete requests to the Processing Framework.

Security Logging Plugin

Logs operations to the security log.

Source Metrics Plugin

Captures metrics on catalog operations.

19.1.6. Post-Process Plugins

NOTE

This code is experimental. While this interface is functional and tested, it may change or be removed in a future version of the library.

Post-Process Plugins are invoked after a metocard has been created, updated, or deleted and committed to the Catalog. They are the last plugins to run and are triggered by a **Post-Ingest Plugin**. Post-Process plugins are well-suited for asynchronous tasks. See the [Asynchronous Processing Framework](#) for more information about how **Post-Process Plugins** are used.

19.1.6.1. Available Post-Process Plugins

None.

19.1.7. Pre-Query Plugins

Pre-query plugins are invoked before a query operation is sent to any of the Sources. This is an opportunity to take any action on the query, including but not limited to:

- validation.
- logging.
- auditing.
- optimization.
- security filtering.

19.1.7.1. Available Pre-Query Plugins

Catalog Metrics Plugin

Captures metrics on catalog operations.

Security Logging Plugin

Logs operations to the security log.

Source Metrics Plugin

Captures metrics on catalog operations.

19.1.8. Pre-Federated-Query Plugins

Pre-federated-query plugins are invoked before a federated query operation is sent to any of the Sources. This is an opportunity to take any action on the query, including but not limited to:

- validation.
- logging.
- auditing.
- optimization.
- security filtering.

19.1.8.1. Available Pre-Federated-Query Plugins

Security Logging Plugin

Logs operations to the security log.

Tags Filter Plugin

Updates queries without filters.

19.1.9. Post-Query Plugins

Post-query plugins are invoked after a query has been executed successfully, but before the response is returned to the endpoint. This is an opportunity to take any action on the query response, including but not limited to:

- logging.
- auditing.
- security filtering/redaction.
- deduplication.

19.1.9.1. Available Post-Query Plugins

Catalog Metrics Plugin

Captures metrics on catalog operations.

JPEG2000 Thumbnail Converter

Creates thumbnails for jpeg2000 images.

Metocard Resource Size Plugin

Updates the resource size attribute of a metocard.

Security Logging Plugin

Logs operations to the security log.

Source Metrics Plugin

Captures metrics on catalog operations.

19.1.10. Post-Federated-Query Plugins

Post-federated-query plugins are invoked after a federated query has been executed successfully, but before the response is returned to the endpoint. This is an opportunity to take any action on the query response, including but not limited to:

- logging.
- auditing.
- security filtering/redaction.
- deduplication.

19.1.10.1. Available Post-Federated-Query Plugins

Security Logging Plugin

Logs operations to the security log.

19.1.11. Pre-Resource Plugins

Pre-Resource plugins are invoked before a request to retrieve a resource is sent to a Source. This is an opportunity to take any action on the request, including but not limited to:

- validation.
- logging.
- auditing.
- optimization.
- security filtering.

19.1.11.1. Available Pre-Resource Plugins

Resource Usage Plugin

Monitors and limits system data usage.

Security Logging Plugin

Logs operations to the security log.

19.1.12. Post-Resource Plugins

Post-resource plugins are invoked after a resource has been retrieved, but before it is returned to the endpoint. This is an opportunity to take any action on the response, including but not limited to:

- logging.
- auditing.
- security filtering/redaction.

19.1.12.1. Available Post-Resource Plugins

Catalog Metrics Plugin

Captures metrics on catalog operations.

Resource Usage Plugin

Monitors and limits system data usage.

Security Logging Plugin

Logs operations to the security log.

Source Metrics Plugin

Captures metrics on catalog operations.

19.1.13. Pre-Create Storage Plugins

Pre-Create storage plugins are invoked immediately before an item is created in the content repository.

19.1.13.1. Available Pre-Create Storage Plugins

Checksum Plugin

Creates a unique checksum for ingested resources.

Security Logging Plugin

Logs operations to the security log.

19.1.14. Post-Create Storage Plugins

Post-Create storage plugins are invoked immediately after an item is created in the content repository.

19.1.14.1. Available Post-Create Storage Plugins

Security Logging Plugin

Logs operations to the security log.

Video Thumbnail Plugin

Generates thumbnails for video files.

19.1.15. Pre-Update Storage Plugins

Pre-Update storage plugins are invoked immediately before an item is updated in the content repository.

19.1.15.1. Available Pre-Update Storage Plugins

Checksum Plugin

Creates a unique checksum for ingested resources.

Security Logging Plugin

Logs operations to the security log.

19.1.16. Post-Update Storage Plugins

Post-Update storage plugins are invoked immediately after an item is updated in the content repository.

19.1.16.1. Available Post-Update Storage Plugins

Security Logging Plugin

Logs operations to the security log.

Video Thumbnail Plugin

Generates thumbnails for video files.

19.1.17. Pre-Subscription Plugins

Pre-subscription plugins are invoked before a Subscription is activated by an Event Processor. This is an opportunity to take any action on the Subscription, including but not limited to:

- validation.
- logging.
- auditing.
- optimization.
- security filtering.

19.1.17.1. Available Pre-Subscription Plugins

None.

19.1.18. Pre-Delivery Plugins

Pre-delivery plugins are invoked before a Delivery Method is invoked on a Subscription. This is an opportunity to take any action before event delivery, including but not limited to:

- logging.
- auditing.
- security filtering/redaction.

19.1.18.1. Available Pre-Delivery Plugins

None.

19.2. Catalog Plugin Details

Installation and configuration details listed by plugin name.

19.2.1. Catalog Backup Plugin

The Catalog Backup Plugin is used to enable data backup of the catalog and the metacards it contains.

WARNING

Catalog Backup Plugin Considerations

Using this plugin may impact performance negatively.

19.2.1.1. Installing the Catalog Backup Plugin

The Catalog Backup Plugin is installed by default with a standard installation in the Catalog application.

19.2.1.2. Configuring the Catalog Backup Plugin

To configure the Catalog Backup Plugin:

1. Navigate to the **Admin Console**.
2. Select **Catalog** application.
3. Select **Configuration** tab.
4. Select **Backup Post-Ingest Plugin**.

See [Catalog Backup Plugin configurations](#) for all possible configurations.

19.2.1.3. Usage Limitations of the Catalog Backup Plugin

- May affect performance.
- Must be installed prior to ingesting any content.
- Once enabled, disabling *may* cause incomplete backups.

19.2.2. Catalog Metrics Plugin

The Catalog Metrics Plugin captures metrics on catalog operations. These metrics can be viewed and analyzed using the [Metrics Reporting Application](#) in the Admin Console.

19.2.2.1. Related Components to the Source Metrics Plugin

- [Source Metrics Plugin](#).

19.2.2.2. Installing the Catalog Metrics Plugin

The Catalog Metrics Plugin is installed by default with a standard installation in the Catalog application.

19.2.2.3. Configuring the Catalog Metrics Plugin

The Catalog Metrics Plugin has no configurable properties.

19.2.3. Catalog Policy Plugin

The Catalog Policy Plugin configures the attributes required for users to perform Create, Read, Update, and Delete operations on the catalog.

19.2.3.1. Installing the Catalog Policy Plugin

The Catalog Policy Plugin is installed by default with a standard installation in the Catalog application.

19.2.3.2. Configuring the Catalog Policy Plugin

To configure the Catalog Policy Plugin:

1. Navigate to the **Admin Console**.
2. Select Catalog application.
3. Select **Configuration** tab.
4. Select **Catalog Policy Plugin**.

See [Catalog Policy Plugin configurations](#) for all possible configurations.

19.2.4. Checksum Plugin

The Checksum plugin creates a unique checksum for resources input into the system to identify updated content.

19.2.4.1. Installing the Checksum Plugin

The Checksum is installed by default with a standard installation in the Catalog application.

19.2.4.2. Configuring the Checksum Plugin

The Checksum Plugin has no configurable properties.

19.2.5. Client Info Plugin

The client info plugin injects request-specific network information into request properties, such as Remote IP Address, Remote Host Name, Servlet Scheme, and Servlet Context.

19.2.5.1. Related Components to the Client Info Plugin

- Client info filter
- [Metocard Ingest Network Plugin](#)

19.2.5.2. Installing the Client Info Plugin

The Client Info Plugin is installed by default with a standard installation in the Catalog application.

19.2.5.3. Configuring the Client Info Plugin

The Client Info Plugin has no configurable properties.

19.2.6. Content URI Access Plugin

The Content URI Access Plugin prevents a Metocard's resource URI from being overridden by an incoming UpdateRequest.

19.2.6.1. Installing the Content URI Access Plugin

The Content URI Access Plugin is installed by default with a standard installation in the Catalog application.

19.2.6.2. Configuring the Content URI Access Plugin

The Content URI Access Plugin has no configurable properties.

19.2.7. Event Processor

The Event Processor creates, updates, and deletes subscriptions for event notification. These subscriptions optionally specify a filter criteria so that only events of interest to the subscriber are posted for notification.

As metacards are created, updated, and deleted, the Catalog's Event Processor is invoked (as a post-ingest plugin) for each of these events. The Event Processor applies the filter criteria for each registered subscription to each of these ingest events to determine if they match the criteria.

For more information on creating subscriptions, see [Creating a Subscription](#).

19.2.7.1. Installing the Event Processor

The Event Processor is installed by default with a standard installation in the Catalog application.

19.2.7.2. Configuring the Event Processor

The Event Processor has no configurable properties.

19.2.7.3. Usage Limitations of the Event Processor

The Standard Event processor currently broadcasts federated events and should not. It should only broadcast events that were generated locally, all other events should be dropped. See [DDF-3151](#) for status.

19.2.8. Expiration Date Pre-Ingest Plugin

The Expiration Date plugin adds or updates expiration dates which can be used later for archiving old data.

19.2.8.1. Installing the Expiration Date Pre-Ingest Plugin

The Expiration Date Pre-Ingest Plugin is not installed by default with a standard installation. To install:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Configuration** tab.
4. Select the **Expiration Data Pre-Ingest Plugin**.

19.2.8.2. Configuring the Expiration Date Pre-Ingest Plugin

To configure the Expiration Date Pre-Ingest Plugin:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Configuration** tab.
4. Select the **Expiration Date Pre-Ingest Plugin**.

See [Expiration Date Plugin configurations](#) for all possible configurations.

19.2.9. Filter Plugin

The Filter Plugin performs filtering on query responses as they pass through the framework.

Each metocard result can contain security attributes that are pulled from the metadata record after being processed by a **PolicyPlugin** that populates this attribute. The security attribute is a Map containing a set of keys that map to lists of values. The metocard is then processed by a filter plugin that creates a **KeyValueCollectionPermission** from the metocard's security attribute. This permission is then checked against the user subject to determine if the subject has the correct claims to view

that metocard. The decision to filter the metocard eventually relies on the installed [Policy Decision Point](#) (PDP). The PDP that is being used returns a decision, and the metocard will either be filtered or allowed to pass through.

How a metocard gets filtered is left up to any number of FilterStrategy implementations that might be installed. Each FilterStrategy will return a result to the filter plugin that says whether or not it was able to process the metocard, along with the metocard or response itself. This allows a metocard or entire response to be partially filtered to allow some data to pass back to the requester. This could also include filtering any products sent back to a requester.

The security attributes populated on the metocard are completely dependent on the type of the metocard. Each type of metocard must have its own [PolicyPlugin](#) that reads the metadata being returned and then returns the appropriate attributes.

Example (represented as simple XML for ease of understanding):

```
<metocard>
  <security>
    <map>
      <entry assertedAttribute1="A,B" />
      <entry assertedAttribute2="X,Y" />
      <entry assertedAttribute3="USA,GBR" />
      <entry assertedAttribute4="USA,AUS" />
    </map>
  </security>
</metocard>
```

```
<user>
  <claim name="subjectAttribute1">
    <value>A</value>
    <value>B</value>
  </claim>
  <claim name="subjectAttribute2">
    <value>X</value>
    <value>Y</value>
  </claim>
  <claim name="subjectAttribute3">
    <value>USA</value>
  </claim>
  <claim name="subjectAttribute4">
    <value>USA</value>
  </claim>
</user>
```

In the above example, the user's claims are represented very simply and are similar to how they would actually appear in a SAML 2 assertion. Each of these user (or subject) claims will be converted to a [KeyValuePermission](#) object. These permission objects will be implied against the permission object generated from the metocard record. In this particular case, the metocard might be allowed if the policy is configured appropriately because all of the permissions line up correctly.

19.2.9.1. Installing the Filter Plugin

The Filter Plugin is installed by default with a standard installation in the Catalog application.

19.2.9.2. Configuring the Filter Plugin

The Filter Plugin has no configurable properties.

19.2.10. GeoCoder Plugin

The GeoCoder Plugin is a pre-ingest plugin that is responsible for populating the Metacard's `Location.COUNTRY_CODE` attribute if the Metacard has an associated location. If there is a valid country code for the Metacard, it will be in ISO 3166-1 alpha-3 format. If the metacard's country code is already populated, the plugin will **not** override it. The GeoCoder relies on either the WebService or [Offline Gazetteer](#) to retrieve country code information.

WARNING

For a polygon or polygons, this plugin takes the center point of the bounding box to assign the country code.

19.2.10.1. Installing the GeoCoder Plugin

The GeoCoder Plugin is installed by default with the Spatial application, when the WebService or Offline Gazetteer is started.

19.2.10.2. Configuring the GeoCoder Plugin

To configure the GeoCoder Plugin:

1. Navigate to the **Admin Console**.
2. Select **Spatial** application.
3. Select **Configuration** tab.
4. Select **Backup Post-Ingest Plugin**.

These are the available configurations:

See [GeoCoder Plugin configurations](#) for all possible configurations.

19.2.11. Historian Policy Plugin

The Historian Policy Plugin protects metacard history from being edited or deleted by users without the history role (a <http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role> of `system-history`).

19.2.11.1. Installing the Historian Policy Plugin

The Historian is installed by default with a standard installation in the Catalog application.

19.2.11.2. Configuring the Historian Policy Plugin

The Historian Policy Plugin has no configurable properties.

19.2.12. Identification Plugin

The Identification Plugin assigns IDs to registry metacards and adds/updates IDs on create and update.

19.2.12.1. Installing the Identification Plugin

The Identification Plugin is not installed by default in a standard installation. It is installed by default with the [Registry](#) application.

19.2.12.2. Configuring the Identification Plugin

The Identification Plugin has no configurable properties.

19.2.13. JPEG2000 Thumbnail Converter

The JPEG2000 Thumbnail converter creates thumbnails from images ingested in jpeg2000 format.

19.2.13.1. Installing the JPEG2000 Thumbnail Converter

The JPEG2000 Thumbnail Converter is installed by default with a standard installation in the Catalog application.

19.2.13.2. Configuring the JPEG2000 Thumbnail Converter

The JPEG2000 Thumbnail Converter has no configurable properties.

19.2.14. Metocard Attribute Security Policy Plugin

The Metocard Attribute Security Policy Plugin combines existing metocard attributes to make new attributes and adds them to the metocard. For example, if a metocard has two attributes, `sourceattribute1` and `sourceattribute2`, the values of the two attributes could be combined into a new attribute, `destinationattribute1`. The `sourceattribute1` and `sourceattribute2` are the *source attributes* and `destinationattribute1` is the *destination attribute*.

There are two way to combine the values of source attributes. The first, and most common, is to take all of the attribute values and put them together. This is called the union. For example, if the source attributes `sourceattribute1` and `sourceattribute2` had the values:

`sourceattribute1 = MASK, VESSEL`

`sourceattribute2 = WIRE, SACK, MASK`

...the **union** would result in the new attribute **destinationattribute1**:

```
destinationattribute1 = MASK, VESSEL, WIRE, SACK
```

The other way to combine attributes is use the values common to all of the attributes. This is called the **intersection**. Using our previous example, the **intersection** of **sourceattribute1** and **sourceattribute2** would create the new attribute **destinationattribute1**

```
destinationattribute1 = MASK
```

because only **MASK** is common to all of the source attributes.

The policy plugin could also be used to rename attributes. If there is only one source attribute, and the combination policy is union, then the attribute's values are effectively renamed to the destination attribute.

19.2.14.1. Installing the Metocard Attribute Security Policy Plugin

The Metocard Attribute Security Policy Plugin is installed by default with a standard installation in the Catalog application.

19.2.14.2. Configuring the Metocard Attribute Security Policy Plugin

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application tile
3. Select the **Configuration** tab
4. Select the **Metocard Attribute Security Policy Plugin**.

See [Metocard Attribute Security Policy Plugin configurations](#) for all possible configurations.

19.2.14.3. Sample configuration of the Metocard Attribute Security Policy Plugin

To configure the plugin to combine the attributes **sourceattribute1** and **sourceattribute2** into a new attribute **destinationattribute1** using the union, enter these two lines under the title **Metocard Union Attributes**

Metocard Union Attributes

```
sourceattribute1=destinationattribute1
sourceattribute2=destinationattribute1
```

19.2.15. Metocard Backup File Storage Provider

The Metocard Backup File Storage Provider is a storage provider that will store backed-up metacards in a specified file system location.

19.2.15.1. Installing the Metocard Backup File Storage Provider

To install the Metocard Backup File Storage Provider

1. Navigate to the **Admin Console**.
2. Select Catalog application.
3. Select **Features** tab.
4. Install the `catalog-metocard-backup-filestorage` feature.

19.2.15.2. Configuring the Metocard Backup File Storage Provider

To configure the Metocard Backup File Storage Provider

1. Navigate to the **Admin Console**.
2. Select Catalog application.
3. Select **Configuration** tab.
4. Select **Metocard Backup File Storage Provider**.

See [Metocard Backup File Storage Provider configurations](#) for all possible configurations.

19.2.16. Metocard Backup S3 Storage Provider

The Metocard Backup S3 Storage Provider is a storage provider that will store backed up metacards in the specified S3 bucket and key.

19.2.16.1. Installing the Metocard S3 File Storage Provider

To install the Metocard Backup File Storage Provider

1. Navigate to the **Admin Console**.
2. Select Catalog application.
3. Select **Features** tab.
4. Install the `catalog-metocard-backup-s3storage` feature.

19.2.16.2. Configuring the Metocard S3 File Storage Provider

To configure the Metocard Backup S3 Storage Provider:

1. Navigate to the **Admin Console**.
2. Select Catalog application.
3. Select **Configuration** tab.
4. Select **Metocard Backup S3 Storage Provider**.

See [Metocard Backup S3 Storage Provider configurations](#) for all possible configurations.

19.2.17. Metacard Groomer

The Metacard Groomer Pre-Ingest plugin makes modifications to [CreateRequest](#) and [UpdateRequest](#) metacards.

Use this pre-ingest plugin as a convenience to apply basic rules for your metacards.

This plugin makes the following modifications when metacards are in a [CreateRequest](#):

- Overwrites the [Metacard.ID](#) field with a generated, unique, 32 character hexadecimal value if missing or if the resource URI is not a catalog resource URI.
- Sets [Metacard.CREATED](#) to the current time stamp if not already set.
- Sets [Metacard.MODIFIED](#) to the current time stamp if not already set.
- Sets [Core.METACARD_CREATED](#) to the current time stamp if not present.
- Sets [Core.METACARD_MODIFIED](#) to the current time stamp.

In an [UpdateRequest](#), the same operations are performed as a [CreateRequest](#), except:

- If no value is provided for [Metacard.ID](#) in the new metocard, it will be set using the [UpdateRequest](#) ID if applicable.

19.2.17.1. Installing the Metacard Groomer

The Metacard Groomer is installed by default with a standard installation in the Catalog application.

19.2.17.2. Configuring the Metacard Groomer

The Metacard Groomer has no configurable properties.

19.2.18. Metacard Resource Size Plugin

This post-query plugin updates the resource size attribute of each metocard in the query results if there is a cached file for the product and it has a size greater than zero; otherwise, the resource size is unmodified and the original result is returned.

Use this post-query plugin as a convenience to return query results with accurate resource sizes for cached products.

19.2.18.1. Installing the Metacard Resource Size Plugin

The Metacard Resource Size Plugin is installed by default with a standard installation.

19.2.18.2. Configuring the Metacard Resource Size Plugin

The Metacard Resource Size Plugin has no configurable properties.

19.2.19. Metocard Validity Filter Plugin

The Metocard Validity Filter Plugin determines whether metacards with validation errors or warnings are filtered from query results.

19.2.19.1. Related Components to the Metocard Validity Filter Plugin

- [Metocard Validity Marker](#).

19.2.19.2. Installing the Metocard Validity Filter Plugin

The Metocard Validity Filter Plugin is installed by default with a standard installation in the Catalog application.

19.2.19.3. Configuring the Metocard Validity Filter Plugin

The Metocard Validity Filter Plugin can be configured:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Configuration** tab.
4. Select the **Metocard Validity Filter Plugin**.

See [Metocard Validity Filter Plugin configurations](#) for all possible configurations.

19.2.20. Metocard Validity Marker

The Metocard Validity Marker Pre-Ingest plugin modifies the metacards contained in create and update requests.

The plugin runs each metocard in the `CreateRequest` and `UpdateRequest` against each registered `MetocardValidator` service.

NOTE This plugin can make it seem like ingested products are not successfully ingested if a user does not have permissions to access invalid metacards. If an ingest did not fail, there are no errors in the ingest log, but the expected results do not show up after a query, verify either that the ingested data is valid or that the [Metocard Validity Filter Plugin](#) is configured to show warnings and/or errors.

19.2.20.1. Related Components to the Metocard Validity Marker

- [Metocard Validity Filter Plugin](#).

19.2.20.2. Installing Metocard Validity Marker

This plugin is installed by default with a standard installation in the Catalog application.

19.2.20.3. Configuring Metacard Validity Marker

By default, it will mark metacards with validation errors and warnings as they are reported by each metacard validator and then allow the ingest. To prevent the ingest of certain invalid metacards, the **Metacard Validity Marker** plugin can be configured to "enforce" one or more validators. Metacards that are invalid according to an "enforced" validator will not be ingested.

See [Metacard Validity Marker Plugin configurations](#) for all possible configurations.

19.2.20.4. Using Metacard Validity Marker

Use this pre-ingest plugin to validate metacards against metacard validators, which can check schemas, schematron, or any other logic.

19.2.21. Metacard Ingest Network Plugin

The Metacard Ingest Network Plugin allows the conditional insertion of new attributes on metacards during ingest based on network information from the ingest request; including IP address and hostname.

For the extent of this section, a 'rule' will refer to a configured, single instance of this plugin.

19.2.21.1. Related Components to the Metacard Ingest Network Plugin

- [Client Info Plugin](#)

19.2.21.2. Installing the Metacard Ingest Network Plugin

The Metacard Ingest Network Plugin is installed by default during a standard installation in the Catalog application.

19.2.21.3. Configuring the Metacard Ingest Network Plugin

To configure the Metacard Ingest Network Plugin:

- Navigate to the **Admin Console**.
- Select the Catalog application.
- Select the **Configuration** tab.
- Select the label *Metacard Ingest Network Plugin* to setup a network rule.

See [Metacard Ingest Network Plugin configurations](#) for all possible configurations.

Multiple instances of the plugin can be configured by clicking on its configuration title within the configuration tab of the Catalog app. Each instance represents a conditional statement, or a 'rule', that gets evaluated for each ingest request. For any request that meets the configured criteria of a rule, that rule will attempt to transform its list of key-value pairs to become new attributes on all metacards in that request.

The rule is divided into two fields: "Criteria" and "Expected Value". The "Criteria" field features a drop-down list containing the four elements for which equality can be tested:

- IP Address of where the ingest request came from
- Host Name of where the ingest request came from
- Scheme that the ingest request arrived on, for example, *http* vs *https*
- Context Path that the ingest request arrived on, for example, */services/catalog*

In order for a rule to evaluate to true and the attributes be applied, the value in the "Expected Value" field must be an exact match to the actual value of the selected criteria. For example, if the selected criteria is "IP Address" with an expected value of "192.168.0.1", the rule only evaluates to true for ingest requests coming from "192.168.0.1" and nowhere else.

Check for IPv6

IMPORTANT

Verify your system's IP configuration. Rules using "IP Address" may need to be written in IPv6 format.

The key-value pairs within each rule should take the following form: "key = value" where the "key" is the name of the attribute and the "value" is the value assigned to that attribute. Whitespace is ignored unless it is within the key or value. Multi-valued attributes can be expressed in comma-separated format if necessary.

Examples of Valid Attribute Assignments

```
contact.contributor-name = John Doe
contact.contributor-email = john.doe@example.net
language = English
language = English, French, German
security.access-groups = SJ202, SR 101, JS2201
```

19.2.21.3.1. Useful Attributes

The following table provides some useful attributes that may commonly be set by this plugin:

Table 57. Useful Attributes

Attribute Name	Expected Format	Multi-Valued
expiration	ISO DateTime	no
description	Any String	no
metocard.owner	Any String	no
language	Any String	yes
security.access-groups	Any String	yes
security.access-individuals	Any String	yes

19.2.21.4. Usage Limitations of the Metocard Ingest Network Plugin

- This plugin only works for ingest (create requests) performed over a network; data ingested via command line does not get processed by this plugin.
 - Any attribute that is already set on the metocard will not be overwritten by the plugin.
 - The order of execution is not guaranteed. For any rule configuration where two or more rules add different values for the same attribute, it is undefined what the final value for that attribute will be in the case where more than one of those rules evaluates to true.
-

19.2.22. Operation Plugin

The operation plugin validates the subject's security attributes to ensure they are adequate to perform the operation.

19.2.22.1. Installing the Operation Plugin

The Operation Plugin is installed by default with a standard installation in the Catalog application.

19.2.22.2. Configuring the Operation Plugin

The Operation Plugin has no configurable properties.

19.2.23. Point of Contact Policy Plugin

The Point of Contact Policy Plugin is a PreUpdate plugin that will check if the point-of-contact attribute has changed. If it does, then it adds a policy to that metocard's policy map that cannot be implied. This will deny such an update request, which essentially makes the point-of-contact attribute read-only.

19.2.23.1. Related Components to Point of Contact Policy Plugin

[Point of Contact Update Plugin](#)

19.2.23.2. Installing the Point of Contact Policy Plugin

The Point of Contact Policy Plugin is installed by default with a standard installation in the Catalog application.

19.2.23.3. Configuring the Point of Contact Policy Plugin

The Point of Contact Policy Plugin has no configurable properties.

19.2.24. Processing Post-Ingest Plugin

The Processing Post Ingest Plugin is responsible for submitting catalog Create, Update, and Delete

(CUD) requests to the [Processing Framework](#).

19.2.24.1. Related Components to Processing Post-Ingest Plugin

None.

19.2.24.2. Installing the Processing Post-Ingest Plugin

The Processing Post-Ingest Plugin is not installed by default with a standard installation, but is installed by default when the in-memory Processing Framework is installed.

19.2.24.3. Configuring the Processing Post-Ingest Plugin

The Processing Post-Ingest Plugin has no configurable properties.

19.2.25. Registry Policy Plugin

The Registry Policy Plugin defines the policies for user access to registry entries and operations.

19.2.25.1. Installing the Registry Policy Plugin

The Registry Policy Plugin is not installed by default on a standard installation. It is installed with the [Registry](#) application.

19.2.25.2. Configuring the Registry Policy Plugin

The Registry Policy Plugin can be configured from the Admin Console:

1. Navigate to the [Admin Console](#).
2. Select the [Registry](#) application.
3. Select the [Configuration](#) tab.
4. Select [Registry Policy Plugin](#).

See [Registry Policy Plugin configurations](#) for all possible configurations.

19.2.26. Resource URI Policy Plugin

The Resource URI Policy Plugin configures the attributes required for users to set the resource URI when creating a metocard or alter the resource URI when updating an existing metocard in the catalog.

19.2.26.1. Installing the Resource URI Policy Plugin

The Resource URI Policy Plugin is installed by default with a standard installation in the Catalog application.

19.2.26.2. Configuring the Resource URI Policy Plugin

To configure the Resource URI Policy Plugin:

1. Navigate to the **Admin Console**.
2. Select Catalog application.
3. Select **Configuration** tab.
4. Select **Resource URI Policy Plugin**.

See [Resource URI Policy Plugin configurations](#) for all possible configurations.

19.2.27. Resource Usage Plugin

The Resource Usage Plugin monitors and limits data usage, and enables cancelling long-running queries.

19.2.27.1. Installing the Resource Usage Plugin

The Resource Usage Plugin is not installed by default with a standard installation. It is installed with the Resource Management application.

19.2.27.2. Configuring the Resource Usage Plugin

The Resource Usage Plugin can be configured from the Admin Console:

1. Navigate to the **Admin Console**.
2. Select the **Resource Management** application.
3. Select the **Configuration** tab.
4. Select **Data Usage**.

See [Resource Usage Plugin configurations](#) for all possible configurations.

19.2.28. Security Audit Plugin

The Security Audit Plugin is used to allow the auditing of specific metocard attributes. Any time a metocard attribute listed in the configuration is updated, a log will be generated in the security log.

19.2.28.1. Installing the Security Audit Plugin

The Security Audit Plugin is installed by default with a standard installation in the Catalog application.

19.2.28.2. Configuring the Security Audit Plugin

To configure the Security Audit Plugin:

1. Navigate to the **Admin Console**.
2. Select **Catalog** application.
3. Select **Configuration** tab.
4. Select **Security Audit Plugin**.

Add the desired metocard attributes that will be audited when modified.

These configurations are available:

See [Security Audit Plugin configurations](#) for all possible configurations.

19.2.29. Security Logging Plugin

The Security Logging Plugin logs operations to the security log.

19.2.29.1. Installing Security Logging Plugin

The Security Logging Plugin is installed by default in a standard installation in the Security application.

19.2.29.2. Configuring the Security Logging Plugin

The Security Logging Plugin has no configurable properties.

19.2.30. Security Plugin

The Security Plugin identifies the subject for an operation.

19.2.30.1. Installing the Security Plugin

The Security Plugin is installed by default with a standard installation in the Catalog application.

19.2.30.2. Configuring the Security Plugin

The Security Plugin has no configurable properties.

19.2.31. Source Metrics Plugin

The Source Metrics Plugin captures metrics on catalog operations. These metrics can be viewed and analyzed using the [Metrics Reporting Application](#) in the Admin Console.

19.2.31.1. Related Components to the Source Metrics Plugin

- [Catalog Metrics Plugin](#).

19.2.31.2. Installing the Source Metrics Plugin

The Source Metrics Plugin is installed by default with a standard installation in the Catalog application.

19.2.31.3. Configuring the Source Metrics Plugin

The Source Metrics Plugin has no configurable properties.

19.2.32. Tags Filter Plugin

The Tags Filter Plugin updates queries without filters for tags, and adds a default tag of [resource](#). For backwards compatibility, a filter will also be added to include metacards without any tags attribute.

19.2.32.1. Related Components to Tags Filter Plugin

None.

19.2.32.2. Installing the Tags Filter Plugin

The Tags Filter Plugin is installed by default with a standard installation in the Catalog application.

19.2.32.3. Configuring the Tags Filter Plugin

The Tags Filter Plugin has no configurable properties.

19.2.33. Video Thumbnail Plugin

The Video Thumbnail Plugin provides the ability to generate thumbnails for video files stored in the Content Repository.

It is an implementation of both the [PostCreateStoragePlugin](#) and [PostUpdateStoragePlugin](#) interfaces. When installed, it is invoked by the Catalog Framework immediately after a content item has been created or updated by the Storage Provider.

This plugin uses a custom 32-bit LGPL build of FFmpeg (a video processing program) to generate thumbnails. When this plugin is installed, it places the FFmpeg executable appropriate for the current operating system in [`<DDF_HOME>/bin_third_party/ffmpeg`](#). When invoked, this plugin runs the FFmpeg binary in a separate process to generate the thumbnail. The [`<DDF_HOME>/bin_third_party/ffmpeg`](#) directory is deleted when the plugin is uninstalled.

NOTE Prebuilt FFmpeg binaries are provided for Linux, Mac, and Windows only.

19.2.33.1. Installing the Video Thumbnail Plugin

The Video Thumbnail Plugin is installed by default with a standard installation in the Catalog

application.

19.2.33.2. Configuring the Video Thumbnail Plugin

To configure the Video Thumbnail Plugin:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Configuration** tab.
4. Select the **Video Thumbnail Plugin**.

See [Video Thumbnail Plugin configurations](#) for all possible configurations.

19.2.34. Workspace Access Plugin

The Workspace Access Plugin prevents non-owner users from changing workspace permissions.

19.2.34.1. Related Components to The Workspace Access Plugin

- [Workspace Sharing Policy Plugin](#).
- [Workspace Pre-Ingest Plugin](#).
- Workspace Extension.

19.2.34.2. Installing the Workspace Access Plugin

The Workspace Access Plugin is installed by default with a standard installation in the Catalog application.

19.2.34.3. Configuring the Workspace Access Plugin

The Workspace Access Plugin has no configurable properties.

19.2.35. Workspace Pre-Ingest Plugin

The Workspace Pre-Ingest Plugin verifies that a workspace has an associated email to enable sharing and assigns that email as "owner".

19.2.35.1. Related Components to The Workspace Pre-Ingest Plugin

- [Workspace Sharing Policy Plugin](#).
- [Workspace Access Plugin](#).
- Workspace Extension.

19.2.35.2. Installing the Workspace Pre-Ingest Plugin

The Workspace Pre-Ingest Plugin is installed by default with a standard installation in the Catalog application.

19.2.35.3. Configuring the Workspace Pre-Ingest Plugin

The Workspace Pre-Ingest Plugin has no configurable properties.

19.2.36. Workspace Sharing Policy Plugin

The Workspace Sharing Policy Plugin collects attributes for a workspace to identify the appropriate policy to apply to allow sharing.

19.2.36.1. Related Components to The Workspace Sharing Policy Plugin

- [Workspace Access Plugin](#).
- [Workspace Pre-Ingest Plugin](#).
- [Workspace Extension](#).

19.2.36.2. Installing the Workspace Sharing Policy Plugin

The Workspace Sharing Policy Plugin is installed by default with a standard installation in the Catalog application.

19.2.36.3. Configuring the Workspace Sharing Policy Plugin

The Workspace Sharing Policy Plugin has no configurable properties.

19.2.37. XML Attribute Security Policy Plugin

The XML Attribute Security Policy Plugin parses XML metadata contained within a metocard for security attributes on any number of XML elements in the metadata. The configuration for the plugin contains one field for setting the XML elements that will be parsed for security attributes and the other two configurations contain the XML attributes that will be pulled off of those elements. The **Security Attributes (union)** field will compute the union of values for each attribute defined and the **Security Attributes (intersection)** field will compute the intersection of values for each attribute defined.

19.2.37.1. Installing the XML Attribute Security Policy Plugin


The XML Attribute Security Policy Plugin is installed by default with a standard installation in the Security application.

19.2.37.2. Configuring the XML Attribute Security Policy Plugin

1. Navigate to the Admin Console.
2. Select the **Catalog** application tile.
3. Select the **Configuration** tab.
4. Select the **XML Attribute Security Policy Plugin** configuration.

See [XML Attribute Security Policy Plugin configurations](#) for all possible configurations.

20. Operations




The Catalog provides the capability to query, create, update, and delete metacards; retrieve resources; and retrieve information about the sources in the enterprise.

Each of these operations follow a request/response paradigm. The request is the input to the operation and contains all of the input parameters needed by the Catalog Framework's operation to communicate with the Sources. The response is the output from the execution of the operation that is returned to the client, which contains all of the data returned by the sources. For each operation there is an associated request/response pair, e.g., the `QueryRequest` and `QueryResponse` pair for the Catalog Framework's query operation.

All of the request and response objects are extensible in that they can contain additional key/value properties on each request/response. This allows additional capability to be added without changing the Catalog API, helping to maintain backwards compatibility.


21. Resources



Resources Architecture

Resources are the data that is represented by the cataloged metadata in DDF.

Metacards are used to describe those resources through metadata. This metadata includes the time the resource was created, the location where the resource was created, etc. A DDF Metacard contains the `getResourceUri` method, which is used to locate and retrieve its corresponding resource.



Content Data Component Architecture

21.1. Content Item

ContentItem is the domain object populated by the Storage Provider that represents the information about the content to be stored or content that has been stored in the Storage Provider. A ContentItem encapsulates the content's globally unique ID, mime type, and input stream (i.e., the actual content). The unique ID of a ContentItem will always correspond to a Metocard ID.

21.1.1. Retrieving Resources

When a client attempts to retrieve a resource, it must provide a metocard ID or URI corresponding to a unique resource. As mentioned above, the resource URI is obtained from a Metocard's 'getResourceUri' method. The CatalogFramework has three methods that can be used by clients to obtain a resource: `getEnterpriseResource`, `getResource`, and `getLocalResource`. The `getEnterpriseResource` method invokes the `retrieveResource` method on a local ResourceReader as well as all the Federated and Connected Sources in the DDF enterprise. The second method, `getResource`, takes in a source ID as a parameter and only invokes `retrieveResource` on the specified Source. The third method invokes `retrieveResource` on a local ResourceReader.

The parameter for each of these methods in the CatalogFramework is a ResourceRequest. DDF includes two implementations of ResourceRequest: `ResourceRequestById` and `ResourceRequestByProductUri`. Since these implementations extend `OperationImpl`, they can pass a Map of generic properties through the CatalogFramework to customize how the resource request is carried out. One example of this is explained in the [Retrieving Resource Options](#) section below. The following is a basic example of how to create a ResourceRequest and invoke the CatalogFramework resource retrieval methods to

process the request.

Retrieve Resource Example

```
Map<String, Serializable> properties = new HashMap<String, Serializable>();
properties.put("PropertyKey1", "propertyA"); //properties to customize Resource
retrieval
ResourceRequestById resourceRequest = new ResourceRequestById(
"0123456789abcdef0123456789abcdef", properties); //object containing ID of Resource to
be retrieved
String sourceName = "LOCAL_SOURCE"; //the Source ID or name of the local Catalog or a
Federated Source
ResourceResponse resourceResponse; //object containing the retrieved Resource and the
request that was made to get it.
resourceResponse = catalogFramework.getResource(resourceRequest, sourceName);
//Source-based retrieve Resource request
Resource resource = resourceResponse.getResource(); //actual Resource object
containing InputStream, mime type, and Resource name
```

`DDF.catalog.resource.ResourceReader` instances can be discovered via the OSGi Service Registry. The system can contain multiple `ResourceReaders`. The `CatalogFramework` determines which one to call based on the scheme of the resource's URI and what schemes the `ResourceReader` supports. The supported schemes are obtained by a `ResourceReader`'s `'getSupportedSchemes` method. As an example, one `ResourceReader` may know how to handle file-based URIs with the scheme `file`, whereas another `ResourceReader` may support HTTP-based URIs with the scheme `http`.

The `ResourceReader` or `Source` is responsible for locating the resource, reading its bytes, adding the binary data to a `Resource` implementation, then returning that `Resource` in a `ResourceResponse`. The `ResourceReader` or `Source` is also responsible for determining the `Resource`'s name and mime type, which it sends back in the `'Resource` implementation.

21.1.1.1. BinaryContent

`BinaryContent` is an object used as a container to store translated or transformed DDF components. `Resource` extends `BinaryContent` and includes a `getName` method. `' BinaryContent`` has methods to get the `InputStream`, byte array, MIME type, and size of the represented binary data. An implementation of `BinaryContent` (`BinaryContentImpl`) can be found in the Catalog API in the `DDF.catalog.data` package.

21.1.2. Retrieving Resource Options

Options can be specified on a retrieve resource request made through any of the supporting endpoint. To specify an option for a retrieve resource request, the endpoint needs to first instantiate a `ResourceRequestByProductUri` or a `ResourceRequestById`. Both of these `ResourceRequest` implementations allow a `Map` of properties to be specified. Put the specified option into the `Map` under the key `RESOURCE_OPTION`.

Retrieve Resource with Options

```
Map<String, Serializable> properties = new HashMap<String, Serializable>();
properties.put("RESOURCE_OPTION", "OptionA");
ResourceRequestById resourceRequest = new ResourceRequestById(
"0123456789abcdef0123456789abcdef", properties);
```

Depending on the support that the `ResourceReader` or `Source` provides for options, the `properties`Map` will be checked for the `RESOURCE_OPTION` entry. If that entry is found, the option will be handled. If the `ResourceReader` or `Source` does not support options, that entry will be ignored.

A new `ResourceReader` or `Source` implementation can be created to support options in a way that is most appropriate. Since the option is passed through the catalog framework as a property, the `ResourceReader` or `Source` will have access to that option as long as the endpoint supports options.

21.1.3. Storing Resources

Resources are saved using a `ResourceWriter`. `DDF.catalog.resource.ResourceWriter` instances can be discovered via the OSGi Service Registry. Once retrieved, the `ResourceWriter` instance provides clients with a way to store resources and get a corresponding URI that can be used to subsequently retrieve the resource via a `ResourceReader`. Simply invoke either of the `storeResource` methods with a resource and any potential arguments. The `ResourceWriter` implementation is responsible for determining where the resource is saved and how it is saved. This allows flexibility for a resource to be saved in any one of a variety of data stores or file systems. The following is an example of how to use a generic implementation of `ResourceWriter`.

Using a `ResourceWriter`

```
InputStream inputStream = <Video_Input_Stream>; //InputStream of raw Resource data
MimeType mimeType = new MimeType("video/mpeg"); //Mime Type or content type of
Resource
String name = "Facility_Video"; //Descriptive Resource name
Resource resource = new ResourceImpl(inputStream, mimeType, name);
Map<String, Object> optionalArguments = new HashMap<String, Object>();
ResourceWriter writer = new ResourceWriterImpl();
URI resourceUri; //URI that can be used to retrieve Resource
resourceUri = writer.storeResource(resource, optionalArguments); //Null can be passed
in here
```

21.2. Resource Components

Resource components are used when working with resources

A resource is a URI-addressable entity that is represented by a metocard. Resources may also be known as **products** or **data**.

Resources may exist either locally or on a remote data store.

Examples of resources include:

- NITF image
- MPEG video
- Live video stream
- Audio recording
- Document

A resource object in DDF contains an [InputStream](#) with the binary data of the resource. It describes that resource with a name, which could be a file name, URI, or another identifier. It also contains a mime type or content type that a client can use to interpret the binary data.

21.3. Resource Readers

A resource reader retrieves resources associated with metacards via URIs. Each resource reader must know how to interpret the resource's URI and how to interact with the data store to retrieve the resource.

There can be multiple resource readers in a Catalog instance. The [Catalog Framework](#) selects the appropriate resource reader based on the scheme of the resource's URI.

In order to make a resource reader available to the Catalog Framework, it must be exported to the OSGi Service Registry as a [DDF.catalog.resource.ResourceReader](#).

21.3.1. URL Resource Reader

The [URLResourceReader](#) is an implementation of [ResourceReader](#) which is included in the DDF Catalog. It obtains a resource given an http, https, or file-based URL. The [URLResourceReader](#) will connect to the provided Resource URL and read the resource's bytes into an [InputStream](#).

WARNING

When a resource linked using a file-based URL is in the product cache, the [URLResourceReader](#)'s `rootResourceDirectories` is not checked when downloading the product. It is downloaded from the product cache which bypasses the [URLResourceReader](#). For example, if path `/my/valid/path` is configured in the [URLResourceReader](#)'s `rootResourceDirectories` and one downloads the product with resource-uri `file:///my/valid/path/product.txt` and then one removes `/my/valid/path` from the [URLResourceReader](#)'s `rootResourceDirectories` configuration, the product will still be accessible via the product cache.

21.3.1.1. Installing the URL Resource Reader

The [URLResourceReader](#) is installed by default with a standard installation in the Catalog application.

21.3.1.2. Configuring the URL Resource Reader

Configure the URL Resource Reader from the Admin Console.

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Configuration** tab.
4. Select the **URL Resource Reader**.

See [URL Resource Reader configurations](#) for all possible configurations.

21.3.1.3. Using the URL Resource Reader

`URLResourceReader` will be used by the Catalog Framework to obtain a resource whose metocard is cataloged in the local data store. This particular `ResourceReader` will be chosen by the `CatalogFramework` if the requested resource's URL has a protocol of `http`, `https`, or `file`.

For example, requesting a resource with the following URL will make the Catalog Framework invoke the `URLResourceReader` to retrieve the product.

Example

```
file:///home/users/DDF_user/data/example.txt
```

If a resource was requested with the URL `udp://123.45.67.89:80/SampleResourceStream`, the `URLResourceReader` would *not* be invoked.

Supported Schemes:

- `http`
- `https`
- `file`

NOTE

If a file-based URL is passed to the `URLResourceReader`, that file path needs to be accessible by the DDF instance.

21.3.1.4. Usage Limitations of the URL Resource Reader

None.

21.4. Resource Writers

A resource writer stores a resource and produces a URI that can be used to retrieve the resource at a later time. The resource URI uniquely locates and identifies the resource. Resource writers can interact with an underlying data store and store the resource in the proper place. Each implementation can do this differently, providing flexibility in the data stores used to persist the resources.

Resource Writers should be used within the Content Framework if and when implementing a custom Storage Provider to store the product. The default Storage Provider that comes with the DDF writes the products to the file system.

21.5. Content Directory Monitor

The Content Directory Monitor (CDM) provides the capability to easily add content and metacards into the Catalog by placing a file in a directory.

21.5.1. Installing the Content Directory Monitor

The Content Directory Monitor is installed by default with a standard installation of the Catalog application.

21.5.2. Configuring the Content Directory Monitor

Configure the CDM from the Admin Console:

1. Navigate to the **Admin Console**.
2. Select the **Catalog** application.
3. Select the **Configuration** tab.
4. Select **Catalog Content Directory Monitor**.

See [Content Directory Monitor configurations](#) for all possible configurations.

21.5.3. Using the Content Directory Monitor

The CDM processes files in a directory, and all of its sub-directories. The CDM offers three options:

- Delete
- Move
- Monitor in place

Regardless of the option, the DDF takes each file in a monitored directory structure and creates a metacard for it. The metacard is linked to the file. The behavior of each option is given below.

Delete

- Copies the file into the Content Repository.
- Creates a metacard in the Catalog from the file.
- **Erases** the original file from the monitored directory.

Move

- Copies the file into the directory **.\ingested** (this will double the disk space used)
- Copies the file into the Content Repository.
- Creates a metacard in the Catalog from the file.
- **Erases** the original file from the monitored directory.

Monitor in place

- Creates a metacard in the Catalog from the file.

- Creates a reference from the metocard to the original file in the monitored directory.
- If the original file is deleted, the metocard is removed from the Catalog.
- If the original file is modified, the metocard is updated to reflect the new content.
- If the original file is renamed, the old metocard is deleted and a new metocard is created.

Parallel Processing

The CDM supports parallel processing of files (up to 8 files processed concurrently). This is configured by setting the number of **Maximum Concurrent Files** in the configuration. A maximum of 8 is imposed to protect system resources.

Read Lock

When the CDM is set up, the directory specified is continuously scanned, and files are locked for processing based on the **ReadLock Time Interval**. This does not apply to the **Monitor in place** processing directive. Files will not be ingested without having a ReadLock that has observed no change in the file size. This is done so that files that are in transit will not be ingested prematurely. The interval should be dependent on the speed of the copy to the directory monitor (ex. network drive vs local disk). For local files, the default value of 500 milliseconds is recommended. The recommended interval for network drives is 1000 - 2000 milliseconds. If the value provided is less than 100, 100 milliseconds will be used. It is also recommended that the **ReadLock Time Interval** be set to a lower amount of time when the **Maximum Concurrent Files** is set above 1 so that files are locked in a timely manner and processed as soon as possible. When a higher **ReadLock Time Interval** is set, the time it takes for files to be processed is increased.

Attribute Mapper

The CDM supports setting metocard attributes directly when DDF ingests a file. Custom mappings are entered in the form:

attribute-name=attribute-value

For example, to set the contact email for all metacards, add the custom mapping:

contact.point-of-contact-email=doctor@clinic.com

Each mapping sets the value of a single metocard attribute. To set the value of an additional attribute, select the "plus" icon in the UI. This creates an empty line for the entry.

To set multi-valued attributes, use a separate override for each value. For example, to add the keywords *PPI* and *radiology* to each metocard, add the custom attribute mappings:

topic.keyword=PPI
topic.keyword=radiology

Blacklist

The CDM blacklist uses the "bad.files" and "bad.file.extensions" properties from the system.properties file in "etc/" in order to prevent malicious or unwanted data from being ingested into DDF. While the CDM automatically omits hidden files, this is particularly useful when an operating system automatically generates files that should not be ingested. One such example of this is "thumbs.db" in Windows. This file type and any temporary files are included in the blacklist.

Errors

If the CDM fails to read the file, an error will be logged in the ingest log. If the directory monitor is configured to **Delete** or **Move**, the original file is also moved to the `\.errors` directory.

Other

- Multiple directories can be monitored. Each directory has an independent configuration.
- To support the monitoring in place behavior, DDF indexes the files to track their names and modification timestamps. This enables the Content Directory Monitor to take appropriate action when files are changed or deleted.
- The Content Directory Monitor recursively processes all subdirectories.

22. Queries

Clients use `ddf.catalog.operation.Query` objects to describe which metacards are needed from **Sources**.

Query objects have two major components:

- [Filters](#)
- [Query Options](#)

A Source uses the Filter criteria constraints to find the requested set of metacards within its domain of metacards. The Query Options are used to further restrict the Filter's set of requested metacards.

22.1. Filters

An OGC Filter is a [Open Geospatial Consortium \(OGC\) standard](#) that describes a query expression in terms of Extensible Markup Language (XML) and key-value pairs (KVP). The OGC Filter is used to represent a query to be sent to sources and the Catalog Provider, as well as to represent a Subscription. The OGC Filter provides support for expression processing, such as adding or dividing expressions in a query, but that is not the intended use for DDF.

The Catalog Framework does not use the XML representation of the OGC Filter standard. DDF instead uses the Java implementation provided by [GeoTools](#). GeoTools provides Java equivalent classes for OGC Filter XML elements. GeoTools originally provided the standard Java classes for the OGC Filter Encoding 1.0 under the package name `org.opengis.filter`. The same package name is used today and is currently used by DDF. Java developers do not parse or view the XML representation of a Filter in DDF. Instead, developers use only the Java objects to complete query tasks.

Note that the `ddf.catalog.operation.Query` interface extends the `org.opengis.filter.Filter` interface, which means that a Query object is an OGC Java Filter with Query Options.

```
public interface Query extends Filter
```

22.1.1. FilterBuilder API

To avoid the complexities of working with the Filter interface directly and implementing the DDF Profile of the Filter specification, the Catalog includes an API, primarily in `DDF.filter`, to build Filters using a fluent API.

To use the `FilterBuilder` API, an instance of `DDF.filter.FilterBuilder` should be used via the OSGi registry. Typically, this will be injected via a dependency injection framework. Once an instance of `FilterBuilder` is available, methods can be called to create and combine Filters.

TIP The fluent API is best accessed using an IDE that supports code-completion. For additional details, refer to the [Catalog API Javadoc].

22.1.2. Boolean Operators

Filters use a number of boolean operators.

`FilterBuilder.allOf(Filter ...)`

creates a new Filter that requires all provided Filters are satisfied (Boolean AND), either from a List or Array of Filter instances.

`FilterBuilder.anyOf(Filter ...)`

creates a new Filter that requires at least one of the provided Filters are satisfied (Boolean OR), either from a List or Array of Filter instances.

`FilterBuilder.not(Filter filter)`

creates a new Filter that requires the provided Filter must not match (Boolean NOT).

22.1.3. Attribute

Filters can be based on specific attributes.

`FilterBuilder.attribute(String attributeName)`:: begins a fluent API for creating an Attribute-based Filter, i.e., a Filter that matches on Metacards with Attributes of a particular value.

22.1.4. XPath

Filters can be based on XML attributes.

`FilterBuilder.xpath(String xpath)`:: begins a fluent API for creating an XPath-based Filter, i.e., a Filter that matches on Metacards with Attributes of type XML that match when evaluating a provided XPath selector.

```
FilterBuilder.attribute(attributeName).is().like().text(String contextualSearchPhrase );
FilterBuilder.attribute(attributeName).is().like().caseSensitiveText(String caseSensitiveContextualSearchPhrase);
FilterBuilder.attribute(attributeName).is().like().fuzzyText(String fuzzySearchPhrase );
```

23. Metrics

DDF includes a system of data-collection to enable monitoring system health, user interactions, and overall system performance: **Metrics Collection**.

The Metrics Collection Application collects data for all of the pre-configured metrics in DDF and stores them in custom JMX Management Bean (MBean) attributes. Samples of each metric's data is collected every 60 seconds and stored in the `<DDF_HOME>/data/metrics` directory with each metric stored in its own `.rrd` file. Refer to the Metrics Reporting Application for how the stored metrics data can be viewed.

Do not remove the `<DDF_HOME>/data/metrics` directory or any files in it. If this is done, all existing metrics data will be permanently lost.

WARNING

Also note that if DDF is uninstalled/re-installed that all existing metrics data will be permanently lost.

Types of Metrics Collected

Catalog Metrics

Metrics collected about the catalog status.

Source Metrics

Metrics collected per source.

23.1. Metrics Collection Application

The Metrics Collection Application is responsible for collecting both Catalog and Source metrics.

Use Metrics Collection to collect historical metrics data, such as catalog query metrics, message latency, or individual sources' metrics type of data.

23.1.1. Installing Metrics Collection

The Metrics Collection application is installed by default with a standard installation.

The catalog-level metrics are packaged as the `catalog-core-metricsplugin` feature, and the source-level metrics are packaged as the `catalog-core-sourcemetricsplugin` feature.

23.1.2. Configuring Metrics Collection

No configuration is made for the Metrics Collection application. All metrics collected are either pre-configured in DDF or dynamically created as sources are created or deleted.

23.1.3. Catalog Metrics

Table 58. Catalog Metrics Collected

Metric	JMX MBean Name	MBean Attribute Name	Description
Catalog Exceptions	ddf.metrics.catalog:name=Exceptions	Count	The number of exceptions, of all types, thrown across all catalog queries executed.
Catalog Exceptions Federation	ddf.metrics.catalog:name=Exceptions.Federation	Count	The total number of Federation exceptions thrown across all catalog queries executed.
Catalog Exceptions Source Unavailable	ddf.metrics.catalog:name=Exceptions.SourceUnavailable	Count	The total number of SourceUnavailable exceptions thrown across all catalog queries executed. These exceptions occur when the source being queried is currently not available.
Catalog Exceptions Unsupported Query	ddf.metrics.catalog:name=Exceptions.UnsupportedQuery	Count	Total number of UnsupportedQuery exceptions thrown across all catalog queries executed. These exceptions occur when the query being executed is not supported or is invalid.
Catalog Ingest Created	ddf.metrics.catalog:name=Ingest.Created	Count	The number of catalog entries created in the Metadata Catalog.
Catalog Ingest Deleted	ddf.metrics.catalog:name=Ingest.Deleted	The number of catalog entries deleted from the Metadata Catalog.	Count
Catalog Ingest Updated	ddf.metrics.catalog:name=Ingest.Updated	Count	The number of catalog entries updated in the Metadata Catalog.
Catalog Queries	ddf.metrics.catalog:name=Queries	Count	The number of queries attempted.

Metric	JMX MBean Name	MBean Attribute Name	Description
Catalog Queries Comparison	ddf.metrics.catalog:name=Queries.Comparison	Count	The number of queries attempted that included a string comparison criteria as part of the search criteria, e.g., <code>PropertyIsLike</code> , <code>PropertyIsEqualTo</code> , etc.
Catalog Queries Federated	ddf.metrics.catalog:name=Queries.Federated	Count	The number of federated queries attempted.
Catalog Queries Fuzzy	ddf.metrics.catalog:name=Queries.Fuzzy	Count	The number of queries attempted that included a string comparison criteria with fuzzy searching enabled as part of the search criteria.
Catalog Queries Spatial	ddf.metrics.catalog:name=Queries.Spatial	Count	The number of queries attempted that included a spatial criteria as part of the search criteria.
Catalog Queries Temporal	ddf.metrics.catalog:name=Queries.Temporal	Count	The number of queries attempted that included a temporal criteria as part of the search criteria.
Catalog Queries Total Results	ddf.metrics.catalog:name=Queries.TotalResults	Mean	The average of the total number of results returned from executed queries. This total results data is averaged over the metric's sample rate.
Catalog Queries Xpath	ddf.metrics.catalog:name=Queries.Xpath	Count	The number of queries attempted that included a Xpath criteria as part of the search criteria.
Catalog Resource Retrieval	ddf.metrics.catalog:name=Resource	Count	The number of resources retrieved.
Services Latency	ddf.metrics.services:name=Latency	Mean	The response time (in milliseconds) from receipt of the request at the endpoint until the response is about to be sent to the client from the endpoint. This response time data is averaged over the metric's sample rate.

23.1.4. Source Metrics

Metrics are also collected on a per source basis for each configured [Federated Source](#) and [Catalog Provider](#). When the source is configured, the metrics listed in the table below are automatically created. Metrics are collected for each request (whether enterprise query or a source-specific query). When the source is deleted (or renamed), the associated metrics' MBeans and Collectors are also deleted. However, the RRD file in the `data/metrics` directory containing the collected metrics

remain indefinitely and remain accessible from the **Metrics** tab in the Admin Console.

In the table below, the metric name is based on the Source's ID (indicated by `<sourceId>`).

Table 59. Source Metrics Collected

Metric	JMX MBean Name	MBean Attribute Name	Description
Source <sourceId> Exceptions	<code>ddf.metrics.catalog.source:name=<sourceId>.Exceptions</code>	Count	A count of the total number of exceptions, of all types, thrown from catalog queries executed on this source.
Source <sourceId> Queries	<code>ddf.metrics.catalog.source:name=<sourceId>.Queries</code>	Count	A count of the number of queries attempted on this source.
Source <sourceId> Queries Total Results	<code>ddf.metrics.catalog.source:name=<sourceId>.Queries.TotalResults</code>	Mean	An average of the total number of results returned from executed queries on this source. This total results data is averaged over the metric's sample rate.

For example, if a Federated Source was created with a name of `fs-1`, then the following metrics would be created for it:

- `Source Fs1 Exceptions`
- `Source Fs1 Queries`
- `Source Fs1 Queries Total Results`

If this federated source is then renamed to `fs-1-rename`, the MBeans and Collectors for the `fs-1` metrics are deleted, and new MBeans and Collectors are created with the new names:

- `Source Fs1 Rename Exceptions`
- `Source Fs1 Rename Queries`
- `Source Fs1 Rename Queries Total Results`

Note that the metrics with the previous name remain on the Metrics tab because the data collected while the Source had this name remains valid and thus needs to be accessible. Therefore, it is possible to access metrics data for sources renamed months ago, i.e., until DDF is reinstalled or the metrics data is deleted from the `<DDF_HOME>/data/metrics` directory. Also note that the source metrics' names are modified to remove all non-alphanumeric characters and renamed in camelCase.

23.2. Metrics Reporting Application

The DDF Metrics Reporting Application provides access to historical data in several formats: a graphic, a comma-separated values file, a spreadsheet, a PowerPoint file, XML, and JSON formats for system metrics collected while DDF is running. Aggregate reports (weekly, monthly, and yearly) are also provided where all collected metrics are included in the report. Aggregate reports are

available in Excel and PowerPoint formats.

To use the Metrics Reporting Application:

1. Navigate to the **Admin Console**.
2. Select the **Platform** Application.
3. Select the **Metrics** tab.

With each metric in the list, a set of hyperlinks is displayed under each column. Each column's header is displayed with the available time ranges. The time ranges currently supported are 15 minutes, 1 hour, 1 day, 1 week, 1 month, 3 months, 6 months, and 1 year, measured from the time that the hyperlink is clicked.

All metrics reports are generated by accessing the collected metric data stored in the `<DDF_HOME>/data/metrics` directory. All files in this directory are generated by the JmxCollector using RRD4J, a Round Robin Database for a Java open source product. All files in this directory will have the `.rrd` file extension and are binary files, hence they cannot be opened directly. These files should only be accessed using the Metrics tab's hyperlinks. There is one RRD file per metric being collected. Each RRD file is sized at creation time and will never increase in size as data is collected. One year's worth of metric data requires approximately 1 MB file storage.

Do not remove the `<DDF_HOME>/data/metrics` directory or any files in the directory. If this is done, all existing metrics data will be permanently lost.

WARNING

Also note that if DDF is uninstalled/re-installed, all existing metrics data will be permanently lost.

Hyperlinks are provided for each metric and each format in which data can be displayed. For example, the PNG hyperlink for 15m for the Catalog Queries metric maps to `https://[FQDN]:[PORT]/services/internal/metrics/catalogQueries.png?dateOffset=900`, where the `dateOffset=900` indicates the previous 900 seconds (15 minutes) to be graphed.

Note that the date format will vary according to the regional/locale settings for the server.

All of the metric graphs displayed are in PNG format and are displayed on their own page. The user may use the back button in the browser to return to the Admin Console, or, when selecting the hyperlink for a graph, they can use the right mouse button in the browser to display the graph in a separate browser tab or window, which will keep the Admin Console displayed. The user can also specify custom time ranges by adjusting the URL used to access the metric's graph. The Catalog Queries metric data may also be graphed for a specific time range by specifying the `startDate` and `endDate` query parameters in the URL.

For example, to map the Catalog Queries metric data for March 31, 6:00 am, to April 1, 2013, 11:00 am, (Arizona timezone, which is -07:00) the URL would be:

```
\https://[FQDN]:[PORT]/services/internal/metrics/catalogQueries.png?startDate=2013-03-31T06:00:00-07:00&endDate=2013-04-01T11:00:00-07:00
```

Or to view the last 30 minutes of data for the Catalog Queries metric, a custom URL with a `dateOffset=1800` (30 minutes in seconds) could be used:

```
\https://{{FQDN}}:{{PORT}}/services/internal/metrics/catalogQueries.png?dateOffset=1800
```

23.2.1. Metrics Aggregate Reports

The Metrics tab also provides aggregate reports for the collected metrics. These are reports that include data for all of the collected metrics for the specified time range.

The aggregate reports provided are:

- Weekly reports for each week up to the past four **complete** weeks from current time. A complete week is defined as a week from Monday through Sunday. For example, if current time is Thursday, April 11, 2013, the past complete week would be from April 1 through April 7.
- Monthly reports for each month up to the past 12 **complete** months from current time. A complete month is defined as the full month(s) preceding current time. For example, if current time is Thursday, April 11, 2013, the past complete 12 months would be from April 2012 through March 2013.
- Yearly reports for the past **complete** year from current time. A complete year is defined as the full year preceding current time. For example, if current time is Thursday, April 11, 2013, the past complete year would be 2012.

An aggregate report in XLS format would consist of a single workbook (spreadsheet) with multiple worksheets in it, where a separate worksheet exists for each collected metric's data. Each worksheet would display:

- the metric's name and the time range of the collected data,
- two columns: Timestamp and Value, for each sample of the metric's data that was collected during the time range, and
- a total count (if applicable) at the bottom of the worksheet.

An aggregate report in PPT format would consist of a single slideshow with a separate slide for each collected metric's data. Each slide would display:

- a title with the metric's name.
- the PNG graph for the metric's collected data during the time range.
- a total count (if applicable) at the bottom of the slide.

Hyperlinks are provided for each aggregate report's time range in the supported display formats, which include Excel (XLS) and PowerPoint (PPT). Aggregate reports for custom time ranges can also be accessed directly via the URL:

```
\https://{{FQDN}}:{{PORT}}/services/internal/metrics/report.<format>?startDate=<start_date_value>&endDate=<end_date_value>
```

where `<format>` is either `xls` or `ppt` and the `<start_date_value>` and `<end_date_value>` specify the custom time range for the report.

These example reports represent custom aggregate reports. NOTE: all example URLs begin with "`https://{FQDN}:{PORT}`", which is omitted in the table for brevity.

Table 60. Example Aggregate Reports

Description	URL
XLS aggregate report for March 15, 2013 to April 15, 2013	<code>/services/internal/metrics/report.xls?startDate=2013-03-15T12:00:00-07:00&endDate=2013-04-15T12:00:00-07:00</code>
XLS aggregate report for last 8 hours	<code>/services/internal/metrics/report.xls?dateOffset=28800</code>
PPT aggregate report for March 15, 2013 to April 15, 2013	<code>/services/internal/metrics/report.ppt?startDate=2013-03-15T12:00:00-07:00&endDate=2013-04-15T12:00:00-07:00</code>
PPT aggregate report for last 8 hours	<code>/services/internal/metrics/report.ppt?dateOffset=28800</code>

23.2.2. Viewing Metrics

The Metrics Viewer has reports in various formats.

1. Navigate to the **Admin Console**.
2. Select the **Platform** application.
3. Select the **Metrics** tab.

Reports are organized by timeframe and output format.

Standard time increments: * `15m`: 15 Minutes * `1h`: 1 Hour * `1d`: 1 Day * `1w`: 1 Week * `1M`: 1 Month * `3M`: 3 Month * `6M`: 6 Month * `1y`: 1 Year

Custom timeframes are also available via the selectors at the bottom of the page.

Output formats: * `PNG` * `CSV` (Comma-separated values) * `XLS`

NOTE

Based on the browser's configuration, either the `.xls` file will be downloaded or automatically displayed in Excel.

24. Action Framework

The Action Framework was designed as a way to limit dependencies between applications (apps) in a system. For instance, a feature in an app, such as an Atom feed generator, might want to include an external link as part of its feed's entries. That feature does not have to be coupled to a REST endpoint to work, nor does it have to depend on a specific implementation to get a link. In reality, the feature does not identify how the link is generated, but it does identify whether the link works or does not work when retrieving the intended entry's metadata. Instead of creating its own mechanism or adding an unrelated feature, it could use the Action Framework to query the OSGi

container for any service that can provide a link. This does two things: it allows the feature to be independent of implementations, and it encourages reuse of common services.

The Action Framework consists of two major Java interfaces in its API:

1. `ddf.action.Action`
2. `ddf.action.ActionProvider`

Actions

Specific tasks that can be performed as services.

Action Providers

Lists of related actions that a service is capable of performing.

24.1. Action Providers

Included Action Providers

Download Resource ActionProvider

Downloads a resource to the local product cache.

IdP Logout Action Provider

Identity Provider Logout.

Karaf Logout Action

Local Logout.

LDAP Logout Action

Ldap Logout.

Overlay ActionProvider

Provides a metocard URL that transforms the metocard into a geographically aligned image (suitable for overlaying on a map).

View Metocard ActionProvider

Provides a URL to a metocard.

Metocard Transformer ActionProvider

Provides a URL to a metocard that has been transformed into a specified format.

25. Asynchronous Processing Framework

NOTE


This code is experimental. While this interface is functional and tested, it may change or be removed in a future version of the library.

The **Asynchronous Processing Framework** is a way to run plugins asynchronously. Generally,

plugins that take a significant amount of processing time and whose results are not immediately required are good candidates for being asynchronously processed. A **Processing Framework** can either be run on the local or remote system. Once the **Processing Framework** finishes processing incoming requests, it may submit ([Create|Update|Delete](#))Requests to the Catalog. The type of plugins that a **Processing Framework** runs are the **Post-Process Plugins**. The **Post-Process Plugins** are triggered by the **Processing Post Ingest Plugin**, which is a **Post-Ingest Plugin**. **Post-Ingest Plugins** are run after the metocard has been ingested into the Catalog. This feature is uninstalled by default.

WARNING


The **Processing Framework** does not support partial updates to the Catalog. This means that if any changes are made to a metocard in the Catalog between the time asynchronous processing starts and ends, those changes will be overwritten by the **ProcessingFramework** updates sent back to the Catalog. This feature should be used with caution.



Processing Framework Architecture

The Asynchronous Processing Framework API Interfaces

1. [org.codice.ddf.catalog.async.processingframework.api.internal.ProcessingFramework](#)
2. [org.codice.ddf.catalog.async.plugin.api.internal.PostProcessPlugin](#)
3. [org.codice.ddf.catalog.async.data.api.internal.ProcessItem](#)
4. [org.codice.ddf.catalog.async.data.api.internal.ProcessCreateItem](#)
5. [org.codice.ddf.catalog.async.data.api.internal.ProcessUpdateItem](#)
6. [org.codice.ddf.catalog.async.data.api.internal.ProcessDeleteItem](#)
7. [org.codice.ddf.catalog.async.data.api.internal.ProcessRequest](#)
8. [org.codice.ddf.catalog.async.data.api.internal.ProcessResource](#)
9. [org.codice.ddf.catalog.async.data.api.internal.ProcessResourceItem](#)



Processing Framework Interface Diagram

ProcessingFramework

The `ProcessingFramework` is responsible for processing incoming `ProcessRequests` that contain a `ProcessItem`. A `ProcessingFramework` should never block. It receives its `ProcessRequests` from a `PostIngestPlugin` on all CUD operations to the Catalog. In order to determine whether or not asynchronous processing is required by the `ProcessingFramework`, the `ProcessingFramework` should mark any request it has submitted back the Catalog, otherwise a processing loop may occur. For example, the default **In-Memory Processing Framework** adds a `POST_PROCESS_COMPLETE` flag to the Catalog CUD request after processing. This flag is checked by the `ProcessingPostIngestPlugin` before a `ProcessRequest` is sent to the `ProcessingFramework`. For an example of a `ProcessingFramework`, please refer to the `org.codice.ddf.catalog.async.processingframework.impl.InMemoryProcessingFramework`.

ProcessRequest

A `ProcessRequest` contains a list of `ProcessItems` for the `ProcessingFramework` to process. Once a `ProcessRequest` has been processed by a `ProcessingFramework`, the `ProcessingFramework` should mark the `ProcessRequest` as already been processed, so that it does not process it again.

PostProcessPlugin

The `PostProcessPlugin` is a plugin that will be run by the `ProcessingFramework`. It is capable of processing `ProcessCreateItems`, `ProcessUpdateItems`, and `ProcessDeleteItem`s.

ProcessItem

WARNING

Do not implement `ProcessItem` directly; it is intended for use only as a common base interface for `ProcessResourceItem` and `ProcessDeleteItem`.

The `ProcessItem` is contained by a `ProcessRequest`. It can be either a `ProcessCreateItem`, `ProcessUpdateItem`, or `ProcessDeleteItem`.

ProcessResource

The `ProcessResource` is a piece of content that is attached to a metocard. The piece of content can be either local or remote.

ProcessResourceItem

The `ProcessResourceItem` indicates that the item being processed may have a `ProcessResource` associated with it.

ProcessResourceItem Warning

WARNING Do not implement `ProcessResourceItem` directly; it is intended for use only as a common base interface for `ProcessCreateItem` and `ProcessUpdateItem`.

ProcessCreateItem

The `ProcessCreateItem` is an item for a metocard that has been created in the Catalog. It contains the created metocard and, optionally, a `ProcessResource`.

ProcessUpdateItem

The `ProcessUpdateItem` is an item for a metocard that has been updated in the Catalog. It contains the original metocard, the updated metocard and, optionally, a `ProcessResource`.

ProcessDeleteItem

The `ProcessDeleteItem` is an item for a metocard that has been deleted in the Catalog. It contains the deleted metocard.

26. Migration API

NOTE

This code is experimental. While the interfaces and classes provided are functional and tested, they may change or be removed in a future version of the library.

DDF currently has an experimental API for making bundles migratable. Interfaces and classes in `platform/migration/platform-migratable-api` are used by the system to identify bundles that provide implementations for export and import operations.

The migration API provides a mechanism for bundles to handle exporting data required to clone or backup/restore a DDF system. The migration process is meant to be flexible, so an implementation of `org.codice.ddf.migration.Migratable` can handle exporting data for a single bundle or groups of bundles such as applications. For example, the `org.codice.ddf.platform.migratable.impl.PlatformMigratable` handles exporting core system files for the Platform application. Each migratable must provide a unique identifier via its `getId()` method used by the migration API to uniquely identify the migratable between exports and imports.

DDF defines migratables of its own to export/import all configurations stored in

`org.osgi.service.cm.ConfigurationAdmin`.

These do not need to be handled by implementations of `org.codice.ddf.migration.Migratable`.

An export and an import operation can be performed through the Command Console.

When an export operation is processed, the migration API will do a look-up for all registered OSGi services that are implementing `Migratable` and call their `doExport()` method. As part of the exported data, information about the migratable as required by the `org.codice.ddf.platform.services.common.Describable` interface will be included. In particular the version string returned will help the migration API identify the version of the exported data from the corresponding migratable and must be provided as a non-blank string.

When an import operation is processed, the migration API will do another look-up for all registered OSGi services that are implementing `Migratable` and call their `doImport()` or `doIncompatibleImport()` methods based on whether the recorded version string at export time is equal to the version string currently provided by the migratable or not. The `doMissingImport()` method will be called instead of one of the other two methods when the migration API detects that the corresponding migratable data is missing from the exported data.

The services that implement the migratable interface will be called one at a time based on their service ranking order, and do not need to be thread safe. A bundle or a feature can have as many services implementing the interfaces as needed.

26.1. The Migration API Interfaces and Classes

1. `org.codice.ddf.migration.Migratable`
2. `org.codice.ddf.migration.MigrationContext`
3. `org.codice.ddf.migration.ExportMigrationContext`
4. `org.codice.ddf.migration.ImportMigrationContext`
5. `org.codice.ddf.migration.MigrationEntry`
6. `org.codice.ddf.migration.ExportMigrationEntry`
7. `org.codice.ddf.migration.ImportMigrationEntry`
8. `org.codice.ddf.migration.MigrationOperation`
9. `org.codice.ddf.migration.MigrationReport`
10. `org.codice.ddf.migration.MigrationMessage`
11. `org.codice.ddf.migration.MigrationException`
12. `org.codice.ddf.migration.MigrationWarning`
13. `org.codice.ddf.migration.MigrationInformation`
14. `org.codice.ddf.migration.MigrationSuccessfulInformation`

26.1.1. Migratable

The contract for a migratable is stored here. This is the only interface that should be implemented by implementers and registered as an OSGi service. All other interfaces will be implemented by the migration API that provides support for migratables.

The `org.codice.ddf.migration.Migratable` interface defines these methods:

- `String getId()`
- `String getVersion()`
- `String getTitle()`
- `String getDescription()`
- `String getOrganization()`
- `void doExport(ExportMigrationContext context)`
- `void doImport(ImportMigrationContext context)`
- `void doIncompatibleImport(ImportMigrationContext context)`
- `void doMissingImport(ImportMigrationContext context)`

The `getId()` method returns a unique identifier for this migratable that must remain constant between the export and the import operations in order for the migration API to correlate the exported data with the migratable during the import operation. It must be unique across all migratables.

The `getVersion()` method returns a unique version string which is meant to identify the version of the data exported or supported at import time by the migratable. It cannot be blank and its format is left to the migratable. The only noticeable requirement is that when the string compares equal using the `String.equals()` method, the migration API will call `doImport()` instead of `doIncompatibleImport()` to restore previously exported data for the migratable.

The `getTitle()` method returns a simple title for the migratable.

The `getDescription()` method returns a short description of the type of data exported by the migratable.

The `getOrganization()` method provides the name of the organization responsible for the migratable.

The `doExport()` method is called by the migration API along with a context for the current export operation to store data.

The `doImport()` method is called by the migration API along with a context for the current import operation when the version of exported data matches the current version reported by the migratable. This method can be used to restore previously exported data.

The `doIncompatibleImport()` method is called to restore incompatible data which might require transformation. It is provided a context for the current import operation and the previously exported version. It can then proceed with restoring incompatible data which might require transformation.

Finally, the `doMissingImport()` method will be called along with the context for the current import operation when data had not been exported for the corresponding migratable. This will be the case when a migratable is later introduced in the software distribution.

In order to create a `Migratable` for a module of the system, the `org.codice.ddf.migration.Migratable` interface must be implemented and the implementation must be registered under the `org.codice.ddf.migration.Migratable` interface as an OSGi service in the OSGi service registry. Creating an OSGi service allows for the migration API to lookup all implementations of `org.codice.ddf.migration.Migratable` and command them to export or import.

26.1.2. MigrationContext

The `org.codice.ddf.migration.MigrationContext` provides contextual information about an operation in progress for a given migratable. This is a sort of sandbox that is unique to each migratable. This interface defines the following methods:

- `MigrationReport getReport()`
- `String getId()`

The `getReport()` method returns a migration report that can be used to record messages while processing an export or an import operation.

The `getId()` method returns the identifier for the currently processing migratable.

26.1.3. ExportMigrationContext

The export migration context provides methods for creating new migration entries and system property referenced migration entries to track exported migration files for a given migratable while processing an export migration operation. It defines the following methods:

- `Optional<ExportMigrationEntry> getSystemPropertyReferencedEntry(String name)`
- `Optional<ExportMigrationEntry> getSystemPropertyReferencedEntry(String name, BiPredicate<MigrationReport, String> validator)`
- `ExportMigrationEntry getEntry(Path path)`
- `Stream<ExportMigrationEntry> entries(Path path)`
- `Stream<ExportMigrationEntry> entries(Path path, PathMatcher filter)`

The `getSystemPropertyReferencedEntry()` methods create a migration entry to track a file referenced by a given system property value.

The `getEntry()` method creates a migration entry given the path for a specific file or directory.

The `entries()` methods create multiple entries corresponding to all files recursively located underneath a given path with an optional path matcher to filter which files to create entries for.

Once an entry is created, it is not stored with the exported data. It is the migratable's responsibility to store the data using one of the entry's provided methods. Entries are uniquely identified using a relative path and are specific to each migratable meaning that an entry with the same path in two migratables will not conflict with each other. Each migratable is given its own context (a.k.a. sandbox) to work with.

26.1.4. ImportMigrationContext

The import migration context provides methods for retrieving migration entries and system property referenced migration entries corresponding to exported files for a given migratable while processing an import migration operation. It defines the following methods:

- `Optional<ImportMigrationEntry> getSystemPropertyReferencedEntry(String name)`
- `ImportMigrationEntry getEntry(Path path)`
- `Stream<ImportMigrationEntry> entries(Path path)`
- `Stream<ImportMigrationEntry> entries(Path path, PathMatcher filter)`

The `getSystemPropertyReferencedEntry()` method retrieves a migration entry for a file that was referenced by a given system property value.

The `getEntry()` method retrieves a migration entry given the path for a specific file or directory.

The `entries()` methods retrieve multiple entries corresponding to all exported files recursively located underneath a given relative path with an optional path matcher to filter which files to retrieve entries for.

Once an entry is retrieved, its exported data is not restored. It is the migratable's responsibility to restore the data using one of the entry's provided methods. Entries are uniquely identified using a relative path and are specific to each migratable meaning that an entry with the same path in two migratables will not conflict with each other. Each migratable is given its own context (a.k.a. sandbox) to work with.

26.1.5. MigrationEntry

This interface provides supports for exported files or directories. It defines the following methods:

- `MigrationReport getReport()`
- `String getId()`
- `String getName()`
- `Path getPath()`
- `boolean isDirectory()`
- `boolean isFile()`
- `long getLastModifiedTime()`

The `getReport()` method provides access to the associated migration report where messages can be recorded.

The `getId()` method returns the identifier for the migratable responsible for this entry.

The `getName()` method provides the unique name for this entry in an OS-independent way.

The `getPath()` method provides the unique path to the corresponding file for this entry in an OS-specific way.

The `isDirectory()` method indicates if the entry represents a directory. The `isFile()` method indicates if the entry represents a file. The `getLastModifiedTime()` method provides the last modification time for the corresponding file or directory as available when the file or directory is exported.

26.1.6. ExportMigrationEntry

The export migration entry provides additional methods available for entries created at export time. It defines the following methods:

- `Optional<ExportMigrationEntry> getPropertyReferencedEntry(String name)`
- `Optional<ExportMigrationEntry> getPropertyReferencedEntry(String name, BiPredicate<MigrationReport, String> validator)`
- `boolean store()`
- `boolean store(boolean required)`
- `boolean store(PathMatcher filter)`
- `boolean store(boolean required, PathMatcher filter)`

- `boolean store(BiThrowingConsumer<MigrationReport, OutputStream, IOException> consumer)`
- `OutputStream getOutputStream() throws IOException`

The `getPropertyReferencedEntry()` methods create another migration entry for a file that was referenced by a given property value in the file represented by this entry.

The `store()` and `store(boolean required)` methods will automatically copy the content of the corresponding file as part of the export making sure the file exists (if required) on disk otherwise an error will be recorded. If the path represents a directory then all files recursively found under the path will be automatically exported.

The `store(PathMatcher filter)` and `store(boolean required, PathMatcher filter)` methods will automatically copy the content of the corresponding file if it matches the filter as part of the export making sure the file exists (if required) on disk otherwise an error will be recorded. If the path represents a directory then all matching files recursively found under the path will be automatically exported.

The `store(BiThrowingConsumer<MigrationReport, OutputStream, IOException> consumer)` method allows the migratable to control the export process by specifying a callback consumer that will be called back with an output stream where the data can be written to instead of having a file on disk being copied by the migration API. The `OutputStream getOutputStream()` method provides access to the low-level output stream where the migratable can write data directly as opposed to having a file on disk copied automatically.

26.1.7. ImportMigrationEntry

The import migration entry provides additional methods available for entries retrieved at import time. It defines the following methods:

- `Optional<ImportMigrationEntry> getPropertyReferencedEntry(String name)`
- `boolean restore()`
- `boolean restore(boolean required)`
- `boolean restore(PathMatcher filter)`
- `boolean restore(boolean required, PathMatcher filter)`
- `boolean restore(BiThrowingConsumer<MigrationReport, Optional<InputStream>, IOException> consumer)`
- `Optional<InputStream> getInputStream() throws IOException`

The `getPropertyReferencedEntry()` method retrieves another migration entry for a file that was referenced by a given property value in the file represented by this entry.

The `restore()` and `restore(boolean required)` methods will automatically copy the exported content of the corresponding file back to disk if it was exported; otherwise an error will be recorded. If the path represents a directory then all file entries originally recursively exported under this entry's path will be automatically imported. If the directory had been completely exported using one of the `store()` or `store(boolean required)` methods then in addition to restoring all entries recursively, calling this method will also remove any existing files or directories that were not on the original system.

The `restore(PathMatcher filter)` and `restore(boolean required, PathMatcher filter)` methods will automatically copy the exported content of the corresponding file if it matches the filter back to disk if it was exported; otherwise an error will be recorded. If the path represents a directory then all matching file entries originally recursively exported under this entry's path will be

automatically imported.

The `restore(BiThrowingConsumer<MigrationReport, Optional<InputStream>, IOException> consumer)` method allows the migratable to control the import process by specifying a callback consumer that will be called back with an optional input stream (empty if the data was not exported) where the data can be read from instead of having a file on disk being created or updated by the migration API.

The `Optional<InputStream> getInputStream()` method provides access to the optional low-level input stream (empty if the data was not exported) where the migratable can read data directly as opposed to having a file on disk created or updated automatically.

26.1.8. MigrationOperation

The `org.codice.ddf.migration.MigrationOperation` provides a simple enumeration for identifying the various migration operations available.

26.1.9. MigrationReport

The `org.codice.ddf.migration.MigrationReport` interface provides information about the execution of a migration operation. It defines the following methods:

- `MigrationOperation getOperation()`
- `Instant getStartTime()`
- `Optional<Instant> getEndTime()`
- `MigrationReport record(String msg)`
- `MigrationReport record(String format, @Nullable Object... args)`
- `MigrationReport record(MigrationMessage msg)`
- `MigrationReport doAfterCompletion(Consumer<MigrationReport> code)`
- `Stream<MigrationMessage> messages()`
- `default Stream<MigrationException> errors()`
- `Stream<MigrationWarning> warnings()`
- `Stream<MigrationInformation> infos()`
- `boolean wasSuccessful()`
- `boolean wasSuccessful(@Nullable Runnable code)`
- `boolean wasIOSuccessful(@Nullable ThrowingRunnable<IOException> code) throws IOException`
- `boolean hasInfos()`
- `boolean hasWarnings()`
- `boolean hasErrors()`
- `void verifyCompletion()`

The `getOperation()` method provides the type of migration operation (i.e. export or import) currently in progress.

The `getStartTime()` method provides the time at which the corresponding operation started.

The `getEndTime()` method provides the optional time at which the corresponding operation ended. The time is only available if the operation has ended.

The `record()` methods enable messages to be recorded with the report. Messages are displayed on the console for the administrator.

The `doAfterCompletion()` methods enable code to be registered such that it is invoked at the end before a successful result is returned. Such code can still affect the result of the operation.

The `messages()` method provides access to all recorded messages so far.

The `errors()` method provides access to all recorded error messages so far.

The `warnings()` method provides access to all recorded warning messages so far.

The `infos()` method provides access to all recorded informational messages so far.

The `wasSuccessful()` method provides a quick check to see if the report is successful. A successful report might have warnings recorded but cannot have errors recorded.

The `wasSuccessful(Runnable code)` method allows code to be executed. It will return true if no new errors are recorded as a result of executing the provided code.

The `'wasIOSuccessful(ThrowingRunnable<IOException> code)` method allows code to be executed which can throw I/O exceptions which are automatically recorded as errors. It will return true if no new errors are recorded as a result of executing the provided code.

The `'hasInfos()` method will return true if at least one information message has been recorded so far.

The `hasWarnings()` method will return true if at least one warning message has been recorded so far.

The `hasErrors()` method will return true if at least one error message has been recorded so far.

The `'verifyCompletion()` method will verify if the report is successful and if not, it will throw back the first recorded exception and attach as suppressed exceptions all other recorded exceptions.

26.1.10. MigrationMessage

The `'org.codice.ddf.migration.MigrationException` is defined as a base class for all recordable messages during migration operations. It defines the following methods:

- `String getMessage()`

The `getMessage()` method provides a message for the corresponding exception, warning, or info that will be displayed to the administrator on the console.

26.1.11. MigrationException

An `org.codice.ddf.migration.MigrationException` should be thrown when an unrecoverable exception occurs that prevents the export or the import operation from continuing. It is also possible to simply record one or many exception(s) with the migration report in order to fail the export or import operation while not aborting it right away. This provides for the ability to record as many errors as possible and report all of them back to the administrator. All migration exception messages are displayed to the administrator.

26.1.12. MigrationWarning

An `org.codice.ddf.migration.MigrationWarning` should be used when a migratable wants to warn the administrator that certain aspects of the export or the import may cause problems. For example, if an absolute path is encountered, that path may not exist on the target system and cause the installation to fail. All migration warning messages are displayed to the administrator.

26.1.13. MigrationInformation

An `org.codice.ddf.migration.MigrationInformation` should be used when a migratable simply wants to provide useful information to the administrator. All migration information messages are displayed to the administrator.

26.1.14. MigrationSuccessfulInformation

The `org.codice.ddf.migration.MigrationSuccessfulInformation` can be used to further qualify an information message as representing the success of an operation.

27. Application Service API

The Application service has multiple interfaces which are exposed on top of the OSGi runtime for other applications to use. For more information on these interfaces, see [Application Service Interfaces](#).

27.1. JMX Managed Bean

Some of the Application service API is exposed via JMX. It can either be accessed using the JMX API or from a REST-based interface created by [Jolokia](#) that comes with DDF. Here are the interfaces that are exposed in the Managed Bean:

getApplicationTree

Creates an application hierarchy tree that shows relationships between applications.

startApplication

Starts an application with the given name.

stopApplication

Stops an application with the given name.

addApplications

Adds a list of applications that are specified by their URL.

27.2. Initial Application Installation

In DDF, an **application** is a collection of one or more **features**.

WARNING

This application list configuration file is only read the first time the server is started after being unzipped.

To minimize the chance of accidentally installing and uninstalling applications, the configuration file for installing the initial applications is only read the first time that DDF is started. The only way to change what applications are active on startup is to use the console commands. Operations can

also be done with the Admin Console that comes with DDF using the **Features** tab and installing the main feature for the desired application.

The application list file is located at `<DDF_HOME>/etc/org.codice.ddf.admin.applicationlist.properties`

Applications should be defined in a `<name>=<location>` syntax where location may be empty for applications that have already been added to DDF or were prepackaged with the distribution.

Examples:

```
# Local application:  
opendj-embedded  
  
# Application installed into a local maven repository:  
opendj-embedded=mvn:org.codice.opendj.embedded/opendj-embedded-app/1.0.1-  
SNAPSHOT/xml/features  
  
# Application located on the file system:  
opendj-embedded=file:/location/to/opendj-embedded-app-1.0.1-SNAPSHOT.kar
```

Applications will be started in the order they are listed in the file. If an application is listed, DDF will also attempt to install all dependencies for that application.

27.3. Application Service Interfaces

The interfaces used by that application service subsystem are `ApplicationService`, `Application`, `ApplicationStatus`, and `ApplicationNode`. The purpose of these interfaces and their public methods are explained below.

27.3.1. ApplicationService Interface

The `ApplicationService` interface is the main class that is used to operate on applications.

`getApplications()`

Returns a set of all applications that are installed on the system. Callers can then use the `Application` handle to get the name and any underlying features and bundles that this application contains.

`getApplication(String applicationName)`

Returns the application that has the given name.

`startApplication(Application application)`

Starts an application, including any defined dependencies in the application.

`stopApplication(Application application)`

Stops an application, does not include any external transitive dependencies as they may be needed by other applications.

`addApplication(URI applicationURL)`

Adds a new application to the application list. **NOTE: This does NOT start the application.**

`removeApplication(URI applicationURL)`

Removes an application that has the given URI.

`isApplicationStarted(Application application)`

Returns a boolean value relating whether a given application is started or not. This method is generally called after retrieving a list of applications in the first method.

`getApplicationStatus(Application application)`

Returns the full status of an application. This status contains detailed information about the health of the application and is described in the [ApplicationStatus](#) interface section.

`getApplicationTree()`

Creates a hierarchy tree of application nodes that show the relationship between applications.

`findFeature(Feature feature)`

Determine which application contains a certain feature.

27.3.2. Application Interface

`getName()`

Returns the name of the application, which should be unique among applications.

`getFeatures()`

Returns all of the features that this application contains regardless if they are required.

`getBundles()`

Returns all of the bundles that are defined by the features and included in this application.

27.3.3. ApplicationStatus Interface

`getApplication`

Returns the application that is associated with this status.

`getState`

Returns the application's state as defined by `ApplicationState`.

`getErrorFeatures`

Returns a set of Features that were required for this application but did not start correctly.

`getErrorBundles`

Returns a set of Bundles that were required for this application but did not start correctly.

27.3.4. ApplicationNode Interface

`getApplication()`

Returns the application for a node reference.

`getStatus()`

Returns the status for the referenced application.

`getParent()`

Returns the parent of the application.

`getChildren()`

Returns the children of this application. That is, the applications that depend on this application.

28. Security Framework

The DDF Security Framework utilizes [Apache Shiro](#) as the underlying security framework. The classes mentioned in this section will have their full package name listed, to make it easy to tell which classes come with the core Shiro framework and which are added by DDF.

28.1. Subject

`ddf.security.Subject <extends> org.apache.shiro.subject.Subject`

The Subject is the key object in the security framework. Most of the workflow and implementations revolve around creating and using a Subject. The Subject object in DDF is a class that encapsulates all information about the user performing the current operation. The Subject can also be used to perform permission checks to see if the calling user has acceptable permission to perform a certain action (e.g., calling a service or returning a metocard). This class was made DDF-specific because the Shiro interface cannot be added to the Query Request property map.

Table 61. Implementations of Subject:

Classname	Description
<code>ddf.security.impl.SubjectImpl</code>	Extends <code>org.apache.shiro.subject.support.DelegatingSubject</code>

28.1.1. Security Manager

`ddf.security.service.SecurityManager`

The Security Manager is a service that handles the creation of Subject objects. A proxy to this service should be obtained by an endpoint to create a Subject and add it to the outgoing [QueryRequest](#). The Shiro framework relies on creating the subject by obtaining it from the current thread. Due to the multi-threaded and stateless nature of the DDF framework, utilizing the Security Manager interface makes retrieving Subjects easier and safer.

Table 62. Implementations of Security Managers:

Classname	Description
<code>ddf.security.service.SecurityManagerImpl</code>	This implementation of the Security Manager handles taking in both <code>org.apache.shiro.authc.AuthenticationToken</code> and <code>org.apache.cxf.ws.security.tokenstore.SecurityToken</code> objects.

28.1.2. Realms

DDF uses [Apache Shiro](#) for the concept of [Realms](#) for Authentication and Authorization. Realms are components that access security data such as users or permissions.

28.1.2.1. Authenticating Realms

[org.apache.shiro.realm.AuthenticatingRealm](#)

Authenticating Realms are used to authenticate an incoming authentication token and create a [Subject](#) on successful authentication. A Subject is an application user and all available security-relevant information about that user.

Table 63. Implementations of Authenticating Realms in DDF:

Classname	Description
ddf.security.realm.sts.StsRealm	This realm delegates authentication to the Secure Token Service (STS). It creates a RequestSecurityToken message from the incoming Authentication Token and converts a successful STS response into a Subject.

28.1.2.2. Authorizing Realms

[org.apache.shiro.realm.AuthorizingRealm](#)

Authorizing Realms are used to perform authorization on the current Subject. These are used when performing both service authorization and filtering. They are passed in the [AuthorizationInfo](#) of the Subject along with the permissions of the object wanting to be accessed. The response from these realms is a true (if the Subject has permission to access) or false (if the Subject does not).

Table 64. Other implementations of the Security API within DDF

Classname	Description
org.codice.ddf.platform.filter.delegate.DelegateServletFilter	The DelegateServletFilter detects any servlet filters that have been exposed as OSGi services implementing org.codice.ddf.platform.filter.SecurityFilter and places them in-order in front of any servlet or web application running on the container.
org.codice.ddf.security.filter.websso.WebSSOFilter	This filter is the main security filter that works with a number of handlers to protect a variety of web contexts, each using different authentication schemes and policies.
org.codice.ddf.security.handler.saml.SAMLAssertionHandler	This handler is executed by the WebSSOFilter for any contexts configured to use it. This handler should always come first when configured in the Web Context Policy Manager, as it provides a caching capability to web contexts that use it. The handler will first check for the existence of an HTTP Authorization header of type SAML, whose value is a Base64 + deflate SAML assertion. If that is not found, then the handler will check for the existence of the deprecated org.codice.websso.saml.token cookie with the same value. Failing that, it will check for a JSESSIONID cookie to use as a reference to a cached assertion. If the JSESSIONID is valid, the SecurityToken will be retrieved from the cache.

Classname	Description
<code>org.codice.ddf.security.handler.basic.BasicAuthenticationHandler</code>	Checks for basic authentication credentials in the http request header. If they exist, they are retrieved and passed to the <code>LoginFilter</code> for exchange.
<code>org.codice.ddf.security.handler.pki.PKIHandler</code>	Handler for PKI based authentication. X509 chain will be extracted from the HTTP request and converted to a <code>BinarySecurityToken</code> .
<code>org.codice.ddf.security.handler.guest.GuestHandler</code>	Handler that allows guest user access via a guest user account. The guest account credentials are configured via the <code>org.codice.ddf.security.claims.guest.GuestClaimsHandler</code> . The <code>GuestHandler</code> also checks for the existence of basic auth credentials or PKI credentials that might be able to override the use of the guest user.
<code>org.codice.ddf.security.filter.login.LoginFilter</code>	This filter runs immediately after the <code>WebSSOFilter</code> and exchanges any authentication information found in the request with a <code>Subject</code> via Shiro.
<code>org.codice.ddf.security.filter.authorization.AuthorizationFilter</code>	This filter runs immediately after the <code>LoginFilter</code> and checks any permissions assigned to the web context against the attributes of the user via Shiro.
<code>org.apache.shiro.realm.AuthenticatingRealm</code>	This is an abstract authenticating realm that exchanges an <code>org.apache.shiro.authc.AuthenticationToken</code> for a <code>ddf.security.Subject</code> in the form of an <code>org.apache.shiro.authc.AuthenticationInfo</code>
<code>ddf.security.realm.sts.StsRealm</code>	This realm is an implementation of <code>org.apache.shiro.realm.AuthenticatingRealm</code> and connects to an STS (configurable) to exchange the authentication token for a <code>Subject</code> .
<code>ddf.security.service.AbstractAuthorizingRealm</code>	This is an abstract authorizing realm that takes care of caching and parsing the <code>Subject</code> 's <code>AuthorizingInfo</code> and should be extended to allow the implementing realm to focus on making the decision.
<code>ddf.security.pdp.realm.AuthZRealm</code>	This realm performs the authorization decision and may or may not delegate out to the external XACML processing engine. It uses the incoming permissions to create a decision. However, it is possible to extend this realm using the <code>ddf.security.policy.extension.PolicyExtension</code> interface. This interface allows an integrator to add additional policy information to the PDP that can't be covered via its generic matching policies. This approach is often easier to configure for those that are not familiar with XACML.

Classname	Description
<code>org.codice.ddf.security.validator.*</code>	A number of STS validators are provided for X.509 (BinarySecurityToken), UsernameToken, SAML Assertion, and DDF custom tokens. The DDF custom tokens are all <code>BinarySecurityTokens</code> that may have PKI or username/password information as well as an authentication realm (correlates to JAAS realms installed in the container). The authentication realm allows an administrator to restrict which services they wish to use to authenticate users. For example: installing the <code>security-sts-ldaplogin</code> feature will enable a JAAS realm with the name "ldap". This realm can then be specified on any context using the Web Context Policy Manager. That realm selection is then passed via the token sent to the STS to determine which validator to use.

NOTE

Using the SAML Web SSO Identity Provider for authentication will ignore any realm settings and simply use all configured JAAS realms.

WARNING

An update was made to the SAML Assertion Handler to pass SAML assertions through the Authorization HTTP header. Cookies *are* still accepted and processed to maintain legacy federation compatibility, but assertions are sent in the header on outbound requests. While a machine's identity will still federate between versions, a user's identity will ONLY be federated when a DDF version 2.7.x server communicates with a DDF version 2.8.x+ server, or between two servers whose versions are 2.8.x or higher.

28.2. Security Core

The Security Core application contains all of the necessary components that are used to perform security operations (authentication, authorization, and auditing) required in the framework.

28.2.1. Security Core API

The Security Core API contains all of the DDF APIs that are used to perform security operations within DDF.

28.2.1.1. Installing the Security Core API

The Security Services App installs the Security Core API by default. Do not uninstall the Security Core API as it is integral to system function and all of the other security services depend upon it.

28.2.1.2. Configuring the Security Core API

The Security Core API has no configurable properties.

28.2.2. Security Core Implementation

The Security Core Implementation contains the reference implementations for the Security Core API interfaces that come with the DDF distribution.

28.2.2.1. Installing the Security Core Implementation

The Security Core app installs this bundle by default. It is recommended to use this bundle as it contains the reference implementations for many classes used within the Security Framework.

28.2.2.2. Configuring the Security Core Implementation

The Security Core Implementation has no configurable properties.

28.2.3. Security Core Commons

The Security Core Commons bundle contains helper and utility classes that are used within DDF to help with performing common security operations. Most notably, this bundle contains the `ddf.security.common.audit.SecurityLogger` class that performs the security audit logging within DDF.

28.2.3.1. Configuring the Security Core Commons

The Security Core Commons bundle has no configurable properties.

28.3. Security Encryption

The Security Encryption application offers an encryption framework and service implementation for other applications to use. This service is commonly used to encrypt and decrypt default passwords that are located within the metatype and Admin Console.

The encryption service and encryption command, which are based on [Keyczar](#), provide an easy way for developers to add encryption capabilities to DDF.

28.3.1. Security Encryption API

The Security Encryption API bundle provides the framework for the encryption service. Applications that use the encryption service should use the interfaces defined within it instead of calling an implementation directly.

28.3.1.1. Installing Security Encryption API

This bundle is installed by default as part of the `security-encryption` feature. Many applications that come with DDF depend on this bundle and it should not be uninstalled.

28.3.1.2. Configuring the Security Encryption API

The Security Encryption API has no configurable properties.

28.3.2. Security Encryption Implementation

The Security Encryption Implementation bundle contains all of the service implementations for the Encryption Framework and exports those implementations as services to the OSGi service registry.

28.3.2.1. Installing Security Encryption Implementation

This bundle is installed by default as part of the `security-encryption` feature. Other projects are dependent on the services this bundle exports and it should not be uninstalled unless another security service implementation is being added.

28.3.2.2. Configuring Security Encryption Implementation

The Security Encryption Implementation has no configurable properties.

28.3.3. Security Encryption Commands

The Security Encryption Commands bundle enhances the DDF system console by allowing administrators and integrators to encrypt and decrypt values directly from the console.

The `security:encrypt` command allows plain text to be encrypted using HMAC + AES for encryption with a randomly generated key that is created when the system is installed. This is useful when displaying password fields in a GUI.

Below is an example of the `security:encrypt` command used to encrypt the plain text "myPasswordToEncrypt". The output, `bR9mJpDVo8bTRwqGwIFxHJ5yFJzatKwjXjIo/8USWm8=`, is the encrypted value.

```
ddf@local>security:encrypt myPasswordToEncrypt
bR9mJpDVo8bTRwqGwIFxHJ5yFJzatKwjXjIo/8USWm8=
```

28.3.3.1. Installing the Security Encryption Commands

This bundle is installed by default with the `security-encryption` feature. This bundle is tied specifically to the DDF console and can be uninstalled if not needed. When uninstalled, however, administrators will not be able to encrypt and decrypt data from the console.

28.3.3.2. Configuring the Security Encryption Commands

The Security Encryption Commands have no configurable properties.

28.4. Security LDAP

The DDF LDAP application allows the user to configure either an embedded or a standalone LDAP server. The provided features contain a default set of schemas and users loaded to help facilitate authentication and authorization testing.

28.4.1. Embedded LDAP Server

DDF includes an embedded LDAP server (OpenDJ) for testing and demonstration purposes.

WARNING

The embedded LDAP server is intended for testing purposes only and is not recommended for production use.

28.4.1.1. Installing the Embedded LDAP Server

The embedded LDAP server is not installed by default with a standard installation.

1. Navigate to the Admin Console.
2. Select **Manage Applications**.
3. Install the **OpenDJ Embedded** application.

28.4.1.2. Configuring the Embedded LDAP

Configure the Embedded LDAP from the Admin Console:

1. Navigate to the **Admin Console**.
2. Select **Manage Applications**.
3. Start the **OpenDJ Embedded** application.
4. Once started, select the **OpenDJ** application.
5. Select the **Configuration** tab.

Table 65. OpenDJ Embedded Configurable Properties

Configurat ion Name	Description
LDAP Port	Sets the port for LDAP (plaintext and startTLS). 0 will disable the port.
LDAPS Port	Sets the port for LDAPS. 0 will disable the port.
Base LDIF File	Location on the server for a LDIF file. This file will be loaded into the LDAP and overwrite any existing entries. This option should be used when updating the default groups/users with a new LDIF file for testing. The LDIF file being loaded may contain any LDAP entries (schemas, users, groups, etc.). If the location is left blank, the default base LDIF file will be used that comes with DDF.

28.4.1.3. Connecting to Standalone LDAP Servers

DDF instances can connect to external LDAP servers by installing and configuring the **security-sts-ldaplogin** and **security-sts-ldapclaimshandler** features detailed here.

In order to connect to more than one LDAP server, configure these features for each LDAP server.

28.4.1.4. Embedded LDAP Configuration

The Embedded LDAP application contains an LDAP server (OpenDJ version 2.6.2) that has a default set of schemas and users loaded to help facilitate authentication and authorization testing.

Table 66. Embedded LDAP Default Ports Settings

Protocol	Default Port
LDAP	1389
LDAPS	1636
StartTLS	1389

Table 67. Embedded LDAP Default Users

Username	Password	Groups	Description
testuser1	password1		General test user for authentication
testuser2	password2		General test user for authentication
nromanova	password1	avengers	General test user for authentication
lcage	password1	admin, avengers	General test user for authentication, Admin user for karaf
jhowlett	password1	admin, avengers	General test user for authentication, Admin user for karaf
pparker	password1	admin, avengers	General test user for authentication, Admin user for karaf
jdrew	password1	admin, avengers	General test user for authentication, Admin user for karaf
tstark	password1	admin, avengers	General test user for authentication, Admin user for karaf
bbanner	password1	admin, avengers	General test user for authentication, Admin user for karaf
srogers	password1	admin, avengers	General test user for authentication, Admin user for karaf
admin	admin	admin	Admin user for karaf

Table 68. Embedded LDAP Default Admin User Settings

Username	Password	Groups	Attributes	Description
admin	secret			Administrative User for LDAP

28.4.1.5. Schemas

The default schemas loaded into the LDAP instance are the same defaults that come with OpenDJ.

Table 69. Embedded LDAP Default Schemas

Schema File Name	Schema Description
00-core.ldif	This file contains a core set of attribute type and objectclass definitions from several standard LDAP documents, including draft-ietf-boreham-nmsubordinates , draft-findlay-ldap-groupofentries , draft-furuseth-ldap-untypedobject , draft-good-ldap-changelog , draft-ietf-ldup-subentry , draft-wahl-ldap-adminaddr , RFC 1274, RFC 2079, RFC 2256, RFC 2798, RFC 3045, RFC 3296, RFC 3671, RFC 3672, RFC 4512, RFC 4519, RFC 4523, RFC 4524, RFC 4530, RFC 5020, and X.501.

Schema File Name	Schema Description
<code>01-pwpolicy.ldif</code>	This file contains schema definitions from draft-behera-ldap-password-policy , which defines a mechanism for storing password policy information in an LDAP directory server.
<code>02-config.ldif</code>	This file contains the attribute type and <code>objectclass</code> definitions for use with the directory server configuration.
<code>03-changelog.ldif</code>	This file contains schema definitions from draft-good-ldap-changelog , which defines a mechanism for storing information about changes to directory server data.
<code>03-rfc2713.ldif</code>	This file contains schema definitions from RFC 2713, which defines a mechanism for storing serialized Java objects in the directory server.
<code>03-rfc2714.ldif</code>	This file contains schema definitions from RFC 2714, which defines a mechanism for storing CORBA objects in the directory server.
<code>03-rfc2739.ldif</code>	This file contains schema definitions from RFC 2739, which defines a mechanism for storing calendar and vCard objects in the directory server. Note that the definition in RFC 2739 contains a number of errors, and this schema file has been altered from the standard definition in order to fix a number of those problems.
<code>03-rfc2926.ldif</code>	This file contains schema definitions from RFC 2926, which defines a mechanism for mapping between Service Location Protocol (SLP) advertisements and LDAP.
<code>03-rfc3112.ldif</code>	This file contains schema definitions from RFC 3112, which defines the authentication password schema.
<code>03-rfc3712.ldif</code>	This file contains schema definitions from RFC 3712, which defines a mechanism for storing printer information in the directory server.
<code>03-uddiv3.ldif</code>	This file contains schema definitions from RFC 4403, which defines a mechanism for storing UDDIV3 information in the directory server.
<code>04-rfc2307bis.ldif</code>	This file contains schema definitions from the draft-howard-rfc2307bis specification, used to store naming service information in the directory server.
<code>05-rfc4876.ldif</code>	This file contains schema definitions from RFC 4876, which defines a schema for storing Directory User Agent (DUA) profiles and preferences in the directory server.
<code>05-samba.ldif</code>	This file contains schema definitions required when storing Samba user accounts in the directory server.
<code>05-solaris.ldif</code>	This file contains schema definitions required for Solaris and OpenSolaris LDAP naming services.
<code>06-compat.ldif</code>	This file contains the attribute type and <code>objectclass</code> definitions for use with the directory server configuration.

28.4.1.6. Starting and Stopping the Embedded LDAP

The embedded LDAP application installs a feature with the name `ldap-embedded`. Installing and uninstalling this feature will start and stop the embedded LDAP server. This will also install a fresh instance of the server each time. If changes need to persist, stop then start the `embedded-ldap-opendj`

bundle (rather than installing/uninstalling the feature).

All settings, configurations, and changes made to the embedded LDAP instances are persisted across DDF restarts. If DDF is stopped while the LDAP feature is installed and started, it will automatically restart with the saved settings on the next DDF start.

28.4.1.7. Limitations of the Embedded LDAP

Current limitations for the embedded LDAP instances include:

- Inability to store the LDAP files/storage outside of the DDF installation directory. This results in any LDAP data (i.e., LDAP user information) being lost when the `ldap-embedded` feature is uninstalled.
- Cannot be run standalone from DDF. In order to run `embedded-ldap`, the DDF must be started.

28.4.1.8. External Links for the Embedded LDAP

Location to the default base LDIF file in the DDF [source code](#).

[OpenDJ documentation](#)

28.4.1.9. LDAP Administration

OpenDJ provides a number of tools for LDAP administration. Refer to the [OpenDJ Admin Guide](#).

28.4.1.10. Downloading the Admin Tools

Download [OpenDJ \(Version 2.4.6\)](#) and the included tool suite.

28.4.1.11. Using the Admin Tools

The admin tools are located in `<opendj-installation>/bat` for Windows and `<opendj-installation>/bin` for nix. These tools can be used to administer both local and remote LDAP servers by setting the `*host` and `port` parameters appropriately.

In this example, the user **Bruce Banner (uid=bbanner)** is disabled using the **manage-account** command on Windows. Run **manage-account --help** for usage instructions.

Example Commands for Disabling/Enabling a User's Account

```
D:\OpenDJ-2.4.6\bat>manage-account set-account-is-disabled -h localhost -p 4444 -0
true
-D "cn=admin" -w secret -b "uid=bbanner,ou=users,dc=example,dc=com"
The server is using the following certificate:
  Subject DN: CN=Win7-1, O=Administration Connector Self-Signed Certificate
  Issuer DN: CN=Win7-1, O=Administration Connector Self-Signed Certificate
  Validity: Wed Sep 04 15:36:46 MST 2013 through Fri Sep 04 15:36:46 MST 2015
Do you wish to trust this certificate and continue connecting to the server?
Please enter "yes" or "no":yes
Account Is Disabled: true
```

Notice **Account Is Disabled: true** in the listing:

Verifying an Account is Disabled

```
D:\OpenDJ-2.4.6\bat>manage-account get-all -h localhost -p 4444 -D "cn=admin" -w
secret
-b "uid=bbanner,ou=users,dc=example,dc=com"
The server is using the following certificate:
  Subject DN: CN=Win7-1, O=Administration Connector Self-Signed Certificate
  Issuer DN: CN=Win7-1, O=Administration Connector Self-Signed Certificate
  Validity: Wed Sep 04 15:36:46 MST 2013 through Fri Sep 04 15:36:46 MST 2015
Do you wish to trust this certificate and continue connecting to the server?
Please enter "yes" or "no":yes
Password Policy DN: cn=Default Password Policy,cn=Password Policies,cn=config
Account Is Disabled: true
Account Expiration Time:
Seconds Until Account Expiration:
Password Changed Time: 19700101000000.000Z
Password Expiration Warned Time:
Seconds Until Password Expiration:
Seconds Until Password Expiration Warning:
Authentication Failure Times:
Seconds Until Authentication Failure Unlock:
Remaining Authentication Failure Count:
Last Login Time:
Seconds Until Idle Account Lockout:
Password Is Reset: false
Seconds Until Password Reset Lockout:
Grace Login Use Times:
Remaining Grace Login Count: 0
Password Changed by Required Time:
Seconds Until Required Change Time:
Password History:
```

Enabling an Account

```
D:\OpenDJ-2.4.6\bat>manage-account clear-account-is-disabled -h localhost -p 4444 -D
"cn=admin" -w secret -b "uid=bbanner,ou=users,dc=example,dc=com"
The server is using the following certificate:
  Subject DN: CN=Win7-1, O=Administration Connector Self-Signed Certificate
  Issuer DN: CN=Win7-1, O=Administration Connector Self-Signed Certificate
  Validity: Wed Sep 04 15:36:46 MST 2013 through Fri Sep 04 15:36:46 MST 2015
Do you wish to trust this certificate and continue connecting to the server?
Please enter "yes" or "no":yes
Account Is Disabled: false
```

Notice **Account Is Disabled: false** in the listing.

Verifying an Account is Enabled

```
D:\OpenDJ-2.4.6\bat>manage-account get-all -h localhost -p 4444 -D "cn=admin" -w
secret
-b "uid=bbanner,ou=users,dc=example,dc=com"
The server is using the following certificate:
  Subject DN: CN=Win7-1, O=Administration Connector Self-Signed Certificate
  Issuer DN: CN=Win7-1, O=Administration Connector Self-Signed Certificate
  Validity: Wed Sep 04 15:36:46 MST 2013 through Fri Sep 04 15:36:46 MST 2015
Do you wish to trust this certificate and continue connecting to the server?
Please enter "yes" or "no":yes
Password Policy DN: cn=Default Password Policy,cn=Password Policies,cn=config
Account Is Disabled: false
Account Expiration Time:
Seconds Until Account Expiration:
Password Changed Time: 19700101000000.000Z
Password Expiration Warned Time:
Seconds Until Password Expiration:
Seconds Until Password Expiration Warning:
Authentication Failure Times:
Seconds Until Authentication Failure Unlock:
Remaining Authentication Failure Count:
Last Login Time:
Seconds Until Idle Account Lockout:
Password Is Reset: false
Seconds Until Password Reset Lockout:
Grace Login Use Times:
Remaining Grace Login Count: 0
Password Changed by Required Time:
Seconds Until Required Change Time:
Password History:
```

28.5. Security PDP

The Security Policy Decision Point (PDP) module contains services that are able to perform authorization decisions based on configurations and policies. In the Security Framework, these components are called realms, and they implement the `org.apache.shiro.realm.Realm` and `org.apache.shiro.authz.Authorizer` interfaces. Although these components perform decisions on access control, enforcement of this decision is performed by components within the notional PEP application.

28.5.1. Security PDP AuthZ Realm

The Security PDP AuthZ Realm exposes a realm service that makes decisions on authorization requests using the attributes stored within the metocard to determine if access should be granted. This realm can use XACML and will delegate decisions to an external processing engine if internal processing fails. Decisions are first made based on the "match-all" and "match-one" logic. Any attributes listed in the "match-all" or "match-one" sections will not be passed to the XACML processing engine and they will be matched internally. It is recommended to list as many attributes

as possible in these sections to avoid going out to the XACML processing engine for performance reasons. If it is desired that all decisions be passed to the XACML processing engine, remove all of the "match-all" and "match-one" configurations. The configuration below provides the mapping between user attributes and the attributes being asserted - one map exists for each type of mapping (each map may contain multiple values).

Match-All Mapping:: This mapping is used to guarantee that all values present in the specified metocard attribute exist in the corresponding user attribute. Match-One Mapping:: This mapping is used to guarantee that at least one of the values present in the specified metocard attribute exists in the corresponding user attribute.

28.5.1.1. Configuring the Security PDP AuthZ Realm

1. Navigate to the **Admin Console**.
2. Select **Security** Application.
3. Select **Configuration** tab.
4. Select **Security AuthZ Realm**.

See [Security AuthZ Realm](#) for all possible configurations.

28.5.2. Guest Interceptor

The goal of the **GuestInterceptor** is to allow non-secure clients (such as SOAP requests without security headers) to access secure service endpoints.

All requests to secure endpoints must satisfy the WS-SecurityPolicy that is included in the WSDL.

Rather than reject requests without user credentials, the guest interceptor detects the missing credentials and inserts an assertion that represents the "guest" user. The attributes included in this guest user assertion are configured by the administrator to represent any unknown user on the current network.

28.5.2.1. Installing Guest Interceptor

The **GuestInterceptor** is installed by default with Security Application.

28.5.2.2. Configuring Guest Interceptor

Configure the Guest Interceptor from the Admin Console:

1. Navigate to the **Admin Console** at <https://{FQDN}:{PORT}/admin>
2. Select the **Security** application.
3. Select the **Configuration** tab.
4. Select the **Security STS Guest Claims Handler** configuration.
5. Select the **+** next to Attributes to add a new attribute.
6. Add any additional attributes that will apply to every user.
7. Select **Save changes**.


Once these configurations have been added, the GuestInterceptor is ready for use. Both secure and non-secure requests will be accepted by all secure DDF service endpoints.

28.6. Web Service Security Architecture

The Web Service Security (WSS) functionality that comes with DDF is integrated throughout the system. This is a central resource describing how all of the pieces work together and where they are located within the system.

DDF comes with a **Security Framework** and **Security Services**. The Security Framework is the set of APIs that define the integration with the DDF framework and the Security Services are the reference implementations of those APIs built for a realistic end-to-end use case.

28.6.1. Securing REST




The Delegate Servlet Filter is the topmost filter of all web contexts. It initializes all Security Filters and runs them in order according to service ranking:

1. The **Web SSO Filter** reads from the web context policy manager and functions as the first decision point. If the request is from a whitelisted context, no further authentication is needed and the request goes directly to the desired endpoint. If the context is not on the whitelist, the filter will attempt to get a claims handler for the context. The filter loops through all configured context handlers until one signals that it has found authentication information that it can use to build a token. This configuration can be changed by modifying the web context policy manager configuration. If unable to resolve the context, the filter will return an authentication error and the process stops. If a handler is successfully found, an auth token is assigned and the request continues to the login filter.
2. The **Login Filter** receives a token and returns a subject. To retrieve the subject, the token is sent through Shiro to the STS Realm where the token will be exchanged for a SAML assertion through a SOAP call to an STS server.
3. If the Subject is returned, the request moves to the **AuthZ Filter** to check permissions on the user. If the user has the correct permissions to access that web context, the request is allowed to

hit the endpoint.

NOTE This diagram does not yet include the SAML 2.0 Web SSO integration.

28.6.2. Securing SOAP



28.6.2.1. SOAP Secure Client

When calling to an endpoint from a SOAP secure client, it first requests the WSDL from the endpoint and the SOAP endpoint returns the WSDL. The client then calls to STS for authentication token to proceed. If the client receives the token, it makes a secure call to the endpoint and receives results.

28.6.2.2. Policy-unaware SOAP Client

If calling an endpoint from a non-secure client, at the point of the initial call, the Guest Interceptor catches the request and prepares it to be accepted by the endpoint.

First, the interceptor reads the configured policy, builds a security header, and gets an anonymous SAML assertion. Using this, it makes a `getSubject` call which is sent through Shiro to the STS realm. Upon success, the STS realm returns the subject and the call is made to the endpoint.

28.7. Security PEP

The Security Policy Enforcement Point (PEP) application contains bundles that allow for policies to be enforced at various parts of the system, for example: to reach contexts, view metacards, access catalog operations, and others.

28.7.1. Security PEP Interceptor

The Security PEP Interceptor bundle contains the `ddf.security.pep.interceptor.PEPAuthorizingInterceptor` class. This class uses CXF to intercept incoming SOAP messages and enforces service authorization policies by sending the service request to the security framework.

28.7.1.1. Installing the Security PEP Interceptor

This bundle is not installed by default but can be added by installing the `security-pep-serviceauthz` feature.

WARNING

To perform service authorization within a default install of DDF, this bundle MUST be installed.

28.7.1.2. Configuring the Security PEP Interceptor

The Security PEP Interceptor has no configurable properties.

28.8. Filtering

Metocard filtering is performed by the [Filter Plugin](#) after a query has been performed, but before the results are returned to the requestor.

Each metocard result will contain security attributes that are populated by the CatalogFramework based on the PolicyPlugins (Not provided! You must create your own plugin for your specific metadata!) that populates this attribute. The security attribute is a HashMap containing a set of keys that map to lists of values. The metocard is then processed by a filter plugin that creates a `KeyValueCollectionPermission` from the metocard's security attribute. This permission is then checked against the user subject to determine if the subject has the correct claims to view that metocard. The decision to filter the metocard eventually relies on the PDP (`feature:install security-pdp-authz`). The PDP returns a decision, and the metocard will either be filtered or allowed to pass through.

The security attributes populated on the metocard are completely dependent on the type of the metocard. Each type of metocard must have its own PolicyPlugin that reads the metadata being returned and returns the metocard's security attribute. If the subject permissions are missing during filtering, all resources will be filtered.

Example (represented as simple XML for ease of understanding):

```
<metocard>
  <security>
    <map>
      <entry key="entry1" value="A,B" />
      <entry key="entry2" value="X,Y" />
      <entry key="entry3" value="USA,GBR" />
      <entry key="entry4" value="USA,AUS" />
    </map>
  </security>
</metocard>
```

```
<user>
  <claim name="claim1">
    <value>A</value>
    <value>B</value>
  </claim>
  <claim name="claim2">
    <value>X</value>
    <value>Y</value>
  </claim>
  <claim name="claim3">
    <value>USA</value>
  </claim>
  <claim name="claim4">
    <value>USA</value>
  </claim>
</user>
```

In the above example, the user's claims are represented very simply and are similar to how they would actually appear in a SAML 2 assertion. Each of these user (or subject) claims will be converted to a KeyValuePermission object. These permission objects will be implied against the permission object generated from the metocard record. In this particular case, the metocard might be allowed if the policy is configured appropriately because all of the permissions line up correctly.

To enable filtering on a new type of record, implement a PolicyPlugin that is able to read the string metadata contained within the metocard record. Note that, in DDF, there is no default plugin that parses a metocard. A plugin must be created to create a policy for the metocard.

28.9. Expansion Service

The Expansion Service and its corresponding expansion-related commands provide an easy way for developers to add expansion capabilities to DDF during user attribute and metadata card processing. In addition to these two defined uses of the expansion service, developers are free to utilize the service in their own implementations.

Each instance of the expansion service consists of a collection of rulesets. Each ruleset consists of a

key value and its associated set of rules. Callers of the expansion service provide a key and a value to be expanded. The expansion service then looks up the set of rules for the specified key. The expansion service cumulatively applies each of the rules in the set, starting with the original value. The result is returned to the caller.

Table 70. Expansion Service Ruleset Format

Key (Attribute)	Rules (original → new)	
key1	value1	replacement1
	value2	replacement2
	value3	replacement3
key2	value1	replacement1
	value2	replacement2

The examples below use the following collection of rulesets:

Table 71. Expansion Service Example Ruleset

Key (Attribute)	Rules (original → new)	
Location	Goodyear	Goodyear AZ
	AZ	AZ USA
	CA	CA USA
Title	VP-Sales	VP-Sales VP Sales
	VP-Engineering	VP-Engineering VP Engineering

Note that the rules listed for each key are processed in order, so they may build upon each other, i.e., a new value from the new replacement string may be expanded by a subsequent rule. In the example `Location:Goodyear` would expand to `Goodyear AZ USA` and `Title:VP-Sales` would expand to `VP-Sales VP Sales`.

28.10. Security Token Service

The Security Token Service (STS) is a service running in DDF that generates SAML v2.0 assertions. These assertions are then used to authenticate a client allowing them to issue other requests, such as ingest or queries to DDF services.

The STS is an extension of Apache CXF-STS. It is a SOAP web service that utilizes WS-Trust. The generated SAML assertions contain attributes about a user and is used by the Policy Enforcement Point (PEP) in the secure endpoints. Specific configuration details on the bundles that come with DDF can be found on the Security STS application page. This page details all of the STS components that come out of the box with DDF, along with configuration options, installation help, and which services they import and export.

The STS server contains validators, claim handlers, and token issuers to process incoming requests. When a request is received, the validators first ensure that it is valid. The validators verify authentication against configured services, such as LDAP, DIAS, PKI. If the request is found to be invalid, the process ends and an error is returned. Next, the claims handlers determine how to handle the request, adding user attributes or properties as configured. The token issuer creates a

SAML 2.0 assertion and associates it with the subject. The STS server sends an assertion back to the requestor, which is used to authenticate and authorize subsequent SOAP and REST requests.

The STS can be used to generate SAML v2.0 assertions via a SOAP web service request. Out of the box, the STS supports authentication from existing SAML tokens, CAS proxy tickets, username/password, and x509 certificates. It also supports retrieving claims using LDAP and properties files.

28.10.1. STS Claims Handlers

Claims handlers are classes that convert the incoming user credentials into a set of attribute claims that will be populated in the SAML assertion. An example in action would be the `LDAPClaimsHandler` that takes in the user's credentials and retrieves the user's attributes from a backend LDAP server. These attributes are then mapped and added to the SAML assertion being created. Integrators and developers can add more claims handlers that can handle other types of external services that store user attributes.

29. Developing DDF Components

Create custom implementations of DDF components.

29.1. Developing Complementary Catalog Frameworks

DDF and the underlying OSGi technology can serve as a robust infrastructure for developing frameworks that complement the Catalog.

29.1.1. Simple Catalog API Implementations

The Catalog API implementations, which are denoted with the suffix of `Impl` on the Java file names, have multiple purposes and uses:

- First, they provide a good starting point for other developers to extend functionality in the framework. For instance, extending the `MetocardImpl` allows developers to focus less on the inner workings of DDF and more on the developer's intended purposes and objectives.
- Second, the Catalog API Implementations display the proper usage of an interface and an interface's intentions. Also, they are good code examples for future implementations. If a developer does not want to extend the simple implementations, the developer can at least have a working code reference on which to base future development.

29.1.2. Use of the Whiteboard Design Pattern

The Catalog makes extensive use of the Whiteboard Design Pattern. Catalog Components are registered as services in the OSGi Service Registry, and the Catalog Framework or any other clients tracking the OSGi Service Registry are automatically notified by the OSGi Framework of additions and removals of relevant services.

The Whiteboard Design Pattern is a common OSGi technique that is derived from a

technical whitepaper provided by the OSGi Alliance in 2004. It is recommended to use the Whiteboard pattern over the Listener pattern in OSGi because it provides less complexity in code (both on the client and server sides), fewer deadlock possibilities than the Listener pattern, and closely models the intended usage of the OSGi framework.

29.1.3. Recommendations for Framework Development

- Provide extensibility similar to that of the Catalog.
 - Provide a stable API with interfaces and simple implementations (refer to http://www.ibm.com/developerworks/websphere/techjournal/1007_charters/1007_charters.html).
- Make use of the Catalog wherever possible to store, search, and transform information.
- Utilize OSGi standards wherever possible.
 - ConfigurationAdmin
 - MetaType
- Utilize the sub-frameworks available in DDF.
 - Karaf
 - CXF
 - PAX Web and Jetty

29.1.4. Catalog Framework Reference

The Catalog Framework can be requested from the OSGi Service Registry.

Blueprint Service Reference

```
<reference id="catalogFramework" interface="DDF.catalog.CatalogFramework" />
```

29.1.4.1. Methods

The **CatalogFramework** provides convenient methods to transform **Metacards** and **QueryResponses** using a reference to the **CatalogFramework**.

29.1.4.1.1. Create, Update, and Delete Methods

Create, Update, and Delete (CUD) methods add, change, or remove stored metadata in the local Catalog Provider.

Example Create, Update, Delete Methods

```
public CreateResponse create(CreateRequest createRequest) throws IngestException,  
SourceUnavailableException;  
public UpdateResponse update(UpdateRequest updateRequest) throws IngestException,  
SourceUnavailableException;  
public DeleteResponse delete(DeleteRequest deleteRequest) throws IngestException,  
SourceUnavailableException;
```

CUD operations process **PolicyPlugin**, **AccessPlugin**, and **PreIngestPlugin** instances before execution and **PostIngestPlugin** instances after execution.

29.1.4.1.2. Query Methods

Query methods search metadata from available Sources based on the **QueryRequest** properties and Federation Strategy. Sources could include Catalog Provider, Connected Sources, and Federated Sources.

Example Query Methods

```
public QueryResponse query(QueryRequest query) throws UnsupportedQueryException  
, SourceUnavailableException, FederationException;  
public QueryResponse query(QueryRequest queryRequest, FederationStrategy strategy)  
throws SourceUnavailableException, UnsupportedQueryException, FederationException;
```

Query requests process **PolicyPlugin**, **AccessPlugin**, and **PreQueryPlugin** instances before execution and **PolicyPlugin**, **AccessPlugin**, and **PostQueryPlugin** instances after execution.

29.1.4.1.3. Resource Methods

Resource methods retrieve products from Sources.

Example Resource Methods

```
public ResourceResponse getEnterpriseResource(ResourceRequest request)  
throws IOException, ResourceNotFoundException, ResourceNotSupportedException;  
public ResourceResponse getLocalResource(ResourceRequest request) throws IOException,  
ResourceNotFoundException, ResourceNotSupportedException;  
public ResourceResponse getResource(ResourceRequest request, String resourceSiteName)  
throws IOException, ResourceNotFoundException, ResourceNotSupportedException;
```

Resource requests process `PreResourcePlugin`'s before execution and `PostResourcePlugin`'s after execution.

29.1.4.1.4. Source Methods

Source methods can get a list of Source identifiers or request descriptions about Sources.

Example Source Methods

```
public Set<String> getSourceIds();
public SourceInfoResponse getSourceInfo(SourceInfoRequest sourceInfoRequest) throws
SourceUnavailableException;
```

29.1.4.1.5. Transform Methods

Transform methods provide convenience methods for using Metocard Transformers and Query Response Transformers.

Transform Methods

```
// Metocard Transformer
public BinaryContent transform(Metocard metocard, String transformerId, Map<String
, Serializable> requestProperties) throws CatalogTransformerException;

// Query Response Transformer
public BinaryContent transform(SourceResponse response, String transformerId, Map
<String, Serializable> requestProperties) throws CatalogTransformerException;
```

29.1.4.2. Implementing Catalog Methods

Query Response Transform Example

```
// inject CatalogFramework instance or retrieve an instance
private CatalogFramework catalogFramework;

public RSSEndpoint(CatalogFramework catalogFramework)
{
    this.catalogFramework = catalogFramework ;
    // implementation
}

// Other implementation details ...

private void convert(QueryResponse queryResponse ) {
    // ...
    String transformerId = "rss";

    BinaryContent content = catalogFramework.transform(queryResponse, transformerId,
null);

    // ...
}
```

29.1.4.3. Dependency Injection

Using Blueprint or another injection framework, transformers can be injected from the OSGi Service Registry.

Blueprint Service Reference

```
<reference id="[[Reference Id]]" interface="DDF.catalog.transform.[[Transformer Interface Name]]" filter="(shortname=[[Transformer Identifier]])" />
```

Each transformer has one or more `transform` methods that can be used to get the desired output.

Input Transformer Example

```
DDF.catalog.transform.InputTransformer inputTransformer = retrieveInjectedInstance() ;  
Metocard entry = inputTransformer.transform(messageInputStream);
```

Metocard Transformer Example

```
DDF.catalog.transform.MetocardTransformer metocardTransformer =  
retrieveInjectedInstance() ;  
BinaryContent content = metocardTransformer.transform(metocard, arguments);
```

Query Response Transformer Example

```
DDF.catalog.transform.QueryResponseTransformer queryResponseTransformer =  
retrieveInjectedInstance() ;  
BinaryContent content = queryResponseTransformer.transform(sourceResponse, arguments);
```

29.1.4.4. OSGi Service Registry

IMPORTANT

In the vast majority of cases, working with the OSGi Service Reference directly should be avoided. Instead, dependencies should be injected via a dependency injection framework like Blueprint.

Transformers are registered with the OSGi Service Registry. Using a `BundleContext` and a filter, references to a registered service can be retrieved.

OSGi Service Registry Reference Example

```
ServiceReference[] refs =  
bundleContext.getServiceReferences(DDF.catalog.transform.InputTransformer.class  
.getName(), "(shortname=" + transformerId + ")");  
InputTransformer inputTransformer = (InputTransformer) context.getService(refs[0]);  
Metocard entry = inputTransformer.transform(messageInputStream);
```

29.2. Developing Metacard Types

Create custome Metacard types with Metacard Type definition files.

29.2.1. Metacard Type Definition File

To define Metacard Types, the definition file must have a `metacardTypes` key in the root object.

```
{  
  "metacardTypes": [...]  
}
```

The value of `metacardTypes` must be an array of Metacard Type Objects, which are composed of the `type` and `attributes` keys.

```
{  
  "metacardTypes": [  
    {  
      "type": "my-metacard-type",  
      "attributes": {...}  
    }  
  ]  
}
```

The value of the `type` key is the name of the metacard type being defined.

The value of the `attributes` key is a map where each key is the name of an attribute type to include in this metacard type and each value is a map with a single key named `required` and a boolean value. Required attributes are used for metacard validation - metacards that lack required attributes will be flagged with validation errors.

```
{  
  "metacardTypes": [  
    {  
      "type": "my-metacard-type",  
      "attributes": {  
        "resolution": {  
          "required": true  
        },  
        "target-areas": {  
          "required": false  
        },  
        "expiration": {  
          "required": false  
        },  
        "point-of-contact": {  
          "required": true  
        }  
      }  
    }  
  ]  
}
```

NOTE

The DDF basic metocard attribute types are added to custom metocard types by default. If any attribute types are required by a metocard type, just include them in the `attributes` map and set `required` to `true`, as shown in the above example with `point-of-contact`.

```
{  
  "metocardTypes": [  
    {  
      "type": "my-metocard-type",  
      "attributes": {  
        "resolution": {  
          "required": true  
        },  
        "target-areas": {  
          "required": false  
        }  
      }  
    },  
    {  
      "type": "another-metocard-type",  
      "attributes": {  
        "effective": {  
          "required": true  
        },  
        "resolution": {  
          "required": false  
        }  
      }  
    }  
  ]  
}
```

29.3. Developing Global Attribute Validators

29.3.1. Global Attribute Validators File

To define Validators, the definition file must have a `validators` key in the root object.

```
{  
  "validators": {...}  
}
```

The value of `validators` is a map of the attribute name to a list of validators for that attribute.

```
{  
  "validators": {  
    "point-of-contact": [...]  
  }  
}
```

Each object in the list of validators is the validator name and list of arguments for that validator.

```
{  
  "validators": {  
    "point-of-contact": [  
      {  
        "validator": "pattern",  
        "arguments": [".*regex.+\\s"]  
      }  
    ]  
  }  
}
```

WARNING

The value of the `arguments` key must always be an array of strings, even for numeric arguments, e.g. `["1", "10"]`

The `validator` key must have a value of one of the following:

2. **validator** Possible Values

- **size** (validates the size of Strings, Arrays, Collections, and Maps)
 - **arguments**: (2) [integer: lower bound (inclusive), integer: upper bound (inclusive)]
 - lower bound must be greater than or equal to zero and the upper bound must be greater than or equal to the lower bound
- **pattern**
 - **arguments**: (1) [regular expression]
- **pastdate**
 - **arguments**: (0) [NO ARGUMENTS]
- **futuredate**
 - **arguments**: (0) [NO ARGUMENTS]
- **range**
 - (2) [number (decimal or integer): inclusive lower bound, number (decimal or integer): inclusive upper bound]
 - uses a default epsilon of 1E-6 on either side of the range to account for floating point representation inaccuracies
 - (3) [number (decimal or integer): inclusive lower bound, number (decimal or integer): inclusive upper bound, decimal number: epsilon (the maximum tolerable error on either side of the range)]
- **enumeration**
 - **arguments**: (unlimited) [list of strings: each argument is one case-sensitive, valid enumeration value]
- **match_any**
 - **validators**: (unlimited) [list of previously defined validators: valid if any validator succeeds]

Example Validator Definition

```
{  
  "validators": {  
    "title": [  
      {  
        "validator": "size",  
        "arguments": ["1", "50"]  
      },  
      {  
        "validator": "pattern",  
        "arguments": ["\\D+"]  
      }  
    ],  
    "created": [  
      {  
        "validator": "pastdate",  
        "arguments": ["2012-01-01"]  
      }  
    ]  
  }  
}
```

```

        "arguments": []
    }
],
"expiration": [
    {
        "validator": "futurereadate",
        "arguments": []
    }
],
"page-count": [
    {
        "validator": "range",
        "arguments": ["1", "500"]
    }
],
"temperature": [
    {
        "validator": "range",
        "arguments": ["12.2", "19.8", "0.01"]
    }
],
"resolution": [
    {
        "validator": "enumeration",
        "arguments": ["1080p", "1080i", "720p"]
    }
],
"datatype": [
    {
        "validator": "match_any",
        "validators": [
            {
                "validator": "range",
                "arguments": ["1", "25"]
            },
            {
                "validator": "enumeration",
                "arguments": ["Collection", "Dataset", "Event"]
            }
        ]
    }
]
}
}

```

29.4. Developing Attribute Types

Create custom attribute types with Attribute Type definition files.

29.4.1. Attribute Type Definition File

To define Attribute Types, the definition file must have an `attributeTypes` key in the root object.

```
{  
  "attributeTypes": {...}  
}
```

The value of `attributeTypes` must be a map where each key is the attribute type's name and each value is a map that includes the data type and whether the attribute type is stored, indexed, tokenized, or multi-valued.

Attribute Types

```
{  
  "attributeTypes": {  
    "temperature": {  
      "type": "DOUBLE_TYPE",  
      "stored": true,  
      "indexed": true,  
      "tokenized": false,  
      "multivalued": false  
    }  
  }  
}
```

The attributes `stored`, `indexed`, `tokenized`, and `multivalued` must be included and must have a boolean value.

3. Required Attribute Definitions

`stored`

If true, the value of the attribute should be stored in the underlying datastore. Some attributes may only be indexed or used in transit and do not need to be persisted.

`indexed`

If true, then the value of the attribute should be included in the datastore's index and therefore be part of query evaluation.

`tokenized`

Only applicable to STRING_TYPE attributes, if true then stopwords and punctuation will be stripped prior to storing and/or indexing. If false, only an exact string will match.

`multi-valued`

If true, then the attribute values will be Lists of the attribute type rather than single values.

The `type` attribute must also be included and must have one of the allowed values:

4. `type` Attribute Possible Values

- `DATE_TYPE`
- `STRING_TYPE`
- `XML_TYPE`
- `LONG_TYPE`
- `BINARY_TYPE`
- `GEO_TYPE`
- `BOOLEAN_TYPE`
- `DOUBLE_TYPE`
- `FLOAT_TYPE`
- `INTEGER_TYPE`
- `OBJECT_TYPE`
- `SHORT_TYPE`

An example with multiple attributes defined:

Multiple Attributes Defined

```
{  
  "attributeTypes": {  
    "resolution": {  
      "type": "STRING_TYPE",  
      "stored": true,  
      "indexed": true,  
      "tokenized": false,  
      "multivalued": false  
    },  
    "target-areas": {  
      "type": "GEO_TYPE",  
      "stored": true,  
      "indexed": true,  
      "tokenized": false,  
      "multivalued": true  
    }  
  }  
}
```

29.5. Developing Default Attribute Types

Create custom default attribute types.

29.5.1. Default Attribute Values

To define default attribute values, the definition file must have a `defaults` key in the root object.

```
{  
  "defaults": [...]  
}
```

The value of `defaults` is a list of objects where each object contains the keys `attribute`, `value`, and optionally `metocardTypes`.

```
{  
  "defaults": [  
    {  
      "attribute": ...,  
      "value": ...,  
      "metocardTypes": [...]  
    }  
  ]  
}
```

The value corresponding to the `attribute` key is the name of the attribute to which the default value will be applied. The value corresponding to the `value` key is the default value of the attribute.

NOTE

The attribute's default value must be of the same type as the attribute, but it has to be written as a string (i.e., enclosed in quotation marks) in the JSON file.

Dates must be UTC datetimes in the ISO 8601 format, i.e., `yyyy-MM-ddTHH:mm:ssZ`

The `metocardTypes` key is optional. If it is left out, then the default attribute value will be applied to every metocard that has that attribute. It can be thought of as a 'global' default value. If the `metocardTypes` key is included, then its value must be a list of strings where each string is the name of a metocard type. In this case, the default attribute value will be applied only to metacards that match one of the types given in the list.

NOTE

In the event that an attribute has a 'global' default value as well as a default value for a specific metocard type, the default value for the specific metocard type will be applied (i.e., the more specific default value wins).

Example:

```
{
  "defaults": [
    {
      "attribute": "title",
      "value": "Default Title"
    },
    {
      "attribute": "description",
      "value": "Default video description",
      "metacardTypes": ["video"]
    },
    {
      "attribute": "expiration",
      "value": "2020-05-06T12:00:00Z",
      "metacardTypes": ["video", "nitf"]
    },
    {
      "attribute": "frame-rate",
      "value": "30"
    }
  ]
}
```

29.6. Developing Attribute Injections

Attribute injections are defined attributes that will be injected into all metocard types or into specific metocard types. This capability allows metocard types to be extended with new attributes.

29.6.1. Attribute Injection Definition

To define attribute injections, the definition file must have an `inject` key in the root object.

```
{
  "inject": [...]
}
```

The value of `inject` is simply a list of objects where each object contains the key `attribute` and optionally `metacardTypes`.

```
{
  "inject": [
    {
      "attribute": ....,
      "metocardTypes": [...]
    }
  ]
}
```

The value corresponding to the `attribute` key is the name of the attribute to inject.

The `metocardTypes` key is optional. If it is left out, then the attribute will be injected into every metocard type. In that case it can be thought of as a 'global' attribute injection. If the `metocardTypes` key is included, then its value must be a list of strings where each string is the name of a metocard type. In this case, the attribute will be injected only into metocard types that match one of the types given in the list.

Example:

```
{
  "inject": [
    {
      "attribute": "rating"
    },
    {
      "attribute": "cloud-cover",
      "metocardTypes": ["nitf", "video"]
    }
  ]
}
```

NOTE Attributes must be registered in the attribute registry (see the [AttributeRegistry](#) interface) in order to be injected into metocard types. For example, attributes defined in JSON definition files are placed in the registry, so they can be injected.

29.7. Developing Endpoints

Custom endpoints can be created, if necessary. See [Available Endpoints](#) for descriptions of provided endpoints.

Complete the following procedure to create an endpoint.

1. Create a Java class that implements the endpoint's business logic. Example: Creating a web service that external clients can invoke.
2. Add the endpoint's business logic, invoking [CatalogFramework](#) calls as needed.
3. Import the DDF packages to the bundle's manifest for run-time (in addition to any other

required packages):

`Import-Package: ddf.catalog, ddf.catalog.*`

4. Retrieve an instance of `CatalogFramework` from the OSGi registry. (Refer to [Working with OSGi - Service Registry](#) for examples.)
5. Deploy the packaged service to DDF. (Refer to [Working with OSGi - Bundles](#).)

NOTE

It is recommended to use the maven bundle plugin to create the Endpoint bundle's manifest as opposed to directly editing the manifest file.

No implementation of an interface is required

TIP

Unlike other DDF components that require you to implement a standard interface, no implementation of an interface is required in order to create an endpoint.

Table 72. Common Endpoint Business Logic

Methods	Use
<code>Ingest</code>	Add, modify, and remove metadata using the ingest-related <code>CatalogFramework</code> methods: <code>create, update, and delete.</code>
<code>Query</code>	Request metadata using the <code>query</code> method.
<code>Source</code>	Get available <code>Source</code> information.
<code>Resource</code>	Retrieve products referenced in Metacards from Sources.
<code>Transform</code>	Convert common Catalog Framework data types to and from other data formats.

29.8. Developing Input Transformers

DDF supports the creation of custom [input transformers](#) for use cases not covered by the included implementations.

Creating a custom input Transformer:

1. Create a new Java class that implements `ddf.catalog.transform.InputTransformer`.

```
public class SampleInputTransformer implements ddf.catalog.transform.InputTransformer
```
2. Implement the transform methods.

```
public Metocard transform(InputStream input) throws IOException,  
CatalogTransformerException  
public Metocard transform(InputStream input, String id) throws IOException,  
CatalogTransformerException
```
3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog,ddf.catalog.transform
```
4. Create an OSGi descriptor file to communicate with the OSGi Service Registry (described in the [Working with OSGi](#) section). Export the service to the OSGi Registry and declare service properties.

Input Transformer Blueprint Descriptor Example

```
...
<service ref="[[SampleInputTransformer]]" interface=
"ddf.catalog.transform.InputTransformer">
  <service-properties>
    <entry key="shortname" value="[[sampletransform]]" />
    <entry key="title" value="[[Sample Input Transformer]]" />
    <entry key="description" value="[[A new transformer for metocard input.]]" />
  </service-properties>
</service>
...
```

Table 73. Input Transformer Variable Descriptions / Blueprint Service Properties

Key	Description of Value	Example
shortname	(Required) An abbreviation for the return-type of the <code>BinaryContent</code> being sent to the user.	atom
title	(Optional) A user-readable title that describes (in greater detail than the shortname) the service.	Atom Entry Transformer Service
description	(Optional) A short, human-readable description that describes the functionality of the service and the output.	This service converts a single metocard xml document to an atom entry element.

5. Deploy OSGi Bundle to OSGi runtime.


29.8.1. Create an XML Input Transformer using SaxEventHandlers

For a transformer to transform XML, (as opposed to JSON or a Word document, for example) there is a simpler solution than fully implementing a `MetacardValidator`. DDF includes an extensible, configurable `XmlInputTransformer`. This transformer can be instantiated via blueprint as a managed service factory and configured via metatype. The `XmlInputTransformer` takes a configuration of `SaxEventHandlers`. A `SaxEventHandler` is a class that handles SAX Events (a very fast XML parser) to parse metadata and create metacards. Any number of `SaxEventHandlers` can be implemented and included in the `XmlInputTransformer` configuration. See the `catalog-transformer-streaming-impl` bundle for examples (`XmlSaxEventHandlerImpl` which parses the DDF Metacard XML Metadata and the `GmlHandler` which parses GML 2.0) Each `SaxEventHandler` implementation has a `SaxEventHandlerFactory` associated with it. The `SaxEventHandlerFactory` is responsible for instantiating new `SaxEventHandlers` - each transform request gets a new instance of `XmlInputTransformer` and set of `SaxEventHandlers` to be *thread- and state-safe*.


The following diagrams intend to clarify implementation details:

The `XmlInputTransformer` Configuration diagram shows the `XmlInputTransformer` configuration, which is configured using the metatype and has the `SaxEventHandlerFactory` ids. Then, when a transform request is received, the `ManagedServiceFactory` instantiates a new `XmlInputTransformer`. This `XmlInputTransformer` then instantiates a new `SaxEventHandlerDelegate` with the configured

`SaxEventHandlersFactory` ids. The factories all in turn instantiate a `SaxEventHandler`. Then, the `SaxEventHandlerDelegate` begins parsing the XML input document, handing the SAX Events off to each `SaxEventHandler`, which handle them if they can. After parsing is finished, each `SaxEventHandler` returns a list of `Attributes` to the `SaxEventHandlerDelegate` and `XmlInputTransformer` which add the attributes to the metocard and then return the fully constructed metocard.



`XMLInputTransformer Configuration`



`XMLInputTransformer SaxEventHandlerDelegate Configuration`

For more specific details, see the Javadoc for the `org.codice.ddf.transformer.xml.streaming.*` package. Additionally, see the source code for the `org.codice.ddf.transformer.xml.streaming.impl.GmlHandler.java`, `org.codice.ddf.transformer.xml.streaming.impl.GmlHandlerFactory`, `org.codice.ddf.transformer.xml.streaming.impl.XmlInputTransformerImpl`, and `org.codice.ddf.transformer.xml.streaming.impl.XmlInputTransformerImplFactory`.

- NOTE**
1. The `XmlInputTransformer` & `SaxEventHandlerDelegate` create and configure themselves based on String matches of the configuration ids with the `SaxEventHandlerFactory` ids, so ensure these match.
 2. The `XmlInputTransformer` uses a `DynamicMetacardType`. This is pertinent because a metacards attributes are only stored in Solr if they are declared on the `MetacardType`. Since the `DynamicMetacardType` is constructed dynamically, attributes are declared by the `SaxEventHandlerFactory` that parses them, as opposed to the `MetacardType`. See `org.codice.ddf.transformer.xml.streaming.impl.XmlSaxEventHandlerFactoryImpl.java` vs `ddf.catalog.data.impl.BasicTypes.java`

29.8.2. Create an Input Transformer Using Apache Camel

Alternatively, make an Apache Camel route in a blueprint file and deploy it using a feature file or via hot deploy.

29.8.2.1. Input Transformer Design Pattern (Camel)

Follow this design pattern for compatibility:

From

When using `from`, `catalog:inputtransformer?id=text/xml`, an Input Transformer will be created and registered in the OSGi registry with an id of `text/xml`.

To

When using `to`, `catalog:inputtransformer?id=text/xml`, an Input Transformer with an id matching `text/xml` will be discovered from the OSGi registry and invoked.

Table 74. InputTransformer Message Formats

Exchange Type	Field	Type
Request (comes from <code><from></code> in the route)	body	<code>java.io.InputStream</code>
Response (returned after called via <code><to></code> in the route)	body	<code>ddf.catalog.data.Metacard</code>

TIP Its always a good idea to wrap the `MimeType` value with the `RAW` parameter as shown in the example above. This will ensure that the value is taken exactly as is, and is especially useful when you are using special characters.

InputTransformer Creation Example

```
<blueprint xmlns="http://www.osgi.org/xmlns/blueprint/v1.0.0">
  <camelContext xmlns="http://camel.apache.org/schema/blueprint">
    <route>
      <from uri="catalog:inputtransformer?mimeType=RAW(id=text/xml;id=vehicle)" />
      <to uri="xslt:vehicle.xslt" /> <!-- must be on classpath for this bundle
-->
      <to uri=
"catalog:inputtransformer?mimeType=RAW(id=application/json;id=geojson)" />
    </route>
  </camelContext>
</blueprint>
```

InputTransformer Creation Details

1. Defines this as an Apache Aries blueprint file.
2. Defines the Apache Camel context that contains the route.
3. Defines start of an Apache Camel route.
4. Defines the endpoint/consumer for the route. In this case it is the DDF custom catalog component that is an `InputTransformer` registered with an id of `text/xml;id=vehicle` meaning it can transform an `InputStream` of vehicle data into a metocard. **Note that the specified XSL stylesheet must be on the classpath of the bundle that this blueprint file is packaged in.**
5. Defines the XSLT to be used to transform the vehicle input into GeoJSON format using the Apache Camel provided XSLT component.
6. Defines the route node that accepts GeoJSON formatted input and transforms it into a Metocard, using the DDF custom catalog component that is an `InputTransformer` registered with an id of `application/json;id=geojson`.

NOTE An example of using an Apache Camel route to define an `InputTransformer` in a blueprint file and deploying it as a bundle to an OSGi container can be found in the DDF SDK examples at [DDF/sdk/sample-transformers/xslt-identity-input-transformer](#)

29.9. Developing Metocard Transformers

In general, a `MetocardTransformer` is used to transform a `Metocard` into some desired format useful to the end user or as input to another process. Programmatically, a `MetocardTransformer` transforms a `Metocard` into a `BinaryContent` instance, which translates the `Metocard` into the desired final format. Metocard transformers can be used through the Catalog Framework `transform` convenience method or requested from the OSGi Service Registry by endpoints or other bundles.

29.9.1. Creating a New Metocard Transformer

Existing metocard transformers are written as Java classes, and these steps walk through the steps to create a custom metocard transformer.

1. Create a new Java class that implements `ddf.catalog.transform.MetocardTransformer`.


```
public class SampleMetocardTransformer implements ddf.catalog.transform.MetocardTransformer
```
2. Implement the `transform` method.


```
public BinaryContent transform(Metocard metocard, Map<String, Serializable> arguments)
throws CatalogTransformerException
```

 - a. `transform` must return a `Metocard` or throw an exception. It cannot return null.
3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).


```
Import-Package: ddf.catalog,ddf.catalog.transform
```
4. Create an OSGi descriptor file to communicate with the OSGi Service registry (described in the [Working with OSGi section](#)). Export the service to the OSGi registry and declare service properties.

Metocard Transformer Blueprint Descriptor Example

```
...
<service ref="[[SampleMetocardTransformer]]" interface=
"ddf.catalog.transform.MetocardTransformer">
  <service-properties>
    <entry key="shortname" value="[[sampletransform]]" />
    <entry key="title" value="[[Sample Metocard Transformer]]" />
    <entry key="description" value="[[A new transformer for metacards.]]" />
  </service-properties>
</service>
...
```

5. Deploy OSGi Bundle to OSGi runtime.

Table 75. Metocard Transformer Blueprint Service Properties / Variable Descriptions

Key	Description of Value	Example
<code>shortname</code>	(Required) An abbreviation for the return type of the <code>BinaryContent</code> being sent to the user.	atom
<code>title</code>	(Optional) A user-readable title that describes (in greater detail than the <code>shortname</code>) the service.	Atom Entry Transformer Service
<code>description</code>	(Optional) A short, human-readable description that describes the functionality of the service and the output.	This service converts a single metocard xml document to an atom entry element.

29.10. Developing Query Response Transformers

A `QueryResponseTransformer` is used to transform a List of Results from a `SourceResponse`. Query

Response Transformers can be used through the Catalog transform convenience method or requested from the OSGi Service Registry by endpoints or other bundles.

1. Create a new Java class that implements `ddf.catalog.transform.QueryResponseTransformer`.

```
public class SampleResponseTransformer implements  
ddf.catalog.transform.QueryResponseTransformer
```

2. Implement the `transform` method.

```
public BinaryContent transform(SourceResponse upstreamResponse, Map<String, Serializable>  
arguments) throws CatalogTransformerException
```

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog, ddf.catalog.transform
```

4. Create an OSGi descriptor file to communicate with the OSGi Service Registry (described in [Working with OSGi](#)). Export the service to the OSGi registry and declare service properties.

5. Deploy OSGi Bundle to OSGi runtime.

Query Response Transformer Blueprint Descriptor Example


```
...  
<service ref="[[SampleResponseTransformer]]" interface=  
"ddf.catalog.transform.QueryResponseTransformer">  
    <service-properties>  
        <entry key="shortname" value="[[sampletransform]]" />  
        <entry key="title" value="[[Sample Response Transformer]]" />  
        <entry key="description" value="[[A new transformer for response queues.]]"  
    />  
    </service-properties>  
</service>  
...
```

Table 76. Query Response Transformer Blueprint Service Properties / Variable Descriptions

Key	Description of Value	Example
<code>shortname</code>	An abbreviation for the return type of the <code>BinaryContent</code> being sent to the user.	atom
<code>title</code>	A user-readable title that describes (in greater detail than the <code>shortname</code>) the service.	Atom Entry Transformer Service
<code>description</code>	A short, human-readable description that describes the functionality of the service and the output.	<i>This service converts a single metocard xml document to an atom entry element.</i>

29.11. Developing Sources

Sources are components that enable DDF to talk to back-end services. They let DDF perform query and ingest operations on catalog stores and query operations on federated sources.



Source Architecture

29.11.1. Implement a Source Interface

There are three types of sources that can be created to perform query operations. All of these sources must also be able to return their availability and the list of content types currently stored in their back-end data stores.

Catalog Provider

`ddf.catalog.source.CatalogProvider` is used to communicate with back-end storage and allows for Query and Create/Update/Delete operations.

Federated Source

`ddf.catalog.source.FederatedSource` is used to communicate with remote systems and only allows query operations.

Connected Source

`ddf.catalog.source.ConnectedSource` is similar to a Federated Source with the following exceptions:

- Queried on all local queries
- `SiteName` is hidden (masked with the DDF sourceId) in query results
- `SiteService` does not show this Source's information separate from DDF's.

Catalog Store

`catalog.store.interface` is used to store data.

The procedure for implementing any of the source types follows a similar format:

1. Create a new class that implements the specified Source interface, the [ConfiguredService](#) and the required methods.
2. Create an OSGi descriptor file to communicate with the OSGi registry. (Refer to [OSGi Services](#).)
 - a. Import DDF packages.
 - b. Register source class as service to the OSGi registry.
3. Deploy to DDF.

IMPORTANT

The `factory-pid` property of the metatype must contain one of the following in the name: service, Service, source, Source

NOTE

Remote sources currently extend the `ResourceReader` interface. However, a `RemoteSource` is not treated as a `ResourceReader`. The `getSupportedSchemes()` method should never be called on a `RemoteSource`, thus the suggested implementation for a `RemoteSource` is to return an empty set. The `retrieveResource(...)` and `getOptions(...)` methods will be called and MUST be properly implemented by a `RemoteSource`.

29.11.1.1. Developing Catalog Providers

Create a custom implementation of a catalog provider.

1. Create a Java class that implements `CatalogProvider`.

```
public class TestCatalogProvider implements ddf.catalog.source.CatalogProvider
```

2. Implement the required methods from the `ddf.catalog.source.CatalogProvider` interface.

```
public CreateResponse create(CreateRequest createRequest) throws IngestException; public UpdateResponse update(UpdateRequest updateRequest) throws IngestException; public DeleteResponse delete(DeleteRequest deleteRequest) throws IngestException;
```

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog, ddf.catalog.source
```

4. Export the service to the OSGi registry.

Catalog Provider Blueprint example

```
<service ref="[[TestCatalogProvider]]" interface="ddf.catalog.source.CatalogProvider"
/>
```

See the [existing Catalog Provider list](#) for examples of Catalog Providers included in DDF.

29.11.1.2. Developing Federated Sources

1. Create a Java class that implements `FederatedSource` and `ConfiguredService`.

```
public class TestFederatedSource implements ddf.catalog.source.FederatedSource,
ddf.catalog.service.ConfiguredService
```

2. Implement the required methods of the `ddf.catalog.source.FederatedSource` and `ddf.catalog.service.ConfiguredService` interfaces.

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

Import-Package: `ddf.catalog, ddf.catalog.source`

4. Export the service to the OSGi registry.

Federated Source Blueprint example

```
<service ref="[[TestFederatedSource]]" interface="ddf.catalog.source.FederatedSource">
```

NOTE

A code example of a Federated Source delivered with DDF can be found in [Solr Cache Source](#)

29.11.1.3. Developing Connected Sources

Create a custom implementation of a connected source.

1. Create a Java class that implements `ConnectedSource` and `ConfiguredService`.

```
public class TestConnectedSource implements ddf.catalog.source.ConnectedSource,  
ddf.catalog.service.ConfiguredService
```

2. Implement the required methods of the `ddf.catalog.source.ConnectedSource` and `ddf.catalog.service.ConfiguredService` interfaces.

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

Import-Package: `ddf.catalog, ddf.catalog.source`

4. Export the service to the OSGi registry.

Connected Source Blueprint example

```
<service ref="[[TestConnectedSource]]" interface="ddf.catalog.source.ConnectedSource">
```

IMPORTANT

In some Providers that are created, there is a need to make Web Service calls through JAXB clients. It is best to NOT create a JAXB client as a global variable. There may be intermittent failures with the creation of Providers and federated sources when clients are created in this manner. To avoid this issue, create any JAXB within the methods requiring it.

29.11.1.4. Exception Handling

In general, sources should only send information back related to the call, not implementation details.

29.11.1.4.1. Exception Examples

Follow these guidelines for effective exception handling:

- Use a "Site XYZ not found" message rather than the full stack trace with the original site not found exception.
- If the caller issues a malformed search request, return an error describing the right form, or specifically what was not recognized in the request. Do not return the exception and stack trace where the parsing broke.
- If the caller leaves something out, do not return the null pointer exception with a stack trace, rather return a generic exception with the message "xyz was missing."

29.11.1.4.2. External Resources for Developing Sources

- [Three Rules for Effective Exception Handling](#)

29.12. Developing Catalog Plugins

Plugins extend the functionality of the Catalog Framework by performing actions at specified times during a transaction. Plugin interfaces are located in the Catalog Core API. By implementing a plugin interface, actions can be performed at the desired time.

The following types of plugins can be created:

Table 77. Plugin Interfaces

Plugin Type	Plugin Interface	Invocation Order
Pre-Authorization	<code>ddf.catalog.plugin.PreAuthorizationPlugin</code>	Before any security rules are applied.
Policy	<code>ddf.catalog.plugin.PolicyPlugin</code>	After pre-authorization plugins, but before other catalog plugins to establish the policy for requests/responses.
Access	<code>ddf.catalog.plugin.AccessPlugin</code>	Directly after any policy plugins
Pre-Ingest	<code>ddf.catalog.plugin.PreIngestPlugin</code>	Before the Create/Update/Delete method is sent to the Catalog Provider.
Post-Ingest	<code>ddf.catalog.plugin.PostIngestPlugin</code>	After the Create/Update/Delete method is sent to the Catalog Provider.
Pre-Query	<code>ddf.catalog.plugin.PreQueryPlugin</code>	Prior to the Query/Read method being sent to the Source.
Post-Query	<code>ddf.catalog.plugin.PostQueryPlugin</code>	After results have been retrieved from the query but before they are posted to the Endpoint.

Plugin Type	Plugin Interface	Invocation Order
Pre-Federated-Query	<code>ddf.catalog.plugin.PreFederatedQueryPlugin</code>	Before a federated query is executed.
Post-Federated-Query	<code>ddf.catalog.plugin.PostFederatedQueryPlugin</code>	After a federated query has been executed.
Pre-Resource	<code>ddf.catalog.plugin.PreResourcePlugin</code>	Prior to a Resource being retrieved.
Post-Resource	<code>ddf.catalog.plugin.PostResourcePlugin</code>	After a Resource is retrieved, but before it is sent to the Endpoint.
Pre-Create Storage	<code>ddf.catalog.content.plugin.PreCreateStoragePlugin</code>	Experimental Before an item is created in the content repository.
Post-Create Storage	<code>ddf.catalog.content.plugin.PostCreateStoragePlugin</code>	Experimental After an item is created in the content repository.
Pre-Update Storage	<code>ddf.catalog.content.plugin.PreUpdateStoragePlugin</code>	Experimental Before an item is updated in the content repository.
Post-Update Storage	<code>ddf.catalog.content.plugin.PostUpdateStoragePlugin</code>	Experimental After an item is updated in the content repository.
Pre-Subscription	<code>ddf.catalog.plugin.PreSubscriptionPlugin</code>	Prior to a Subscription being created or updated.
Pre-Delivery	<code>ddf.catalog.plugin.PreDeliveryPlugin</code>	Prior to the delivery of a Metocard when an event is posted.

29.12.1. Implementing Catalog Plugins

The procedure for implementing any of the plugins follows a similar format:

1. Create a new class that implements the specified plugin interface.
2. Implement the required methods.
3. Create an OSGi descriptor file to communicate with the OSGi registry.
 - a. Register the plugin class as a service to OSGi registry.
4. Deploy to DDF.

Plugin Performance Concerns

NOTE

Plugins should include a check to determine if requests are local or not. It is usually preferable to take no action on non-local requests.

TIP

Refer to the Javadoc for more information on all Requests and Responses in the `ddf.catalog.operation` and `ddf.catalog.event` packages.

29.12.1.1. Catalog Plugin Failure Behavior

In the event that this Catalog Plugin cannot operate but does not wish to fail the transaction, a `PluginExecutionException` should be thrown. If processing is to be explicitly stopped, a `StopProcessingException` should be thrown. For any other exceptions, the Catalog should "fail fast" and cancel the Operation.

29.12.1.2. Implementing Pre-Ingest Plugins

Develop a custom Pre-Ingest Plugin.

1. Create a Java class that implements `PreIngestPlugin`.

```
public class SamplePreIngestPlugin implements ddf.catalog.plugin.PreIngestPlugin
```

2. Implement the required methods.

- `public CreateRequest process(CreateRequest input) throws PluginExecutionException, StopProcessingException;`
- `public UpdateRequest process(UpdateRequest input) throws PluginExecutionException, StopProcessingException;`
- `public DeleteRequest process(DeleteRequest input) throws PluginExecutionException, StopProcessingException;`

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog,ddf.catalog.plugin
```

4. Export the service to the OSGi registry.

```
Blueprint descriptor example <service ref=" "
interface="ddf.catalog.plugin.PreIngestPlugin" />
```

29.12.1.3. Implementing Post-Ingest Plugins

Develop a custom Post-Ingest Plugin.

1. Create a Java class that implements `PostIngestPlugin`.

```
public class SamplePostIngestPlugin implements ddf.catalog.plugin.PostIngestPlugin
```

2. Implement the required methods.

- `public CreateResponse process(CreateResponse input) throws PluginExecutionException;`
- `public UpdateResponse process(UpdateResponse input) throws PluginExecutionException;`
- `public DeleteResponse process(DeleteResponse input) throws PluginExecutionException;`

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog,ddf.catalog.plugin
```

4. Export the service to the OSGi registry.

```
Blueprint descriptor example <service ref=" "
interface="ddf.catalog.plugin.PostIngestPlugin" />
```

29.12.1.4. Implementing Pre-Query Plugins

Develop a custom Pre-Query Plugin

1. Create a Java class that implements `PreQueryPlugin`.

```
public class SamplePreQueryPlugin implements ddf.catalog.plugin.PreQueryPlugin
```

2. Implement the required method.

```
public QueryRequest process(QueryRequest input) throws PluginExecutionException,  
StopProcessingException;
```

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog,ddf.catalog.plugin
```

4. Export the service to the OSGi registry.

```
<service ref=" " interface="ddf.catalog.plugin.PreQueryPlugin" />
```

29.12.1.5. Implementing Post-Query Plugins

Develop a custom Post-Query Plugin

1. Create a Java class that implements `PostQueryPlugin`.

```
public class SamplePostQueryPlugin implements ddf.catalog.plugin.PostQueryPlugin
```

2. Implement the required method.

```
public QueryResponse process(QueryResponse input) throws PluginExecutionException,  
StopProcessingException;
```

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog,ddf.catalog.plugin
```

4. Export the service to the OSGi registry.

```
<service ref=" " interface="ddf.catalog.plugin.PostQueryPlugin" />
```

29.12.1.6. Implementing Pre-Delivery Plugins

Develop a custom Pre-Delivery Plugin.

1. Create a Java class that implements `PreDeliveryPlugin`.

```
public class SamplePreDeliveryPlugin implements ddf.catalog.plugin.PreDeliveryPlugin
```

2. Implement the required methods.

```
public Metocard processCreate(Metocard metocard) throws PluginExecutionException,  
StopProcessingException; public Update processUpdateMiss(Update update) throws  
PluginExecutionException, StopProcessingException;  
◦ public Update processUpdateHit(Update update) throws PluginExecutionException,  
StopProcessingException;  
◦ public Metocard processCreate(Metocard metocard) throws PluginExecutionException,  
StopProcessingException;
```

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog,ddf.catalog.plugin,ddf.catalog.operation,ddf.catalog.event
```

4. Export the service to the OSGi registry.

Blueprint descriptor example

```
<service ref=" " interface="ddf.catalog.plugin.PreDeliveryPlugin" />
```

29.12.1.7. Implementing Pre-Subscription Plugins

Develop a custom Pre-Subscription Plugin.

1. Create a Java class that implements `PreSubscriptionPlugin`.

```
`public class SamplePreSubscriptionPlugin implements  
ddf.catalog.plugin.PreSubscriptionPlugin
```

2. Implement the required method.

- `public Subscription process(Subscription input) throws PluginExecutionException,
StopProcessingException;`

29.12.1.8. Implementing Pre-Resource Plugins

Develop a custom Pre-Resource Plugin.

1. Create a Java class that implements `PreResourcePlugin`. `public class SamplePreResourcePlugin
implements ddf.catalog.plugin.PreResourcePlugin`

2. Implement the required method.

- `public ResourceRequest process(ResourceRequest input) throws PluginExecutionException,
StopProcessingException;`

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

`Import-Package: ddf.catalog,ddf.catalog.plugin,ddf.catalog.operation`

4. Export the service to the OSGi registry. `.Blueprint descriptor example`

```
<service ref="[[SamplePreResourcePlugin]]"  
interface="ddf.catalog.plugin.PreResourcePlugin" />
```

29.12.1.9. Implementing Post-Resource Plugins

Develop a custom Post-Resource Plugin.

1. Create a Java class that implements `PostResourcePlugin`.

```
public class SamplePostResourcePlugin implements ddf.catalog.plugin.PostResourcePlugin
```

2. Implement the required method.

- `public ResourceResponse process(ResourceResponse input) throws PluginExecutionException,
StopProcessingException;`

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

`Import-Package: ddf.catalog,ddf.catalog.plugin,ddf.catalog.operation`

4. Export the service to the OSGi registry.

Blueprint descriptor example

```
<]]> inter"[[SamplePostResourcePlugin]]"
interface="ddf.catalog.plugin.PostResourcePlugin" />
```

29.12.1.10. Implementing Policy Plugins

Develop a custom Policy Plugin.

1. Create a Java class that implements `PolicyPlugin`.

```
public class SamplePolicyPlugin implements ddf.catalog.plugin.PolicyPlugin
```

2. Implement the required methods.

- `PolicyResponse processPreCreate(Metocard input, Map<String, Serializable> properties) throws StopProcessingException;`
- `PolicyResponse processPreUpdate(Metocard input, Map<String, Serializable> properties) throws StopProcessingException;`
- `PolicyResponse processPreDelete(String attributeName, List<Serializable> attributeValues, Map<String, Serializable> properties) throws StopProcessingException;`
- `PolicyResponse processPreQuery(Query query, Map<String, Serializable> properties) throws StopProcessingException;`
- `PolicyResponse processPostQuery(Result input, Map<String, Serializable> properties) throws StopProcessingException;`

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog,ddf.catalog.plugin,ddf.catalog.operation
```

4. Export the service to the OSGi registry.

Blueprint descriptor example

```
<]]> inter" " interface="ddf.catalog.plugin.PolicyPlugin" />
```

29.12.1.11. Implementing Access Plugins

Develop a custom Access Plugin.

1. Create a Java class that implements `AccessPlugin`.

```
public class SamplePostResourcePlugin implements ddf.catalog.plugin.AccessPlugin
```

2. Implement the required methods.

- `CreateRequest processPreCreate(CreateRequest input) throws StopProcessingException;`
- `UpdateRequest processPreUpdate(UpdateRequest input) throws StopProcessingException;`
- `DeleteRequest processPreDelete(DeleteRequest input) throws StopProcessingException;`
- `QueryRequest processPreQuery(QueryRequest input) throws StopProcessingException;`
- `QueryResponse processPostQuery(QueryResponse input) throws StopProcessingException;`

3. Import the DDF interface packages to the bundle manifest (in addition to any other required packages).

```
Import-Package: ddf.catalog,ddf.catalog.plugin,ddf.catalog.operation
```

4. Export the service to the OSGi registry.

Blueprint descriptor example

```
<]]> inter" " interface="ddf.catalog.plugin.AccessPlugin" />
```

29.13. Developing Token Validators

Token validators are used by the Security Token Service (STS) to validate incoming token requests. The `TokenValidator` CXF interface must be implemented by all custom token validators. The `canHandleToken` and `validateToken` methods must be overridden. The `canHandleToken` method should return true or false based on the `ValueType` value of the token that the validator is associated with. The validator may be able to handle any number of different tokens that you specify. The `validateToken` method returns a `TokenValidatorResponse` object that contains the `Principal` of the identity being validated and also validates the `ReceivedToken` object collected from the RST (`RequestSecurityToken`) message.

29.14. Developing STS Claims Handlers

Develop a custom claims handler to retrieve attributes from an external attribute store.

A claim is an additional piece of data about a subject that can be included in a token along with basic token data. A claims manager provides hooks for a developer to plug in claims handlers to ensure that the STS includes the specified claims in the issued token.

The following steps define the procedure for adding a custom claims handler to the STS.

1. The new claims handler must implement the `org.apache.cxf.sts.claims.ClaimsHandler` interface.

```
/**
 * Licensed to the Apache Software Foundation (ASF) under one
 * or more contributor license agreements. See the NOTICE file
 * distributed with this work for additional information
 * regarding copyright ownership. The ASF licenses this file
 * to you under the Apache License, Version 2.0 (the
 * "License"); you may not use this file except in compliance
 * with the License. You may obtain a copy of the License at
 *
 * http://www.apache.org/licenses/LICENSE-2.0
 *
 * Unless required by applicable law or agreed to in writing,
 * software distributed under the License is distributed on an
 * "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY
 * KIND, either express or implied. See the License for the
 * specific language governing permissions and limitations
 * under the License.
 */
```

```
package org.apache.cxf.sts.claims;

import java.net.URI;
import java.util.List;

/**
 * This interface provides a pluggable way to handle Claims.
 */
public interface ClaimsHandler {

    List<URI> getSupportedClaimTypes();

    ClaimCollection retrieveClaimValues(RequestClaimCollection claims,
    ClaimsParameters parameters);

}
```

2. Expose the new claims handler as an OSGi service under the `org.apache.cxf.sts.claims.ClaimsHandler` interface.

```

<?xml version="1.0" encoding="UTF-8"?>
<blueprint xmlns="http://www.osgi.org/xmlns/blueprint/v1.0.0">

    <bean id="CustomClaimsHandler" class=
"security.sts.claimsHandler.CustomClaimsHandler" />

    <service ref="customClaimsHandler" interface=
"org.apache.cxf.sts.claims.ClaimsHandler"/>

</blueprint>

```

3. Deploy the bundle.

If the new claims handler is hitting an external service that is secured with SSL/TLS, a developer may need to add the root CA of the external site to the DDF trustStore and add a valid certificate into the DDF keyStore. For more information on certificates, refer to [\[Configuring a Java Keystore for Secure Communications\]](#).

NOTE

This XML file is found inside of the STS bundle and is named `ws-trust-1.4-service.wsdl`.

STS WS-Trust WSDL Document

```

<?xml version="1.0" encoding="UTF-8"?>
<wsdl:definitions xmlns:tns="http://docs.oasis-open.org/ws-sx/ws-trust/200512/" xmlns:wstrust="http://docs.oasis-open.org/ws-sx/ws-trust/200512/" xmlns:wsdl="http://schemas.xmlsoap.org/wsdl/" xmlns:soap="http://schemas.xmlsoap.org/wsdl/soap/" xmlns:wsap10="http://www.w3.org/2006/05/addressing/wsdl" xmlns:wsu="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-utility-1.0.xsd" xmlns: wsp="http://www.w3.org/ns/ws-policy" xmlns:wst="http://docs.oasis-open.org/ws-sx/ws-trust/200512" xmlns:xs="http://www.w3.org/2001/XMLSchema" xmlns:wsam="http://www.w3.org/2007/05/addressing/metadata" targetNamespace="http://docs.oasis-open.org/ws-sx/ws-trust/200512">
    <wsdl:types>
        <xsschema elementFormDefault="qualified" targetNamespace="http://docs.oasis-open.org/ws-sx/ws-trust/200512">
            <xselement name="RequestSecurityToken" type="wst:AbstractRequestSecurityTokenType"/>
            <xselement name="RequestSecurityTokenResponse" type="wst:AbstractRequestSecurityTokenType"/>
            <xsccomplexType name="AbstractRequestSecurityTokenType">
                <xsssequence>
                    <xssany namespace="##any" processContents="lax" minOccurs="0" maxOccurs="unbounded"/>
                </xsssequence>
                <xssattribute name="Context" type="xs:anyURI" use="optional"/>
                <xssanyAttribute namespace="##other" processContents="lax"/>
            </xsccomplexType>
            <xselement name="RequestSecurityTokenCollection" type="

```

```

"wst:RequestSecurityTokenCollectionType"/>
    <xs:complexType name="RequestSecurityTokenCollectionType">
        <xs:sequence>
            <xs:element name="RequestSecurityToken" type=
"wst:AbstractRequestSecurityTokenType" minOccurs="2" maxOccurs="unbounded"/>
        </xs:sequence>
    </xs:complexType>
    <xs:element name="RequestSecurityTokenResponseCollection" type=
"wst:RequestSecurityTokenResponseCollectionType"/>
    <xs:complexType name="RequestSecurityTokenResponseCollectionType">
        <xs:sequence>
            <xs:element ref="wst:RequestSecurityTokenResponse" minOccurs="1"
maxOccurs="unbounded"/>
        </xs:sequence>
        <xs:anyAttribute namespace="#other" processContents="lax"/>
    </xs:complexType>
</xs:schema>
</wsdl:types>
<!-- WS-Trust defines the following GEDs -->
<wsdl:message name="RequestSecurityTokenMsg">
    <wsdl:part name="request" element="wst:RequestSecurityToken"/>
</wsdl:message>
<wsdl:message name="RequestSecurityTokenResponseMsg">
    <wsdl:part name="response" element="wst:RequestSecurityTokenResponse"/>
</wsdl:message>
<wsdl:message name="RequestSecurityTokenCollectionMsg">
    <wsdl:part name="requestCollection" element=
"wst:RequestSecurityTokenCollection"/>
</wsdl:message>
<wsdl:message name="RequestSecurityTokenResponseCollectionMsg">
    <wsdl:part name="responseCollection" element=
"wst:RequestSecurityTokenResponseCollection"/>
</wsdl:message>
<!-- This portType an example of a Requestor (or other) endpoint that
     Accepts SOAP-based challenges from a Security Token Service -->
<wsdl:portType name="WSSecurityRequestor">
    <wsdl:operation name="Challenge">
        <wsdl:input message="tns:RequestSecurityTokenResponseMsg"/>
        <wsdl:output message="tns:RequestSecurityTokenResponseMsg"/>
    </wsdl:operation>
</wsdl:portType>
<!-- This portType is an example of an STS supporting full protocol -->
<wsdl:portType name="STS">
    <wsdl:operation name="Cancel">
        <wsdl:input wsam:Action="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RST/Cancel" message="tns:RequestSecurityTokenMsg"/>
        <wsdl:output wsam:Action="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RSTR/CancelFinal" message="tns:RequestSecurityTokenResponseMsg"/>
    </wsdl:operation>
    <wsdl:operation name="Issue">
        <wsdl:input wsam:Action="http://docs.oasis-open.org/ws-sx/ws-

```

```

trust/200512/RST/Issue" message="tns:RequestSecurityTokenMsg"/>
    <wsdl:output wsam:Action="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RSTRC/IssueFinal" message="tns:RequestSecurityTokenResponseCollectionMsg
"/>
    </wsdl:operation>
    <wsdl:operation name="Renew">
        <wsdl:input wsam:Action="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RST/Renew" message="tns:RequestSecurityTokenMsg"/>
        <wsdl:output wsam:Action="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RSTR/RenewFinal" message="tns:RequestSecurityTokenResponseMsg"/>
    </wsdl:operation>
    <wsdl:operation name="Validate">
        <wsdl:input wsam:Action="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RST/Validate" message="tns:RequestSecurityTokenMsg"/>
        <wsdl:output wsam:Action="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RSTR/ValidateFinal" message="tns:RequestSecurityTokenResponseMsg"/>
    </wsdl:operation>
    <wsdl:operation name="KeyExchangeToken">
        <wsdl:input wsam:Action="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RST/KET" message="tns:RequestSecurityTokenMsg"/>
        <wsdl:output wsam:Action="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RSTR/KETFinal" message="tns:RequestSecurityTokenResponseMsg"/>
    </wsdl:operation>
    <wsdl:operation name="RequestCollection">
        <wsdl:input message="tns:RequestSecurityTokenCollectionMsg"/>
        <wsdl:output message="tns:RequestSecurityTokenResponseCollectionMsg"/>
    </wsdl:operation>
</wsdl:portType>
<!-- This portType is an example of an endpoint that accepts
    Unsolicited RequestSecurityTokenResponse messages -->
<wsdl:portType name="SecurityTokenResponseService">
    <wsdl:operation name="RequestSecurityTokenResponse">
        <wsdl:input message="tns:RequestSecurityTokenResponseMsg"/>
    </wsdl:operation>
</wsdl:portType>
<wsdl:binding name="STS_Binding" type="wstrust:STS">
    <wsp:PolicyReference URI="#STS_policy"/>
    <soap:binding style="document" transport="
http://schemas.xmlsoap.org/soap/http"/>
    <wsdl:operation name="Issue">
        <soap:operation soapAction="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RST/Issue"/>
        <wsdl:input>
            <soap:body use="literal"/>
        </wsdl:input>
        <wsdl:output>
            <soap:body use="literal"/>
        </wsdl:output>
    </wsdl:operation>
    <wsdl:operation name="Validate">
        <soap:operation soapAction="http://docs.oasis-open.org/ws-sx/ws-

```

```

trust/200512/RST/Validate" />
    <wsdl:input>
        <soap:body use="literal"/>
    </wsdl:input>
    <wsdl:output>
        <soap:body use="literal"/>
    </wsdl:output>
</wsdl:operation>
<wsdl:operation name="Cancel">
    <soap:operation soapAction="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RST/Cancel" />
    <wsdl:input>
        <soap:body use="literal"/>
    </wsdl:input>
    <wsdl:output>
        <soap:body use="literal"/>
    </wsdl:output>
</wsdl:operation>
<wsdl:operation name="Renew">
    <soap:operation soapAction="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RST/Renew" />
    <wsdl:input>
        <soap:body use="literal"/>
    </wsdl:input>
    <wsdl:output>
        <soap:body use="literal"/>
    </wsdl:output>
</wsdl:operation>
<wsdl:operation name="KeyExchangeToken">
    <soap:operation soapAction="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RST/KeyExchangeToken" />
    <wsdl:input>
        <soap:body use="literal"/>
    </wsdl:input>
    <wsdl:output>
        <soap:body use="literal"/>
    </wsdl:output>
</wsdl:operation>
<wsdl:operation name="RequestCollection">
    <soap:operation soapAction="http://docs.oasis-open.org/ws-sx/ws-
trust/200512/RST/RequestCollection" />
    <wsdl:input>
        <soap:body use="literal"/>
    </wsdl:input>
    <wsdl:output>
        <soap:body use="literal"/>
    </wsdl:output>
</wsdl:operation>
</wsdl:binding>
<wsp:Policy wsu:Id="STS_policy">
    <wsp:ExactlyOne>

```

```

<wsp:All>
  <wsap10:UsingAddressing/>
  <wsp:ExactlyOne>
    <sp:TransportBinding xmlns:sp="http://docs.oasis-open.org/ws-
sx/ws-securitypolicy/200702">
      <wsp:Policy>
        <sp:TransportToken>
          <wsp:Policy>
            <sp:HttpsToken>
              <wsp:Policy/>
            </sp:HttpsToken>
          </wsp:Policy>
        </sp:TransportToken>
        <sp:AlgorithmSuite>
          <wsp:Policy>
            <sp:Basic128/>
          </wsp:Policy>
        </sp:AlgorithmSuite>
        <sp:Layout>
          <wsp:Policy>
            <sp:Lax/>
          </wsp:Policy>
        </sp:Layout>
        <sp:IncludeTimestamp/>
      </wsp:Policy>
    </sp:TransportBinding>
  </wsp:ExactlyOne>
  <sp:Wss11 xmlns:sp="http://docs.oasis-open.org/ws-sx/ws-
securitypolicy/200702">
    <wsp:Policy>
      <sp:MustSupportRefKeyIdentifier/>
      <sp:MustSupportRefIssuerSerial/>
      <sp:MustSupportRefThumbprint/>
      <sp:MustSupportRefEncryptedKey/>
    </wsp:Policy>
  </sp:Wss11>
  <sp:Trust13 xmlns:sp="http://docs.oasis-open.org/ws-sx/ws-
securitypolicy/200702">
    <wsp:Policy>
      <sp:MustSupportIssuedTokens/>
      <sp:RequireClientEntropy/>
      <sp:RequireServerEntropy/>
    </wsp:Policy>
  </sp:Trust13>
  </wsp:All>
  </wsp:ExactlyOne>
</wsp:Policy>
<wsp:Policy wsu:Id="Input_policy">
  <wsp:ExactlyOne>
    <wsp:All>
      <sp:SignedParts xmlns:sp="http://docs.oasis-open.org/ws-sx/ws-

```

```

securitypolicy/200702">
    <sp:Body/>
    <sp:Header Name="To" Namespace=
"http://www.w3.org/2005/08/addressing"/>
    <sp:Header Name="From" Namespace=
"http://www.w3.org/2005/08/addressing"/>
    <sp:Header Name="FaultTo" Namespace=
"http://www.w3.org/2005/08/addressing"/>
    <sp:Header Name="ReplyTo" Namespace=
"http://www.w3.org/2005/08/addressing"/>
    <sp:Header Name="MessageID" Namespace=
"http://www.w3.org/2005/08/addressing"/>
    <sp:Header Name="RelatesTo" Namespace=
"http://www.w3.org/2005/08/addressing"/>
    <sp:Header Name="Action" Namespace=
"http://www.w3.org/2005/08/addressing"/>
    </sp:SignedParts>
    <sp:EncryptedParts xmlns:sp="http://docs.oasis-open.org/ws-sx/ws-
securitypolicy/200702">
        <sp:Body/>
        </sp:EncryptedParts>
    </wsp:All>
    </wsp:ExactlyOne>
</wsp:Policy>
<wsp:Policy wsu:Id="Output_policy">
    <wsp:ExactlyOne>
        <wsp:All>
            <sp:SignedParts xmlns:sp="http://docs.oasis-open.org/ws-sx/ws-
securitypolicy/200702">
                <sp:Body/>
                <sp:Header Name="To" Namespace=
"http://www.w3.org/2005/08/addressing"/>
                <sp:Header Name="From" Namespace=
"http://www.w3.org/2005/08/addressing"/>
                <sp:Header Name="FaultTo" Namespace=
"http://www.w3.org/2005/08/addressing"/>
                <sp:Header Name="ReplyTo" Namespace=
"http://www.w3.org/2005/08/addressing"/>
                <sp:Header Name="MessageID" Namespace=
"http://www.w3.org/2005/08/addressing"/>
                <sp:Header Name="RelatesTo" Namespace=
"http://www.w3.org/2005/08/addressing"/>
                <sp:Header Name="Action" Namespace=
"http://www.w3.org/2005/08/addressing"/>
            </sp:SignedParts>
            <sp:EncryptedParts xmlns:sp="http://docs.oasis-open.org/ws-sx/ws-
securitypolicy/200702">
                <sp:Body/>
                </sp:EncryptedParts>
            </wsp:All>
            </wsp:ExactlyOne>
        </wsp:All>
    </sp:SignedParts>
    <sp:EncryptedParts xmlns:sp="http://docs.oasis-open.org/ws-sx/ws-
securitypolicy/200702">
        <sp:Body/>
        </sp:EncryptedParts>
    </wsp:All>
    </wsp:ExactlyOne>
</wsp:Policy>

```

```

</wsp:Policy>
<wsdl:service name="SecurityTokenService">
    <wsdl:port name="STS_Port" binding="tns:STS_Binding">
        <soap:address location=
"\http://{FQDN}:{PORT}/services/SecurityTokenService"/>
    </wsdl:port>
</wsdl:service>
</wsdl:definitions>

```

29.14.1. Example Requests and Responses for SAML Assertions

A client performs a RequestSecurityToken operation against the STS to receive a SAML assertion. The DDF STS offers several different ways to request a SAML assertion. For help in understanding the various request and response formats, samples have been provided. The samples are divided out into different request token types.

29.14.2. BinarySecurityToken (CAS) SAML Security Token Samples

Most endpoints in DDF require the X.509 PublicKey SAML assertion.

BinarySecurityToken (CAS) SAML Security Token Sample Request

```

<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/">
    <soap:Header>
        <Action xmlns="http://www.w3.org/2005/08/addressing">http://docs.oasis-
open.org/ws-sx/ws-trust/200512/RST/Issue</Action>
        <MessageID xmlns="http://www.w3.org/2005/08/addressing">urn:uuid:60652909-
faca-4e4a-a4a7-8a5ce243a7cb</MessageID>
        <To xmlns="http://www.w3.org/2005/08/addressing">
https://server:8993/services/SecurityTokenService</To>
        <ReplyTo xmlns="http://www.w3.org/2005/08/addressing">
            <Address>http://www.w3.org/2005/08/addressing/anonymous</Address>
        </ReplyTo>
        <wsse:Security xmlns:wsse="http://docs.oasis-open.org/wss/2004/01/oasis-
200401-wss-wssecurity-secext-1.0.xsd" xmlns:wsu="http://docs.oasis-
open.org/wss/2004/01/oasis-200401-wss-wssecurity-utility-1.0.xsd" soap:mustUnderstand
="1">
            <wsu:Timestamp wsu:Id="TS-1">
                <wsu:Created>2013-04-29T18:35:10.688Z</wsu:Created>
                <wsu:Expires>2013-04-29T18:40:10.688Z</wsu:Expires>
            </wsu:Timestamp>
        </wsse:Security>
    </soap:Header>
    <soap:Body>
        <wst:RequestSecurityToken xmlns:wst="http://docs.oasis-open.org/ws-sx/ws-
trust/200512">
            <wst:RequestType>http://docs.oasis-open.org/ws-sx/ws-
trust/200512/Issue</wst:RequestType>
            <wsp:AppliesTo xmlns:wsp="http://schemas.xmlsoap.org/ws/2004/09/policy">
                <wsa:EndpointReference xmlns:wsa="

```

```

http://www.w3.org/2005/08/addressing">
    <wsa:Address>
        https://server:8993/services/SecurityTokenService</wsa:Address>
            </wsa:EndpointReference>
            </wsp:AppliesTo>
            <wst:Claims xmlns:ic="http://schemas.xmlsoap.org/ws/2005/05/identity"
                xmlns:wst="http://docs.oasis-open.org/ws-sx/ws-trust/200512" Dialect=
                "http://schemas.xmlsoap.org/ws/2005/05/identity">
                <ic:ClaimType xmlns:ic="
                    http://schemas.xmlsoap.org/ws/2005/05/identity" Optional="true" Uri=
                    "http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier"/>
                    <ic:ClaimType xmlns:ic="
                        http://schemas.xmlsoap.org/ws/2005/05/identity" Optional="true" Uri=
                        "http://schemas.xmlsoap.org/ws/2005/05/identity/claims/emailaddress"/>
                        <ic:ClaimType xmlns:ic="
                            http://schemas.xmlsoap.org/ws/2005/05/identity" Optional="true" Uri=
                            "http://schemas.xmlsoap.org/ws/2005/05/identity/claims/surname"/>
                            <ic:ClaimType xmlns:ic="
                                http://schemas.xmlsoap.org/ws/2005/05/identity" Optional="true" Uri=
                                "http://schemas.xmlsoap.org/ws/2005/05/identity/claims/givenname"/>
                                <ic:ClaimType xmlns:ic="
                                    http://schemas.xmlsoap.org/ws/2005/05/identity" Optional="true" Uri=
                                    "http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role"/>
                                    </wst:Claims>
                                    <wst:OnBehalfOf>
                                        <BinarySecurityToken ValueType="#CAS" EncodingType="http://docs.oasis-
                                            open.org/wss/2004/01/oasis-200401-wss-soap-message-security-1.0#Base64Binary" ns1:Id=
                                            "CAS" xmlns="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-
                                            secext-1.0.xsd" xmlns:ns1="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-
                                            wssecurity-utility-1.0.xsd"
                                            >U1QtMTQtYUtmcDYxcFRtS0FxZG1pVDMzOWMtY2FzfGh0dHBz0i8vdG9rZW5pc3N1ZXI60Dk5My9zZXJ2aWNlc
                                            y9TzWN1cm10eVRva2VuU2VydmljZQ==</BinarySecurityToken>
                                        </wst:OnBehalfOf>
                                        <wst:TokenType>http://docs.oasis-open.org/wss/oasis-wss-saml-token-
                                            profile-1.1#SAMLV2.0</wst:TokenType>
                                        <wst:KeyType>http://docs.oasis-open.org/ws-sx/ws-
                                            trust/200512/PublicKey</wst:KeyType>
                                        <wst:UseKey>
                                            <ds:KeyInfo xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
                                                <ds:X509Data>
                                                    <ds:X509Certificate>
MIIC5DCCAk2gAwIBAgIJAKj7ROPHjo1yMA0GCSqGSIb3DQEBCwUAMIGKMQswCQYDVQQGEwJVUzEQ
MA4GA1UECAwHQXJpem9uYTERMA8GA1UEBwwIR29vZHl1YXIxGDAwBgNVBAoMD0xvY2toZWVkIE1h
cnRpbjENMAssGA1UECwwESTRDRTEPMA0GA1UEAwwGY2xpZW50MRwwGgYJKoZIhvcNAQkBFg1pNGNl
QGxtY28uY29tMB4XDTEyMDYyMDE5NDMwOVoXTDTiyMDYxODE5NDMwOVoWwgYoxCzAJBgNVBAYTA1VT
MRAwDgYDVQQIDA0Bcm16b25hMREwDwYDVQQHDAhHb29keWVhcjEYMBYGA1UECgwPTG9ja2h1ZWQg
TWFydGluMQ0wCwYDVQQLDARJNENFMQ8wDQYDVQQDDAZjbG1lbnQxHDAaBqkqhkiG9w0BCQEWDWk0
Y2VabG1jby5jb20wgZ8wDQYJKoZIhvcNAQEBBQADgY0AMIGJAoGBAIpHxCBLYE7xfDLcITS9SsPG
4Q04Z6S32/+TriGsRgpGTj/7GuMG7oJ98m6Ws5cTYl7nyunyHTkZuP7rBzy4esDIHheyx18EgdSJ
vvACgGVcNEmHndkf9bWU1A0fNaxW+vZwljUkRUVdkhPbPdPw0cMdKg/SsLSNjZfsQIjoWd4rAgMB
AAGjUDBOMB0GA1UdDgQWBBQx11VLtYXLvFGpFdHnh1NW9+1xBDAfBqNVHSMEGDAwBQx11VLtYXL
                                            </ds:X509Certificate>
                                        </ds:X509Data>
                                    </ds:KeyInfo>
                                </wst:UseKey>
                            </wst:TokenType>
                        </wst:OnBehalfOf>
                    </ic:Claims>
                </wst:SecurityTokenRequest>
            </wsp:AppliesTo>
        </wsp:Policy>
    </wsp:PolicyReference>
</wsp:PolicySet>

```

```

vFGpFdHnhLNW9+lxBDAMBgNVHRMEBTADAQH/MA0GCSqGSIb3DQEBCwUAA4GBAHYs20I0K6yVXzyS
sKcv2fmfw6XCICGTnyA7B0dAjYoqq6wD+33dHJUCFDqye7AWdcivuc7RWJt9jnlfJZKIm2BHcDTR
Hhk6CvjJ14Gf40WQdeMHoX8U8b0diq7Iy5Ravx+zRg7SdiyJUqFYjRh/05tywXRT1+freI3bwAN0
L6tQ
</ds:X509Certificate>
    </ds:X509Data>
    </ds:KeyInfo>
</wst:UseKey>
<wst:Renewing/>
</wst:RequestSecurityToken>
</soap:Body>
</soap:Envelope>

```

BinarySecurityToken (CAS) SAML Security Token Sample Response

```

<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/">
    <soap:Header>
        <Action xmlns="http://www.w3.org/2005/08/addressing">http://docs.oasis-
open.org/ws-sx/ws-trust/200512/RSTRC/IssueFinal</Action>
        <MessageID xmlns="http://www.w3.org/2005/08/addressing">urn:uuid:7a6fde04-
9013-41ef-b08b-0689ffa9c93e</MessageID>
        <To xmlns="http://www.w3.org/2005/08/addressing">
            >http://www.w3.org/2005/08/addressing/anonymous</To>
        <RelatesTo xmlns="http://www.w3.org/2005/08/addressing">urn:uuid:60652909-
faca-4e4a-a4a7-8a5ce243a7cb</RelatesTo>
        <wsse:Security xmlns:wsse="http://docs.oasis-open.org/wss/2004/01/oasis-
200401-wss-wssecurity-secext-1.0.xsd" xmlns:wsu="http://docs.oasis-
open.org/wss/2004/01/oasis-200401-wss-wssecurity-utility-1.0.xsd" soap:mustUnderstand
="1">
            <wsu:Timestamp wsu:Id="TS-2">
                <wsu:Created>2013-04-29T18:35:11.459Z</wsu:Created>
                <wsu:Expires>2013-04-29T18:40:11.459Z</wsu:Expires>
            </wsu:Timestamp>
        </wsse:Security>
    </soap:Header>
    <soap:Body>
        <RequestSecurityTokenResponseCollection xmlns="http://docs.oasis-open.org/ws-
sx/ws-trust/200512" xmlns:ns2="http://docs.oasis-open.org/wss/2004/01/oasis-200401-
wss-wssecurity-utility-1.0.xsd" xmlns:ns3="http://docs.oasis-
open.org/wss/2004/01/oasis-200401-wss-wssecurity-secext-1.0.xsd" xmlns:ns4=
"http://www.w3.org/2005/08/addressing" xmlns:ns5="http://docs.oasis-open.org/ws-sx/ws-
trust/200802">
            <RequestSecurityTokenResponse>
                <TokenType>http://docs.oasis-open.org/wss/oasis-wss-saml-token-
profile-1.1#SAMLV2.0</TokenType>
                <RequestedSecurityToken>
                    <saml2:Assertion xmlns:saml2=
"urn:oasis:names:tc:SAML:2.0:assertion" xmlns:xs="http://www.w3.org/2001/XMLSchema"
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" ID=
"_BDC44EB8593F47D1B213672605113671" IssueInstant="2013-04-29T18:35:11.370Z" Version=
"2.0" xsi:type="saml2:AssertionType">

```

```

<saml2:Issuer>tokenissuer</saml2:Issuer>
<ds:Signature xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
  <ds:SignedInfo>
    <ds:CanonicalizationMethod Algorithm="http://www.w3.org/2001/10/xml-exc-c14n#" />
    <ds:SignatureMethod Algorithm="http://www.w3.org/2000/09/xmldsig#rsa-sha1" />
    <ds:Reference URI="#_BDC44EB8593F47D1B213672605113671" />
    <ds:Transforms>
      <ds:Transform Algorithm="http://www.w3.org/2000/09/xmldsig#enveloped-signature" />
      <ds:Transform Algorithm="http://www.w3.org/2001/10/xml-exc-c14n#" />
      <ec:InclusiveNamespaces xmlns:ec="http://www.w3.org/2001/10/xml-exc-c14n#" PrefixList="xs" />
      </ds:Transform>
    </ds:Transforms>
    <ds:DigestMethod Algorithm="http://www.w3.org/2000/09/xmldsig#sha1" />
    <ds:DigestValue>6wnWbft6Pz5X0F5Q9AG59gcGwLY=</ds:DigestValue>
  </ds:Reference>
</ds:SignedInfo>

<ds:SignatureValue>h+NvkgXGdQtca3/eKebhAKg638tHp3i2n5uLLy8xXXIg02qyKgEP0FCowp2LiYlsQU9YjKfSwCUbH3WR6jhAv9zj29CE+ePfEny7MeXvgN13wId+vcHqtI/DGGhhgt02Mbx/tyX1BhHQwKR1chajxHeecwmvV7D85NMdV48tI=</ds:SignatureValue>
<ds:KeyInfo>
  <ds:X509Data>

<ds:X509Certificate>MIIDmjCCAwOgAwIBAgIBBDANBgkqhkiG9w0BAQQFADB1MQswCQYDVQQGEwJVUzEQMA4GA1UECBMHQXJpem9uYTERMA8GA1UEBxMIR29vZH11YXIxEDAOBgNVBAoTB0V4YW1wbGUxEDAOBgNVBAoTB0V4YW1wbGUxEDAOBgNVBAsTB0V4YW1wbGUxCzAJBgNVBAMTAkNBMB4XDTEzMDQwOTE4MzcxMVoXTDIZMDQwNzE4MzcxMVoWgaYxCzAJBgNVBAYTA1VTMRAwDgYDVQQIEwdBcm16b25hMREwDwYDVQQHEwhHb29keWVhcjEQMA4GA1UEChMHRXhhbXBsZTEQMA4GA1UEChMHRXhhbXBsZTEQMA4GA1UECxMHRXhhbXBsZTEUMBIGA1UEAxMLdG9rZW5pc3N1ZXIxJjAkBgkqhkiG9w0BCQEWF3Rva2VuaXNzdWVyQGV4YW1wbGUuY29tMIGfMA0GCSqGSIb3DQEBAQUAA4GNADCBiQKBgQDDfktpA8Lrp9rTfRibKdgtxtN9uB44diiIqq3J0zDGfDh6Lu6mjpuH01hrKItv42hB0hhmH71S9ipiaQCIpVfgIG63MB7fa5dBrfGF669vFrU1Lf17IvsVVsvNsrtAEQ1j0Mmw9sxS3SUsRQX+bD8jq7Uj1hpoF7DdqpV8Kb0C00GwIDAQABo4IBBjCCAQIwCQYDVR0TBAIwADAsBglghkgBvhCAQ0EHxYdT3B1b1NTTCBHZW51cmF0ZWQgQ2VydG1maWNhdGUwHQYDVR0OBBYEFD1mHviop2Tc4HaNu8yPXR6GqWP1MIGnBgnVHSMEgZ8wgZyAFBcnen6/j05DzaVw0RwrtEkc7TZo0XmkdzB1MQswCQYDVQQGEwJVUzEQMA4GA1UECBMHQXJpem9uYTERMA8GA1UEBxMIR29vZH11YXIxEDAOBgNVBAoTB0V4YW1wbGUxEDAOBgNVBAoTB0V4YW1wbGUxEDAOBgnVBAAsTB0V4YW1wbGUxCzAJBgNVBAMTAkNBggkAwXk70cw07gwwDQYJKoZIhvcNAQEEBQADgYEAPiTDX5kYXwdhmijutSkr0bKpRbQkvkkzcyZ106VrAxRQ+eFeN6NyuyhgYy5K61/sIWdaGou5iJOQx2pQYWx1v8Klyl0W22IfEAXYv/epi089hpdaCryuDjpioXI/X8TAwvRwLKL21Dk3k2b+eyCgA00++HM0dPfiQLQ99ElWkv/0=</ds:X509Certificate>
  </ds:X509Data>
</ds:KeyInfo>

```

```

</ds:Signature>
<saml2:Subject>
  <saml2:NameID Format="urn:oasis:names:tc:SAML:1.1:nameid-format:unspecified" NameQualifier="http://cxf.apache.org/sts">srogers</saml2:NameID>
  <saml2:SubjectConfirmation Method="urn:oasis:names:tc:SAML:2.0:cm:holder-of-key">
    <saml2:SubjectConfirmationData xsi:type="saml2:KeyInfoConfirmationDataType">
      <ds:KeyInfo xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
        <ds:X509Data>

<ds:X509Certificate>MIIC5DCCAk2gAwIBAgIJAKj7ROPHjo1yMA0GCSqGSIb3DQEBCwUAMIGKMQswCQYDVQ
QGEwJVUzEQ
MA4GA1UECAwHQXJpem9uYTERMA8GA1UEBwwIR29vZHl1YXIxGDAWBgNVBAoMD0xvY2toZWVkIE1h
cnRpbjENMAssGA1UECwwESTRDRTEPMA0GA1UEAwwGY2xpZW50MRwwGgYJKoZIhvcNAQkBFg1pNGN1
QGxtY28uY29tMB4XDTEyMDYyMDE5NDMwOVoXDTIyMDYxODE5NDMwOvowgYoxCzAJBgNVBAYTA1VT
MRAwDgYDVQQIDAdBcm16b25hMREwDwYDVQQHDAhB29keWVhcjEYMBYGA1UECgwPTG9ja2h1ZWQg
TWFydGluMQ0wCwYDVQQLDARJNENFMQ8wDQYDVQQDAZjbG1lbnQxHDAaBqkqhkIG9w0BCQEWDWk0
Y2VAbG1jby5jb20wgZ8wDQYJKoZIhvcNAQEBBQADgY0AMIGJAoGBAiPhxCBLYE7xfFDLcITS9SsPG
4Q04Z6S32/+TriGsRgpGTj/7GuMG7oJ98m6Ws5cTY17nyunyHTkZuP7rBzy4esDIhheyx18EgdSJ
vvACgGVcNEmHndkf9bWU1AoNaxW+vZwljUkRUVdkhPbPdPwOcMdKg/SsLSNjZfsQIjoWd4rAgMB
AAGjUDBOMB0GA1UdDgQWBBQx11VLtYXLvFGpFdHnh1NW9+lxBDAfBgNVHSMEGDAwgBQx11VLtYXL
vFGpFdHnh1NW9+lxBDAMBgNVHRMEBTADAQH/MA0GCSqGSIb3DQEBCwUAA4GBAHYs20I0K6yVXzyS
sKcv2fmfw6XCIcGTnyA7B0dAjYqq6wD+33dHJUCFDqye7AWdcivuc7RWJt9jn1fJZKIm2BHcDTR
Hhk6CvjJ14Gf40WQdeMHoX8U8b0diq7Iy5Ravx+zRg7SdiyJUqFYjRh/05tywXRT1+freI3bwAN0
L6tQ</ds:X509Certificate>
        </ds:X509Data>
      </ds:KeyInfo>
    </saml2:SubjectConfirmationData>
  </saml2:SubjectConfirmation>
</saml2:Subject>
<saml2:Conditions NotBefore="2013-04-29T18:35:11.407Z"
NotOnOrAfter="2013-04-29T19:05:11.407Z">
  <saml2:AudienceRestriction>

<saml2:Audience>https://server:8993/services/SecurityTokenService</saml2:Audience>
  </saml2:AudienceRestriction>
</saml2:Conditions>
<saml2:AuthnStatement AuthnInstant="2013-04-29T18:35:11.392Z">
  <saml2:AuthnContext>

<saml2:AuthnContextClassRef>urn:oasis:names:tc:SAML:2.0:ac:classes:unspecified</saml2:
AuthnContextClassRef>
  </saml2:AuthnContext>
</saml2:AuthnStatement>
<saml2:AttributeStatement>
  <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
    <saml2:AttributeValue xsi:type="xs:string">
```

```

srogers</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/emailaddress" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
        <saml2:AttributeValue xsi:type="xs:string">
>srogers@example.com</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/surname" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
        <saml2:AttributeValue xsi:type="xs:string">
srogers</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/givenname" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
        <saml2:AttributeValue xsi:type="xs:string">Steve
Rogers</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
        <saml2:AttributeValue xsi:type="xs:string">
avengers</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
        <saml2:AttributeValue xsi:type="xs:string">
admin</saml2:AttributeValue>
    </saml2:Attribute>
    </saml2:AttributeStatement>
    </saml2:Assertion>
</RequestedSecurityToken>
<RequestedAttachedReference>
    <ns3:SecurityTokenReference xmlns:wsse11="http://docs.oasis-
open.org/wss/oasis-wss-wssecurity-secext-1.1.xsd" wsse11:TokenType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0">
        <ns3:KeyIdentifier ValueType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLID"
>_BDC44EB8593F47D1B213672605113671</ns3:KeyIdentifier>
        </ns3:SecurityTokenReference>
    </RequestedAttachedReference>
    <RequestedUnattachedReference>
        <ns3:SecurityTokenReference xmlns:wsse11="http://docs.oasis-
open.org/wss/oasis-wss-wssecurity-secext-1.1.xsd" wsse11:TokenType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0">
            <ns3:KeyIdentifier ValueType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLID"
>_BDC44EB8593F47D1B213672605113671</ns3:KeyIdentifier>

```

```

        </ns3:SecurityTokenReference>
    </RequestedUnattachedReference>
    <wsp:AppliesTo xmlns:wsp="http://schemas.xmlsoap.org/ws/2004/09/policy" xmlns:wst="http://docs.oasis-open.org/ws-sx/ws-trust/200512">
        <wsa:EndpointReference xmlns:wsa="http://www.w3.org/2005/08/addressing">
            <wsa:Address>
                https://server:8993/services/SecurityTokenService</wsa:Address>
            </wsa:EndpointReference>
        </wsp:AppliesTo>
        <Lifetime>
            <ns2:Created>2013-04-29T18:35:11.444Z</ns2:Created>
            <ns2:Expires>2013-04-29T19:05:11.444Z</ns2:Expires>
        </Lifetime>
    </RequestSecurityTokenResponse>
</RequestSecurityTokenResponseCollection>
</soap:Body>
</soap:Envelope>

```

To obtain a SAML assertion to use in secure communication to DDF, a RequestSecurityToken (RST) request has to be made to the STS.

A Bearer SAML assertion is automatically trusted by the endpoint. The client doesn't have to prove it can own that SAML assertion. It is the simplest way to request a SAML assertion, but many endpoints won't accept a KeyType of Bearer.

29.14.3. UsernameToken Bearer SAML Security Token Sample

- WS-Addressing header with Action, To, and Message ID
- Valid, non-expired timestamp
- Username Token containing a username and password that the STS will authenticate
- Issued over HTTPS
- KeyType of <http://docs.oasis-open.org/ws-sx/ws-trust/200512/Bearer>
- Claims (optional): Some endpoints may require that the SAML assertion include attributes of the user, such as an authenticated user's role, name identifier, email address, etc. If the SAML assertion needs those attributes, the [RequestSecurityToken](#) must specify which ones to include.

UsernameToken Bearer SAML Security Token Sample Request

```

<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/">
    <soap:Header xmlns:wsa="http://www.w3.org/2005/08/addressing">
        <wsse:Security xmlns:wsse="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-secext-1.0.xsd" xmlns:wsu="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-utility-1.0.xsd" soap:mustUnderstand="1">
            <wsu:Timestamp wsu:Id="TS-1">
                <wsu:Created>2013-04-29T17:47:37.817Z</wsu:Created>

```

```

        <wsu:Expires>2013-04-29T17:57:37.817Z</wsu:Expires>
    </wsu:Timestamp>
    <wsse:UsernameToken wsu:Id="UsernameToken-1">
        <wsse:Username>srogers</wsse:Username>
        <wsse:Password Type="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-username-token-profile-1.0#PasswordText">password1</wsse:Password>
    </wsse:UsernameToken>
    </wsse:Security>
    <wsa:Action>http://docs.oasis-open.org/ws-sx/ws-trust/200512/RST/Issue</wsa:Action>
        <wsa:MessageID>uuid:a1bba87b-0f00-46cc-975f-001391658cbe</wsa:MessageID>
        <wsa:To>https://server:8993/services/SecurityTokenService</wsa:To>
    </soap:Header>
    <soap:Body>
        <wst:RequestSecurityToken xmlns:wst="http://docs.oasis-open.org/ws-sx/ws-trust/200512">
            <wst:SecondaryParameters>
                <t:TokenType xmlns:t="http://docs.oasis-open.org/ws-sx/ws-trust/200512">http://docs.oasis-open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0</t:TokenType>
                <t:KeyType xmlns:t="http://docs.oasis-open.org/ws-sx/ws-trust/200512">http://docs.oasis-open.org/ws-sx/ws-trust/200512/Bearer</t:KeyType>
                <t:Claims xmlns:ic="http://schemas.xmlsoap.org/ws/2005/05/identity" xmlns:t="http://docs.oasis-open.org/ws-sx/ws-trust/200512" Dialect="http://schemas.xmlsoap.org/ws/2005/05/identity">
                    <!--Add any additional claims you want to grab for the service-->
                    <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/uid"/>
                    <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role"/>
                    <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier"/>
                    <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/emailaddress"/>
                    <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/surname"/>
                    <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/givenname"/>
                </t:Claims>
            </wst:SecondaryParameters>
            <wst:RequestType>http://docs.oasis-open.org/ws-sx/ws-trust/200512/Issue</wst:RequestType>
                <wsp:AppliesTo xmlns:wsp="http://schemas.xmlsoap.org/ws/2004/09/policy">
                    <wsa:EndpointReference xmlns:wsa="http://www.w3.org/2005/08/addressing">
                        <wsa:Address>https://server:8993/services/QueryService</wsa:Address>
                    </wsa:EndpointReference>
                </wsp:AppliesTo>
                <wst:Renewing/>
            </wst:RequestSecurityToken>
        
```

```
</soap:Body>  
</soap:Envelope>
```

This is the response from the STS containing the SAML assertion to be used in subsequent requests to QCRUD endpoints:

The **saml2:Assertion** block contains the entire SAML assertion.

The **Signature** block contains a signature from the STS's private key. The endpoint receiving the SAML assertion will verify that it trusts the signer and ensure that the message wasn't tampered with.

The **AttributeStatement** block contains all the Claims requested.

The **Lifetime** block indicates the valid time interval in which the SAML assertion can be used.

UsernameToken Bearer SAML Security Token Sample Response

```
<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/">  
  <soap:Header>  
    <Action xmlns="http://www.w3.org/2005/08/addressing">http://docs.oasis-open.org/ws-sx/ws-trust/200512/RSTR/IssueFinal</Action>  
    <MessageID xmlns="http://www.w3.org/2005/08/addressing">urn:uuid:eee4c6ef-ac10-4cbc-a53c-13d960e3b6e8</MessageID>  
    <To xmlns="http://www.w3.org/2005/08/addressing">  
      <RelatesTo xmlns="http://www.w3.org/2005/08/addressing">uuid:a1bba87b-0f00-46cc-975f-001391658cbe</RelatesTo>  
      <wsse:Security xmlns:wsse="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-secext-1.0.xsd" xmlns:wsu="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-utility-1.0.xsd" soap:mustUnderstand="1">  
        <wsu:Timestamp wsu:Id="TS-2">  
          <wsu:Created>2013-04-29T17:49:12.624Z</wsu:Created>  
          <wsu:Expires>2013-04-29T17:54:12.624Z</wsu:Expires>  
        </wsu:Timestamp>  
      </wsse:Security>  
    </soap:Header>  
    <soap:Body>  
      <RequestSecurityTokenResponseCollection xmlns="http://docs.oasis-open.org/ws-sx/ws-trust/200512" xmlns:ns2="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-utility-1.0.xsd" xmlns:ns3="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-secext-1.0.xsd" xmlns:ns4="http://www.w3.org/2005/08/addressing" xmlns:ns5="http://docs.oasis-open.org/ws-sx/ws-trust/200802">  
        <RequestSecurityTokenResponse>  
          <TokenType>http://docs.oasis-open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0</TokenType>  
          <RequestedSecurityToken>  
            <saml2:Assertion xmlns:saml2="urn:oasis:names:tc:SAML:2.0:assertion" xmlns:xs="http://www.w3.org/2001/XMLSchema">
```

```

xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance" ID=
"_7437C1A55F19AFF22113672577526132" IssueInstant="2013-04-29T17:49:12.613Z" Version=
"2.0" xsi:type="saml2:AssertionType">
    <saml2:Issuer>tokenissuer</saml2:Issuer>
    <ds:Signature xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
        <ds:SignedInfo>
            <ds:CanonicalizationMethod Algorithm=
"http://www.w3.org/2001/10/xml-exc-c14n#" />
            <ds:SignatureMethod Algorithm=
"http://www.w3.org/2000/09/xmldsig#rsa-sha1" />
            <ds:Reference URI="#_7437C1A55F19AFF22113672577526132
">
                <ds:Transforms>
                    <ds:Transform Algorithm=
"http://www.w3.org/2000/09/xmldsig#enveloped-signature" />
                    <ds:Transform Algorithm=
"http://www.w3.org/2001/10/xml-exc-c14n#" />
                    <ec:InclusiveNamespaces xmlns:ec=
"http://www.w3.org/2001/10/xml-exc-c14n#" PrefixList="xs" />
                    </ds:Transform>
                </ds:Transforms>
                <ds:DigestMethod Algorithm=
"http://www.w3.org/2000/09/xmldsig#sha1" />
                <ds:DigestValue>
                    Re0qEbGZlyplW5kqiyX0jPnVEA=</ds:DigestValue>
                    </ds:Reference>
                </ds:SignedInfo>

                <ds:SignatureValue>X5Kzd54PrK1lGVV2XxzCmWFRzHRoybF7hU6zxbEhSLMR0AWS9R7Me3epq91Xqe0wvID
DbwmE/oJNC7vI0fIw/rqXkx4aZsY5a5nbAs7f+aXF9T6dk82x2eNhNGYpViq0YZJfsJ5WSyMtG8w5nRekmDMy9
oTLsHG+Y/0hJDEwq58=</ds:SignatureValue>
                <ds:KeyInfo>
                    <ds:X509Data>

                    <ds:X509Certificate>MIIDmjCCAwOgAwIBAgIBBDANBgkqhkiG9w0BAQQFADB1MQswCQYDVQQGEwJVUzEQMA
4GA1UECBMH
QXJpem9uYTERMA8GA1UEBxMIR29vZHl1YXIxEDAOBgNVBAoTB0V4YW1wbGUxEDAOBgNVBAoTB0V4
YW1wbGUxEDAOBgNVBAoTB0V4YW1wbGUxCzAJBgNVBAMTAkNBMB4XDTEzMDQwOTE4MzcxMVoXDTIz
MDQwNzE4MzcxMVowgaYxCzAJBgNVBAYTA1VTMRAwDgYDVQQIEwdBcm16b25hMREwDwYDVQQHEwhH
b29keWVhcjEQMA4GA1UEChMHRXhhbXBsZTEQMA4GA1UEChMHRXhhbXBsZTEQMA4GA1UECxMHRXhh
bXBsZTEUMBIGA1UEAxMLdG9rZW5pc3N1ZXIxJjAkBgkqhkiG9w0BCQEWF3Rva2VuaXNzdWVyQGV4
YW1wbGUuY29tMIGfMA0GCSqGSIb3DQEBAQUAA4GNADCBiQKBgQDDfktA8Lrp9rTfRibKdgtxtN9
uB44diiIqq3J0zDGfDhGLu6mjpuH01hrKItv42hB0hhmH71S9ipiaQCIpVfgIG63MB7fa5dBrfGF
G69vFrU1Lfi7IvsVVsnrtAEQljoMmw9sxS3SUsRQX+bD8jq7Uj1hpoF7DdqpV8Kb0C00GwIDAQAB
o4IBBjCCAQIwCQYDVROTBAlwADAsBglghkgBvhCAQ0EHxYdT3B1b1NTTCBHZW51cmF0ZWQgQ2Vy
dG1maWNhdGUwHQYDVRO0BBYEFD1mHviop2Tc4HaNu8yPXR6GqWP1MIGnBgNVHSMEgZ8wgZyAFBcn
en6/j05DzaVw0RwrtteKc7TZ0oXmkdzB1MQswCQYDVQQGEwJVUzEQMA4GA1UECBMHQXJpem9uYTER
MA8GA1UEBxMIR29vZHl1YXIxEDAOBgNVBAoTB0V4YW1wbGUxEDAOBgNVBAoTB0V4YW1wbGUxEDAO
BgNVBAoTB0V4YW1wbGUxCzAJBgNVBAMTAkNBggkAwXk70cw07gwwDQYJKoZIhvcNAQEEBQADgYEA
PiTX5kYXwdhmijutSkr0bKpRbQkvkkzcyZl06VrAxRQ+eFeN6NyuyhgYy5K61/sIWdaGou5iJOQx
2pQYWx1v8Klyl0W22IfEAXYv/epi089hpdaCryuDJpioXI/X8TAwvRwLKL21Dk3k2b+eyCgA00++
```

```

HM0dPfiQLQ99ElWkv/0=</ds:X509Certificate>
    </ds:X509Data>
    </ds:KeyInfo>
</ds:Signature>
<saml2:Subject>
    <saml2:NameID Format="urn:oasis:names:tc:SAML:1.1:nameid-format:unspecified" NameQualifier="http://cxf.apache.org/sts">srogers</saml2:NameID>
    <saml2:SubjectConfirmation Method=
"urn:oasis:names:tc:SAML:2.0:cm:bearer"/>
</saml2:Subject>
<saml2:Conditions NotBefore="2013-04-29T17:49:12.614Z"
NotOnOrAfter="2013-04-29T18:19:12.614Z">
    <saml2:AudienceRestriction>

<saml2:Audience>https://server:8993/services/QueryService</saml2:Audience>
    </saml2:AudienceRestriction>
</saml2:Conditions>
<saml2:AuthnStatement AuthnInstant="2013-04-29T17:49:12.613Z">
    <saml2:AuthnContext>

<saml2:AuthnContextClassRef>urn:oasis:names:tc:SAML:2.0:ac:classes:unspecified</saml2:AuthnContextClassRef>
    </saml2:AuthnContext>
    </saml2:AuthnStatement>
    <saml2:AttributeStatement>
        <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
            <saml2:AttributeValue xsi:type="xs:string">
srogers</saml2:AttributeValue>
        </saml2:Attribute>
        <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/emailaddress" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
            <saml2:AttributeValue xsi:type="xs:string">
>srogers@example.com</saml2:AttributeValue>
        </saml2:Attribute>
        <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/surname" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
            <saml2:AttributeValue xsi:type="xs:string">
srogers</saml2:AttributeValue>
        </saml2:Attribute>
        <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/givenname" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
            <saml2:AttributeValue xsi:type="xs:string">Steve
Rogers</saml2:AttributeValue>
        </saml2:Attribute>
        <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role" NameFormat=

```

```

"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
    <saml2:AttributeValue xsi:type="xs:string">
avengers</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
        <saml2:AttributeValue xsi:type="xs:string">
admin</saml2:AttributeValue>
        </saml2:Attribute>
    </saml2:AttributeStatement>
</saml2:Assertion>
</RequestedSecurityToken>
<RequestedAttachedReference>
    <ns3:SecurityTokenReference xmlns:wsse11="http://docs.oasis-
open.org/wss/oasis-wss-wssecurity-secext-1.1.xsd" wsse11:TokenType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0">
        <ns3:KeyIdentifier ValueType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLID"
>_7437C1A55F19AFF22113672577526132</ns3:KeyIdentifier>
        </ns3:SecurityTokenReference>
    </RequestedAttachedReference>
    <RequestedUnattachedReference>
        <ns3:SecurityTokenReference xmlns:wsse11="http://docs.oasis-
open.org/wss/oasis-wss-wssecurity-secext-1.1.xsd" wsse11:TokenType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0">
            <ns3:KeyIdentifier ValueType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLID"
>_7437C1A55F19AFF22113672577526132</ns3:KeyIdentifier>
            </ns3:SecurityTokenReference>
        </RequestedUnattachedReference>
        <wsp:AppliesTo xmlns:wsp="

http://schemas.xmlsoap.org/ws/2004/09/policy" xmlns:wst="http://docs.oasis-
open.org/ws-sx/ws-trust/200512">
            <wsa:EndpointReference xmlns:wsa=
"http://www.w3.org/2005/08/addressing">
                <wsa:Address>
https://server:8993/services/QueryService</wsa:Address>
                </wsa:EndpointReference>
            </wsp:AppliesTo>
            <Lifetime>
                <ns2:Created>2013-04-29T17:49:12.620Z</ns2:Created>
                <ns2:Expires>2013-04-29T18:19:12.620Z</ns2:Expires>
            </Lifetime>
            </RequestSecurityTokenResponse>
        </RequestSecurityTokenResponseCollection>
    </soap:Body>
</soap:Envelope>

```

In order to obtain a SAML assertion to use in secure communication to DDF, a [RequestSecurityToken](#)

(RST) request has to be made to the STS.

An endpoint's policy will specify the type of security token needed. Most of the endpoints that have been used with DDF require a SAML v2.0 assertion with a required KeyType of <http://docs.oasis-open.org/ws-sx/ws-trust/200512/PublicKey>. This means that the SAML assertion provided by the client to a DDF endpoint must contain a SubjectConfirmation block with a type of "holder-of-key" containing the client's public key. This is used to prove that the client can possess the SAML assertion returned by the STS.

29.14.4. X.509 PublicKey SAML Security Token Sample

The STS that comes with DDF requires the following to be in the RequestSecurityToken request in order to issue a valid SAML assertion. See the request block below for an example of how these components should be populated.

- WS-Addressing header containing Action, To, and MessageID blocks
- Valid, non-expired timestamp
- Issued over HTTPS
- TokenType of <http://docs.oasis-open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0>
- KeyType of <http://docs.oasis-open.org/ws-sx/ws-trust/200512/PublicKey>
- X509 Certificate as the Proof of Possession or POP. This needs to be the certificate of the client that will be both requesting the SAML assertion and using the SAML assertion to issue a query
- Claims (optional): Some endpoints may require that the SAML assertion include attributes of the user, such as an authenticated user's role, name identifier, email address, etc. If the SAML assertion needs those attributes, the RequestSecurityToken must specify which ones to include.
 - UsernameToken: If Claims are required, the RequestSecurityToken security header must contain a UsernameToken element with a username and password.

X.509 PublicKey SAML Security Token Sample Request

```
<soapenv:Envelope xmlns:ns="http://docs.oasis-open.org/ws-sx/ws-trust/200512"
  xmlns:soapenv="http://schemas.xmlsoap.org/soap/envelope/">
  <soapenv:Header xmlns:wsa="http://www.w3.org/2005/08/addressing">
    <wsa:Action>http://docs.oasis-open.org/ws-sx/ws-trust/200512/RST/Issue</wsa:Action>
    <wsa:MessageID>uuid:527243af-94bd-4b5c-a1d8-024fd7e694c5</wsa:MessageID>
    <wsa:To>https://server:8993/services/SecurityTokenService</wsa:To>
    <wsse:Security soapenv:mustUnderstand="1" xmlns:wsse="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-secext-1.0.xsd" xmlns:wsu="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-utility-1.0.xsd">
      <wsu:Timestamp wsu:Id="TS-17">
        <wsu:Created>2014-02-19T17:30:40.771Z</wsu:Created>
        <wsu:Expires>2014-02-19T19:10:40.771Z</wsu:Expires>
      </wsu:Timestamp>
      <!-- OPTIONAL: Only required if the endpoint that the SAML assertion will be sent to requires claims. -->
    </wsse:Security>
  </soapenv:Header>
</soapenv:Envelope>
```

```

<wsse:UsernameToken wsu:Id="UsernameToken-16">
    <wsse:Username>pparker</wsse:Username>
    <wsse:Password Type="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-username-token-profile-1.0#PasswordText">password1</wsse:Password>
    <wsse:Nonce EncodingType="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-soap-message-security-1.0#Base64Binary">
LCTD+5Y7h1WIP6SpsEg9XA==</wsse:Nonce>
    <wsu:Created>2014-02-19T17:30:37.355Z</wsu:Created>
    </wsse:UsernameToken>
    </wsse:Security>
</soapenv:Header>
<soapenv:Body>
    <wst:RequestSecurityToken xmlns:wst="http://docs.oasis-open.org/ws-sx/ws-trust/200512">
        <wst:TokenType>http://docs.oasis-open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0</wst:TokenType>
        <wst:KeyType>http://docs.oasis-open.org/ws-sx/ws-trust/200512/PublicKey</wst:KeyType>
        <!-- OPTIONAL: Only required if the endpoint that the SAML assertion will be sent to requires claims. -->
        <wst:Claims Dialect="http://schemas.xmlsoap.org/ws/2005/05/identity" xmlns:ic="http://schemas.xmlsoap.org/ws/2005/05/identity">
            <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role"/>
            <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier"/>
            <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/emailaddress"/>
            <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/surname"/>
            <ic:ClaimType Optional="true" Uri="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/givenname"/>
        </wst:Claims>
        <wst:RequestType>http://docs.oasis-open.org/ws-sx/ws-trust/200512/Issue</wst:RequestType>
        <wsp:AppliesTo xmlns:wsp="http://schemas.xmlsoap.org/ws/2004/09/policy">
            <wsa:EndpointReference xmlns:wsa="http://www.w3.org/2005/08/addressing">
                <wsa:Address>https://server:8993/services/QueryService</wsa:Address>
            </wsa:EndpointReference>
        </wsp:AppliesTo>
        <wst:UseKey>
            <ds:KeyInfo xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
                <ds:X509Data>
<ds:X509Certificate>MIIFGDCCBACgAwIBAgICJe0wDQYJKoZIhvcNAQEFBQAwXDELMAkGA1UEBhMCVVMxGD
AWBgNVBAoT
D1UuUy4gR292ZXJubWVudDEMMAoGA1UECxMDRG9EMQwwCgYDVQQLEwNQS0kxFzAVBgNVBAMTDkRP
RCBKSVRDIENBLTI3MB4XDTEzMDUwNzAwMjU00VoXDTE2MDUwNzAwMjU00VowaTELMAkGA1UEBhMC
VVMxGDAWBgNVBAoTD1UuUy4gR292ZXJubWVudDEMMAoGA1UECxMDRG9EMQwwCgYDVQQLEwNQS0kxF
zARBgNVBAsTCKNPTlRSQUNUT1IxDzANBgNVBAMTBmNsawWVudDCCASIwDQYJKoZIhvcNAQEBBQAD

```

```

ggEPADCCAQoCggEBA0q6L1/jjZ5cyhjhHEbOHr5WQpboKACYbrsn8lg85LGNoAfcwImr9KBm0xGb
ZCxHYIhkW7pJ+kppYH8DbbbDMviIvvdkvrAIU0180BRn2wReCBGQ01Imdc3+WzFF2svW75d6wi2ZVd
eMvU015p/pAD/sdIfXmAfyu8+tqi08KVZGkTnlg3AMzfeSrkci5UHMVwj0qUSuzLk9SAg/9STgb
Kf2xBpHUYecWFSB+dTpdzN2pC85tj9xIoWGH5dFWG1fPcYRgzGPxsybiG0ylbJ7rHDJuL7IIIyx5
EnkCuxmQwoQ6xQAh iWRGyP1Y08w1LZixI2v+Cv/ZjUfIHv49I9P4Mt8CAwEAAaOCAdUwggHRMB8G
A1UdIwQYMBaAFCMUNCBNxY43NZLBB1nDjDp1NZJoMB0GA1UdDgQWBBRPGiX6zZzKTqQSx/tjg6hx
9opDoTAOBgNVHQ8BAf8EBAMCBaAwgdoGA1UdHwSB0jCBzzA2oDSgMoYwaHR0cDovL2NybC5nZHMu
bm10LmRp2EubWlsL2NybC9ET0RKSVDQ0FfMjcuY3J5MIGUoIGRoIGOh0GLbGRhcDovL2NybC5n
ZHMubm10LmRp2EubWlsL2NuJTNkRE9EJTIwSk1LUQyUyMENBLTI3JTJjb3U1M2RQS0k1MmNvdSUz
ZERvRCUyY281M2RVL1MuJTIwR292ZXJubWVudCuyY2M1M2RVUz9jZXJ0aWZpY2F0ZXJ1dm9jYXRp
b25saXN002JpbmFyeTAjBjNVHSAEHDaaMASGCWCGSAF1AgELBTALBglghkgBZQIBCxIwfQYIKwYB
BQUHAQEEcTBvMD0GCCsGAQUBFBzAChjFodHRw0i8vY3J5Lmdkcy5uaXQuZG1zYS5taWwvc2lnbi9E
T0RKSVDQ0FfMjcuY2VvMC4GCCsGAQUBFBzABhiJodHRw0i8vb2NzcC5uc24wLnJjdMubm10LmRp
c2EubWlsMA0GCSqGSIB3DQEBBQUAA4IBAQCGUJPGh4iGCbr2xCMqCq04SFQ+iaLmTIFAxZPFvup1
4E9Ir6CSDalpF9eBx9fS+Z2xuesKyM/g3YqWU1LtfWGRRIxzEujaC4YpwHuffkx9QqkwSkXXIsim
EhmzSgznT4Q9X8WalqVYOfNZ6sSLZ8qPPFrLHkkw/zIFRzo62wXLu0tfcp0r+iaJBhyDRinIhr
hwtE3xo6qQRRWl03/c1C4RnTev1crFVJQVF3yfpRu8udJ2S0GdqU0vjUSu1h7aMkYJMH1u08Whj
8KASjJBFeHPirMV1oddJ5ydZCQ+Jmnpbwq+XsCxg1LjC4dmbjKVr9s4QK+/JLNjxD8IkJiZE</ds:X509Certificate>
    </ds:X509Data>
    </ds:KeyInfo>
</wst:UseKey>
</wst:RequestSecurityToken>
</soapenv:Body>
</soapenv:Envelope>

```

29.14.5. X.509 PublicKey SAML Security Token Sample

This is the response from the STS containing the SAML assertion to be used in subsequent requests to QCRUD endpoints.

The **saml2:Assertion** block contains the entire SAML assertion.

The **Signature** block contains a signature from the STS's private key. The endpoint receiving the SAML assertion will verify that it trusts the signer and ensure that the message wasn't tampered with.

The **SubjectConfirmation** block contains the client's public key, so the server can verify that the client has permission to hold this SAML assertion. The **AttributeStatement** block contains all of the claims requested.

X.509 PublicKey SAML Security Token Sample Response

```

<soap:Envelope xmlns:soap="http://schemas.xmlsoap.org/soap/envelope/">
  <soap:Header>
    <Action xmlns="http://www.w3.org/2005/08/addressing">http://docs.oasis-open.org/ws-sx/ws-trust/200512/RSTR/IssueFinal</Action>
    <MessageID xmlns="http://www.w3.org/2005/08/addressing">urn:uuid:b46c35ad-3120-4233-ae07-b9e10c7911f3</MessageID>
    <To xmlns="http://www.w3.org/2005/08/addressing">

```

```

>http://www.w3.org/2005/08/addressing/anonymous</To>
  <RelatesTo xmlns="http://www.w3.org/2005/08/addressing">uuid:527243af-94bd-4b5c-
a1d8-024fd7e694c5</RelatesTo>
    <wsse:Security soap:mustUnderstand="1" xmlns:wsse="http://docs.oasis-
open.org/wss/2004/01/oasis-200401-wss-wssecurity-secext-1.0.xsd" xmlns:wsu=
"http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-wssecurity-utility-1.0.xsd">
      <wsu:Timestamp wsu:Id="TS-90DBA0754E55B4FE7013928310431357">
        <wsu:Created>2014-02-19T17:30:43.135Z</wsu:Created>
        <wsu:Expires>2014-02-19T17:35:43.135Z</wsu:Expires>
      </wsu:Timestamp>
    </wsse:Security>
  </soap:Header>
  <soap:Body>
    <ns2:RequestSecurityTokenResponseCollection xmlns="http://docs.oasis-
open.org/ws-sx/ws-trust/200802" xmlns:ns2="http://docs.oasis-open.org/ws-sx/ws-
trust/200512" xmlns:ns3="http://docs.oasis-open.org/wss/2004/01/oasis-200401-wss-
wssecurity-utility-1.0.xsd" xmlns:ns4="http://docs.oasis-open.org/wss/2004/01/oasis-
200401-wss-wssecurity-secext-1.0.xsd" xmlns:ns5="http://www.w3.org/2005/08/addressing
">
      <ns2:RequestSecurityTokenResponse>
        <ns2:TokenType>http://docs.oasis-open.org/wss/oasis-wss-saml-token-
profile-1.1#SAMLV2.0</ns2:TokenType>
        <ns2:RequestedSecurityToken>
          <saml2:Assertion ID="_90DBA0754E55B4FE7013928310431176" IssueInstant=
"2014-02-19T17:30:43.117Z" Version="2.0" xsi:type="saml2:AssertionType" xmlns:saml2=
"urn:oasis:names:tc:SAML:2.0:assertion" xmlns:xs="http://www.w3.org/2001/XMLSchema"
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">
            <saml2:Issuer>tokenissuer</saml2:Issuer>
            <ds:Signature xmlns:ds="http://www.w3.org/2000/09/xmldsig#">
              <ds:SignedInfo>
                <ds:CanonicalizationMethod Algorithm=
"http://www.w3.org/2001/10/xml-exc-c14n#" />
                <ds:SignatureMethod Algorithm=
"http://www.w3.org/2000/09/xmldsig#rsa-sha1" />
                <ds:Reference URI="#_90DBA0754E55B4FE7013928310431176">
                  <ds:Transforms>
                    <ds:Transform Algorithm=
"http://www.w3.org/2000/09/xmldsig#enveloped-signature" />
                    <ds:Transform Algorithm="http://www.w3.org/2001/10/xml-
exc-c14n#" />
                    <ec:InclusiveNamespaces PrefixList="xs" xmlns:ec=
"http://www.w3.org/2001/10/xml-exc-c14n#" />
                  </ds:Transforms>
                <ds:DigestMethod Algorithm=
"http://www.w3.org/2000/09/xmldsig#sha1" />
                <ds:DigestValue>
                  /bEGqsRGHVJbx298WPmGd8I53zs=</ds:DigestValue>
                </ds:Reference>
              </ds:SignedInfo>
              <ds:SignatureValue>

```

mYR7w1/dnuh8Z7t9xjCb4XkYQLshj+UuYlG0uTwDYSUPcS2qI0nAgMD1VsDP7y1fDJxeqsq7HYhFKsnqRfebMM
 4WLH1D/1J4rD4U0+i9l3tu1Hm17SN24WM1/b0qfDUCoDqmwG8afUJ3r4vmTNPxfwfOss8BZ/80DgZzm08nd1kx
 DfvcN70rExbV/3/45JwF/MMPZoqvi2MJGfx56E9fErJNuzezpWnRqP01WPxyffKMA1VaB9zF6gvVnUqcW2k/Z8
 X91N705jouBI281ZnIfsIPuBJERFtYNVDHsIXM1pJnrY6FlKIa0si55LQu3Ruir/n82pU7BT5aWtxwrn7akBg=

=

</ds:SignatureValue>

<ds:KeyInfo>

<ds:X509Data>

<ds:X509Certificate>MIIFHTCCBAwIBAgICJe8wDQYJKoZIhvcNAQEFBQAwXDELMAkGA1UEBhMCVVMxGD
 AWBgNVBAoT
 D1UuUy4gR292ZXJubWVudDEMMAoGA1UECxMDRG9EMQwwCgYDVQQLEwNQS0kxFzAVBqNVBAMTDkRP
 RCBKSVRDIENBLT13MB4XDTEzMDUwNzAwMjYzN1oXDTE2MDUwNzAwMjYzN1owbjELMAkGA1UEBhMC
 VVMxGDAwBqNVBAoTD1UuUy4gR292ZXJubWVudDEMMAoGA1UECxMDRG9EMQwwCgYDVQQLEwNQS0kx
 EzARBgNVBAsTCkNPTlRSQUNUT1IxFDASBqNVBAMTC3Rva2VuaXNzdWVyMIIBIjANBgkqhkiG9w0B
 AQEFAAACQ8AMIIBCgKCAQEAx01/U4M1wG+wL1JxX2RL1glj101FkJXMK3Kft3zD//N8x/Dcwvvs
 ngCQjXrV6YhbB2V7scHwnThPv3RSwYYi062z+g6ptfBbKGGBLSZ0zLe3fyJR4Rxb1FKsELFgPHfX
 vgUHS/keG5uSRk9S/0kqps/yxKB7+Z1xeFxsIz5QywXvBpMiXtc2zF+M7BsbSIDSx5LcPcDFBwjF
 c66rE3/y/25VMht9EZX1QoKr7f8rWD4xgd5J6DYMFWEcniCz4BDJH9sfTw+n1P+CYgrhws1WGqxt
 cDME9t6SWR3GLT4Sdtr8ziIM5uUteEhPIV3rVC3/u23JbYEeS8mpnp0bxt5eHQIDAQABo4IB1TCC
 AdEwHwYDVR0jBBgwFoAUIxQ0IE1fLjc1ksEGWcOM0mU1kmgwHQYDVR0OBBYEFGBjdKdey+bMHMhC
 Z7gwIq/mJf5VMA4GA1UdDwEB/wQEawIFoDCB2gYDVR0fBIHSMIHPMDagNKAyhjBodHRw0i8vY3Js
 Lmdkcy5uaXQuZG1zYS5taWwvY3JsL0RPREpJVENDQV8yNy5jcmwwgZSggZGggY6GgYtsZGFw0i8v
 Y3JsLmdkcy5uaXQuZG1zYS5taWwvY241M2RET0Q1MjBKSVDJTIwQ0EtMjc1MmNvdSUzZFBLSUy
 Y291JTNkRG9EJTjbyUzZFuuUy41MjBhb3Zlcm5tZW50JTjYyUzZFTP2N1cnRpZmljYXRlcwV2
 b2NhdG1vbmxpc3Q7YmluYXJ5MCMGA1UdIAQcMBowCwYJYIZIAWUCAQsFMAsgCwCGSAFlAgELEjB9
 BggRbgEFBQcBAQRxMG8wPQYIKwYBBQUHMAKGMWh0dHA6Ly9jcmwuZ2RzLm5pdC5kaXNhLm1pbC9z
 aWduL0RPREpJVENDQV8yNy5jZXIwLgYIKwYBBQUHMAGGImh0dHA6Ly9vY3NwLm5zbjAucmN2cy5u
 aXQuZG1zYS5taWwvDQYJKoZIhvcNAQEFBQADggEBAlHZQTINU3bMpJ/PkwTYLWPmwCqAYgEUzSYx
 bNcVY5MWD8b4XCdw5nM3GnFl0qr4IrHeyyOzsEbIebTe3bv0l1pHx0Uyj059nAhx/AP8DjVtuRU1
 /Mp4b6uJ/4yaoMjIGceqBzHqhHIJinG0Y2azua7eM9hVbWZsa912ihbiupCq22mYuHFP7NUNzBvV
 j03YUcsy/sES5sRx9Rops/CBN+LUUY0dJ0xYWxo8oAbtF8ABE5ATLAwqz4ttsToKPUYh1sxdx5Ef
 APeZ+wYDmMu40fLckwnCKZgkEtJ0xXpdIJHY+VmyZtQSB0LkR5toeH/ANV4259Ia5ZT8h2/vIJBg
 6B4=</ds:X509Certificate>

</ds:X509Data>

</ds:KeyInfo>

</ds:Signature>

<saml2:Subject>

<saml2:NameID Format="urn:oasis:names:tc:SAML:1.1:nameid-format:unspecified" NameQualifier="http://cxf.apache.org/sts">pparker</saml2:NameID>

<saml2:SubjectConfirmation Method="urn:oasis:names:tc:SAML:2.0:cm:holder-of-key">

<saml2:SubjectConfirmationData xsi:type="saml2:KeyInfoConfirmationDataType">

<ds:KeyInfo xmlns:ds="http://www.w3.org/2000/09/xmldsig#">

<ds:X509Data>

<ds:X509Certificate>MIIFGDCCBAwIBAgICJe0wDQYJKoZIhvcNAQEFBQAwXDELMAkGA1UEBhMCVVMxGD
 AWBgNVBAoT
 D1UuUy4gR292ZXJubWVudDEMMAoGA1UECxMDRG9EMQwwCgYDVQQLEwNQS0kxFzAVBqNVBAMTDkRP
 RCBKSVRDIENBLT13MB4XDTEzMDUwNzAwMjU00VoXDTE2MDUwNzAwMjU00VowaTELMAkGA1UEBhMC
 VVMxGDAwBqNVBAoTD1UuUy4gR292ZXJubWVudDEMMAoGA1UECxMDRG9EMQwwCgYDVQQLEwNQS0kx

EzARBgNVBAsTCKNPTlRSQUNUT1IxDzANBgNVBAMTBmNsawVudDCCASIwDQYJKoZIhvcNAQEBBQADggEPADCCAQoCggEBA0q6L1/jjZ5cyhjhHEb0Hr5WQpboKACYbrsn8lg85LGNoAfcwImr9Kbm0xGbzCxCYIhkW7pJ+kppYH8bbbviIvvdkvrAIU0180BRn2wReCBGQ01Imdc3+WzFF2svW75d6wi2ZVdeMvU015p/pAD/sdIfXmAfyu8+tqi08KVZGkTnlg3AMzfeSrkci5UHMVWj0qUSuzLk9SAg/9STgbKf2xBpHUYecWFSB+dTpZN2pC85tj9xIoWGH5dFWG1fPcYRgzGPxsybiG0ylbJ7rHDJuL7IIIyx5EnkCuxmQwoQ6XQAh iWRGyP1Y08w1LZixI2v+Cv/ZjUfIHv49I9P4Mt8CAwEAAaOCAdUwggHRMB8GA1UdIwQYMBaAFCMUNCBNXy43NZLBB1nDjDp1NZJ0MB0GA1UdDgQWBBRPGiX6zZzKTqQSx/tjg6hx9opDoTAOBgNVHQ8BAf8EBAMCBaAwgdoGA1UdHwSB0jCBzzA2oDSgMoYwaHR0cDovL2NybC5nZHMubm10LmRp2EubWlsL2NybC9ET0RKSVDQ0FfMjcuY3J sMIGUoIGRoIG0hoGLbGRhcDovL2NybC5nZHMubm10LmRp2EubWlsL2NuJTNkRE9EJTIwSk1LUQyUyMENBLTI3JTJjb3U1M2RQS0k1MmNvdSUzZERvRCUyY281M2RVL1MuJTIwR292ZJubWVudCUyY2M1M2RVUz9jZXJ0aWZpY2F0ZXJ1dm9jYXRpb25saXN002JpbmFyeTAjBqNVHSAEHDAAmAsgCWCgsAF1AgELBTALBglghkgBZQIBCxIwfQYIKwYBbQUHAQEEcTBvMD0GCCsGAQUBFzAChjFodHRw0i8vY3J sLmdkcy5uaXQuZG1zYS5taWwvc2lnbi9ET0RKSVDQ0FfMjcuY2VyMC4GCCsGAQUBFzABhiJodHRw0i8vb2NzcC5uc24wLnJjdMubm10LmRp c2EubWlsMA0GCSqGSIb3DQEBBQUAA4IBAQCGUJPGh4iGCr2xCMqCq04SFQ+iaLmTIFAxZPFvup14E9Ir6CSDalpF9eBx9fS+Z2xuesKyM/g3YqWU1LtfWGRRIxzEujaC4YpwHuffkx9QqkwSkXXIsimEhmzSgznT4Q9X8WwalqVYOfNZ6sSLZ8qPPFrLHkkw/zIFRzo62wXLu0tfcp0r+iaJBhyDRinIHrhwtE3x06qQRRWl03/c1C4RnTev1crFVJQVB3yfpRu8udJ2S0GdqU0v jUSu1h7aMkYJMH1u08Whj8KASjJBFeHPirMV1oddJ5ydZCQ+Jmnpbwq+XsCxg1LjC4dmbjKVr9s4QK+/JLNjxD8IkJiZE</ds:X509Certificate>

```
        </ds:X509Data>
        </ds:KeyInfo>
    </saml2:SubjectConfirmationData>
    </saml2:SubjectConfirmation>
    </saml2:Subject>
    <saml2:Conditions NotBefore="2014-02-19T17:30:43.119Z" NotOnOrAfter="2014-02-19T18:00:43.119Z"/>
    <saml2:AuthnStatement AuthnInstant="2014-02-19T17:30:43.117Z">
        <saml2:AuthnContext>

<saml2:AuthnContextClassRef>urn:oasis:names:tc:SAML:2.0:ac:classes:unspecified</saml2:AuthnContextClassRef>
        </saml2:AuthnContext>
    </saml2:AuthnStatement>

        <!-- This block will only be included if Claims were requested in the RST. -->
    <saml2:AttributeStatement>
        <saml2:Attribute Name="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier" NameFormat="urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
            <saml2:AttributeValue xsi:type="xs:string">pparker</saml2:AttributeValue>
        </saml2:Attribute>
        <saml2:Attribute Name="http://schemas.xmlsoap.org/ws/2005/05/identity/claims/emailaddress" NameFormat="urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
            <saml2:AttributeValue xsi:type="xs:string">pparker@example.com</saml2:AttributeValue>
        </saml2:Attribute>
        <saml2:Attribute Name="
```

```

"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/surname" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
    <saml2:AttributeValue xsi:type="xs:string">
parker</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/givenname" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
    <saml2:AttributeValue xsi:type="xs:string">Peter
Parker</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
    <saml2:AttributeValue xsi:type="xs:string">
users</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
    <saml2:AttributeValue xsi:type="xs:string">
users</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
    <saml2:AttributeValue xsi:type="xs:string">
avengers</saml2:AttributeValue>
    </saml2:Attribute>
    <saml2:Attribute Name=
"http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role" NameFormat=
"urn:oasis:names:tc:SAML:2.0:attrname-format:unspecified">
    <saml2:AttributeValue xsi:type="xs:string">
admin</saml2:AttributeValue>
    </saml2:Attribute>
    </saml2:AttributeStatement>
    </saml2:Assertion>
</ns2:RequestedSecurityToken>
<ns2:RequestedAttachedReference>
    <ns4:SecurityTokenReference wsse11:TokenType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0" xmlns:wsse11=
"http://docs.oasis-open.org/wss/oasis-wss-wssecurity-secext-1.1.xsd">
        <ns4:KeyIdentifier ValueType="http://docs.oasis-open.org/wss/oasis-
wss-saml-token-profile-1.1#SAMLID">
_90DBA0754E55B4FE7013928310431176</ns4:KeyIdentifier>
        </ns4:SecurityTokenReference>
    </ns2:RequestedAttachedReference>
    <ns2:RequestedUnattachedReference>
        <ns4:SecurityTokenReference wsse11:TokenType="http://docs.oasis-
open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0" xmlns:wsse11=
"http://docs.oasis-open.org/wss/oasis-wss-wssecurity-secext-1.1.xsd">

```

```

<ns4:KeyIdentifier ValueType="http://docs.oasis-open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLID">
_90DBA0754E55B4FE7013928310431176</ns4:KeyIdentifier>
</ns4:SecurityTokenReference>
</ns2:RequestedUnattachedReference>
<ns2:Lifetime>
<ns3:Created>2014-02-19T17:30:43.119Z</ns3:Created>
<ns3:Expires>2014-02-19T18:00:43.119Z</ns3:Expires>
</ns2:Lifetime>
</ns2:RequestSecurityTokenResponse>
</ns2:RequestSecurityTokenResponseCollection>
</soap:Body>
</soap:Envelope>

```

29.15. Developing Registry Clients

Registry Clients create Federated Sources using the OSGi Configuration Admin. Developers should reference an individual [Source's \(Federated, Connected, or Catalog Provider\)](#) documentation for the Configuration properties (such as a Factory PID, addresses, intervals, etc) necessary to establish that 'Source in the framework.

Creating a Source Configuration

```

org.osgi.service.cm.ConfigurationAdmin configurationAdmin = getConfigurationAdmin() ;
org.osgi.service.cm.Configuration currentConfiguration = configurationAdmin
.createFactoryConfiguration(getFactoryPid(), null);
Dictionary properties = new Dictionary() ;
properties.put(QUERY_ADDRESS_PROPERTY,queryAddress);
currentConfiguration.update( properties );

```

Note that the `QUERY_ADDRESS_PROPERTY` is specific to this Configuration and might not be required for every [Source](#). The properties necessary for creating a Configuration are different for every [Source](#).

29.16. Developing Resource Readers

A [ResourceReader](#) is a class that retrieves a resource or product from a native/external source and returns it to DDF. A simple example is that of a File [ResourceReader](#). It takes a file from the local file system and passes it back to DDF. New implementations can be created in order to support obtaining Resources from various Resource data stores.

29.16.1. Creating a New [ResourceReader](#)

Complete the following procedure to create a [ResourceReader](#).

1. Create a Java class that implements the `DDF.catalog.resource.ResourceReader` interface.
2. Deploy the OSGi bundled packaged service to the DDF run-time.

29.16.1.1. Implementing the `ResourceReader` Interface

```
public class TestResourceReader implements DDF.catalog.resource.ResourceReader
```

`ResourceReader` has a couple of key methods where most of the work is performed.

NOTE

URI

It is recommended to become familiar with the Java API `URI` class in order to properly build a `ResourceReader`. Furthermore, a `URI` should be used according to its [specification](#).

29.16.1.2. `retrieveResource`

```
public ResourceResponse retrieveResource( URI uri, Map<String, Serializable> arguments  
) throws IOException, ResourceNotFoundException, ResourceNotSupportedException;
```

This method is the main entry to the `ResourceReader`. It is used to retrieve a `Resource` and send it back to the caller (generally the `CatalogFramework`). Information needed to obtain the entry is contained in the `URI` reference. The `URI` Scheme will need to match a scheme specified in the `getSupportedSchemes` method. This is how the `CatalogFramework` determines which `ResourceReader` implementation to use. If there are multiple `ResourceReaders` supporting the same scheme, these `ResourceReaders` will be invoked iteratively. Invocation of the `ResourceReaders` stops once one of them returns a `Resource`.

Arguments are also passed in. These can be used by the `ResourceReader` to perform additional operations on the resource.

The `URLResourceReader` is an example `ResourceReader` that reads a file from a `URI`.

NOTE

The `Map<String, Serializable> arguments` parameter is passed in to support any options or additional information associated with retrieving the resource.

29.16.1.3. Implement `retrieveResource()`

1. Define supported schemes (e.g., file, http, etc.).
2. Check if the incoming `URI` matches a supported scheme. If it does not, throw `ResourceNotSupportedException`.

Example:

```
if ( !uri.getScheme().equals("http") )  
{  
    throw new ResourceNotSupportedException("Unsupported scheme received, was expecting  
http")  
}
```

1. Implement the business logic.
2. For example, the `URLResourceReader` will obtain the resource through a connection:

```
URL url = uri.toURL();
URLConnection conn = url.openConnection();
String mimeType = conn.getContentType();
if ( mimeType == null ) {
    mimeType = URLConnection.guessContentTypeFromName( url.getFile() );
}
InputStream is = conn.getInputStream();
```

NOTE The `Resource` needs to be accessible from the DDF installation (see the `rootResourceDirectories` property of the `URLResourceReader`). This includes being able to find a file locally or reach out to a remote URI. This may require Internet access, and DDF may need to be configured to use a proxy (`http.proxyHost` and `http.proxyPort` can be added to the system properties on the command line script).

1. Return `Resource` in `ResourceResponse`.

For example:

```
return ResourceResponseImpl( new ResourceImpl( new BufferedInputStream( is ), new
MimeType( mimeType ), url.getFile() ) );
```

If the `Resource` cannot be found, throw a `ResourceNotFoundException`.

29.16.1.4. `getSupportedSchemes`

```
public Set<String> getSupportedSchemes();
```

This method lets the `ResourceReader` inform the `CatalogFramework` about the type of URI scheme that it accepts and should be passed. For single-use `ResourceReaders` (like a `URLResourceReader`), there may be only one scheme that it can accept while others may understand more than one. A `ResourceReader` must, at minimum, accept one qualifier. As mentioned before, this method is used by the `CatalogFramework` to determine which `ResourceReader` to invoke.

ResourceReader extends Describable

NOTE Additionally, there are other methods that are used to uniquely describe a `ResourceReader`. The `describe` methods are straight-forward and can be implemented with guidance from the Javadoc.

29.16.1.5. Export to OSGi Service Registry

In order for the `ResourceReader` to be used by the `CatalogFramework`, it should be exported to the OSGi Service Registry as a `DDF.catalog.resource.ResourceReader`.

See the XML below for an example:

Blueprint example

```
<bean id="[[customResourceReaderId]]" class=
"[[example.resource.reader.impl.CustomResourceReader]]" />
<service ref="[[customResourceReaderId]]" interface="
DDF.catalog.source.ResourceReader" />
```

29.17. Developing Resource Writers

A `ResourceWriter` is an object used to store or delete a `Resource`. `ResourceWriter` objects should be registered within the OSGi Service Registry, so clients can retrieve an instance when they need to store a `Resource`.

29.17.1. Create a New `ResourceWriter`

Complete the following procedure to create a `ResourceWriter`.

1. Create a Java class that implements the `DDF.catalog.resource.ResourceWriter` interface.

ResourceWriter Implementation Skeleton

```
import java.io.IOException;
import java.net.URI;
import java.util.Map;
import DDF.catalog.resource.Resource;
import DDF.catalog.resource.ResourceNotFoundException;
import DDF.catalog.resource.ResourceNotSupportedException;
import DDF.catalog.resource.ResourceWriter;

public class SampleResourceWriter implements ResourceWriter {

    @Override
    public void deleteResource(URI uri, Map<String, Object> arguments) throws
ResourceNotFoundException, IOException {
        // WRITE IMPLEMENTATION
    }

    @Override
    public URI storeResource(Resource resource, Map<String, Object> arguments) throws
ResourceNotSupportedException, IOException {
        // WRITE IMPLEMENTATION
        return null;
    }

    @Override
    public URI storeResource(Resource resource, String id, Map<String, Object>
arguments) throws ResourceNotSupportedException, IOException {
        // WRITE IMPLEMENTATION
        return null;
    }

}
```

1. Register the implementation as a Service in the OSGi Service Registry.

Blueprint Service Registration Example

```
...
<service ref="[[ResourceWriterReference]]" interface=
"DDF.catalog.resource.ResourceWriter" />
...
```

1. Deploy the OSGi bundled packaged service to the DDF run-time (Refer to the Working with OSGI - Bundles section.)

ResourceWriter Javadoc

TIP Refer to the Catalog API Javadoc for more information about the methods required for implementing the interface.

29.18. Developing Filters

The common way to create a [Filter](#) is to use the GeoTools [FilterFactoryImpl](#) object, which provides Java implementations for the various types of filters in the Filter Specification. Examples are the easiest way to understand how to properly create a [Filter](#) and a [Query](#).

NOTE Refer to the [GeoTools javadoc](#) for more information on [FilterFactoryImpl](#).

WARNING Implementing the Filter interface directly is only for extremely advanced use cases and is highly discouraged. Instead, use of the DDF-specific [FilterBuilder](#) API is recommended.

Developers create a [Filter](#) object in order to filter or constrain the amount of records returned from a [Source](#). The OGC Filter Specification has several types of filters that can be combined in a tree-like structure to describe the set of metacards that should be returned.

Categories of Filters

- Comparison Operators
- Logical Operators
- Expressions
- Literals
- Functions
- Spatial Operators
- Temporal Operators

29.18.1. Units of Measure

According to the [OGC Filter Specifications: 09-026r1](#) and [OGC Filter Specifications: 04-095](#), units of measure can be expressed as a URI. To fulfill that requirement, DDF utilizes the GeoTools class [org.geotools.styling.UomOgcMapping](#) for spatial filters requiring a standard for units of measure for scalar distances. Essentially, the [UomOgcMapping](#) maps the [OGC Symbology Encoding](#) standard URIs to Java Units. This class provides three options for units of measure:

- FOOT
- METRE
- PIXEL

DDF only supports FOOT and METRE since they are the most applicable to scalar distances.

29.18.2. Filter Examples

The example below illustrates creating a query, and thus an OGC Filter, that does a case-insensitive search for the phrase "mission" in the entire metacard's text. Note that the OGC [PropertyIsLike](#) Filter is used for this simple contextual query.

Simple Contextual Search

```
org.opengis.filter.FilterFactory filterFactory = new FilterFactoryImpl() ;
boolean isCaseSensitive = false ;

String wildcardChar = "*" ; // used to match zero or more characters
String singleChar = "?" ; // used to match exactly one character
String escapeChar = "\\" ; // used to escape the meaning of the wildCard, singleChar,
and the escapeChar itself

String searchPhrase = "mission" ;
org.opengis.filter.Filter propertyIsLikeFilter =
    filterFactory.like(filterFactory.property(Metacard.ANY_TEXT), searchPhrase,
    wildcardChar, singleChar, escapeChar, isCaseSensitive);
DDF.catalog.operation.QueryImpl query = new QueryImpl( propertyIsLikeFilter );
```

The example below illustrates creating an absolute temporal query, meaning the query is searching for Metacards whose modified timestamp occurred during a specific time range. Note that this query uses the **During** OGC Filter for an absolute temporal query.

Absolute Temporal Search

```
org.opengis.filter.FilterFactory filterFactory = new FilterFactoryImpl() ;
org.opengis.temporal.Instant startInstant = new org.geotools.temporal.object
.DefaultInstant(new DefaultPosition(start));

org.opengis.temporal.Instant endInstant = new org.geotools.temporal.object
.DefaultInstant(new DefaultPosition(end));

org.opengis.temporal.Period period = new org.geotools.temporal.object.DefaultPeriod
(startInstant, endInstant);

String property = Metacard.MODIFIED ; // modified date of a metacard

org.opengis.filter.Filter filter = filterFactory.during( filterFactory.property
(property), filterFactory.literal(period) ) ;

DDF.catalog.operation.QueryImpl query = new QueryImpl(filter) ;
```

29.18.2.1. Contextual Searches

Most contextual searches can be expressed using the **PropertyIsLike** filter. The special characters that have meaning in a **PropertyIsLike** filter are the wildcard, single wildcard, and escape characters (see Example Creating-Filters-1).

Table 78. **PropertyIsLike** Special Characters

Character	Description
Wildcard	Matches zero or more characters.

Character	Description
Single Wildcard	Matches exactly one character.
Escape	Escapes the meaning of the Wildcard, Single Wildcard, and the Escape character itself

Characters and words, such as **AND**, **&**, **and**, **OR**, **|**, **or**, **NOT**, **~**, **not**, **{**, and **}**, are treated as literals in a **PropertyIsLike** filter. In order to create equivalent logical queries, a developer must instead use the Logical Operator filters **{AND, OR, NOT}**. The Logical Operator filters can be combined together with **PropertyIsLike** filters to create a tree that represents the search phrase expression.

Creating the search phrase "mission and planning"

```
org.opengis.filter.FilterFactory filterFactory = new FilterFactoryImpl() ;

boolean isCaseSensitive = false ;


String wildcardChar = "*" ; // used to match zero or more characters
String singleChar = "?" ; // used to match exactly one character
String escapeChar = "\\" ; // used to escape the meaning of the wildCard, singleChar,
and the escapeChar itself

Filter filter =
    filterFactory.and(
        filterFactory.like(filterFactory.property(Metacard.METADATA), "mission" ,
wildcardChar, singleChar, escapeChar, isCaseSensitive),
        filterFactory.like(filterFactory.property(Metacard.METADATA), "planning" ,
wildcardChar, singleChar, escapeChar, isCaseSensitive)
    );

DDF.catalog.operation.QueryImpl query = new QueryImpl( filter );
```

29.18.2.1.1. Tree View of Creating Filters

Filters used in DDF can always be represented in a tree diagram.



Filter Example Tree Diagram

29.18.2.1.2. XML View of Creating Filters

Another way to view this type of Filter is through an XML model, which is shown below.

Pseudo XML of Example Creating-Filters-3

```
<Filter>
  <And>
    <PropertyIsLike wildCard="*" singleChar="?" escapeChar="\">
      <PropertyName>metadata</PropertyName>
      <Literal>mission</Literal>
    </PropertyIsLike>
    <PropertyIsLike wildCard="*" singleChar="?" escapeChar="\">
      <PropertyName>metadata</PropertyName>
      <Literal>planning</Literal>
    </PropertyIsLike>
  <And>
</Filter>
```

Using the Logical Operators and `PropertyIsLike` filters, a developer can create a whole language of search phrase expressions.

29.18.2.2. Fuzzy Operations

DDF only supports one custom function. The Filter specification does not include a fuzzy operator, so a Filter function was created to represent a fuzzy operation. The function and class is called `FuzzyFunction`, which is used by clients to notify the Sources to perform a fuzzy search. The syntax expected by providers is similar to the Fuzzy Function. Refer to the example below.

```
String wildcardChar = "*" ; // used to match zero or more characters
String singleChar = "?" ; // used to match exactly one character
String escapeChar = "\\" ; // used to escape the meaning of the wildCard, singleChar

boolean isCaseSensitive = false ;

Filter fuzzyFilter = filterFactory.like(
  new DDF.catalog.impl.filter.FuzzyFunction(
    Arrays.asList(Expression (filterFactory.property(Metocard.ANY_TEXT))),
    filterFactory.literal("")),
  searchPhrase,
  wildcardChar,
  singleChar,
  escapeChar,
  isCaseSensitive);

QueryImpl query = new QueryImpl(fuzzyFilter);
```

29.18.3. Parsing Filters


According to the [OGC Filter Specification 04-095](#): a "(filter expression) representation can be ... parsed and then transformed into whatever target language is required to retrieve or modify object instances stored in some persistent object store." Filters can be thought of as the **WHERE** clause for a SQL SELECT statement to "fetch data stored in a SQL-based relational database."

Sources can parse OGC Filters using the **FilterAdapter** and **FilterDelegate**. See [Developing a Filter Delegate](#) for more details on implementing a new **FilterDelegate**. This is the preferred way to handle OGC Filters in a consistent manner.

Alternately, `org.opengis.filter.Filter` implementations can be parsed using implementations of the interface `org.opengis.filter.FilterVisitor`. The `FilterVisitor` uses the [Visitor pattern](#). Essentially, `FilterVisitor` instances "visit" each part of the `Filter` tree allowing developers to implement logic to handle the filter's operations. GeoTools 8 includes implementations of the `FilterVisitor` interface. The `DefaultFilterVisitor`, as an example, provides only business logic to visit every node in the `Filter` tree. The `DefaultFilterVisitor` methods are meant to be overwritten with the correct business logic. The simplest approach when using `FilterVisitor` instances is to build the appropriate query syntax for a target language as each part of the `Filter` is visited. For instance, when given an incoming `Filter` object to be evaluated against a RDBMS, a `CatalogProvider` instance could use a 'FilterVisitor' to interpret each filter operation on the `Filter` object and translate those operations into SQL. The `FilterVisitor` may be needed to support `Filter` functionality not currently handled by the `FilterAdapter` and `FilterDelegate` reference implementation.

29.18.3.1. Interpreting a Filter to Create SQL

If the `FilterAdapter` encountered or "visited" a `PropertyIsLike` filter with its property assigned as `title` and its literal expression assigned as `mission`, the `FilterDelegate` could create the proper SQL syntax similar to `title LIKE mission`.




Parsing Filters Tree Diagram

29.18.3.2. Interpreting a Filter to Create XQuery

If the `FilterAdapter` encountered an `OR` filter, such as in Figure [Parsing-Filters2](#) and the target language was XQuery, the `FilterDelegate` could yield an expression such as

```
ft:query(/inventory:book/@subject,'math') union
ft:query(/inventory:book/@subject,'science').
```



Parsing Filters XQuery

29.18.3.2.1. FilterAdapter/Delegate Process for Figure Parsing

1. **FilterAdapter** visits the **OR** filter first.
2. **OR** filter visits its children in a loop.
3. The first child in the loop that is encountered is the LHS **PropertyIsLike**.
4. The **FilterAdapter** will call the **FilterDelegate** `PropertyIsLike` method with the LHS property and literal.
5. The LHS **PropertyIsLike** delegate method builds the XQuery syntax that makes sense for this particular underlying object store. In this case, the *subject* property is specific to this XML database, and the business logic maps the *subject* property to its index at `//inventory:book/@subject`. Note that `ft:query` in this instance is a custom XQuery module for this specific XML database that does full text searches.
6. The **FilterAdapter** then moves back to the **OR** filter, which visits its second child.
7. The **FilterAdapter** will call the **FilterDelegate** **PropertyIsLike** method with the RHS property and literal.
8. The RHS **PropertyIsLike** delegate method builds the XQuery syntax that makes sense for this particular underlying object store. In this case, the *subject* property is specific to this XML database, and the business logic maps the *subject* property to its index at `//inventory:book/@subject`. Note that `ft:query` in this instance is a custom XQuery module for this specific XML database that does full text searches. . The **FilterAdapter** then moves back to its **OR Filter** which is now done with its children.
9. It then collects the output of each child and sends the list of results to the **FilterDelegate** **OR** method.

10. The final result object will be returned from the `FilterAdapter adapt` method.

29.18.3.2.2. FilterVisitor Process for Figure Parsing

1. FilterVisitor visits the `OR` filter first.
2. `OR` filter visits its children in a loop.
3. The first child in the loop that is encountered is the LHS `PropertyIsLike`.
4. The LHS `PropertyIsLike` builds the XQuery syntax that makes sense for this particular underlying object store. In this case, the `subject` property is specific to this XML database, and the business logic maps the `subject` property to its index at `//inventory:book/@subject`. Note that `ft:query` in this instance is a custom XQuery module for this specific XML database that does full text searches.
5. The FilterVisitor then moves back to the `OR` filter, which visits its second child.
6. The RHS `PropertyIsLike` builds the XQuery syntax that makes sense for this particular underlying object store. In this case, the `subject` property is specific to this XML database, and the business logic maps the `subject` property to its index at `//inventory:book/@subject`. Note that `ft:query` in this instance is a custom XQuery module for this specific XML database that does full text searches.
7. The FilterVisitor then moves back to its `OR` filter, which is now done with its children. It then collects the output of each child and could potentially execute the following code to produce the above expression.

```
public visit( Or filter, Object data) {  
    ...  
    /* the equivalent statement for the OR filter in this domain (XQuery) */  
    xQuery = childFilter1Output + " union " + childFilter2Output;  
    ...  
}
```

29.18.4. Filter Profile

The filter profile maps filters to metocard types.

29.18.4.1. Role of the OGC Filter

Both Queries and Subscriptions extend the OGC GeoAPI Filter interface.

The Filter Builder and Adapter do not fully implement the OGC Filter Specification. The filter support profile contains suggested filter to metocard type mappings. For example, even though a Source could support a `PropertyIsGreaterThan` filter on `XML_TYPE`, it would not likely be useful.

29.18.4.2. Catalog Filter Profile

The following table displays the common metocard attributes with their respective types for reference.

Table 79. Metocard Attribute To Type Mapping

Metocard Attribute	Metocard Type
ANY_DATE	DATE_TYPE
ANY_GEO	GEO_TYPE
ANY_TEXT	STRING_TYPE
CONTENT_TYPE	STRING_TYPE
CONTENT_TYPE_VERSION	STRING_TYPE
CREATED	DATE_TYPE
EFFECTIVE	DATE_TYPE
GEOGRAPHY	GEO_TYPE
ID	STRING_TYPE
METADATA	XML_TYPE
MODIFIED	DATE_TYPE
RESOURCE_SIZE	STRING_TYPE
RESOURCE_URI	STRING_TYPE
SOURCE_ID	STRING_TYPE
TARGET_NAMESPACE	STRING_TYPE
THUMBNAIL	BINARY_TYPE
TITLE	STRING_TYPE

29.18.4.2.1. Comparison Operators

Comparison operators compare the value associated with a property name with a given Literal value. Endpoints and sources should try to use metocard types other than the object type. The object type only supports backwards compatibility with `java.net.URI`. Endpoints that send other objects will not be supported by standard sources. The following table maps the metocard types to supported comparison operators.

Table 80. Metocard Types to Comparison Operators

PropertyIs	Between	EqualTo	GreaterThanOrEqual	GreaterThanOrEqual	OrEqualTo	LessThan	LessThan	OrEqualTo	Like	NotEqualTo	Null
BINARY_TYPE		X									
BOOLEAN_TYPE		X									
DATE_TYPE	X	X	X	X	X	X	X	X	X	X	
DOUBLE_TYPE	X	X	X	X	X	X	X	X	X	X	

PropertyIs	Between	Equal To	Greater Than	Greater Than	OrEqualTo	LessThan	LessThan	OrEqualTo	Like	NotEqualTo	Null
FLOAT_TYPE	X	X	X	X	X	X	X	X		X	X
GEO_TYPE											X
INTEGER_TYPE	X	X	X	X	X	X	X	X		X	X
LONG_TYPE	X	X	X	X	X	X	X	X		X	X
OBJECT_TYPE	X	X	X	X	X	X	X	X		X	X
SHORT_TYPE	X	X	X	X	X	X	X	X		X	X
STRING_TYPE	X	X	X	X	X	X	X	X	X	X	X
XML_TYPE		X							X		X

Table 81. Comparison Operators

Operator	Description
PropertyIsBetween	Lower \Leftarrow Property \Leftarrow Upper
PropertyIsEqualTo	Property == Literal
PropertyIsGreaterThan	Property > Literal
PropertyIsGreaterThanOrEqualTo	Property \geq Literal
PropertyIsLessThan	Property < Literal
PropertyIsLessThanOrEqualTo	Property \leq Literal
PropertyIsLike	Property LIKE Literal
	Equivalent to SQL "like"
PropertyIsNotEqualTo	Property \neq Literal
PropertyIsNull	Property == null

29.18.4.2.2. Logical Operators

Logical operators apply Boolean logic to one or more child filters.

Table 82. Supported Logical Operators

	And	Not	Or
Supported Filters	X	X	X

29.18.4.2.3. Temporal Operators

Temporal operators compare a date associated with a property name to a given Literal date or date range.

Table 83. Supported Temporal Operators

	After	AnyInteracts	Before	Begins	Begun By	During	Ended By	Meets	MetBy	OverlappedBy	TContains
DATE_TYPE	X		X			X					

Literal values can be either date instants or date periods.

Table 84. Temporal Operator Descriptions

Operator	Description
After	Property > (Literal Literal.end)
Before	Property < (Literal Literal.start)
During	Literal.start < Property < Literal.end

29.18.4.2.4. Spatial Operators

Spatial operators compare a geometry associated with a property name to a given Literal geometry.

Table 85. Supported Spatial Operators.

BBox	Beyond	Contains	Crosses	Disjoint	Equals	DWithin	Intersects	Overlaps	Touches	Within
GEO_TYPE		X	X	X	X		X	X	X	

Geometries are usually represented as Well-Known Text (WKT).

Table 86. Spatial Operator Descriptions

Operator	Description
Beyond	Property geometries beyond given distance of Literal geometry
Contains	Property geometry contains Literal geometry
Crosses	Property geometry crosses Literal geometry
Disjoint	Property geometry direct positions are not interior to Literal geometry
DWithin	Property geometry lies within distance to Literal geometry

Operator	Description
Intersects	Property geometry intersects Literal geometry; opposite to the Disjoint operator
Overlaps	Property geometry interior overlaps Literal geometry interior somewhere
Touches	Property geometry touches but does not overlap Literal geometry
Within	Property geometry completely contains Literal geometry

29.19. Developing Filter Delegates

Filter Delegates help reduce the complexity of parsing OGC Filters. The reference Filter Adapter implementation contains the necessary boilerplate visitor code and input normalization to handle commonly supported OGC Filters.

29.19.1. Creating a New Filter Delegate

A Filter Delegate contains the logic that converts normalized filter input into a form that the target data source can handle. Delegate methods will be called in a depth first order as the Filter Adapter visits filter nodes.

29.19.1.1. Implementing the Filter Delegate

1. Create a Java class extending `FilterDelegate`.

```
public class ExampleDelegate extends DDF.catalog.filter.FilterDelegate<ExampleReturnObjectType> {
```

2. `FilterDelegate` will throw an appropriate exception for all methods not implemented. Refer to the DDF JavaDoc for more details about what is expected of each `FilterDelegate` method.

NOTE

A code example of a Filter Delegate can be found in `DDF.catalog.filter.proxy.adapter.test` of the `filter-proxy` bundle.

29.19.1.2. Throwing Exceptions

Filter delegate methods can throw `UnsupportedOperationException` run-time exceptions. The `GeotoolsFilterAdapterImpl` will catch and re-throw these exceptions as `UnsupportedQueryExceptions`.

29.19.1.3. Using the Filter Adapter

The FilterAdapter can be requested from the OSGi registry.

```
<reference id="filterAdapter" interface="DDF.catalog.filter.FilterAdapter" />
```

The Query in a QueryRequest implements the Filter interface. The Query can be passed to a `FilterAdapter` and `FilterDelegate` to process the Filter.

```

@Override
public DDF.catalog.operation.QueryResponse query(DDF.catalog.operation.QueryRequest
queryRequest)
    throws DDF.catalog.source.UnsupportedQueryException {

    DDF.catalog.operation.Query query = queryRequest.getQuery();

    DDF.catalog.filter.FilterDelegate<ExampleReturnObjectType> delegate = new
ExampleDelegate();

    // DDF.catalog.filter.FilterAdapter adapter injected via Blueprint
    ExampleReturnObjectType result = adapter.adapt(query, delegate);
}

```

Import the Catalog API Filter package and the reference implementation package of the Filter Adapter in the bundle manifest (in addition to any other required packages).

Import-Package: `DDF.catalog, DDF.catalog.filter, DDF.catalog.source`

29.19.1.4. Filter Support

Not all OGC Filters are exposed at this time. If demand for further OGC Filter functionality is requested, it can be added to the Filter Adapter and Delegate so sources can support more complex filters. The following OGC Filter types are currently available:

Logical

And

Or

Not

Include

Exclude

Property Comparison

`PropertyIsBetween`

`PropertyIsEqualTo`

`PropertyIsGreater Than`

`PropertyIsGreater ThanOrEqual To`

`PropertyIsLess Than`

`PropertyIsLess ThanOrEqual To`

`PropertyIsLike`

`PropertyIsNotEqual To`

`PropertyIsNull`

Spatial

Definition

`Beyond`

True if the geometry being tested is beyond the stated distance of the geometry provided.

Spatial	Definition
Contains	True if the second geometry is wholly inside the first geometry.
Crosses	True if: * the intersection of the two geometries results in a value whose dimension is less than the geometries * the maximum dimension of the intersection value includes points interior to both the geometries * the intersection value is not equal to either of the geometries.
Disjoint	True if the two geometries do not touch or intersect.
DWithin	True if the geometry being tested is within the stated distance of the geometry provided.
Intersects	True if the two geometries intersect. This is a convenience method as <code>Not Disjoint(A,B)</code> gets the same result.
Overlaps	True if the intersection of the geometries results in a value of the same dimension as the geometries that is different from both of the geometries.
Touches	True if and only if the only common points of the two geometries are in the union of the boundaries of the geometries.
Within	True if the first geometry is wholly inside the second geometry.

Temporal
After
Before
During

29.20. Developing Action Components

To provide a service, such as a link to a metocard, the `ActionProvider` interface should be implemented. An `ActionProvider` essentially provides a List of `Actions` when given input that it can recognize and handle. For instance, if a REST endpoint ActionProvider was given a metocard, it could provide a link based on the metocard's ID. An Action Provider performs an action when given a subject that it understands. If it does not understand the subject or does not know how to handle the given input, it will return `Collections.emptyList()`. An Action Provider is required to have an `ActionProvider` id. The Action Provider must register itself in the OSGi Service Registry with the `ddf.action.ActionProvider` interface and must also have a service property value for `id`. An action is a URL that, when invoked, provides a resource or executes intended business logic.

29.20.1. Action Component Naming Convention

For each Action, a title and description should be provided to describe what the action does. The recommended naming convention is to use the verb 'Get' when retrieving a portion of a metocard, such as the metadata or thumbnail, or when downloading a product. The verb 'Export' or the expression 'Export as' is recommended when the metocard is being exported in a different format or presented after going some transformation.

29.20.1.1. Action Component Taxonomy

An Action Provider registers an `id` as a service property in the OGSI Service Registry based on the type of service or action that is provided. Regardless of implementation, if more than one Action Provider provides the same service, such as providing a URL to a thumbnail for a given metocard, they must both register under the same `id`. Therefore, Action Provider implementers must follow an Action Taxonomy.

The following is a sample taxonomy:

1. `catalog.data.metocard` shall be the grouping that represents Actions on a Catalog metocard.
 - a. `catalog.data.metocard.view`
 - b. `catalog.data.metocard.thumbnail`
 - c. `catalog.data.metocard.html`
 - d. `catalog.data.metocard.resource`
 - e. `catalog.data.metocard.metadata`

Table 87. Action ID Service Descriptions

ID	Required Action	Naming Convention
<code>catalog.data.metocard.view</code>	Provides a valid URL to view a metocard. Format of data is not specified; i.e. the representation can be in XML, JSON, or other.	Export as ...
<code>catalog.data.metocard.thumbnail</code>	Provides a valid URL to the bytes of a thumbnail (<code>Metocard.THUMBNAIL</code>) with MIME type image/jpeg.	Export as Thumbnail
<code>catalog.data.metocard.map.overlay.thumbnail</code>	Provides a metocard URL that translates the metocard into a geographically aligned image (suitable for overlaying on a map).	Export as Thumbnail Overlay
<code>catalog.data.metocard.html</code>	Provides a valid URL that, when invoked, provides an HTML representation of the metocard.	Export as HTML
<code>catalog.data.metocard.xml</code>	Provides a valid URL that, when invoked, provides an XML representation of the metocard.	Export as XML
<code>catalog.data.metocard.geojson</code>	Provides a valid URL that, when invoked, provides an XML representation of the metocard.	Export as GeoJSON
<code>catalog.data.metocard.resource</code>	Provides a valid URL that, when invoked, provides the underlying resource of the metocard.	Export as Resource
<code>catalog.data.metocard.metadata</code>	Provides a valid URL to the XML metadata in the metocard (<code>Metocard.METADATA</code>).	Export as Metadata

29.21. Developing Query Options

The easiest way to create a Query is to use the `ddf.catalog.operation.QueryImpl` object. It is first necessary to create an OGC Filter object then set the Query Options after `QueryImpl` has been constructed.

QueryImpl Example

```
/*
Builds a query that requests a total results count and
that the first record to be returned is the second record found from
the requested set of metacards.
*/
String property = ...;

String value = ...;

org.geotools.filter.FilterFactoryImpl filterFactory = new FilterFactoryImpl() ;

QueryImpl query = new QueryImpl( filterFactory.equals(filterFactory.property(
property),
filterFactory.literal(value))) ;

query.setstartIndex(2) ;

query.setRequestsTotalResultsCount(true);
```

29.21.1. Evaluating a query

Every Source must be able to evaluate a Query object. Nevertheless, each Source could evaluate the Query differently depending on what that Source supports as to properties and query capabilities. For instance, a common property all Sources understand is id, but a Source could possibly store frequency values under the property name "frequency." Some Sources may not support frequency property inquiries and will throw an error stating it cannot interpret the property. In addition, some Sources might be able to handle spatial operations, while others might not. A developer should consult a Source's documentation for the limitations, capabilities, and properties that a Source can support.

Table 88. Query Options

Option	Description
startIndex	1-based index that states which metocard the Source should return first out of the requested metacards.
pageSize	Represents the maximum amount of metacards the Source should return.
sortBy	Determines how the results are sorted and on which property.
requestsTotalResultsCount	Determines whether the total number of results should be returned.

Option	Description
<code>TimeoutMillis</code>	The amount of time in milliseconds before the query is to be abandoned. If a zero or negative timeout is set, the catalog framework will default to a value configurable via the Admin UI under Catalog → Configuration → Query Operations.

29.21.2. Commons-DDF Utilities

The `commons-DDF` bundle provides utilities and functionality commonly used across other DDF components, such as the endpoints and providers.

29.21.2.1. FuzzyFunction

`DDF.catalog.impl.filter.FuzzyFunction` class is used to indicate that a `PropertyIsLike` filter should interpret the search as a fuzzy query.

29.21.2.2. XPathHelper

`DDF.util.XPathHelper` provides convenience methods for executing XPath operations on XML. It also provides convenience methods for converting XML as a `String` from a `org.w3c.dom.Document` object and vice versa.

29.22. Configuring Managed Service Factory Bundles

29.22.1. Configuring Managed Service Factory Bundles

Services that are created using a Managed Service Factory can be configured using `.config` files as well. These configuration files, however, follow a different naming convention than `.cfg` files. The filenames must start with the Managed Service Factory PID, be followed by a dash and a unique identifier, and have a `.config` extension. For instance, assuming that the Managed Service Factory PID is `org.codice.ddf.factory.pid` and two instances of the service need to be configured, files `org.codice.ddf.factory.pid-<UNIQUE_ID_1>.config` and `org.codice.ddf.factory.pid-<UNIQUE_ID_2>.config` should be created and added to `<DDF_HOME>/etc`.

The unique identifiers used in the file names have no impact on the order in which the configuration files are processed. No specific processing order should be assumed. Also, a new service will be created and configured every time a configuration file matching the Managed Service Factory PID is added to the directory, regardless of the *unique id* used.

Any `service.factoryPid` and `service.pid` values in these `.config` files will be overridden by the values parsed from the file name, so `.config` files should not contain these properties.

29.22.1.1. File Format

The basic syntax of the `.config` configuration files is similar to the older `.cfg` files but introduces support for lists and types other than simple strings. The type associated with a property must match the type attribute used in the corresponding `metatype.xml` file when applicable.

The following table shows the format to use for each property type supported.

Table 89. Property Formats

Type	Format (see details below for variations)	Example
String	name="value"	name="John"
Boolean	name=B"true false"	authorized=B"true"
Integer	name=I"value"	timeout=I"10"
Long	name=L"value"	diameter=L"100"
Float	name=F"value"	cost=F"1093140480"
Double	name=D"value"	latitude=D"4636745974857667812"
List of Strings	name=["value1", "value2", ...]	<pre>complexStringArray=[\ "{\"url\": \"http://test.sample.com\"\ \"layers\": [\"0\"]\\" VERSION\"\ \"1.1 1.2\"\\ \"image/png\"\\"}\ \"beta\"\ \"1\", \ "{\"url\": \"http://test.sample.com\"\ \"0.5\", \ \"/solr\\"=SAML PKI basic\", \ \"/security-config\\"=SAML basic\", \ \"\"}]</pre>
List of Booleans	name=B["true false", "true false", ...]	<pre>authorizedList=B[\ \"true\", \ \"false\", \ \"\"]</pre>
List of Integers	name=I["value1", "value2", ...]	<pre>sizes=I[\ \"10\", \ \"20\", \ \"30\", \ \"\"]</pre>

Type	Format (see details below for variations)	Example
List of Longs	name=L["value1","value2",...]	<pre>sizes=L[\ "100", \ "200", \ "300", \]</pre>
List of Floats	name=F["value1","value2",...]	<pre>sizes=F[\ "1066192077", \ "1074580685", \ "1079194419", \]</pre>
List of Doubles	name=D["value1","value2",...]	<pre>sizes=D[\ "4607736361554183979", \ "4612212939583790252", \ "4614714689176794563", \]</pre>

- Values with types other than String must be prefixed with a lower-case or upper-case character. See the examples in the table.
 - Boolean: `B` or `b`
 - Integer: `I` or `i`
 - Long: `L` or `l`
 - Float: `F` or `f`
 - Double: `D` or `d`
- Equal signs (`=`), double quotes (`"`), and spaces within values must be escaped using a backslash (`\`).
- When properties are split over multiple lines for readability, end of lines must be specified with a backslash (`\`). See the examples for lists in the table.
- A comma (,) after the last value in a list is optional.

NOTE

- Surrounding the equal signs (`=`) with spaces for properties is optional. Because there is a known issue when using OPS4J Pax Exam 4.11.0 and modifying `.config` files that include spaces, all default `.config` files that may be modified in OPS4J Pax Exam 4.11.0 tests should not include spaces.
- Boolean values will default to `false` if any value other than `true` is provided.
- Float values must be represented in the IEEE 754 floating-point "single format" bit layout, preserving Not-a-Number (NaN) values. For example, `F"1093140480"` corresponds to `F"10.5"`. See the documentation for `java.lang.Integer#parseInt(java.lang.String)` and `java.lang.Float#intBitsToFloat(int)` for more details.
- Double values must be represented in the IEEE 754 floating-point "double format" bit layout, preserving Not-a-Number (NaN) values. For example, `D"4636745974857667812"` corresponds to `D"100.1234"`. See the documentation for `java.lang.Long#parseLong(java.lang.String)` and `java.lang.Double#longBitsToDouble` for more details.

```
authenticationTypes=[ \
  "/\=SAML|GUEST", \
  "/admin\=SAML|basic", \
  "/system\=basic", \
  "/solr\=SAML|PKI|basic", \
  "/sources\=SAML|basic", \
  "/security-config\=SAML|basic", \
  "/search\=basic", \
]
realms=[ \
  "/\=karaf", \
]
requiredAttributes=[ \
  "/\=", \
  "/admin\={http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role\=admin}", \
  "/solr\={http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role\=admin}", \
  "/system\={http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role\=admin}", \
  "/security- \
config\={http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role\=admin}", \
]
whiteListContexts=[ \
  "/services/SecurityTokenService", \
  "/services/internal/metrics", \
  "/services/saml", \
  "/proxy", \
  "/services/csw", \
]
```

29.23. Developing XACML Policies

This document assumes familiarity with the XACML schema and does not go into detail on the XACML language. When creating a policy, a target is used to indicate that a certain action should be run only for one type of request. Targets can be used on both the main policy element and any individual rules. Targets are geared toward the actions that are set in the request. These actions generally consist of the standard CRUD operations (create, read, update, delete) or a SOAPAction if the request is coming through a SOAP endpoint.

NOTE These are only the action values that are currently created by the components that come with DDF. Additional components can be created and added to DDF to identify specific actions.

In the examples below, the policy has specified targets for the above type of calls. For the Filtering code, the target was set for "filter", and the Service validation code targets were geared toward two services: `query` and `LocalSiteName`. In a production environment, these actions for service authorization will generally be full URNs that are described within the SOAP WSDL.

29.23.1. XACML Policy Attributes

Attributes for the XACML request are populated with the information in the calling subject and the resource being checked.

29.23.2. XACML Policy Subject

The attributes for the subject are obtained from the SAML claims and populated within the XACML policy as individual attributes under the `urn:oasis:names:tc:xacml:1.0:subject-category:access-subject` category. The name of the claim is used for the `AttributeId` value. Examples of the items being populated are available at the end of this page.

29.23.3. XACML Policy Resource

The attributes for resources are obtained through the permissions process. When checking permissions, the XACML processing engine retrieves a list of permissions that should be checked against the subject. These permissions are populated outside of the engine and should be populated with the attributes that should be asserted against the subject. When the permissions are of a key-value type, the key being used is populated as the `AttributeId` value under the `urn:oasis:names:tc:xacml:3.0:attribute-category:resource` category.

29.23.4. Using a XACML Policy

To use a XACML policy, copy the XACML policy into the `<DDF_HOME>/etc/pdp/policies` directory.

29.24. Assuring Authenticity of Bundles and Applications

DDF Artifacts in the JAR file format (such as bundles or DDF applications packaged as KAR files) can be signed and verified using the tools included as part of the Java Runtime Environment.

29.24.1. Prerequisites

To work with Java signatures, a keystore/truststore is required. For testing or trial purposes DDF can sign and validate using a self-signed certificate, generated with the keytool utility. In an actual installation, a certificate issued from a trusted Certificate Authority will be used.

Additional documentation on keytool can be found at [Keytool home](#).

Using keytool to generate a self-signed certificate keystore

```
~ $ keytool -genkey -keyalg RSA -alias selfsigned -keystore keystore.jks -storepass
password -validity 360 -keysize 2048
What is your first and last name?
[Unknown]: Nick Fury
What is the name of your organizational unit?
[Unknown]: Marvel
What is the name of your organization?
[Unknown]: SHIELD
What is the name of your City or Locality?
[Unknown]: New York
What is the name of your State or Province?
[Unknown]: NY
What is the two-letter country code for this unit?
[Unknown]: US
Is CN=Nick Fury, OU=SHIELD, O=Marvel, L="New York", ST=NY, C=US correct?
[no]: yes
Enter key password for <selfsigned>
      (RETURN if same as keystore password):
Re-enter new password:
```

29.24.2. Signing a JAR/KAR

Once a keystore is available, the JAR can be signed using the [jarsigner](#) tool.

Additional documentation on jarsigner can be found at [Jarsigner](#).

Using jarsigner to sign a KAR

```
~ $ jarsigner -keystore keystore.jks -keypass shield -storepass password catalog-app-
2.5.1.kar selfsigned
```

29.24.2.1. Verifying a JAR/KAR

The jarsigner utility is also used to verify a signature in a JAR-formatted file.

Using jarsigner to verify a file

```
~ $ jarsigner -verify -verbose -keystore keystore.jks catalog-app-2.5.1.kar
    9447 Mon Oct 06 17:05:46 MST 2014 META-INF/MANIFEST.MF
    9503 Mon Oct 06 17:05:46 MST 2014 META-INF/SELFSIGN.SF

[... section abbreviated for space]

smk      6768 Wed Sep 17 17:13:58 MST 2014 repository/ddf/catalog/security/catalog-
security-logging/2.5.1/catalog-security-logging-2.5.1.jar
s = signature was verified
m = entry is listed in manifest
k = at least one certificate was found in keystore
i = at least one certificate was found in identity scope
jar verified.
```

Note the last line: *jar verified*. This indicates that the signatures used to sign the JAR (or in this case, KAR) were valid according to the trust relationships specified by the keystore.

29.25. WFS Services

The Web Feature Service (WFS) is an [Open Geospatial Consortium \(OGC\)](#) Specification. DDF supports the ability to integrate WFS 1.0 and WFS 2.0 Web Services.

NOTE

DDF does not include a supported WFS Web Service (Endpoint) implementation. Therefore, federation for 2 DDF instances is not possible via WFS.

WFS Features

When a query is issued to a WFS server, the output of the query is an XML document that contains a collection of feature member elements. Each WFS server can have one or more feature types with each type being defined by a schema that extends the WFS `featureMember` schema. The schema for each type can be discovered by issuing a `DescribeFeatureType` request to the WFS server for the feature type in question. The WFS source handles WFS capability discovery and requests for feature type description when an instance of the WFS source is configured and created.

See the [WFS v1.0.0 Source](#) or [WFS v2.0.0 Source](#) for more information about how to configure a WFS source.

Converting a WFS Feature

In order to expose WFS features to DDF clients, the WFS feature must be converted into the common data format of the DDF, a metocard. The OGC package contains a `GenericFeatureConverter` that attempts to populate mandatory metocard fields with properties from the WFS feature XML. All properties will be mapped directly to new attributes in the metocard. However, the `GenericFeatureConverter` may not be able to populate the default metocard fields with properties from the feature XML.

Creating a Custom Converter

To more accurately map WFS feature properties to fields in the metocard, a custom converter can

be created. The OGC package contains an interface, `FeatureConverter`, which extends the <http://xstream.codehaus.org/javadoc/com/thoughtworks/xstream/converters/Converter.htmlConverter> interface provided by the `XStream` project. XStream is an open source API for serializing XML into Java objects and vice-versa. Additionally, a base class, `AbstractFeatureConverter`, has been created to handle the mapping of many fields to reduce code duplication in the custom converter classes.

1. Create the `CustomConverter` class extending the `ogc.catalog.common.converter.AbstractFeatureConverter` class.

```
public class CustomConverter extends ogc.catalog.common.converter
    .AbstractFeatureConverter
```

2. Implement the `FeatureConverterFactory` interface and the `createConverter()` method for the `CustomConverter`. +

```
public class CustomConverterFactory implements FeatureConverterFactory {
    private final featureType;
    public CustomConverterFactory(String featureType) {
        this.featureType = featureType;
    }
    public FeatureConverter createConverter() {
        return new CustomConverter();
    }
    public String getFeatureType() {
        return featureType;
    }
}
```

+ . Implement the `unmarshal` method required by the `FeatureConverter` interface. The `createMetocardFromFeature(reader, metacardType)` method implemented in the `AbstractFeatureConverter` is recommended.

+

```
public Metocard unmarshal(HierarchicalStreamReader reader, UnmarshallingContext ctx) {
    MetocardImpl mc = createMetocardFromFeature(reader, metacardType);
    //set your feature specific fields on the metocard object here
    //
    //if you want to map a property called "beginningDate" to the Metocard.createdDate
    //field
    //you would do:
    mc.setCreatedDate(mc.getAttribute("beginningDate").getValue());
}
```

+ . Export the `ConverterFactory` to the OSGi registry by creating a `blueprint.xml` file for its bundle. The bean id and argument value must match the WFS Feature type being converted.

```

<?xml version="1.0" encoding="UTF-8"?>
<blueprint xmlns="http://www.osgi.org/xmlns/blueprint/v1.0.0" xmlns:cm=
"http://aries.apache.org/blueprint/xmlns/blueprint-cm/v1.1.0">
  <bean id="custom_type" class="com.example.converter.factory.CustomConverterFactory">
    <argument value="custom_type"/>
  </bean>
  <service ref="custom_type" interface=
"orgc.catalog.common.converter.factory.FeatureConverterFactory"/>
</blueprint>

```

30. Development Guidelines

30.1. Contributing

The Distributed Data Framework is free and open-source software offered under the GNU Lesser General Public License. The DDF is managed under the guidance of the [Codice Foundation](#). Contributions are welcomed and encouraged. Please visit the [Codice DDF Contributor Guidelines](#) and the [DDF source code repository](#) for more information.

30.2. OSGi Basics

DDF runs on top of an OSGi framework, a Java virtual machine (JVM), several choices of operating systems, and the physical hardware infrastructure. The items within the dotted line represent the standard DDF components.

DDF is a customized and branded distribution of [Apache Karaf](#). DDF could also be considered to be a more lightweight OSGi distribution, as compared to Apache ServiceMix, FUSE ESB, or Talend ESB, all of which are also built upon Apache Karaf. Similar to its peers, DDF incorporates ([additional upstream dependencies](#)).

The DDF framework hosts DDF applications, which are extensible by adding components via OSGi. The best example of this is the DDF Catalog (API), which offers extensibility via several types of Catalog Components. The DDF Catalog API serves as the foundation for several applications and resides in the applications tier.

The Catalog Components consist of [Endpoints](#), [Plugins](#), [Catalog Frameworks](#), [Sources](#), and [Catalog Providers](#). Customized components can be added to DDF.

Capability

A general term used to refer to an ability of the system.

Component

Represents a portion of an Application that can be extended.

Bundle

Java Archives (JARs) with special OSGi manifest entries.

Feature

One or more bundles that form an installable unit; defined by Apache Karaf but portable to other OSGi containers.

Application

One or more features that together form a cohesive collection of capabilities.

30.2.1. Packaging Capabilities as Bundles

Services and code are physically deployed to DDF using bundles. The bundles within DDF are created using the maven bundle plug-in. Bundles are Java JAR files that have additional metadata in the [MANIFEST.MF](#) that is relevant to an OSGi container.

The best resource for learning about the structure and headers in the manifest definition is in section 3.6 of the [OSGi Core Specification](#). The bundles within DDF are created using the [maven bundle plug-in](#), which uses the [BND tool](#).

Alternative Bundle Creation Methods

TIP

Using Maven is not necessary to create bundles. Many alternative tools exist, and OSGi manifest files can also be created by hand, although hand-editing should be avoided by most developers.

30.2.1.1. Creating a Bundle

30.2.1.1.1. Bundle Development Recommendations

Avoid creating bundles by hand or editing a manifest file

Many tools exist for creating bundles, notably the Maven Bundle plugin, which handle the details of OSGi configuration and automate the bundling process including generation of the manifest file.

Always make a distinction on which imported packages are [optional](#) or [required](#)

Requiring every package when not necessary can cause an unnecessary dependency ripple effect among bundles.

Embedding is an implementation detail

Using the [Embed-Dependency](#) instruction provided by the [maven-bundle-plugin](#) will insert the specified jar(s) into the target archive and add them to the [Bundle-ClassPath](#). These jars and their contained packages/classes are not for public consumption; they are for the internal implementation of this service implementation only.

Bundles should never be embedded

Bundles expose service implementations; they do not provide arbitrary classes to be used by other bundles.

Bundles should expose service implementations

This is the corollary to the previous rule. Bundles should not be created when arbitrary concrete classes are being extracted to a library. In that case, a library/jar is the appropriate module packaging type.

Bundles should generally *only* export service packages

If there are packages internal to a bundle that comprise its implementation but not its public manifestation of the API, they should be excluded from export and kept as private packages.

Concrete objects that are not loaded by the root classloader should not be passed in or out of a bundle

This is a general rule with some exceptions (JAXB generated classes being the most prominent example). Where complex objects need to be passed in or out of a service method, an interface should be defined in the API bundle.

Bundles separate contract from implementation and allow for modularized development and deployment of functionality. For that to be effective, they must be defined and used correctly so inadvertent coupling does not occur. Good bundle definition and usage leads to a more flexible environment.

30.2.1.1.2. Maven Bundle Plugin

Below is a code snippet from a Maven `pom.xml` for creating an OSGi Bundle using the Maven Bundle plugin.

```
...
<packaging>bundle</packaging>
...
<build>
...
<plugin>
  <groupId>org.apache.felix</groupId>
  <artifactId>maven-bundle-plugin</artifactId>
  <configuration>
    <instructions>
      <Bundle-Name>${project.name}</Bundle-Name>
      <Export-Package />
      <Bundle-SymbolicName>${project.groupId}.${project.artifactId}</Bundle-
SymbolicName>
      <Import-Package>
        ddf.catalog,
        ddf.catalog.-
      </Import-Package>
    </instructions>
  </configuration>
</plugin>
...
</build>
...
```

30.2.1.2. Third Party and Utility Bundles

It is recommended to avoid building directly on included third party and utility bundles. These components do provide utility and reuse potential; however, they may be upgraded or even replaced at anytime as bug fixes and new capabilities dictate. For example, web services may be built using CXF. However, the distributions frequently upgrade CXF between releases to take advantage of new features. If building on these components, be aware of the version upgrades with each distribution release.

Instead, component developers should package and deliver their own dependencies to ensure future compatibility. For example, if re-using a bundle, the specific bundle version that you are depending on should be included in your packaged release, and the proper versions should be referenced in your bundle(s).

30.2.1.3. Deploying a Bundle

A bundle is typically installed in one of two ways:

1. Installed as a feature
2. Hot deployed in the `/deploy` directory

The fastest way to deploy a created bundle during development is to copy it to the `/deploy` directory of a running DDF. This directory checks for new bundles and deploys them immediately. According

to Karaf documentation, "Karaf supports hot deployment of OSGi bundles by monitoring JAR files inside the `[home]/deploy` directory. Each time a JAR is copied in this folder, it will be installed inside the runtime. It can be updated or deleted and changes will be handled automatically. In addition, Karaf also supports exploded bundles and custom deployers (Blueprint and Spring DM are included by default)." Once deployed, the bundle should come up in the Active state, if all of the dependencies were properly met. When this occurs, the service is available to be used.

30.2.1.4. Verifying Bundle State

To verify if a bundle is deployed and running, go to the running command console and view the status.

- Execute the `list` command.
- If the name of the bundle is known, the `list` command can be piped to the `grep` command to quickly find the bundle.

The example below shows how to verify if a Client is deployed and running.

Verifying with grep

```
ddf@local>list | grep -i example
[ 162] [Active] [ ] [ ] [ 80] DDF :: Registry :: example Client (2.0.0)
```

The state is `Active`, indicating that the bundle is ready for program execution.

30.2.2. Working with Features Files

Features XML files group other features and/or bundle(s) for ease of installation/uninstallation.

In order to ensure that the installer can correctly interpret and display application details, there are several guidelines that should be followed when creating the features file for the application.

- Ensure only one feature in the `features.xml` has the same name as the project itself. (e.g. If the project name is `spatial-app`, there should be one feature named `spatial-app`) This is the feature that the installer displays to the user (name, description, version, etc.) and provides the description of what the application contains. This feature should only declare the `prerequisite` features of the app and any bundles or features that are required by all of the features within the app.
- Be sure all features that should be installed by default include the `auto-install` flag.

```
install='auto'
```

- Each feature should declare its `prerequisite` features. These are the features that are required to be installed prior to this feature starting up. At a minimum this should include the application feature, so that the application prerequisites and required dependencies are satisfied before this feature starts.

```
<feature prerequisite="true">feature1</feature>
```

30.2.2.1. Auto-starting an Application Feature

Within the `features.xml` file for an application, the feature that is named the same as the project will have the `install` attribute set to `auto`. Within this feature, refer to any prerequisites of the application as well as any features or bundles within the application that are required by all features. Features that should start automatically should have `install` set to `auto`. Other features should have `install` set to `manual`.

The following example demonstrates configuring features to be auto-started. The naming convention for this feature is typically “application name” + “`-app`,” as shown.

Auto-start Features

```
<feature name="example-app" install="auto">
  <feature prerequisite="true">prereq-feature</feature>
  <bundle>mvn:com.example/example.dependency/1.3</bundle>
</feature>
<feature name="auto-start-feature" install="auto">
  <feature prerequisite="true">example-app</feature>
  <bundle>mvn:com.example/example.auto/1.3</bundle>
</feature>
<feature name="manual-start-feature" install="manual">
  <feature prerequisite="true">example-app</feature>
  <bundle>mvn:com.example/example.manual/1.3</bundle>
</feature>
```

30.3. High Availability Guidance

Capabilities that need to function in a Highly Available Cluster should have one of the two below properties.

Stateless

Stateless capabilities will function in an Highly Available Cluster because no synchronization between DDF nodes is necessary.

Common storage

If a capability must store data or share state with another node, then the data or shared state must be accessible to all nodes in the Highly Available Cluster. For example, the Catalog’s storage provider must be accessible to all DDF nodes.

Appendices

Appendix A: Application References

Installation and configuration details by application.

A.1. Admin Application Reference

The Admin Application contains components that are integral for the installation and configuration of DDF applications.

It contains various services and interfaces that allow administrators more control over their systems and enhances administrative capabilities when installing and managing DDF.

The Admin application contains an application service that handles all operations that are performed on applications. This includes adding, removing, starting, stopping, and showing status.

A.1.1. Admin Application Prerequisites

None.

A.1.2. Installing the Admin Application

The Admin application is installed by default with a standard installation.

A.1.3. Configuring the Admin Application

To configure the Admin Application:

1. Navigate to the Admin Console.
2. Select the **Admin** application.
3. Select the **Configuration** tab.

Table 90. Admin Available Configurations

Name	Property	Description
Admin Configuration Policy	org.codice.ddf.admin.config.policy.AdminConfigPolicy	Admin Configuration Policy configurations.
Admin UI	org.codice.admin.ui.configuration	Admin UI configurations.

Table 91. Admin Configuration Policy

Name	Id	Type	Description	Default Value	Required
Feature and App Permissions	featurePolicies	String	When enabled, the desired features or apps will only be modifiable and viewable to users with the set attributes. The entry should be the format of: <code>feature name/app name = "user attribute name=user attribute value"</code>		false
Configuration Permissions	servicePolicies	String	When enabled, the desired service will only be modifiable and viewable to users with the set attributes. The entry should be the format of: <code>configuration ID = "user attribute name=user attribute value"</code>	null	false

Table 92. Admin UI

Name	Id	Type	Description	Default Value	Required
Enable System Usage message	systemUsageEnabled	Boolean	Turns on a system usage message, which is shown when the Admin Application is opened.	false	true
System Usage Message Title	systemUsageTitle	String	A title for the system usage message when the application is opened.		true
System Usage Message	systemUsageMessage	String	A system usage message to be displayed to the user each time the user opens the application.		true
Show System Usage Message once per session	systemUsageOncePerSession	Boolean	With this selected, the system usage message will be shown once for each browser session. Uncheck this to have the usage message appear every time the admin page is opened or refreshed.	true	true
Header	header	String	Specifies the header text to be rendered on the Administrator Console.		true
Footer	footer	String	Specifies the footer text to be rendered on the Administrator Console.		true
Style	style	String	Specifies the style of the Header and Footer.	green	true
Text Color	textColor	String	Specifies the text color of the Header and Footer.	banner-text-white	true

Name	Id	Type	Description	Default Value	Required
Ignored Installer Applications	disabledInStallerApps	String	Comma delimited list (appName, appName2, ...appNameN) of applications that will be disabled in the installer.	admin-app,platform-app	null

A.2. Message Broker Application Reference

The Message Broker application gives an administrator the ability to configure and control the behavior of the Message Broker. These configurations will include aspects like the graceful shutdown period of components, names of queues and topics, and routing of messages.

A.2.1. Prerequisites for Message Broker Application

None.

A.2.2. Installing Message Broker Application

The Message Broker app is uninstalled by default and requires manual installation.

To install:

1. Navigate to the **Admin Console**.
2. Select **Manage Applications**.
3. Select the **Message Broker** application.

A.2.3. Configuring the Message Broker Application

The standard installation of the Message Broker application has no configurable properties.

A.2.3.1. Configuring the Message Broker for a Highly Available Cluster

Prior to making these configuration changes, follow the instructions in [Installing DDF](#) to install DDF on two physically-separate hosts.

1. Configure each of the DDF installations to point to each other in a live/backup server configuration. One server will have an additional step to be designated as the backup.
2. Modify system.properties:
The `<DDF_HOME>/etc/system.properties` in each of the installations needs to be updated so that the servers know about each other. The following properties need to have the values on the right side of the `=` updated.

```
artemis.live.host=<Hostname.or.ip.here>
artemis.backup.host=<Hostname.or.ip.here>
artemis.network.iplist=<Comma, separated, IPs>
artemis.cluster.password=<Common password across all nodes>
```

IMPORTANT

Using a Non-Local IP or Host

`artemis.network.iplist` should contain a list of non-local IPs or host names that are not hosted on the same physical machine as either the live or backup machines. These IP addresses are pinged in the event of a network outage. If the backup cannot reach the live server but can successfully ping one of these hosts it will then take over as the live server. If the host list is incorrectly configured with a local IP it could break the cluster by causing both servers to go live. It is also recommended that the live server have the backups server's IP in its list and the backup server have the live server's IP in its list.

3. Configure a Backup Broker:

The installation that is going to be used as the backup needs to have an additional configuration change made so that it knows it's the backup. The `<DDF_HOME>/etc/org.apache.activemq.artemis.cfg` should be modified to point to the provided `artemis-backup.xml` instead of `artemis.xml`. Once updated, the `config` value should look like this:

```
config=file:etc/artemis-backup.xml
```

4. Restart Servers:

If the DDF instances are currently running, stop and restart the backup and then the live server. Making sure the backup starts first ensures that the live server doesn't have any issues establishing the backup due to the backup being busy initializing. See [Starting DDF](#) for detailed steps.

5. Verify Cluster Replication:

Once both servers are started, the following command can be run using curl or a [browser](#) to verify that the servers have successfully synced.

Server Cluster Verification Command

```
sh
curl
\https://{FQDN}:{PORT}/admin/jolokia/read/org.apache.activemq.artemis:broker=%22<FQ
DN>%22/ReplicaSync --user admin:admin --insecure
```

Example ReplicaSync JSON Response

```
{  
  "request": {  
    "mbean": "org.apache.activemq.artemis:broker=localhost",  
    "attribute": "ReplicaSync",  
    "type": "read"  
  },  
  "value": true,  
  "timestamp": 1485967446,  
  "status": 200  
}
```

IMPORTANT

If LDAP has been configured then the admin user and password for the above command will need to be changed.

IMPORTANT

Note the `"value":true` field: if it is false, then the replication is still in progress or the logs should be consulted to see if there was an issue establishing a connection between the live and backup servers.

Additionally, for more details about the health of the cluster, the following command can be run using

curl or a

`https://{FQDN}:{PORT}/admin/jolokia/read/org.apache.activemq.artemis:broker=%22localhost%22,component=cluster-connections,name=%22my-cluster%22/Topologybrowser.`

Server Health Status Command

```
sh  
curl  
\https://{FQDN}:{PORT}/admin/jolokia/read/org.apache.activemq.artemis:broker=%22<FQDN>%22,component=cluster-connections,name=%22my-cluster%22/Topology --user admin:admin  
--insecure
```

This endpoint returns diagnostic info about the cluster that can be used for troubleshooting. Values of interest in the response are the `node=2` value which is a count of the nodes in the cluster and the port/host values for each node.

Example Topology JSON Response for a Cluster of 2

```
{  
  "request": {  
    "mbean": "org.apache.activemq.artemis:broker=localhost", component=cluster-  
connections, name=my-cluster,  
    "attribute": "Topology",  
    "type": "read"  
  },  
  "value": "topology on Topology@750c2a56[owner=ClusterConnectionImpl@228651110[nod  
eUUID=17b48db9-e7ee-11e6-9d56-38c986025a6f, connector=TransportConfiguration(name=net  
ty-connector, factory=org.apache.activemq.artemis.core.remoting.impl.netty.NettyConne  
ctorFactory) ?port=5672&host=10-101-3-185, address=jms, server=ActiveMQServerImpl::se  
rverUUID=17b48db9-e7ee-11e6-9d56-38c986025a6f]]:\n\t17b48db9-e7ee-11e6-9d56-38c986025  
a6f => TopologyMember[id = 17b48db9-e7ee-11e6-9d56-38c986025a6f, connector=Pair[a=Tra  
nsportConfiguration(name=netty-connector, factory=org.apache.activemq.artemis.core.re  
moting.impl.netty.NettyConnectorFactory) ?port=5672&host=10-101-3-185, b=TransportCon  
figuration(name=netty-connector, factory=org.apache.activemq.artemis.core.remoting.im  
pl.netty.NettyConnectorFactory) ?port=5672&host=10-101-2-97], backupGroupName=null, s  
caleDownGroupName=null]\n\ttnodes=2\tmembers=1",  
    "timestamp": 1485971158,  
    "status": 200  
  }  
}
```

A.2.3.2. Securing the Message Broker Application

DDF can be configured to use Artemis to perform authentication and authorization against an LDAP server.

Artemis provides the ability to apply role-based security to queues based on addresses (see [the Artemis documentation](#) for details). It can be configured to use an LDAP server to perform authentication and authorization for users who connect to it.

IMPORTANT

If you are setting up multiple DDF instances in a cluster for high availability, then you will need to perform these steps on each instance.

The Security STS LDAP Login and Security STS LDAP Claims Handler bundles are responsible for authenticating and authorizing users with your LDAP server. To configure them for your LDAP server, follow the instructions in [Security STS LDAP Login](#) and [Security STS LDAP Claims Handler](#).

Once the STS LDAP Login and Claims Handlers are configured, update `<DDF_HOME>/etc/org.apache.activemq.artemis.cfg` to use the `ldap` realm (just change `domain=karaf` to `domain=ldap`):

```
<DDF_HOME>/etc/org.apache.activemq.artemis.cfg
```

```
domain=ldap
```

DDF uses two roles in the security settings for Artemis: `manager` and `broker-client`.

```

<security-setting match="#">
  <permission type="createNonDurableQueue" roles="manager,broker-client"/>
  <permission type="deleteNonDurableQueue" roles="manager,broker-client"/>
  <permission type="createDurableQueue" roles="manager"/>
  <permission type="deleteDurableQueue" roles="manager"/>
  <permission type="consume" roles="manager,broker-client"/>
  <permission type="browse" roles="manager,broker-client"/>
  <permission type="send" roles="manager,broker-client"/>
  <permission type="manage" roles="manager"/>
</security-setting>

```

Users with the role `manager` have full permissions, but users with the role `broker-client` cannot create or delete durable queues or invoke management operations.

Your LDAP should have groups that correspond to these roles so that members of those groups will have the correct permissions when connecting to Artemis to send or consume messages. Alternatively, you can choose roles other than `manager` and `broker-client`, which may be useful if your LDAP already has groups that you would like to use as Artemis roles. If you wish to use different roles, just replace `manager` and/or `broker-client` in the `<security-setting>` in `artemis.xml` with the roles you would like to use.

A.2.3.3. Artemis Broker Connection Configuration

The [Artemis Broker Connection Configuration](#) manages the parameters for DDF's connection to Artemis. The username and password in the [Artemis Broker Connection Configuration](#) need to be updated so that they correspond to a user in your LDAP. If possible, this user should have the `manager` role (or the role that is being used in place of `manager` if the default Artemis role has been changed).

To update the username and password:

1. Navigate to the [Admin Console](#)
2. Select the [Broker App](#) application.
3. Select the [Configuration](#) tab.
4. Select the [Artemis Broker Connection Configuration](#).
5. Enter the username and password and select [Save changes](#).

A.2.4. Using the Message Broker Application

The Message Broker app can be used through the Admin Console. See [the Route Manager](#) and [the Undelivered Messages UI](#) for more information.

A.2.4.1. Undelivered Messages UI

The Undeliverable Messages tab gives an administrator the ability to view undeliverable messages and then decide whether to resend or delete those messages.

The Undelivered Messages UI is installed as a part of the Message Broker.

To view undelivered messages, an administrator can use the "retrieve" button, which makes an immediate call to the backend and displays all the messages. Alternatively, the "start polling" button makes calls to the backend every 5 seconds and updates the display accordingly.

An administrator can select messages by clicking anywhere in the row of the message. Multiple messages can be selected simply by clicking multiple messages or by clicking the "Select all" option at the head of the table. Deselecting is done by clicking a message again or clicking the "Deselect all" option, next to the "Select all" option.

To attempt to resend messages, select the messages, and then click the "resend" button. Currently, there is no way to identify if a message was successfully redelivered.

To delete messages, select the messages, and then click the "delete" button.

NOTE

Only 200 messages can be viewed at a time, even though there may be more than 200 undelivered messages

Known issues with the Undelivered Messages UI:

- If attempting to resend a message, but the listener is no longer available, the message will be "successfully" resent and removed from the UI and the Artemis DLQ but will not be successfully redelivered.

A.2.4.2. Route Manager

The Route Manager gives an administrator the ability to configure and deploy Camel routes, queues, and topics dynamically. The `sjms` component is available by default. If a need arises for a new route, an administrator can easily develop a new route and deploy it to satisfy the requirement, rather than spending the time to develop, compile, and test new code.

The Route Manager is installed as a part of the Message Broker application.

The route shutdown timeout can be configured.

To deploy a new route, simply place a route `.XML` file in the `<DDF_HOME>/etc/routes` directory of DDF. To remove a route (or set of routes), delete the `.XML` file.

There are example routes in the `<DDF_HOME>/etc/routes` directory by default.

A.3. Catalog Application Reference

The Catalog provides a framework for storing, searching, processing, and transforming information.

Clients typically perform create, read, update, and delete (CRUD) operations against the Catalog.

At the core of the Catalog functionality is the Catalog Framework, which routes all requests and responses through the system, invoking additional processing per the system configuration.

A.3.1. Catalog Application Prerequisites

To use the Catalog Application, the following applications/features must be installed:

- Platform

A.3.2. Installing the Catalog Application

The Catalog is pre-installed with a standard installation.

A.3.3. Configuring the Catalog Application

To configure the Catalog Application:

1. Navigate to the Admin Console.
2. Select the **Catalog** application.
3. Select the **Configuration** tab.

Table 93. Catalog Available Configurations

Name	Property	Description
Catalog Federation Strategy	ddf.catalog.federation.impl.CachingFederationStrategy	Catalog Federation Strategy.
Catalog Backup Plugin	ddf.catalog.backup.CatalogBackupPlugin	Catalog Backup Plugin configurations.
Catalog Standard Framework	ddf.catalog.CatalogFrameworkImpl	Catalog Standard Framework configurations.
Confluence Federated Source	Confluence_Federated_Source	Confluence Federated Source.
Content Directory Monitor	org.codice.ddf.catalog.content.monitor.ContentDirectoryMonitor	Content Directory Monitor configurations.
Content File System Storage Provider	org.codice.ddf.catalog.content.impl.FileSystemStorageProvider	Content File System Storage Provider.
CSW Connected Source	Csw_Connected_Source	CSW Connected Source.
Expiration Date Pre-Ingest Plugin	org.codice.ddf.catalog.plugin.expiration.ExpirationDatePlugin	Catalog pre-ingest plugin to set an expiration date on metacards.
FTP Endpoint	ddf.catalog.ftp.FtpServerStarter	FTP Endpoint configurations.
Historian	ddf.catalog.history.Historian	Enables versioning of both metacards and content.
Metocard Attribute Security Policy Plugin	org.codice.ddf.catalog.security.policy.metocard.MetocardAttributeSecurityPolicyPlugin	Metocard Attribute Security Policy Plugin.
Catalog Metocard Ingest Network Plugin	org.codice.ddf.catalog.plugin.metocard.MetocardIngestNetworkPlugin	Catalog Metocard Ingest Network Plugin.

Name	Property	Description
Metocard Validation Filter Plugin	ddf.catalog.metocard.validation.MetocardValidityFilterPlugin	Metocard Validation Filter Plugin.
Metocard Backup File Storage Provider	Metocard_File_Storage_Route	Enable data backup of metacards using a configurable transformer.
Resource Download Settings	Metocard_S3_Storage_Route	Resource Download Configuration.
Catalog OpenSearch Federated Source	OpenSearchSource	Catalog OpenSearch Federated Source.
Resource Download Settings	ddf.catalog.resource.download.ReliableResourceDownloadManager	Resource Download configurations.
Schematron Validation Services	ddf.services.schematron.SchematronValidationService	Schematron Validation Services configurations.
Security Audit Plugin	Conorg.codice.ddf.catalog.plugin.security.audit.SecurityAuditPlugininfluence_Federated_Source	Security Audit Plugin.
Tika Input Transformer	ddf.catalog.transformer.input.tika.TikaInputTransformer	Tika Input Transformer.
URL Resource Reader	ddf.catalog.resource.impl.URLResourceReader	URL Resource Reader
Video Thumbnail Plugin	org.codice.ddf.catalog.content.plugin.video.VideoThumbnailPlugin	Video Thumbnail Plugin.
XML Attribute Security Policy Plugin	org.codice.ddf.catalog.security.policy.xml.XmlAttributeSecurityPolicyPlugin	XML Attribute Security Policy Plugin.
Xml Query Transformer	ddf.catalog.transformer.xml.XmlResponseQueueTransformer	Xml Response Query Transformer.
Catalog Policy Plugin	org.codice.ddf.catalog.security.CatalogPolicy	Catalog Policy Plugin.
Resource URI Policy Plugin	org.codice.ddf.catalog.security.ResourceUriPolicy	Resource URI Policy Plugin.
Source Poller	ddf.catalog.util.impl.SourcePoller	Source Poller configurations.

Table 94. Catalog Federation Strategy

Name	<code>Id</code>	Type	Description	Default Value	Required
Maximum start index	<code>maxStartIndex</code>	Integer	Sets a limit on the number of results this sorted federation strategy can handle from each federated source. A large start index in conjunction with several federated sources could yield a large result set, which the sorted federation strategy has a limited ability to do. The admin can make a rough calculation to decide what maximum start index to use based on the amount of memory in the system, the amount of federated sources, the number of threads, and the expected amount of query results requested ((average # of threads) * (maximum # of federated sources) * (maxStartIndex + maximumQueryResults)) must fit into the allocated memory of the running distribution. This field will be removed when sorted federation strategy has the ability to sort a larger amount of results.	50000	true
Expiration Interval	<code>expirationIntervalInMinutes</code>	Long	Interval that Solr Cache checks for expired documents to remove.	10	true
Expiration Age	<code>expirationAgeInMinutes</code>	Long	The number of minutes a document will remain in the cache before it will expire. Default is 7 days.	10080	true
Cache Everything	<code>cachingEverything</code>	Boolean	Cache all results unless configured as native	false	true
Cache Remote Ingests	<code>cacheRemoteIngests</code>	Boolean	Cache remote ingest results	false	true
Show Validation Errors	<code>showErrors</code>	Boolean	Show metacards with validation errors in search results	false	true
Show Validation Warnings	<code>showWarnings</code>	Boolean	Show metacards with validation warnings in search results	true	true

Table 95. Catalog Backup Plugin

Name	Id	Type	Description	Default Value	Required
Root backup directory path	rootBackupDir	String	Root backup directory for Metacards. A relative path is relative to ddf.home.	data/backup	true
Subdirectory levels	subDirLevels	Integer	Number of subdirectory levels to create. Two characters from the ID will be used to name each subdirectory level.	2	true

Table 96. Catalog Standard Framework

Name	Id	Type	Description	Default Value	Required
Enable Fanout Proxy	fanoutEnabled	Boolean	When enabled the Framework acts as a proxy, federating requests to all available sources. All requests are executed as federated queries and resource retrievals, allowing the framework to be the sole component exposing the functionality of all of its Federated Sources.	false	true
Enable Notifications	notificationEnabled	Boolean	Check to enable notifications.	true	false
Fanout tag blacklist	fanoutTagBlacklist	String	Ingest operations with tags in this list will be rejected.		true

Table 97. Confluence Federated Source

Name	Property	Type	Description	Default Value	Required
Source Name	shortname	String			Yes
Confluence Rest URL	endpointUrl	String	The Confluence Rest API endpoint URL. Example: https://{{FQDN}}:{{PORT}}/rest/api/content		Yes
Username	username	String	Username to use with HTTP Basic Authentication. This auth info will overwrite any federated auth info. Only set this if the Confluence endpoint requires basic authentication.		No

Name	Property	Type	Description	Default Value	Required
Password	password	Password	Password to use with HTTP Basic Authentication. This auth info will overwrite any federated auth info. Only set this if the Confluence endpoint requires basic authentication.		No
Include Page Contents In Results	includePageContent	Boolean	Flag indicating if Confluence page contents should be included in the returned results.	false	No
Include Archived Spaces	includeArchivedSpaces	Boolean	Flag indicating if archived confluence spaces should be included in search results.	false	No
Exclude Confluence Spaces	excludeSpaces	Boolean	Flag indicating if the list of Confluence Spaces should be excluded from searches instead of included.	false	No
Confluence Spaces	confluenceSpaces	String cardinality =1000	The confluence spaces to include/exclude from searches. If no spaces are specified, all visible spaces will be searched.		No
Additional Attributes	additionalAttributes	String cardinality =100	Additional attributes to add to confluence metacards returned from this source.		No
Availability Poll Interval	availabilityPollInterval	Long	Availability polling interval in milliseconds.	60000	No

Table 98. Catalog Content Directory Monitor

Name	Id	Type	Description	Default Value	Required
Directory Path	monitoredDirectoryPath	String	"Specifies the directory to be monitored, can be a filesystem path or webdav address (only supported for Monitor in place)"	false	true
Maximum Concurrent Files	numThreads	Integer	Specifies the maximum number of concurrent files to be processed within a directory (maximum of 8). If this number exceeds 8, 8 will be used in order to preserve system resources. Make sure that your system has enough memory to support the number of concurrent processing threads across all directory monitors.	1	true

Name	Id	Type	Description	Default Value	Required
ReadLock Time Interval	readLockIntervalMilliseconds	Integer	Specifies the time to wait (in milliseconds) before acquiring a lock on a file in the monitored directory. This interval is used for sleeping between attempts to acquire the read lock on a file to be ingested. The default value of 100 milliseconds is recommended.	100	true
Processing Mechanism	processingMechanism	String	Choose what happens to the content item after it is ingested. Delete will remove the original file after storing it in the content store. Move will store the item in the content store, and a copy under ./ingested, then remove the original file. (NOTE: this will double the amount of disk space used.) Monitor in place will index the file and serve it from its original location.	in_place	false
Attribute Overrides	attributeOverrides	String	Optional: Metocard attribute overrides (Key-Value pairs) that can be set on the content monitor. If an attribute is specified here, it will overwrite the metocard's attribute that was created from the content directory. The format should be 'key=value'. The maximum allowed size of an attribute override is 65,535 bytes.	null	false

Table 99. Content File System Storage Provider

Name	Id	Type	Description	Default Value	Required
Content Repository File Path	baseContentDirectory	String	Specifies the directory to use for the content repository. A shutdown of the server is necessary for this property to take effect. If a filepath is provided with directories that don't exist, File System Provider will attempt to create them.	<DDF_HOME>/data/content/store	true

Table 100. CSW Connected Source

Name	Id	Type	Description	Default Value	Required
Source ID	id	String	The unique name of the Source.	CSW	true

Name	<code>Id</code>	Type	Description	Default Value	Required
CSW URL	<code>cswUrl</code>	String	URL to the endpoint implementing the Catalogue Service for Web (CSW) spec.	null	true
Event Service Address	<code>eventServiceAddress</code>	String	DDF Event Service endpoint. Do NOT include .wsdl or ?wsdl.	null	false
Register for Events	<code>registerForEvents</code>	Boolean	Check to register for events from this connected source.	false	false
Username	<code>username</code>	String	Username for CSW Service (optional).	null	false
Password	<code>password</code>	String	Password for CSW Service (optional).	null	false
Disable CN Check	<code>disableCnCheck</code>	Boolean	Disable CN check for the server certificate. This should only be used when testing.	false	true
Force Longitude/ Latitude coordinate order	<code>isLonLatOrder</code>	Boolean	Force Longitude/Latitude coordinate order.	false	true
Use posList in LinearRing	<code>usePosList</code>	Boolean	Use a <posList> element rather than a series of <pos> elements when issuing geospatial queries containing a LinearRing.	false	false
Metocard Mappings	<code>metocardMappings</code>	String	Mapping of the Metocard Attribute names to their CSW property names. The format should be ' <code>title=dc:title</code> '.	effective=created, created=dateSubmitted, modified=modified, thumbnail=reference, content-type=type, id=identifier, resource-uri=source	false
Poll Interval	<code>pollInterval</code>	Integer	Poll Interval to Check if the Source is available (in minutes - minimum 1).	5	true
Connection Timeout	<code>connectionTimeout</code>	Integer	Amount of time to attempt to establish a connection before timing out, in milliseconds.	30000	true

Name	Id	Type	Description	Default Value	Required
Receive Timeout	<code>receiveTimeout</code>	Integer	Amount of time to wait for a response before timing out, in milliseconds.	60000	true
Output Schema	<code>outputSchema</code>	String	Output Schema	http://www.opengis.net/cat/csw/2.0.2	true
Query Type Name	<code>queryTypeName</code>	String	Qualified Name for the Query Type used in the CSW GetRecords request.	csw:Record	true
Query Type Namespace	<code>queryTypeNamespace</code>	String	Namespace prefix for the Query Type used in the CSW GetRecords request.	http://www.opengis.net/cat/csw/2.0.2	true
Force CQL Text as the Query Language	<code>isCqlForce</code>	Boolean	Force CQL Text.	false	true
Forced Spatial Filter Type	<code>forceSpatialFilter</code>	String	Force only the selected Spatial Filter Type as the only available Spatial Filter.	NO_FILTER	false

Table 101. Expiration Date Pre-Ingest Plugin

Name	Id	Type	Description	Default Value	Required
Overwrite If Empty	<code>overwriteIfBlank</code>	Boolean	If this is checked, overwrite all blank expiration dates in metacards. If this is not checked, leave metacards with blank expiration dates as-is.	false	true
Overwrite If Exists	<code>overwriteIfExists</code>	Boolean	If this is checked, overwrite all existing non-empty expiration dates in metacards with a new date. If this is not checked, leave metacards with an existing expiration date.	false	true
Offset from Created Date (in days)	<code>offsetFromCreatedDate</code>	Integer	A metacard's new expiration date is calculated by adding this value (in days) to its created date.	30	true

Table 102. FTP Endpoint

Name	Id	Type	Description	Default Value	Required
FTP Port Number	<code>port</code>	Integer	The port number for the FTP server to listen on.	8021	true

Name	Id	Type	Description	Default Value	Required
Client Authentication	clientAuth	String	Whether or not client authentication is required or wanted. A value of "Need" requires client auth, a value of "Want" leaves it up to the client.	want	true

Table 103. Historian

Name	Id	Type	Description	Default Value	Required
Enable Versioning	historyEnabled	Boolean	Enables versioning of both metacards and content.	true	true

Table 104. Metocard Attribute Security Policy Plugin

Name	Id	Type	Description	Default Value	Required
Metocard Intersect Attributes:	intersectMetocardAttributes	List of rules	Each line item in the configuration is a rule. The format of a rule is the name of a single source attribute, followed by an equals sign, followed by the destination attribute. For example: source_attribute1=destination_attribute. The plugin gathers the source attributes that have a common destination. It takes the combined values of the source attributes and makes them the values of a (new) metocard attribute, the destination attribute. The strategy for combining the values is intersection, which means only the values common to all source attribute are added to the destination attribute. Note: Do not use the same destination attributes in both the Intersect and Union rule sets. The plugin will behave unpredictably.	none	false

Name	Id	Type	Description	Default Value	Required
Metocard Union Attributes:	unionMetocardAttributes	List of rules	<p>Each line item in the configuration is a rule. The format of a rule is the name of a single source attribute, followed by an equals sign, followed by the destination attribute. For example:</p> <p>source_attribute1=destination_attribute. The plugin gathers the source attributes that have a common destination. It takes the combined values of the source attributes and makes them the values of a (new) metocard attribute, the destination attribute. The strategy for combining the values is union, which means only all the values of the source attribute are added to the destination attribute (excluding duplicates) Note: Do not use the same destination attributes in both the Intersect and Union rule sets. The plugin will behave unpredictably.</p>	none	false

Table 105. Catalog Metocard Ingest Network Plugin

Name	Id	Type	Description	Default Value	Required	Criteria
criteriaKey	String	Specifies the criteria for the test of equality; which value will be tested? IP Address? Hostname?	remoteAddr	true	Expected Value	expected Value

Name	Id	Type	Description	Default Value	Required	Criteria
String	The value that the criteria must equate to for the attribute overrides to occur.		true	New Attributes	newAttributes	String"

Table 106. Metocard Validation Filter Plugin

Name	Id	Type	Description	Default Value	Required
Attribute map	attributeMap	String	Mapping of Metocard SECURITY attribute to user attribute.	invalid-state=system-admin	false
Filter errors	filterErrors	Boolean	Sets whether metacards with validation errors are filtered.	true	false
Filter warnings	filterWarnings	Boolean	Sets whether metacards with validation warnings are filtered.	false	false

Table 107. Metocard Backup File Storage Provider

Name	Id	Type	Description	Default Value	Required
Keep Deleted Metocard	keepDeletedMetacards	Boolean	Should backups for deleted metacards be kept or removed.	false	true
Metocard Transformer Id	metocardTransformerId	String	ID of the metocard transformer to use to serialize metocard for backup.	metocard	true
Backup Invalid Metacards	keepDeletedMetacards	Boolean	Keep backups for metacards that fail validation with warnings or errors.	true	true
Metocard Backup Output Provider(s)	metocardOutputProviderIds	Comma delimited list of metocard output provider IDs.	Metocard Backup Provider IDs to use for this backup plugin.	fileStorage Provider	true

Table 108. Metocard Backup S3 Storage Provider

Name	Id	Type	Description	Default Value	Required
Keep Deleted Metocard	keepDeleteMetacards	Boolean	Should backups for deleted metacards be kept or removed.	false	true
Metocard Transformer Id	metocardTransformerId	String	ID of the metocard transformer to use to serialize metocard for backup.	metocard	true
Backup Invalid Metacards	keepDeleteMetacards	Boolean	Keep backups for metacards that fail validation with warnings or errors.	true	true
Metocard Tags	backupMetacardTags	String	Backup only metacards with one of the tags specified.	resource	true
S3 Access Key	s3AccessKey	String	The access key to use for S3. Leave blank if on an EC2 host with roles assigned.	""	true
S3 Secret Key	s3SecretKey	Password	The secret key to use for S3. Leave blank if on an EC2 host with roles assigned.		true
S3 Bucket	s3Bucket	String	The S3 Bucket in which to store the backed up metocard data.		true
S3 Endpoint	s3Endpoint	String	The endpoint for the region in which the bucket is located.		true
Object Template	objectTemplate	String	Template specifying the S3 object key for the metocard data. The template uses handlebars syntax. Use [] to reference dotted attributes e.g. {{[attribute.name]}}. If you wish to include date, you would use {{dateFormat created yyyy-MM-dd}}	data/backups/metocard/{{substr string id 0 3}}/{{substr string id 3 6}}/Metocard_S3_Storage_Route.xml	true

Table 109. Catalog OpenSearch Federated Source

Name	Id	Type	Description	Default Value	Required
Source Name	shortname	String	null	DDF-OS	true

Name	Id	Type	Description	Default Value	Required
OpenSearch service URL	<code>endpointUrl</code>	String	The OpenSearch endpoint URL or DDF's OpenSearch endpoint (https://{FQDN}:{PORT}/services/catalog/query)	<code> \${org.codice.ddf.system.protocol} \${org.codice.ddf.system.hostname}: \${org.codice.ddf.system.port} \${org.codice.ddf.system.rootContext}/catalog/query</code>	true
Username	<code>username</code>	String	Username to use with HTTP Basic Authentication. This auth info will overwrite any federated auth info. Only set this if the OpenSearch endpoint requires basic authentication.		false
Password	<code>password</code>	Password	Password to use with HTTP Basic Authentication. This auth info will overwrite any federated auth info. Only set this if the OpenSearch endpoint requires basic authentication.		false
OpenSearch query parameters	<code>parameters</code>	String	Query parameters to use with the OpenSearch connection.	<code>q,src,mr,start,count,mt,dn,lat,lon, radius,bbx,polygon,dtstart,dte nd,dateName,filter,sort</code>	true
Always perform local query	<code>localQueryOnly</code>	Boolean	When federating with other DDFs, keep this checked. If checked, this source performs a local query on the remote site (by setting <code>src=local</code> in endpoint URL), as opposed to an enterprise search.	true	true
Convert to BBox	<code>shouldConvertToBBox</code>	Boolean	Converts Polygon and Point-Radius searches to a Bounding Box for compatibility with older interfaces. Generated bounding box is a very rough representation of the input geometry.	true	true

Table 110. Resource Download Settings

Name	Property	Type	Description	Default Value	Required
Product Cache Directory	productCacheDirectory	String	Directory where retrieved products will be cached for faster, future retrieval. If a directory path is specified with directories that do not exist, Product Download feature will attempt to create those directories. Without configuration, the product cache directory is <DDF_HOME>/data/product-cache. If a relative path is provided it will be relative to the <DDF_HOME>. It is recommended to enter an absolute directory path such as /opt/product-cache in Linux or C:\product-cache in Windows.		false
Enable Product Caching	cacheEnabled	Boolean	Check to enable caching of retrieved products.	true	false
Delay (in seconds) between product retrieval retry attempts	delayBetweenRetryAttempts	Integer	The time to wait (in seconds) between attempting to retry retrieving a product.	10	false
Max product retrieval retry attempts	maxRetryAttempts	Integer	The maximum number of attempts to retry retrieving a product.	3	false
Product Retrieval Monitor Period	retrievalMonitorPeriod	Integer	How many seconds to wait and not receive product data before retrying to retrieve a product.	5	false
Always Cache Product	cacheWhenCanceled	Boolean	Check to enable caching of retrieved products even if client cancels the download. Note: this has no effect if product caching is disabled.	false	false

Table 111. Schematron Validation Services

Name	Id	Type	Description	Default Value	Required
Ruleset Name	id	String	Give this ruleset a name	null	true

Name	Id	Type	Description	Default Value	Required
Root Namespace	<code>namespace</code>	String	The root namespace of the XML	null	true
Schematron Files	<code>schematronFileNames</code>	String	Schematron files (*.sch) to be validated against	null	true

Table 112. Security Audit Plugin

Name	Id	Type	Description	Default Value	Required
Security attributes to audit	<code>auditAttributes</code>	String	List of security attributes to audit when modified	security.access-groups,security.access-individuals	true

Table 113. Tika Input Transformer

Name	Id	Type	Description	Default Value	Required
Use Resource Title	<code>useResourceTitleAsTitle</code>	Boolean	Use the resource's metadata to determine the metocard title. If this is not enabled, the metocard title will be the file name.	false	true

Table 114. URL Resource Reader

Name	Property	Type	Description	Default Value
Follow Server Redirects	<code>followRedirects</code>	Boolean	Check the box if you want the Resource Reader to automatically follow server issued redirects (HTTP Response Code 300 series).	false
Root Resource Directories	<code>rootResourceDirectories</code>	String	List of root resource directories. A relative path is relative to ddf.home. Specifies the only directories the URLResourceReader has access to when attempting to download resources linked using file-based URLs.	data/products

Table 115. Video Thumbnail Plugin

Name	Property	Type	Description	Default Value	Required
Maximum video file size to process (Megabytes)	maxFileSizeMB	Long	Maximum video file size in Megabytes for which to create a thumbnail. Default is 120 Megabytes. Processing large videos may affect system performance.	120	false

Table 116. XML Attribute Security Policy Plugin

Name	Id	Type	Description	Default Value	Required
XML Elements:	xmlElements	String	XML elements within the metadata that will be searched for security attributes. If these elements contain matching attributes, the values of the attributes will be combined.		true
Security Attributes (union):	securityAttributeUnions	String	Security Attributes. These attributes, if they exist on any of the XML elements listed above, will have their values extracted and the union of all of the values will be saved to the metocard. For example: if element1 and element2 both contain the attribute 'attr' and that attribute has values X,Y and X,Z, respectively, then the final result will be the union of those values: X,Y,Z. The X,Y,Z value will be the value that is placed within the security attribute on the metocard.		false
Security Attributes (intersection):	securityAttributeIntersections	String	Security Attributes. These attributes, if they exist on any of the XML elements listed above, will have their values extracted and the intersection of all of the values will be saved to the metocard. For example: if element1 and element2 both contain the attribute 'attr' and that attribute has values X,Y and X,Z, respectively, then the final result will be the intersection of those values: X. The X value will be the value that is placed within the security attribute on the metocard.	null	false

Table 117. Xml Query Transformer

Name	Id	Type	Description	Default Value	Required
Parallel Marshalling Threshold	<code>threshold</code>	Integer	Response size threshold above which marshalling is run in parallel	50	true

Table 118. Catalog Policy Plugin

Name	Id	Type	Description	Default Value	Required
Create Required Attributes	<code>createPermissions</code>	String	Roles/attributes required for the create operations. Example: role=role1,role2	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest	true
Update Required Attributes	<code>updatePermissions</code>	String	Roles/attributes required for the update operation. Example: role=role1,role2	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest	true
Delete Required Attributes	<code>deletePermissions</code>	String	Roles/attributes required for the delete operation. Example: role=role1,role2	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest	true
Read Required Attributes	<code>readPermissions</code>	String	Roles/attributes required for the read operations (query and resource). Example: role=role1,role2	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest	true

Table 119. Resource URI Policy Plugin

Name	Id	Type	Description	Default Value	Required
Permit Resource URI on Creation	<code>createPermissions</code>	String	Allow users to provide a resource URI when creating a metocard	<code>http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest</code>	true
Permit Resource URI on Update	<code>updatePermissions</code>	String	Allow users to provide a resource URI when updating a metocard	<code>http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest</code>	true

Table 120. Source Poller

Name	Id	Type	Description	Default Value	Required
Interval	<code>interval</code>	Integer	The interval (in minutes) at which to execute successive source polling. Must be a positive, nonzero integer.	1	true

A.4. GeoWebCache Application Reference

GeoWebCache enables a server providing a map tile cache and tile service aggregation.

WARNING

The GeoWebCache application is currently in an EXPERIMENTAL status and should not be installed on a security-hardened installation.

GeoWebCache enables a server providing a tile cache and tile service aggregation. See ([GeoWebCache](#)) for more information. This application also provides an administrative plugin for the management of GeoWebCached layers. GeoWebCache also provides a user interface for previewing, truncating, or seeding layers at <https://{FQDN}:{PORT}/geowebcache/>.

A.4.1. GeoWebCache Application Prerequisites

None.

A.4.2. Installing GeoWebCache

The GeoWebCache application is **not** installed by default.

To install:

1. Navigate to the Admin Console.
2. Select **Manage**.

3. Select the GeoWebCache application **Start** button. The application will move to **Active Applications** on startup.
4. Select **Back** in Admin Console.

A.4.3. Configuring GeoWebCache

GeoWebCache can be configured to cache layers locally, using the following procedures.

A.4.3.1. Adding GeoWebCache Layers

Add layers to the local cache:

1. Navigate to the Admin Console.
2. Select the GeoWebCache Application.
3. Select the **GeoWebCache Layers** tab.
4. Click the **Add** button.
5. Enter the data in the fields provided.
6. If necessary, click the **Add** button to add additional MIME types.
7. If necessary, click the **Add** button to add additional WMS Layer Names.

Table 121. Add Layer

Name	Property	Type	Description	Default Value	Required
Name		String			no
Mime Formats		String			yes
URL		URI			yes
WMS Layer Name		String			yes

A.4.3.2. Editing GeoWebCache Layers

1. Navigate to the Admin Console.
2. Select the GeoWebCache application.
3. Navigate to the **GeoWebCache Layers** tab.
4. Click the **Name** field of the layer to edit.

A.4.3.3. Removing GeoWebCache Layers

1. Click the **Delete** icon at the end of the row of the layer to be deleted.

A.4.3.4. Configuring GWC Disk Quota

Storage usage for a GeoWebCache server is managed by a `diskquota.xml` file with configuration

details to prevent image-intensive data from filling the available storage.

To view the disk quota XML representative: [https://\[{FQDN}\]:{PORT}/geowebcache/rest/diskquota.xml](https://[{FQDN}]:{PORT}/geowebcache/rest/diskquota.xml)

To update the disk quota, a client can post a new XML configuration: `curl -v -k -XPUT -H "Content-type: text/xml" -d @diskquota.xml "\https://[{FQDN}]:{PORT}/geowebcache/rest/diskquota.xml"`

Example `diskquota.xml`

```
<gwcQuotaConfiguration>
  <enabled>true</enabled>
  <diskBlockSize>2048</diskBlockSize>
  <cacheCleanUpFrequency>5</cacheCleanUpFrequency>
  <cacheCleanUpUnits>SECONDS</cacheCleanUpUnits>
  <maxConcurrentCleanUps>5</maxConcurrentCleanUps>
  <globalExpirationPolicyName>LFU</globalExpirationPolicyName>
  <globalQuota>
    <value>100</value>
    <units>GiB</units>
  </globalQuota>
  <layerQuotas/>
</gwcQuotaConfiguration>
```

See [Disk Quotas](#) for more information on configuration options for disk quota.

A.4.3.5. Configuring the Standard Search UI for GeoWebCache

Add a new Imagery Provider in the Admin Console:

1. Navigate to the **Admin Console**.
2. Select **Configuration** tab.
3. Select **Standard Search UI** configuration.
4. Click the **Add** button next to **Imagery Providers**
5. Enter configuration for Imagery Provider in new textbox:
6. `{"type": "WMS", "url": "\https://[{FQDN}]:{PORT}/geowebcache/service/wms", "layers": ["states"], "parameters": {"FORMAT": "image/png"}, "alpha": 0.5}`
7. Set the Map Projection to **EPSG:900913** or **EPSG:4326**. (GeoWebCache supports either of these projections.)

NOTE

Currently, GeoWebCache only supports WMS 1.1.1 and below. If the version number is not specified in the imagery provider, DDF will default to version **1.3.0**, and OpenLayers will not project the image tiles properly. Thus, the version **1.1.1** must be specified when using **EPSG:4326** projections.

```
{"type": "WMS", "url": "\https://[{FQDN}]:{PORT}/geowebcache/service/wms", "layers": ["states"], "parameters": {"FORMAT": "image/png", "VERSION": "1.1.1"}, "alpha": 0.5}
```

A.5. Intrigue Application Reference

The Intrigue application provides metadata search and discovery, resource retrieval and workspace management with a 3d or optional 2d map visualization.

A.5.1. Intrigue Prerequisites

To use the Intrigue application, the following applications/features must be installed:

- Platform
- Catalog

A.5.2. Installing Intrigue

Intrigue is included with a standard installation.

A.5.3. Configuring the Intrigue Application

To configure the Intrigue Application:

1. Navigate to the Admin Console.
2. Select the **Intrigue** application.
3. Select the **Configuration** tab.

Table 122. Intrigue Available Configurations

Name	Property	Description
Email Notifier	org.codice.ddf.catalog.ui.query.monitor.email.EmailNotifier	Email Notifier.
Workspace Query Monitor	org.codice.ddf.catalog.ui.query.monitor.impl.WorkspaceQuery Service	Workspace Query Monitor.
Intrigue Search	Intrigue_Search	Intrigue Search.
Workspace Security	org.codice.ddf.catalog.ui.security	Workspace Security.

Table 123. Email Notifier

Name	Id	Type	Description	Default Value	Required
Subject	subjectTemplate	String	Set the subject line template.	Workspace '%[attribute=title]' notification	true

Name	Id	Type	Description	Default Value	Required
Body	bodyTemplate	String	Set the body template.	The workspace '%[attribute=title]' contains up to %[hitCount] results. Log in to see results https://{{FQDN}}:{{PORT}}/search/catalog/#workspaces/{{attribute=id}} .	true
Mail Server	mailHost	String	Set the hostname of the mail server.	localhost	true
From Address	fromEmail	String	Set the 'from' email address.	donotreply@test.com	true

Table 124. Workspace Query Monitor

Name	Id	Type	Description	Default Value	Required
Query Timeout	queryTimeoutMinutes	Long	Set the number of minutes to wait for query to complete.	5	true
Notification Time Interval	queryTimeInterval	Integer	Set the Relative Time Search (past X minutes up to 24 hours). Note: This will query for results from the interval to the time the query is sent out.	1440	true
Email Subscription Interval	cronString	String	Email Subscription Interval (Cron Expression)	0 0 0 * * ?	true

Table 125. Intrigue Search

Name	Id	Type	Description	Default Value	Required
Result Count	resultCount	Integer	Specifies the number of results to request from each source	250	true

Name	Id	Type	Description	Default Value	Required
Imagery Providers	imageryProviders	String	<p>List of imagery providers to use. Valid types are: OSM (OpenStreetMap), AGM (ArcGisMap), BM (BingMap), WMS (WebMapService), WMT (WebMapTile), TMS (TileMapService), GE (GoogleEarth).</p> <p>WMS example: <code>[{"name": "Example WMS", "show": true, "type": "WMS", "url": "http://suite.opengeo.org/geoserver/gwc/service/wms", "layers": ["opengeo:countries"], "parameters": {"FORMAT": "image/png", "VERSION": "1.1.1"}, "alpha":1}]</code></p> <p>OSM example: <code>[{"name": "Example OSM", "show": true, "type": "OSM", "url": "http://a.tile.openstreetmap.org", "fileExtension": "png", "alpha": 1.0}]</code></p> <p>AGM example: <code>[{"name": "Example AGM", "show": true, "type": "AGM", "url": "https://server.arcgisonline.com/arcgis/rest/services/World_Imagery/MapServer"}]</code></p> <p>Multiple layer example: <code>[{"name": "Example AGM", "show": true, "type": "AGM", "url": "https://server.arcgisonline.com/arcgis/rest/services/World_Imagery/MapServer"}, {"name": "Example AGM 2", "show": true, "type": "AGM", "url": "https://server.arcgisonline.com/arcgis/rest/services/World_Street_Map/MapServer"}]</code></p> <p>WMT example: <code>[{"name": "Example WMT", "show": true, "type": "WMT", "url": "http://suite.opengeo.org/geoserver/gwc/service/wmts", "layer": "opengeo:countries", "style": "", "format": "image/jpeg", "tileMatrixSetID": "EPSG:4326", "tileMatrixLabels": ['EPSG:4326:0', 'EPSG:4326:1', 'EPSG:4326:2', 'EPSG:4326:3', 'EPSG:4326:4', 'EPSG:4326:5', 'EPSG:4326:6', 'EPSG:4326:7', 'EPSG:4326:8', 'EPSG:4326:9', 'EPSG:4326:10', 'EPSG:4326:11', 'EPSG:4326:12', 'EPSG:4326:13', 'EPSG:4326:14', 'EPSG:4326:15']}]</code></p>		false

Name	Id	Type	Description	Default Value	Required
Terrain Provider	terrainProvider	String	Terrain provider to use for height data. Valid types are: CT (CesiumTerrain), AGS (ArcGisImageServer), VRW (VRTheWorld). Example: {"type": "CT", "url": "http://example.com"}	{ "type": "CT", "url": "http://assets.agi.com/stk-terrain/tilesets/world/tiles" }	false
Map Projection	projection	String	Projection of imagery providers. Valid values are EPSG:4326 and EPSG:3857 .	EPSG:4326	false
Bing Maps Key	bingKey	String	Bing Maps API key. This should only be set if you are using Bing Maps Imagery or Terrain Providers.		false
Connection Timeout	timeout	Integer	Specifies the client-side connection timeout in milliseconds.	300000	false
Source Poll Interval	sourcePollInterval	Integer	Specifies the interval to poll for sources in milliseconds.	60000	true
Show Sign In	signIn	Boolean	Allow Sign In to Search UI and welcome notice. Enable this if the Search UI is protected.	true	false
Show Tasks	task	Boolean	Show task menu area for long running actions.	false	false
Show Gazetteer	gazetteer	Boolean	Show gazetteer for searching place names.	true	false
Use Online Gazetteer	onlineGazetteer	Boolean	Should the online gazetteer be used? If unchecked a local gazetteer service will be used. This only applies to the search gazetteer in Intrigue.	true	false
Show Uploader	ingest	Boolean	Show upload menu for adding new record.	true	false
Use External Authentication	externalAuthentication	Use an external authentication point, such as IdP	Boolean	false	false
Disable Cache	cacheDisabled	Locally cached results will not be returned in search results.	Boolean	false	false

Name	Id	Type	Description	Default Value	Required
Type Name Mapping	typeNameMapping	String	Use an external authentication point, such as IdP.		false
Read Only Metocard Attributes	readOnly	String	List of metocard attributes that are read-only. NOTE: the provided values will be evaluated as JavaScript regular expressions when matched against metocard attributes.	^checksum\$, ^checksum-algorithm\$, ^id\$, ^resource-download-url\$, ^resource-uri\$, ^resource-derived-uri\$, ^resource-derived-download-url\$, ^modified\$, ,	false
Summary Metocard Attributes	summaryShow	String	List of metocard attributes to display in the summary view.		false
Result Preview Metocard Attributes	resultShow	String	List of metocard attributes to display in the result preview.		false
Attribute Aliases	attributeAliases	String	List of attribute aliases. Example 'title=Title'		false

Name	Id	Type	Description	Default Value	Required
Hidden Attributes	hiddenAttributes	String	List of attributes to be hidden. NOTE: the provided values will be evaluated as JavaScript regular expressions when matched against metocard attributes.	^sortOrder\$, ^sortField\$, ^cql\$, ^polling\$, ^cached\$	false
Query Schedule Frequencies	scheduleFrequencyList	Long	Custom list of schedule frequencies in seconds. This will override the frequency list in the query schedule tab. Leave this empty to use the frequency list on Intrigue.	1800,3600, 7200,14400 ,28800,576 00,86400	true
Allow Editing	isEditingAllowed	Boolean	Allow editing capability to be visible in the UI	true	true
Result Page Size	resultPageSize	Integer	Specifies the number of results allowed per page on the client-side.	25	true
Enable Query Feedback	queryFeedbackEnabled	Boolean	Enable the query comments option.	true	true
Feedback Email Subject Template	queryFeedbackEmailSubjectTemplate	String	Query Feedback template for the email subject.	Query Feedback from {{username}}	true

Name	Id	Type	Description	Default Value	Required
Feedback Email Body Template	queryFeedbackEmailBodyTemplate	String	Query Feedback template for the email body.	Query Feedback\\n\\nAuthenticated User: {{auth_use rname}}\\n User: {{username}}\\nEmail : {{email}}\\nWorkspace: {{workspace_name}} ({{{workspace_id}}})\\n Query: {{{query}}}\\nQuery time: {{query_initiated_time }}\\nQuery status: {{query_st atus}}\\nQu ery_results : {{query_re sults}}\\nC omments: {{comment s}}	true
Feedback Email Destination	queryFeedbackEmailDestinat ion	String	Email destination to send Query Feedback results.		true
Maximum Endpoint Upload Size	maximumUploadSize	Integer	The maximum size of data that the Intrigue REST endpoints will receive. Please note this is not related to the product upload size, this only refers to the json data sent and received from the Intrigue endpoints.	1048576	false

Table 126. Workspace Security

Name	Id	Type	Description	Default Value	Required
System User Attribute	systemUserAttribute	String	The name of the attribute to determine the system user.	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role	true
System User Attribute Value	systemUserAttributeValue	String	The value of the attribute to determine the system user.	admin	true

A.6. Platform Application Reference

The Platform application is considered to be a core application of the distribution. The Platform application provides the fundamental building blocks that the distribution needs to run. These building blocks include subsets of:

- [Karaf](#)
- [CXF](#)
- [Camel](#)

A [Command Scheduler](#) is also included as part of the Platform application to allow users to schedule Command Line Shell Commands.

A.6.1. Platform Application Prerequisites

None.

A.6.2. Installing Platform

The Platform application is installed by default with a standard installation.

A.6.3. Configuring the Platform Application

To configure the Platform Application:

1. Navigate to the Admin Console.
2. Select the **Platform** application.
3. Select the **Configuration** tab.

Table 127. Platform Available Configurations

Name	Property	Description
MIME Custom Types	DDF_Custom_Mime_Type_Resolver	DDF Custom Mime Types.

Name	Property	Description
Logging Service	org.codice.ddf.platform.logging.LoggingService	Logging Service configurations.
Metrics Reporting	MetricsReporting	Metrics Reporting.
HTTP Response Security	org.codice.ddf.security.response.filter.ResponseHeaderConfig	HTTP Response Security response configurations.
Email Service	org.codice.ddf.platform.email.impl.SmtpClientImpl	Email Service configurations.
Landing Page	org.codice.ddf.distribution.landingpage.properties	Starting page for users to interact with DDF.
Platform Command Scheduler	ddf.platform.scheduler.Command	Platform Command Scheduler.
Platform UI	ddf.platform.ui.config	Platform UI configurations.

Table 128. MIME Custom Types

Name	Id	Type	Description	Default Value	Required
Resolver Name	name	String	null	DDF Custom Resolver	false
Priority	priority	Integer	null	10	true
File Extensions to Mime Types	customMimeTypes	String	List of key/value pairs where key is the file extension and value is the mime type, e.g., nitf=image/nitf	null	true

Table 129. Logging Service

Name	Id	Type	Description	Default Value	Required
Max Log Events	maxLogEvents	Integer	The maximum number of log events stored for display in the Admin Console. This must be greater than 0 and must not exceed 5000.	500	true

Table 130. Metrics Reporting

Name	Property	Type	Description	Default Value	Required
Metrics Max Threshold	metricsMaxThreshold	Double	Max value a data sample can be for any metric (used to suppress spike data on metrics graphs)	400000000 0.0	true

Table 131. HTTP Response Security

Name	<code>Id</code>	Type	Description	Default Value	Required
Content Security Policy	<code>xContentSecurityPolicy</code>	String	Instructions for the client browser detailing which location and/or which type of resources may be loaded.		true
X-Frame-Options	<code>xFrameOptions</code>	String	The X-Frame-Options HTTP response header can be used to indicate whether or not a browser may render a page in a frame, iframe or object.		true
X-XSS-Protection	<code>xXssProtection</code>	String	The HTTP X-XSS-Protection response header is a feature that stops pages from loading when they detect reflected cross-site scripting (XSS) attacks.		true

Table 132. Email Service

Name	<code>Property</code>	Type	Description	Default Value	Required
Host	<code>hostName</code>	String	Mail server hostname (must be resolvable by DNS) or IP address.		Yes
Port	<code>portNumber</code>	Integer	Mail server port number.	25	Yes
User Name	<code>userName</code>	String	Mail server user name used only for authenticated connections over TLS.		No
Password	<code>password</code>	Password	Mail server password used only for authenticated connections over TLS.		No

Table 133. Landing Page

Name	<code>Id</code>	Type	Description	Default Value	Required
Description	<code>description</code>	String	Specifies the description to display on the landing page.	As a common data layer, DDF provides secure enterprise-wide data access for both users and systems.	true
Phone Number	<code>phone</code>	String	Specifies the phone number to display on the landing page.		true
Email Address	<code>email</code>	String	Specifies the email address to display on the landing page.		true

Name	Id	Type	Description	Default Value	Required
External Web Site	<code>externalURL</code>	String	Specifies the external web site URL to display on the landing page.		true
Announcements	<code>announcements</code>	String	Announcements that will be displayed on the landing page.	null	true
Branding Background	<code>background</code>	String	Specifies the landing page background color. Use html css colors or #rrggbb.		true
Branding Foreground	<code>foreground</code>	String	Specifies the landing page foreground color. Use html css colors or #rrggbb.		true
Branding Logo	<code>logo</code>	String	Specifies the landing page logo. Use a base64 encoded image.		true
Additional Links	<code>links</code>	String	Additional links to be displayed on the landing page. Use the format <text>,<link> (e.g. <code>example</code> , http://www.example.com). Empty entries are ignored.		yes

Table 134. Platform Command Scheduler

Name	Property	Type	Description	Default Value	Required
Command	<code>command</code>	String	Shell command to be used within the container. For example, <code>log:set DEBUG"></code>		true
Interval	<code>intervalString</code>	String	The Interval String for each execution. Based on the Interval Type, this will either be a Cron String or a Second Interval. (e.x. '0 0 0 1/1 * ? *' or '12')		true
Interval Type	<code>intervalType</code>	String	Interval Type	cronString	true

Table 135. Platform UI Configuration

Name	Id	Type	Description	Default Value	Required
Enable System Usage Message	<code>systemUsageEnabled</code>	Boolean	Turns on a system usage message, which is shown when the Search Application is opened	false	true
System Usage Message Title	<code>systemUsageTitle</code>	String	A title for the system usage Message when the application is opened	true	

Name	Id	Type	Description	Default Value	Required
System Usage Message	<code>systemUsageMessage</code>	String	A system usage message to be displayed to the user each time the user opens the application	true	
Show System Usage Message once per session	<code>systemUsageOncePerSession</code>	With this selected, the system usage message will be shown once for each browser session. Uncheck this to have the usage message appear every time the search window is opened or refreshed.	Boolean	true	true
Header	<code>header</code>	Specifies the header text to be rendered on all pages.	String	true	
Footer	<code>footer</code>	Specifies the footer text to be rendered on all pages.	String	true	
Text Color	<code>color</code>	Specifies the Text Color of the Header and Footer. Use html css colors or <code>#rrggb</code> .	String	true	

Name	Id	Type	Description	Default Value	Required
Background Color	background	Specifies the Background Color of the Header and Footer. Use html css colors or #rrggb.	String	true	

A.7. Registry Application Reference

Registry contains the base registry components, plugins, sources, and interfaces needed for DDF to function as a registry connecting multiple nodes.

A.7.1. Registry Prerequisites

To use the Registry, the following apps/features must be installed:

- Catalog
- Admin
- Spatial
- Platform
- Security

A.7.2. Installing Registry

Registry is not installed by default with a standard installation.

To install:

1. Navigate to the **Admin Console**.
2. Click the **Manage** applications button.
3. Click the **Install** icon.
4. Registry will move to **Active Applications** upon startup.

A.7.3. Customizing Registry Fields

All the fields that appear in a registry node are customizable. This is done through a JSON configuration file located at `<DDF_HOME>/etc/registry/registry-custom-slots.json` that defines the registry fields. In this file there are JSON objects that relate to each part of the edit registry modal. These objects are

- General

- Service
 - ServiceBinding
- Organization
- Person (Contact)
- Content (Content Collection)

Each of the objects listed above is a JSON array of field objects that can be modified. There are some other objects in the JSON file like **PersonName**, **Address**, **TelephoneNumber**, and **EmailAddress** that should not be modified.

Table 136. Field Properties

Property Key	Required	Property Value
key	yes	The string value that will be used to identify this field. Must be unique within field grouping array. This value is what will show up in the generated EBRIM xml.
displayName	yes	The string name that will be displayed in the edit node dialog for this field
description	yes	A brief description of what the field represents or is used for. Shown when user hovers or click the question mark icon for the field.
value	no	The initial or default value of the field. For most cases this should be left as an empty array or string.
type	yes	Identifies what type of field this is. Value must be one of string , date , number , boolean , point , or bounds
required	no	Indicates if this field must be filled out. Default is false . If true an asterisk will be displayed next to the field name.
possibleValues	no	An array of values that could be used for this field. If multiValued=true this list will be used for suggestions for autocomplete. If multiValued=false this list will be used to populate a dropdown.

Property Key	Required	Property Value
multiValued	no	Flag indicating if this field accepts multiple values or not. Default is false.
isSlot	no	Indicates that this field represents a slot value in the EBRIM document. If this is false the key must match a valid EBRIM attribute for the parent object. Default is true.
advanced	no	A flag indicating if this field should be placed under the Advanced section of the edit modal ui. Default is false.
regex	no	A regular expression for validating users input.
regexMessage	no	A message to show the user if the regular expression test fails.
isGroup, constructTitle	N/A	These fields are used for nesting objects and should not be modified

A.7.4. Using Registry

The **Node Information** and **Remote Registries** tabs appear in both the Registry application and the Catalog application.

A.7.4.1. Configuring Identity Node

The "Identity Node" is the local DDF instance. Configure the information to share with other registries/nodes.

1. Navigate to **Registry** (or **Catalog**) application.
2. Navigate to **Node Information** tab.
3. Click the name of the identity node.
4. Complete all *required* and any desired *optional* fields.
 - a. Add any desired **service bindings** under the **Services** tab.
5. Click **Save**.

Table 137. General Information Tab

Field	Description	Type	Required
Node Name	This node's name as it should appear to external systems	string	yes
Node Description	Short description for this node	string	yes

Field	Description	Type	Required
Node Version	This node's Version	string	yes
Security Attributes	Security attributes associated with this node.	String	
Last Updated	Date this entry's data was last updated	Date	
Live Date	Date indicating when this node went live or operational	Date	
Custom Fields	click Add button to add custom fields	Configurable	no
Associations	click Add button to add associations	Configurable	no

Table 138. Services

Field	Description	Type	Required
Service Name	This service name	string	
Service Description	Short description for this service	string	
Service Version	This service version	string	
Service Type	Identifies the type of service this is by a URN.	string	
Bindings (Click Add to add a service binding)			
Binding Name	This binding name	String	yes
Binding Description	Short description for this binding	String	
Binding Version	This binding version		
Access URL	The url used to access this binding		
Service Binding Type	The binding type for the service		
URL Property Key	Property that the accessURI value should be put into for source creation		
Custom Fields	click Add button to add custom fields	Configurable	no
Associations	click Add button to add associations	Configurable	no

Table 139. Organizations Tab (click **Add** to add an organization)

Field	Description	Type	Required
Organization Name	This organization's name	string	yes
Address	This organization's primary address	Expand to enter address information	yes

Field	Description	Type	Required
TelephoneNumber	Primary contact number for this organization		no
Email	Primary contact email for this organization		no
Custom Fields	click Add button to add custom fields	Configurable	no
Associations	click Add button to add associations	Configurable	no

Table 140. Contacts (click **Add** button to add contact info)

Field	Description	Type	Required
Contact Title	Contact Title	String	yes
Contact First Name	Contact First Name	String	yes
Contact Last Name	Contact Last Name	String	yes
Address	Address for listed contact	String	minimum one
Phone number	Contact phone number		minimum one
Email	Contact email	String	minimum one
Custom Fields	click Add button to add custom fields	Configurable	no
Associations	click Add button to add associations	Configurable	no

Table 141. Collections (Click **Add** to add Content Collection(s))

Field	Description	Type	Required
Content Name	Name for this metadata content	string	yes
Content Description	Short description for this metadata content	string	no
Content Object Type	The kind of content object this will be. Default value should be used in most cases.	string	yes
Custom Fields	click Add button to add custom fields	Configurable	no
Associations	click Add button to add associations	Configurable	no

A.7.4.1.1. Adding a Service Binding to a Node

Advertise the methods other nodes use to connect to the local DDF instance.

1. Navigate to Admin Console.
2. Select Registry or Catalog.
 - a. (**Node Information** tab is editable from either application.)

3. Click the name of the desired local node.
4. Click the **Services** tab.
5. Click **Add** to add a service.
6. Expand new Service.
7. Enter Service name and details.
8. Click **Add** to add binding.
9. Select Service Binding type.
 - a. Select one of the defaults or *empty* for a custom service binding.
 - b. If selecting *empty*, fill in all required fields.
10. Click Save.

A.7.4.2. Publishing to Others

Send details about the local DDF instance to other nodes.

1. Navigate to the **Remote Registries** tab in either Registry or Catalog application.
2. Click **Add** to add a remote registry.
3. Enter Registry Service (CSW) Url.
4. Confirm **Allow Push** is checked.
5. Click **Add** to save the changes.
6. Navigate to the **Sources** Tab in Catalog App
7. Click desired node to be published.
8. Under **Operations**, click the *Publish to ... * link that corresponds to the desired registry.

A.7.4.3. Subscribing to Another Node

Receive details about another node.

1. Navigate to the **Remote Registries** tab in either Registry or Catalog application.
2. Click **Add** to add a remote registry.
3. Add the URL to access node.
4. Enter any needed credentials in the Username/password fields.
5. Click **Save/Add**.

A.7.4.3.1. Editing a Subscription

Update the configuration of an existing subscription.

1. Navigate to the **Remote Registries** tab in either Registry or Catalog application.
2. Click the name of the desired subscription.
3. Make changes.

4. Click **Save**.

A.7.4.3.2. Deleting a Subscription

Remove a subscription.

1. Click the **Delete** icon at the top of the **Remote Registries** tab.
2. Check the boxes of the Registry Nodes to be deleted.
3. Select the **Delete** button.

A.7.5. Configuring the Registry Application

To configure the Registry Application:

1. Navigate to the Admin Console.
2. Select the **Registry** application.
3. Select the **Configuration** tab.

Table 142. Registry Available Configurations

Name	Property	Description
CSW Registry Store	Csw_Registry_Store	Registry CSW Store.
Registry Policy Plugin	org.codice.ddf.registry.policy.RegistryPolicyPlugin	Registry Policy Plugin.
Registry Source Configuration Handler	Registry_Configuration_Event_Handler	Registry Source Configuration Handler configurations.

Table 143. CSW Registry Store

Name	Id	Type	Description	Default Value	Required
Registry ID	<code>id</code>	String	The unique name of the store	null	true
Registry Service URL	<code>cswUrl</code>	String	URL to the endpoint implementing CSW spec capable of returning ebrim formatted records	null	true
Username	<code>username</code>	String	Username for CSW Service (optional)	null	false
Password	<code>password</code>	Password	Password for CSW Service (optional)	null	false
Allow Push	<code>pushAllowed</code>	Boolean	Enable push (write) to this registry	true	true
Allow Pull	<code>pullAllowed</code>	Boolean	Enable pull (read) from this registry	true	true
Push Identity Node	<code>autoPush</code>	Boolean	Enable an automatic publish from the local identity node to this registry. Setting this to Off will have the effect of unpublishing the identity from this registry.	true	true

Table 144. Registry Policy Plugin

Name	Id	Type	Description	Default Value	Required
Registry Create Attributes	<code>createAccessPolicyStrings</code>	String	Roles/attributes required for Create operations on registry entries. Example: {role=role1;type=type1}	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest	true
Registry Update Attributes	<code>updateAccessPolicyStrings</code>	String	Roles/attributes required for Update operations on registry entries. Example: {role=role1;type=type1}	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest	true
Registry Delete Attributes	<code>deleteAccessPolicyStrings</code>	String	Roles/attributes required for Delete operations on registry entries. Example: {role=role1;type=type1}	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest	true
Registry Read Attributes	<code>readAccessPolicyStrings</code>	String	Roles/attributes required for reading registry entries. Example: {role=role1;type=type1}	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest	true
Registry Admin Attributes	<code>registryByPassPolicyStrings</code>	String	Roles/attributes required for an admin to bypass all filtering/access controls. Example: {role=role1;type=type1}	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=system-admin	true
Disable Registry Write Access	<code>registryDisabled</code>	Boolean	Disables all write access to registry entries in the catalog. Only users with Registry Admin Attributes will be able to write registry entries	null	false
Entries are White List	<code>whiteList</code>	Boolean	A flag indicating whether or not the Registry Entry Ids represent a 'white list' (allowed - checked) or a 'black list' (blocked - unchecked) ids	null	false

Name	Id	Type	Description	Default Value	Required
Registry Entries Ids	<code>registryEntryIds</code>	String	List of registry entry ids to be used in the white/black list.	null	false

Table 145. Registry Source Configuration Handler

Name	Id	Type	Description	Default Value	Required
Url Binding Name	<code>urlBindingName</code>	String	The url name for communicating with the specific instance.	<code>urlBindingName</code>	true
BindingType to Factory PID	<code>bindingTypeFactoryPid</code>	String	Key/Value mappings of binding type to factory PID	<code>CSW_2.0.2=Csw_Federated_Source,WFS_1.0.0=Wfs_v1_0_0_Federated_Source,OpenSearch_1.0.0=OpenSearchSource</code>	true
Remove Configurations on Metocard Delete	<code>cleanUponDelete</code>	Boolean	Flag used to determine if configurations should be deleted when the metocard is deleted.	false	true
Activate Configurations	<code>activateConfigurations</code>	Boolean	Flag used to determine if a configuration should be activated on creation	false	true
Preserve Active Configuration	<code>preserveActiveConfigurations</code>	Boolean	Flag used to determine if configurations should be preserved. If true will only allow auto activation on creation. If false auto activation will happen on updates as well. Only applicable if activateConfigurations is true.	true	true
Source Activation Priority Order	<code>sourceActivationPriorityOrder</code>	String	This is the priority list used to determine which source should be activated on creation	<code>100CSW_2.0.2,WFS_1.0.0,OpenSearch_1.0.0</code>	true

A.8. Resource Management Application Reference

The Resource Management Application provides administrative functionality to monitor and manage data usage on the system. This application allows an administrator to:

- View data usage.

- Set limits on users.
- View and terminate searches that are in progress.

Components of the Resource Management application include:

Resource Management Data Usage Tab

View data usage and configure users' data limits and reset times for those limits.

Resource Management Queries Tab

View and cancel actively running queries.

A.8.1. Resource Management Prerequisites

To use the Resource Management Application, the following apps/features must be installed:

- Platform
- Security
- Admin
- Catalog

A.8.2. Installing Resource Management

The Resource Management application is not installed by default with a standard installation.

To Install:

1. Navigate to Admin Console.
2. Select **Manage** applications.
3. Select **Resource Management** application.
4. Select **Start**.
5. Resource Management application moves to **Active Applications** upon start.

A.8.3. Configuring the Resource Management Application

To configure the Resource Management Application:

1. Navigate to the Admin Console.
2. Select the **Resource Management** application.
3. Select the **Configuration** tab.

Table 146. Resource Management Available Configurations

Name	Property	Description
Data Usage	org.codice.ddf.resourcemanage ment.usage	Data Usage configurations.

Table 147. Data Usage

Name	Id	Type	Description	Default Value	Required
Monitor Local Sources	monitorLocalSources	Boolean	When checked, the Data Usage Plugin will also consider data usage from local sources.	false	true

A.9. Security Application Reference

The Security application provides authentication, authorization, and auditing services for the DDF. These services comprise both a framework that developers and integrators can extend as well as a reference implementation that meets security requirements.

This section documents the installation, maintenance, and support of this application.

Applications Included in Security

- Security CAS
- Security Core
- Security Encryption
- Security IdP
- Security PEP
- Security PDP
- Security STS

A.9.1. Security Prerequisites

To use the Security application, the following applications/features must be installed:

- Platform

A.9.2. Installing Security

Security is included with a standard installation.

A.9.3. Configuring the Security Application

To configure the Security Application:

1. Navigate to the Admin Console.
2. Select the **Security** application.
3. Select the **Configuration** tab.

Table 148. Security Available Configurations

Name	Property	Description
Security STS LDAP and Roles Claims Handler	Claims_Handler_Manager	STS Ldap and Roles Claims Handler Configuration.
Security SOAP Guest Interceptor	org.codice.ddf.security.interceptor.GuestInterceptor	Security SOAP Guest Interceptor.
IdP Client	org.codice.ddf.security.idp.client.IdpMetadata	IdP Client configurations.
Logout Page	org.codice.ddf.security.idp.client.LogoutRequestService	Logout Page configurations.
Web Context Policy Manager	org.codice.ddf.security.policy.context.impl.PolicyManager	Web Context Security Policies.
File Based Claims Handler	org.codice.ddf.security.sts.claims.property.PropertyFileClaimsHandler	File Based Claims Handler.
Session	org.codice.ddf.security.filter.login.Session	Session configurations.
Login Page	org.codice.ddf.security.handler.guest.configuration	Login Page configurations.
IdP Handler	org.codice.ddf.security.idp.client.IdpHandler	IdP Handler configurations.
Security AuthZ Realm	ddf.security.pdp.realm.AuthzRealm	AuthZ Security configurations.
SAML NameID Policy	ddf.security.service.SecurityManager	SAML NameID Policy.
Security STS Address Provider	ddf.security.sts.address.provider	STS Address Provider.
Security STS Server	ddf.security.sts	STS configurations.
Security STS Client	ddf.security.sts.client.configuration	STS Client configurations.
Security STS Guest Claims Handler	ddf.security.sts.guestclaims	Guest Claims Handler configurations.
Guest Validator	ddf.security.sts.guestvalidator	Security STS Guest Validator configurations.
Security STS WSS	ddf.security.sts.wss.configuration	STS WSS configurations.
Security STS PKI Token Validator	org.codice.ddf.security.validator.pki	STS PKI Token Validator configurations.

Table 149. Security STS LDAP and Roles Claims Handler

Name	Property	Type	Description	Default Value
LDAP URL	url	String	LDAP or LDAPS server and port	ldaps://\${org.codice.ddf stem.hostname}:1636
LDAP Bind Method	bindMethod	String	LDAP Bind Method	Simple
Realm (GSSAPI and Digest MD5 SASL and only)	realm	String	Realm to use during bind. This applies only to Kerberos and Digest MD5 authentication.	
KDC Address (GSSAPI SASL only)	kdcAddress	String	KDC address to use during bind. This applies only to Kerberos authentication.	
StartTLS	startTls	Boolean	Determines whether or not to use StartTLS when connecting via the ldap protocol. This setting is ignored if the URL uses ldaps.	false
LDAP Bind User DN	ldapBindUserDn	String	DN of the user to bind with LDAP. This user should have the ability to verify passwords and read attributes for any user.	cn=admin
LDAP Bind User Password	password	Password	Password used to bind user with LDAP.	secret
LDAP Group User Membership Attribute	membershipUserAttribute	String	Attribute used as the membership attribute for the user in the group. Usually this is uid, cn, or something similar.	uid
LDAP User Login Attribute	loginUserAttribute	String	Attribute used as the login username. Usually this is uid, cn, or something similar.	uid
LDAP Base User DN	userBaseDn	String	Full LDAP path to where users can be found.	ou=users\,dc=example\,dc=com
Override User Certificate DN	overrideCertDn	Boolean	When checked, this setting will ignore the DN of a user and instead use the LDAP Base User DN value.	false
LDAP Group ObjectClass	objectClass	String	ObjectClass that defines structure for group membership in LDAP. Usually this is groupOfNames or groupOfUniqueNames.	groupOfNames

Name	Property	Type	Description	Default Value
LDAP Membership Attribute	memberNameAttribute	String	Attribute used to designate the user's name as a member of the group in LDAP. Usually this is member or uniqueMember.	member
LDAP Base Group DN	groupBaseDn	String	Full LDAP path to where groups can be found.	ou=groups\,dc=example\,dc=com
Attribute Map File	propertyFileLocation	String	Location of the file which contains user attribute maps to use.	etc/ws-security/attributeMap.properties

Table 150. Security SOAP Guest Interceptor

Name	Id	Type	Description	Default Value	Required
Deny Guest Access	guestAccessDenied	Boolean	If set to true, no guest access will be allowed via this guest interceptor. If set to false, this interceptor will generate guest tokens for incoming requests that lack a WS-Security header.	false	false

Table 151. IdP Client

Name	Id	Type	Description	Default Value
IdP Metadata	metadata	String	Refer to metadata by HTTPS URL (https://), file URL (file:), or an XML block(<md:EntityDescriptor>...</md:EntityDescriptor>).	https://\${org.codice.ddf.system.hostname}:\${org.codice.ddf.system.httpsPort}/services/idp/login/metadata
Perform User-Agent Check	userAgentCheck	Boolean	If selected, this will allow clients that do not support ECP and are not browsers to fall back to PKI, BASIC, and potentially GUEST authentication, if enabled.	true

Table 152. Logout Page

Name	Id	Type	Description	Default Value
Log Out Page Time Out	logOutPageTimeOut	Long	This is the time limit that the IDP client will wait for a user to click log out on the log out page. Any requests that take longer than this time for the user to submit will be rejected."/>	3600000

Table 153. Web Context Policy Manager

Name	Id	Type	Description	Default Value	Required
Context Traversal Depth	traversalDepth	Integer	Depth to which paths will be traversed. Any value greater than 500 will be set to 500.	20	true
Context Realms	realms	String	List of realms supporting each context. Karaf is provided by default. Example: /=karaf	/=karaf	true
Authentication Types	authenticationTypes	String	List of authentication types required for each context. List of default valid authentication types are: IDP, SAML, BASIC, PKI, GUEST. Example: /context=AUTH1	AUTH2	AUTH3
/=IDP	GUEST,/solr=SAML	PKI	BASIC	true	Required Attributes
requiredAttributes	String	List of attributes required for each Web Context. Example: /context={role=role1;type=type1}	/=,/admin={http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=system-admin},/solr={http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=system-admin},/system={http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=system-admin},/security-config={http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=system-admin}	true	White Listed Contexts

Table 154. File Based Claims Handler

Name	Id	Type	Description	Default Value	Required
Role Claim Type	roleClaimType	String	Role claim URI.	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role	true

Name	Id	Type	Description	Default Value	Required
ID Claim Type	<code>idClaimType</code>	String	ID claim URI.	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier	true
User Role File	<code>propertyFileLocation</code>	String	Location of the file which maps roles to users.	etc/users.properties	true
User Attribute File	<code>attributeFileLocation</code>	String	Location of the file which maps attributes to users.	etc/users.attributes	true

Table 155. Session

Name	Id	Type	Description	Default Value	Required
Session Timeout (in minutes)	<code>expirationTime</code>	Integer	The number of minutes after a session has been inactive that it should be invalidated.	31	true

Table 156. Login Page

Name	Id	Type	Description	Default Value	Required
Header	<code>header</code>	String	Specifies the header text to be rendered on the Login page		true
Footer	<code>footer</code>	String	Specifies the footer text to be rendered on the Login page		true
Style	<code>style</code>	String	Specifies the style of the Header and Footer.	green	true
Text Color	<code>textColor</code>	String	Specifies the text color of the Header and Footer.	banner-text-white	true

Table 157. IdP Handler

Name	Id	Type	Description	Default Value
Authentication Context Class	authContextClasses	String	Authentication Context Classes that are considered acceptable means of authentication by the IdP Client.	urn:oasis:names:tc:SAML:2.0:ac:classes:Password,urn:oasis:names:tc:SAML:2.0:ac:classes:PasswordProtectedTransport,urn:oasis:names:tc:AML:2.0:ac:classes:X509,urn:oasis:names:tc:SAML:2.0:ac:classes:SmartcardPKI,urn:oasis:names:tc:SAML:2.0:ac:classes:SoftwarePKI,urn:oasis:names:tc:AML:2.0:ac:classes:SPKI,urn:oasis:names:tc:SAML:2.0:ac:classes:TLSClient

Table 158. Security AuthZ Realm

Name	Id	Type	Description	Default Value	Required
Match-All Mappings	matchAllMappings	String	List of 'Match-All' subject attribute to Metocard attribute mapping. All values of this metocard key must be present in the corresponding subject key values. Format is subjectAttrName=metocardAttrName.		false
Match-One Mappings	matchOneMappings	String	List of 'Match-One' subject attribute to Metocard attribute mapping. One value of this metocard key must be present in the corresponding subject key values. Format is subjectAttrName=metocardAttrName.		false

Name	Id	Type	Description	Default Value	Required
Environment Attributes	environmentAttributes	String	List of environment attributes to pass to the XACML engine. Format is attributeId=attributeValue1,attributeValue2.		false

Table 159. SAML NameID Policy

Name	Id	Type	Description	Default Value	Required
SAML NameID Policy	usernameAttributeList	String	List of attributes that are considered for replacing the username of the logged in user. If any of these attributes match any of the attributes within the SecurityAssertion, the value of the first matching attribute will be used as the username. (Does not apply when NameIDFormat is of the following: X509, persistent, kerberos or unspecified, and the username is not empty).	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier , uid	true

Table 160. Security STS Address Provider

Name	Id	Type	Description	Default Value	Required
Use WSS STS	useWss	Boolean	If you have a WSS STS configured, you may prefer to use it for services that need the STS address, such as SOAP sources.	false	true

Table 161. Security STS Server

Name	Id	Type	Description	Default Value	Required
SAML Assertion Lifetime	lifetime	Long	Set the number of seconds that an issued SAML assertion will be good for.	1800	true
Token Issuer	issuer	String	The name of the server issuing tokens. Generally this is the cn or hostname of this machine on the network.	<code>#{org.codice.ddf.system.hostname}</code>	true
Signature Username	signatureUsername	String	Alias of the private key in the STS Server's keystore used to sign messages.	<code>#{org.codice.ddf.system.hostname}</code>	true

Name	Id	Type	Description	Default Value	Required
Encryption Username	encryptionUsername	String	Alias of the private key in the STS Server's keystore used to encrypt messages.	\${org.codice.ddf.system.hostname}	true

Table 162. Security STS Client

Name	Id	Type	Description	Default Value	Required
SAML Assertion Type	assertionType	String	The version of SAML to use. Most services require SAML v2.0. Changing this value from the default could cause services to stop responding.	http://docs.oasis-open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0	true
SAML Key Type	keyType	String	The key type to use with SAML. Most services require Bearer. Changing this value from the default could cause services to stop responding.	http://docs.oasis-open.org/ws-sx/ws-trust/200512/Bearer	true
SAML Key Size	keySize	String	The key size to use with SAML. The default key size is 256 and this is fine for most applications. Changing this value from the default could cause services to stop responding.	256	true
Use Key	useKey	Boolean	Signals whether or not the STS Client should supply a public key to embed as the proof key. Changing this value from the default could cause services to stop responding.	true	true

Name	Id	Type	Description	Default Value	Required
STS WSDL Address	<code>address</code>	String	STS WSDL Address	<code> \${org.codice.ddf.system.protocol }\${org.codice.ddf.system.hostname}: \${org.codice.ddf.system.port }\${org.codice.ddf.system.rootContext }/SecurityTokenService?wsdl</code>	true
STS Endpoint Name	<code>endpointName</code>	String	STS Endpoint Name.	<code>{http://docs.oasis-open.org/wss-sx/ws-trust/200512/}STS_Port</code>	false
STS Service Name	<code>serviceName</code>	String	STS Service Name.	<code>{http://docs.oasis-open.org/wss-sx/ws-trust/200512/}SecurityTokenService</code>	false
Signature Properties	<code>signatureProperties</code>	String	Path to Signature crypto properties. This path can be part of the classpath, relative to ddf.home, or an absolute path on the system.	<code>etc/ws-security/server/signature.properties</code>	true
Encryption Properties	<code>encryptionProperties</code>	String	Path to Encryption crypto properties file. This path can be part of the classpath, relative to ddf.home, or an absolute path on the system.	<code>etc/ws-security/server/encryption.properties</code>	true
STS Properties	<code>tokenProperties</code>	String	Path to STS crypto properties file. This path can be part of the classpath, relative to ddf.home, or an absolute path on the system.	<code>etc/ws-security/server/signature.properties</code>	true

Name	Id	Type	Description	Default Value	Required
Claims	claims	String	List of claims that should be requested by the STS Client.	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier,http://schemas.xmlsoap.org/ws/2005/05/identity/claims/emailaddress,http://schemas.xmlsoap.org/ws/2005/05/identity/claims/surname,http://schemas.xmlsoap.org/ws/2005/05/identity/claims/givenname,http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role	true

Table 163. Security STS Guest Claims Handler

Name	Id	Type	Description	Default Value	Required
Attributes	attributes	String	The attributes to be returned for any Guest user.	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier=guest , http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role=guest	true

Table 164. Guest Validator

Name	Id	Type	Description	Default Value	Required
Supported Realms	supportedRealms	String	The realms that this validator supports.	karaf,ldap	true

Table 165. Security STS WSS

Name	Id	Type	Description	Default Value	Required
SAML Assertion Type	assertionType	String	The version of SAML to use. Most services require SAML v2.0. Changing this value from the default could cause services to stop responding.	http://docs.oasis-open.org/wss/oasis-wss-saml-token-profile-1.1#SAMLV2.0	true
SAML Key Type	keyType	String	The key type to use with SAML. Most services require Bearer. Changing this value from the default could cause services to stop responding.	http://docs.oasis-open.org/ws-sx/ws-trust/200512/Bearer	true
SAML Key Size	keySize	String	The key size to use with SAML. The default key size is 256 and this is fine for most applications. Changing this value from the default could cause services to stop responding.	256	true

Name	Id	Type	Description	Default Value	Required
Use Key	<code>useKey</code>	Boolean	Signals whether or not the STS Client should supply a public key to embed as the proof key. Changing this value from the default could cause services to stop responding.	true	true
STS WSDL Address	<code>address</code>	String	STS WSDL Address	<code> \${org.codice.ddf.system.protocol }\${org.codice.ddf.system.hostname}: \${org.codice.ddf.system.httpsPort} \${org.codice.ddf.system.rootContext }/SecurityTokenService?wsdl</code>	true
STS Endpoint Name	<code>endpointName</code>	String	STS Endpoint Name.	<code>{http://docs.oasis-open.org/wss-sx/ws-trust/200512/}STS_Port</code>	false
STS Service Name	<code>serviceName</code>	String	STS Service Name.	<code>{http://docs.oasis-open.org/wss-sx/ws-trust/200512/}SecurityTokenService</code>	false
Signature Properties	<code>signatureProperties</code>	String	Path to Signature crypto properties. This path can be part of the classpath, relative to ddf.home, or an absolute path on the system.	<code>etc/ws-security/server/signature.properties</code>	true
Encryption Properties	<code>encryptionProperties</code>	String	Path to Encryption crypto properties file. This path can be part of the classpath, relative to ddf.home, or an absolute path on the system.	<code>etc/ws-security/server/encryption.properties</code>	true

Name	Id	Type	Description	Default Value	Required
STS Properties	tokenProperties	String	Path to STS crypto properties file. This path can be part of the classpath, relative to ddf.home, or an absolute path on the system.	etc/ws-security/server/signature.properties	true
Claims	claims	String	Comma-delimited list of claims that should be requested by the STS.	http://schemas.xmlsoap.org/ws/2005/05/identity/claims/nameidentifier , http://schemas.xmlsoap.org/ws/2005/05/identity/claims/emailaddress , http://schemas.xmlsoap.org/ws/2005/05/identity/claims/surname , http://schemas.xmlsoap.org/ws/2005/05/identity/claims/givenname , http://schemas.xmlsoap.org/ws/2005/05/identity/claims/role	true

Table 166. Security STS PKI Token Validator

Name	Id	Type	Description	Default Value	Required
Realms	realms	String	The realms to be validated by this validator.	karaf	true
Do Full Path Validation	pathValidation	Boolean	Validate the full certificate path. Uncheck to only validate the subject cert. (RFC5280 6.1)	true	true

A.10. Solr Catalog Application Reference

DDF uses [Solr](#) for data storage, by default.

A.10.1. Solr Catalog Prerequisites

To use the Solr Catalog Application, the following apps/features must be installed:

- Platform
- Catalog

A.10.2. Installing Solr Catalog

The Solr Catalog Application is installed by default with a standard installation.

A.10.3. Configuring the Solr Catalog Application

To configure the Solr Catalog Application:

1. Navigate to the Admin Console.
2. Select the **Solr Catalog** application.
3. Select the **Configuration** tab.

Table 167. Solr Catalog Available Configurations

Name	Property	Description
Solr Catalog Provider	ddf.catalog.solr.provider.SolrCatalogProvider	Solr Catalog Provider.

Table 168. Solr Catalog Provider

Name	Property	Type	Description	Default Value	Required
Force Auto Commit	forceAutoCommit	Boolean	WARNING: Performance Impact. Only in special cases should auto-commit be forced. Forcing auto-commit makes the search results visible immediately.	false	true
Disable Text Path indexing	disableTextPath	Boolean	Disables the ability to make Text Path queries by disabling the Text Path index. Disabling Text Path indexing typically increases ingest performance.	false	true

A.11. Spatial Application Reference

The Spatial Application provides KML transformer and a KML network link endpoint that allows a user to generate a View-based KML Query Results Network Link.

A.11.1. Offline Gazetteer Service

In the Spatial Application, the `offline-gazetteer` is installed by default. This feature enables you to use an offline index of GeoNames data (as an alternative to the GeoNames Web service enabled by the `webservice-gazetteer` feature) to perform searches via the gazetteer search box in the Search UI.

By default a small index set is included with the offline gazetteer. This index can be expanded or updated by using the `geonames:update` command.

A.11.1.1. Spatial GeoNames Console Commands

The `geonames` commands provide the ability to interact with a local GeoNames index. The local GeoNames index is used by the `offline-gazetteer` feature, which is an optional feature available in this application and is explained above. Note that these commands are only available if the `offline-gazetteer` feature is installed.

Table 169. *GeoNames Command Descriptions*

Command	Description	Usage
<code>update</code>	<p>Adds new entries to an existing local GeoNames index. Entries can be manually downloaded from http://download.geonames.org/export/dump, where the absolute path of the file would be passed as an argument to the command (ex. /Users/johndoe/Downloads/AU.zip). Currently .txt and .zip files are supported for manual entries. Entries can also be automatically downloaded from http://download.geonames.org/export/dump by passing the country code as an argument to the command (ex. AU) which will add the country to the local GeoNames index. The full list of country codes available can be found in http://download.geonames.org/export/dump/countryInfo.txt. Using the argument "all" will download all of the current country codes (this process may take some time). In addition to country codes, GeoNames also provides entries for cities based on their population sizes. The arguments "cities1000", "cities5000", and "cities15000" will add cities to the index that have at least 1000, 5000, or 15000 people respectively.</p> <p>The index location can be configured via the Admin Console. By default, the index location is <code>data/geonames-index</code>. If you specify a relative path, it is relative to the location of the unzipped DDF distribution. You may specify an absolute path if you want the index to be located somewhere else.</p> <p>The <code>-c</code> or <code>--create</code> flag can be added to create a new GeoNames index. This will overwrite any existing index at the location specified in the Admin Console. The new index will be filled with the entries in the file you pass to the command. You must create an index before you can add additional entries to it (i.e. running the command without the <code>-c</code> or <code>--create</code> flag).</p>	<code>geonames:update</code>

A.11.2. Spatial Prerequisites

To use the Spatial Application, the following apps/features must be installed:

- Platform
- Catalog

A.11.3. Installing Spatial

The Spatial application is installed by default with a standard installation.

A.11.4. Configuring the Spatial Application

To configure the Spatial Application:

1. Navigate to the Admin Console.
2. Select the **Spatial** application.
3. Select the **Configuration** tab.

Table 170. Spatial Available Configurations

Name	Property	Description
CSW Specification Profile Federated Source	Csw_Federated_Source	CSW Specification Profile Federated Source should be used when federating to an external CSW service.
CSW Federation Profile Source	Csw_Federation_Profile_Source	DDF's full fidelity CSW Federation Profile. Use this when federating to a DDF-based system.
CSW Transactional Profile Federated Source	Csw_Transactional_Federated_Source	CSW Federated Source that supports transactions (create, update, delete).
GeoCoder Plugin	org.codice.ddf.spatial.geocoding.plugin.GeoCoderPlugin	GeoCoder Plugin.
GMD CSW ISO Federated Source	Gmd_Csw_Federated_Source	CSW Federated Source using the Geographic MetaData (GMD) format (ISO 19115:2003).
Spatial KML Endpoint	org.codice.ddf.spatial.kml.endpoint.KmlEndpoint	Spatial KML Endpoint.
Metocard to WFS Feature Map	org.codice.ddf.spatial.ogc.wfs.catalog.mapper.MetocardMapper	Metocard to WFS Feature Map.
WFS 2.0.0 Connected Source	Wfs_v2_0_0_Connected_Source	WFS 2.0.0 Connected Source.
WFS v1.0.0 Federated Source	Wfs_v1_0_0_Federated_Source	WFS v1.0.0 Federated Source.
WFS 2.0.0 Connected Source	Wfs_v2_0_0_Connected_Source	WFS 2.0.0 Connected Source.
WFS 2.0.0 Federated Source	Wfs_v2_0_0_Federated_Source	WFS 2.0.0 Federated Source.

Name	Property	Description
Spatial KML Style Map Entry	org.codice.ddf.spatial.kml.style	Spatial KML Style Map Entry.

Table 171. CSW Specification Profile Federated Source

Name	Id	Type	Description	Default Value	Required
Source ID	<code>id</code>	String	The unique name of the Source	null	true
CSW URL	<code>cswUrl</code>	String	URL to the endpoint implementing the Catalogue Service for Web (CSW) spec	<code> \${org.codice.ddf.system.protocol }\${org.codice.ddf.system.hostname}:\${org.codice.ddf.system.port}\${org.codice.ddf.system.rootContext}/csw</code>	true
Event Service Address	<code>eventServiceAddress</code>	String	DDF Event Service endpoint.	<code> \${org.codice.ddf.system.protocol }\${org.codice.ddf.system.hostname}:\${org.codice.ddf.system.port}\${org.codice.ddf.system.rootContext}/csw/subscription</code>	false
Register for Events	<code>registerForEvents</code>	Boolean	Check to register for events from this source.	false	false
Username	<code>username</code>	String	Username for CSW Service (optional)	null	false
Password	<code>password</code>	Password	Password for CSW Service (optional)	null	false
Disable CN Check	<code>disableCnCheck</code>	Boolean	Disable CN check for the server certificate. This should only be used when testing.	false	true
Coordinate Order	<code>coordinateOrder</code>	String	Coordinate order that remote source expects and returns spatial data in	LON_LAT	true

Name	Id	Type	Description	Default Value	Required
Use posList in LinearRing	usePosList	Boolean	Use a <posList> element rather than a series of <pos> elements when issuing geospatial queries containing a LinearRing	false	false
Metocard Mappings	metocardMappings	String	Mapping of the Metocard Attribute names to their CSW property names. The format should be 'title=dc:title'.	effective=created,created=dateSubmitted,modified=modified,thumbnail=references,content-type=type,id=identifier,resource-uri=source	false
Poll Interval	pollInterval	Integer	Poll Interval to Check if the Source is available (in minutes - minimum 1).	5	true
Connection Timeout	connectionTimeout	Integer	Amount of time to attempt to establish a connection before timing out,in milliseconds.	30000	true
Receive Timeout	receiveTimeout	Integer	Amount of time to wait for a response before timing out,in milliseconds.	60000	true
Output Schema	outputSchema	String	Output Schema	http://www.opengis.net/cat/csw/2.0.2	true
Query Type Name	queryTypeName	String	Qualified Name for the Query Type used in the CSW GetRecords request	csw:Record	true
Query Type Namespace	queryTypeNameSpace	String	Namespace for the Query Type used in the CSW GetRecords request	http://www.opengis.net/cat/csw/2.0.2	true
Force CQL Text as the Query Language	isCqlForce	Boolean	Force CQL Text	false	true
Forced Spatial Filter Type	forceSpatialFilter	String	Force only the selected Spatial Filter Type as the only available Spatial Filter.	NO_FILTER	false
Security Attributes	securityAttributeStrings	String	Security attributes for this source	null	true

Table 172. CSW Federation Profile Source

Name	Id	Type	Description	Default Value	Required
Source ID	id	String	The unique name of the Source	CSW	true
CSW URL	csgetUrl	String	URL to the endpoint implementing the Catalogue Service for Web (CSW) spec	\${org.codice.ddf.system.protocol }\${org.codice.ddf.system.hostname}: \${org.codice.ddf.system.port} \${org.codice.ddf.system.rootContext}/csw	true
CSW Event Service Address	eventServiceAddress	String	CSW Event Service endpoint.	\${org.codice.ddf.system.protocol }\${org.codice.ddf.system.hostname}: \${org.codice.ddf.system.port} \${org.codice.ddf.system.rootContext}/csw/subscription	false
Register for Events	registerForEvents	Boolean	Check to register for events from this connected source.	false	false
Username	username	String	Username for CSW Service (optional)	null	false
Password	password	String	Password for CSW Service (optional)	null	false
Connection Timeout	connectionTimeout	Integer	Amount of time to attempt to establish a connection before timing out,in milliseconds.	30000	true
Receive Timeout	receiveTimeout	Integer	Amount of time to wait for a response before timing out,in milliseconds.	60000	true

Table 173. CSW Transactional Profile Federated Source

Name	Id	Type	Description	Default Value	Required
Source ID	<code>id</code>	String	The unique name of the Source		true
CSW URL	<code>cswUrl</code>	String	URL to the endpoint implementing the Catalogue Service for Web (CSW) spec	<code> \${variable-name}org.codice.ddf.system.protocol}\${variable-name}org.codice.ddf.system.hostName}:\${variable-name}org.codice.ddf.system.port }\${variable-name}org.codice.ddf.system.rootContext}/cs</code>	true
Event Service Address	<code>eventServiceAddress</code>	String	Event Service endpoint.	<code> \${variable-name}org.codice.ddf.system.protocol}\${variable-name}org.codice.ddf.system.hostName}:\${variable-name}org.codice.ddf.system.port }\${variable-name}org.codice.ddf.system.rootContext}/csw/subscription</code>	false
Register for Events	<code>registerForEvents</code>	Boolean	Check to register for events from this source.	false	false
username	<code>Username</code>	String	Username for CSW Service (optional)		false
Password	<code>password</code>	Password	Password for CSW Service (optional)		false

Name	<code>Id</code>	Type	Description	Default Value	Required
Disable CN Check	<code>disableCnCheck</code>	Boolean	Disable CN check for the server certificate. This should only be used when testing.	false	true
Coordinate Order	<code>coordinateOrder</code>	String	Coordinate order expected and returned by remote source	LON_LAT	true
Use posList in LinearRing	<code>usePosList</code>	Boolean	Use a <code><posList></code> element rather than a series of <code><pos></code> elements when issuing geospatial queries containing a LinearRing	false	false
Metocard Mappings	<code>metocardMappings</code>	String	Mapping of the Metocard Attribute names to their CSW property names. The format should be <code>'title=dc:title'</code> .	effective=created,created=dateSubmitted,modified=modified,thumbnail=references,content-type=type,id=identifier,resource-uri=source	false
Poll Interval	<code>pollInterval</code>	Integer	Poll Interval to Check if the Source is available (in minutes - minimum 1).	5	true
Connection Timeout	<code>connectionTimeout</code>	Integer	Amount of time to attempt to establish a connection before timing out, in milliseconds.	30000	true
Receive Timeout	<code>receiveTimeout</code>	Integer	Amount of time to wait for a response before timing out, in milliseconds.	60000	true
Output Schema	<code>outputSchema</code>	String	Output Schema	urn:catalog:metocard	true
Query Type Name	<code>queryTypeName</code>	String	Qualified Name for the Query Type used in the CSW GetRecords request	csw:Record	true
Query Type Namespace	<code>queryTypeNamespace</code>	String	Namespace for the Query Type used in the CSW GetRecords request	http://www.opengis.net/cat/csw/2.0.2	true
Force CQL Text	<code>isCqlForce</code>	Boolean	Force CQL Text as the Query Language	false	true
Forced Spatial Filter Type	<code>forceSpatialFilter</code>	String	Force only the selected Spatial Filter Type as the only available Spatial Filter.	NO_FILTER	false

Name	Id	Type	Description	Default Value	Required
Security Attributes	<code>securityAttributeStrings</code>	String	Security attributes for this source		true

Table 174. GeoCoder Plugin

Title	Property	Type	Description	Default Value	Required	Radius
-------	----------	------	-------------	---------------	----------	--------

Table 175. GMD CSW ISO Federated Source

Name	Id	Type	Description	Default Value	Required
Source ID	<code>id</code>	String	The unique name of the Source		true
CSW URL	<code>cswUrl</code>	String	URL to the endpoint implementing the Catalogue Service for Web (CSW) spec		true
Username	<code>username</code>	String	Username for CSW Service (optional)		false
Password	<code>password</code>	Password	Password for CSW Service (optional)		false
Disable CN Check	<code>disableCnCheck</code>	Boolean	Disable CN check for the server certificate. This should only be used when testing.	false	true
Coordinate Order	<code>coordinateOrder</code>	String	Coordinate order expected and returned by remote source	LON_LAT	true
Use posList in LinearRing	<code>usePosList</code>	Boolean	Use a <posList> element rather than a series of <pos> elements when issuing geospatial queries containing a LinearRing	false	false

Name	Id	Type	Description	Default Value	Required
Metocard Mappings	metocardMappings	String	Mapping of the Metocard Attribute names to their CSW property names. The format should be 'title=dc:title'.	id=apiso:Id entifier,eff ective=apis o:Publicati onDate,cre ated=apiso: CreationDa te,modified =apiso:Revi sionDate,tit le=apiso:Al ternateTitl e,AnyText=ap iso:AnyText,ows :Bou ndingBox=ap iso:Bou ndingBox,la nguage=ap iso:Lang age,langua ge=apiso:R esourceLa nguage,dat atype=apis o:Type,des cription=ap iso:Abstra ct,contact.p oint-of- contact- name=apis o:Organisa tionName,topic .keywo rd=apiso:S ubject,med ia.format=ap iso:Form at,modified =apiso:Mo dified	false
Poll Interval	pollInterval	Integer	Poll Interval to Check if the Source is available (in minutes - minimum 1).	5	true
Connection Timeout	connectionTimeout	Integer	Amount of time to attempt to establish a connection before timing out, in milliseconds.	30000	true

Name	Id	Type	Description	Default Value	Required
Receive Timeout	<code>receiveTimeout</code>	Integer	Amount of time to wait for a response before timing out, in milliseconds.	60000	true
Output Schema	<code>outputSchema</code>	String	Output Schema	http://www.isotc211.org/2005/gmd	true
Query Type Name	<code>queryTypeName</code>	String	Qualified Name for the Query Type used in the CSW GetRecords request	gmd:MD_Metadata	true
Query Type Namespace	<code>queryTypeNamespace</code>	String	Namespace for the Query Type used in the CSW GetRecords request	http://www.isotc211.org/2005/gmd	true
Force CQL Text	<code>isCqlForce</code>	Boolean	Force CQL Text as the Query Language	false	true
Forced Spatial Filter Type	<code>forceSpatialFilter</code>	String	Force only the selected Spatial Filter Type as the only available Spatial Filter.	NO_FILTER	false
Security Attributes	<code>securityAttributeStrings</code>	String	Security attributes for this source		true

Table 176. Spatial KML Endpoint

Name	Id	Type	Description	Default Value	Required
Style Document	<code>styleUrl</code>	String	KML Document containing custom styling. This will be served up by the KmlEndpoint. (e.g. file:///path/to/kml/style/doc.kml)		false
Icons Location	<code>iconLoc</code>	String	Location of icons for the KML endpoint		false
Description	<code>description</code>	String	Description of this NetworkLink. Enter a short description of what this NetworkLink provides.		false
Web Site	<code>webSite</code>	String	URL of the web site to be displayed in the description.		false
Logo	<code>logo</code>	String	URL to the logo to be displayed in the description.		false
Visible By Default	<code>visibleByDefault</code>	Boolean	Check if the source NetworkLinks should be visible by default.	false	false
Max Number of Results	<code>maxResults</code>	Integer	The maximum number of results that should be returned from each layer.	100	false

Table 177. Metocard to WFS Feature Map

Name	Id	Type	Description	Default Value	Required
Feature Type	<code>featureType</code>	String	Feature Type. Format is {URI}local-name		true
Metocard Title to WFS Feature Property Mapping	<code>titleMapping</code>	String	Metocard Title to WFS Feature Property Mapping		false
Metocard Created Date to WFS Feature Property Mapping	<code>createdDateMapping</code>	String	Metocard Created Date to WFS Feature Property Mapping		false
Metocard Modified Date to WFS Feature Property Mapping	<code>modifiedDateMapping</code>	String	Metocard Modified Date to WFS Feature Property Mapping		false
Metocard Effective Date to WFS Feature Property Mapping	<code>effectiveDateMapping</code>	String	Metocard Effective Date to WFS Feature Property Mapping		false
Metocard Expiration Date to WFS Feature Property Mapping	<code>expirationDateMapping</code>	String	Metocard Expiration Date to WFS Feature Property Mapping		false
Metocard Resource URI to WFS Feature Property Mapping	<code>resourceUriMapping</code>	String	Metocard Resource URI to WFS Feature Property Mapping		false

Name	Id	Type	Description	Default Value	Required
Metocard Resource Size to WFS Feature Property Mapping	<code>resourceSizeMapping</code>	String	Metocard Resource Size to WFS Feature Property Mapping		false
The Units of the Feature Property that corresponds to the Metocard Resource Size	<code>dataUnit</code>	String	The Units of the Feature Property that corresponds to the Metocard Resource Size	B	true
Metocard Thumbnail to WFS Feature Property Mapping	<code>thumbnailMapping</code>	String	Metocard Thumbnail to WFS Feature Property Mapping		false
Metocard Geography to WFS Feature Property Mapping	<code>geographyMapping</code>	String	Metocard Geography to WFS Feature Property Mapping		false
Temporal Sort By Feature Property	<code>sortByTemporalFeatureProperty</code>	String	When Sorting Temporally, Sort By This Feature Property.		false
Relevance Sort By Feature Property	<code>sortByRelevanceFeatureProperty</code>	String	When Sorting By Relevance, Sort By This Feature Property.		false
Distance Sort By Feature Property	<code>sortByDistanceFeatureProperty</code>	String	When Sorting By Distance, Sort By This Feature Property.		false

Table 178. WFS v1.0.0 Connected Source

Name	Id	Type	Description	Default Value	Required
Source ID	<code>id</code>	String	The unique name of the Source	WFS	true

Name	Id	Type	Description	Default Value	Required
WFS URL	<code>wfsUrl</code>	String	URL to the endpoint implementing the Web Feature Service (WFS) spec	null	true
Disable CN Check	<code>disableCnCheck</code>	Boolean	Disable CN check for the server certificate. This should only be used when testing.	false	true
Username	<code>username</code>	String	Username for WFS Service (optional)	null	false
Password	<code>password</code>	Password	Password for WFS Service (optional)	null	false
Non Queryable Properties	<code>nonQueryableProperties</code>	String	Properties listed here will NOT be queryable and any attempt to filter on these properties will result in an exception.	null	false
Poll Interval	<code>pollInterval</code>	Integer	Poll Interval to Check if the Source is available (in minutes - minimum 1).	5	true
Forced Spatial Filter Type	<code>forceSpatialFilter</code>	String	Force only the selected Spatial Filter Type as the only available Spatial Filter.	NO_FILTER	false
Connection Timeout	<code>connectionTimeout</code>	Integer	Amount of time to attempt to establish a connection before timing out, in milliseconds.	30000	true
Receive Timeout	<code>receiveTimeout</code>	Integer	Amount of time to wait for a response before timing out, in milliseconds.	60000	true

Table 179. WFS v1.0.0 Federated Source

Name	Id	Type	Description	Default Value	Required
Source ID	<code>id</code>	String	The unique name of the Source	WFS_v1_0_0	true
WFS URL	<code>wfsUrl</code>	String	URL to the endpoint implementing the Web Feature Service (WFS) spec	null	true
Disable CN Check	<code>disableCnCheck</code>	Boolean	Disable CN check for the server certificate. This should only be used when testing.	false	true
Username	<code>username</code>	String	Username for WFS Service (optional)	null	false
Password	<code>password</code>	Password	Password for WFS Service (optional)	null	false
Forced Feature Type	<code>forcedFeatureType</code>	String	Force only a specific FeatureType to be queried instead of all featureTypes	null	false

Name	Id	Type	Description	Default Value	Required
Non Queryable Properties	<code>nonQueryableProperties</code>	String	Properties listed here will NOT be queryable and any attempt to filter on these properties will result in an exception.	null	false
Poll Interval	<code>pollInterval</code>	Integer	Poll Interval to Check if the Source is available (in minutes - minimum 1).	5	true
Forced Spatial Filter Type	<code>forceSpatialFilter</code>	String	Force only the selected Spatial Filter Type as the only available Spatial Filter.	NO_FILTER	false
Connection Timeout	<code>connectionTimeout</code>	Integer	Amount of time to attempt to establish a connection before timing out, in milliseconds.	30000	true
Receive Timeout	<code>receiveTimeout</code>	Integer	Amount of time to wait for a response before timing out, in milliseconds.	60000	true

Table 180. WFS 2.0.0 Connected Source

Name	Id	Type	Description	Default Value	Required
Source ID	<code>id</code>	String	The unique name of the Source	WFS	true
WFS URL	<code>wfsUrl</code>	String	URL to the endpoint implementing the Web Feature Service (WFS) 2.0.0 spec	null	true
Disable CN Check	<code>disableCnCheck</code>	Boolean	Disable CN check for the server certificate. This should only be used when testing.	false	true
Force Longitude/ Latitude coordinate order	<code>isLonLatOrder</code>	Boolean	Force Longitude/Latitude coordinate order	false	true
Disable Sorting	<code>disableSorting</code>	Boolean	When selected, the system will not specify sort criteria with the query. This should only be used if the remote source is unable to handle sorting even when the capabilities states 'ImplementsSorting' is supported.	false	true
Username	<code>username</code>	String	Username for the WFS Service (optional)	null	false
Password	<code>password</code>	Password	Password for the WFS Service (optional)	null	false

Name	Id	Type	Description	Default Value	Required
Non Queryable Properties	nonQueryableProperties	String	Properties listed here will NOT be queryable and any attempt to filter on these properties will result in an exception.	null	false
Poll Interval	pollInterval	Integer	Poll Interval to Check if the Source is available (in minutes - minimum 1).	5	true
Forced Spatial Filter Type	forceSpatialFilter	String	Force only the selected Spatial Filter Type as the only available Spatial Filter.	NO_FILTER	false
Connection Timeout	connectionTimeout	Integer	Amount of time to attempt to establish a connection before timing out, in milliseconds.	30000	true
Receive Timeout	receiveTimeout	Integer	Amount of time to wait for a response before timing out, in milliseconds.	60000	true

Table 181. WFS 2.0.0 Federated Source

Name	Id	Type	Description	Default Value	Required
Source ID	id	String	The unique name of the Source	WFS_v2_0_0	true
WFS URL	wfsUrl	String	URL to the endpoint implementing the Web Feature Service (WFS) 2.0.0 spec	null	true
Disable CN Check	disableCnCheck	Boolean	Disable CN check for the server certificate. This should only be used when testing.	false	true
Coordinate Order	coordinateOrder	String	Coordinate order that remote source expects and returns spatial data in	LAT_LON	true
Forced Feature Type	forcedFeatureType	String	Force only a specific FeatureType to be queried instead of all featureTypes	null	false
Disable Sorting	disableSorting	Boolean	When selected, the system will not specify sort criteria with the query. This should only be used if the remote source is unable to handle sorting even when the capabilities states 'ImplementsSorting' is supported.	false	true
Username	username	String	Username for the WFS Service (optional)	null	false

Name	<code>Id</code>	Type	Description	Default Value	Required
Password	<code>password</code>	Password	Password for the WFS Service (optional)	null	false
Non Queryable Properties	<code>nonQueryableProperties</code>	String	Properties listed here will NOT be queryable and any attempt to filter on these properties will result in an exception.	null	false
Poll Interval	<code>pollInterval</code>	Integer	Poll Interval to Check if the Source is available (in minutes - minimum 1).	5	true
Forced Spatial Filter Type	<code>forceSpatialFilter</code>	String	Force only the selected Spatial Filter Type as the only available Spatial Filter.	NO_FILTER	false
Connection Timeout	<code>connectionTimeout</code>	Integer	Amount of time to attempt to establish a connection before timing out, in milliseconds.	30000	true
Receive Timeout	<code>receiveTimeout</code>	Integer	Amount of time to wait for a response before timing out, in milliseconds.	60000	true

Table 182. Spatial KML Style Map Entry

Name	<code>Id</code>	Type	Description	Default Value	Required
Attribute Name	<code>attributeName</code>	String	The name of the Metocard Attribute to match against. e.g. title, metadata-content-type, etc	null	true
Attribute Value	<code>attributeValue</code>	String	The value of the Metocard Attribute.	null	true
Style URL	<code>styleUrl</code>	String	The full qualified URL to the KML Style. e.g. http://example.com/styles#myStyle	null	true

A.12. Search UI Application Reference

The Standard Search UI is a user interface that enables users to search a catalog and associated sites for content and metadata.

A.12.1. Search UI Prerequisites

To use the Search UI application, the following applications/features must be installed:

- Platform
- Catalog

A.12.2. Installing Search UI

The Search UI application is installed by default with a standard installation.

A.12.3. Configuring the Search UI Application

To configure the Search UI Application:

1. Navigate to the Admin Console.
2. Select the **Search UI** application.
3. Select the **Configuration** tab.

Table 183. Search UI Available Configurations

Name	Property	Description
Search UI Redirect	org.codice.ddf.ui.searchui.filter.RedirectServlet	Search UI redirect.
Simple Search UI	org.codice.ddf.ui.search.simple.properties	Simple Search UI.
Search UI Endpoint	org.codice.ddf.ui.search.standard.endpoint	Search UI Endpoint.
Standard Search UI	org.codice.ddf.ui.search.standard.properties	Standard Search UI.

Table 184. Search UI Redirect

Name	Id	Type	Description	Default Value	Required
Redirect URI	defaultUri	String	Specifies the redirect URI to use when accessing the /search URI.	/search/catalog	true

Table 185. Simple Search UI

Name	Id	Type	Description	Default Value	Required
Header	header	String	Specifies the header text to be rendered on the generated Query Page		true
Footer	footer	String	Specifies the footer text to be rendered on the generated Query Page		true
Text Color	color	String	Specifies the Text Color of the Header and Footer. Use html css colors or #rrggbb.		true
Background Color	background	String	Specifies the Background Color of the Header and Footer. Use html css colors or #rrggbb.		true

Table 186. Search UI Endpoint

Name	Id	Type	Description	Default Value	Required
Disable Cache	<code>cacheDisabled</code>	Boolean	Disables use of cache.	false	false
Disable Normalization	<code>normalizationDisabled</code>	Boolean	Disables relevance and distance normalization.	false	false

Table 187. Standard Search UI

Name	Id	Type	Description	Default Value	Required
Result Count	<code>resultCount</code>	Integer	Specifies the number of results to request from each source (Max 1000)	250	true
Imagery Providers	<code>imageryProviders</code>	String	List of imagery providers to use. Valid types are: OSM (OpenStreetMap), AGM (ArcGisMap), BM (BingMap), WMS (WebMapService), WMT (WebMapTile), TMS (TileMapService), GE (GoogleEarth). Example: {"type" "WMS" "url" "http://example.com" "layers" ["layer1" "layer2"] "parameters" {"FORMAT" "image/png" "VERSION" "1.1.1"} "alpha" 0.5}		false
Terrain Provider	<code>terrainProvider</code>	String	Terrain provider to use for height data. Valid types are: CT (CesiumTerrain), AGS (ArcGisImageServer), VRW (VRTheWorld). Example: {"type" "CT" "url" "http://example.com"}	{"type" "CT" "url" "http://asse ts.agi.com/ stk-terrain/tile sets/world/tiles"}	false
Map Projection	<code>projection</code>	String	Projection of imagery providers	EPSG:4326	false
Bing Maps Key	<code>bingKey</code>	String	Bing Maps API key. This should only be set if you are using Bing Maps Imagery or Terrain Providers.		false
Connection Timeout	<code>timeout</code>	Integer	Specifies the client-side connection timeout in milliseconds.	15000	false
Help Location	<code>helpUrl</code>	String	URL to help documentation	help.html	false
Show Sign In	<code>signIn</code>	Boolean	Allow Sign In to Search UI and welcome notice. Enable this if the Search UI is protected.	true	false

Name	Id	Type	Description	Default Value	Required
Show Tasks	task	Boolean	Show task menu area for long running actions.	false	false
Show Gazetteer	gazetteer	Boolean	Show gazetteer for searching place names.	true	false
Show Uploader	ingest	Boolean	Show upload menu for adding new record.	true	false
Use External Authentication	externalAuthentication	Boolean	Use an external authentication point, such as IdP	false	false
Type Name Mapping	typeNameMapping	String	Mapping of display names to content types in the form name=type.		false

Appendix B: Application Whitelists

Within each DDF application, certain packages are exported for use by third parties.

B.1. Admin Whitelist

The following classes have been exported by the Admin application and are approved for use by third parties:

In package `org.codice.ddf.ui.admin.api.module`

- `AdminModule`

In package `org.codice.ddf.admin.application.service`

- `Application`
- `ApplicationNode`
- `ApplicationService`
- `ApplicationStatus`
- `ApplicationServiceException`

In package `org.codice.ddf.admin.application.plugin`

- `ApplicationPlugin`

B.2. Catalog Whitelist

The following classes have been exported by the Catalog application and are approved for use by third parties:

In package `ddf.catalog`

- [CatalogFramework](#)
- [Constants](#)

In package `ddf.catalog.cache`

- [ResourceCacheInterface](#) Deprecated

In package `ddf.catalog.data`

- [Attribute](#)
- [AttributeDescriptor](#)
- [AttributeType](#)
- [BinaryContent](#)
- [ContentType](#)
- [Metocard](#)
- [MetocardCreationException](#)
- [MetocardType](#)
- [MetocardTypeUnregistrationException](#)
- [Result](#)

In package `ddf.catalog.event`

- [DeliveryException](#)
- [DeliveryMethod](#)
- [EventException](#)
- [EventProcessor](#)
- [InvalidSubscriptionException](#)
- [Subscriber](#)
- [Subscription](#)
- [SubscriptionExistsException](#)
- [SubscriptionNotFoundException](#)

In package `ddf.catalog.federation`

- [Federatable](#)
- [FederationException](#)
- [FederationStrategy](#)

In package `ddf.catalog.filter`

- [AttributeBuilder](#)
- [BufferedSpatialExpressionBuilder](#)
- [ContextualExpressionBuilder](#)
- [EqualityExpressionBuilder](#)
- [ExpressionBuilder](#)
- [FilterAdapter](#)
- [FilterBuilder](#)

- [FilterDelegate](#)
- [NumericalExpressionBuilder](#)
- [NumericalRangeExpressionBuilder](#)
- [SpatialExpressionBuilder](#)
- [TemporalInstantExpressionBuilder](#)
- [TemporalRangeExpressionBuilder](#)
- [XPathBasicBuilder](#)
- [XPathBuilder](#)

In package `ddf.catalog.filter.delegate`

- [CopyFilterDelegate](#)
- [FilterToTextDelegate](#)

In package `ddf.catalog.operation`

- [CreateRequest](#)
- [CreateResponse](#)
- [DeleteRequest](#)
- [DeleteResponse](#)
- [Operation](#)
- [OperationTransaction](#)
- [Pingable](#)
- [ProcessingDetails](#)
- [Query](#)
- [QueryRequest](#)
- [QueryResponse](#)
- [Request](#)
- [ResourceRequest](#)
- [ResourceResponse](#)
- [Response](#)
- [SourceInfoRequest](#)
- [SourceInfoResponse](#)
- [SourceProcessingDetails](#)
- [SourceResponse](#)
- [Update](#)
- [UpdateRequest](#)
- [UpdateResponse](#)

In package `ddf.catalog.plugin`

- [AccessPlugin](#)
- [PluginExecutionException](#)
- [PolicyPlugin](#)
- [PolicyResponse](#)

- [PostFederatedQueryPlugin](#)
- [PostIngestPlugin](#)
- [PostQueryPlugin](#)
- [PostResourcePlugin](#)
- [PreDeliveryPlugin](#)
- [PreFederatedQueryPlugin](#)
- [PreIngestPlugin](#)
- [PreQueryPlugin](#)
- [PreResourcePlugin](#)
- [PreSubscriptionPlugin](#)
- [StopProcessingException](#)

In package `ddf.catalog.resource`

- [DataUsageLimitExceededException](#)
- [Resource](#)
- [ResourceNotFoundException](#)
- [ResourceNotSupportedException](#)
- [ResourceReader](#)
- [ResourceWriter](#)

In package `ddf.catalog.service`

- [ConfiguredService](#)

In package `ddf.catalog.source`

- [CatalogProvider](#)
- [ConnectedSource](#)
- [FederatedSource](#)
- [IngestException](#)
- [InternalIngestException](#)
- [RemoteSource](#)
- [Source](#)
- [SourceDescriptor](#)
- [SourceMonitor](#)
- [SourceUnavailableException](#)
- [UnsupportedQueryException](#)

In package `ddf.catalog.transform`

- [CatalogTransformerException](#)
- [InputCollectionTransformer](#)
- [InputTransformer](#)
- [MetacardTransformer](#)
- [QueryResponseTransformer](#)

In package `ddf.catalog.transformer.api`

- `MetocardMarshaller`
- `PrintWriter`
- `PrintWriterProvider`

In package `ddf.catalog.util`

- `Describable` Deprecated
- `Maskable`

In package `ddf.catalog.validation`

- `MetocardValidator`
- `ValidationException`

In package `ddf.geo.formatter`

- `CompositeGeometry`
- `GeometryCollection`
- `LineString`
- `MultiLineString`
- `MultiPoint`
- `MultiPolygon`
- `Point`
- `Polygon`

In package `ddf.util`

- `InetAddressUtil`
- `NamespaceMapImpl`
- `NamespaceResolver`
- `WktStandard`
- `XPathCache`
- `XPathHelper`
- `XSLTUtil`

B.3. Platform Whitelist

The following classes have been exported by the Platform application and are approved for use by third parties:

In package `ddf.action`

- `Action`
- `ActionProvider`
- `ActionRegistry`

In package `org.codice.ddf.branding`

- [BrandingPlugin](#)
- [BrandingRegistry](#)

In package `org.codice.ddf.configuration`

- [ConfigurationWatcher](#) Deprecated

B.4. Registry Whitelist

The following classes have been exported by the Registry Application and are approved for use by third parties:

None.

B.5. Security Whitelist

The following classes have been exported by the Security application and are approved for use by third parties:

In package `ddf.security`

- [SecurityConstants](#)
- [Subject](#)

In package `ddf.security.assertion`

- [SecurityAssertion](#)

In package `ddf.security.common.util`

- [Security](#) Deprecated
- [SecurityProperties](#)
- [ServiceComparator](#)
- [SortedServiceList](#)

In package `ddf.security.encryption`

- [EncryptionService](#)

In package `ddf.security.expansion`

- [Expansion](#)

In package `ddf.security.http`

- [SessionFactory](#)

In package `ddf.security.service`

- [SecurityManager](#)
- [SecurityServiceException](#)
- [TokenRequestHandler](#)

In package `ddf.security.sts.client.configuration`

- `STSClientConfiguration`

In package `ddf.security.ws.policy`

- `AbstractOverrideInterceptor`
- `PolicyLoader`

In package `ddf.security.ws.proxy`

- `ProxyServiceFactory`

In package `org.codice.ddf.security.handler.api`

- `AuthenticationHandler`

In package `org.codice.ddf.security.policy.context.attributes`

- `ContextAttributeMapping`

In package `org.codice.ddf.security.policy.context`

- `ContextPolicy`
- `ContextPolicyManager`

B.6. Solr Catalog Whitelist

The following classes have been exported by the Solr Catalog application and are approved for use by third parties:

None.

B.7. Spatial Whitelist

The following classes have been exported by the Spatial application and are approved for use by third parties:

In package `org.codice.ddf.spatial.kml.transformer`

- `KmlStyleMapEntry`
- `KMLTransformer`

In package `org.codice.ddf.spatial.ogc.catalog`

- `MetadataTransformer`

In package `org.codice.ddf.spatial.ogc.wfs.catalog.converter`

- `FeatureConverter`
- `FeatureConverterFactory`

In package `org.codice.ddf.spatial.ogc.wfs.catalog.mapper`

- [MetacardMapper](#)

In package [org.codice.ddf.spatial.ogc.wfs.v1_0_0.catalog.converter](#)

- [FeatureConverterFactory](#)

In package [org.codice.ddf.spatial.ogc.wfs.v2_0_0.catalog.converter](#)

- [FeatureConverterFactory](#)

B.8. Search UI Whitelist

The following classes have been exported by the Search UI application and are approved for use by third parties:

None.

B.9. Packages Removed From Whitelist

In the transition of the whitelist from the ambiguous package listing to the new class listing several errors were found. The packages originally listed that were removed either did not exist, contained experimental interfaces, or contained only internal implementations and should have never been included in the whitelist. The following is a list of packages that were listed in error and have been removed from the whitelist.

NOTE

None of the packages in this list have been removed from the distribution. They may however be changed or removed in the future.

Admin

- [org.codice.ddf.ui.admin.api.plugin](#)
- [org.codice.ddf.admin.configuration.plugin](#)

Catalog

- [org.codice.ddf.admin.configuration.plugin](#)
- [ddf.catalog.data.metacardtype](#)
- [ddf.catalog.federation.impl](#)
- [ddf.catalog.plugin.groomer](#)
- [ddf.catalog.pubsub](#)
- [ddf.catalog.pubsub.tracker](#)
- [ddf.catalog.resource.data](#)
- [ddf.catalog.resource.impl](#)
- [ddf.catalog.resourceretriever](#)
- [ddf.catalog.transformer.metacard.geojson](#)
- [ddf.common](#)
- [org.codice.ddf.endpoints](#)
- [org.codice.ddf.endpoints.rest](#)

- [org.codice.ddf.endpoints.rest.action](#)
- [org.codice.ddf.opensearch.query](#)
- [org.codice.ddf.opensearch.query.filter](#)

Platform

- [org.codice.ddf.configuration.admin](#)
- [org.codice.ddf.configuration.migration](#)
- [org.codice.ddf.configuration.persistence](#)
- [org.codice.ddf.configuration.persistence.felix](#)
- [org.codice.ddf.configuration.status](#)
- [org.codice.ddf.parser](#)
- [org.codice.ddf.parser.xml](#)
- [org.codice.ddf.platform.error.handler](#)
- [org.codice.ddf.platform.util](#)

Security

- [ddf.security.assertion.impl](#)
- [ddf.security.common.audit](#)
- [ddf.security.http.impl](#)
- [ddf.security.impl](#)
- [ddf.security.pdp.realm](#)
- [ddf.security.permission](#)
- [ddf.security.principal](#)
- [ddf.security.realm.sts](#)
- [ddf.security.samlp.impl](#)
- [ddf.security.service.impl](#)
- [ddf.security.settings](#)
- [ddf.security.soap.impl](#)
- [ddf.security.sts](#)
- [ddf.security.ws.policy.impl](#)
- [org.codice.ddf.security.certificate.generator](#)
- [org.codice.ddf.security.certificate.keystore.editor](#)
- [org.codice.ddf.security.common](#)
- [org.codice.ddf.security.filter.authorization](#)
- [org.codice.ddf.security.filter.login](#)
- [org.codice.ddf.security.filter.websso](#)
- [org.codice.ddf.security.handler.basic](#)
- [org.codice.ddf.security.handler.guest.configuration](#)
- [org.codice.ddf.security.handler.guest](#)
- [org.codice.ddf.security.handler.pki](#)
- [org.codice.ddf.security.handler.saml](#)
- [org.codice.ddf.security.interceptor](#)

- `org.codice.ddf.security.interceptor`
- `org.codice.ddf.security.policy.context.impl`
- `org.codice.ddf.security.servlet.logout`
- `org.codice.ddf.security.validator.username`

Spatial

- `org.codice.ddf.spatial.geocoder`
- `org.codice.ddf.spatial.geocoder.geonames`
- `org.codice.ddf.spatial.geocoding`
- `org.codice.ddf.spatial.geocoding.context`
- `org.codice.ddf.spatial.kml.endpoint`
- `org.codice.ddf.spatial.ogc.catalog.resource.impl`

Appendix C: Metadata Attributes

DDF extracts basic metadata from the resources ingested. Many file types contain additional [format-specific metadata attributes](#). A neutral [Catalog Taxonomy](#) enables transformation of metadata to other formats. See also a [list of all formats supported](#) for ingest.

C.1. Format-specific Attribute Mappings

Many file formats support additional metadata attributes that DDF is able to extract and make discoverable.

C.1.1. Mp4 Additional Attribute

Mp4 files have an additional attribute:

- `ext.mp4.audio-sample-rate`

C.2. Catalog Taxonomy

To facilitate data sharing while maximizing the usefulness of metadata, the attributes on resources are normalized into a common taxonomy that maps to attributes in the desired output format.

NOTE The taxonomy is presented here for reference only.

C.2.1. Core Attributes Table

Table 188. Core Attributes

Term	Definition	Datatype	Constraints	Example Value
title	A name for the resource. Dublin Core elements-title .	String	< 1024 characters	

Term	Definition	Datatype	Constraints	Example Value
source-id	ID of the source where the Metacard is cataloged. While this cannot be moved or renamed for legacy reasons, it should be treated as non-mappable, since this field is overwritten by the system when federated results are retrieved.	String	< 1024 characters	
metadata-content-type [deprecated] see Media Attributes	Content type of the resource.	String	< 1024 characters	
metadata-content-type-version [deprecated] see Media Attributes	Version of the metadata content type of the resource.	String	< 1024 characters	
metadata-target-namespace [deprecated] see Media Attributes	Target namespace of the metadata.	String	< 1024 characters	
metadata	Additional XML metadata describing the resource.	XML	A valid XML string per RFC 4825 (must be well-formed but not necessarily schema-compliant).	
location	The primary geospatial location of the resource.	Geometry	Valid Well Known Text (WKT) per http://www.opengis.net/standards/wktcrs <i>Coordinates must be in lon-lat coordinate order</i>	POINT(150 30)
expiration	The expiration date of the resource.	Date		

Term	Definition	Datatype	Constraints	Example Value
effective [deprecated]	The effective time of the event or resource represented by the metocard. Deprecated in favor of [created] and [modified] .	Date		
point-of-contact [deprecated]	The name of the point of contact for the resource. This is set internally to the user's subject and should be considered read-only to other DDF components.	String	< 1024 characters	
resource-uri	Location of the resource for the metocard.	String	Valid URI per RFC 2396	
resource-download-url	URL location of the resource for the metocard. This attribute provides a resolvable URL to the download location of the resource.	String	Valid URL per RFC 2396	
resource-size	Size in bytes of resource.	String	Although this type cannot be changed for legacy reasons, its value should always be a parsable whole number.	
thumbnail	The thumbnail for the resource in JPEG format.	Base 64 encoded binary string per RFC 4648	≤ 128 KB	
description	An account of the resource. Dublin Core elements-description .	String		
checksum	Checksum value for the primary resource for the metocard.	String	< 1024 characters	
checksum-algorithm	Algorithm used to calculate the checksum on the primary resource of the metocard.	String	< 1024 characters	
created	The creation date of the resource Dublin Core terms-created .	Date		
modified	The modification date of the resource Dublin Core terms-modified .	Date		

Term	Definition	Datatype	Constraints	Example Value
language	The language(s) of the resource. Dublin Core language .	List of Strings	Alpha-3 language code(s) per ISO_639-2	
resource.derived-download-url	Download URL(s) for accessing the derived formats for the metocard resource.	List of Strings	Valid URL(s) per RFC 2396	
resource.derived-uri	Location(s) for accessing the derived formats for the metocard resource.	List of Strings	Valid URI per RFC 2396	
datatype	The generic type(s) of the resource including the Dublin Core terms-type . DCMI Type term labels are expected here as opposed to term names.	List of Strings	Collection , Dataset , Event , Image , Interactive Resource , Moving Image , Physical Object , Service , Software , Sound , Still Image , and/or Text	

C.2.2. Associations Attributes Table

Table 189. Associations: Attributes in this group represent associations between products.

Term	Definition	Datatype	Constraints	Example Value
metocard.associations.derived	ID of one or more metacards derived from this metocard.	List of Strings	A valid metocard ID (conventionally, a type 4 random UUID with hyphens removed).	70809f17782c42b8ba15747b86b50ebf
metocard.associations.related	ID of one or more metacards related to this metocard.	List of Strings	A valid metocard ID (conventionally, a type 4 random UUID with hyphens removed).	70809f17782c42b8ba15747b86b50ebf
associations.external	One or more URI's identifying external associated resources.	List of Strings	A valid URI.	https://infocorp.org/wikia/reference

C.2.3. Contact Attributes Table

Table 190. Contact: Attributes in this group reflect metadata about different kinds of people/groups/units/organizations that can be associated with a metocard.

Term	Definition	Datatype	Constraints	Example Value
contact.creator-name	The name(s) of this metocard's creator(s).	List of Strings	< 1024 characters per entry	
contact.creator-address	The physical address(es) of this metocard's creator(s).	List of Strings	< 1024 characters per entry	
contact.creator-email	The email address(es) of this metocard's creator(s).	List of Strings	A valid email address per RFC 5322.	
contact.creator-phone	The phone number(s) of this metocard's creator(s).	List of Strings	< 1024 characters per entry	
contact.publisher-name	The name(s) of this metocard's publisher(s).	List of Strings	< 1024 characters per entry	
contact.publisher-address	The physical address(es) of this metocard's publisher(s).	List of Strings	< 1024 characters per entry	
contact.publisher-email	The email address(es) of this metocard's publisher(s).	List of Strings	A valid email address per RFC 5322.	
contact.publisher-phone	The phone number(s) of this metocard's publisher(s).	List of Strings	< 1024 characters per entry	
contact.contributor-name	The name of the contributor(s) to this metocard.	List of Strings	< 1024 characters per entry	
contact.contributor-address	The physical address(es) of the contributor(s) to this metocard.	List of Strings	< 1024 characters per entry	
contact.contributor-email	The email address(es) of the contributor(s) to this metocard.	List of Strings	A valid email address per RFC 5322.	
contact.contributor-phone	The phone number(s) of the contributor(s) to this metocard.	List of Strings	< 1024 characters per entry	
contact.point-of-contact-name	The name(s) of the point(s) of contact for this metocard.	List of Strings	< 1024 characters per entry	

Term	Definition	Datatype	Constraints	Example Value
contact.point-of-contact-address	The physical address(es) of a point(s) of contact for this metocard.	List of Strings	< 1024 characters per entry	
contact.point-of-contact-email	The email address(es) of the point(s) of contact for this metocard.	List of Strings	A valid email address per RFC 5322.	
contact.point-of-contact-phone	The phone number(s) of the point(s) of contact for this metocard.	List of Strings	< 1024 characters per entry	

C.2.4. DateTime Attributes Table

Table 191. *DateTime*: Attributes in this group reflect temporal aspects about the resource.

Term	Definition	Datatype	Constraints	Example Value
datetime.start	Start time(s) for the resource.	List of Dates		
datetime.end	End time(s) for the resource.	List of Dates		
datetime.name	A descriptive name for the corresponding temporal attributes. See datetime.start and datetime.end.	List of Strings	< 1024 characters per entry	

C.2.5. History Attributes Table

Table 192. *History*: Attributes in this group describe the history/versioning of the metocard.

Term	Definition	Datatype	Constraints	Example Value
metocard.version.id	Internal attribute identifier for which metocard this version is representing	String	A valid metocard ID (conventionally, a type 4 random UUID with hyphens removed).	70809f17782c42b8ba15747b86b50ebf
metocard.version.edited-by	Internal attribute identifying the editor of a history metocard.	String	A valid email address per RFC 5322	
metocard.version.versioned-on	Internal attribute for the versioned date of a metocard version.	Date		

Term	Definition	Datatype	Constraints	Example Value
metocard.version.action	Internal attribute for the action associated with a history metocard.	String	One of Deleted , Deleted-Content , Versioned , Versioned-Content	
metocard.version.tags	Internal attribute for the tags that were on the original metocard.	String		
metocard.version.type	Internal attribute for the metocard type of the original metocard.	String		
metocard.version.type-binary	Internal attribute for the serialized metocard type of the original metocard.	Binary		
metocard.version.resource-uri	Internal attribute for the original resource uri.	URI		

C.2.6. Location Attributes Table

Table 193. Location: Attributes in this group reflect location aspects about the resource.

Term	Definition	Datatype	Constraints	Example Value
location.altitude-meters	Altitude of the resource in meters.	List of Doubles	> 0	
location.country-code	One or more country codes associated with the resource.	List of Strings	ISO_3166-1 alpha-3 codes	
location.crs-code	Coordinate reference system code of the resource.	List of Strings	< 1024 characters per entry	EPSG:4326
location.crs-name	Coordinate reference system name of the resource.	List of Strings	< 1024 characters per entry	WGS 84

C.2.7. Media Attributes Table

Table 194. Media: Attributes in this group reflect metadata about media in general.

Term	Definition	Datatype	Constraints	Example Value
media.format	The file format, physical medium, or dimensions of the resource. Dublin Core elements-format	String	< 1024 characters	txt, docx, xml - typically the extension or a more complete name for such, note that this is not the mime type
media.format-version	The file format version of the resource. Note that the syntax can vary widely from format to format.	String	< 1024 characters	POSIX, 2016, 1.0
media.bit-rate	The bit rate of the media, in bits per second.	Double		
media.frame-rate	The frame rate of the video, in frames per second.	Double		
media.frame-center	The center of the video frame.	Geometry	Valid Well Known Text (WKT)	
media.height-pixels	The height of the media resource in pixels.	Integer		
media.width-pixels	The width of the media resource in pixels.	Integer		
media.compression	The type of compression this media uses. EXIF STANAG 4559 NC, NM, C1, M1, I1, C3, M3, C4, M4, C5, M5, C8, M8	String	One of the values defined for EXIF Compression tag.	
media.bits-per-sample	The number of bits per image component.	Integer		
media.type (RFC 2046)	A two-part identifier for file formats and format content.	String	A valid mime-type per https://www.ietf.org/rfc/rfc2046.txt	application/json
media.encoding	The encoding format of the media.	List of Strings	< 1024 characters per entry	MPEG-2, RGB

Term	Definition	Datatype	Constraints	Example Value
media.number-of-bands	The number of spectral bands in the media.	Integer	The significance of this number is instrumentation-specific, but there are eight commonly recognized bands. https://en.wikipedia.org/wiki/Multispectral_image	
media.scanning-mode (MPEG2)	Indicate if progressive or interlaced scans are being applied.	String	PROGRESSIVE, INTERLACED	

C.2.8. Metocard Attributes Table

Table 195. Metocard: Attributes in this group describe the metocard itself.

Term	Definition	Datatype	Constraints	Example Value
metocard.created	The creation date of the metocard.	Date		
metocard.modified	The modified date of the metocard.	Date		
metocard.owner	The email address of the metocard owner.	String	A valid email address per RFC 5322	
metocard-tags	Collections of data that go together, used for filtering query results. NOTE: these are system tags. For descriptive tags, Topic Attributes .	List of Strings	< 1024 characters per entry	

C.2.9. Security Attributes Table

Table 196. Security: Attributes in this group relate to security of the resource and metadata.

Term	Definition	Datatype	Constraints	Example Value
security.access-groups	Attribute name for storing groups to enforce access controls upon.	List of Strings	< 1024 characters per entry	

Term	Definition	Datatype	Constraints	Example Value
security.access-individuals	Attribute name for storing the email addresses of users to enforce access controls upon.	List of Strings	A valid email address per RFC 5322.	

C.2.10. Topic Attributes Table

Table 197. Topic: Attributes in this group describe the topic of the resource.

Term	Definition	Datatype	Constraints	Example Value
topic.category	A category code from a given vocabulary.	List of Strings	A valid entry from the corresponding controlled vocabulary.	
topic.keyword	One or more keywords describing the subject matter of the metocard or resource.	List of Strings	< 1024 characters per entry	
topic.vocabulary	An identifier of a controlled vocabulary from which the topic category is derived.	List of Strings	Valid URI per RFC 2396.	

C.2.11. Validation Attributes Table

Table 198. Validation: Attributes in this group identify validation issues with the metocard and/or resource.

Term	Definition	Datatype	Constraints	Example Value
validation-warnings	Textual description of validation warnings on the resource.	List of Strings	< 1024 characters per entry	
validation-errors	Textual description of validation errors on the resource.	List of Strings	< 1024 characters per entry	

C.3. All Formats Supported

Supported File Types

Using the various input transformers, DDF supports ingest of the following MIME types.

Table 199. Application File Types

activemessage	andrew-inset	applefile
applixware	atom+xml	atomcat+xml
atomicmail	atomsrv+xml	auth-policy+xml
batch-smtp	beep+xml	bizagi-modeler

cals-1840	cbor	ccxml+xml
cea-2018+xml	cellml+xml	cnp+xml
commonground	conference-info+xml	cpl+xml
csta+xml	cstadata+xml	cu-seeme
cybercash	davmount+xml	dca-rft
dec-dx	dialog-info+xml	dicom
dif+xml	dita+xml	dita+xml
dita+xml	dita+xml	dita+xml
dita+xml	dns	dvc
ecmascript	edi-consent	edi-x12
edifact	emma+xml	epp+xml
epub+zip	eshop	example
fastinfoset	fastsoap	fits
font-tdpfr	gzip	h224
http	hyperstudio	ibe-key-request+xml
ibe-pkg-reply+xml	ibe-pp-data	iges
illustrator	im-iscomposing+xml	index
index.cmd	index.obj	index.response
index.vnd	inf	iotp
ipp	isup	java-archive
java-serialized-object	java-vm	javascript
json	kate	kpml-request+xml
kpml-response+xml	lost+xml	mac-binhex40
mac-compactpro	macwriteii	marc
mathematica	mathml+xml	mbms-associated-procedure-description+xml
mbms-deregister+xml	mbms-envelope+xml	mbms-msk+xml
mbms-msk-response+xml	mbms-protection-description+xml	mbms-reception-report+xml
mbms-register+xml	mbms-register-response+xml	mbms-user-service-description+xml
mbox	media_control+xml	mediaservercontrol+xml
mikey	moss-keys	moss-signature
mosskey-data	mosskey-request	mp4
mpeg4-generic	mpeg4-iod	mpeg4-iod-xmt
msword	msword2	msword5
mxf	nasdata	news-checkgroups
news-groupinfo	news-transmission	nss
ocsp-request	ocsp-response	octet-stream
oda	oebps-package+xml	ogg
onenote	parityfec	patch-ops-error+xml
pdf	pgp-encrypted	pgp-keys
pgp-signature	pics-rules	pidf+xml
pidf-diff+xml	pkcs10	pkcs7-mime

pkcs7-signature	pkix-cert	pkix-crl
pkix-pkipath	pkixcmp	pls+xml
poc-settings+xml	postscript	prs.alvestrand.titrax-sheet
prs.cww	prs.nprend	prs.plucker
qsig	quicktime	rdf+xml
reginfo+xml	relax-ng-compact-syntax	remote-printing
resource-lists+xml	resource-lists-diff+xml	riscos
rlmi+xml	rls-services+xml	rsd+xml
rss+xml	rtf	rtx
samlassertion+xml	samlmetadata+xml	sbml+xml
scvp-cv-request	scvp-cv-response	scvp-vp-request
scvp-vp-response	sdp	sereal
sereal	sereal	sereal
set-payment	set-payment-initiation	set-registration
set-registration-initiation	sgml	sgml-open-catalog
shf+xml	sieve	simple-filter+xml
simple-message-summary	simplesymbolcontainer	slate
sldworks	smil+xml	soap+fastinfoset
soap+xml	sparql-query	sparql-results+xml
spirits-event+xml	srgs	srgs+xml
ssml+xml	timestamp-query	timestamp-reply
tve-trigger	ulpfec	vemmi
vividence.scriptfile	vnd.3gpp.bsf+xml	vnd.3gpp.pic-bw-large
vnd.3gpp.pic-bw-small	vnd.3gpp.pic-bw-var	vnd.3gpp.sms
vnd.3gpp2.bcmcsinfo+xml	vnd.3gpp2.sms	vnd.3gpp2.tcap
vnd.3m.post-it-notes	vnd.accpac.simply.aso	vnd.accpac.simply.imp
vnd.acucobol	vnd.acucorp	vnd.adobe.aftereffects.project
vnd.adobe.aftereffects.template	vnd.adobe.air-application-installer-package+zip	vnd.adobe.xdp+xml
vnd.adobe.xfdf	vnd.aether.imp	vnd.airzip.filesecure.azf
vnd.airzip.filesecure.azs	vnd.amazon.ebook	vnd.americandynamics.acc
vnd.amiga.ami	vnd.android.package-archive	vnd.anser-web-certificate-issue-initiation
vnd.anser-web-funds-transfer-initiation	vnd.antix.game-component	vnd.apple.installer+xml
vnd.apple.iwork	vnd.apple.keynote	vnd.apple.numbers
vnd.apple.pages	vnd.arastraw.swi	vnd.audiograph
vnd.autopackage	vnd.avistar+xml	vnd.blueice.multipass
vnd.bluetooth.ep.oob	vnd.bmi	vnd.businessobjects
vnd.cab-jscript	vnd.canon-cpd1	vnd.canon-lips
vnd.cendio.thinlinc.clientconf	vnd.chemdraw+xml	vnd.chipnuts.karaoke-mmd
vnd.cinderella	vnd.cirpack.isdn-ext	vnd.claymore
vnd.clonk.c4group	vnd.commerce-battelle	vnd.commonspace
vnd.contact.cmsg	vnd.cosmocaller	vnd.crick.clicker

vnd.crick.clicker.keyboard	vnd.crick.clicker.palette	vnd.crick.clicker.template
vnd.crick.clicker.wordbank	vnd.criticaltools.wbs+xml	vnd.ctc-posml
vnd.ctct.ws+xml	vnd.cups-pdf	vnd.cups-postscript
vnd.cups-ppd	vnd.cups-raster	vnd.cups-raw
vnd.curl.car	vnd.curl.pcurl	vnd.cybank
vnd.data-vision.rdz	vnd.denovo.fcslayout-link	vnd.dir-bi.plate-dl-nosuffix
vnd.dna	vnd.dolby.mlp	vnd.dolby.mobile.1
vnd.dolby.mobile.2	vnd.dpgraph	vnd.dreamfactory
vnd.dvb.esgcontainer	vnd.dvb.ipcdftnotifaccess	vnd.dvb.ipdcesgaccess
vnd.dvb.ipdcroaming	vnd.dvb.ffmpeg.alfec-base	vnd.dvb.ffmpeg.alfec-enhancement
vnd.dvb.notif-aggregate-root+xml	vnd.dvb.notif-container+xml	vnd.dvb.notif-generic+xml
vnd.dvb.notif-ia-msglist+xml	vnd.dvb.notif-ia-registration-request+xml	vnd.dvb.notif-ia-registration-response+xml
vnd.dvb.notif-init+xml	vnd.dxr	vnd.dynageo
vnd.ecdis-update	vnd.ecowin.chart	vnd.ecowin.filerequest
vnd.ecowin.fileupdate	vnd.ecowin.series	vnd.ecowin.seriesrequest
vnd.ecowin.seriesupdate	vnd.emclient.accessrequest+xml	vnd.enliven
vnd.epson.esf	vnd.epson.msf	vnd.epson.quickanime
vnd.epson.salt	vnd.epson.ssf	vnd.ericsson.quickcall
vnd.eszigno3+xml	vnd.etsi.aoc+xml	vnd.etsi.asic-e+zip
vnd.etsi.asic-s+zip	vnd.etsi.cug+xml	vnd.etsi.ffmpeg+xml
vnd.etsi.iptvdiscovery+xml	vnd.etsi.ffmpegprofile+xml	vnd.etsi.ffmpegsad-bc+xml
vnd.etsi.ffmpegsad-cod+xml	vnd.etsi.ffmpegsad-npvr+xml	vnd.etsi.ffmpegueprofile+xml
vnd.etsi.mcid+xml	vnd.etsi.sci+xml	vnd.etsi.simservs+xml
vnd.eudora.data	vnd.ezpix-album	vnd.ezpix-package
vnd.f-secure.mobile	vnd.fdf	vnd.fdsn.mseed
vnd.fdsn.seed	vnd.ffdns	vnd.fints
vnd.flographit	vnd.fluxtime.clip	vnd.font-fontforge-sfd
vnd.framemaker	vnd.frogans.fnc	vnd.frogans.ltf
vnd.fsc.weblaunch	vnd.fujitsu.oasys	vnd.fujitsu.oasys2
vnd.fujitsu.oasys3	vnd.fujitsu.oasysgp	vnd.fujitsu.oasysprs
vnd.fujixerox.art-ex	vnd.fujixerox.art4	vnd.fujixerox.ddd
vnd.fujixerox.docuworks	vnd.fujixerox.docuworks.binder	vnd.fujixerox.hbpl
vnd.fut-misnet	vnd.fuzzysheet	vnd.genomatix.tuxedo
vnd.geogebra.file	vnd.geogebra.tool	vnd.geometry-explorer
vnd.gmx	vnd.google-earth.kml+xml	vnd.google-earth.kmz
vnd.grafeq	vnd.gridmp	vnd.groove-account
vnd.groove-help	vnd.groove-identity-message	vnd.groove-injector
vnd.groove-tool-message	vnd.groove-tool-template	vnd.groove-vcard
vnd.handheld-entertainment+xml	vnd.hbci	vnd.hcl-bireports
vnd.hhe.lesson-player	vnd.hp-hpcl	vnd.hp-hpid
vnd.hp-hps	vnd.hp-jlyt	vnd.hp-pcl
vnd.hp-pclxl	vnd.httpphone	vnd.hydrostatix.sof-data

vnd.hzn-3d-crossword	vnd.ibm.afplinedata	vnd.ibm.electronic-media
vnd.ibm.minipay	vnd.ibm.modcap	vnd.ibm.rights-management
vnd.ibm.secure-container	vnd.iccprofile	vnd.igloader
vnd.immervision-ivp	vnd.immervision-ivu	vnd.informedcontrol.rms+xml
vnd.informix-visionary	vnd.intercon.formnet	vnd.intertrust.digibox
vnd.intertrust.nncp	vnd.intu.qbo	vnd.intu.qfx
vnd.iptc.g2.conceptitem+xml	vnd.iptc.g2.knowledgeitem+xml	vnd.iptc.g2.newsitem+xml
vnd.iptc.g2.packageitem+xml	vnd.ipunplugged.rcprofile	vnd.irepository.package+xml
vnd.is-xpr	vnd.jam	vnd.japannet-directory-service
vnd.japannet-jpnstore-wakeup	vnd.japannet-payment-wakeup	vnd.japannet-registration
vnd.japannet-registration-wakeup	vnd.japannet-setstore-wakeup	vnd.japannet-verification
vnd.japannet-verification-wakeup	vnd.jcp.javame.midlet-rms	vnd.jisp
vnd.joost.joda-archive	vnd.kahootz	vnd.kde.karbon
vnd.kde.kchart	vnd.kde.kformula	vnd.kde.kivio
vnd.kde.kontour	vnd.kde.kpresenter	vnd.kde.kspread
vnd.kde.kword	vnd.kenameaapp	vnd.kidspiration
vnd.kinar	vnd.koan	vnd.kodak-descriptor
vnd.liberty-request+xml	vnd.llamagraphics.life-balance.desktop	vnd.llamagraphics.life-balance.exchange+xml
vnd.lotus-1-2-3	vnd.lotus-approach	vnd.lotus-freelance
vnd.lotus-notes	vnd.lotus-organizer	vnd.lotus-screencam
vnd.lotus-wordpro	vnd.macports.portpkg	vnd.marlin drm.actiontoken+xml
vnd.marlin drm.conftoken+xml	vnd.marlin drm.license+xml	vnd.marlin drm.mDCF
vnd.mcd	vnd.medcalldata	vnd.mediastation.cdkey
vnd.meridian-slingshot	vnd.mfer	vnd.mfmp
vnd.micrografx.flo	vnd.micrografx.igx	vnd.mif
vnd.mindjet.mindmanager	vnd.minisoft-hp3000-save	vnd.mitsubishi.misty-guard.trustweb
vnd.mobius.daf	vnd.mobius.dis	vnd.mobius.mbk
vnd.mobius.mqy	vnd.mobius.msl	vnd.mobius.plc
vnd.mobius.txf	vnd.mophun.application	vnd.mophun.certificate
vnd.motorola.flexsuite	vnd.motorola.flexsuite.adsi	vnd.motorola.flexsuite.fis
vnd.motorola.flexsuite.gotap	vnd.motorola.flexsuite.kmr	vnd.motorola.flexsuite.ttc
vnd.motorola.flexsuite.wem	vnd.motorola.iprm	vnd.mozilla.xul+xml
vnd.ms-artgalry	vnd.ms-asf	vnd.ms-cab-compressed
vnd.ms-excel	vnd.ms-excel.addin.macroenabled.12	vnd.ms-excel.sheet.2
vnd.ms-excel.sheet.3	vnd.ms-excel.sheet.4	vnd.ms-excel.sheet.binary.macroenable d.12
vnd.ms-excel.sheet.macroenabled.12	vnd.ms-excel.template.macroenabled.12	vnd.ms-excel.workspace.3
vnd.ms-excel.workspace.4	vnd.ms-fontobject	vnd.ms-htmlhelp
vnd.ms-ims	vnd.ms-lrm	vnd.ms-outlook

vnd.ms-outlook-pst	vnd.ms-pki.seccat	vnd.ms-pki.stl
vnd.ms-playready.initiator+xml	vnd.ms-powerpoint	vnd.ms-powerpoint.addin.macroenabled.12
vnd.ms-powerpoint.presentation.macroenabled.12	vnd.ms-powerpoint.slide.macroenabled.12	vnd.ms-powerpoint.slideshow.macroenabled.12
vnd.ms-powerpoint.template.macroenabled.12	vnd.ms-project	vnd.ms-tnef
vnd.ms-visio.drawing	vnd.ms-visio.drawing.macroenabled.12	vnd.ms-visio.stencil
vnd.ms-visio.stencil.macroenabled.12	vnd.ms-visio.template	vnd.ms-visio.template.macroenabled.12
vnd.ms-visio.viewer	vnd.ms-wmdrm.lic-chlg-req	vnd.ms-wmdrm.lic-resp
vnd.ms-wmdrm.meter-chlg-req	vnd.ms-wmdrm.meter-resp	vnd.ms-word.document.macroenabled.12
vnd.ms-word.template.macroenabled.12	vnd.ms-works	vnd.ms-wpl
vnd.ms-xpsdocument	vnd.mseq	vnd.ms-sign
vnd.multiad.creator	vnd.multiad.creator.cif	vnd.music-niff
vnd.musician	vnd.muvee.style	vnd.ncd.control
vnd.ncd.reference	vnd.nervana	vnd.netfp
vnd.neurolanguage.nlu	vnd.noblenet-directory	vnd.noblenet-sealer
vnd.noblenet-web	vnd.nokia.catalogs	vnd.nokia.conml+wbxml
vnd.nokia.conml+xml	vnd.nokia.ipv.config+xml	vnd.nokia.isds-radio-presets
vnd.nokia.landmark+wbxml	vnd.nokia.landmark+xml	vnd.nokia.landmarkcollection+xml
vnd.nokia.n-gage.ac+xml	vnd.nokia.n-gage.data	vnd.nokia.n-gage.symbian.install
vnd.nokia.ncd	vnd.nokia.pcd+wbxml	vnd.nokia.pcd+xml
vnd.nokia.radio-preset	vnd.nokia.radio-presets	vnd.novadigm.edm
vnd.novadigm.edx	vnd.novadigm.ext	vnd.oasis.opendocument.chart
vnd.oasis.opendocument.chart-template	vnd.oasis.opendocument.database	vnd.oasis.opendocument.formula
vnd.oasis.opendocument.formula-template	vnd.oasis.opendocument.graphics	vnd.oasis.opendocument.graphics-template
vnd.oasis.opendocument.image	vnd.oasis.opendocument.image-template	vnd.oasis.opendocument.presentation
vnd.oasis.opendocument.presentation-template	vnd.oasis.opendocument.spreadsheet	vnd.oasis.opendocument.spreadsheet-template
vnd.oasis.opendocument.text	vnd.oasis.opendocument.text-master	vnd.oasis.opendocument.text-template
vnd.oasis.opendocument.text-web	vnd.obn	vnd.olpc-sugar
vnd.oma-scws-config	vnd.oma-scws-http-request	vnd.oma-scws-http-response
vnd.oma.bcast.associated-procedure-parameter+xml	vnd.oma.bcast.drm-trigger+xml	vnd.oma.bcast.imd+xml
vnd.oma.bcast.ltkm	vnd.oma.bcast.notification+xml	vnd.oma.bcast.provisioningtrigger

vnd.oma.bcast.sgboot	vnd.oma.bcast.sgdd+xml	vnd.oma.bcast.sgdu
vnd.oma.bcast.simple-symbol-container	vnd.oma.bcast.smartcard-trigger+xml	vnd.oma.bcast.sprov+xml
vnd.oma.bcast.stkm	vnd.oma.dcd	vnd.oma.dcdc
vnd.oma.dd2+xml	vnd.oma.drm.risd+xml	vnd.oma.group-usage-list+xml
vnd.oma.poc.detailed-progress-report+xml	vnd.oma.poc.final-report+xml	vnd.oma.poc.groups+xml
vnd.oma.poc.invocation-descriptor+xml	vnd.oma.poc.optimized-progress-report+xml	vnd.oma.xcap-directory+xml
vnd.omads-email+xml	vnd.omads-file+xml	vnd.omads-folder+xml
vnd.omaloc-supl-init	vnd.openofficeorg.extension	vnd.openxmlformats-officedocument.presentationml.presentation
vnd.openxmlformats-officedocument.presentationml.slide	vnd.openxmlformats-officedocument.presentationml.slideshow	vnd.openxmlformats-officedocument.presentationml.template
vnd.openxmlformats-officedocument.spreadsheetml.sheet	vnd.openxmlformats-officedocument.spreadsheetml.template	vnd.openxmlformats-officedocument.wordprocessingml.document
vnd.openxmlformats-officedocument.wordprocessingml.template	vnd.osa.netdeploy	vnd.osgi.bundle
vnd.osgi.dp	vnd.otp.sct-kip+xml	vnd.palm
vnd.paos.xml	vnd.pg.format	vnd.pg.osasli
vnd.piaccess.application-licence	vnd.picsel	vnd.poc.group-advertisement+xml
vnd.pocketlearn	vnd.powerbuilder6	vnd.powerbuilder6-s
vnd.powerbuilder7	vnd.powerbuilder7-s	vnd.powerbuilder75
vnd.powerbuilder75-s	vnd.preminet	vnd.previewsystems.box
vnd.proteus.magazine	vnd.publishare-delta-tree	vnd.pvi.ptid1
vnd.pwg-multiplexed	vnd.pwg-xhtml-print+xml	vnd.qualcomm.brew-app-res
vnd.quark.quarkxpress	vnd.rapid	vnd.recordare.musicxml
vnd.recordare.musicxml+xml	vnd.renlearn.rlprint	vnd.rim.cod
vnd.rn-realmedia	vnd.route66.link66+xml	vnd.ruckus.download
vnd.s3sms	vnd.sbm.cid	vnd.sbm.mid2
vnd.scribus	vndsealed.3df	vndsealed.csf
vndsealed.doc	vndsealed.eml	vndsealed.mht
vndsealed.net	vndsealed.ppt	vndsealed.tiff
vndsealed.xls	vndsealedmedia.softseal.html	vndsealedmedia.softseal.pdf
vndseemail	vndsema	vndsemd
vndsemf	vndshana.informed.formdata	vndshana.informed.formtemplate
vndshana.informed.interchange	vndshana.informed.package	vndsimtech-mindmapper
vndsmaf	vndsmart.teacher	vndsoftware602.filler.form+xml
vndsoftware602.filler.form-xml-zip	vndsolent.sdkm+xml	vndspotfire.dxp
vndspotfire.sfs	vndsss-cod	vndsss-dtf

vnd.sss-ntf	vnd.stardivision.calc	vnd.stardivision.draw
vnd.stardivision.impress	vnd.stardivision.math	vnd.stardivision.writer
vnd.stardivision.writer-global	vnd.street-stream	vnd.sun.wadl+xml
vnd.sun.xml.calc	vnd.sun.xml.calc.template	vnd.sun.xml.draw
vnd.sun.xml.draw.template	vnd.sun.xml.impress	vnd.sun.xml.impress.template
vnd.sun.xml.math	vnd.sun.xml.writer	vnd.sun.xml.writer.global
vnd.sun.xml.writer.template	vnd.sus-calendar	vnd.svd
vnd.swiftview-ics	vnd.symbian.install	vnd.syncml+xml
vnd.syncml.dm+wbxml	vnd.syncml.dm+xml	vnd.syncml.dm.notification
vnd.syncml.ds.notification	vnd.tao.intent-module-archive	vnd.tcpdump.pcap
vnd.tmobile-livetv	vnd.trid.tpt	vnd.triscape.mxs
vnd.trueapp	vnd.truedoc	vnd.ufdl
vnd.uiq.theme	vnd.umajin	vnd.unity
vnd.uoml+xml	vnd.uplanet.alert	vnd.uplanet.alert-wbxml
vnd.uplanet.bearer-choice	vnd.uplanet.bearer-choice-wbxml	vnd.uplanet.cacheop
vnd.uplanet.cacheop-wbxml	vnd.uplanet.channel	vnd.uplanet.channel-wbxml
vnd.uplanet.list	vnd.uplanet.list-wbxml	vnd.uplanet.listcmd
vnd.uplanet.listcmd-wbxml	vnd.uplanet.signal	vnd.vcx
vnd.vd-study	vnd.vectorworks	vnd.vidsoft.vidconference
vnd.visio	vnd.visionary	vnd.vividence.scriptfile
vnd.vsf	vnd.wap.sic	vnd.wap.slc
vnd.wap.wbxml	vnd.wap.wmlc	vnd.wap.wmlscriptc
vnd.webturbo	vnd.wfa.wsc	vnd.wmc
vnd.wmf.bootstrap	vnd.wordperfect	vnd.wqd
vnd.wrq-hp3000-labelled	vnd.wt.stf	vnd.wv.csp+wbxml
vnd.wv.csp+xml	vnd.wv.ssp+xml	vnd.xara
vnd.xfdl	vnd.xfdl.webform	vnd.xmi+xml
vnd.xmpie.cpkg	vnd.xmpie.dpkg	vnd.xmpie.plan
vnd.xmpie.ppkg	vnd.xmpie.xlim	vnd.yamaha.hv-dic
vnd.yamaha.hv-script	vnd.yamaha.hv-voice	vnd.yamaha.openscoreformat
vnd.yamaha.openscoreformat.osf pvg+xml	vnd.yamaha.smaf-audio	vnd.yamaha.smaf-phrase
vnd.yellowriver-custom-menu	vnd.zul	vnd.zzazz.deck+xml
voicexml+xml	watcherinfo+xml	whoispp-query
whoispp-response	winhlp	wita
wordperfect5.1	wsdl+xml	wspolicy+xml
x-123	x-7z-compressed	x-abiword
x-ace-compressed	x-adobe-indesign	x-adobe-indesign-interchange
x-apple-diskimage	x-appleworks	x-archive
x-arj	x-authorware-bin	x-authorware-map
x-authorware-seg	x-axcrypt	x-bcpio
x-berkeley-db	x-berkeley-db	x-berkeley-db
x-berkeley-db	x-berkeley-db	x-berkeley-db

x-berkeley-db	x-berkeley-db	x-berkeley-db
x-berkeley-db	x-berkeley-db	x-berkeley-db
x-bibtex-text-file	x-bittorrent	x-bplist
x-bzip	x-bzip2	x-cdlink
x-chat	x-chess-pgn	x-chrome-package
x-compress	x-coredump	x-corelpresentations
x-cpio	x-csh	x-debian-package
x-dex	x-director	x-doom
x-dosexec	x-dtbncx+xml	x-dtbook+xml
x-dtbresource+xml	x-dvi	x-elc
x-elf	x-emf	x-erdas-hfa
x-executable	x-fictionbook+xml	x-filemaker
x-font-adobe-metric	x-font-bdf	x-font-dos
x-font-framemaker	x-font-ghostscript	x-font-libgrx
x-font-linux-psf	x-font-otf	x-font-pcf
x-font-printer-metric	x-font-snf	x-font-speedo
x-font-sunos-news	x-font-ttf	x-font-type1
x-font-vfont	x-foxmail	x-futuresplash
x-gnucash	x-gnumeric	x-grib
x-gtar	x-hdf	x-hwp
x-hwp-v5	x-ibooks+zip	x-isatab
x-isatab-assay	x-isatab-investigation	x-iso9660-image
x-itunes-ipa	x-java-jnilib	x-java-jnlp-file
x-java-pack200	x-kdelnk	x-killustrator
x-latex	x-lha	x-lharc
x-matlab-data	x-matroska	x-mobipocket-ebook
x-ms-application	x-ms-installer	x-ms-wmd
x-ms-wmz	x-ms-xbap	x-msaccess
x-msbinder	x-mscardfile	x-msclip
x-msdownload	x-msdownload	x-msdownload
x-msdownload	x-msdownload	x-msdownload
x-msdownload	x-msmediaview	x-msmetafile
x-msmoney	x-mspublisher	x-msschedule
x-msterminal	x-mswrite	x-mysql-db
x-mysql-misam-compressed-index	x-mysql-misam-data	x-mysql-misam-index
x-mysql-table-definition	x-netcdf	x-object
x-pkcs12	x-pkcs7-certificates	x-pkcs7-certreqresp
x-project	x-prt	x-quattro-pro
x-rar-compressed	x-roxio-toast	x-rpm
x-sas	x-sas-access	x-sas-audit
x-sas-backup	x-sas-catalog	x-sas-data
x-sas-data-index	x-sas-dmdb	x-sas-fdb
x-sas-itemstor	x-sas-mddb	x-sas-program-data

x-sas-utility	x-sas-transport	x-sas-utility
x-sas-view	x-sc	x-sfdu
x-sh	x-shapefile	x-shar
x-sharedlib	x-shockwave-flash	x-silverlight-app
x-snappy-framed	x-sqlite3	x-staroffice-template
x-stuffit	x-stuffitx	x-sv4cpio
x-sv4crc	x-tar	x-tex
x-tex-tfm	x-texinfo	x-tika-iworks-protected
x-tika-java-enterprise-archive	x-tika-java-web-archive	x-tika-msoffice
x-tika-msoffice-embedded	x-tika-msoffice-embedded	x-tika-msoffice-embedded
x-tika-msworks-spreadsheet	x-tika-old-excel	x-tika-ooxml
x-tika-ooxml-protected	x-tika-staroffice	x-tika-unix-dump
x-tika-visio-ooxml	x-uc2-compressed	x-ustar
x-vhd	x-vmdk	x-wais-source
x-wearchive	x-x509-ca-cert	x-xfig
x-xmind	x-xpinstall	x-xz
x-zoo	x400-bp	xcap-att+xml
xcap-caps+xml	xcap-el+xml	xcap-error+xml
xcap-ns+xml	xcon-conference-info+xml	xcon-conference-info-diff+xml
xenc+xml	xhtml+xml	xhtml-voice+xml
xml	xml-dtd	xml-external-parsed-entity
xmpp+xml	xop+xml	xquery
xslfo+xml	xslt+xml	xspf+xml
xv+xml	zip	zlib

Table 200. Audio File Types

32kadpcm	3gpp	3gpp2
ac3	adpcm	amr
amr-wb	amr-wb+	asc
basic	bv16	bv32
clearmode	cn	dat12
dls	dsr-es201108	dsr-es202050
dsr-es202211	dsr-es202212	dvi4
eac3	evrc	evrc-qcp
evrc0	evrc1	evrcb
evrcb0	evrcb1	evrcwb
evrcwb0	evrcwb1	example
g719	g722	g7221
g723	g726-16	g726-24
g726-32	g726-40	g728
g729	g7291	g729d
g729e	gsm	gsm-efr
ilbc	l16	l20

124	18	lpc
midi	mobile-xmf	mp4
mp4a-latm	mpa	mpa-robust
mpeg	mpeg4-generic	ogg
opus	parityfec	pcma
pcma-wb	pcmu	pcmu-wb
prs.sid	qcelp	red
rtp-enc-aescm128	rtp-midi	rtx
smv	smv-qcp	smv0
sp-midi	speex	t140c
t38	telephone-event	tone
ulpfec	vdvi	vmr-wb
vnd.3gpp.iupf	vnd.4sb	vnd.adobe.soundbooth
vnd.audiokoz	vnd.celp	vnd.cisco.nse
vnd.cmles.radio-events	vnd.cns.anp1	vnd.cns.inf1
vnd.digital-winds	vnd.dlna.adts	vnd.dolby.heaac.1
vnd.dolby.heaac.2	vnd.dolby.mlp	vnd.dolby.mps
vnd.dolby.pl2	vnd.dolby.pl2x	vnd.dolby.pl2z
vnd.dts	vnd.dts.hd	vnd.everad.plj
vnd.hns.audio	vnd.lucent.voice	vnd.ms-playready.media.pya
vnd.nokia.mobile-xmf	vnd.nortel.vbk	vnd.nuera.ecelp4800
vnd.nuera.ecelp7470	vnd.nuera.ecelp9600	vnd.octel.sbc
vnd.qcelp	vnd.rhetorex.32kadpcm	vnd.sealedmedia.softseal.mpeg
vnd.vmx.cvsd	vorbis	vorbis-config
x-aac	x-adbcm	x-aiff
x-dec-adbcm	x-dec-basic	x-flac
x-matroska	x-mod	x-mpegurl
x-ms-wax	x-ms-wma	x-oggflac
x-oggpcm	x-pn-realaudio	x-pn-realaudio-plugin
x-wav		

Table 201. Chemical File Types

x-cdx	x-cif	x-cmdf
x-cml	x-csml	x-pdb
x-xyz		

Table 202. Image File Types

bmp	cgm	example
fits	g3fax	gif
icns	ief	jp2
jpeg	jpm	jpx
naplps	nitf	png
prs.btif	prs.pti	svg+xml

t38	tiff	tiff-fx
vnd.adobe.photoshop	vnd.adobe.premiere	vnd.cns.inf2
vnd.djvu	vnd.dwg	vnd.dxb
vnd.dxf	vnd.dxf	vnd.dxf
vnd.fastbidsheet	vnd.fpx	vnd.fst
vnd.fujixerox.edmics-mmr	vnd.fujixerox.edmics-rlc	vnd.globalgraphics.pgb
vnd.microsoft.icon	vnd.mix	vnd.ms-modi
vnd.net-fpx	vnd.radiance	vndsealed.png
vndsealedmedia.softseal.gif	vndsealedmedia.softseal.jpg	vnd.svf
vnd.wap.wbmp	vnd.xiff	webp
x-bpg	x-cmu-raster	x-cmx
x-freehand	x-jp2-codestream	x-jp2-container
x-ms-bmp	x-niff	x-pcx
x-pict	x-portable-anymap	x-portable-bitmap
x-portable-graymap	x-portable-pixmap	x-raw-adobe
x-raw-canon	x-raw-casio	x-raw-epson
x-raw-fuji	x-raw-hasselblad	x-raw-imacon
x-raw-kodak	x-raw-leaf	x-raw-logitech
x-raw-mamiya	x-raw-minolta	x-raw-nikon
x-raw-olympus	x-raw-panasonic	x-raw-pentax
x-raw-phaseone	x-raw-rawzor	x-raw-red
x-raw-sigma	x-raw-sony	x-rgb
x-xbitmap	x-xcf	x-xpixmap
x-xwindowdump		

Table 203. Message File Types

cpim	delivery-status	disposition-notification
example	external-body	global
global-delivery-status	global-disposition-notification	global-headers
http	imdn+xml	news
partial	rfc822	s-http
sip	sipfrag	tracking-status
vnd.si.simp	x-emlx	

Table 204. Model File Types

example	iges	mesh
vnd.dwf	vnd.dwf	vnd.dwf
vnd.dwf	vnd.dwfx+xps	vnd.flatland.3dml
vnd.gdl	vnd.gs-gdl	vnd.gs.gdl
vnd.gtw	vnd.moml+xml	vnd.mts
vnd.parasolid.transmit.binary	vnd.parasolid.transmit.text	vnd.vtu
vrml		

Table 205. Multipart File Types

alternative	appledouble	byteranges
digest	encrypted	example
form-data	header-set	mixed
parallel	related	report
signed	voice-message	

Table 206. Text File Types

asp	aspdotnet	calendar
css	csv	directory
dns	ecmascript	enriched
example	html	iso19139+xml
parityfec	plain	prs.fallenstein.rst
prs.lines.tag	red	rfc822-headers
richtext	rtp-enc-aescm128	rtx
sgml	t140	tab-separated-values
troff	ulpfec	uri-list
vnd.abc	vnd(curl	vnd(curl.dcurl
vnd(curl.mcurl	vnd(curl.scurl	vnd(dmclientscript
vnd.esmertec.theme-descriptor	vnd.fly	vnd.fmi.flexstor
vnd.graphviz	vnd.in3d.3dml	vnd.in3d.spot
vnd.ietf.params	vnd.ietf.params	vnd.ietf.params
vnd.latex-z	vnd.motorola.reflex	vnd.ms-mediapackage
vnd.net2phone.commcenter.command	vnd.si.uricatalogue	vnd.sun.j2me.app-descriptor
vnd.trolltech.linguist	vnd.wap.si	vnd.wap.sl
vnd.wap.wml	vnd.wap.wmlscript	vtt
x-actionscript	x-ada	x-applescript
x-asciidoc	x-aspectj	x-assembly
x-awk	x-basic	x-c++hdr
x-c++src	x-cgi	x-chdr
x-clojure	x-cobol	x-coffeescript
x-coldfusion	x-common-lisp	x-csharp
x-csrc	x-d	x-diff
x-eiffel	x-emacs-lisp	x-erlang
x-expect	x-forth	x-fortran
x-go	x-groovy	x-haml
x-haskell	x-haxe	x-idl
x-ini	x-java-properties	x-java-source
x-jsp	x-less	x-lex
x-log	x-lua	x-matlab
x-ml	x-modula	x-objcsrc
x-ocaml	x-pascal	x-perl

x-php	x-prolog	x-python
x-rexx	x-rsrc	x-rst
x-ruby	x-scala	x-scheme
x-sed	x-setext	x-sql
x-stsrc	x-tcl	x-tika-text-based-message
x-uuencode	x-vbasic	x-vbdotnet
x-vbscript	x-vcalendar	x-vcard
x-verilog	x-vhdl	x-web-markdown
x-yacc	x-yaml	

Table 207. Video File Types

3gpp	3gpp-tt	3gpp2
bmp	bt656	celb
daala	dv	example
h261	h263	h263-1998
h263-2000	h264	jpeg
jpeg2000	mj2	mp1s
mp2p	mp2t	mp4
mp4v-es	mpeg	mpeg4-generic
mpv	nv	ogg
parityfec	pointer	quicktime
raw	rtp-enc-aescm128	rtx
smpete292m	theora	ulpfec
vc1	vnd.cctv	vnd.dlna.mpeg-tts
vnd.fvt	vnd.hns.video	vnd.ffmpeg.1dparityfec-1010
vnd.ffmpeg.1dparityfec-2005	vnd.ffmpeg.2dparityfec-1010	vnd.ffmpeg.2dparityfec-2005
vnd.ffmpeg.ttsavc	vnd.ffmpeg.ttsmpeg2	vnd.motorola.video
vnd.motorola.videop	vnd.mpegurl	vnd.ms-playready.media.pyv
vnd.nokia.interleaved-multimedia	vnd.nokia.videovoip	vnd.objectvideo
vndsealed.mpeg1	vndsealed.mpeg4	vndsealed.swf
vndsealedmedia.softseal.mov	vnd.vivo	webm
x-dirac	x-f4v	x-flc
x-fli	x-flv	x-jng
x-m4v	x-matroska	x-mng
x-ms-asf	x-ms-wm	x-ms-wmv
x-ms-wmx	x-ms-wvx	x-msvideo
x-oggrgb	x-ogguv	x-oggyuv
x-ogm	x-sgi-movie	

Table 208. x-conference File Types

x-cooltalk		
------------	--	--

Appendix D: DDF Dependency List

This list of DDF dependencies is automatically generated:

DDF 2.11.3 Dependency List. published

- antlr:antlr:jar:2.7.7
- c3p0:c3p0:jar:0.9.1.1
- ca.juliusdavies:not-yet-commons-ssl:jar:0.3.11
- cglib:cglib-nodep:jar:3.2.0
- ch.qos.logback:logback-access:jar:1.2.3
- ch.qos.logback:logback-classic:jar:1.2.3
- ch.qos.logback:logback-core:jar:1.2.3
- com.codahale.metrics:metrics-core:jar:3.0.1
- com.connexta.arbitro:arbitro-core:jar:1.0.0
- com.github.drapostolos:type-parser:jar:0.5.0
- com.github.jai-imageio:jai-imageio-core:jar:1.3.1
- com.github.jai-imageio:jai-imageio-jpeg2000:jar:1.3.0
- com.github.jknack:handlebars-jackson2:jar:1.0.0
- com.github.jknack:handlebars:jar:1.1.2
- com.github.jknack:handlebars:jar:2.0.0
- com.github.lookfirst:sardine:jar:5.7
- com.github.rvesse:airline:jar:2.1.0
- com.google.code.findbugs:findbugs:jar:3.0.1
- com.google.code.findbugs:jsr305:jar:3.0.1
- com.google.code.gson:gson:jar:2.7
- com.google.guava:guava:jar:18.0
- com.google.http-client:google-http-client:jar:1.22.0
- com.googlecode.json-simple:json-simple:jar:1.1.1
- com.hazelcast:hazelcast:jar:3.2.1
- com.jayway.restassured:rest-assured:jar:2.9.0
- com.jhlabs:filters:jar:2.0.235-1
- com.lucidworks:banana:war:1.5.1
- com.rometools:rome-utils:jar:1.5.0
- com.rometools:rome:jar:1.5.0
- com.sparkjava:spark-core:jar:2.5.5
- com.sun.mail:javax.mail:jar:1.5.5

- com.sun.xml.bind:jaxb-core:jar:2.2.11
- com.sun.xml.bind:jaxb-impl:jar:2.2.11
- com.thoughtworks.xstream:xstream:jar:1.4.9
- com.unboundid:unboundid-ldapsdk:jar:3.2.1
- com.vividsolutions:jts:jar:1.12
- com.vividsolutions:jts:jar:1.13
- com.xebialabs.restito:restito:jar:0.8.2
- commons-beanutils:commons-beanutils:jar:1.9.2
- commons-beanutils:commons-beanutils:jar:1.9.3
- commons-codec:commons-codec:jar:1.10
- commons-collections:commons-collections:jar:3.2.2
- commons-configuration:commons-configuration:jar:1.10
- commons-digester:commons-digester:jar:1.8.1
- commons-fileupload:commons-fileupload:jar:1.3.2
- commons-io:commons-io:jar:2.1
- commons-io:commons-io:jar:2.4
- commons-io:commons-io:jar:2.5
- commons-lang:commons-lang:jar:2.6
- commons-logging:commons-logging:jar:1.2
- commons-net:commons-net:jar:3.5
- commons-pool:commons-pool:jar:1.6
- commons-validator:commons-validator:jar:1.6
- de.micromata.jak:JavaAPIforKml:jar:2.2.0
- info.ganglia.gmetric4j:gmetric4j:jar:1.0.7
- io.dropwizard.metrics:metrics-core:jar:3.1.2
- io.dropwizard.metrics:metrics-core:jar:3.2.2
- io.dropwizard.metrics:metrics-ganglia:jar:3.2.2
- io.dropwizard.metrics:metrics-graphite:jar:3.2.2
- io.dropwizard.metrics:metrics-jetty9:jar:3.2.2
- io.dropwizard.metrics:metrics-jvm:jar:3.2.2
- io.fastjson:boon:jar:0.30
- io.fastjson:boon:jar:0.33
- io.fastjson:boon:jar:0.34
- io.netty:netty-buffer:jar:4.1.9.Final
- io.netty:netty-codec:jar:4.1.9.Final

- io.netty:netty-common:jar:4.1.9.Final
- io.netty:netty-handler:jar:4.1.9.Final
- io.netty:netty-resolver:jar:4.1.9.Final
- io.netty:netty-transport-native-epoll:jar:4.1.9.Final
- io.netty:netty-transport:jar:4.1.9.Final
- javax.cache:cache-api:jar:1.0.0
- javax.inject:javax.inject:jar:1
- javax.mail:mail:jar:1.4.5
- javax.servlet:javax.servlet-api:jar:3.1.0
- javax.servlet:servlet-api:jar:2.5
- javax.validation:validation-api:jar:1.1.0.Final
- javax.ws.rs:javax.ws.rs-api:jar:2.0-m10
- javax.ws.rs:javax.ws.rs-api:jar:2.0.1
- javax.ws.rs:javax.ws.rs-api:jar:2.0
- joda-time:joda-time:jar:2.9.4
- junit:junit:jar:4.12
- log4j:log4j:jar:1.2.17
- net.iharder:base64:jar:2.3.9
- net.jodah:failsafe:jar:0.9.3
- net.jodah:failsafe:jar:0.9.5
- net.jodah:failsafe:jar:1.0.0
- net.lingala.zip4j:zip4j:jar:1.3.2
- net.markenwerk:commons-nulls:jar:1.0.3
- net.markenwerk:utils-data-fetcher:jar:4.0.1
- net.minidev:asm:jar:1.0.2
- net.minidev:json-smart:jar:2.2.1
- net.sf.saxon:Saxon-HE:jar:9.5.1-3
- net.sf.saxon:Saxon-HE:jar:9.6.0-4
- org.antlr:antlr4-runtime:jar:4.1
- org.antlr:antlr4-runtime:jar:4.3
- org.apache.abdera:abdera-extensions-geo:jar:1.1.3
- org.apache.abdera:abdera-extensions-opensearch:jar:1.1.3
- org.apache.activemq:activemq-all:jar:5.14.5
- org.apache.activemq:artemis-commons:jar:2.1.0
- org.apache.activemq:artemis-jms-client:jar:2.1.0

- org.apache.ant:ant-launcher:jar:1.9.7
- org.apache.ant:ant:jar:1.9.7
- org.apache.aries.jmx:org.apache.aries.jmx.api:jar:1.1.5
- org.apache.aries.jmx:org.apache.aries.jmx.core:jar:1.1.7
- org.apache.aries:org.apache.aries.util:jar:1.1.3
- org.apache.camel:camel-amqp:jar:2.19.0
- org.apache.camel:camel-aws:jar:2.19.0
- org.apache.camel:camel-blueprint:jar:2.19.0
- org.apache.camel:camel-context:jar:2.19.0
- org.apache.camel:camel-core-osgi:jar:2.19.0
- org.apache.camel:camel-core:jar:2.19.0
- org.apache.camel:camel-cxf:jar:2.19.0
- org.apache.camel:camel-http-common:jar:2.19.0
- org.apache.camel:camel-http4:jar:2.19.0
- org.apache.camel:camel-http:jar:2.19.0
- org.apache.camel:camel-quartz:jar:2.19.0
- org.apache.camel:camel-saxon:jar:2.19.0
- org.apache.camel:camel-servlet:jar:2.19.0
- org.apache.camel:camel-sjms:jar:2.19.0
- org.apache.camel:camel-stream:jar:2.19.0
- org.apache.commons:commons-collections4:jar:4.1
- org.apache.commons:commons-compress:jar:1.14
- org.apache.commons:commons-csv:jar:1.4
- org.apache.commons:commons-exec:jar:1.3
- org.apache.commons:commons-lang3:jar:3.0
- org.apache.commons:commons-lang3:jar:3.1
- org.apache.commons:commons-lang3:jar:3.3.2
- org.apache.commons:commons-lang3:jar:3.4
- org.apache.commons:commons-math:jar:2.2
- org.apache.commons:commons-pool2:jar:2.4.2
- org.apache.cxf.services.sts:cxf-services-sts-core:jar:3.1.11
- org.apache.cxf:cxf-core:jar:3.1.11
- org.apache.cxf:cxf-rt-bindings-soap:jar:3.0.4
- org.apache.cxf:cxf-rt-databinding-jaxb:jar:3.0.4
- org.apache.cxf:cxf-rt-frontend-jaxrs:jar:3.1.11

- org.apache.cxf:cxf-rt-frontend-jaxws:jar:3.0.4
- org.apache.cxf:cxf-rt-frontend-jaxws:jar:3.1.11
- org.apache.cxf:cxf-rt-rs-client:jar:3.1.11
- org.apache.cxf:cxf-rt-rs-security-sso-saml:jar:3.1.11
- org.apache.cxf:cxf-rt-rs-security-xml:jar:3.0.4
- org.apache.cxf:cxf-rt-rs-security-xml:jar:3.1.11
- org.apache.cxf:cxf-rt-transports-http:jar:3.1.11
- org.apache.cxf:cxf-rt-ws-policy:jar:3.1.11
- org.apache.cxf:cxf-rt-ws-security:jar:3.1.11
- org.apache.felix:org.apache.felix.configadmin:jar:1.8.14
- org.apache.felix:org.apache.felix.coordinator:jar:1.0.2
- org.apache.felix:org.apache.felix.fileinstall:jar:3.6.0
- org.apache.felix:org.apache.felix.framework:jar:5.6.6
- org.apache.felix:org.apache.felix.utils:jar:1.10.0
- org.apache.ftpserver:ftplet-api:jar:1.0.6
- org.apache.ftpserver:ftpserver-core:jar:1.0.6
- org.apache.geronimo.specs:geronimo-servlet_3.0_spec:jar:1.0
- org.apache.httpcomponents:httpclient:jar:4.5.3
- org.apache.httpcomponents:httpcore:jar:4.4.6
- org.apache.httpcomponents:httpmime:jar:4.5.3
- org.apache.karaf.bundle:org.apache.karaf.bundle.core:jar:4.1.2
- org.apache.karaf.features:org.apache.karaf.features.core:jar:4.1.2
- org.apache.karaf.features:standard:xml:features:4.1.2
- org.apache.karaf.jaas:org.apache.karaf.jaas.boot:jar:4.1.2
- org.apache.karaf.jaas:org.apache.karaf.jaas.config:jar:4.1.2
- org.apache.karaf.jaas:org.apache.karaf.jaas.modules:jar:4.1.2
- org.apache.karaf.shell:org.apache.karaf.shell.console:jar:4.1.2
- org.apache.karaf.shell:org.apache.karaf.shell.core:jar:4.1.2
- org.apache.karaf:apache-karaf:tar.gz:4.1.2
- org.apache.karaf:apache-karaf:zip:4.1.2
- org.apache.karaf:org.apache.karaf.util:jar:4.1.2
- org.apache.logging.log4j:log4j-api:jar:2.4.1
- org.apache.lucene:lucene-analyzers-common:jar:6.6.0
- org.apache.lucene:lucene-core:jar:3.0.2
- org.apache.lucene:lucene-core:jar:6.6.0

- org.apache.lucene:lucene-queries:jar:6.6.0
- org.apache.lucene:lucene-queryparser:jar:6.6.0
- org.apache.lucene:lucene-sandbox:jar:6.6.0
- org.apache.lucene:lucene-spatial-extras:jar:6.6.0
- org.apache.lucene:lucene-spatial3d:jar:6.6.0
- org.apache.lucene:lucene-spatial:jar:6.6.0
- org.apache.maven.shared:maven-invoker:jar:2.2
- org.apache.mina:mina-core:jar:2.0.6
- org.apache.pdfbox:fontbox:jar:2.0.2
- org.apache.pdfbox:pdfbox-tools:jar:2.0.2
- org.apache.pdfbox:pdfbox:jar:2.0.2
- org.apache.poi:poi-ooxml:jar:3.12
- org.apache.poi:poi-scratchpad:jar:3.12
- org.apache.poi:poi:jar:3.12
- org.apache.servicemix.bundles:org.apache.servicemix.bundles.jsr305:jar:1.3.9_1
- org.apache.servicemix.bundles:org.apache.servicemix.bundles.poi:jar:3.16_1
- org.apache.servicemix.specs:org.apache.servicemix.specs.jsr339-api-2.0:jar:2.6.0
- org.apache.shiro:shiro-core:jar:1.3.2
- org.apache.solr:solr-core:jar:6.6.0
- org.apache.solr:solr-solrj:jar:6.6.0
- org.apache.tika:tika-core:jar:1.15
- org.apache.tika:tika-parsers:jar:1.15
- org.apache.ws.commons.axiom:axiom-api:jar:1.2.14
- org.apache.wss4j:wss4j-bindings:jar:2.1.11
- org.apache.wss4j:wss4j-policy:jar:2.1.11
- org.apache.wss4j:wss4j-ws-security-common:jar:2.1.11
- org.apache.wss4j:wss4j-ws-security-dom:jar:2.1.11
- org.apache.wss4j:wss4j-ws-security-policy-stax:jar:2.1.11
- org.apache.wss4j:wss4j-ws-security-stax:jar:2.1.11
- org.asciidoctor:asciidoctorj-diagram:jar:1.5.4.1
- org.asciidoctor:asciidoctorj:jar:1.5.6
- org.assertj:assertj-core:jar:2.1.0
- org.awaitility:awaitility:jar:3.0.0
- org.bouncycastle:bcmail-jdk15on:jar:1.55
- org.bouncycastle:bcpkix-jdk15on:jar:1.55

- org.bouncycastle:bcpprov-jdk15on:jar:1.55
- org.codehaus.groovy:groovy-all:jar:2.4.7
- org.codehaus.jackson:jackson-core-asl:jar:1.9.13
- org.codehaus.jackson:jackson-mapper-asl:jar:1.9.13
- org.codehaus.woodstox:woodstox-core-asl:jar:4.4.1
- org.codice.geowebcache:geowebcache-server-standalone:war:0.7.0
- org.codice.geowebcache:geowebcache-server-standalone:xml:geowebcache:0.7.0
- org.codice.httpproxy:proxy-camel-route:jar:2.11.3
- org.codice.httpproxy:proxy-camel-servlet:jar:2.11.3
- org.codice.opendj.embedded:opendj-embedded-app:xml:features:1.3.3
- org.codice.thirdparty:cas-client-core:jar:3.1.10_1
- org.codice.thirdparty:commons-httpclient:jar:3.1.0_1
- org.codice.thirdparty:ffmpeg:zip:bin:3.1.1_1
- org.codice.thirdparty:geotools-suite:jar:8.4_2
- org.codice.thirdparty:gt-opengis:jar:8.4_1
- org.codice.thirdparty:jts:jar:1.12_1
- org.codice.thirdparty:lucene-core:jar:3.0.2_1
- org.codice.thirdparty:ogc-filter-v_1_1_0-schema:jar:1.1.0_4
- org.codice.thirdparty:picocontainer:jar:1.2_1
- org.codice.thirdparty:tika-bundle:jar:1.14.0_1
- org.codice.thirdparty:vecmath:jar:1.3.2_1
- org.codice.usng4j:usng4j-api:jar:0.1
- org.codice.usng4j:usng4j-impl:jar:0.1
- org.codice.webjars:backbone-poller:jar:1.1.3
- org.codice.webjars:backbone.modelbinder:jar:1.1.0
- org.codice.webjars:jquery-ui-multiselect-widget:jar:1.14
- org.codice.webjars:jqueryui-timepicker-addon:jar:1.4.5
- org.codice.webjars:marionette:jar:1.8.8
- org.codice.webjars:marionette:jar:2.4.1
- org.codice.webjars:openlayers3:jar:3.16.0
- org.codice.webjars:pnotify:jar:1.3.1
- org.codice.webjars:wellknown:jar:0.4.0
- org.codice:lux:jar:1.2
- org.cometd.java:bayeux-api:jar:3.0.9
- org.cometd.java:cometd-java-annotations:jar:3.0.9

- org.cometd.java:cometd-java-client:jar:3.0.7
- org.cometd.java:cometd-java-client:jar:3.0.9
- org.cometd.java:cometd-java-common:jar:3.0.9
- org.cometd.java:cometd-java-server:jar:3.0.9
- org.eclipse.jetty:jetty-http:jar:9.3.14.v20161028
- org.eclipse.jetty:jetty-server:jar:9.3.14.v20161028
- org.eclipse.jetty:jetty-servlet:jar:9.3.14.v20161028
- org.eclipse.jetty:jetty-servlets:jar:9.2.19.v20160908
- org.eclipse.jetty:jetty-util:jar:9.3.14.v20161028
- org.forgerock.commons:forgerock-util:jar:3.0.2
- org.forgerock.commons:i18n-core:jar:1.4.2
- org.forgerock.commons:i18n-slf4j:jar:1.4.2
- org.forgerock.opendj:opendj-core:jar:3.0.0
- org.forgerock.opendj:opendj-grizzly:jar:3.0.0
- org.fusesource.jansi:jansi:jar:1.16
- org.geotools.xsd:gt-xsd-gml3:jar:8.4
- org.geotools:gt-cql:jar:13.0
- org.geotools:gt-cql:jar:14.4
- org.geotools:gt-cql:jar:8.4
- org.geotools:gt-epsg-hsql:jar:8.4
- org.geotools:gt-jts-wrapper:jar:8.4
- org.geotools:gt-main:jar:14.4
- org.geotools:gt-main:jar:8.4
- org.geotools:gt-opengis:jar:14.4
- org.geotools:gt-opengis:jar:8.4
- org.geotools:gt-referencing:jar:8.4
- org.geotools:gt-shapefile:jar:8.4
- org.geotools:gt-xml:jar:8.4
- org.glassfish.grizzly:grizzly-framework:jar:2.3.30
- org.glassfish.grizzly:grizzly-http-server:jar:2.3.25
- org.hamcrest:hamcrest-all:jar:1.3
- org.imgscalr:imgscalr-lib:jar:4.2
- org.jasig.cas:cas-client-core:jar:3.1.10
- org.jasypt:jasypt:jar:1.9.0
- org.jcodec:jcodec:jar:0.2.0_1

- org.jdom:jdom:jar:2.0.2
- org.joda:joda-convert:jar:1.2
- org.jolokia:jolokia-osgi:jar:1.2.3
- org.jruby:jruby-complete:jar:1.7.26
- org.jscience:jscience:jar:4.3.1
- org.jvnet.jaxb2_commons:jaxb2-basics-runtime:jar:0.6.0
- org.jvnet.jaxb2_commons:jaxb2-basics-runtime:jar:0.9.4
- org.jvnet.ogc:filter-v_2_0_0-schema:jar:1.1.0
- org.jvnet.ogc:gml-v_3_1_1-schema:jar:1.0.2
- org.jvnet.ogc:gml-v_3_1_1-schema:jar:1.0.3
- org.jvnet.ogc:gml-v_3_1_1-schema:jar:1.1.0
- org.jvnet.ogc:gml-v_3_2_1-schema:jar:1.1.0
- org.jvnet.ogc:gml-v_3_2_1:pom:1.1.0
- org.jvnet.ogc:ogc-tools-gml-jts:jar:1.0.3
- org.jvnet.ogc:ows-v_1_0_0-schema:jar:1.1.0
- org.jvnet.ogc:ows-v_1_1_0-schema:jar:1.1.0
- org.jvnet.ogc:wcs-v_1_0_0-schema:jar:1.1.0
- org.keyczar:keyczar:jar:0.66
- org.la4j:la4j:jar:0.6.0
- org.locationtech.spatial4j:spatial4j:jar:0.6
- org.noggit:noggit:jar:0.6
- org.objenesis:objenesis:jar:2.1
- org.openexi:nagasena-rta:jar:0000.0002.0049.0
- org.openexi:nagasena:jar:0000.0002.0049.0
- org.opensaml:opensaml-core:jar:3.1.1
- org.opensaml:opensaml-soap-impl:jar:3.1.1
- org.opensaml:opensaml-xmlsec-api:jar:3.1.1
- org.opensaml:opensaml-xmlsec-impl:jar:3.1.1
- org.ops4j.pax.exam:pax-exam-container-karaf:jar:4.11.0
- org.osgi:org.osgi.compendium:jar:4.3.1
- org.osgi:org.osgi.compendium:jar:5.0.0
- org.osgi:org.osgi.core:jar:4.3.1
- org.osgi:org.osgi.core:jar:5.0.0
- org.ow2.asm:asm:jar:5.0.4
- org.parboiled:parboiled-java:jar:1.1.7

- org.quartz-scheduler:quartz-jobs:jar:2.2.3
- org.quartz-scheduler:quartz:jar:2.1.7
- org.quartz-scheduler:quartz:jar:2.2.3
- org.rrd4j:rrd4j:jar:2.2
- org.simplejavamail:simple-java-mail:jar:4.1.3
- org.slf4j:jcl-over-slf4j:jar:1.7.7
- org.slf4j:jul-to-slf4j:jar:1.7.7
- org.slf4j:slf4j-api:jar:1.7.12
- org.slf4j:slf4j-api:jar:1.7.1
- org.slf4j:slf4j-api:jar:1.7.7
- org.slf4j:slf4j-ext:jar:1.7.1
- org.slf4j:slf4j-log4j12:jar:1.7.12
- org.slf4j:slf4j-log4j12:jar:1.7.7
- org.slf4j:slf4j-simple:jar:1.7.1
- org.slf4j:slf4j-simple:jar:1.7.5
- org.springframework.ldap:spring-ldap-core:jar:1.3.2.RELEASE
- org.springframework.osgi:spring-osgi-core:jar:1.2.1
- org.springframework:spring-core:jar:4.3.5.RELEASE
- org.taktik:mpegts-streamer:jar:0.1.0_2
- org.twitter4j:twitter4j-core:jar:4.0.4
- org.webjars.bower:Backbone.Undo:jar:0.2.5
- org.webjars.bower:backbone-associations:jar:0.6.2
- org.webjars.bower:backbone-relational:jar:0.8.8
- org.webjars.bower:backbone:jar:1.1.2
- org.webjars.bower:bootstrap-multiselect:jar:0.9.3
- org.webjars.bower:bootstrap-select:jar:1.6.4
- org.webjars.bower:bootstrap:jar:3.2.0
- org.webjars.bower:bootstrap:jar:3.3.7
- org.webjars.bower:bootswatch:jar:3.2.0
- org.webjars.bower:bootswatch:jar:3.3.7
- org.webjars.bower:cesiumjs:jar:1.22.0
- org.webjars.bower:components-font-awesome:jar:4.7.0
- org.webjars.bower:font-awesome:jar:4.6.3
- org.webjars.bower:font-awesome:jar:4.7.0
- org.webjars.bower:handlebars:jar:2.0.0

- org.webjars.bower:handlebars:jar:4.0.10
- org.webjars.bower:html5shiv:jar:3.7.2
- org.webjars.bower:iframe-resizer:jar:2.6.2
- org.webjars.bower:jquery-file-upload:jar:9.18.0
- org.webjars.bower:jquery-file-upload:jar:9.5.7
- org.webjars.bower:jquery-ui:jar:1.12.1
- org.webjars.bower:jquery:jar:2.2.4
- org.webjars.bower:jquery:jar:3.2.1
- org.webjars.bower:js-cookie:jar:2.1.4
- org.webjars.bower:lodash:jar:3.7.0
- org.webjars.bower:marionette:jar:2.4.5
- org.webjars.bower:moment:jar:2.10.3
- org.webjars.bower:moment:jar:2.18.1
- org.webjars.bower:perfect-scrollbar:jar:0.7.0
- org.webjars.bower:purl:jar:2.3.1
- org.webjars.bower:require-css:jar:0.1.10
- org.webjars.bower:requirejs-plugins:jar:1.0.3
- org.webjars.bower:requirejs:jar:2.1.15
- org.webjars.bower:requirejs:jar:2.3.3
- org.webjars.bower:spectrum:jar:1.8.0
- org.webjars.bower:spin.js:jar:1.3.3
- org.webjars.bower:spin.js:jar:2.3.2
- org.webjars.bower:underscore:jar:1.8.3
- org.webjars.bower:usng.js:jar:0.2.2
- org.webjars.npm:q:jar:1.4.1
- us.bpsm:edn-java:jar:0.4.4
- xalan:serializer:jar:2.7.2
- xalan:xalan:jar:2.7.2
- xerces:xercesImpl:jar:2.11.0
- xerces:xercesImpl:jar:2.9.1
- xml-apis:xml-apis:jar:1.4.01
- xmlpull:xmlpull:jar:1.1.3.1
- xpp3:xpp3:jar:1.1.4c